



**HARVARD**

**School of Engineering  
and Applied Sciences**

# Caching

*CS61, Lecture 12*

*Prof. Stephen Chong*

*October 12, 2010*

# Announcements

- Lab 3 design checkpoint deadline
  - Meet with your TF by Thursday 10pm
  - If you haven't already, arrange a meeting now!

# Topics for today

- The Principle of Locality
- Memory Hierarchies
- Caching Concepts
- Direct-Mapped Cache Organization
- Set-Associative Cache Organization
- Multi-level caches
- Cache writes

# Background: Locality

- Principle of Locality:
  - Programs tend to reuse data and instructions “near” those they have used recently.
- **Temporal locality:** Recently referenced memory addresses are likely to be referenced in the near future.
- **Spatial locality:** Similar memory addresses tend to be referenced close together in time.

# Locality Example

```
sum = 0;
for (i = 0; i < n; i++) {
    sum += a[i];
}
return sum;
```

- Locality of program data
  - Array elements are referenced in succession
  - Location `sum` is referenced on each iteration
- Locality of program code
  - Instructions within loop body executed in sequence
  - Loop is cycled over potentially many times

# Locality Example

- Claim: Being able to look at code and get a qualitative sense of its locality is a key skill for a professional programmer.
- Question: Does this function have good locality?

```
int sum_array_rows(int a[M][N])
{
    int i, j, sum = 0;

    for (j = 0; j < N; j++)
        for (i = 0; i < M; i++)
            sum += a[i][j];
    return sum;
}
```

# Locality Example

- Question: Does this function have good locality?

```
int sum_array_cols(int a[M][N])
{
    int i, j, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < N; j++)
            sum += a[i][j];
    return sum;
}
```

# Locality Example

- Question: Can you permute the loops so that the function scans the 3D array `a[][][]` with a stride-1 reference pattern (and thus has good spatial locality)?

```
int sum_array_3d(int a[M][N][N])
{
    int i, j, k, sum = 0;

    for (i = 0; i < N; i++)
        for (j = 0; j < N; j++)
            for (k = 0; k < M; k++)
                sum += a[k][i][j];

    return sum;
}
```

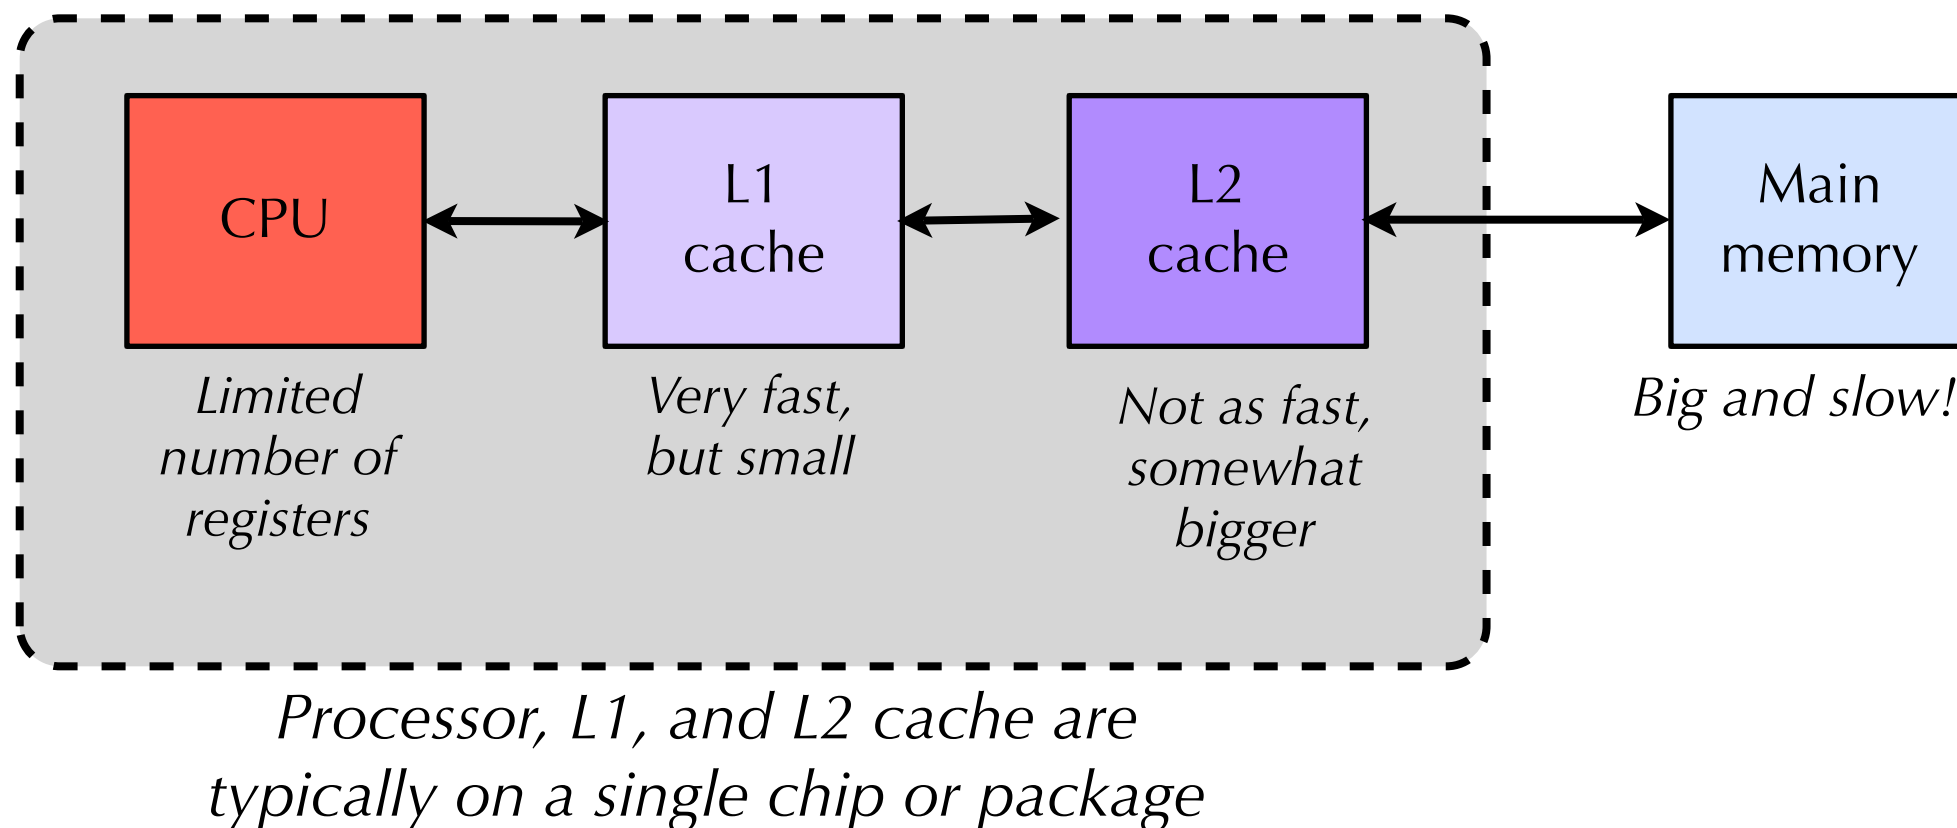


# Memory Hierarchies

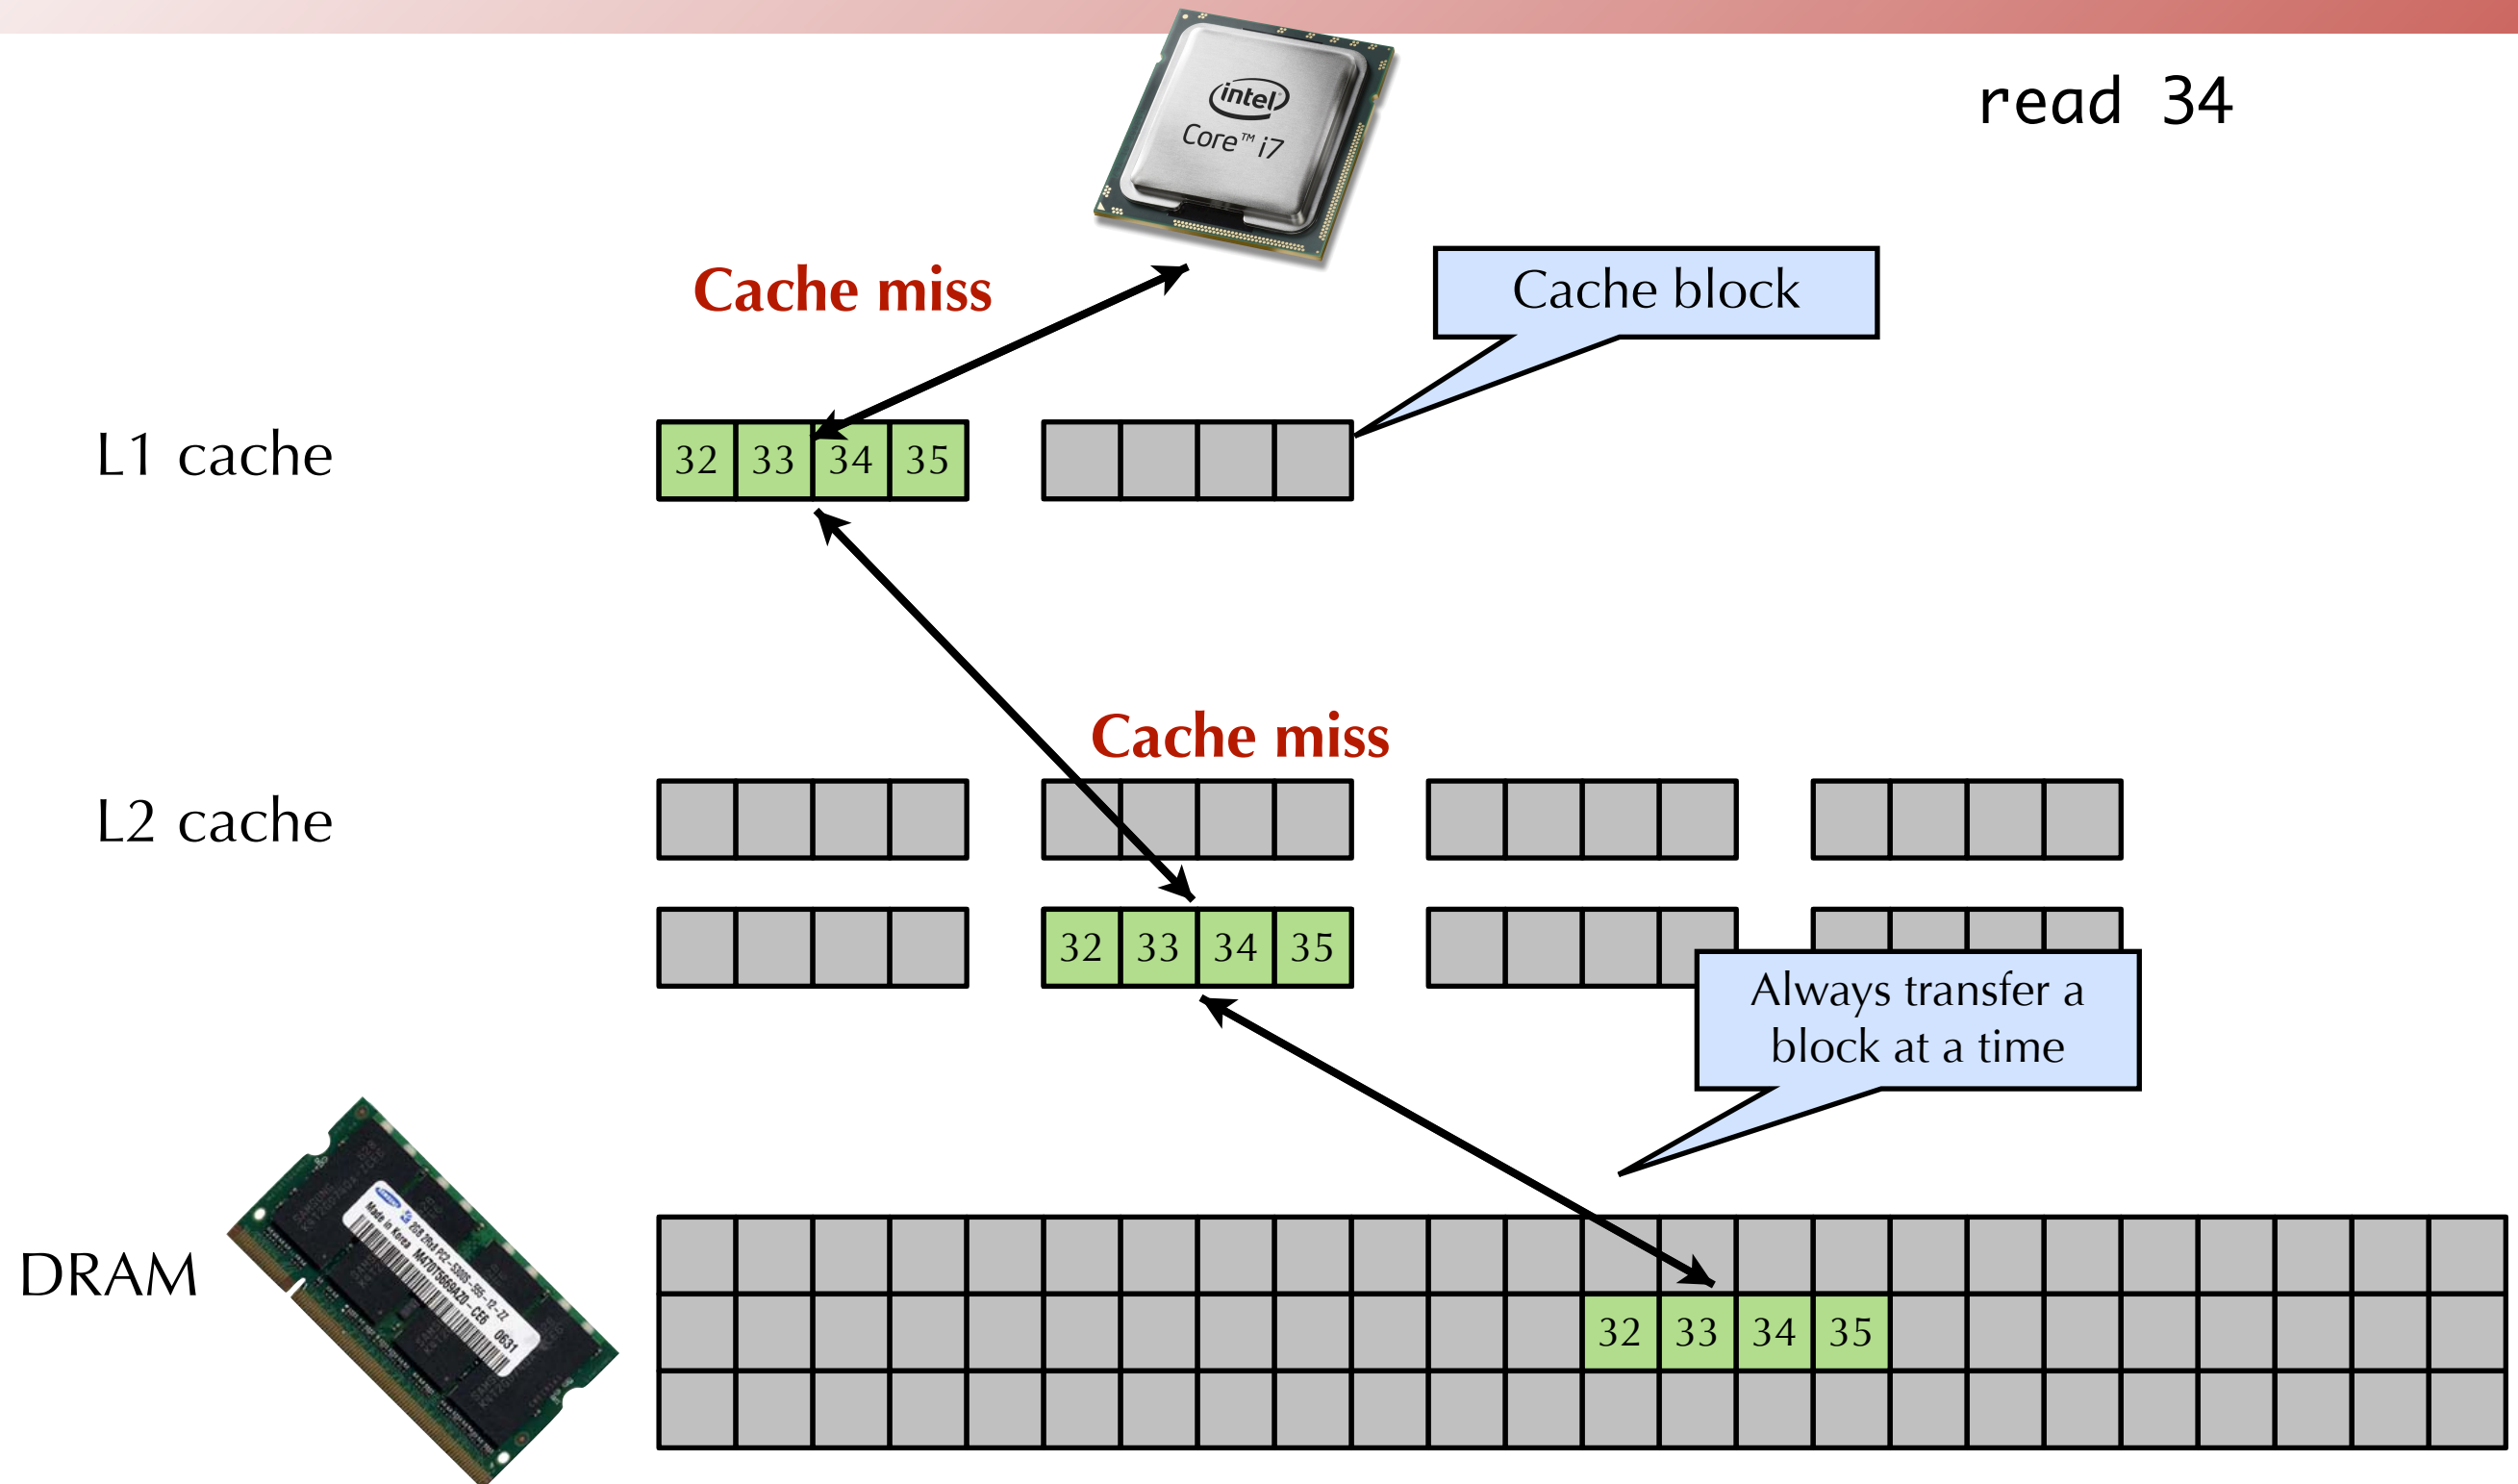
- Some fundamental and enduring properties of hardware and software:
  - Fast storage technologies cost more per byte, have less capacity, and require more power (heat!).
  - The gap between CPU and main memory speed is widening.
  - Well-written programs tend to exhibit good locality.
- These fundamental properties complement each other beautifully.
- They suggest an approach for organizing memory and storage systems known as a **memory hierarchy**.

# Cache Memories

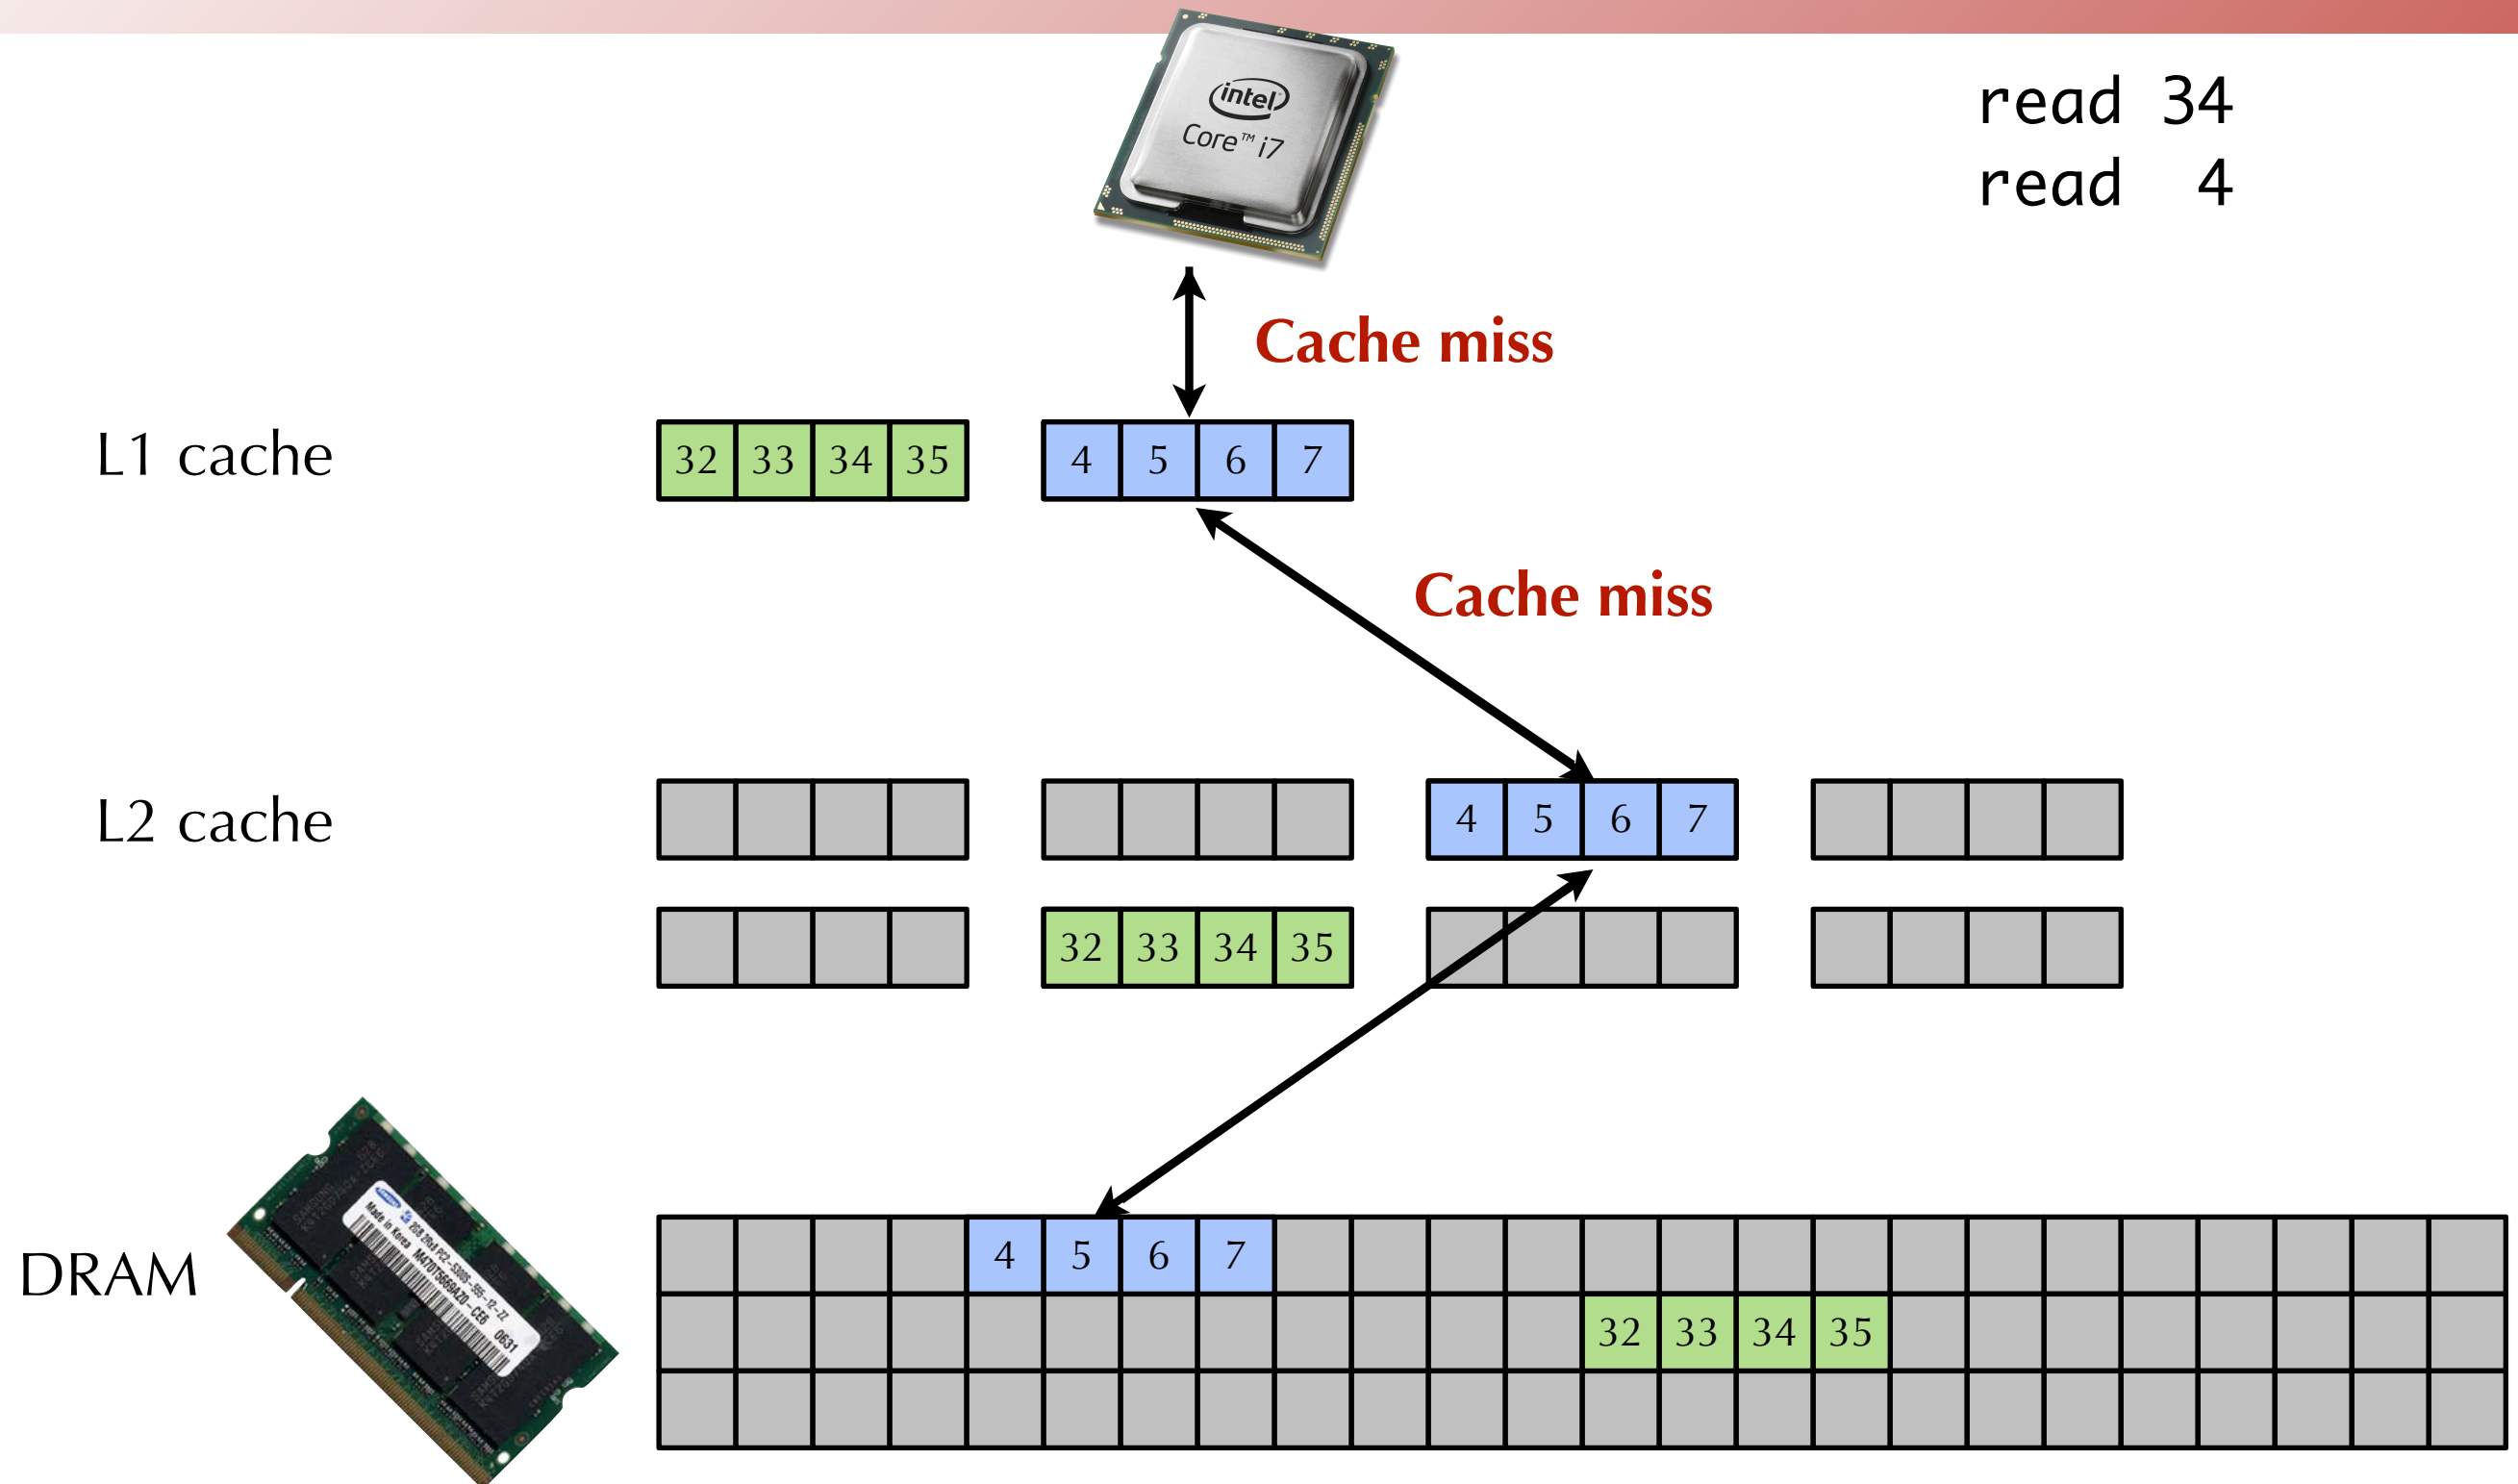
- Cache memories are small, fast memories managed automatically in hardware.
  - Hold frequently accessed blocks of main memory
- CPU looks first for data in L1, then in L2, then in main memory.
- Typical system structure:



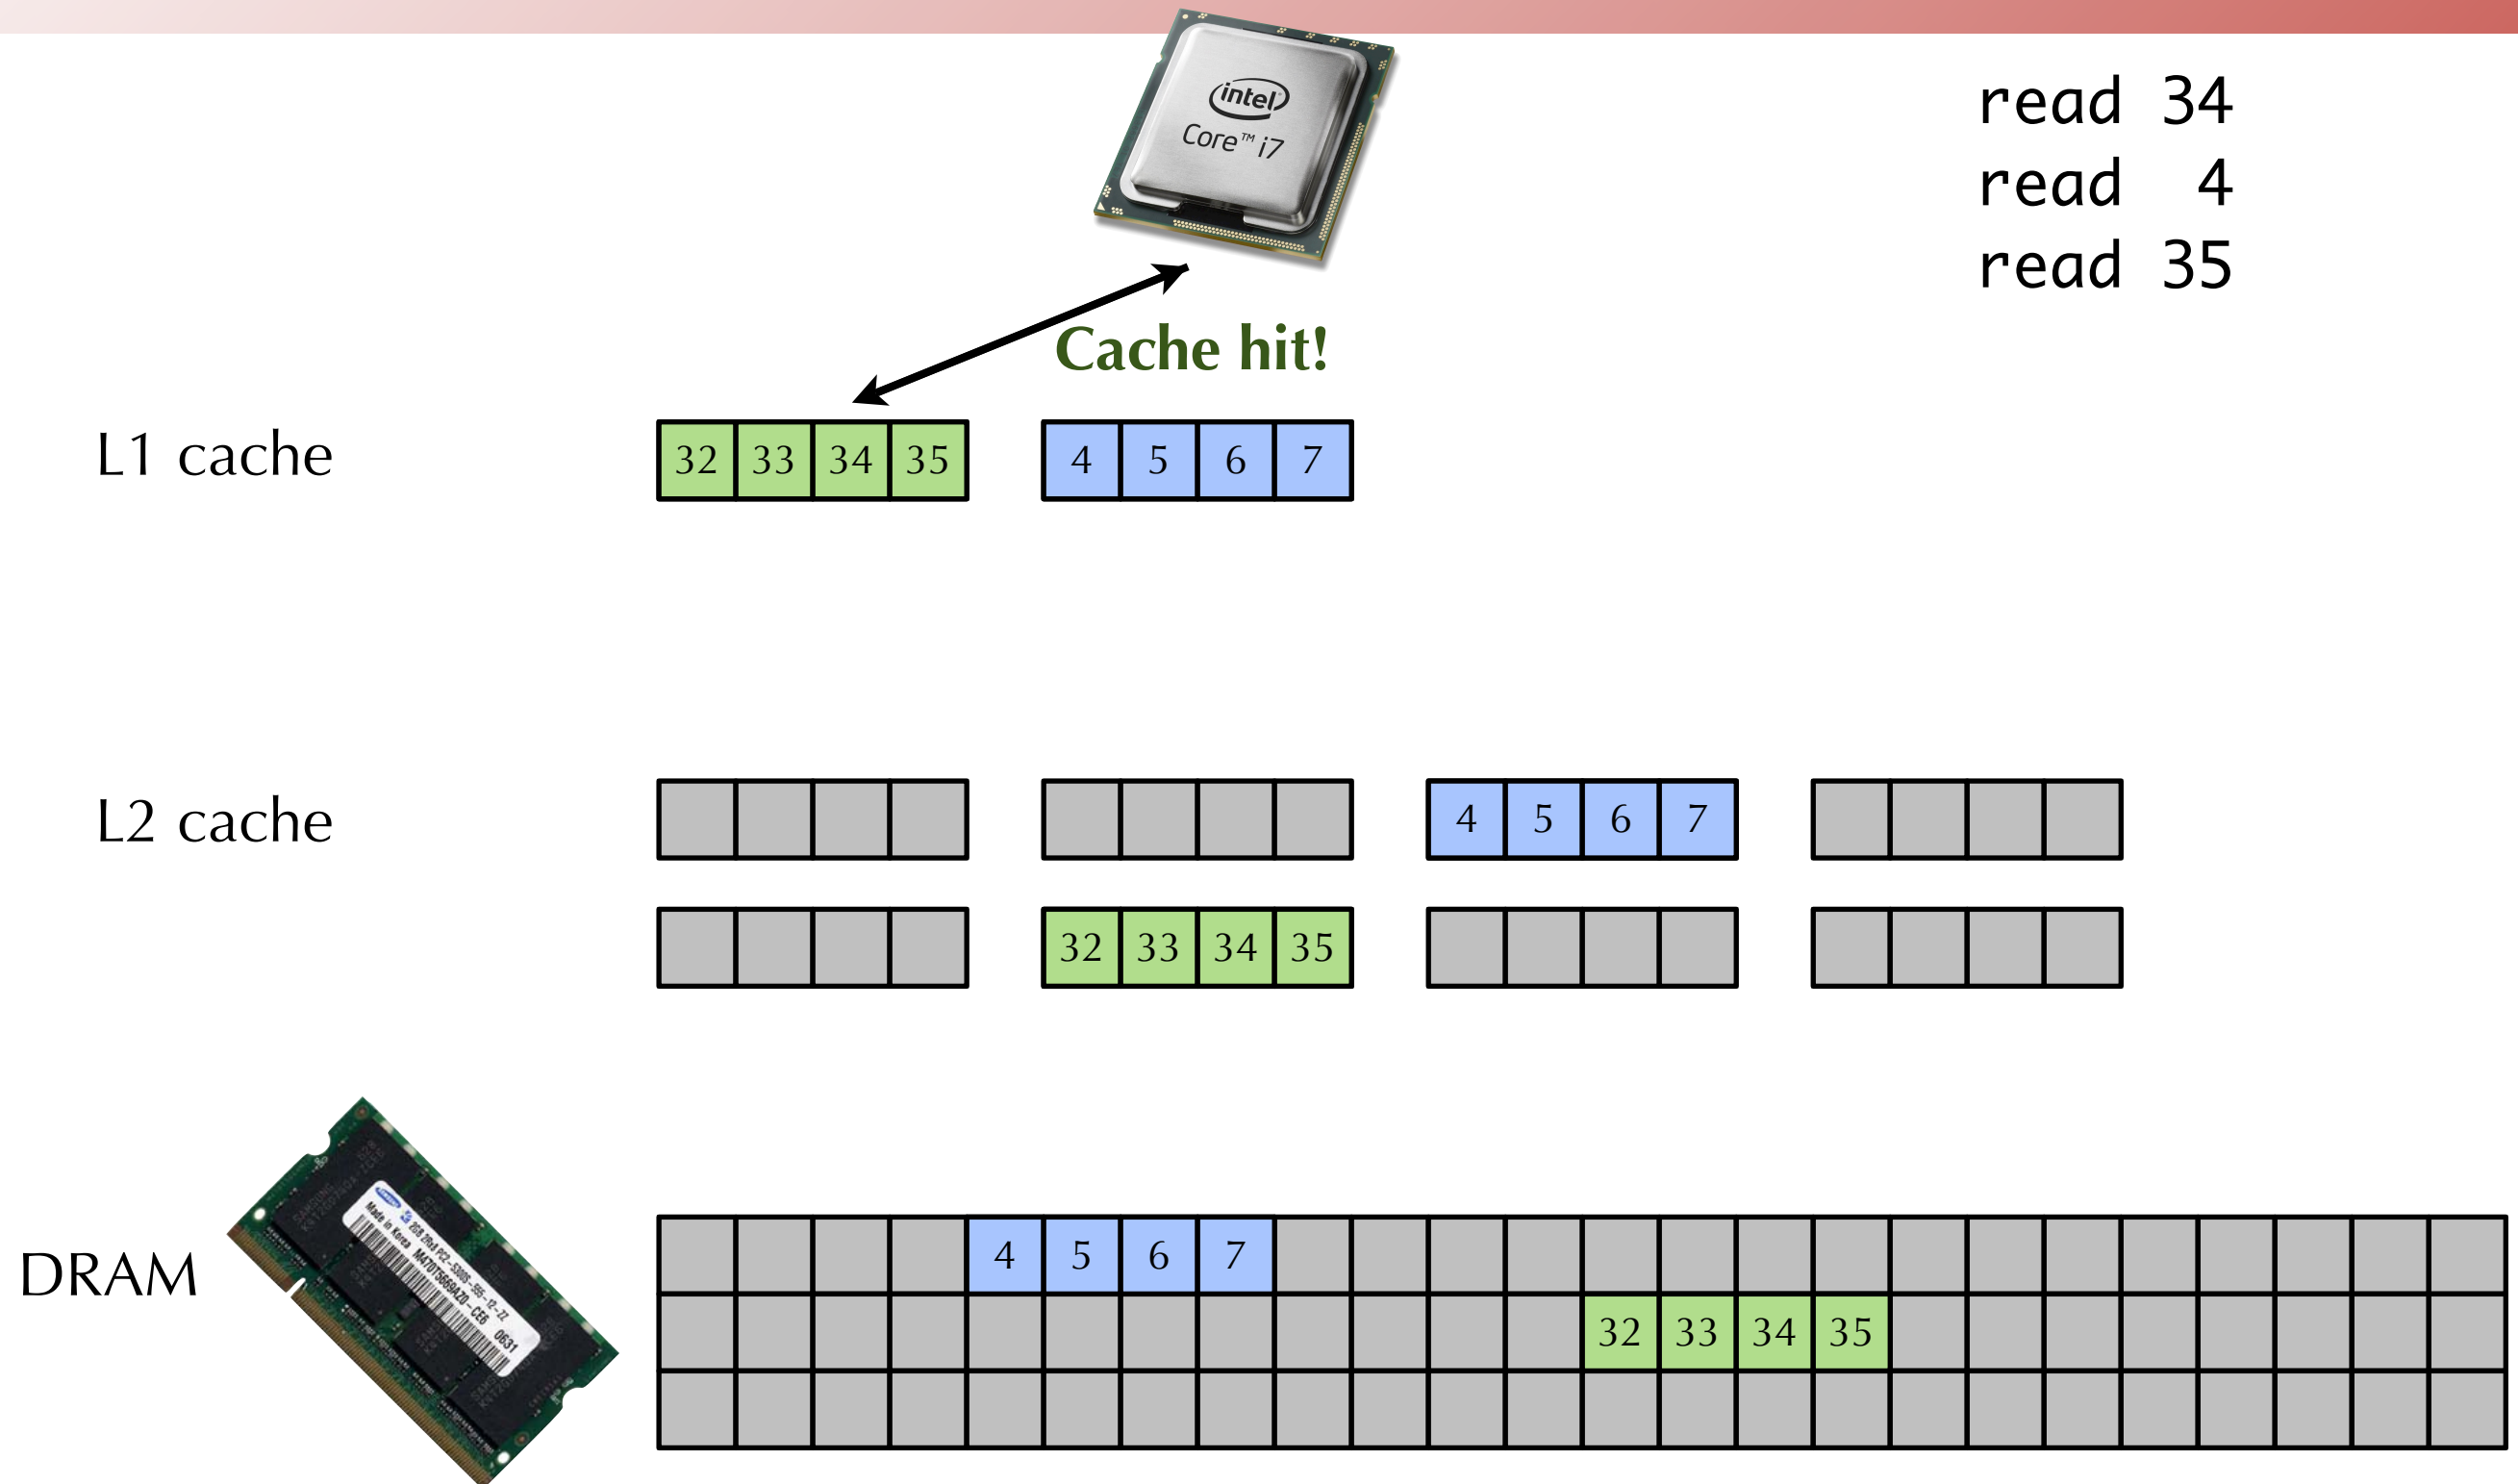
# Caching example



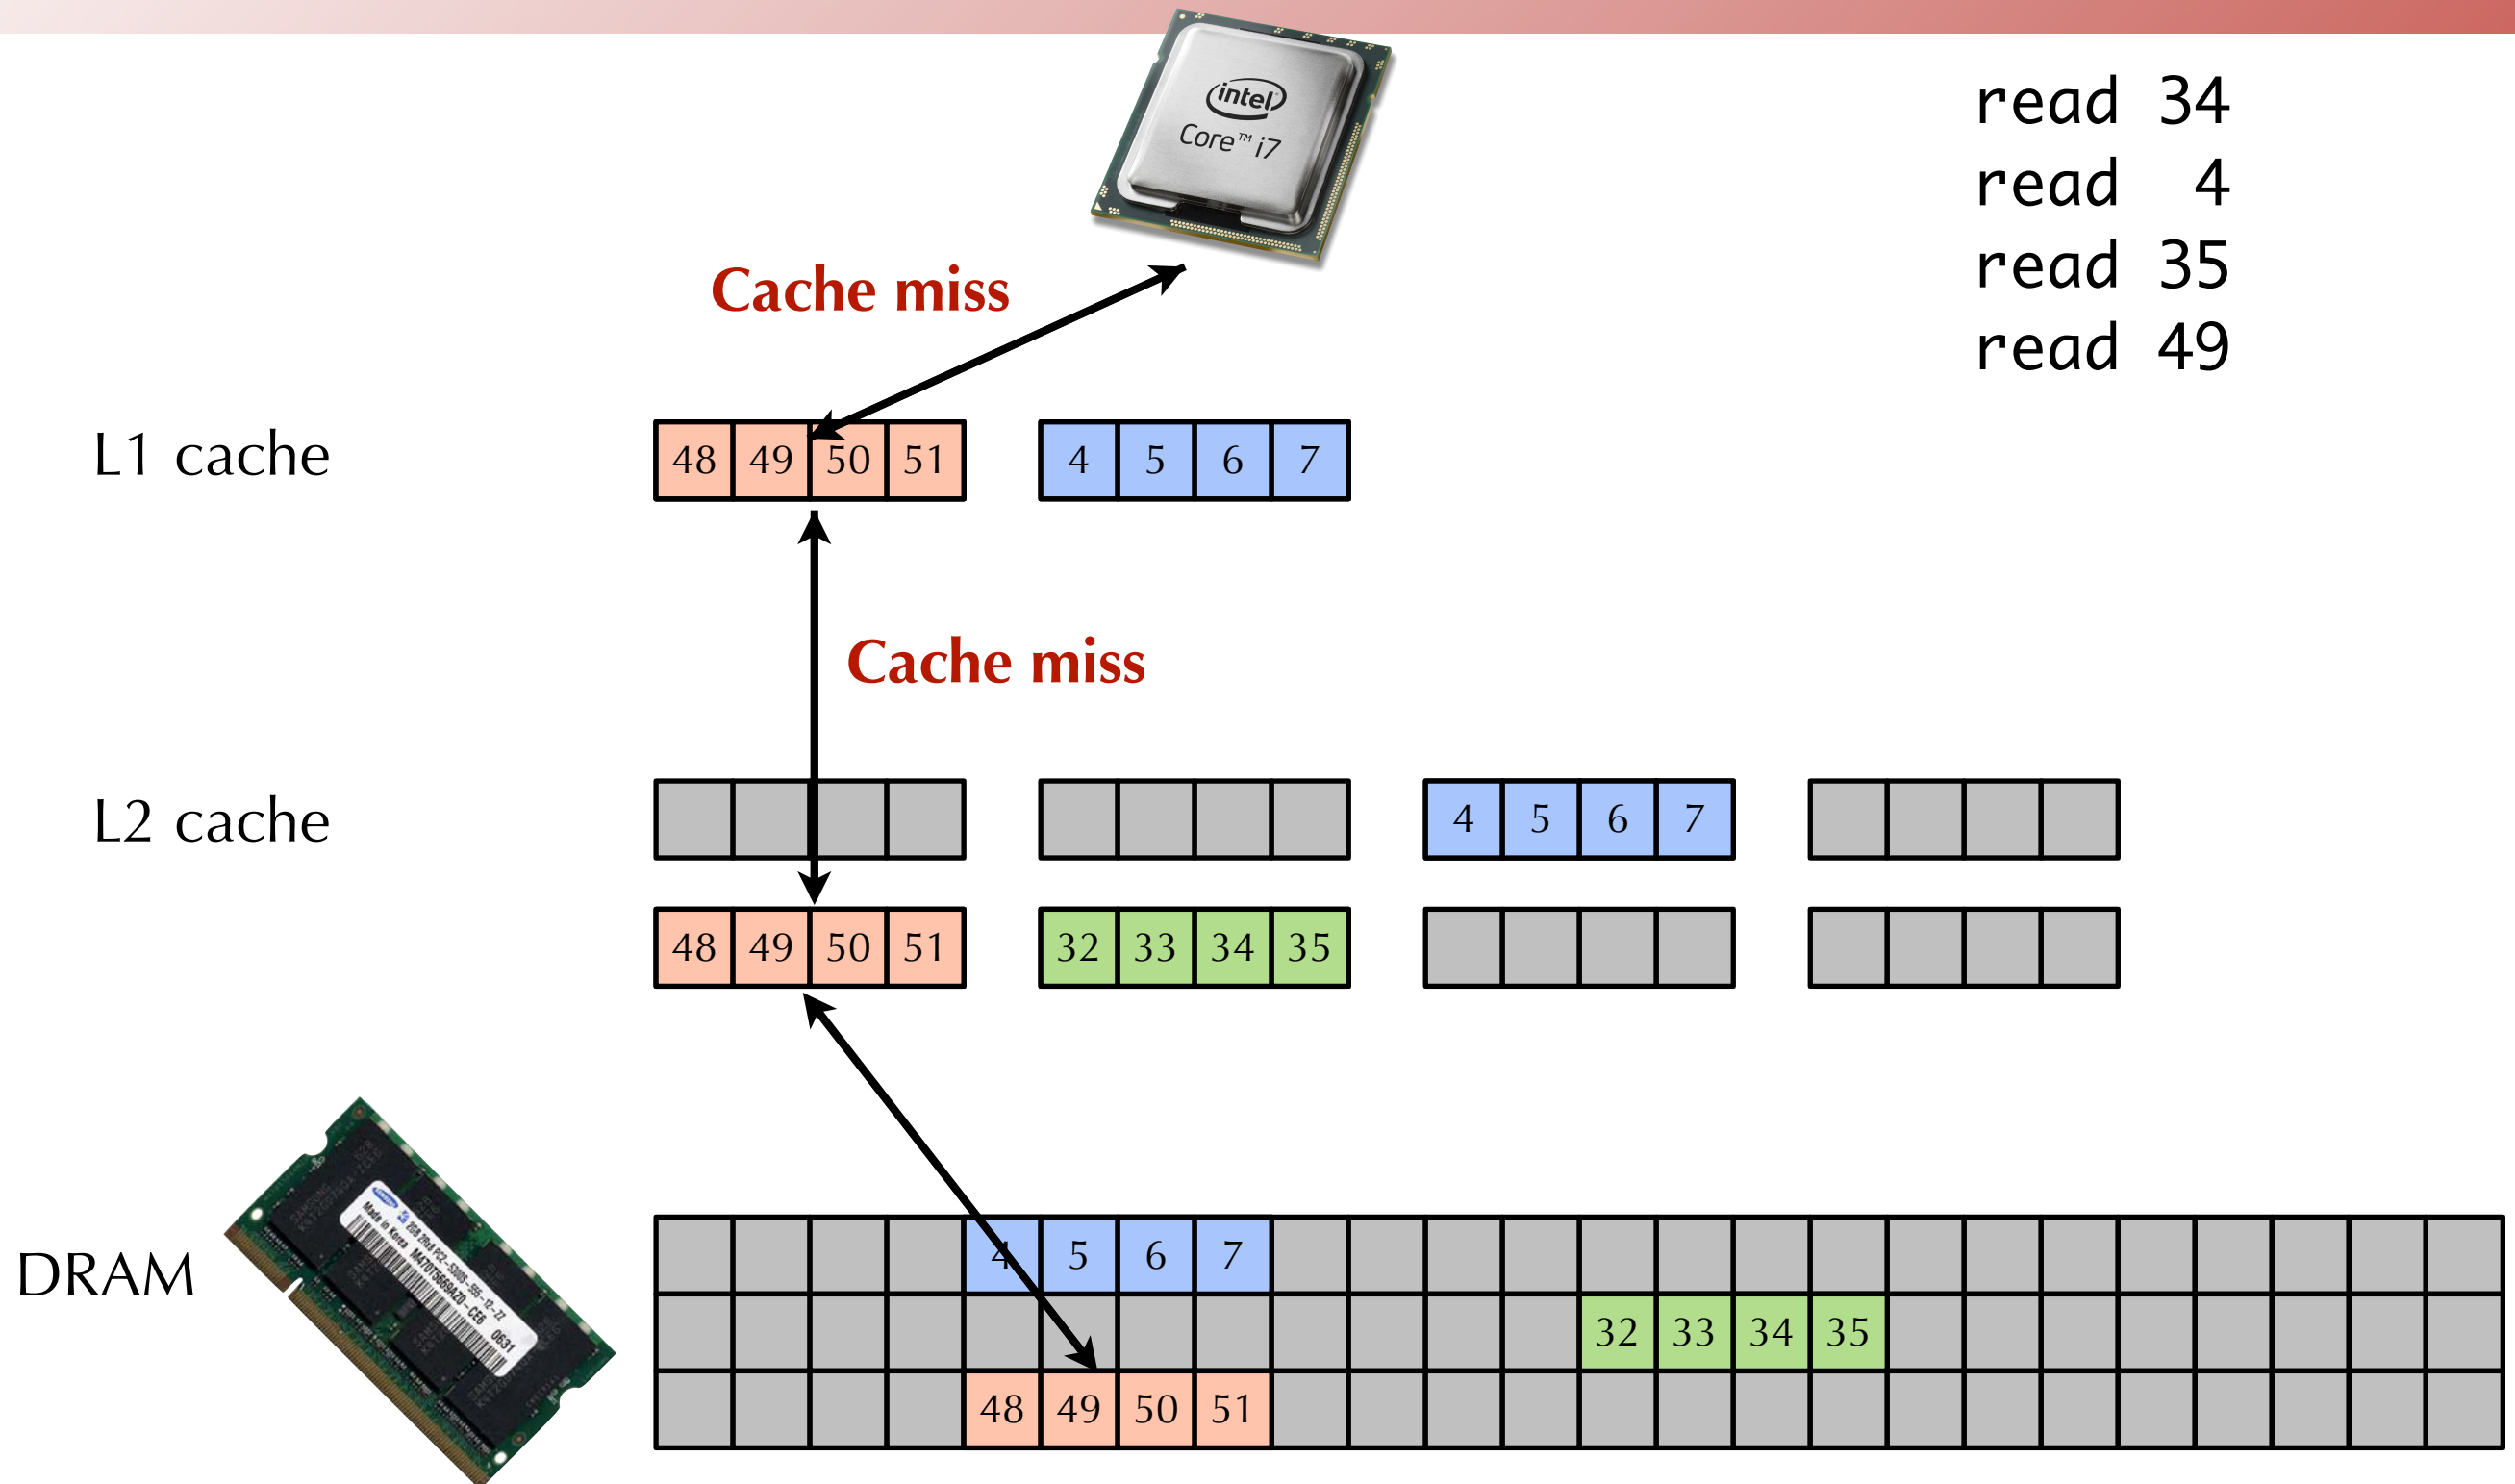
# Caching example



# Caching example

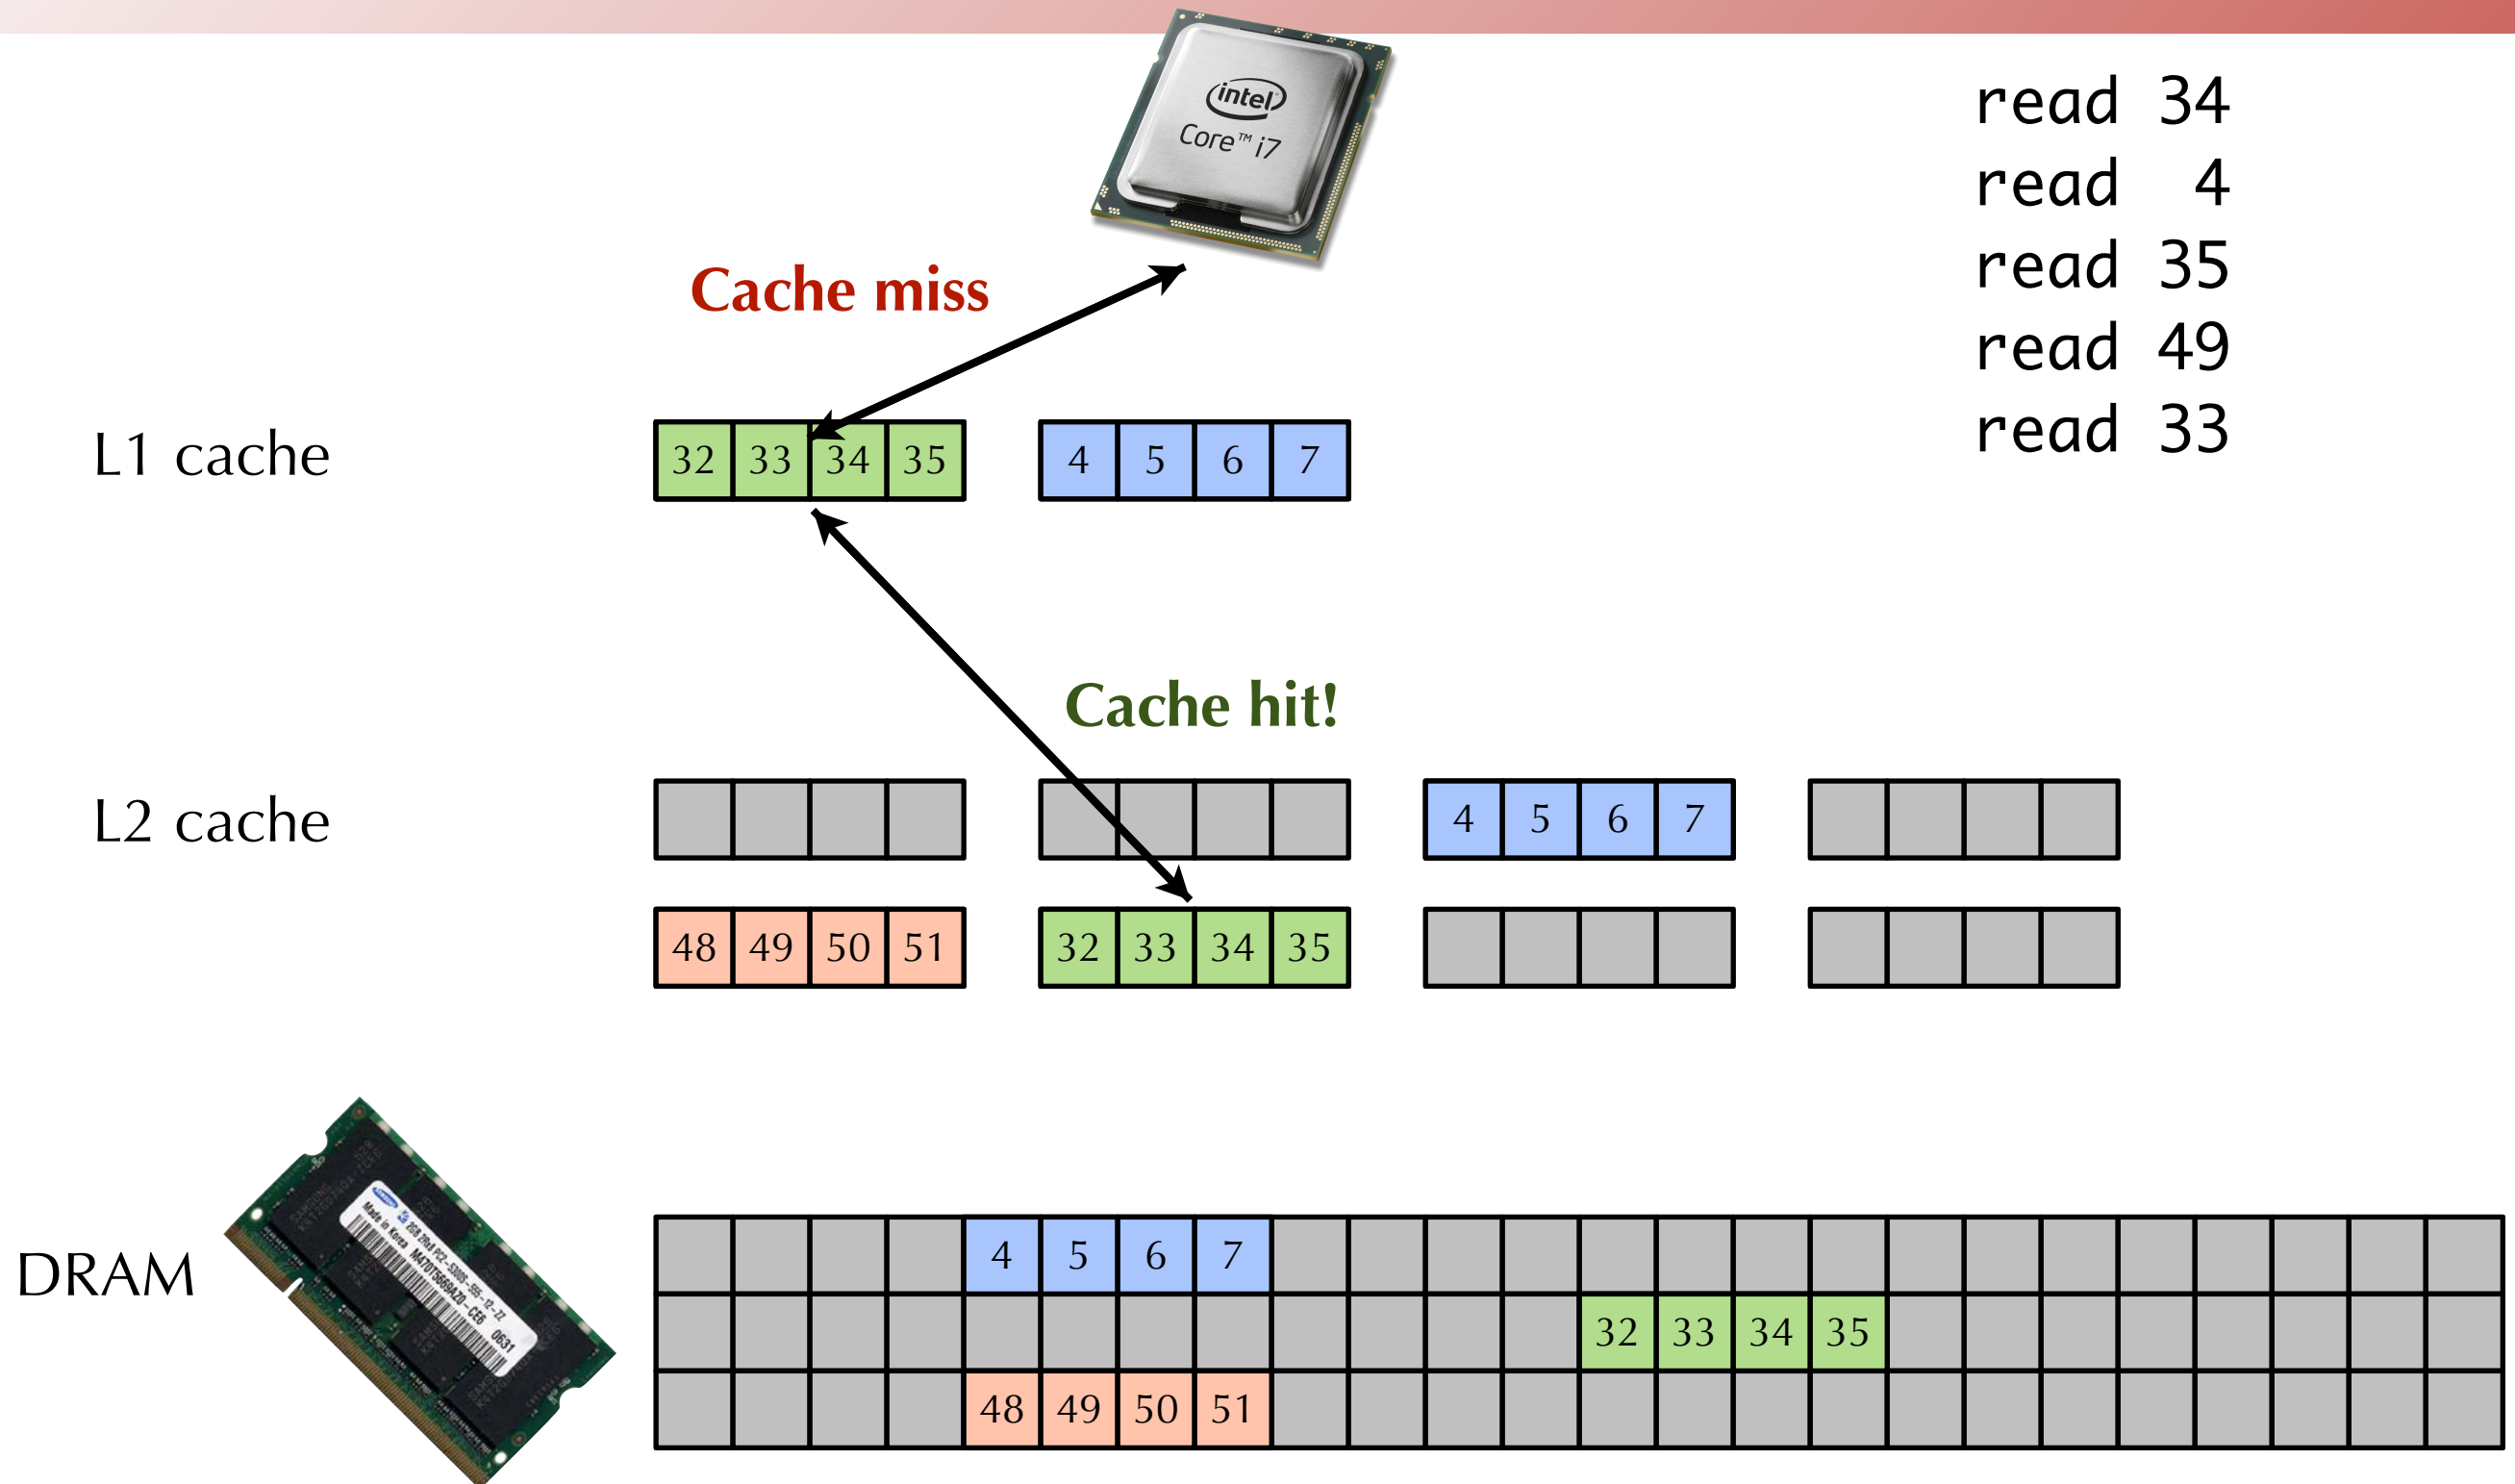


# Caching example





# Caching example



# Caching example: access times

## Register access: ~0.5ns



read 34



read 4

read 35

read 49

read 33

## L1 cache





**L1 cache access: ~1ns**


L2 cache



## L2 cache access: ~3ns



# DRAM



**main memory access: ~60ns**



# A few key points.

- Data transferred from memory to cache (and between cache levels) an entire **cache block** at a time.
  - Amortizes the overhead of a cache miss by getting more data from the lower level.
  - A block typically holds 32 or 64 bytes of data.
- CPU first looks in L1, then L2, then DRAM.
  - Can have a cache miss at one level and a cache hit at another!
- These examples only showed reads from memory.
  - We will talk about writes later.
- Issue #1: When CPU looks in the cache, how does it know what is stored in a given cache block?
- Issue #2: when we have to evict a cache block, which one do we evict?
  - Lots of ways of doing this.

# General Organization of a Cache

Cache is an array of **sets**.

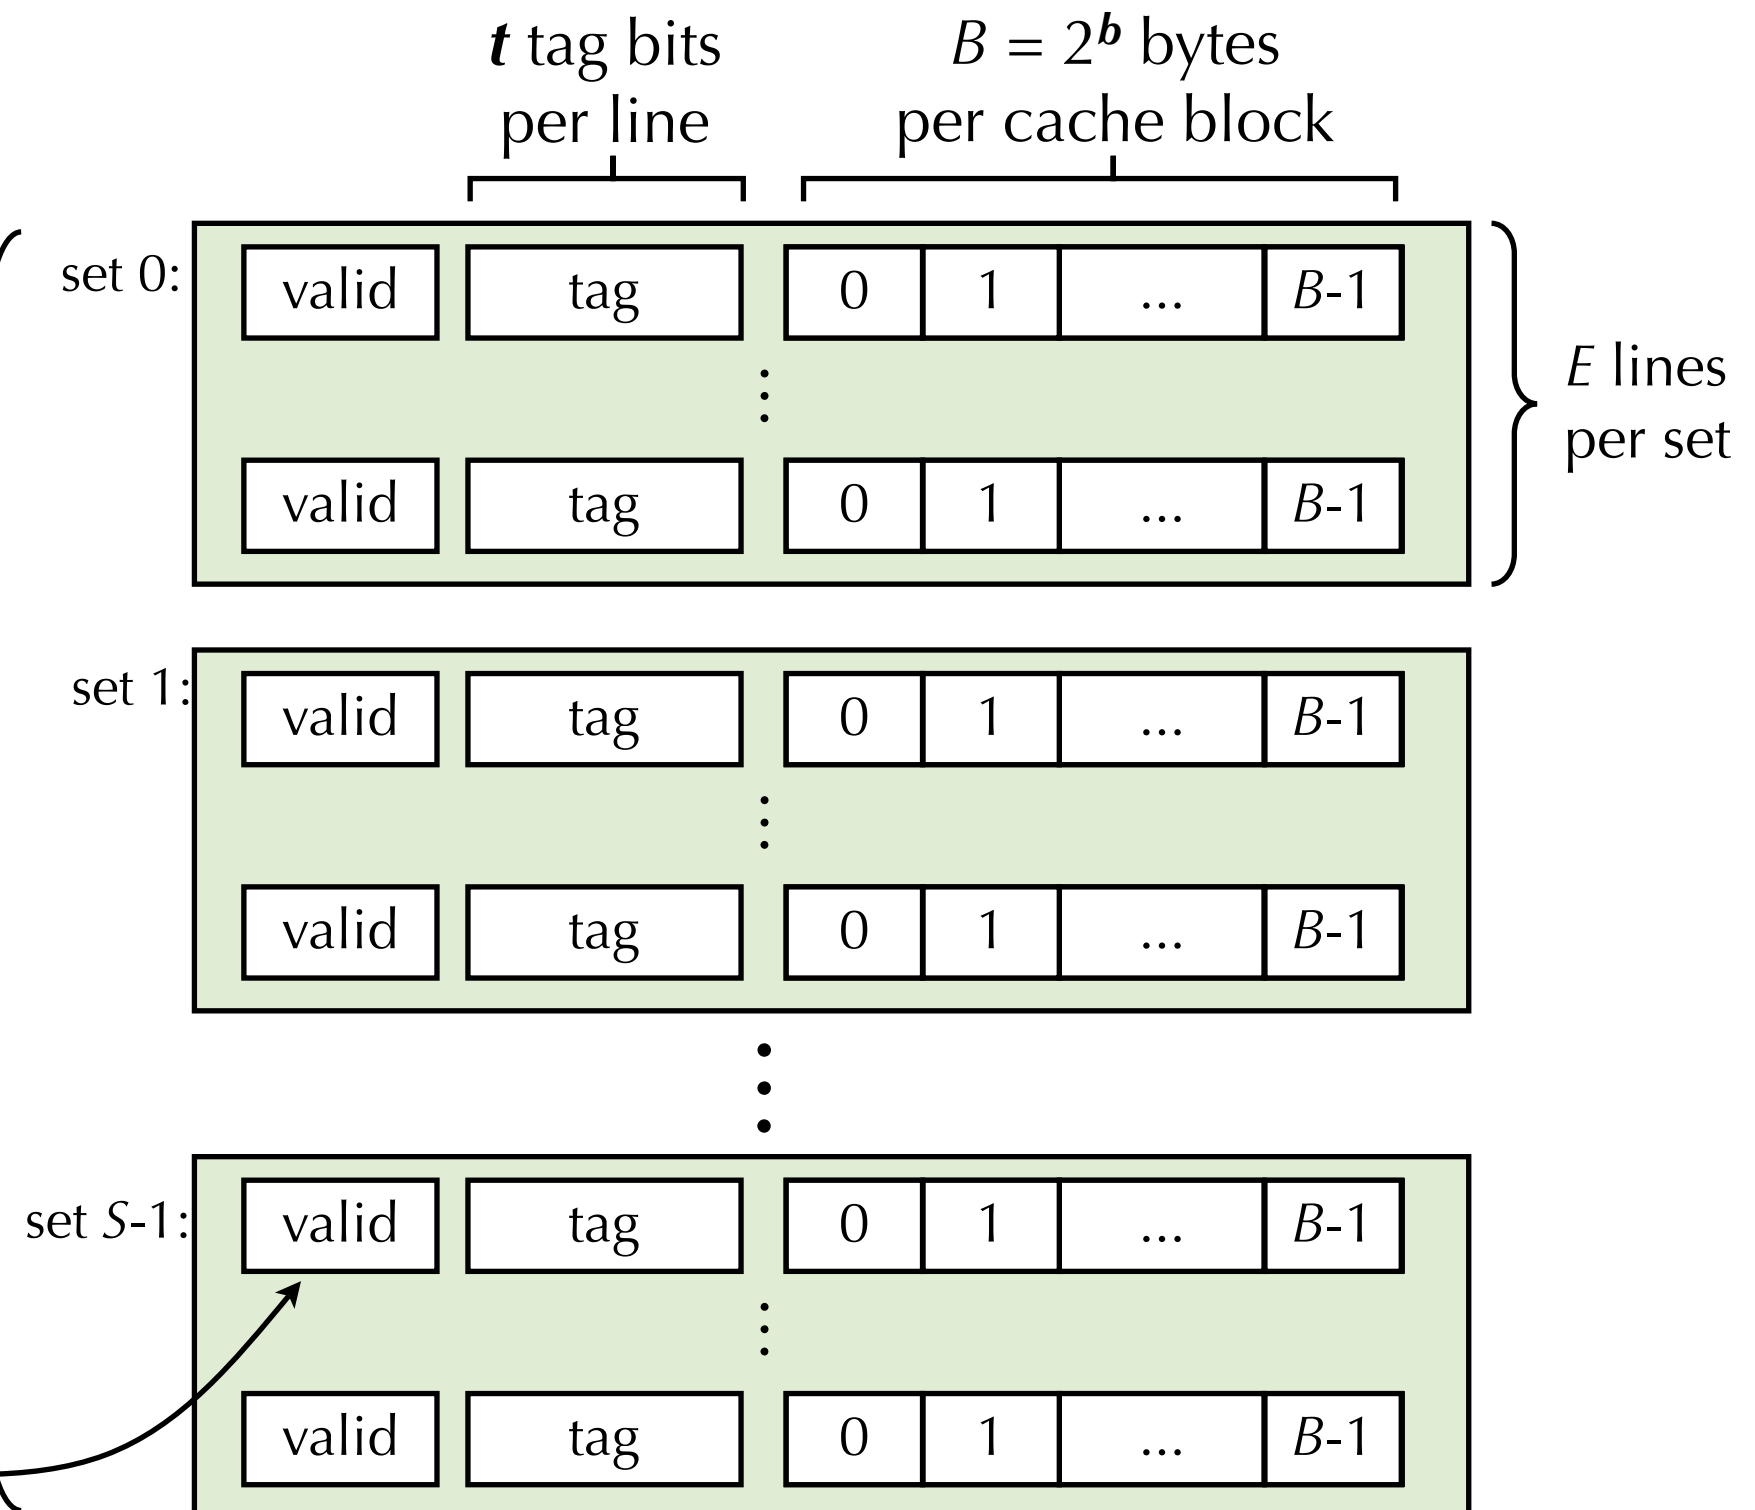
Each set contains one or more **lines**.

Each line holds a **block** of data.

Cache size:

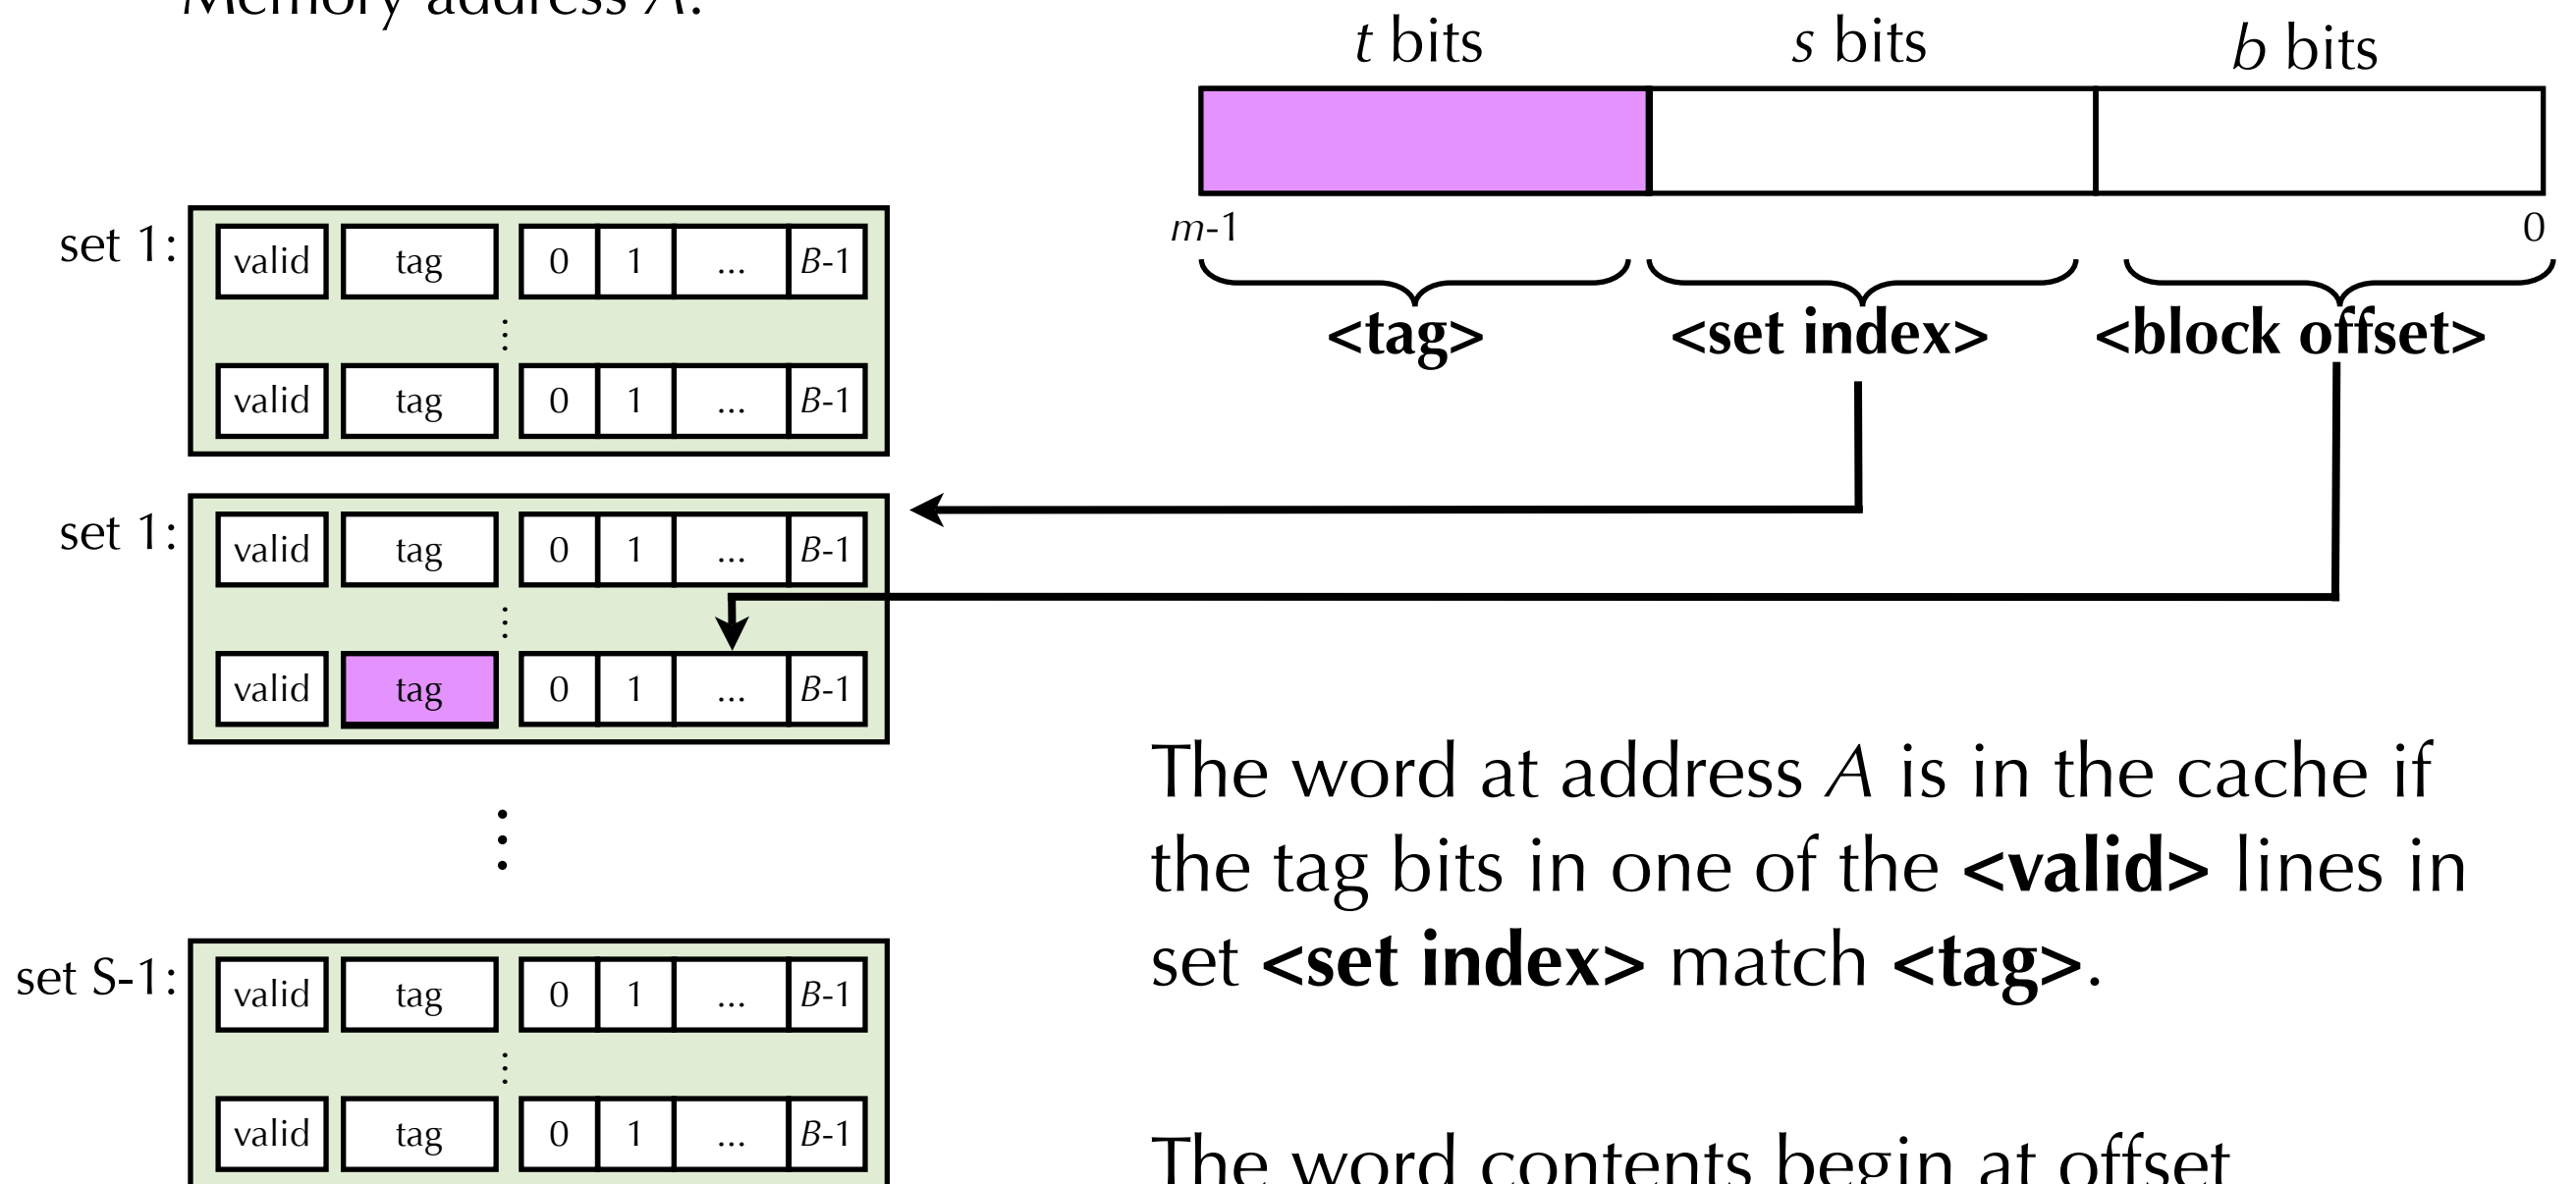
$$C = B \times E \times S \text{ data bytes}$$

1 valid bit per line



# Addressing Caches

Memory address  $A$ :

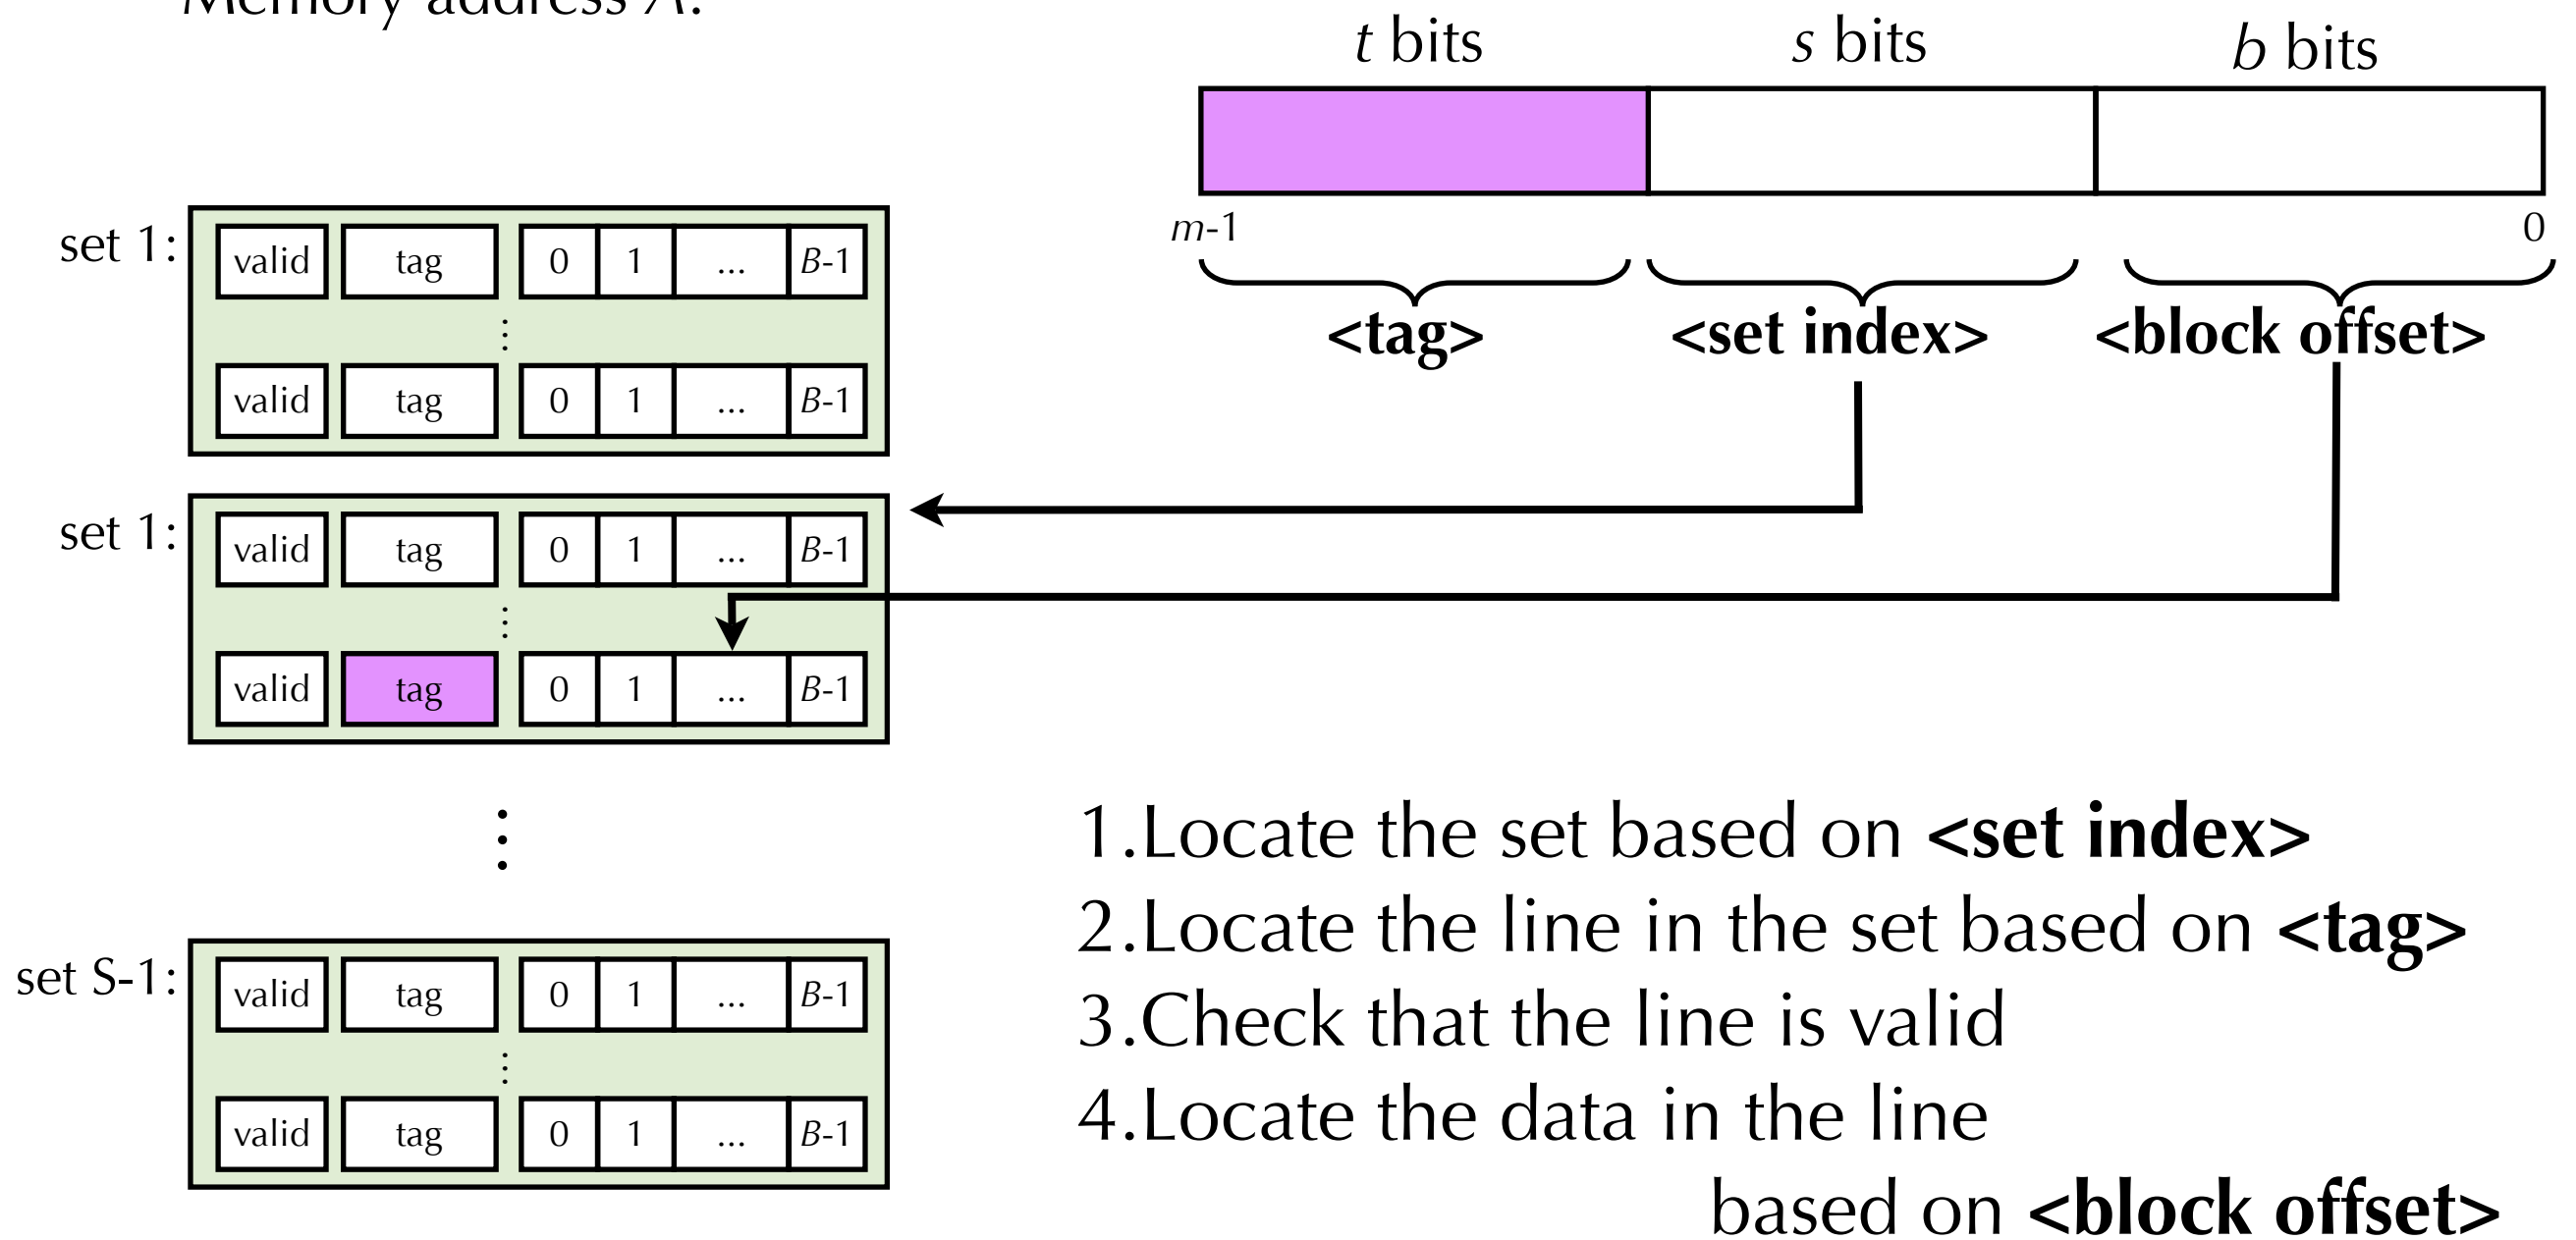


The word at address  $A$  is in the cache if the tag bits in one of the **<valid>** lines in set **<set index>** match **<tag>**.

The word contents begin at offset **<block offset>** bytes from the beginning of the block.

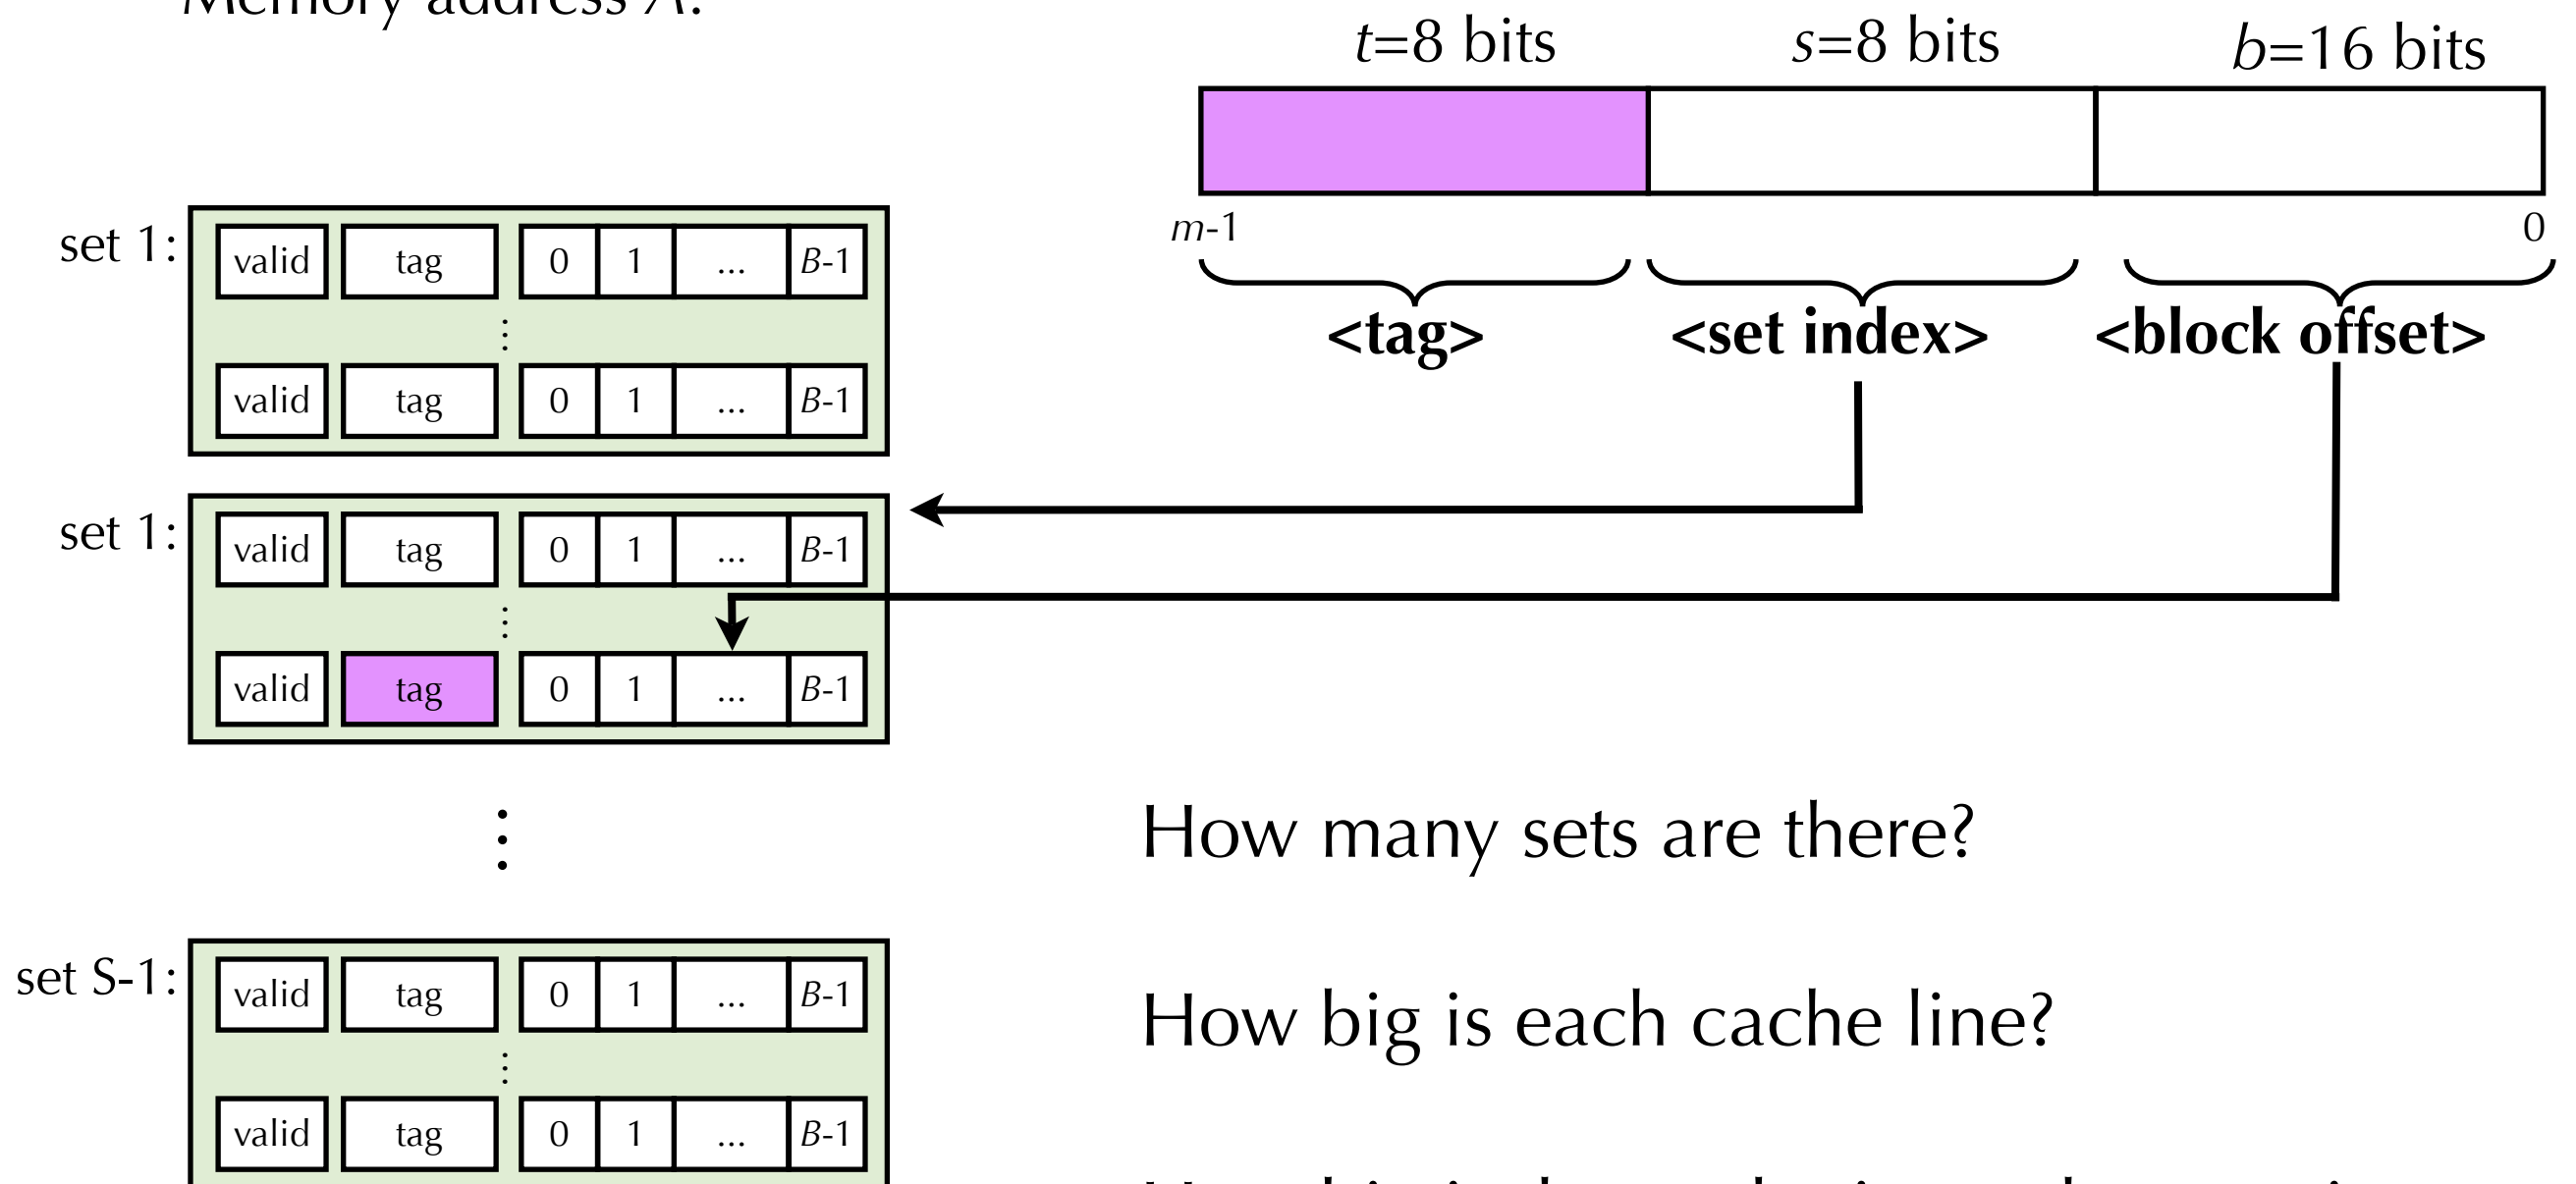
# Addressing Caches

Memory address  $A$ :



# Addressing Caches

Memory address  $A$ :



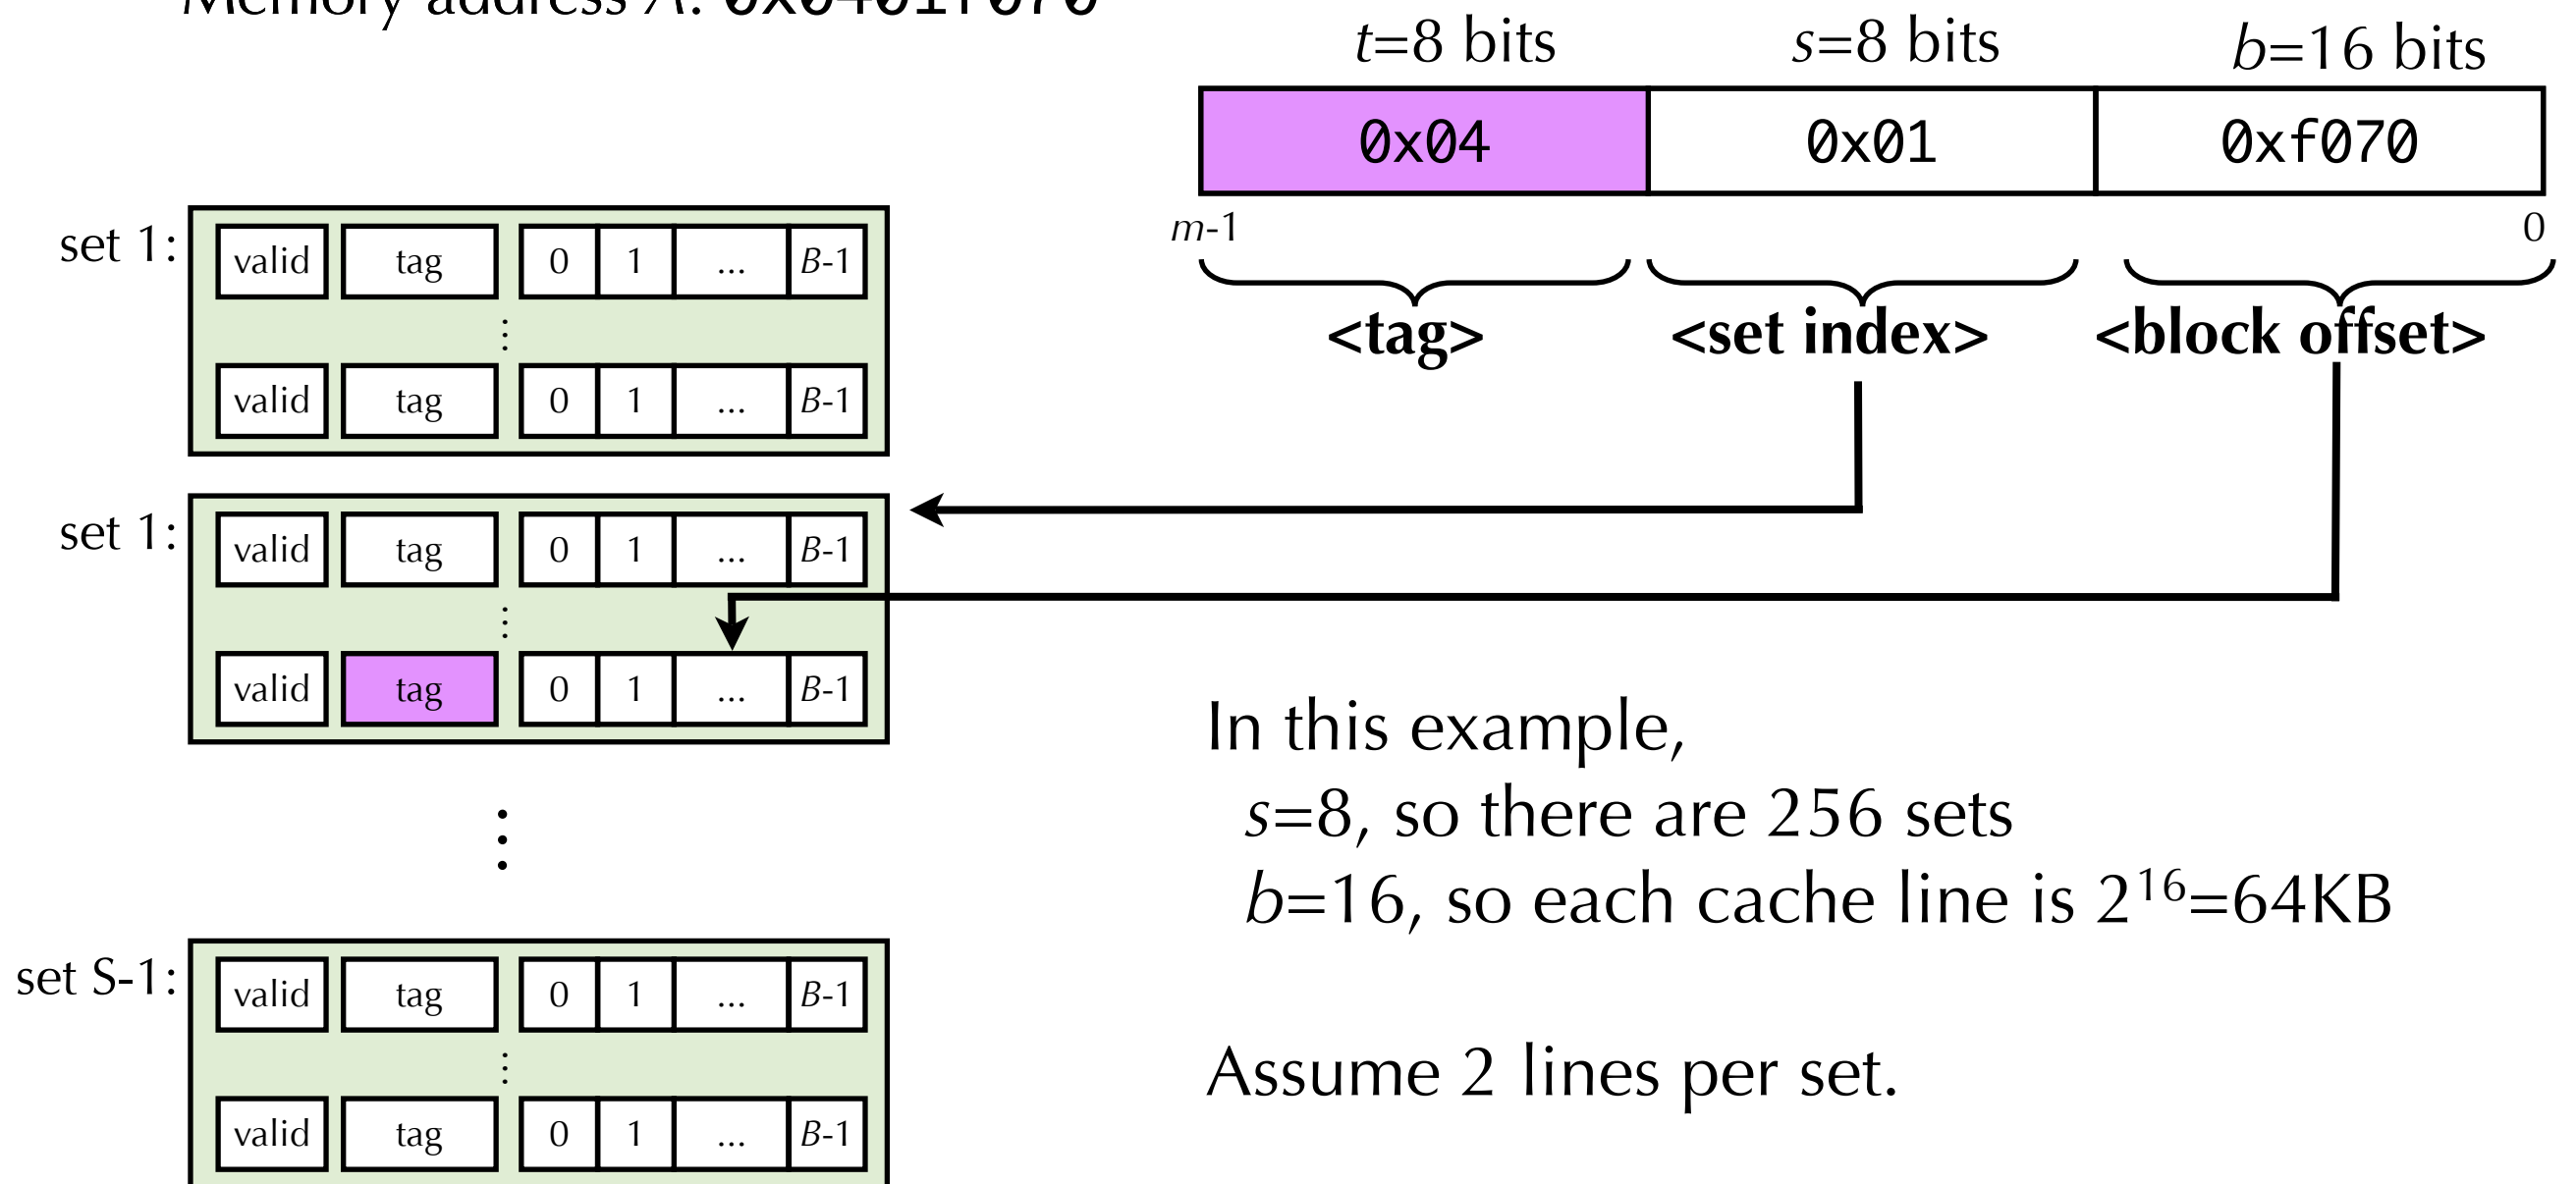
How many sets are there?

How big is each cache line?

How big is the cache in total, assuming two cache lines per set?

# Addressing Caches

Memory address  $A$ : 0x0401f070



In this example,

$s=8$ , so there are 256 sets

$b=16$ , so each cache line is  $2^{16}=64\text{KB}$

Assume 2 lines per set.

Total cache size is  $256 \times 2 \times 64\text{KB} = 32 \text{ MB}$

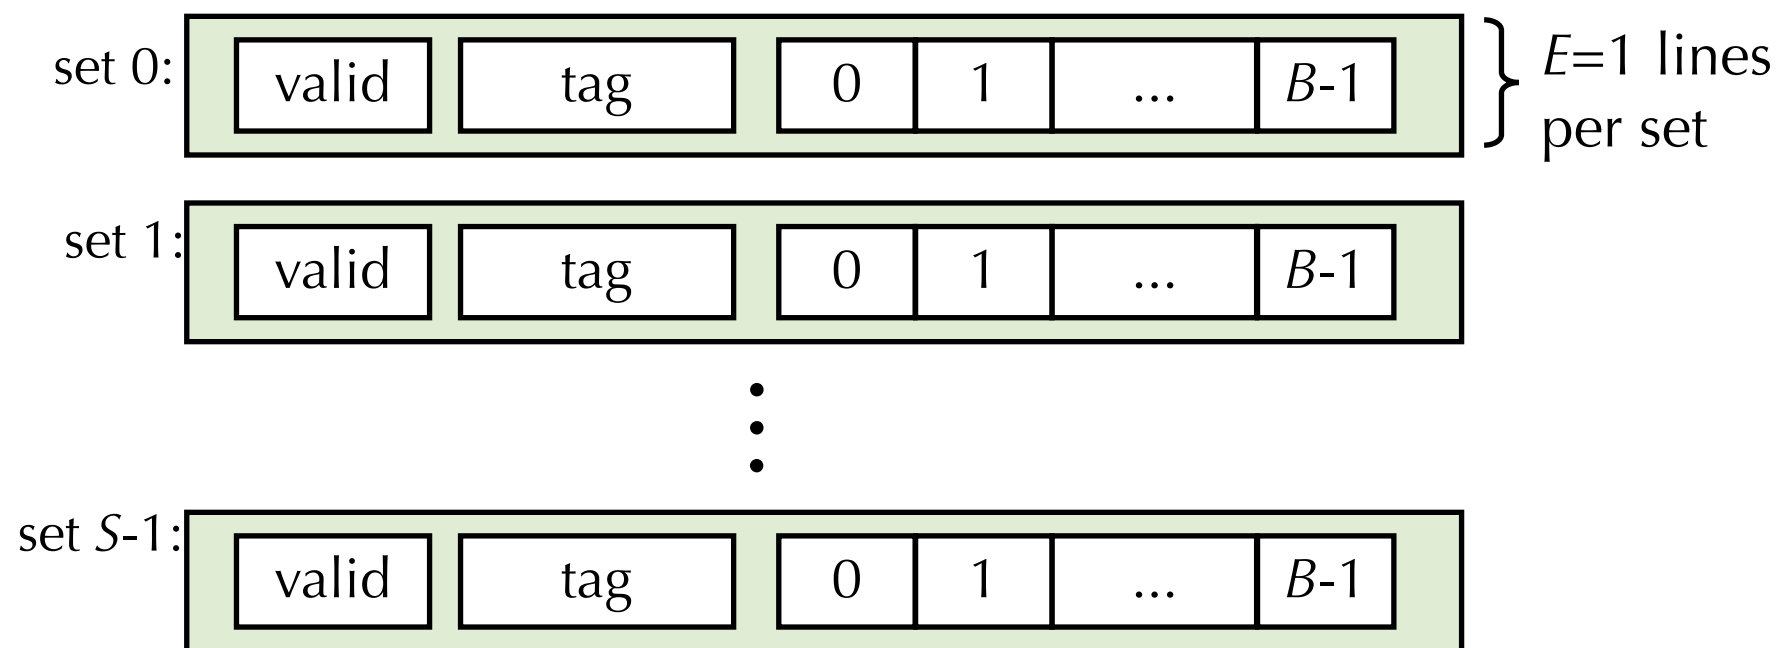
# Topics for today

- The Principle of Locality
- Memory Hierarchies
- Caching Concepts
- Direct-Mapped Cache Organization
- Set-Associative Cache Organization
- Multi-level caches
- Cache writes



# Direct-Mapped Cache

- Each set has exactly one line.
- Simplest kind of cache, and easy to build.
  - Only 1 tag compare required per access

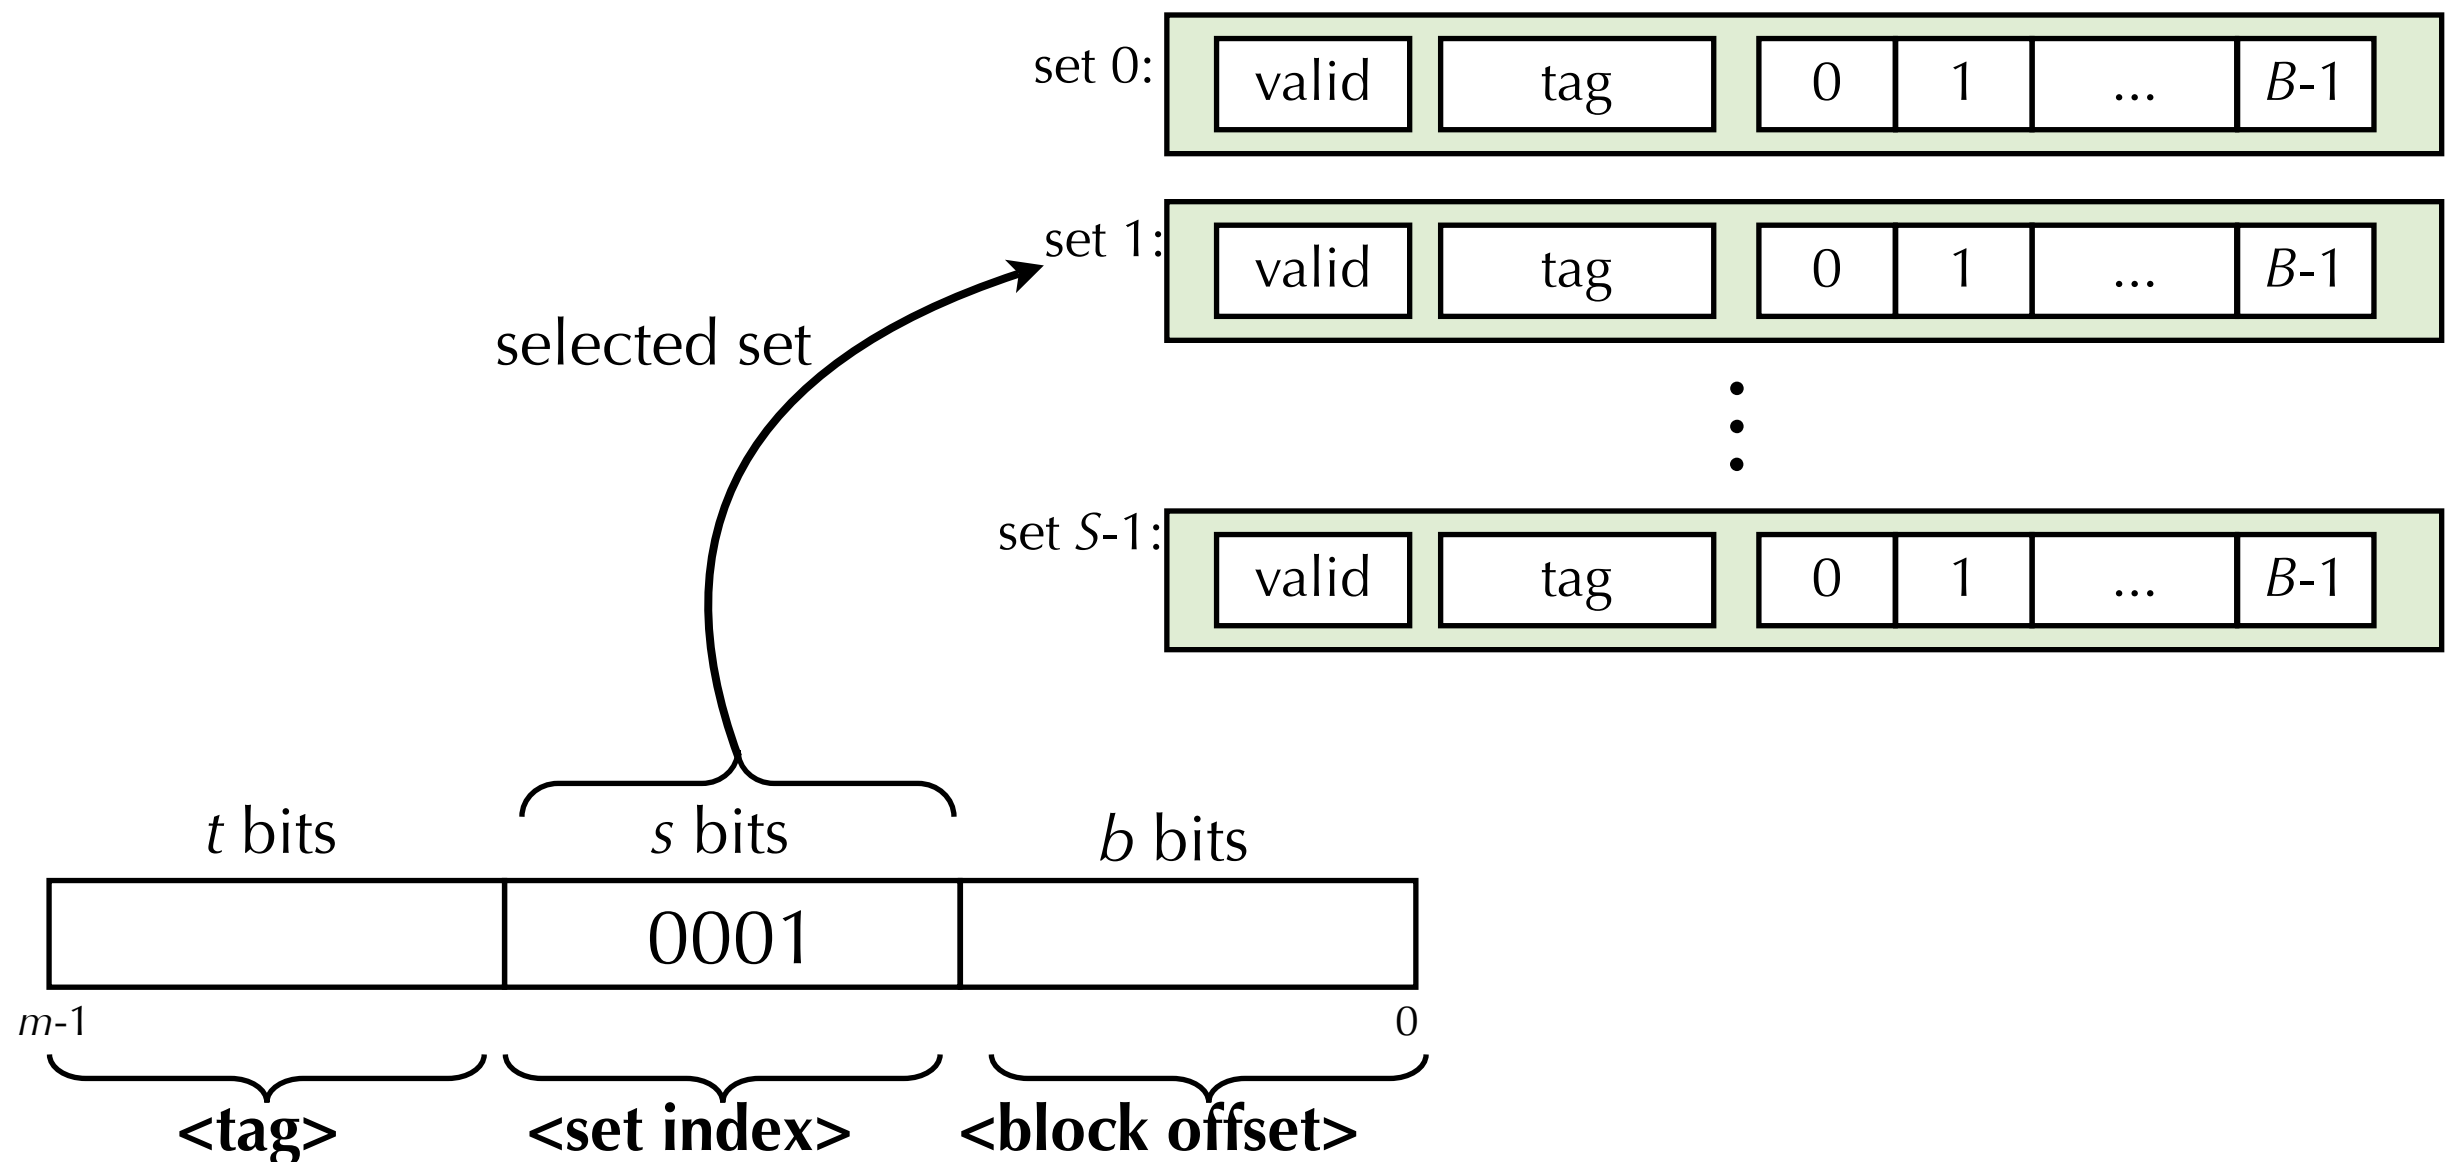


Cache size:  $C = B \times S$  data bytes




# Accessing Direct-Mapped Caches

- Set selection
  - Use the set index bits to determine the set of interest.



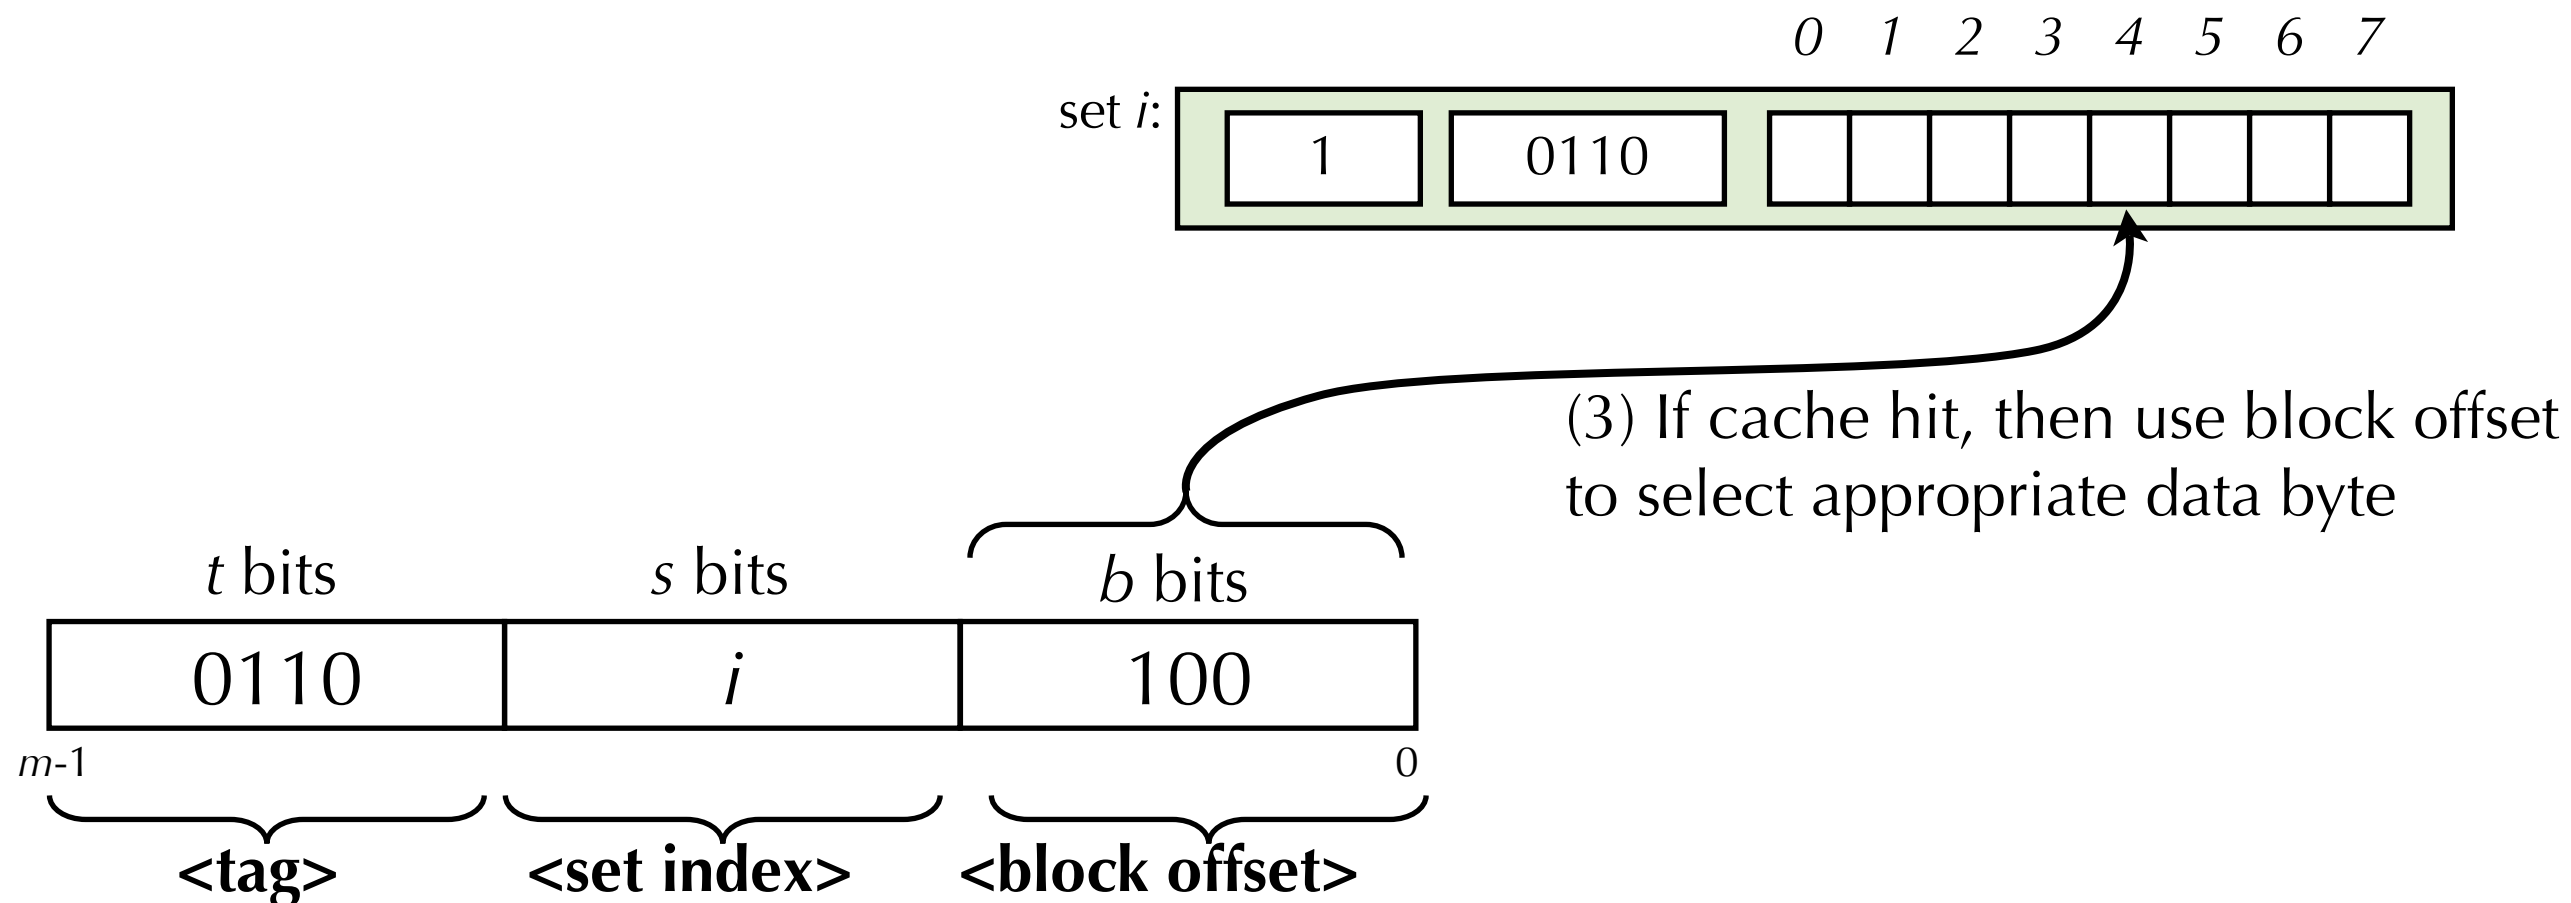
# Accessing Direct-Mapped Caches

- Line matching and word selection
  - **Line matching:** Find a *valid* line in the selected set with a *matching tag*
  - **Byte selection:** Then extract the byte that we want



# Accessing Direct-Mapped Caches

- Line matching and word selection
  - **Line matching:** Find a *valid* line in the selected set with a *matching tag*
  - **Byte selection:** Then extract the byte that we want



# Handling cache misses

- A **cache miss** occurs if:
  - Tag does not match the corresponding line in the cache, or
  - Valid bit in the line is not set

# Handling cache misses

- In event of a cache miss for address  $A$ :
  - 1) Read an entire **block** from next-lowest cache level
    - Take address  $A$ , set the  $b$  least significant bits to zero
    - E.g., suppose  $A = 0x43b7$  and  $b=4$  bits (block size =  $2^4$  bytes = 16 bytes)
      - ▶ Base address of block is  $0x43b0$
      - ▶ Block contains contents of addresses  $0x43b0$  through  $0x43bf$
  - 2) Insert block into the corresponding cache line
    - May require **evicting** contents of cache line, if line was valid already
  - 3) Set the tag bits corresponding to address  $A$
  - 4) Set valid bit to 1

# Types of cache misses

- **Cold (compulsory) miss**

- A miss that occurs because the cache is initially empty.
- Nothing you can do about it: thus, “compulsory”.

- **Conflict miss**

- When reading something into the cache requires that we evict something else.
- BUT... the cache is not full!
- Example: If set index matches a line that is already occupied, have to evict what is already there (even if there are empty blocks in the cache!)

- **Capacity miss**

- Occurs when the cache is too small to store everything currently being accessed by the program.
- **Working set:** Set of “actively used” cache lines
- The working set varies over time, depending on what the program is doing.


# Direct-Mapped Cache Simulation

- Simple example:
  - 4 bit addresses (memory size = 16 bytes)
  - 1 tag bit, 2 set index bits, 1 block offset bit

Address trace (reads):

$$A = 0 = \textcolor{red}{0}\textcolor{orange}{0}\textcolor{orange}{0}\textcolor{blue}{0}_2$$

**cache miss**




|       | valid | tag      | data       |
|-------|-------|----------|------------|
| Set 0 | 1     | <b>0</b> | M[0], M[1] |
| Set 1 | 0     |          |            |
| Set 2 | 0     |          |            |
| Set 3 | 0     |          |            |



# Direct-Mapped Cache Simulation

- Simple example:
  - 4 bit addresses (memory size = 16 bytes)
  - 1 tag bit, 2 set index bits, 1 block offset bit



Address trace (reads):

$$A = 0 = 0000_2$$

cache miss

$$A = 1 = 0001_2$$


cache hit

|       | valid | tag | data       |
|-------|-------|-----|------------|
| Set 0 | 1     | 0   | M[0], M[1] |
| Set 1 | 0     |     |            |
| Set 2 | 0     |     |            |
| Set 3 | 0     |     |            |



# Direct-Mapped Cache Simulation

- Simple example:
  - 4 bit addresses (memory size = 16 bytes)
  - 1 tag bit, 2 set index bits, 1 block offset bit



Address trace (reads):

$$A = 0 = 0000_2$$

cache miss

$$A = 1 = 0001_2$$

cache hit


$$A = 7 = 0111_2$$

cache miss

|       | valid | tag | data       |
|-------|-------|-----|------------|
| Set 0 | 1     | 0   | M[0], M[1] |
| Set 1 | 0     |     |            |
| Set 2 | 0     |     |            |
| Set 3 | 1     | 0   | M[6], M[7] |

# Direct-Mapped Cache Simulation

- Simple example:
  - 4 bit addresses (memory size = 16 bytes)
  - 1 tag bit, 2 set index bits, 1 block offset bit



Address trace (reads):

$$A = 0 = 0000_2$$

cache miss

$$A = 1 = 0001_2$$

cache hit

$$A = 7 = 0111_2$$

cache miss


$$A = 8 = 1000_2$$

cache miss + evict

|       | valid | tag | data       |
|-------|-------|-----|------------|
| Set 0 | 1     | 1   | M[8], M[9] |
| Set 1 | 0     |     |            |
| Set 2 | 0     |     |            |
| Set 3 | 1     | 0   | M[6], M[7] |

# Direct-Mapped Cache Simulation

- Simple example:
  - 4 bit addresses (memory size = 16 bytes)
  - 1 tag bit, 2 set index bits, 1 block offset bit



Address trace (reads):

$$A = 0 = 0000_2$$

cache miss

$$A = 1 = 0001_2$$

cache hit

$$A = 7 = 0111_2$$

cache miss

$$A = 8 = 1000_2$$

cache miss + evict

$$A = 0 = 0000_2$$

cache miss + evict


|       | valid | tag | data       |
|-------|-------|-----|------------|
| Set 0 | 1     | 0   | M[0], M[1] |
| Set 1 | 0     |     |            |
| Set 2 | 0     |     |            |
| Set 3 | 1     | 0   | M[6], M[7] |

# Topics for today

- The Principle of Locality
- Memory Hierarchies
- Caching Concepts
- Direct-Mapped Cache Organization
- Set-Associative Cache Organization
- Multi-level caches
- Cache writes

# Set Associative Caches

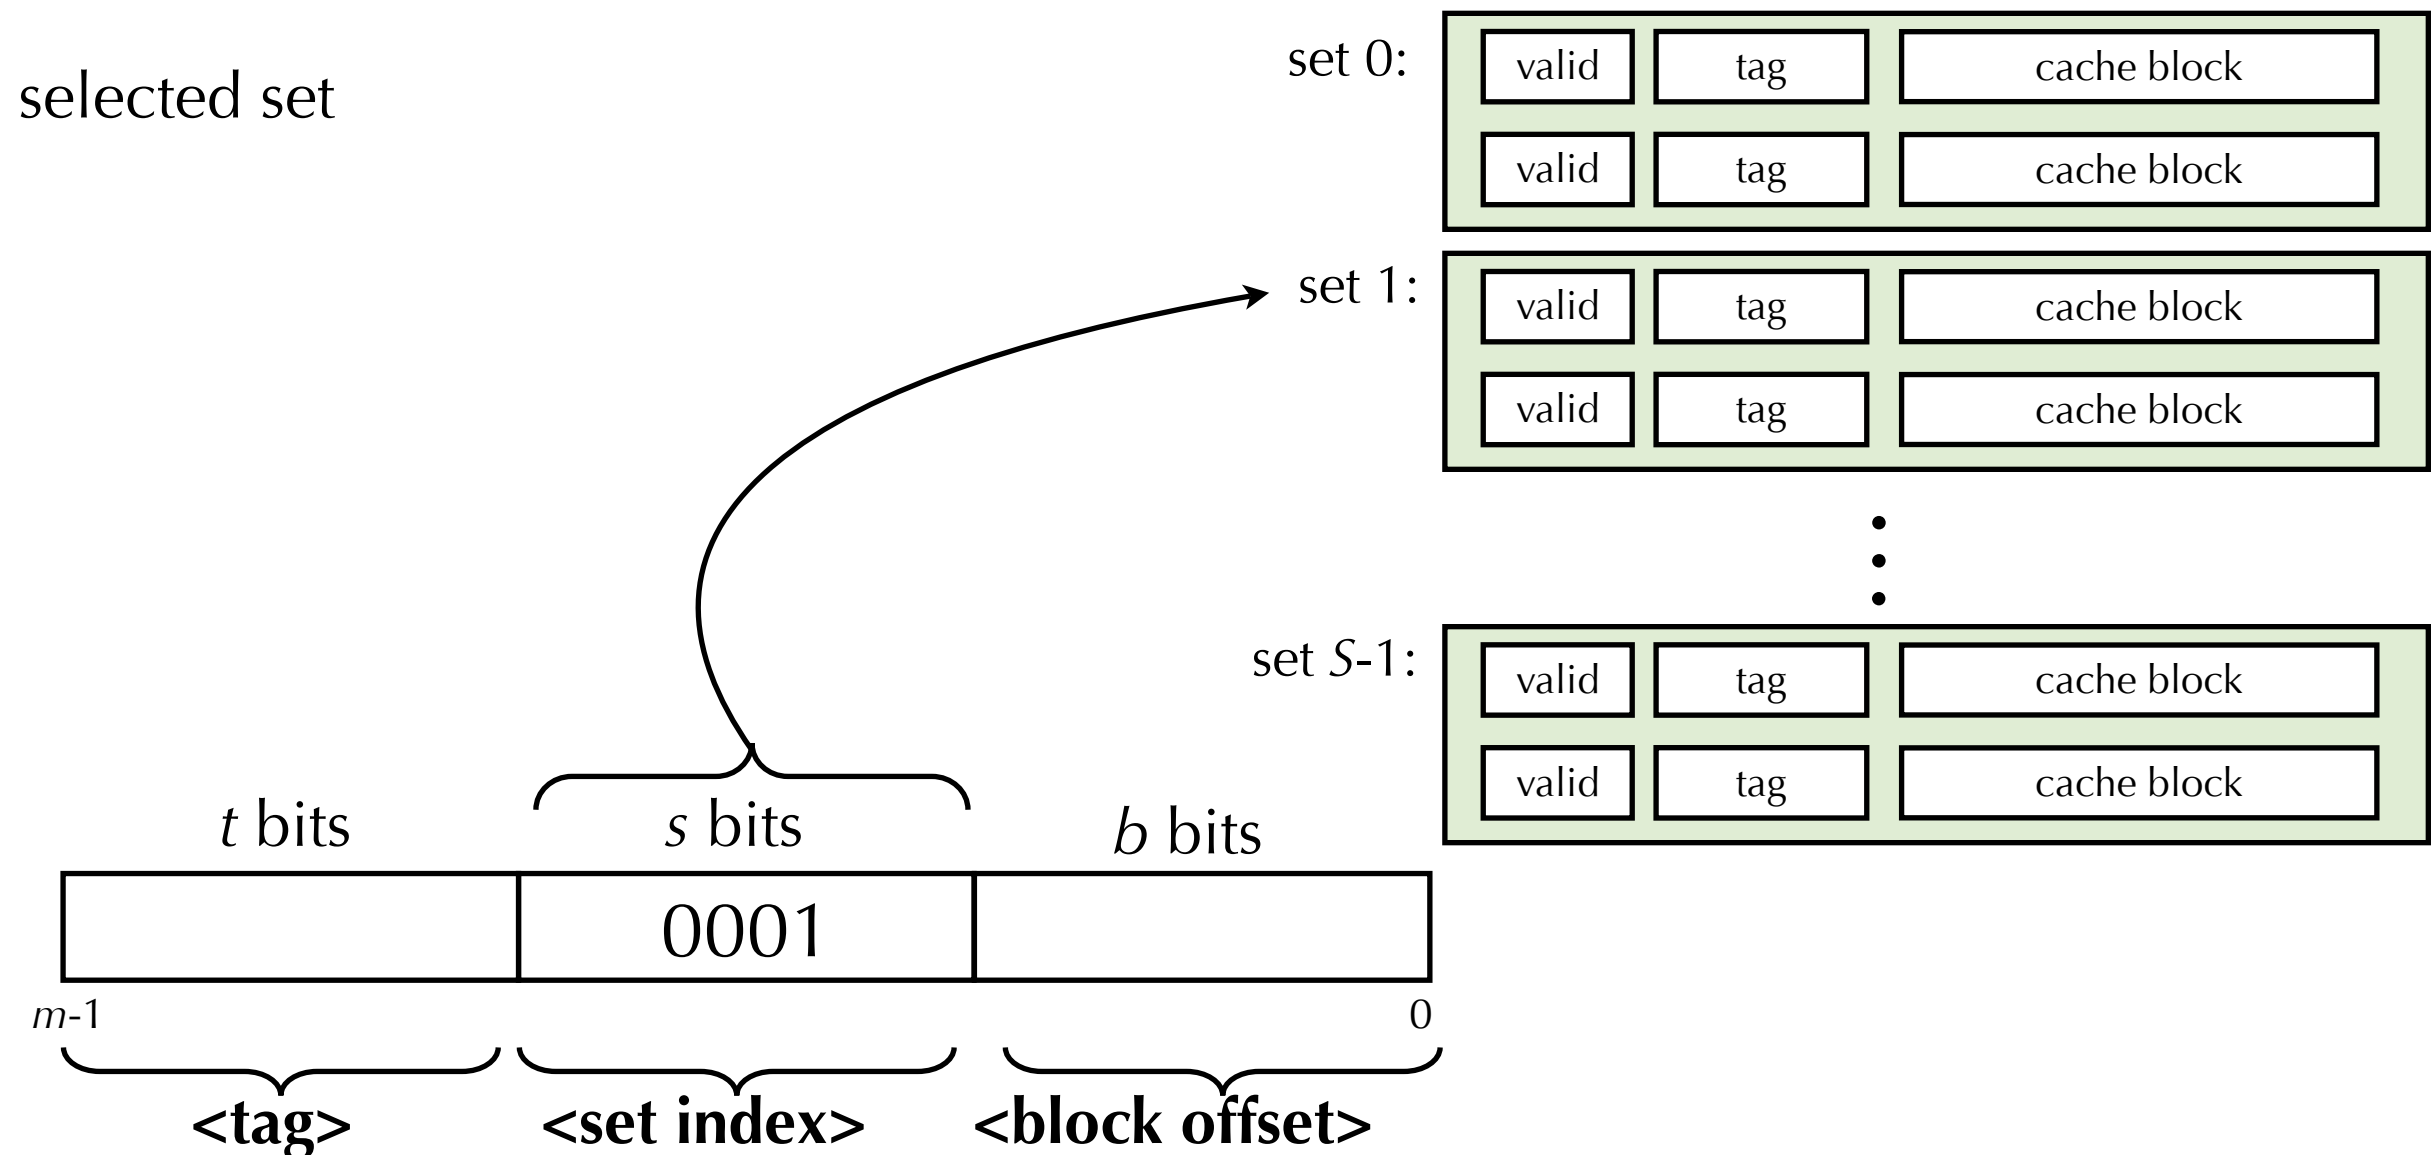
- Characterized by more than one line per set
  - Can use **any line in the set** to store a given cache block.



$E$ -way associative cache


# Accessing Set Associative Caches

- Set selection: same as direct-mapped cache
  - Use the set index bits to determine the set of interest.



# Accessing Set Associative Caches

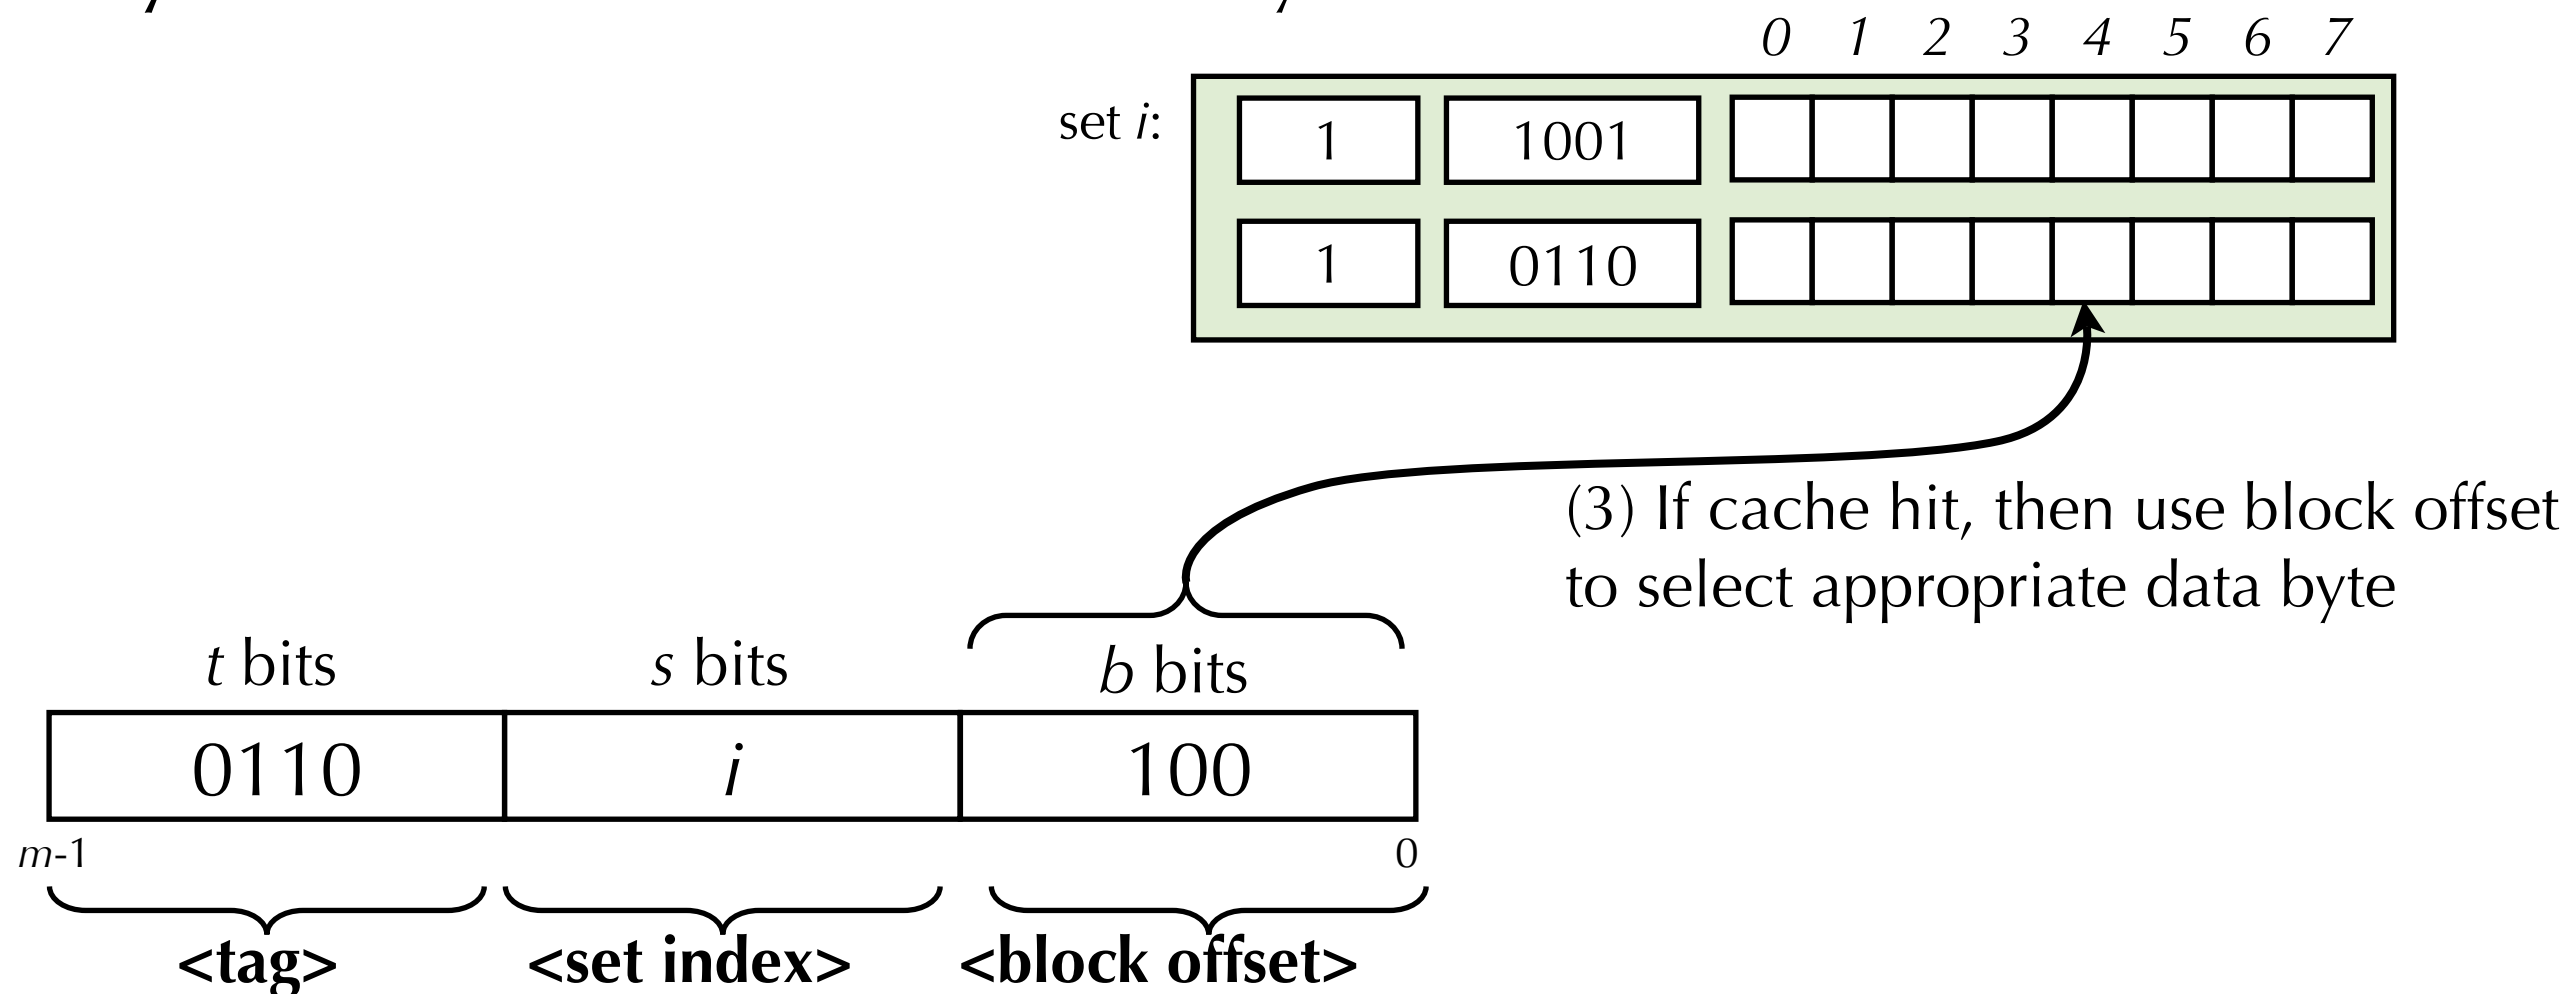
- Line matching and word selection: same as direct-mapped
  - **Line matching:** Find a *valid* line in the selected set with a *matching tag*. Check **all valid lines**.
  - **Byte selection:** Then extract the byte that we want





# Accessing Set Associative Caches

- Line matching and word selection: same as direct-mapped
  - **Line matching:** Find a *valid* line in the selected set with a *matching tag*. Check **all valid lines**.
  - **Byte selection:** Then extract the byte that we want






# 2-Way Associative Cache Simulation

- Simple example:
  - 4 bit addresses
  - 2 tag bits, 1 set index bit, 1 block offset bit

Address trace (reads):

$$A = 0 = \textcolor{red}{00}\textcolor{blue}{00}_2$$

**cache miss**



|       | valid | tag                                 | data       |
|-------|-------|-------------------------------------|------------|
| Set 0 | 1     | <span style="color: red;">00</span> | M[0], M[1] |
|       | 0     |                                     |            |
| Set 1 | 0     |                                     |            |
|       | 0     |                                     |            |

# 2-Way Associative Cache Simulation

- Simple example:
  - 4 bit addresses
  - 2 tag bits, 1 set index bit, 1 block offset bit


Address trace (reads):

$$A = 0 = \textcolor{red}{00}\textcolor{blue}{00}_2$$

**cache miss**

$$A = 1 = \textcolor{red}{00}\textcolor{blue}{01}_2$$

**cache hit**



|       | valid | tag       | data       |
|-------|-------|-----------|------------|
| Set 0 | 1     | <b>00</b> | M[0], M[1] |
|       | 0     |           |            |
| Set 1 | 0     |           |            |
|       | 0     |           |            |

# 2-Way Associative Cache Simulation

- Simple example:
  - 4 bit addresses
  - 2 tag bits, 1 set index bit, 1 block offset bit

Address trace (reads):

$$A = 0 = \textcolor{red}{00}\textcolor{blue}{00}_2$$


**cache miss**

$$A = 1 = \textcolor{red}{00}\textcolor{blue}{01}_2$$

**cache hit**

$$A = 7 = \textcolor{red}{01}\textcolor{blue}{11}_2$$

**cache miss**



|       | valid | tag       | data       |
|-------|-------|-----------|------------|
| Set 0 | 1     | <b>00</b> | M[0], M[1] |
|       | 0     |           |            |
| Set 1 | 1     | <b>01</b> | M[6], M[7] |
|       | 0     |           |            |

# 2-Way Associative Cache Simulation

- Simple example:
  - 4 bit addresses
  - 2 tag bits, 1 set index bit, 1 block offset bit

Address trace (reads):

$$A = 0 = 0000_2$$

cache miss

$$A = 1 = 0001_2$$


cache hit

$$A = 7 = 0111_2$$

cache miss

$$A = 8 = 1000_2$$

cache miss




|       | valid | tag | data       |
|-------|-------|-----|------------|
| Set 0 | 1     | 00  | M[0], M[1] |
|       | 1     | 10  | M[8], M[9] |
| Set 1 | 1     | 01  | M[6], M[7] |
|       | 0     |     |            |

# 2-Way Associative Cache Simulation

- Simple example:
  - 4 bit addresses
  - 2 tag bits, 1 set index bit, 1 block offset bit


Address trace (reads):

|                  |            |
|------------------|------------|
| $A = 0 = 0000_2$ | cache miss |
| $A = 1 = 0001_2$ | cache hit  |
| $A = 7 = 0111_2$ | cache miss |
| $A = 8 = 1000_2$ | cache miss |
| $A = 0 = 0000_2$ | cache hit  |




|       | valid | tag | data       |
|-------|-------|-----|------------|
| Set 0 | 1     | 00  | M[0], M[1] |
|       | 1     | 10  | M[8], M[9] |
| Set 1 | 1     | 01  | M[6], M[7] |
|       | 0     |     |            |

# Why use the middle bits as the set index?



- High-Order Bit Indexing
  - Adjacent memory lines would map to same set
  - Poor use of spatial locality
- Middle-Order Bit Indexing
  - Consecutive memory lines map to different cache lines
  - Can hold  $S \times B \times E$ -byte contiguous region of address space in cache at one time



# Maintaining a Set-Associative Cache

- When a set has multiple lines, which line do we use?
- Several different policies are used:
  - **Least Recently Used** (LRU) – must keep track of last access time for each line
  - **Not recently Used** (NRU) – Similar idea but not exact: Just track whether or not a cache line has been accessed during the last “N” accesses.
  - **Random** – just pick any one of the lines randomly



# Summary of cache concepts


- A **block** is a fixed-sized packet of information that moves back and forth between a cache and a lower level of memory hierarchy.
- A **line** is a container in a cache that stores a block, as well as other info such as valid bit and tag bits.
- A **set** is a collection of one or more lines.
  - Sets in direct-mapped caches have a single line
  - Sets in set-associative caches have multiple lines

# Topics for today

- The Principle of Locality
- Memory Hierarchies
- Caching Concepts
- Direct-Mapped Cache Organization
- Set-Associative Cache Organization
- Multi-level caches
- Cache writes

# Multi-Level Caches

- Options: separate **data** and **instruction caches**, or a **unified cache**



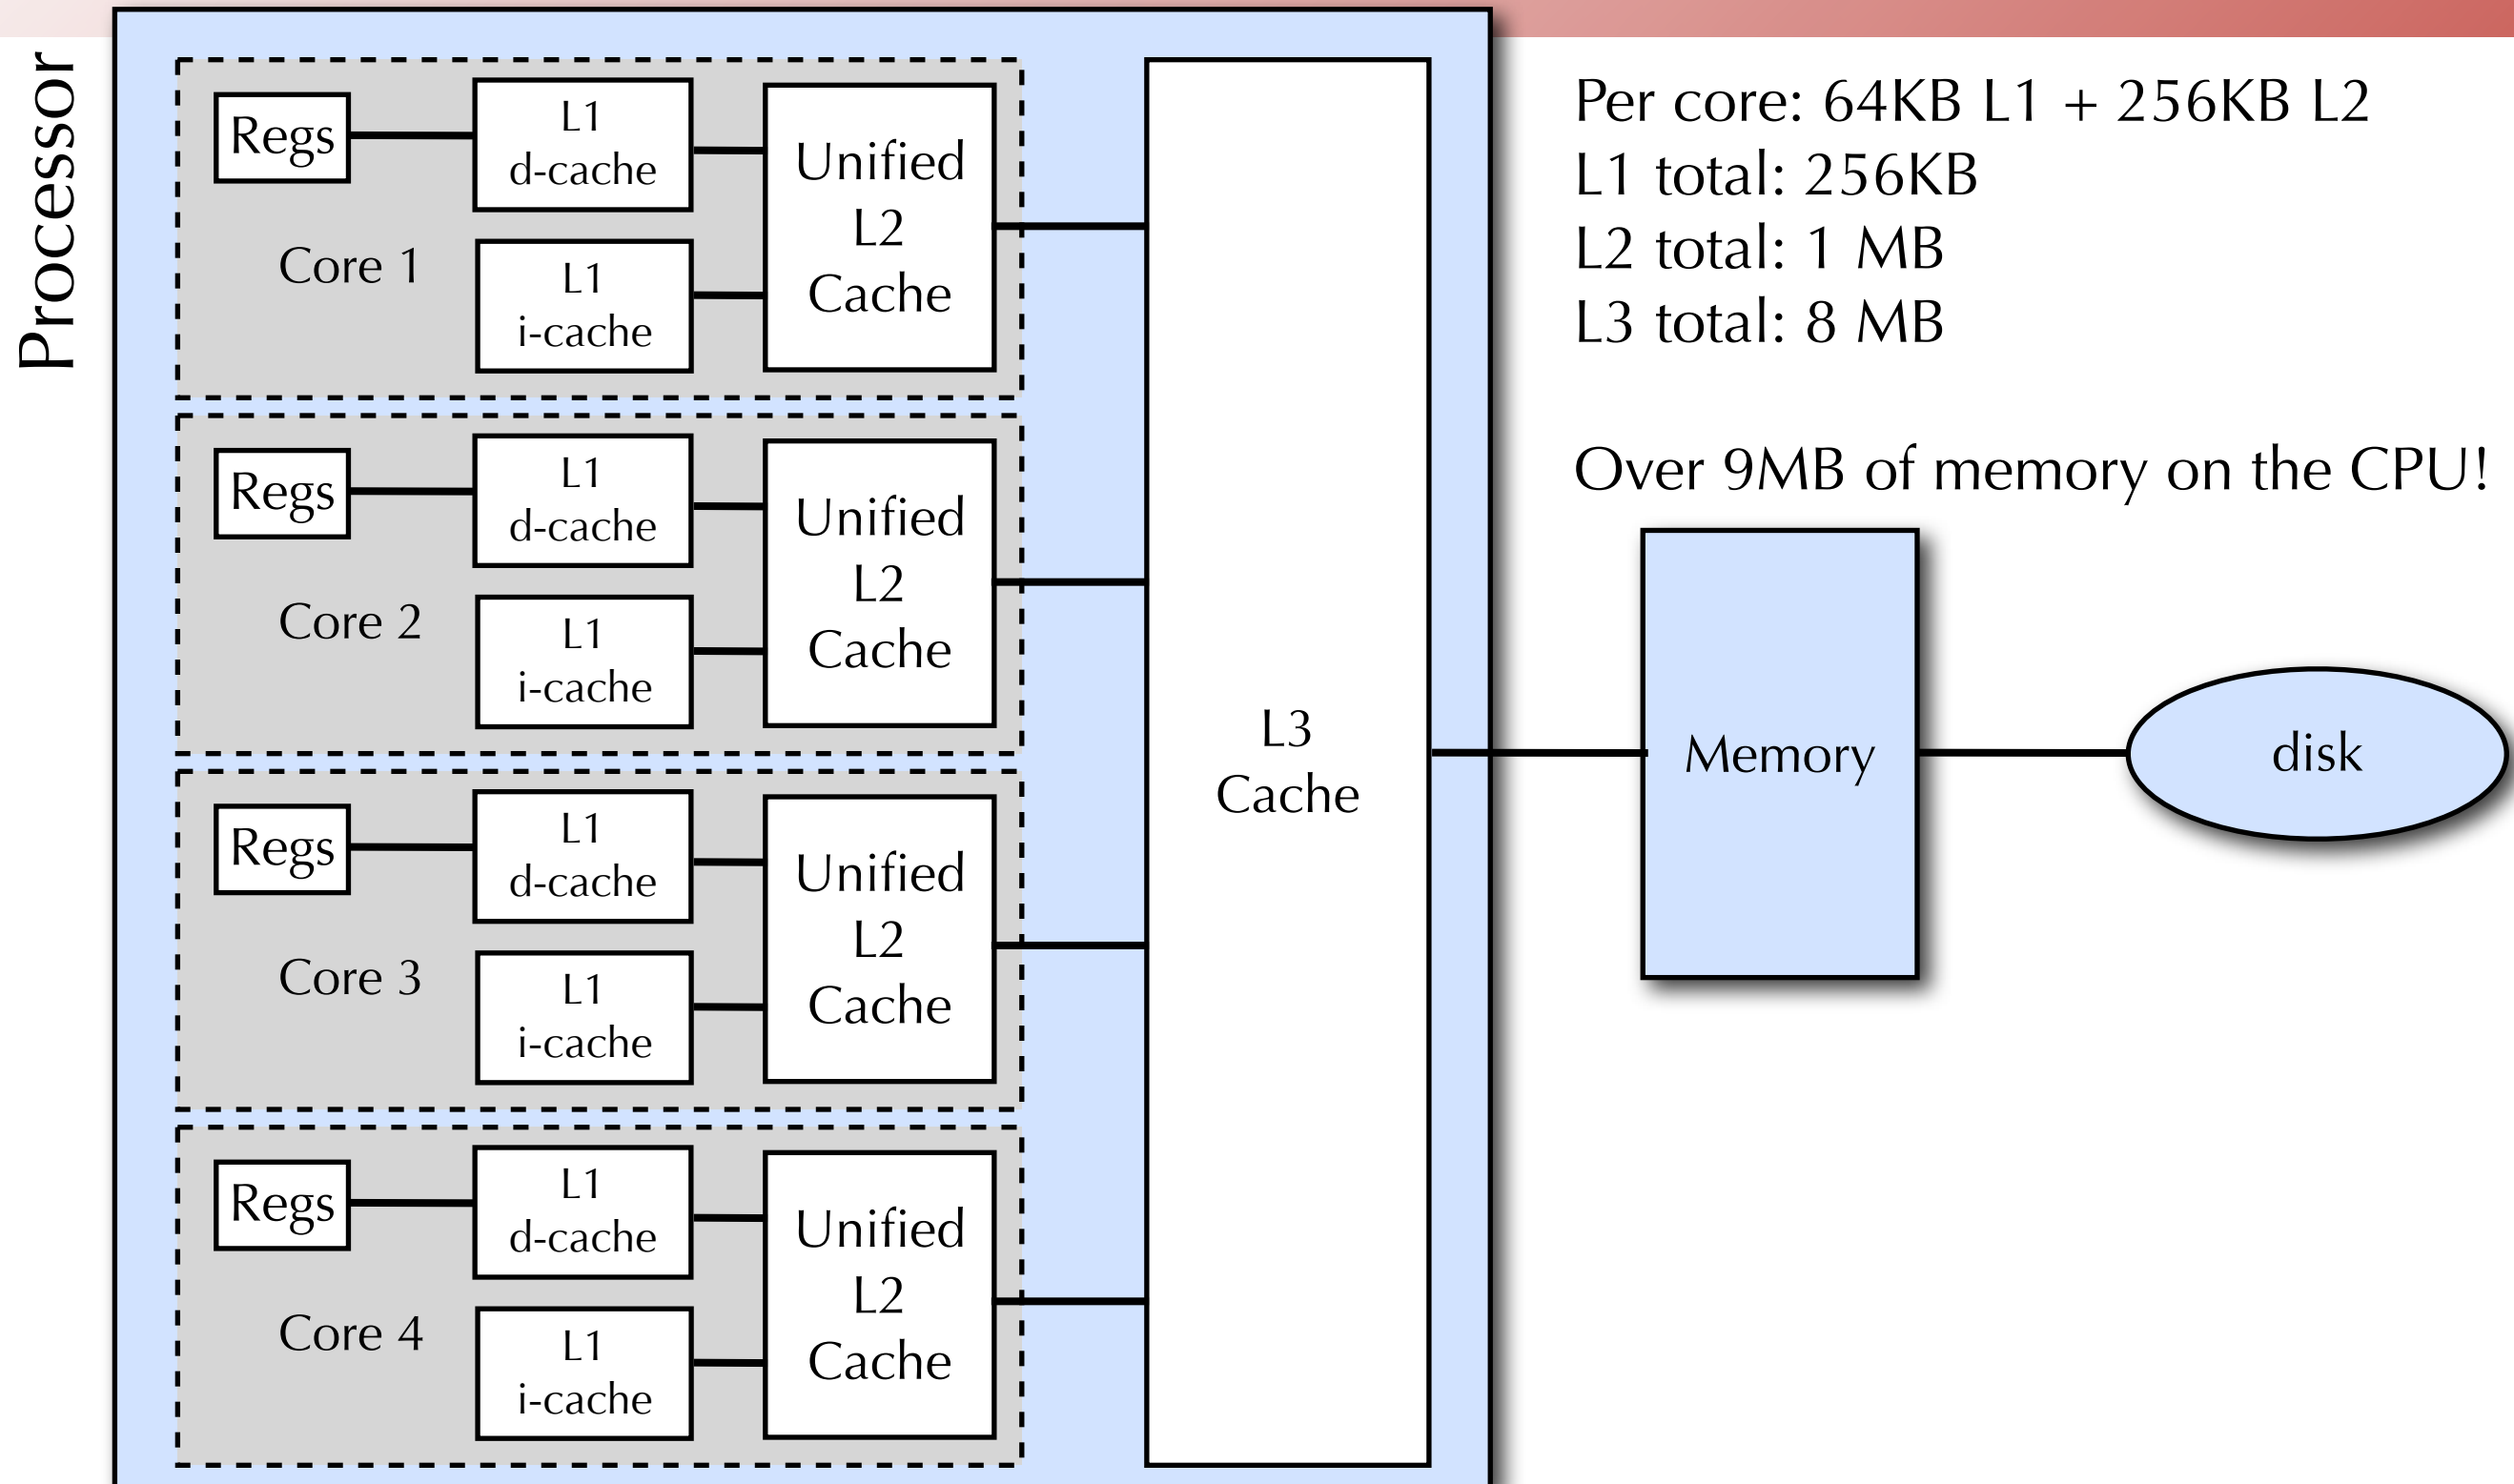
# Intel Core i7-965 processor specs

- 3.2 GHz clock speed, four cores
- L1 cache:
  - 64 KB (32 KB instruction, 32KB data), per core; 8-way set associative, 64 sets
  - 4 CPU cycles to access (~ 1 ns)
- L2 cache
  - 256KB, per core; 8-way set associative, 512 sets
  - 10 CPU cycles to access (~3 ns)
- L3 cache:
  - 8 MB, is shared across all cores; 8-way set associative, 8192 sets
  - 35-40 CPU cycles to access (~12 ns)
- Main memory:
  - Up to 1 TB
  - About 60 ns to access – 60× slower than L1 cache, 20× slower than L2, 5× slower than L3



<http://techreport.com/articles.x/15818>  
<http://ark.intel.com/Product.aspx?id=37149&processor=i7-965&spec-codes=SLBCJ>  
<http://software.intel.com/en-us/articles/who-moved-the-goal-posts-the-rapidly-changing-world-of-cpus/>

# Intel Core i7-965 processor specs




# Topics for today

- The Principle of Locality
- Memory Hierarchies
- Caching Concepts
- Direct-Mapped Cache Organization
- Set-Associative Cache Organization
- Multi-level caches
- Cache writes



# Handling cache writes


- On **reads**, the memory system fetches data from lower levels and moves it into higher levels.
  - For example, moving data from L2 cache to L1 cache, or from DRAM to L2 cache.
- On a **write**, where do we put the modified data?






# Handling cache writes

- Approach #1: **Write through** policy
  - Data written through to lower levels of cache as soon as it is modified.
  - Simple, but causes an (expensive) memory write for each cache write.



# Handling cache writes

- Approach #2: **Write back** policy
  - Only write data to highest-level cache.
  - When cache line is **evicted**, write back to next lower level cache.
  - Caches must keep track of whether each line is **dirty** (needs writing back)



# Next lecture

- Measuring cache size and performance.



- Optimizing programs to exploit cache locality.