



HARVARD

School of Engineering
and Applied Sciences

Cache performance measurement and optimization

CS61, Lecture 13

Prof. Stephen Chong

October 14, 2010

Announcements

- Lab 3
 - Design checkpoint deadline tonight
 - Final submission deadline Tuesday 19th

Topics for today

- Cache performance metrics
- Discovering your cache's size and performance
- The “Memory Mountain”
- Matrix multiply, six ways
- Blocked matrix multiplication
- Exploiting locality in your programs

Cache Performance Metrics

- Miss Rate

- Fraction of memory references not found in cache ($\# \text{ misses} / \# \text{ references}$)
- Typical numbers:
 - 3-10% for L1
 - Can be quite small (e.g., $< 1\%$) for L2, depending on size and locality.

- Hit Time

- Time to deliver a line in the cache to the processor (includes time to determine whether the line is in the cache)
- Typical numbers:
 - 1-2 clock cycles for L1
 - 5-20 clock cycles for L2

- Miss Penalty

- Additional time required because of a miss
 - Typically 50-200 cycles for main memory

Wait, what do those numbers mean?

- Huge difference between a hit and a miss
 - Could be 100x, if just L1 and main memory
- Would you believe 99% hits is twice as good as 97%?
 - Consider:
cache hit time of 1 cycle
miss penalty of 100 cycles
- Average access time:
 - 97% hits: $1 \text{ cycle} + 0.03 * 100 \text{ cycles} = 4 \text{ cycles}$
 - 99% hits: $1 \text{ cycle} + 0.01 * 100 \text{ cycles} = 2 \text{ cycles}$
- This is why “miss rate” is used instead of “hit rate”

Writing Cache Friendly Code

- Repeated references to variables are good (**temporal locality**)
- Stride-1 reference patterns are good (**spatial locality**)
- Examples:
 - cold cache, 4-byte words, 4-word cache blocks

```
int sum_array_rows(int a[M][N]) {  
    int i, j, sum = 0;  
  
    for (i = 0; i < M; i++)  
        for (j = 0; j < N; j++)  
            sum += a[i][j];  
    return sum;  
}
```

Miss rate = $1/4 = 25\%$


```
int sum_array_cols(int a[M][N]) {  
    int i, j, sum = 0;  
  
    for (j = 0; j < N; j++)  
        for (i = 0; i < M; i++)  
            sum += a[i][j];  
    return sum;  
}
```

Miss rate = 100%


Determining cache characteristics

- Say you have a machine but don't know its cache size or speeds.
- How would you figure these values out?
- Idea: Write a program to measure the cache's behavior and performance.
 - Program needs to perform memory accesses with different locality patterns.
- Simple approach:
 - Allocate array of size W words
 - Loop over the array with stride index S and measure speed of memory accesses
 - Vary W and S to estimate cache characteristics

$S = 4$ words



Determining cache characteristics



- What happens as you vary W and S ?
- Changing W varies total amount of memory accessed by program
 - As W gets larger than one cache level, performance of program will drop
- Changing S varies the spatial locality of each access.
 - If S is less than the size of a cache line, sequential accesses will be fast.
 - If S is greater than the size of a cache line, sequential accesses will be slower.
- See end of lecture notes for example C program to do this.

The Memory Mountain

Intel Core i7

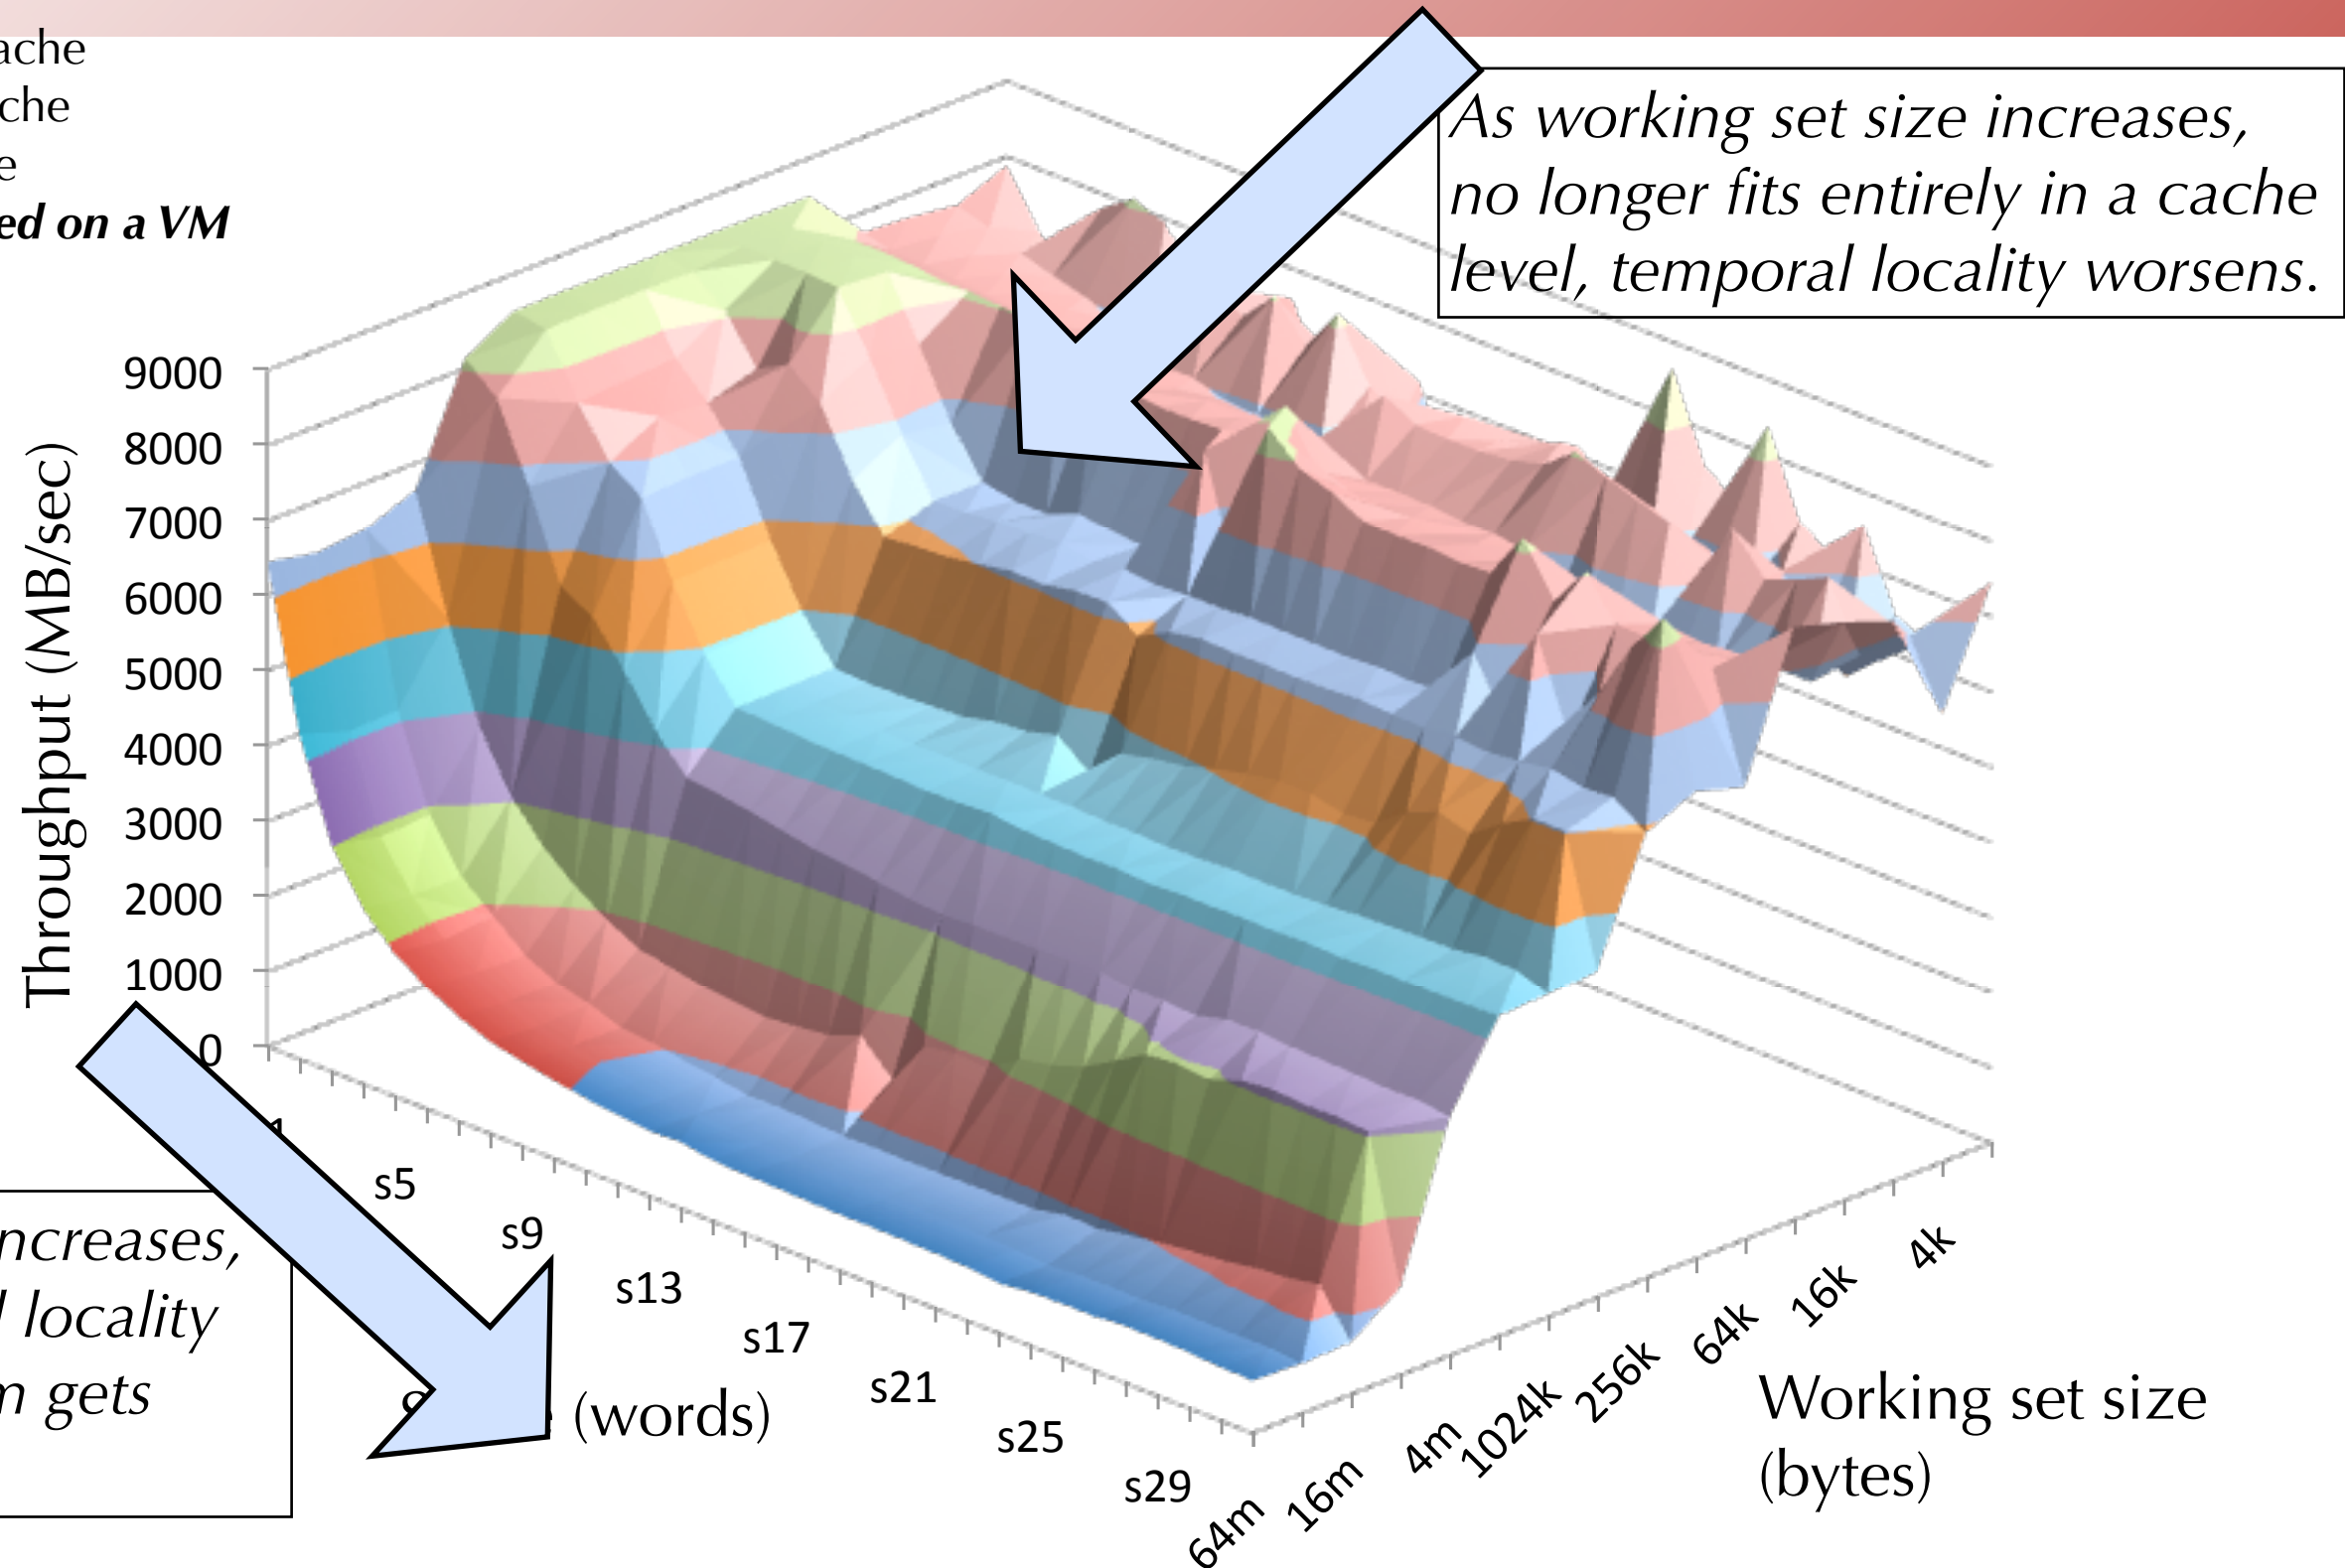
2.7 GHz

32 KB L1 d-cache

256 KB L2 cache


8MB L3 cache

CAVEAT: Tested on a VM




Varying Working Set

- Keep stride constant at $S = 16$ words, and vary W from 1KB to 64MB
 - Shows size and read throughputs of different cache levels and memory




Varying stride

- Keep working set constant at $W = 256$ KB, vary stride from 1-32 words




Core i7




Pentium Xeon

Pentium Nocona Xeon x86-64
3.2 GHz
12 Kuop on-chip L1 trace cache
16 KB on-chip L1 d-cache
1 MB off-chip unified L2 cache



Opteron Memory Mountain




Topics for today

- Cache performance metrics
- Discovering your cache's size and performance
- The “Memory Mountain”
- Matrix multiply, six ways
- Blocked matrix multiplication
- Exploiting locality in your programs

Matrix Multiplication Example

- Matrix multiplication is heavily used in numeric and scientific applications.
 - It's also a nice example of a program that is highly sensitive to cache effects.
- Multiply two $N \times N$ matrices
 - $O(N^3)$ total operations
 - Read N values for each source element
 - Sum up N values for each destination

```
void mmm(double *a, double *b, double *c, int n) {  
    int i, j, k;  
    /* ijk */  
    for (i=0; i<n; i++) {  
        for (j=0; j<n; j++) {  
            sum = 0.0;  
            for (k=0; k<n; k++)  
                sum += a[i][k] * b[k][j];  
            c[i][j] = sum;  
        }  
    }  
}
```



Variable sum
held in register

Matrix Multiplication Example

```
/* ijk */  
for (i=0; i<n; i++) {  
    for (j=0; j<n; j++) {  
        sum = 0.0;  
        for (k=0; k<n; k++)  
            sum += a[i][k] * b[k][j];  
        c[i][j] = sum;  
    }  
}
```

$$4 \times 3 + 2 \times 2 + 7 \times 5 = 51$$

4	2	7
1	8	2
6	0	1

×


3	0	1
2	4	5
5	9	1

=

51		

Miss Rate Analysis for Matrix Multiply

- Assume:
 - Line size = 32B (big enough for four 64-bit “double” values)
 - Matrix dimension N is very large
 - Cache is not big enough to hold multiple rows
- Analysis Method:
 - Look at access pattern of inner loop



Layout of C Arrays in Memory (review)

- C arrays allocated in **row-major** order
 - Each row in contiguous memory locations
- Stepping through columns in one row:
 - `for (i = 0; i < N; i++)`
 `sum += a[0][i];`
 - Accesses successive elements
 - Compulsory miss rate: (8 bytes per double) / (block size of cache)
- Stepping through rows in one column:
 - `for (i = 0; i < n; i++)`
 `sum += a[i][0];`
 - Accesses distant elements — no spatial locality!
 - Compulsory miss rate = 100%

Matrix Multiplication (ijk)

```
/* ijk */
for (i=0; i<n; i++) {
  for (j=0; j<n; j++) {
    sum = 0.0;
    for (k=0; k<n; k++)
      sum += a[i][k] * b[k][j];
    c[i][j] = sum;
  }
}
```

Inner loop:


$(i,*)$



A

↑
Row-wise


$(*,j)$



B

↑
Column-wise

(i,j)



C

↑
Fixed

- 2 loads, 0 stores per iteration
- Assume cache line size of 32 bytes, so 4 doubles per line
- Misses per iteration:

$A = 0.25$

$B = 1$

$C = 0$

Total: 1.25

Cache miss analysis

```
/* ijk */
for (i=0; i<n; i++) {
  for (j=0; j<n; j++) {
    sum = 0.0;
    for (k=0; k<n; k++)
      sum += a[i][k] * b[k][j];
    c[i][j] = sum;
  }
}
```

Inner loop:


$(i,*)$



A

Row-wise


$(*,j)$



B

Column-wise


(i,j)



C

Fixed

First iteration:



A





B



C

After first iteration in cache (schematic):



4 doubles wide

Cache miss analysis

```
/* ijk */
for (i=0; i<n; i++) {
  for (j=0; j<n; j++) {
    sum = 0.0;
    for (k=0; k<n; k++)
      sum += a[i][k] * b[k][j];
    c[i][j] = sum;
  }
}
```

Inner loop:


$(i,*)$



A

Row-wise


$(*,j)$



B

Column-wise


(i,j)




C

Fixed

First iteration:



A






B



C

After first iteration in cache (schematic):



4 doubles wide

Matrix Multiplication (jik)

```
/* jik */
for (j=0; j<n; j++) {
  for (i=0; i<n; i++) {
    sum = 0.0;
    for (k=0; k<n; k++)
      sum += a[i][k] * b[k][j];
    c[i][j] = sum;
  }
}
```

Inner loop:


(i,*)



A

↑
Row-wise


(*,j)



B

↑
Column-
wise

(i,j)



C

↑
Fixed

- Same as ijk, just swapped order of outer loops
- 2 loads, 0 stores per iteration
- Assume cache line size of 32 bytes, so 4 doubles per line
- Misses per iteration:

A = 0.25

B = 1


C = 0

Total: 1.25

Matrix Multiplication (kij)

```
/* kij */
for (k=0; k<n; k++) {
  for (i=0; i<n; i++) {
    r = a[i][k];
    for (j=0; j<n; j++)
      c[i][j] += r * b[k][j];
  }
}
```

Inner loop:



- 2 load, 1 store per iteration
- Assume cache line size of 32 bytes, so 4 doubles per line
- Misses per iteration:

$A = 0$

$B = 0.25$

$C = 0.25$


Total: 0.5

Matrix Multiplication (ikj)

```
/* ikj */
for (i=0; i<n; i++) {
  for (k=0; k<n; k++) {
    r = a[i][k];
    for (j=0; j<n; j++)
      c[i][j] += r * b[k][j];
  }
}
```

Inner loop:

(i,k)



A

↑
Fixed


(k,*)



B

↑
Row-wise

(i,*)



C

↑
Row-wise

- Same as kij, just swapped order of outer loops
- 2 load, 1 store per iteration
- Assume cache line size of 32 bytes, so 4 doubles per line
- Misses per iteration:

A = 0

B = 0.25


C = 0.25

Total: 0.5

Matrix Multiplication (jki)

```
/* jki */
for (j=0; j<n; j++) {
  for (k=0; k<n; k++) {
    r = b[k][j];
    for (i=0; i<n; i++)
      c[i][j] += a[i][k] * r;
  }
}
```

Inner loop:



- 2 load, 1 store per iteration
- Assume cache line size of 32 bytes, so 4 doubles per line
- Misses per iteration:

A = 1

B = 0

C = 1


Total: 2

Matrix Multiplication (kji)

```
/* kji */
for (k=0; k<n; k++) {
  for (j=0; j<n; j++) {
    r = b[k][j];
    for (i=0; i<n; i++)
      c[i][j] += a[i][k] * r;
  }
}
```

Inner loop:


(*,k)



A

Column-
wise


(k,j)



B

Fixed

(*,j)



C

Column-
wise

- Same as kji, just swapped order of outer loops
- 2 load, 1 store per iteration
- Assume cache line size of 32 bytes, so 4 doubles per line
- Misses per iteration:

A = 1

B = 0

C = 1

Total: 2

Summary of Matrix Multiplication

```
for (i=0; i<n; i++) {  
  for (j=0; j<n; j++) {  
    sum = 0.0;  
    for (k=0; k<n; k++)  
      sum += a[i][k] * b[k][j];  
    c[i][j] = sum;  
  }  
}
```

ijk or jik:

2 loads, 0 stores

misses/iter = 1.25

```
for (k=0; k<n; k++) {  
  for (i=0; i<n; i++) {  
    r = a[i][k];  
    for (j=0; j<n; j++)  
      c[i][j] += r * b[k][j];  
  }  
}
```

kij or ikj:

2 loads, 1 store

misses/iter = 0.5


```
for (j=0; j<n; j++) {  
  for (k=0; k<n; k++) {  
    r = b[k][j];  
    for (i=0; i<n; i++)  
      c[i][j] += a[i][k] * r;  
  }  
}
```

jki or kji:

2 loads, 1 store


misses/iter = 2.0

Matrix Multiply Performance



- Each implementation doing same number of arithmetic operations, but ~20× difference!
- Pairs with same number of mem. references and misses per iteration almost identical

Matrix Multiply Performance



- Miss rate better predictor of performance than number of mem. accesses!
- For large N, kij and ikj performance almost constant.
Due to **hardware prefetching**, able to recognize stride-1 patterns.

Topics for today

- Cache performance metrics
- Discovering your cache's size and performance
- The “Memory Mountain”
- Matrix multiply, six ways
- Blocked matrix multiplication
- Exploiting locality in your programs

Using blocking to improve locality

- Blocked matrix multiplication
 - Break matrix into smaller blocks and perform independent multiplications on each block.
 - Improves locality by operating on one block at a time.
 - Best if each block can fit in the cache!
- Example: Break each matrix into four sub-blocks

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \times \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix} = \begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix}$$

Key idea: Sub-blocks (i.e., A_{xy}) can be treated just like scalars.

$$C_{11} = A_{11}B_{11} + A_{12}B_{21} \qquad C_{12} = A_{11}B_{12} + A_{12}B_{22}$$

$$C_{21} = A_{21}B_{11} + A_{22}B_{21} \qquad C_{22} = A_{21}B_{12} + A_{22}B_{22}$$

Blocked Matrix Multiply


```
void bmmm(int n, double a[n][n], double b[n][n], double c[n][n]) {  
    int i, j, k;  
    for (i = 0; i < n; i+=B)  
        for (j = 0; j < n; j+=B)  
            for (k = 0; k < n; k+=B)  
                /* B x B mini matrix multiplications  
                for (i1 = i; i1 < i+B; i++)  
                    for (j1 = j; j1 < j+B; j++)  
                        for (k1 = k; k1 < k+B; k++)  
                            c[i1][j1] += a[i1][k1] * b[k1][j1];  
}
```

Code becomes harder to read!
Is it worth it?
Tradeoff between performance
and maintainability...


- Partition arrays into $b\text{size} \times b\text{size}$ chunks
- Innermost (i1, j1, k1) loop pair multiplies an A chunk by a B chunk and accumulates result in a C chunk

Blocked matrix multiply


- Assume 3 chunks can fit into the cache, i.e., $3b\text{size}^2 < C$
- First block iteration



- After first iteration in cache (schematic)




Cache miss analysis



- Assume 3 chunks can fit into the cache
- Assume $b\text{size}$ is a multiple of 4
- $b\text{size}^2/4$ misses per chunk, so $3/4 * b\text{size}^2$ misses per chunk iteration
- $(n/b\text{size})^3$ chunk iterations
- Total of $(n/b\text{size})^3 * 3/4 * b\text{size}^2$ misses = $n^3 * 3/(4 * b\text{size})$
- Compare with $n^3 * 1/2$ total misses for kij algorithm


Blocked Matrix Multiply Performance



- Blocking (bijk and bikj) improves performance by a factor of two over unblocked versions (ijk and jik)
 - Relatively insensitive to array size.

Blocked Matrix Multiply Performance


Intel Core i7
2.7 GHz
32 KB L1 d-cache
256 KB L2 cache
8MB L3 cache
CAVEAT: Tested on a VM



Blocked Matrix Multiply Performance

Intel Core i7
2.7 GHz
32 KB L1 d-cache
256 KB L2 cache
8MB L3 cache

CAVEAT: Tested on a VM



Exploiting locality in your programs

- Focus attention on inner loops
 - This is where most computation and memory accesses in your program occurs
- Try to maximize spatial locality
 - Read data objects sequentially, with stride 1, in the order they are stored in memory
- Try to maximize temporal locality
 - Use a data object as often as possible once it has been read from memory

Next lecture

- Virtual memory
 - Using memory as a cache for disk

Cache performance test program

```
/* The test function */
void test(int elems, int stride) {
    int i, result = 0;
    volatile int sink;

    for (i = 0; i < elems; i += stride)
        result += data[i];
    sink = result; /* So compiler doesn't optimize away the loop */
}

/* Run test(elems, stride) and return read throughput (MB/s) */
double run(int size, int stride)
{
    uint64_t start_cycles, end_cycles, diff;
    int elems = size / sizeof(int);

    test(elems, stride); /* warm up the cache */
    start_cycles = get_cpu_cycle_counter(); /* Read CPU cycle counter */
    test(elems, stride); /* Run test */
    end_cycles = get_cpu_cycle_counter(); /* Read CPU cycle counter again */
    diff = end_cycles - start_cycles; /* Compute time */
    return (size / stride) / (diff / CPU_MHZ); /* convert cycles to MB/s */
}
```

Cache performance main routine

```
#define CPU_MHZ 2.8 * 1024.0 * 1024.0; /* e.g., 2.8 GHz */
#define MINBYTES (1 << 10) /* Working set size ranges from 1 KB */
#define MAXBYTES (1 << 23) /* ... up to 8 MB */
#define MAXSTRIDE 16 /* Strides range from 1 to 16 */
#define MAXELEMS MAXBYTES/sizeof(int)

int data[MAXELEMS]; /* The array we'll be traversing */

int main()
{
    int size; /* Working set size (in bytes) */
    int stride; /* Stride (in array elements) */

    init_data(data, MAXELEMS); /* Initialize each element in data to 1 */
    for (size = MAXBYTES; size >= MINBYTES; size >>= 1) {
        for (stride = 1; stride <= MAXSTRIDE; stride++)
            printf("%.1f\t", run(size, stride));
        printf("\n");
    }
    exit(0);
}
```