



**HARVARD**

**School of Engineering  
and Applied Sciences**

# Virtual memory

*CS61, Lecture 14*

*Prof. Stephen Chong*

*October 19, 2010*

# Announcements

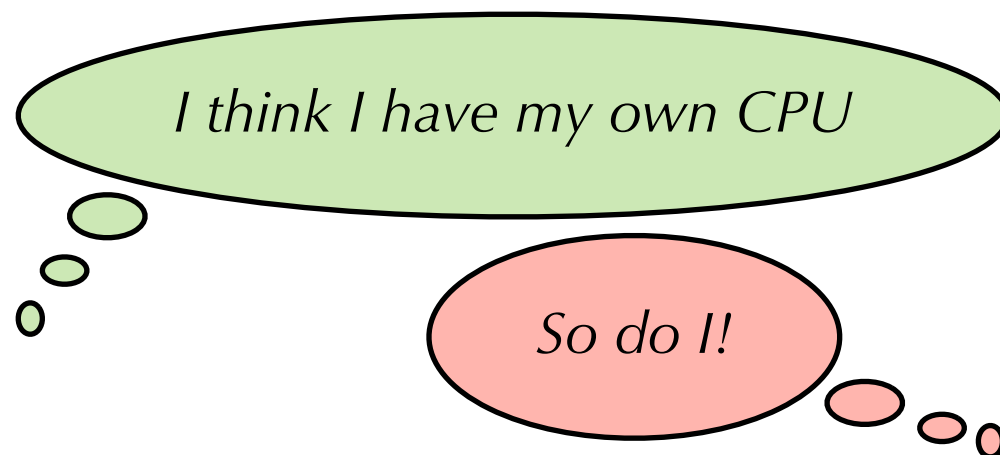
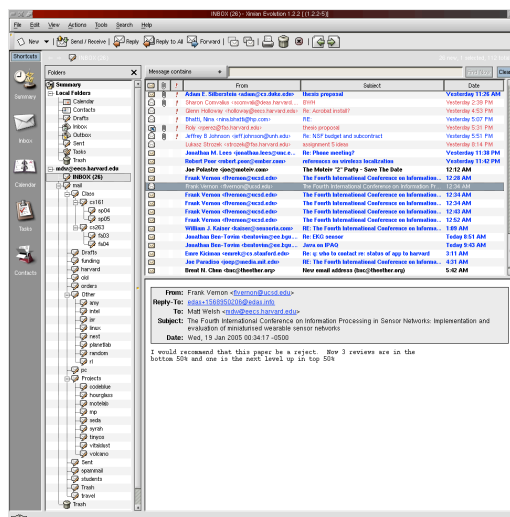
- Lab 3 due tonight!
- In-class **midterm exam**, Thursday 28 Oct
  - Open book, closed note
    - *Note: we will provide you with useful reference material, such as the stack frame cheat sheet, description of assembly commands, etc.*
  - No computers, no smartphones, no internet access
- Midterm review session, Monday 25 Oct
  - 7pm to 8:30pm, location TBA
  - Optional

# Today

- Running multiple programs at once
- Virtual memory
  - Address spaces
  - Benefits
  - Physical memory as a cache for virtual memory
  - Page tables
  - Page hits and page faults
  - Virtual memory as a tool for memory management
  - Virtual memory as a tool for memory protection
  - Address translation
  - Table Lookaside Buffer (TLB)
  - Example

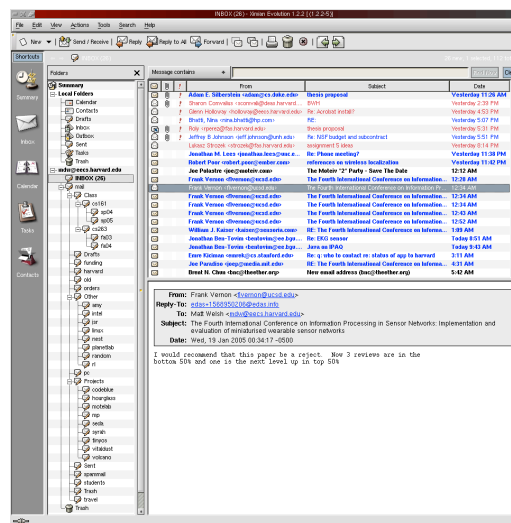
# Running multiple programs at once

- So far, we have discussed memory management for a single program running on a computer.
- Most modern computers run multiple programs simultaneously.
  - This is called **multiprogramming** or **multitasking**.



# Timeslicing

- The OS **timeslices** multiple applications on a single CPU
  - Switches between applications extremely rapidly, i.e., 100 times/sec
  - Each switch between two applications is called a context switch.



Operating System



time

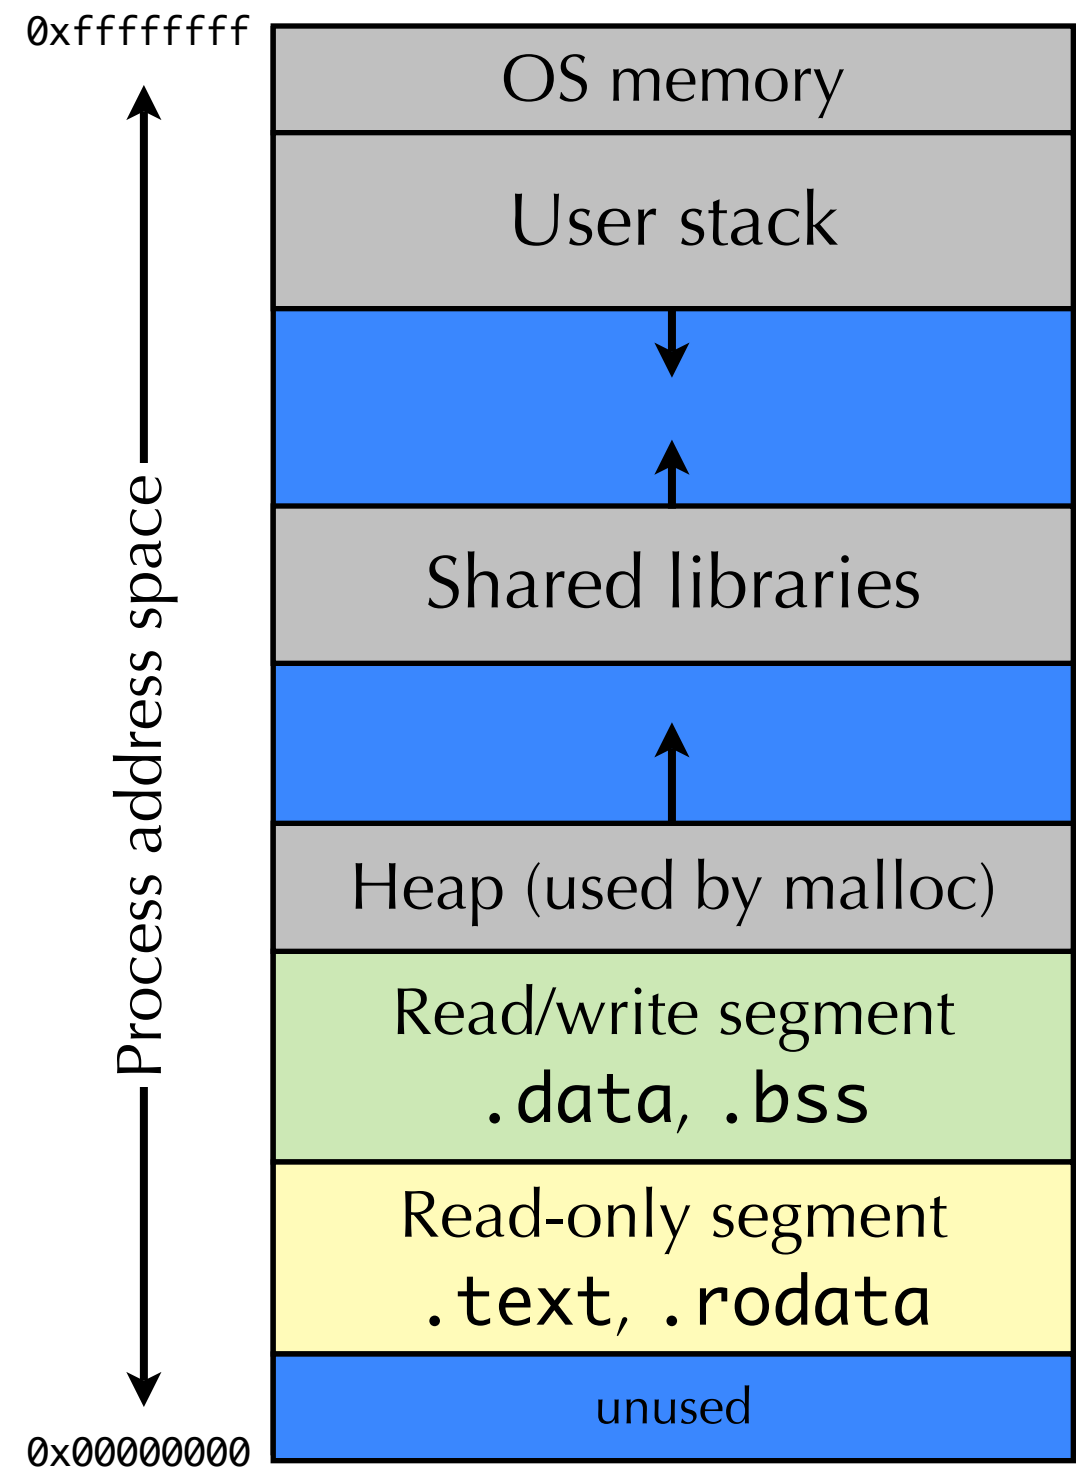
Timeslice on  
single CPU system






# Virtual memory

- OS also gives every program illusion of having it's own memory!
- Each program running on the machine is called a **process**.
  - Each process has a **process ID** (a number) and an **owner** – the user ID running the process.
  - UNIX command “ps” prints out info about the running processes.
- Each process has its own **address space**
  - Contents of that process' memory, visible only to that process.
  - 4GB on 32 bit machine, ~16 TB on 64 bit machine




# Problem 1: How Does Everything Fit?

4GB to 16 EB virtual memory




Physical main memory:  
Few Gigabytes



# Problem 2: Memory Management

Physical main memory

Process 1  
Process 2  
Process 3  
...  
Process n




*What goes  
where?*






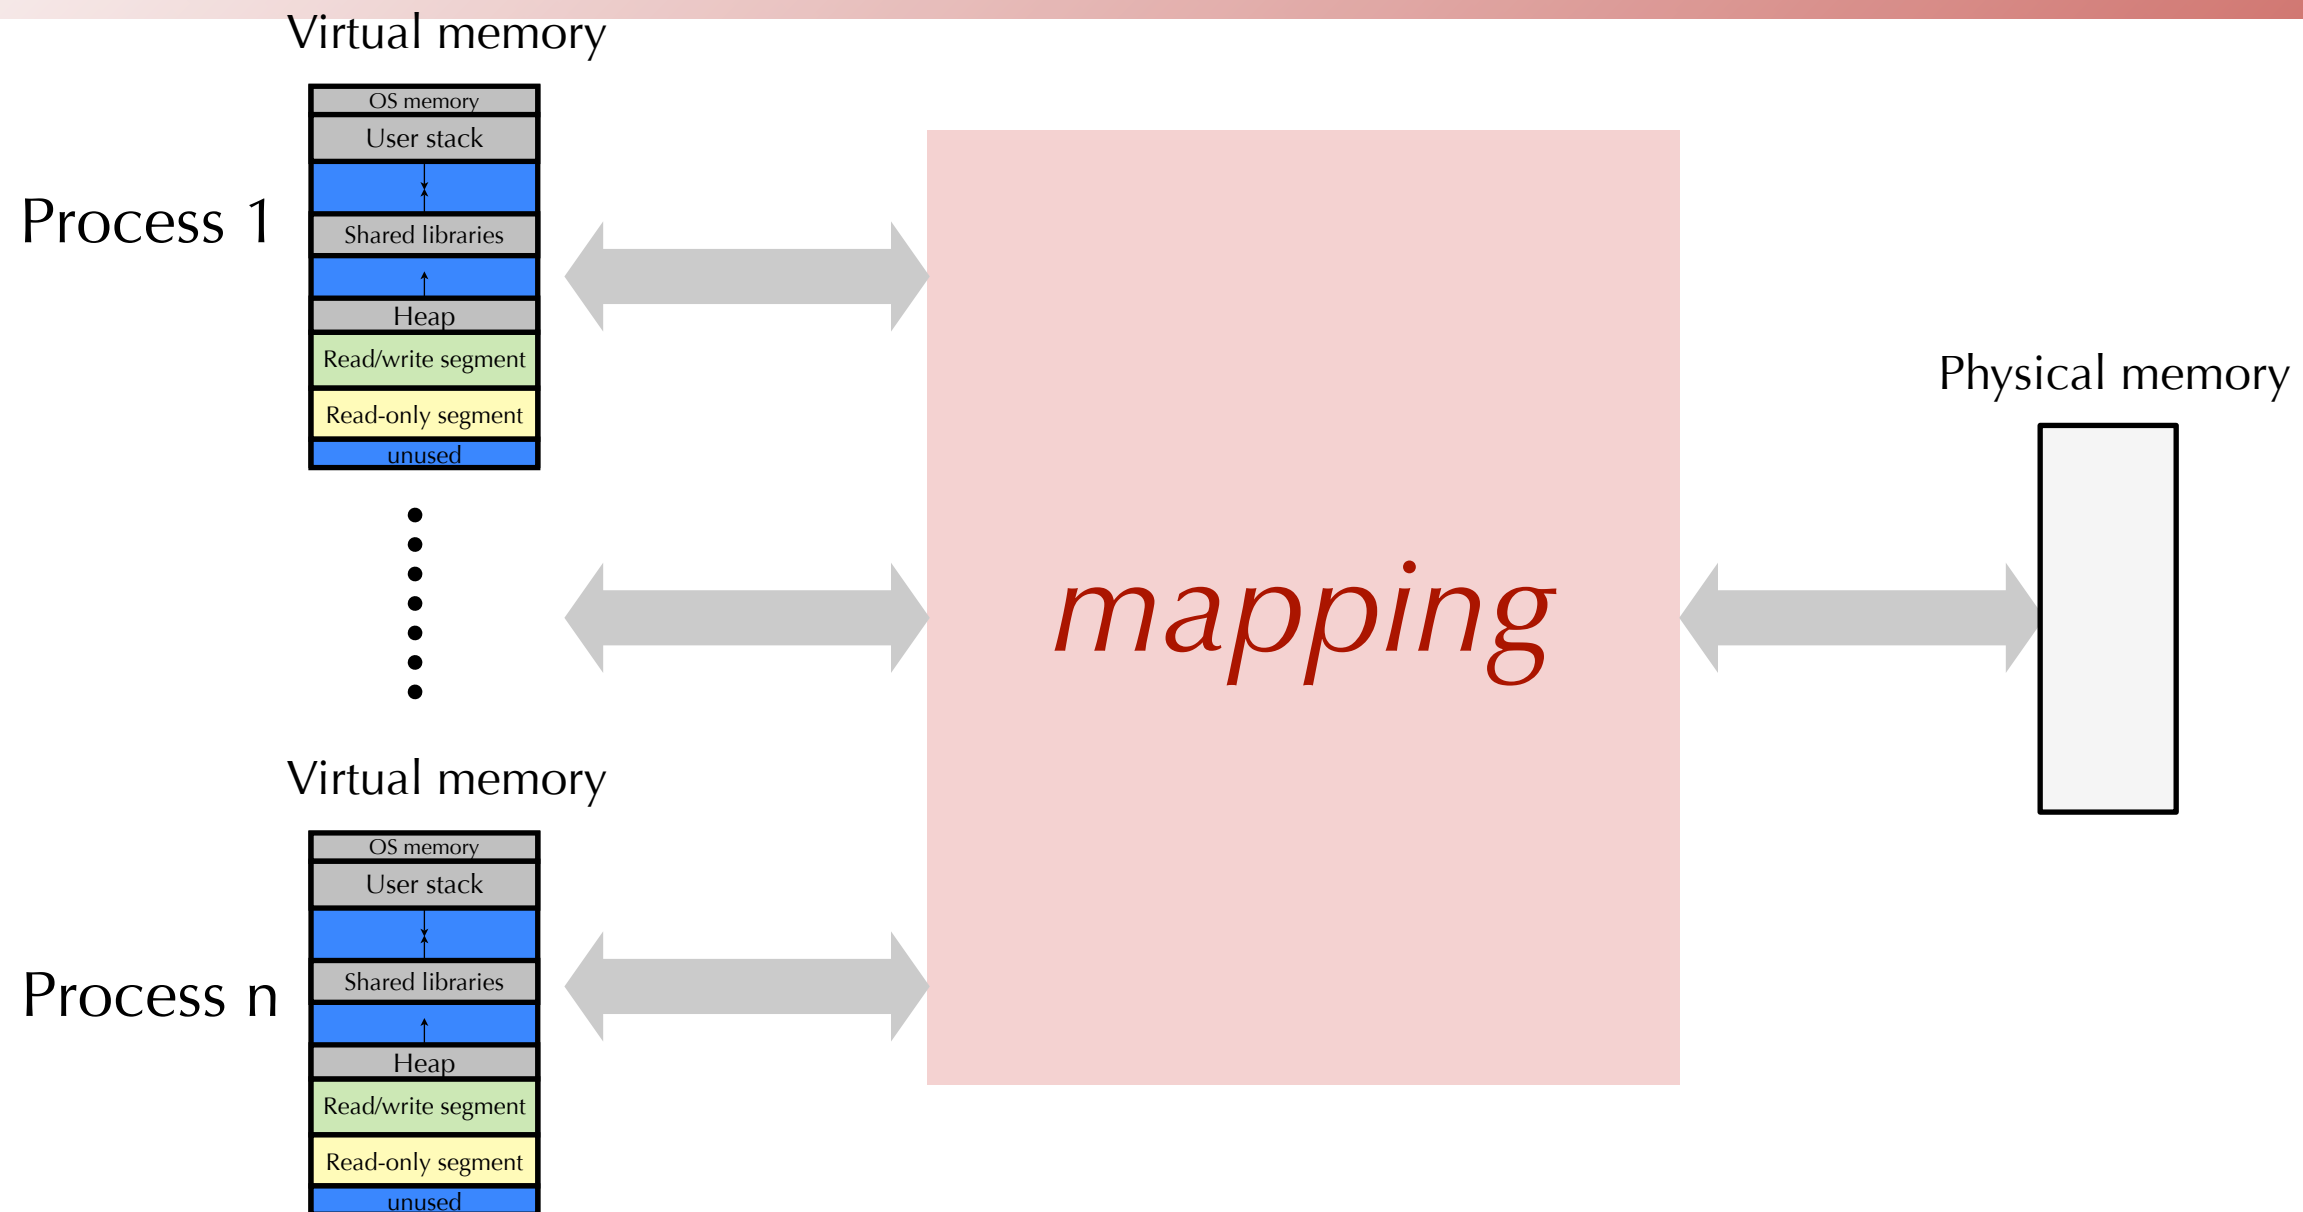
# Problem 3: How To Protect



# Problem 4: How To Share



# Solution: Level Of Indirection

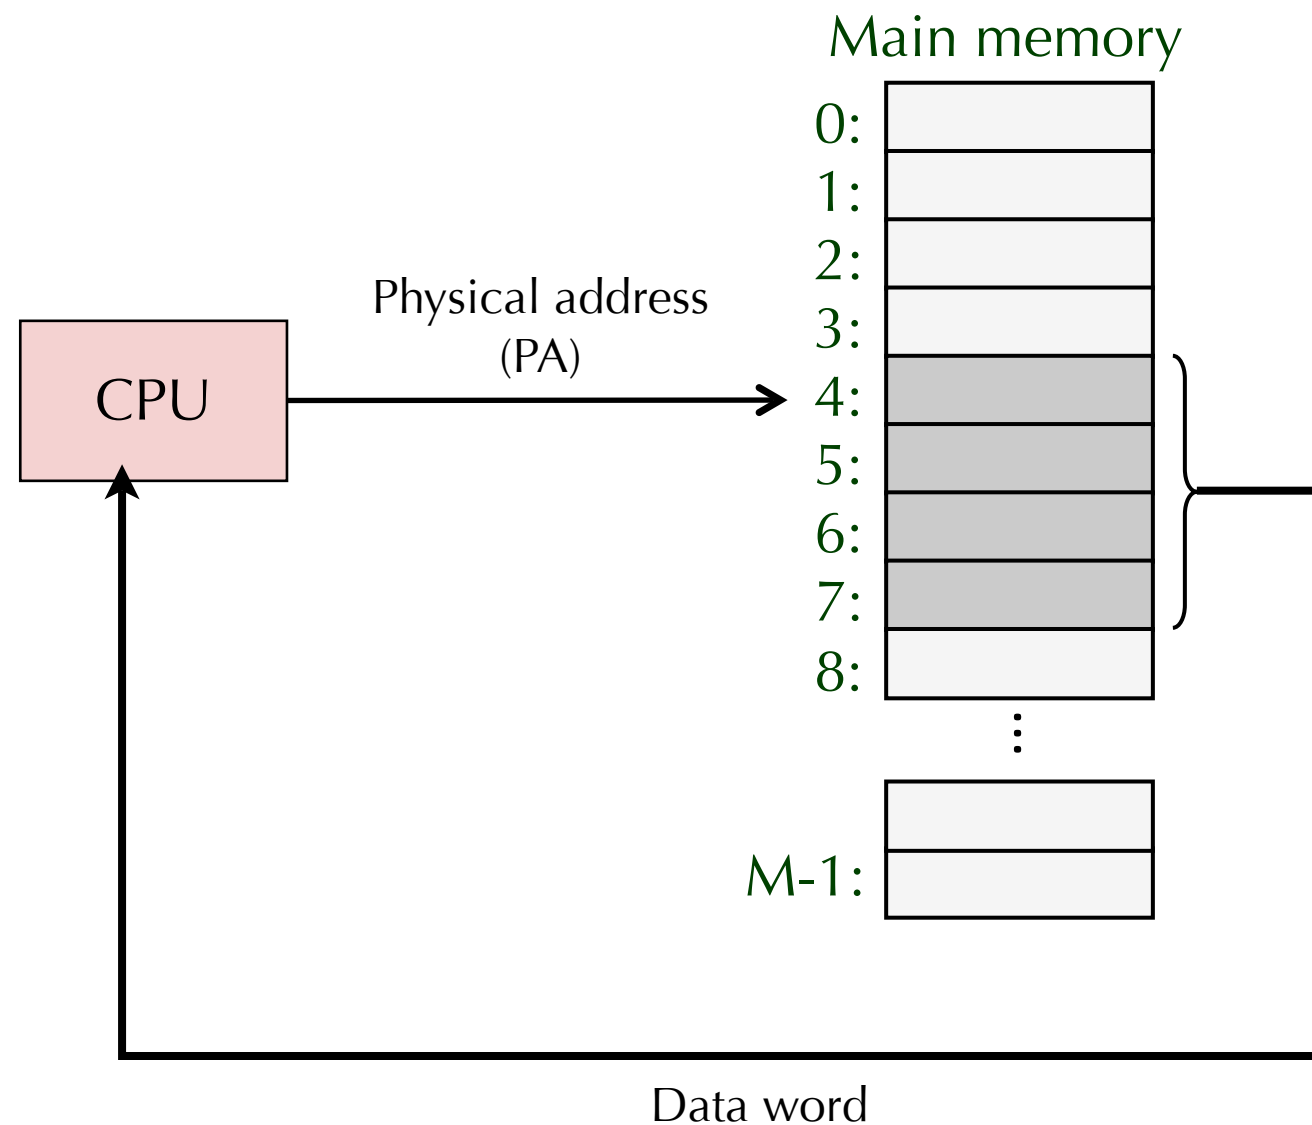


- Each process gets its own private memory space
- Solves the previous problems

# Address Spaces

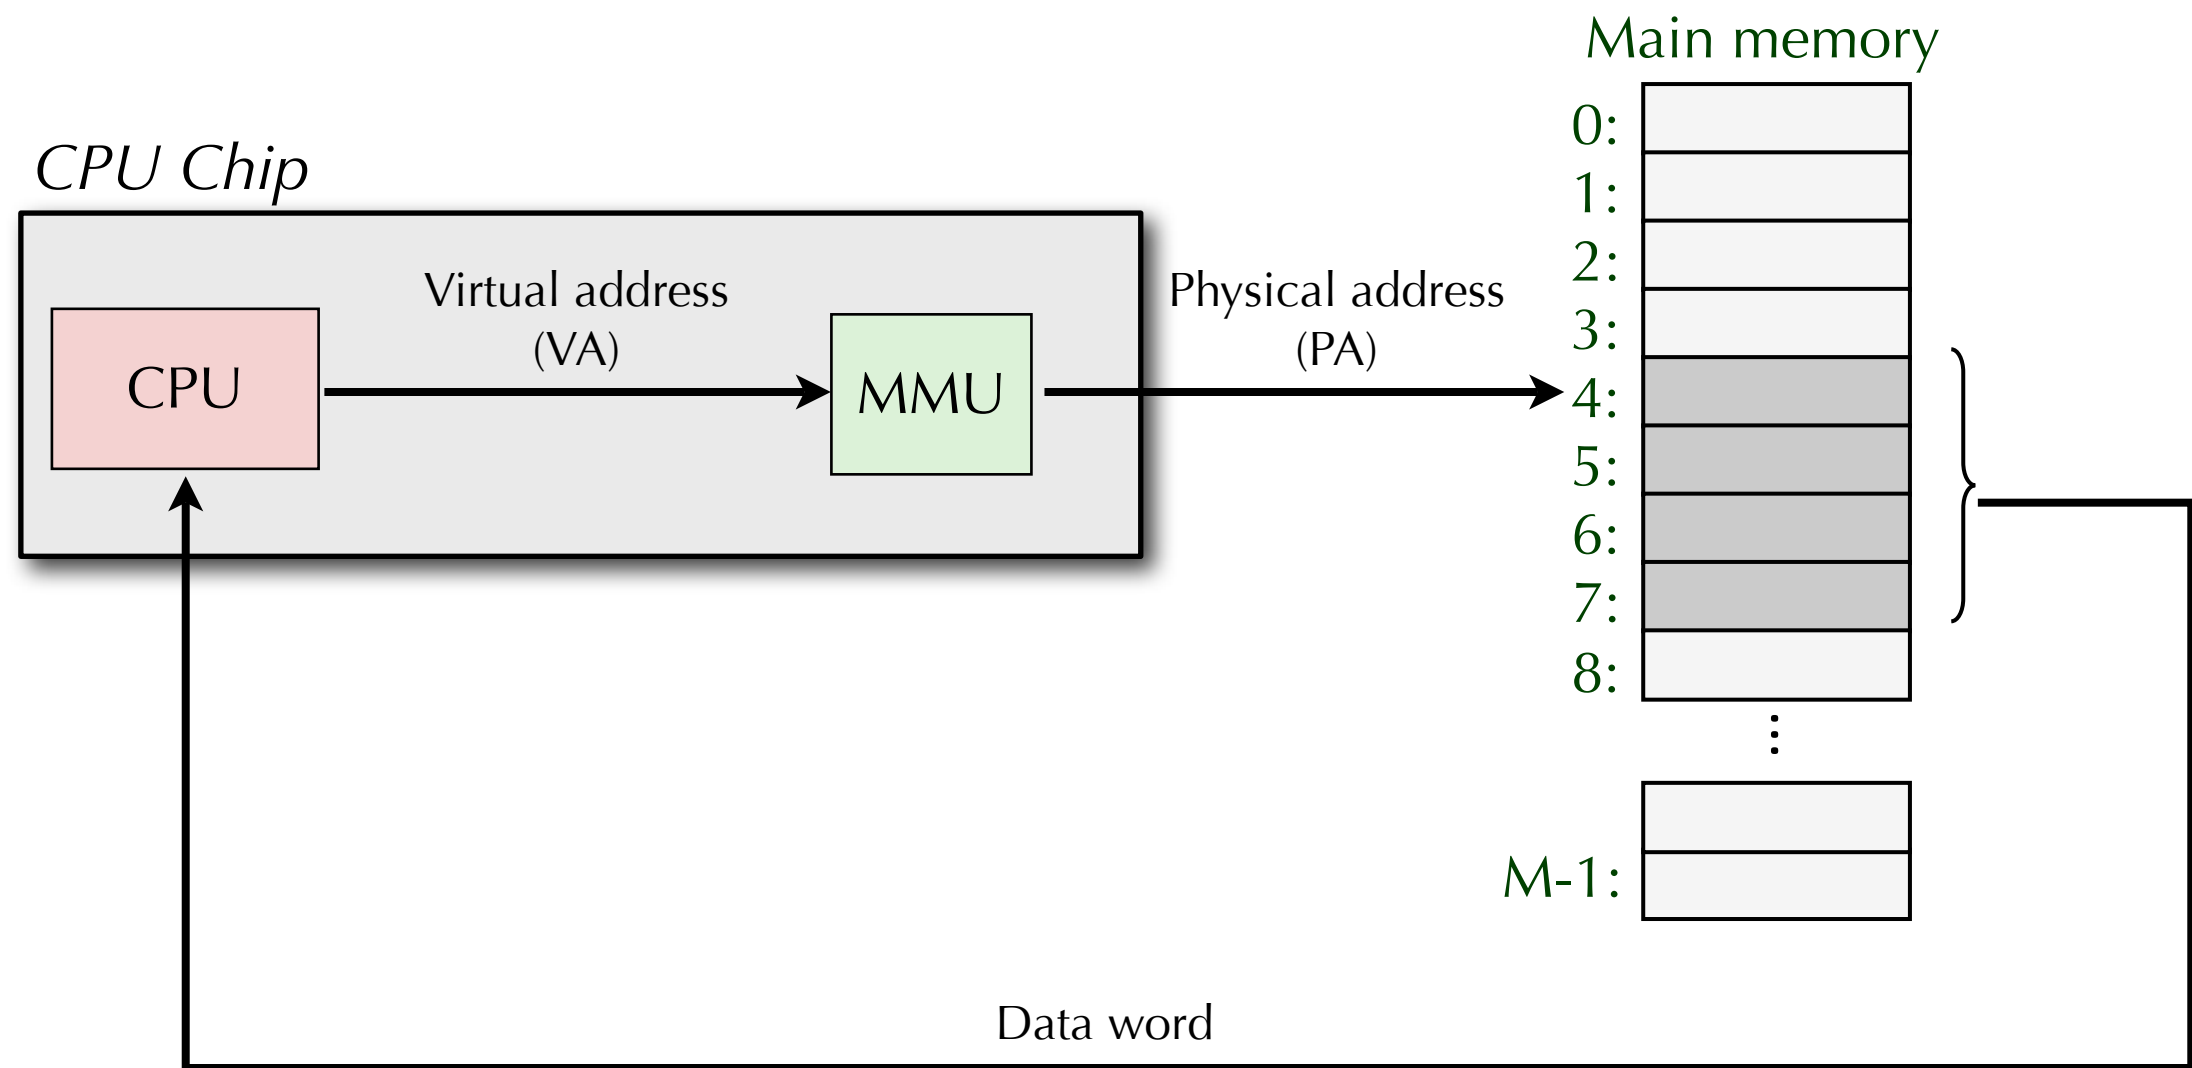
- **Linear address space:**
  - Ordered set of contiguous non-negative integer addresses:  
 $\{0, 1, 2, 3 \dots\}$
- **Virtual address space:**
  - Set of  $N = 2^n$  virtual addresses  
 $\{0, 1, 2, 3, \dots, N-1\}$
- **Physical address space:**
  - Set of  $M = 2^m$  physical addresses  
 $\{0, 1, 2, 3, \dots, M-1\}$
- Clean distinction between data (bytes) and their attributes (addresses)
- Each object can now have multiple addresses
- Every byte in main memory: one physical address, one (or more) virtual addresses

# A System Using Physical Addressing



- Used in “simple” systems like embedded microcontrollers in devices like cars, elevators, and digital picture frames

# A System Using Virtual Addressing



- Used in all modern desktops, laptops, workstations
- One of the great ideas in computer science



# Benefits of Virtual Memory (VM)

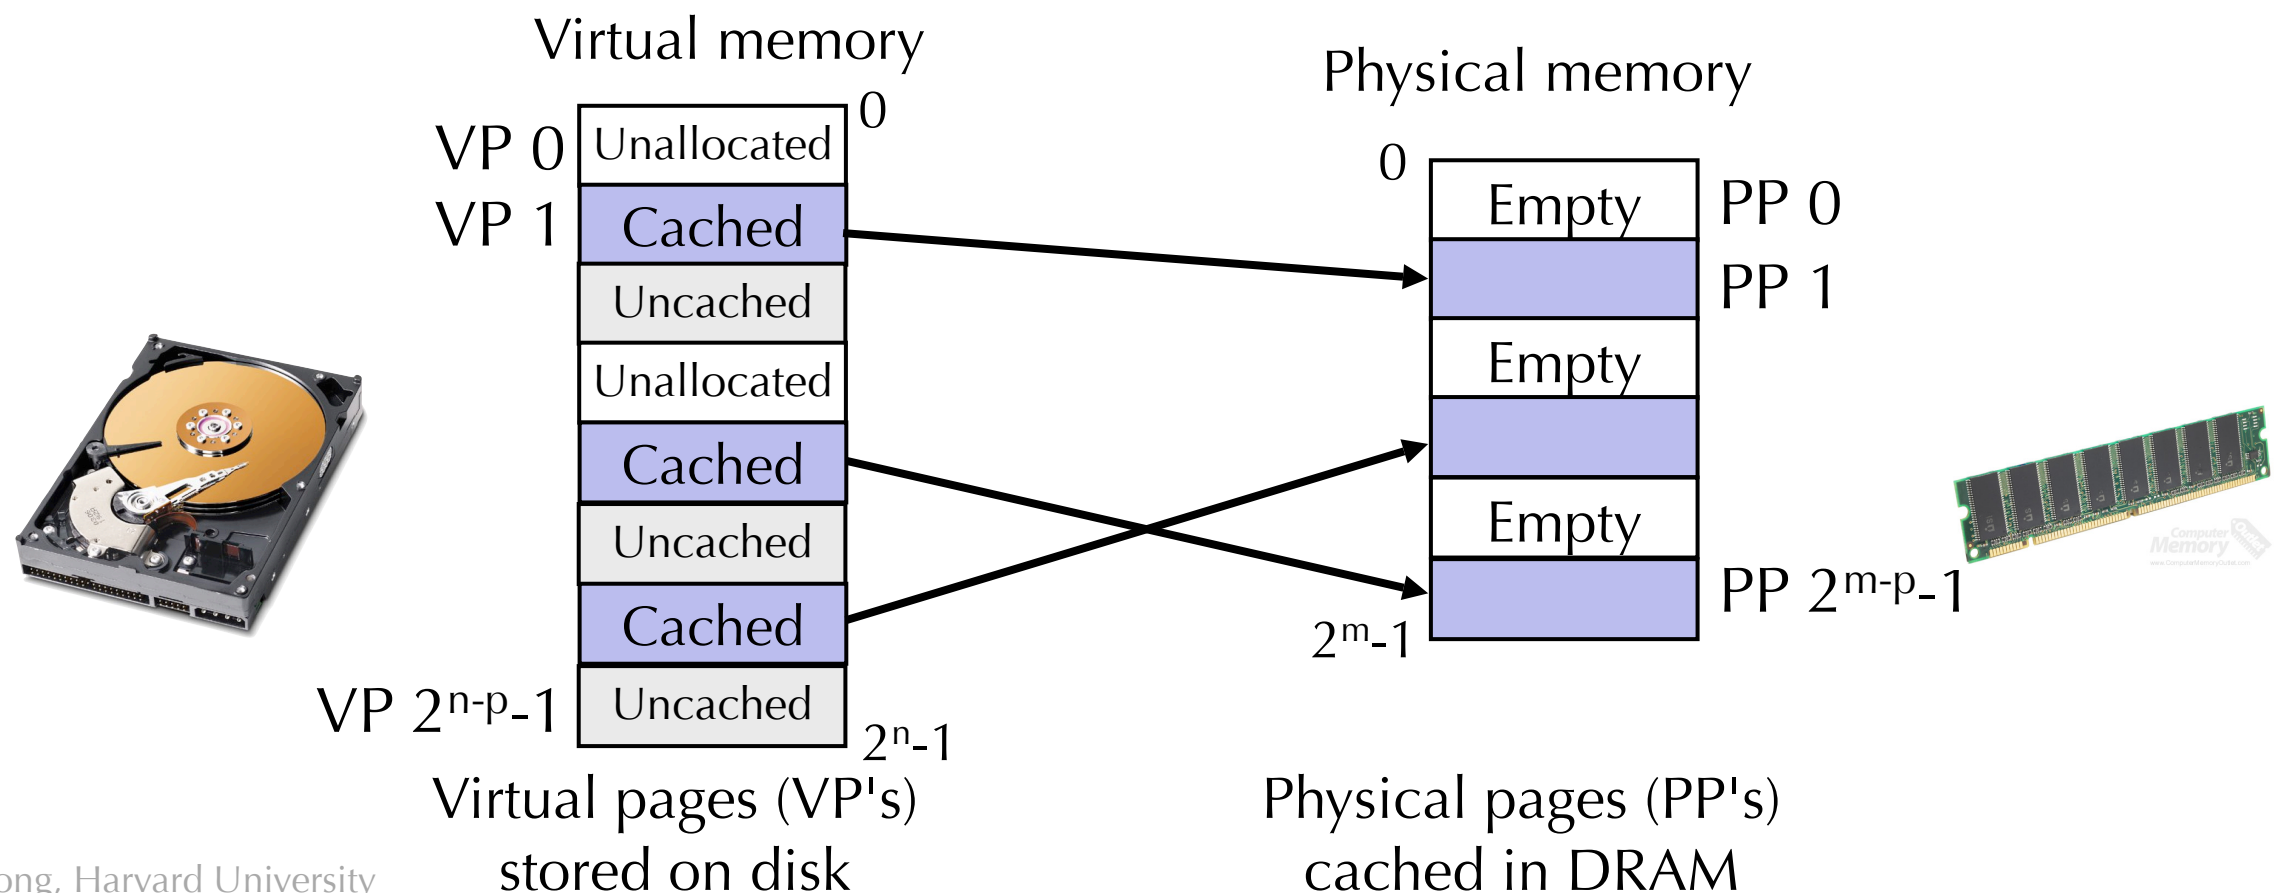
- Efficient use of limited main memory (RAM)
  - Use RAM as a cache for the parts of a virtual address space
    - some non-cached parts stored on disk
    - some (unallocated) non-cached parts stored nowhere
  - Keep only active areas of virtual address space in memory
    - transfer data back and forth as needed
- Simplifies memory management for programmers
  - Each process gets the same full, private linear address space
- Isolates address spaces
  - One process can't interfere with another's memory
    - because they operate in different address spaces
  - User process cannot access privileged information
    - different sections of address spaces have different permissions

# Today

- Running multiple programs at once
- Virtual memory
  - Address spaces
  - Benefits
  - Physical memory as a cache for virtual memory
  - Page tables
  - Page hits and page faults
  - Virtual memory as a tool for memory management
  - Virtual memory as a tool for memory protection
  - Address translation
  - Table Lookaside Buffer (TLB)
  - Example


# VM as a Tool for Caching

- Virtual memory: array of  $N = 2^n$  contiguous bytes
  - think of the array (allocated part) as being stored on disk
- Physical main memory (DRAM) = cache for allocated virtual memory
- Blocks are called **pages**; size =  $2^p$



# Memory hierarchy: Intel Core i7

## Processor



# DRAM Cache Organization

- DRAM cache organization driven by the enormous miss penalty
  - DRAM is about 10x slower than SRAM
  - Disk is about 100,000x slower than DRAM
    - For first byte, faster for next byte
- Consequences
  - Large page (block) size: typically 4-8 KB, sometimes 4 MB
  - **Fully associative**
    - Fully associate cache: A cache with only one set
    - Any virtual page can be placed in any physical page
    - Requires a “large” mapping function – different from CPU caches
  - Highly sophisticated, expensive replacement algorithms
    - Too complicated and open-ended to be implemented in hardware
  - Write-back rather than write-through

# Why does virtual memory work?

- Virtual memory works because of **locality**
- At any point in time, programs tend to access a set of active virtual pages called the working set
  - Programs with better temporal locality will have smaller working sets
- If (working set size  $<$  main memory size)
  - Good performance for one process after compulsory misses
- If (  $SUM(\text{working set sizes}) > \text{main memory size}$  )
  - **Thrashing**: Performance meltdown where pages are swapped (copied) in and out continuously

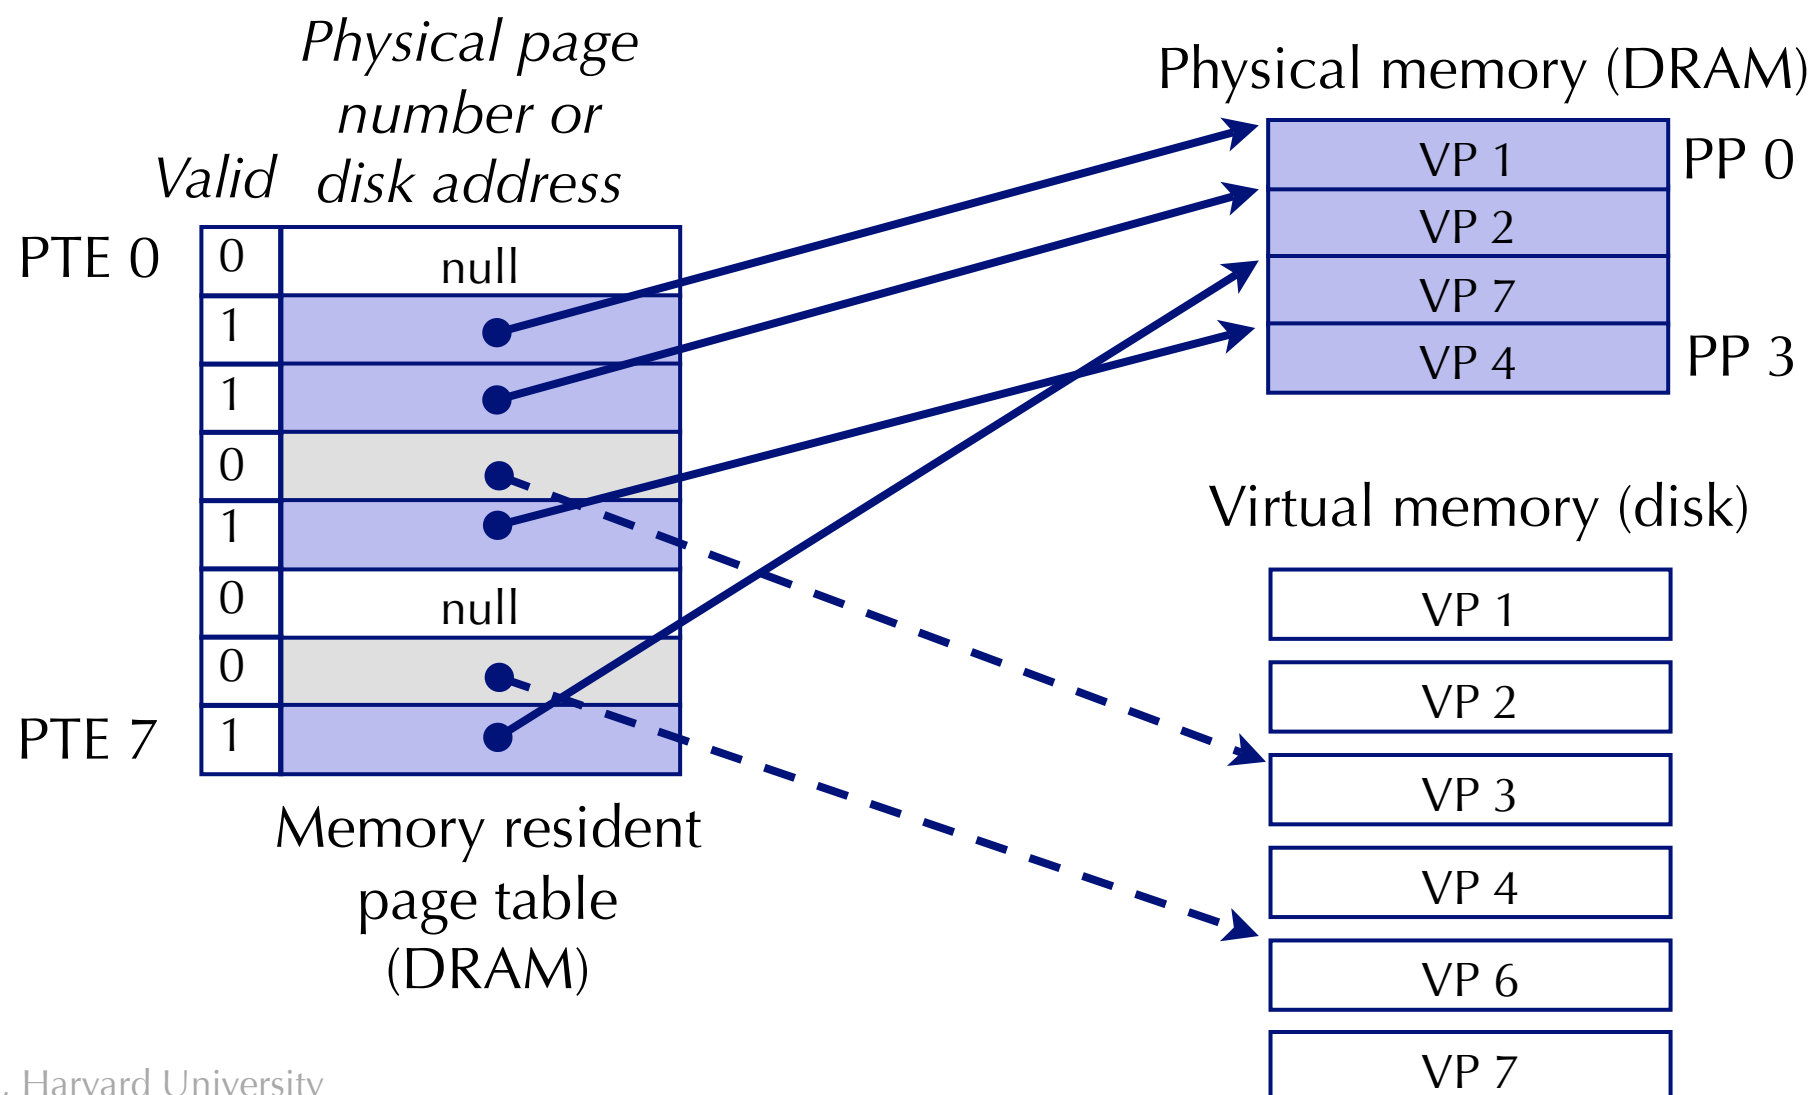


# Today

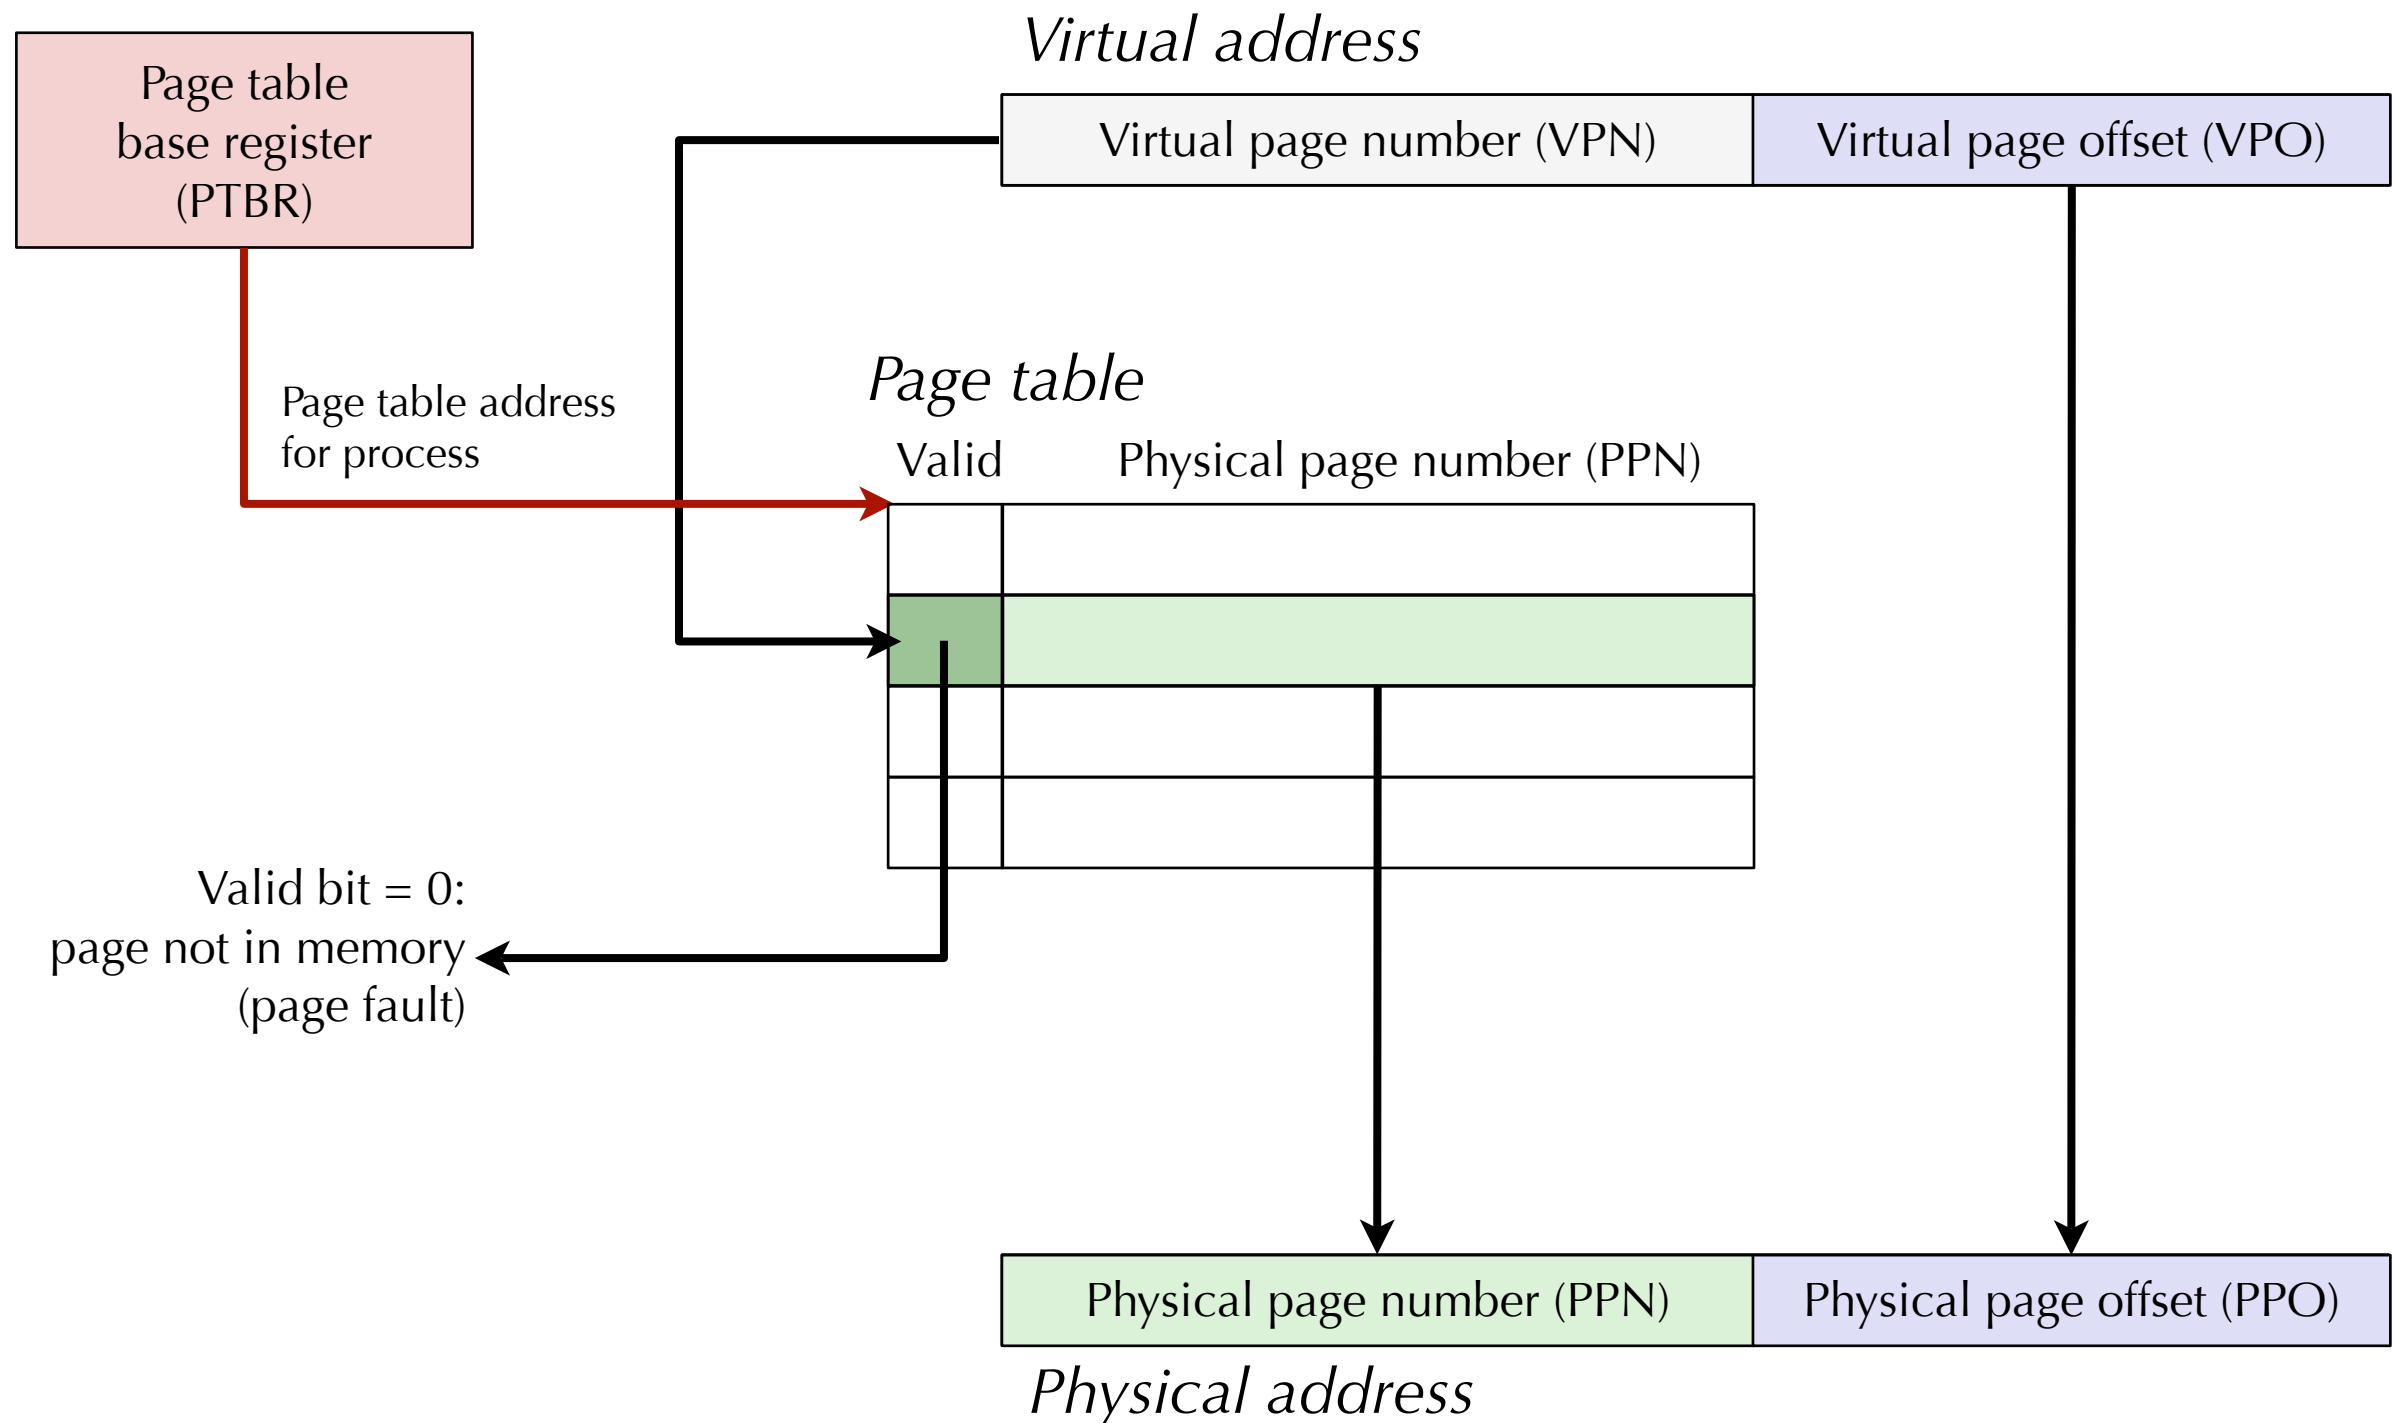
- Running multiple programs at once
- Virtual memory
  - Address spaces
  - Benefits
  - Physical memory as a cache for virtual memory
  - Page tables
  - Page hits and page faults
  - Virtual memory as a tool for memory management
  - Virtual memory as a tool for memory protection
  - Address translation
  - Table Lookaside Buffer (TLB)
  - Example

# Page Tables

- A **page table** is an array of page table entries (PTEs) that maps virtual pages to physical pages. Here: 8 VPs
  - Per-process kernel data structure in DRAM



# Address Translation With a Page Table



# Page table size

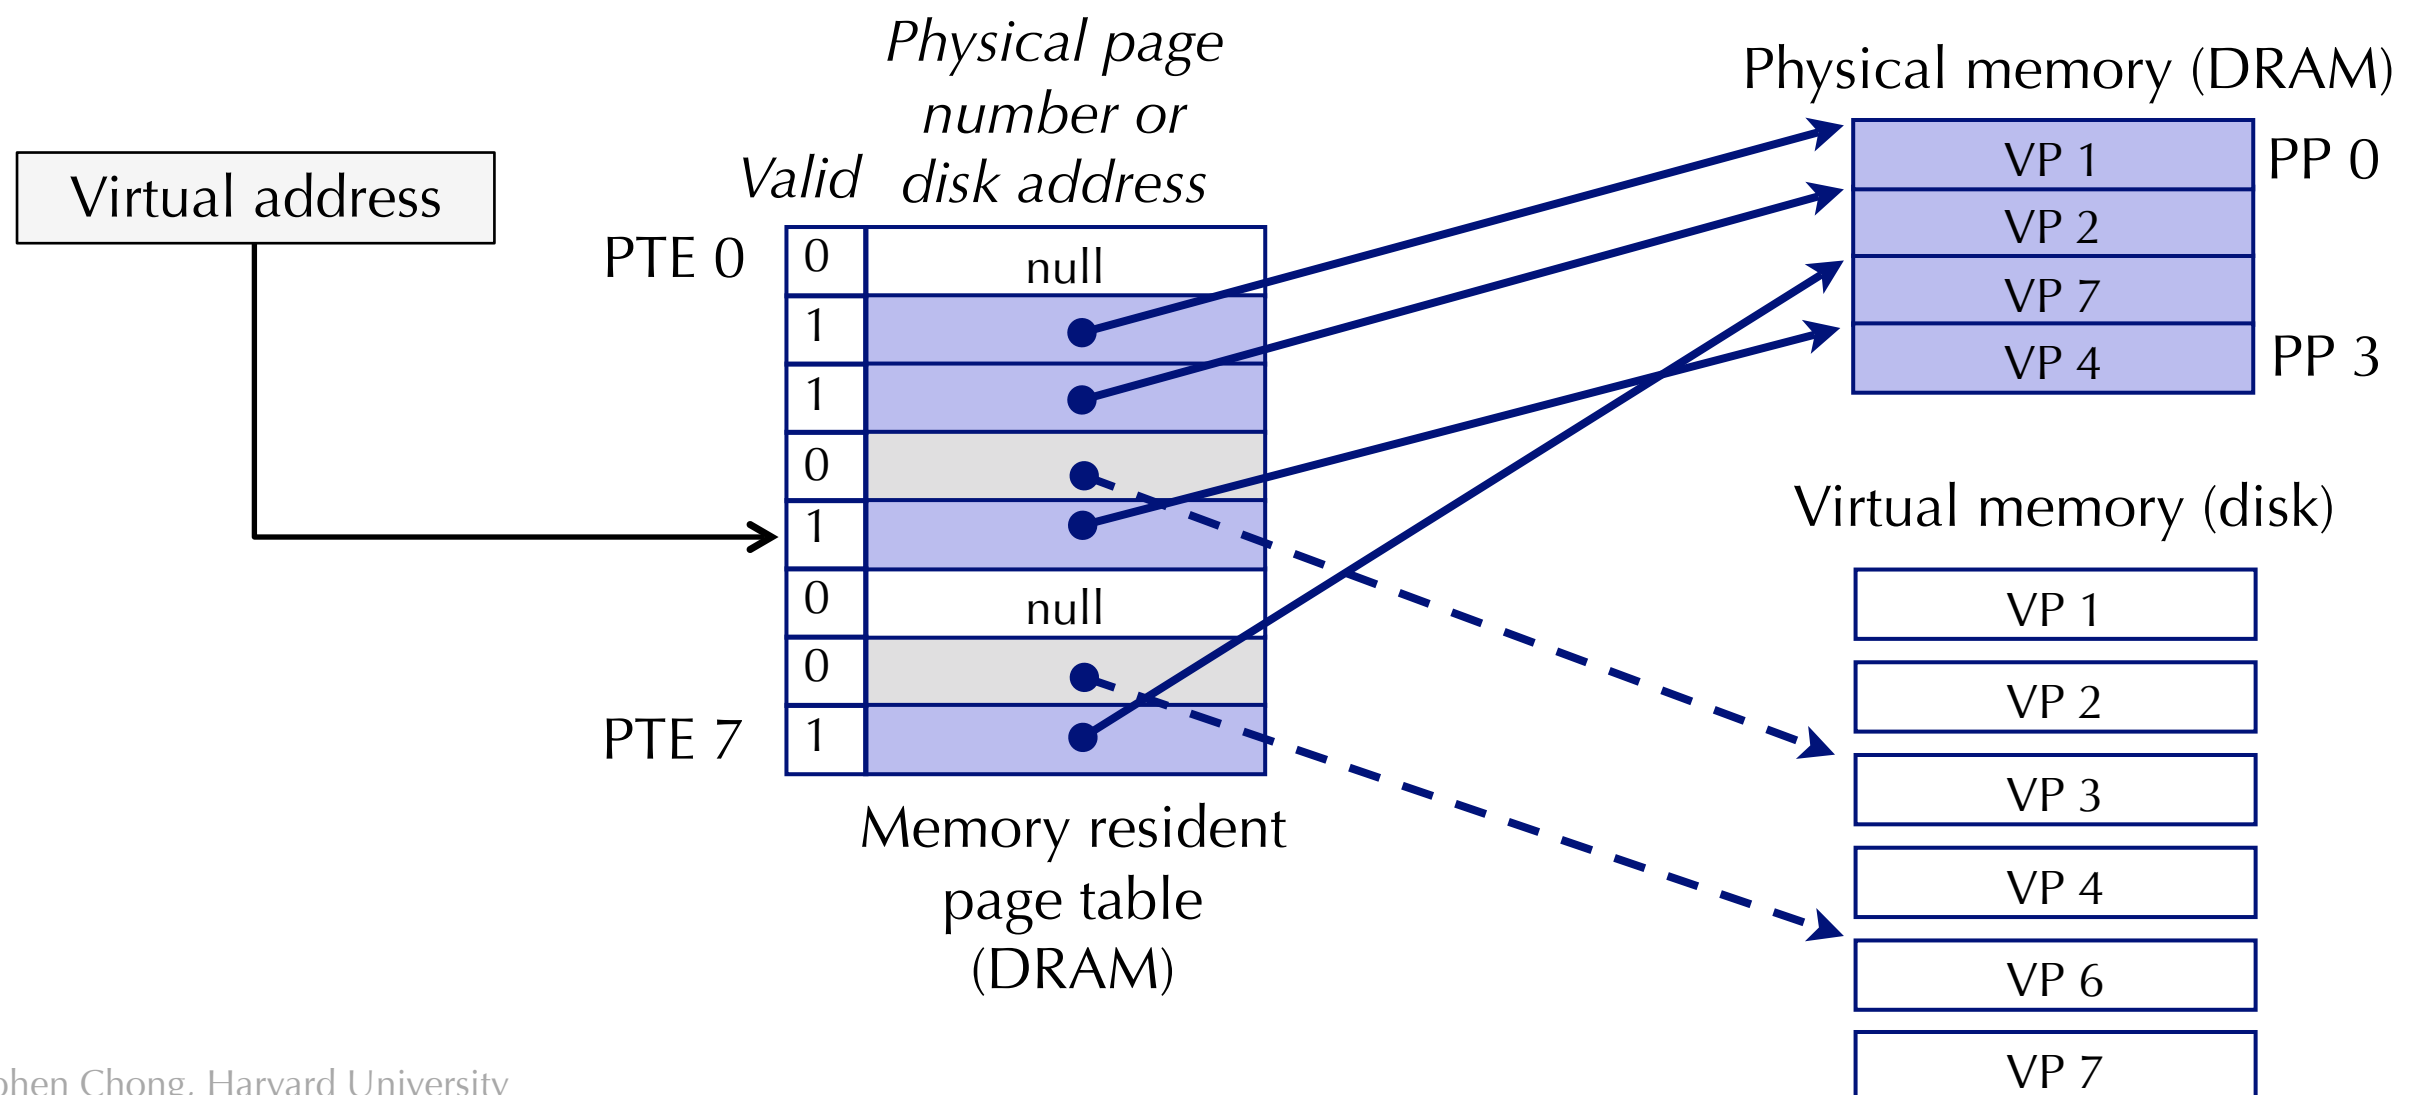
- How big is the page table for a process?
- Well ... we need one PTE per page.
- Say we have a 32-bit address space, and the page size is 4KB
- How many pages?
  - $2^{32} = 4$  GB total memory
  - $4\text{GB} / 4\text{KB} = 1,048,576 (= 1\text{M})$  pages = 1M PTEs
- How big is a PTE?
  - Depends on CPU architecture. On x86 it is 4 bytes
- Page table size for one process:  $1\text{M PTEs} \times 4 \text{ bytes/PTE} = 4 \text{ MB}$
- For 100 processes that's 400 MB of memory just for page tables!!!
- Solution: swap page tables out to disk...

# Today

- Running multiple programs at once
- Virtual memory
  - Address spaces
  - Benefits
  - Physical memory as a cache for virtual memory
  - Page tables
  - Page hits and page faults
  - Virtual memory as a tool for memory management
  - Virtual memory as a tool for memory protection
  - Address translation
  - Table Lookaside Buffer (TLB)
  - Example

# Page Hit

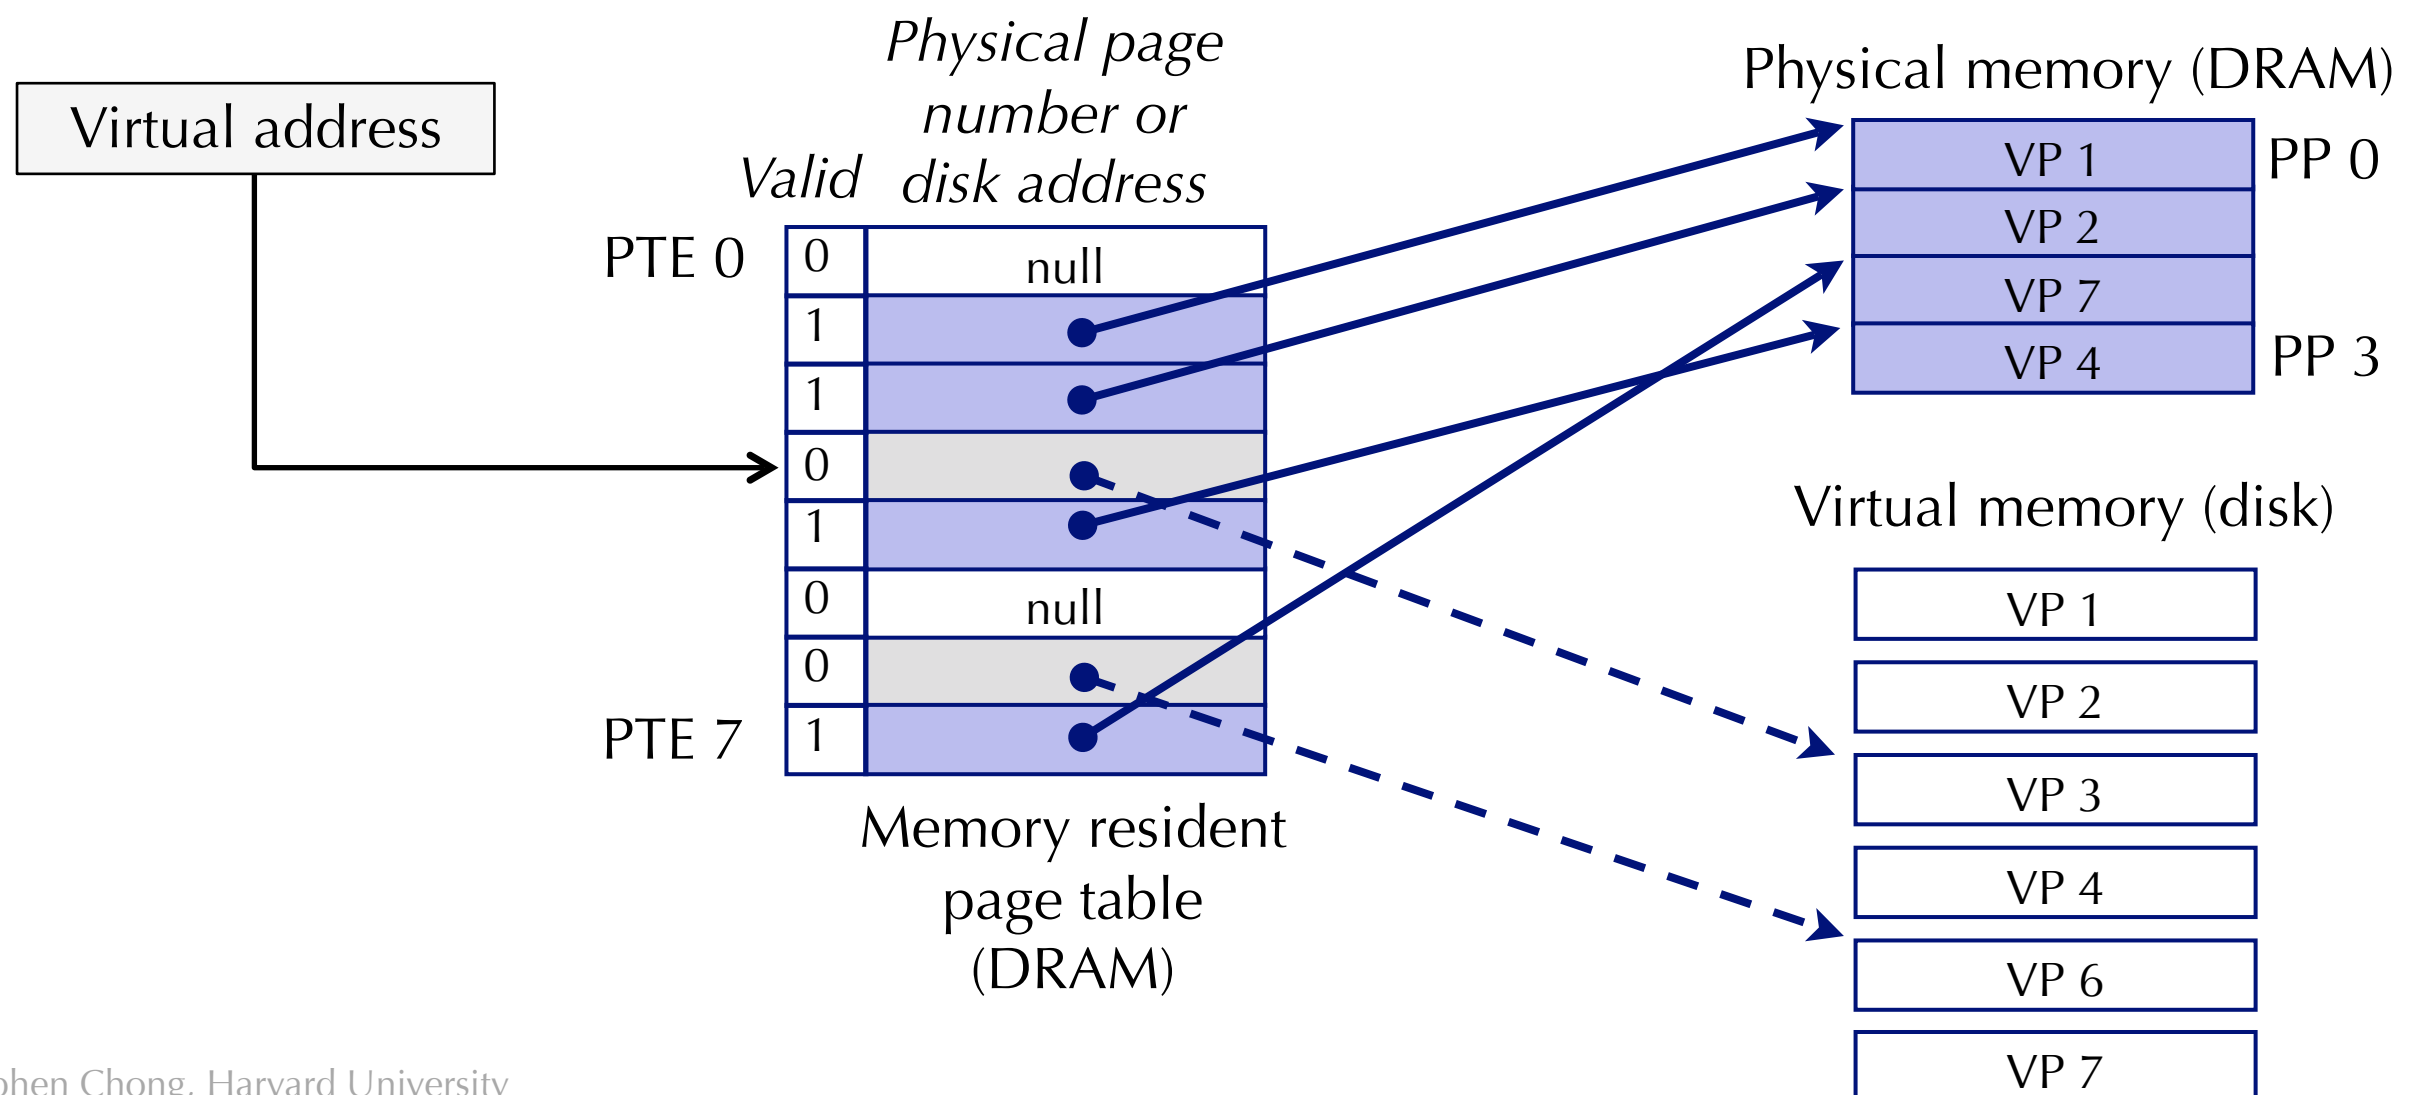
- **Page hit:** reference to VM word that is in physical memory





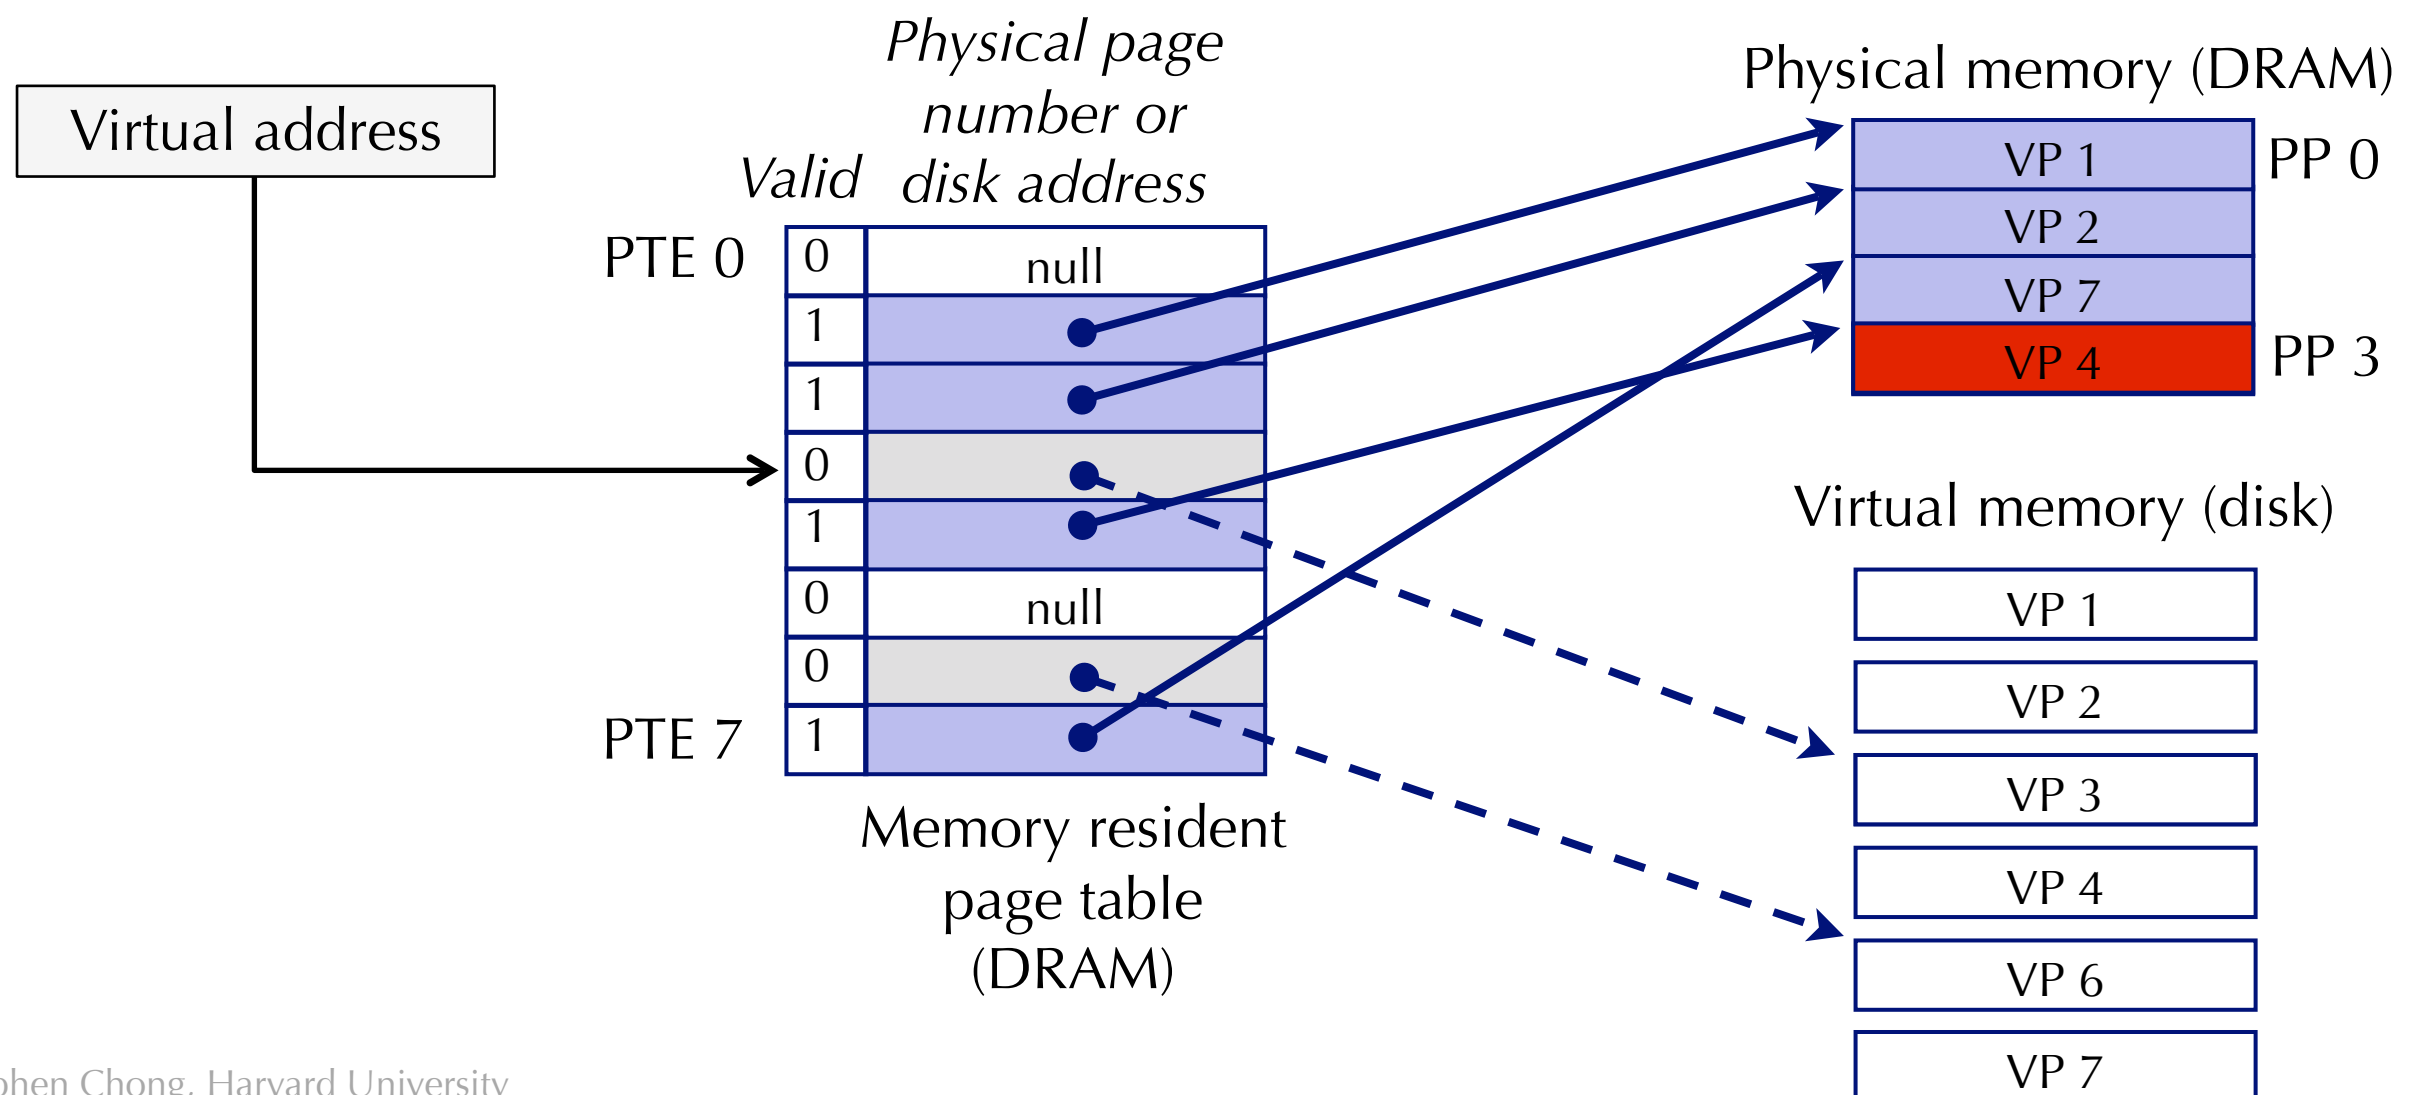
# Page fault

- **Page fault:** reference to VM word that is not in physical memory
  - i.e., invalid entry in page table
  - Could be bad memory address, or page table currently on disk



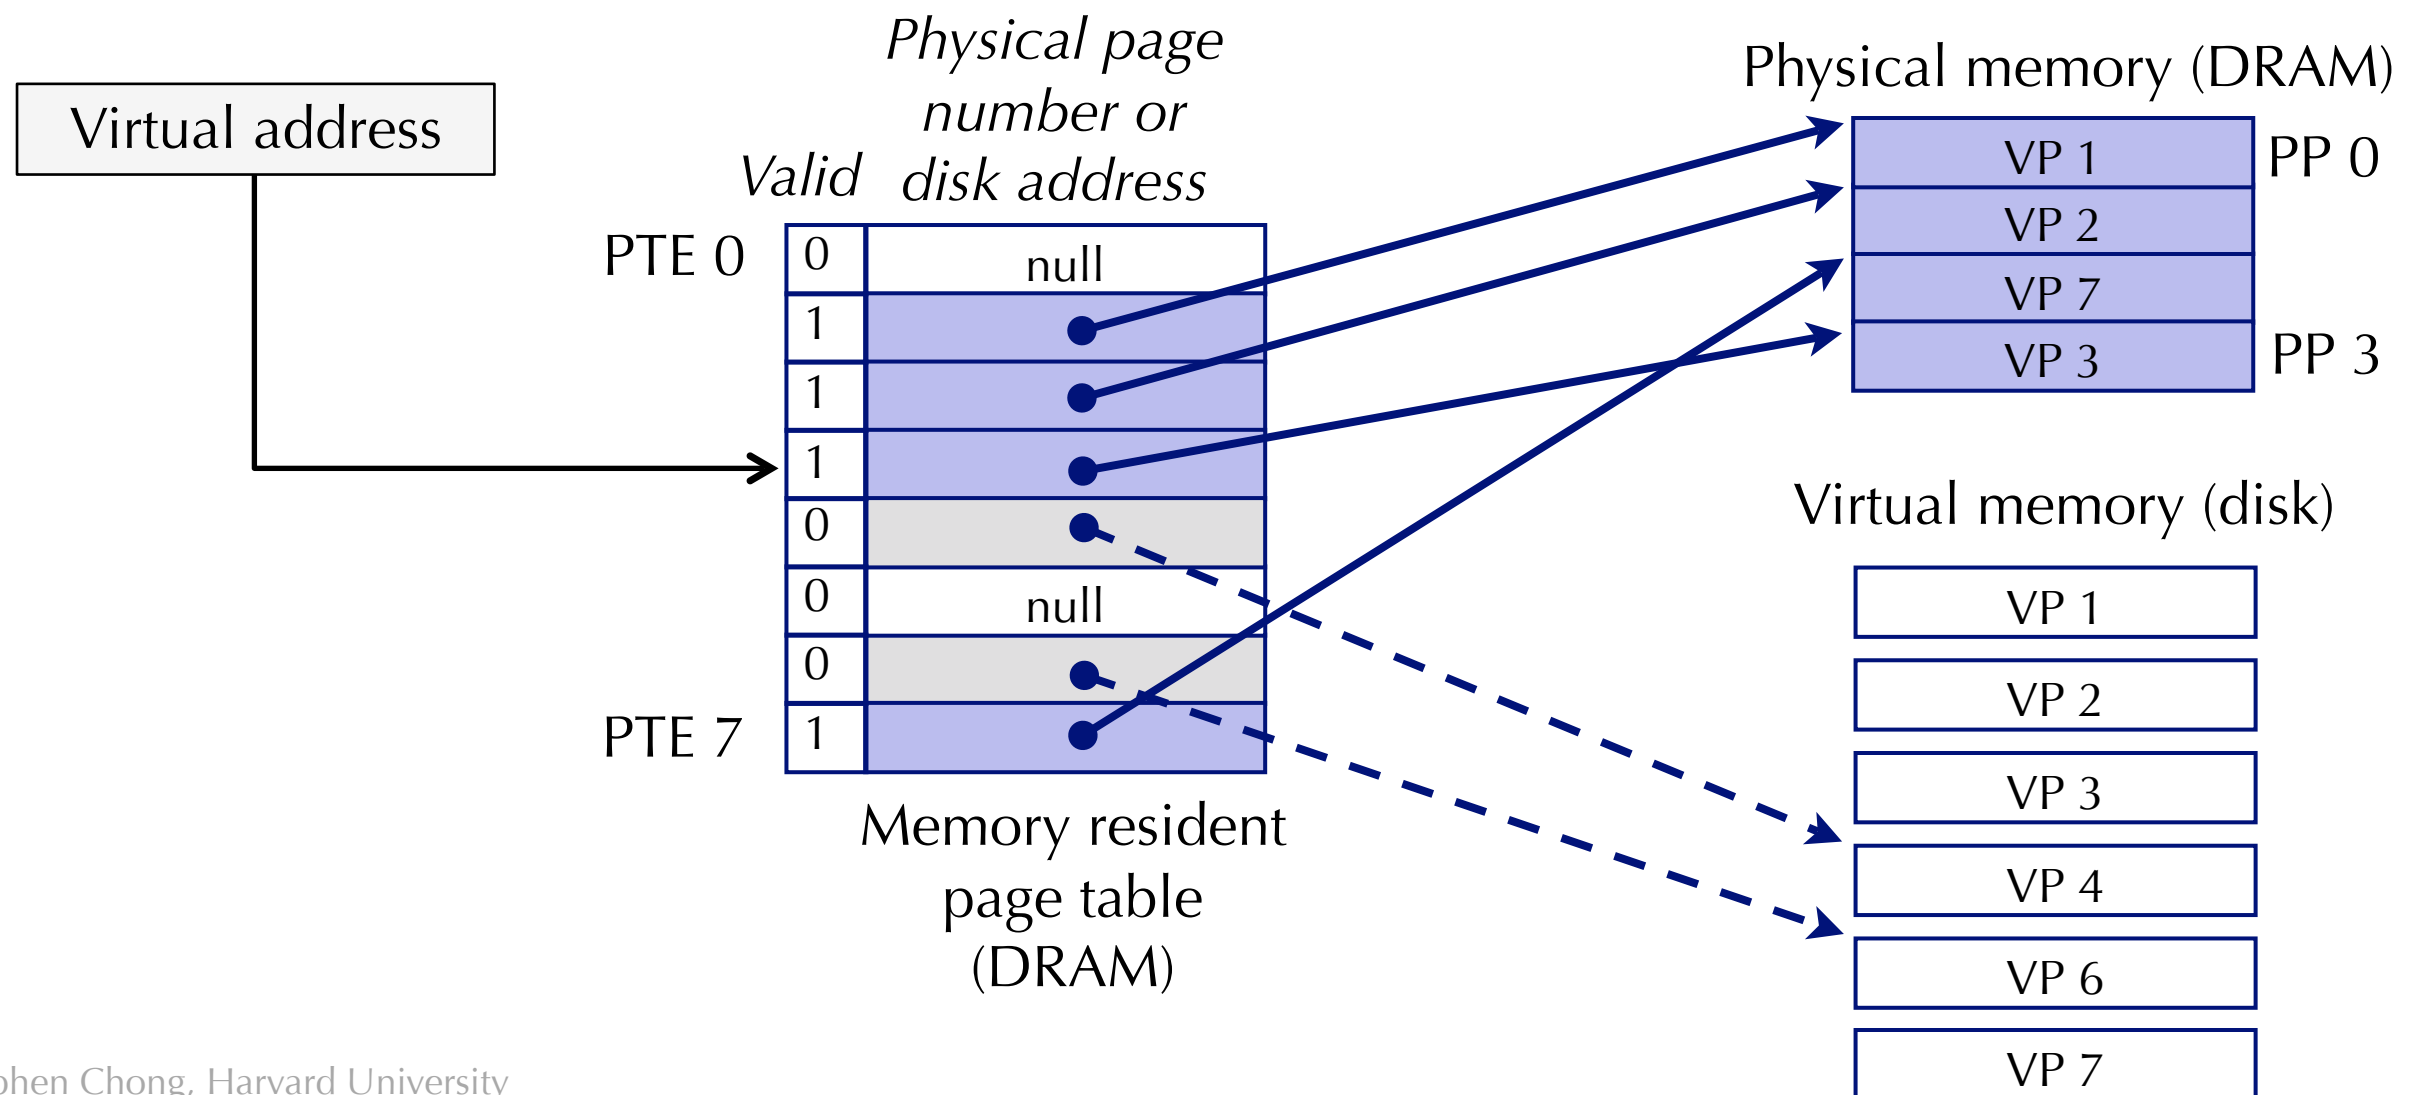
# Handling page faults

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)



# Handling page faults

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)
- Offending instruction is restarted: page hit!

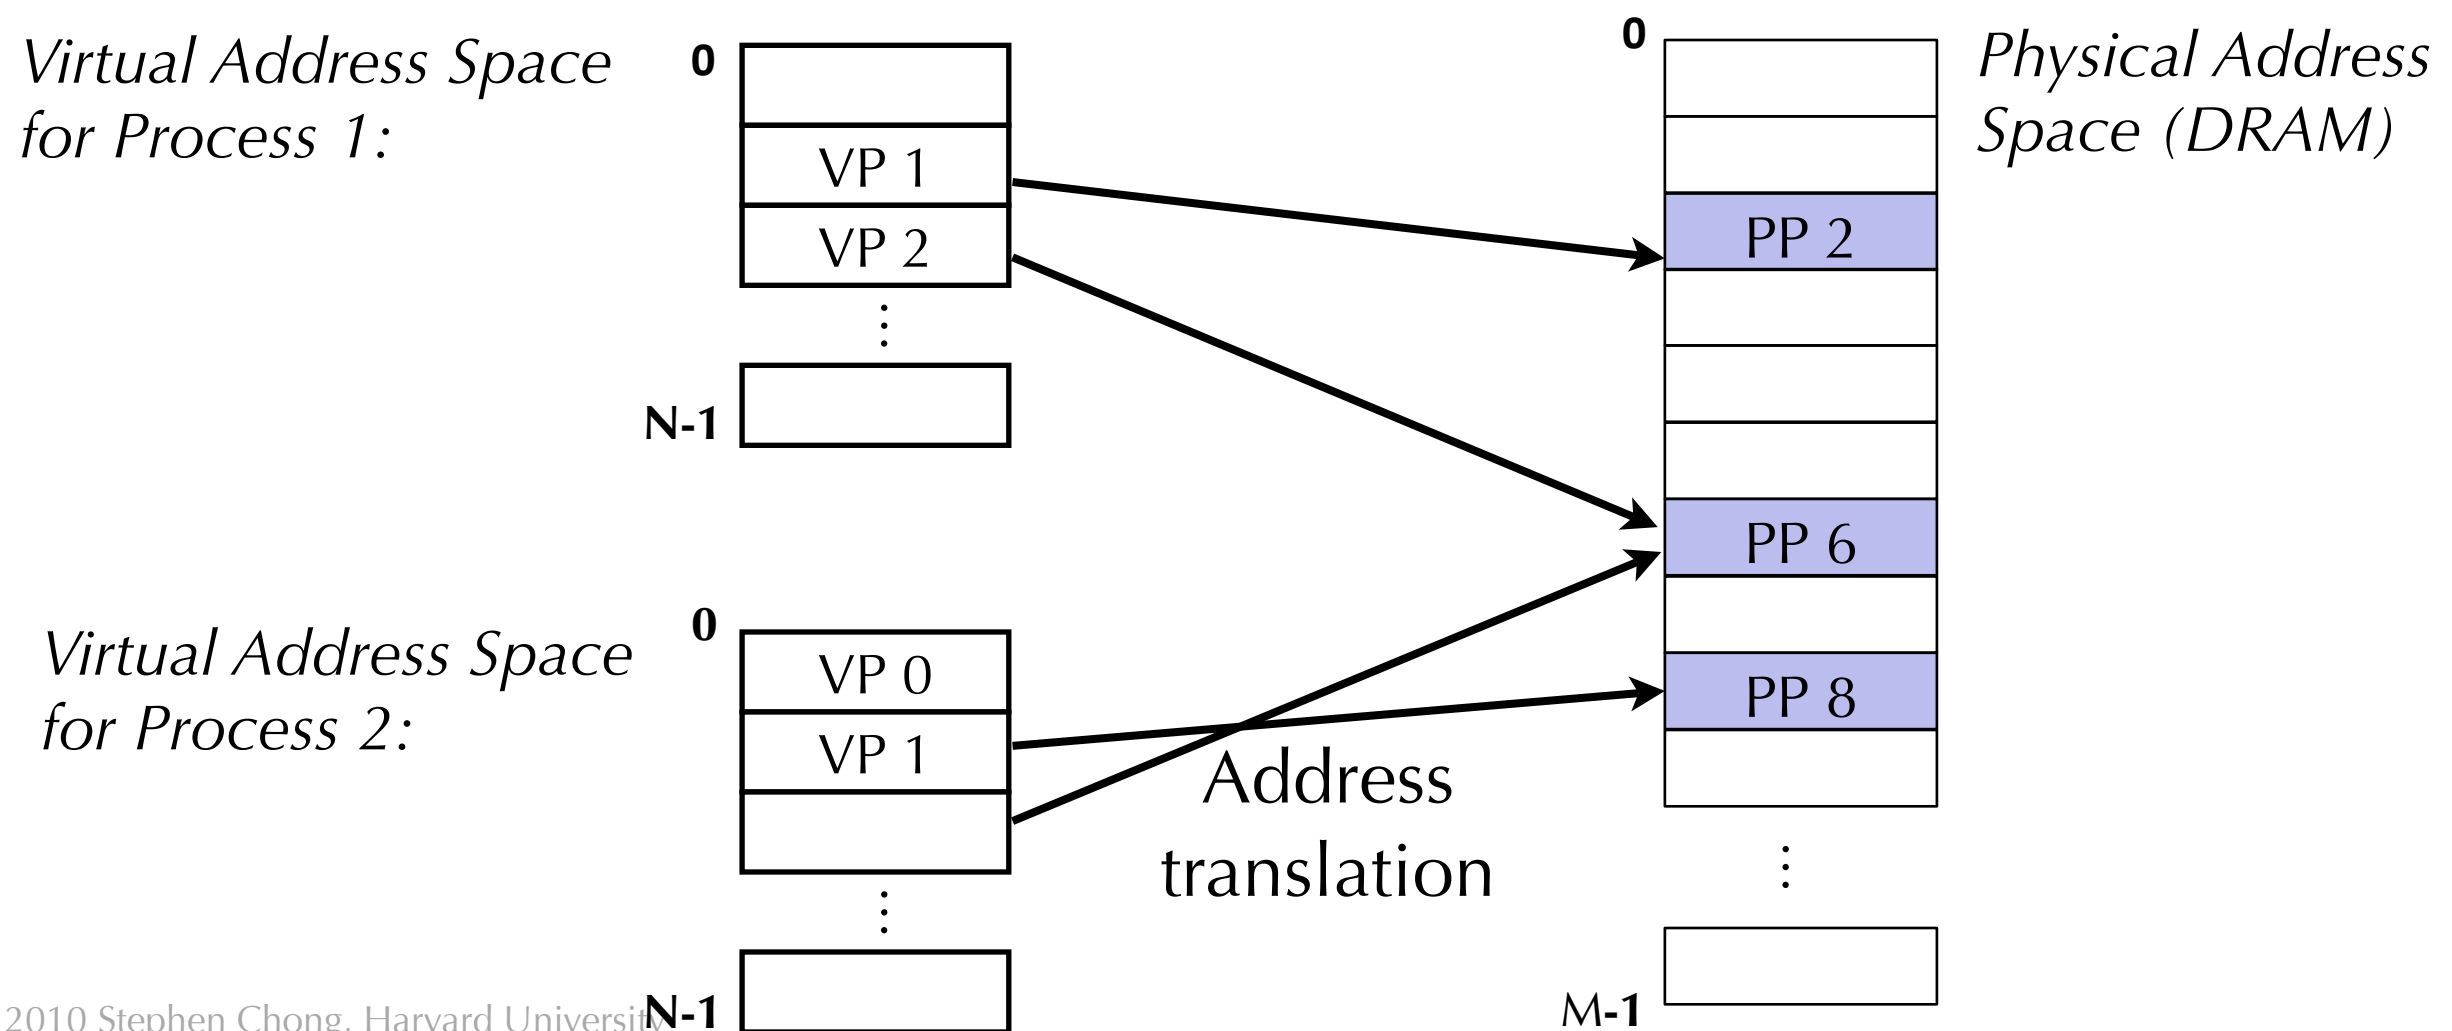


# Today

- Running multiple programs at once
- Virtual memory
  - Address spaces
  - Benefits
  - Physical memory as a cache for virtual memory
  - Page tables
  - Page hits and page faults
  - Virtual memory as a tool for memory management
  - Virtual memory as a tool for memory protection
  - Address translation
  - Table Lookaside Buffer (TLB)
  - Example

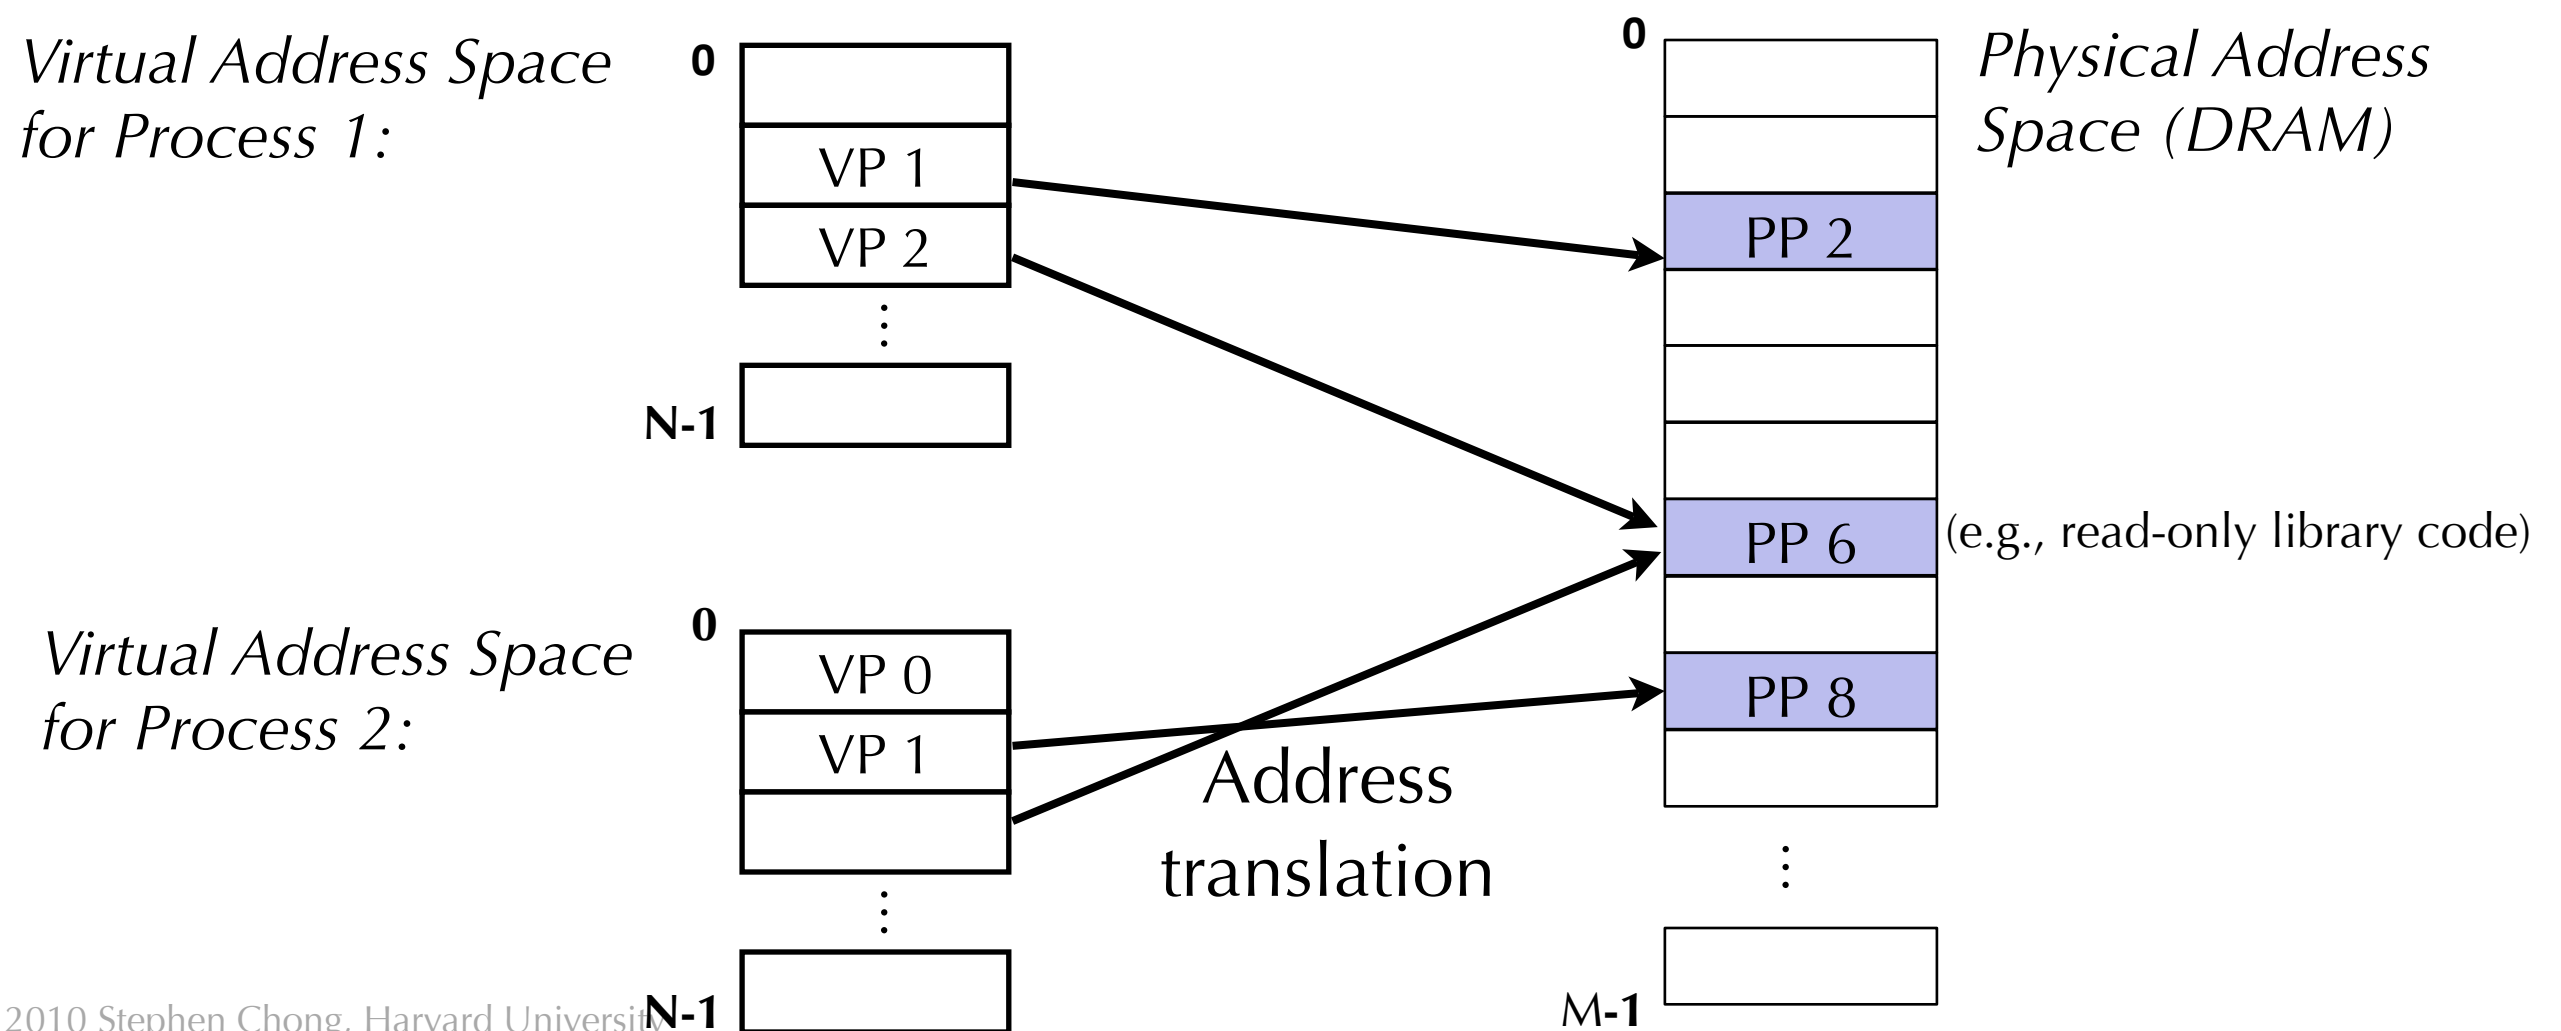
# VM as a Tool for Memory Management

- Key idea: each process has its own virtual address space
  - It can view memory as a simple linear array
  - Mapping function scatters addresses through physical memory
    - Well chosen mappings simplify memory allocation and management



# VM as a Tool for Memory Management

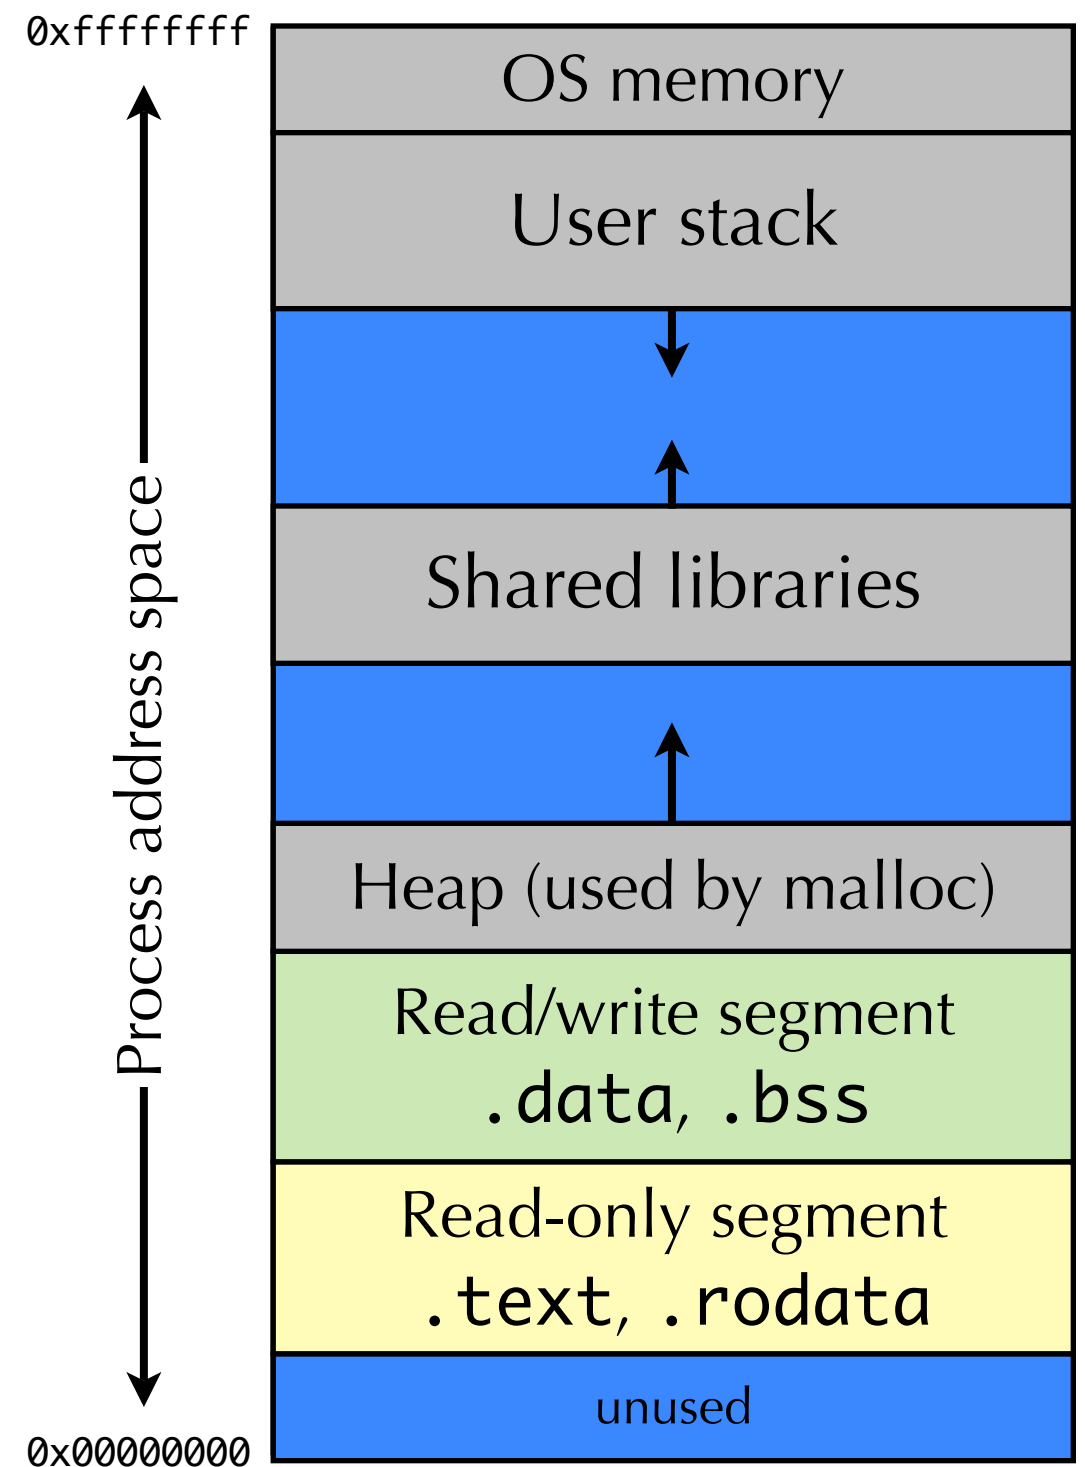
- Memory allocation
  - Each virtual page can be mapped to any physical page
  - A virtual page can be stored in different physical pages at different times
- Sharing code and data among processes
  - Map virtual pages to the same physical page (here: PP 6)






# Simplifying linking

- Linking
  - Each program has similar virtual address space
  - Code, stack, and shared libraries always start at the same address



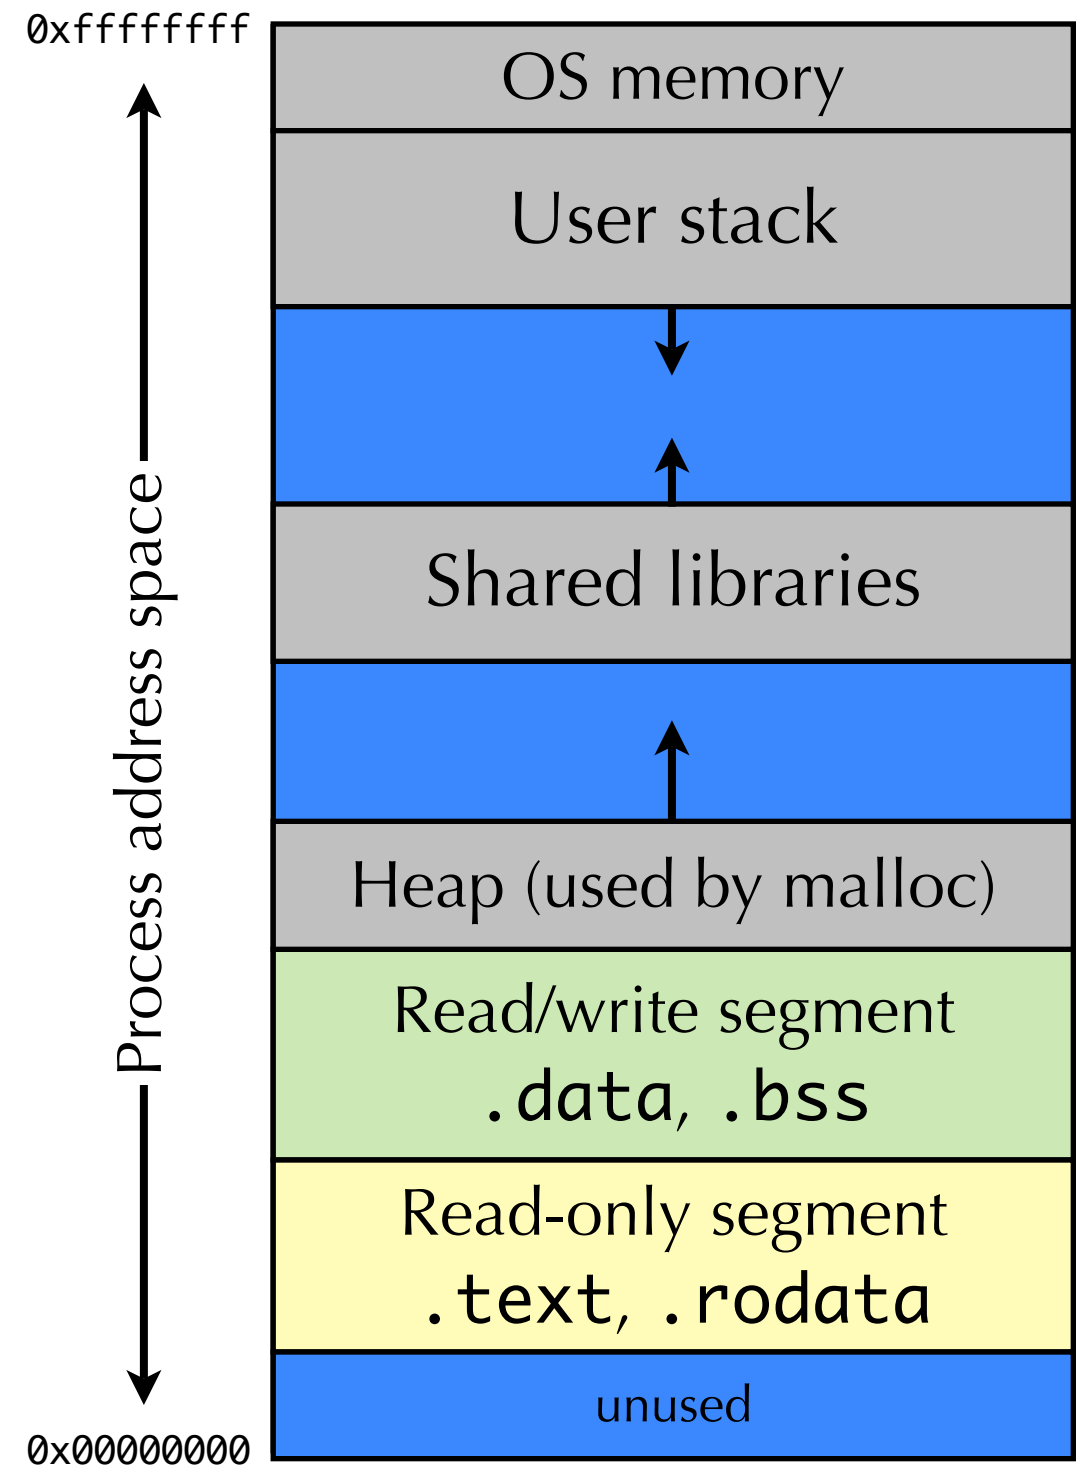
# Simplifying loading

- Doesn't make sense to immediately load all of a program into physical memory
  - Lots of code and data may not be used
  - E.g., rarely used functionality
- Initially, page table maps data, text, rodata, etc., segments to disk
  - i.e., PTE are marked as invalid
- As segments are used (e.g., code accessed for the first time), page fault will bring them into physical memory
- **Demand paging:** pages copied into memory only when they are needed



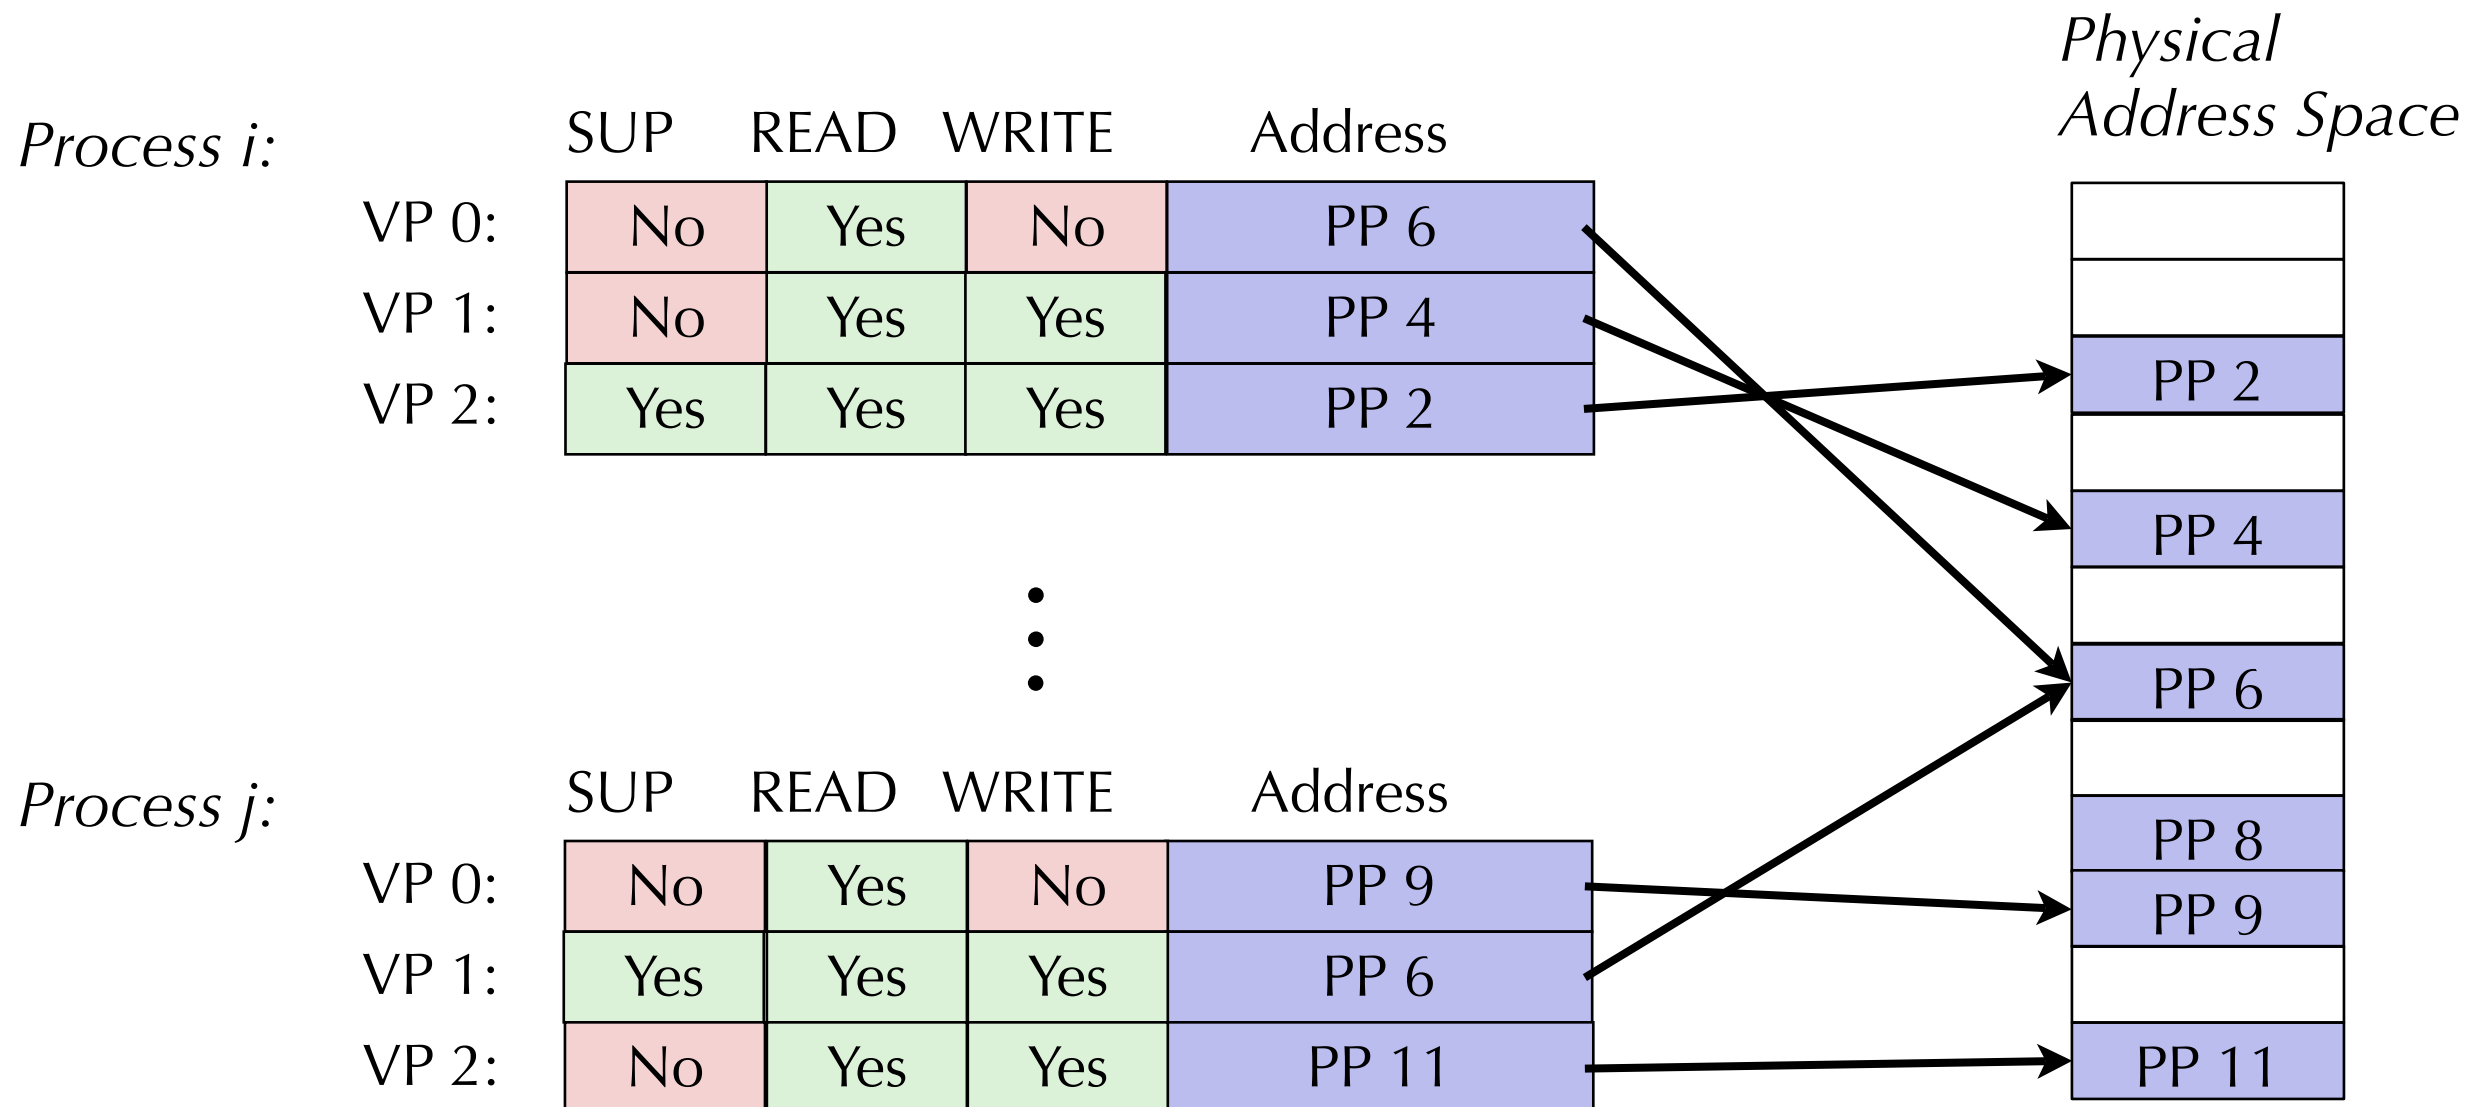
# Demand-zero pages

- When page from bss segment (and new memory for the heap, stack, etc.) faulted for the first time, physical page of all zeros allocated.
  - Once page is written, like any other page.
  - **“Demand-zero” pages**



# VM as a Tool for Memory Protection

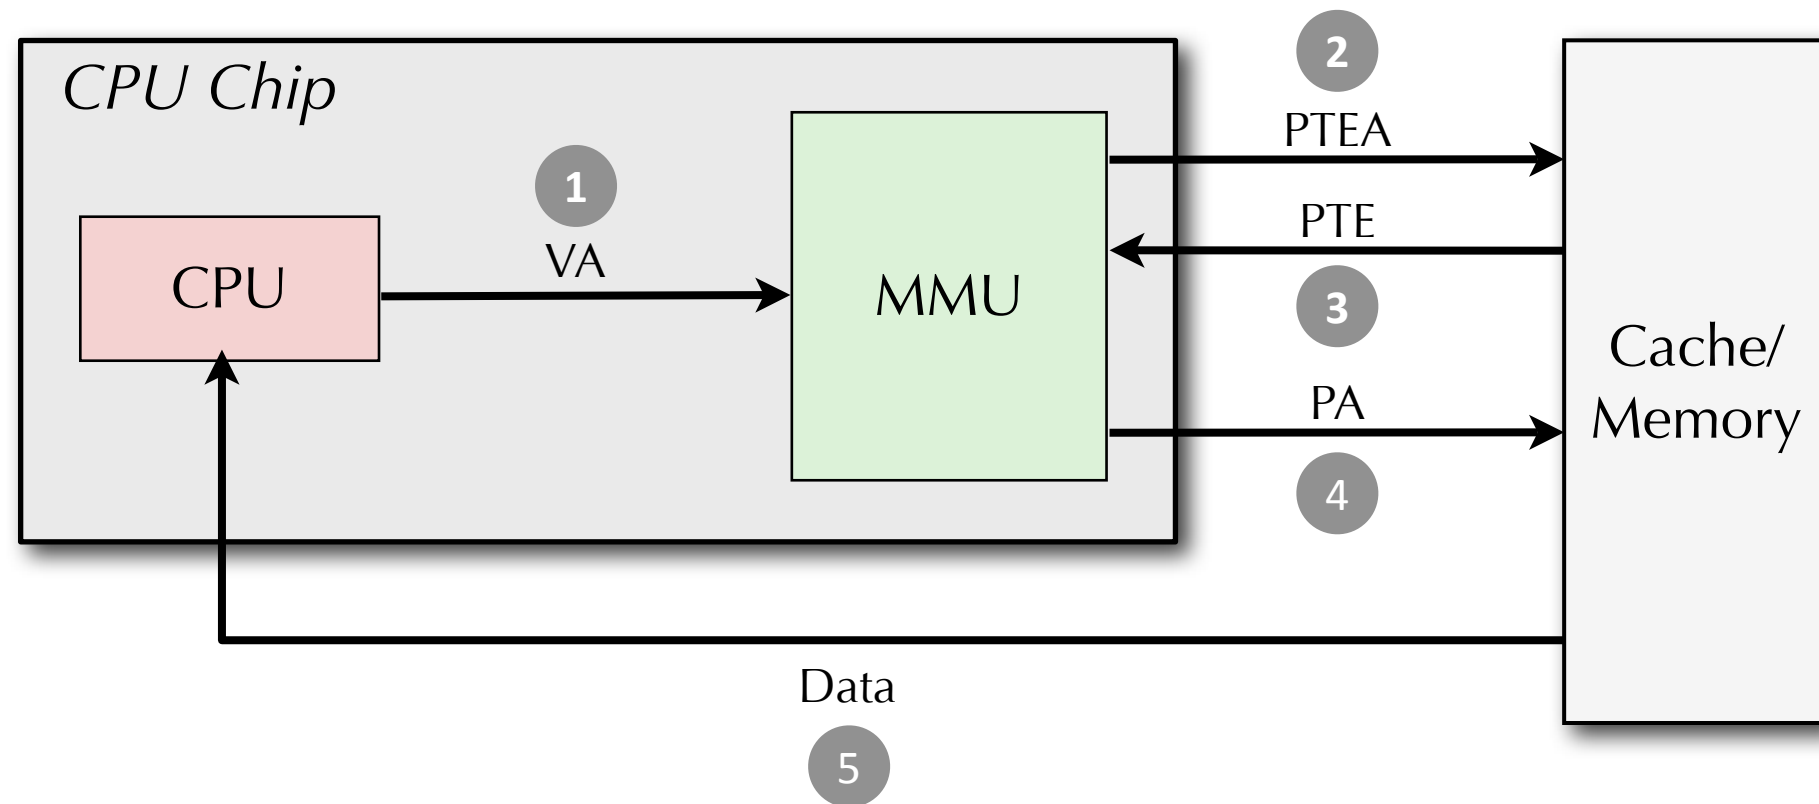
- Extend PTEs with permission bits
- Page fault handler checks these before remapping



# Today

- Running multiple programs at once
- Virtual memory
  - Address spaces
  - Benefits
  - Physical memory as a cache for virtual memory
  - Page tables
  - Page hits and page faults
  - Virtual memory as a tool for memory management
  - Virtual memory as a tool for memory protection
  - Address translation
  - Table Lookaside Buffer (TLB)
  - Example

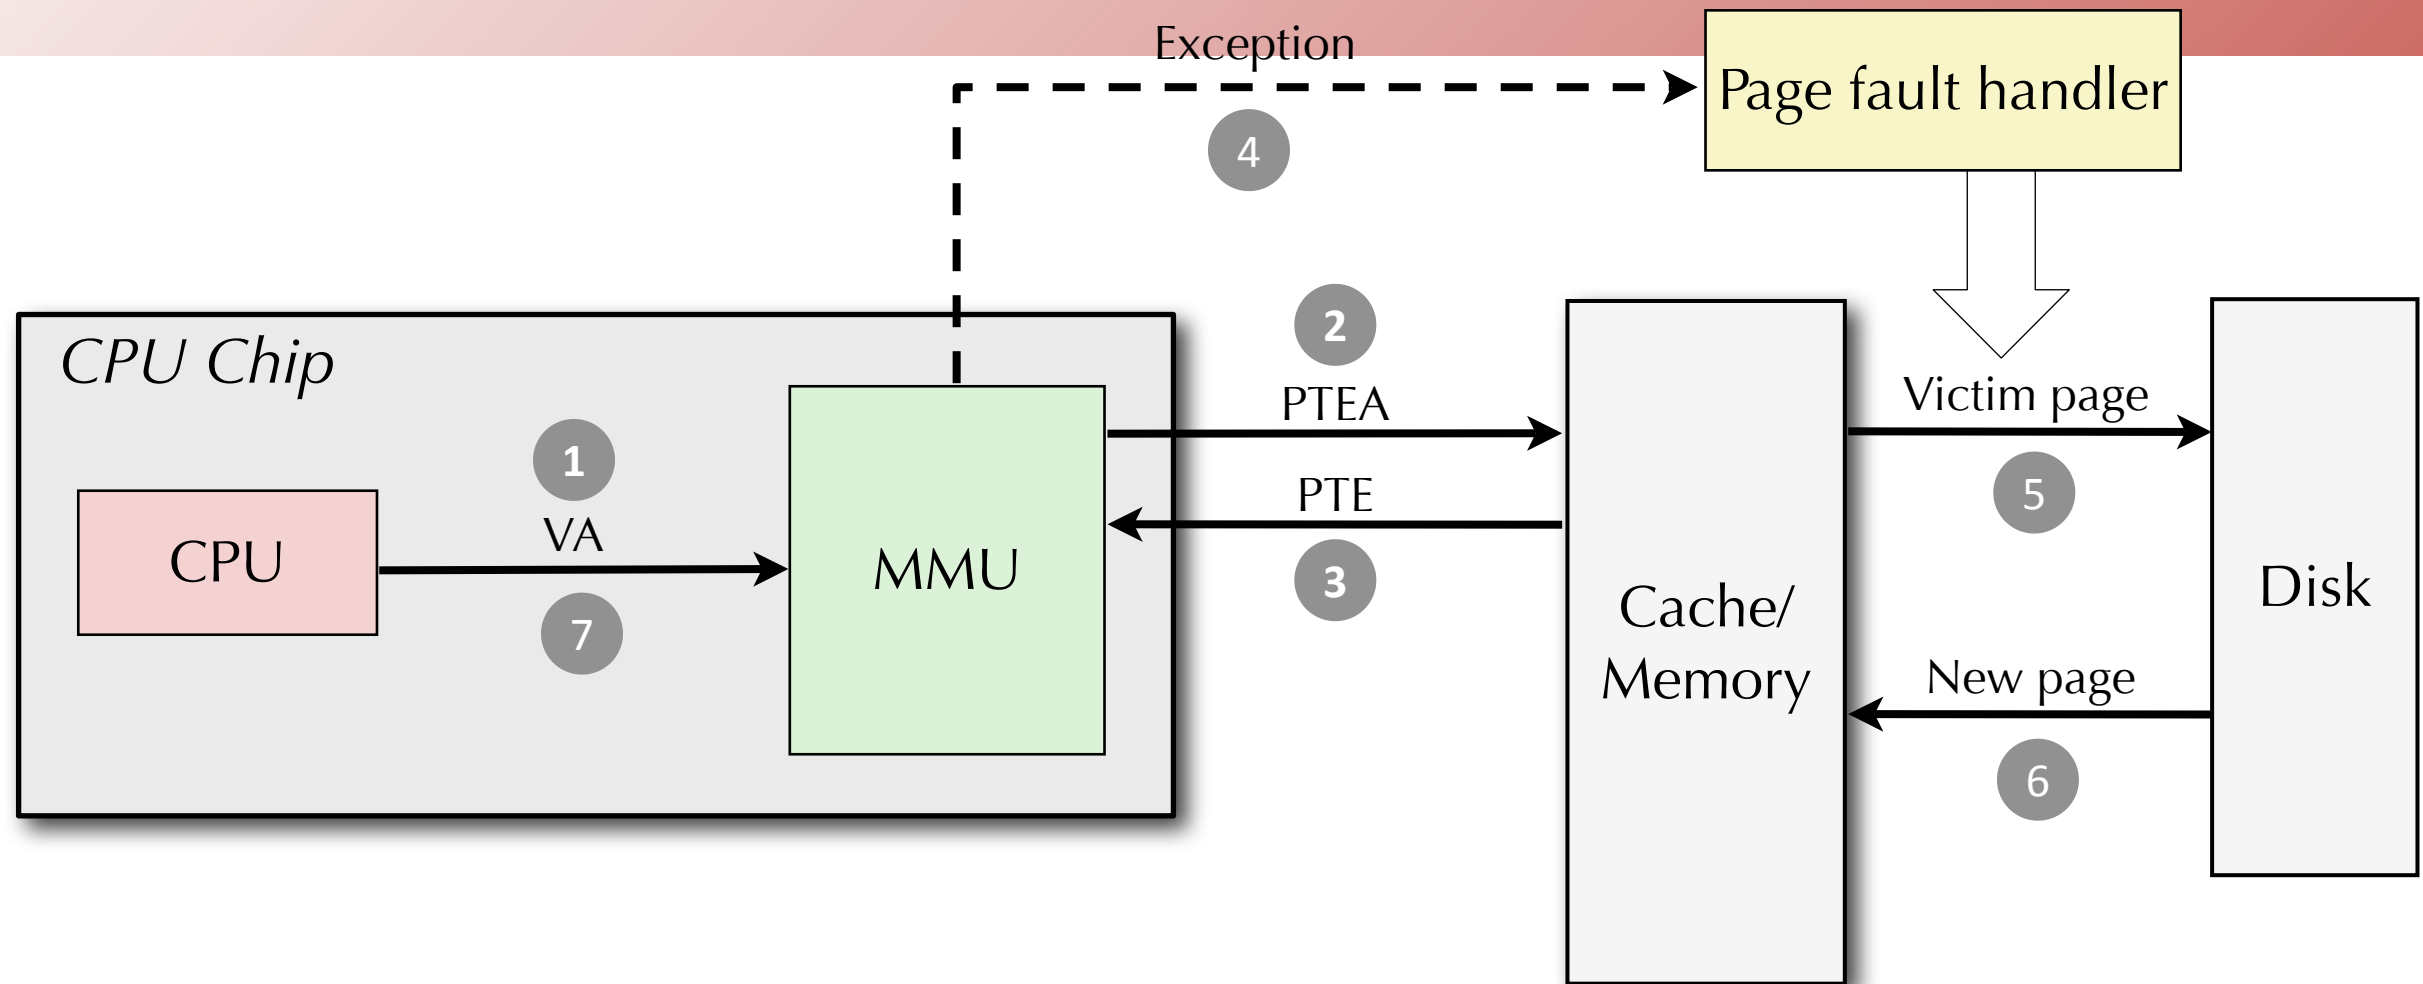
# Address Translation: Page Hit



- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) MMU sends physical address to cache/memory
- 5) Cache/memory sends data word to processor



# Address Translation: Page Fault

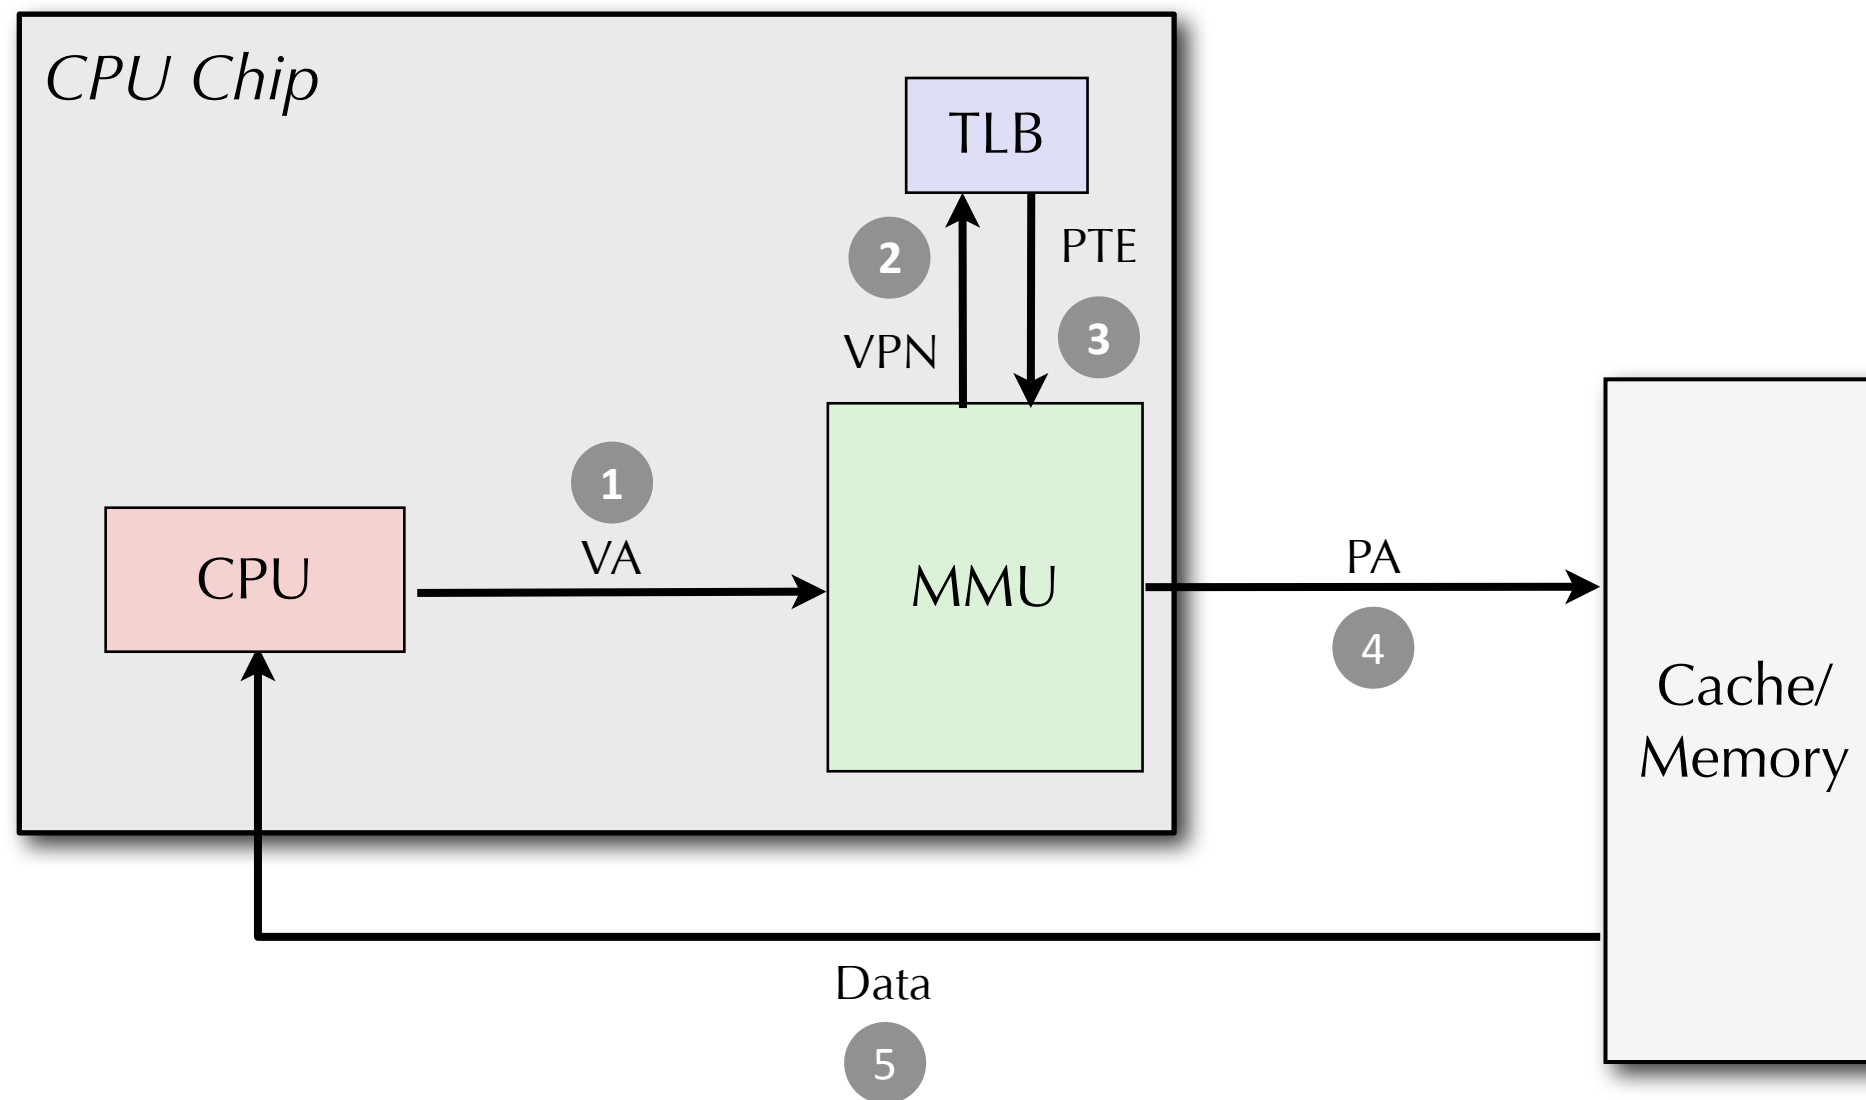


- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) Valid bit is zero, so MMU triggers page fault exception
- 5) Handler identifies victim (and, if dirty, pages it out to disk)
- 6) Handler pages in new page and updates PTE in memory
- 7) Handler returns to original process, restarting faulting instruction

# Speeding up Translation with a TLB

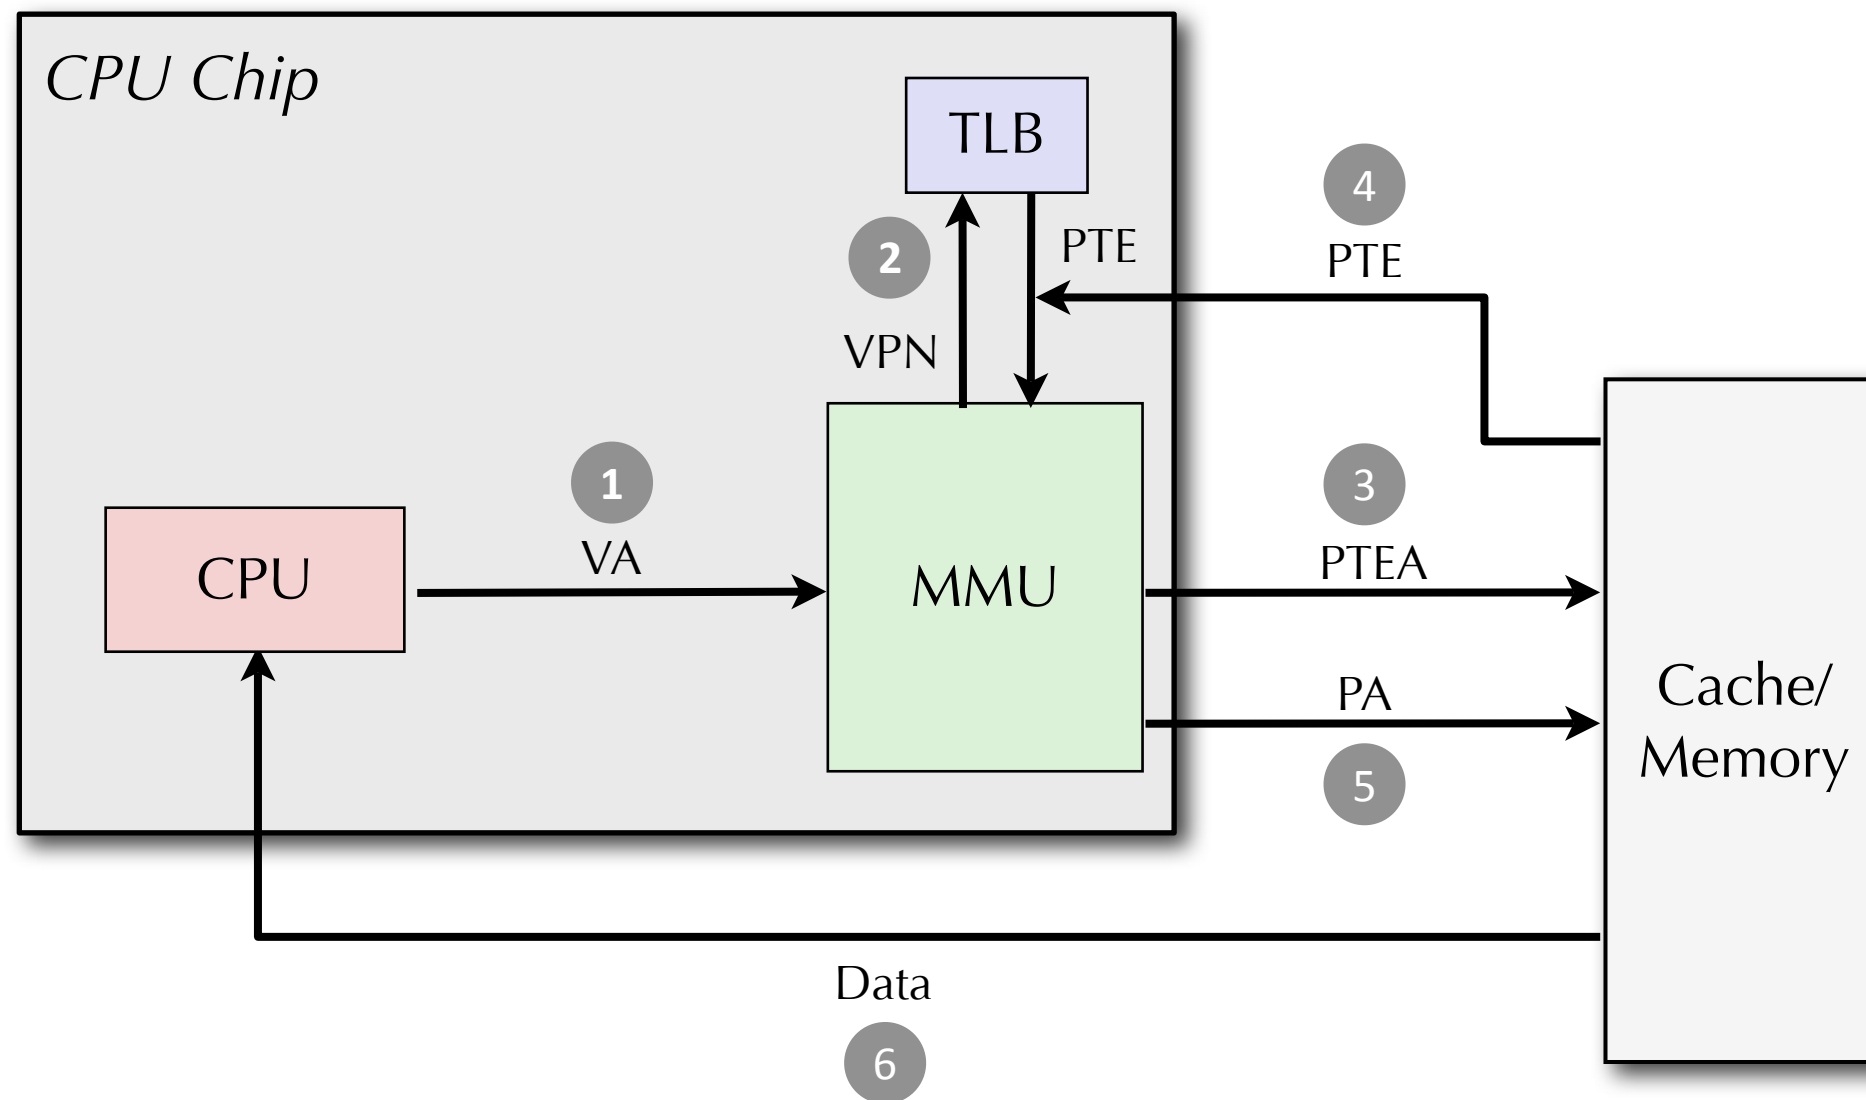
- Page table entries (PTEs) are cached in L1 like any other memory word
  - PTEs may be evicted by other data references
  - PTE hit still requires a 1-cycle delay
- Solution: **Translation Lookaside Buffer (TLB)**
  - Small hardware cache in MMU
  - Maps virtual page numbers to physical page numbers
  - Contains complete page table entries for small number of pages

# TLB Hit



- A TLB hit eliminates a memory access

# TLB Miss




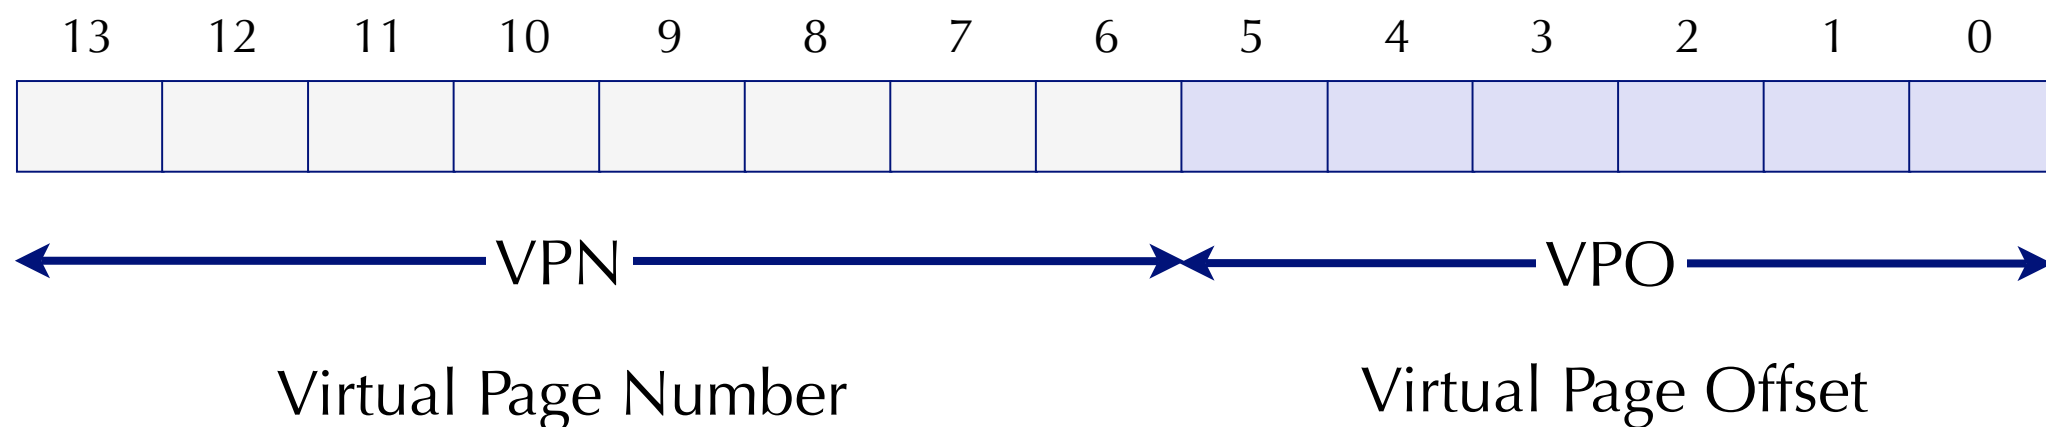
- A TLB miss incurs an additional memory access (the PTE)

# Today

- Running multiple programs at once
- Virtual memory
  - Address spaces
  - Benefits
  - Physical memory as a cache for virtual memory
  - Page tables
  - Page hits and page faults
  - Virtual memory as a tool for memory management
  - Virtual memory as a tool for memory protection
  - Address translation
  - Table Lookaside Buffer (TLB)
  - Example

# Simple Memory System Example

- Addressing
  - 14-bit virtual addresses
  - 12-bit physical address
  - Page size = 64 bytes



# Simple Memory System Page Table

- Only show first 16 entries (out of 256)

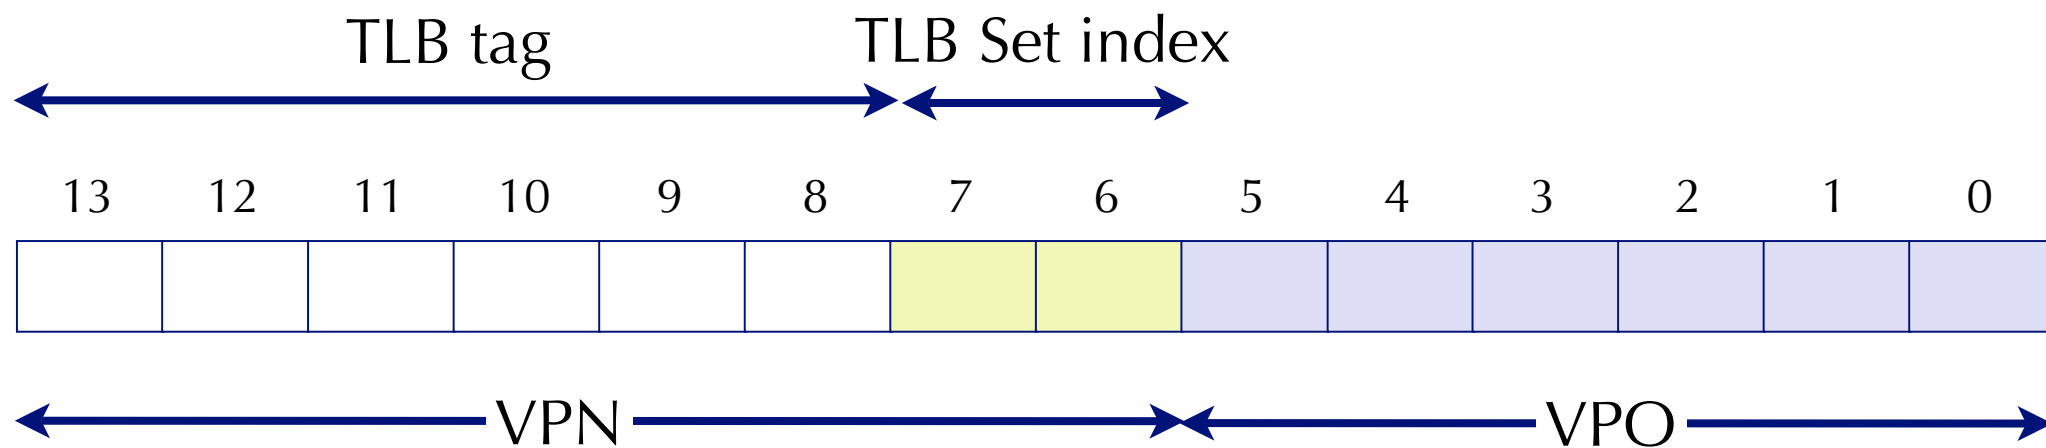
<i>VPN</i>	<i>PPN</i>	<i>Valid</i>
00	28	1
01	—	0
02	33	1
03	02	1
04	—	0
05	16	1
06	—	0
07	—	0

<i>VPN</i>	<i>PPN</i>	<i>Valid</i>
08	13	1
09	17	1
0A	09	1
0B	—	0
0C	—	0
0D	2D	1
0E	11	1
0F	0D	1



# Simple Memory System TLB

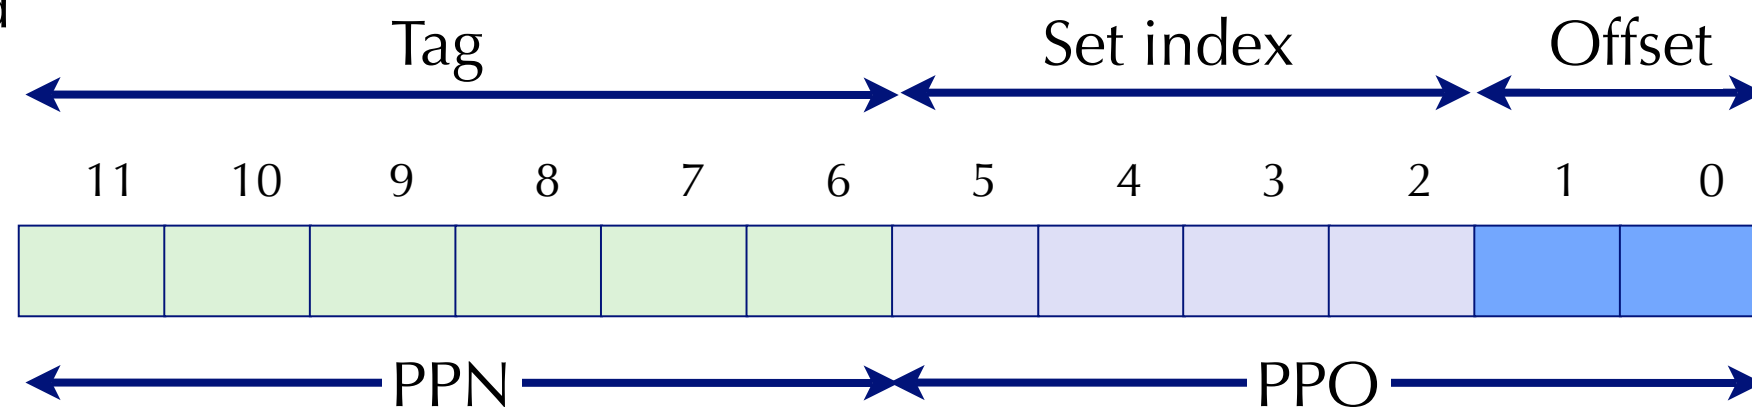
- 16 entries
- 4-way associative



Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	03	—	0	09	0D	1	00	—	0	07	02	1
1	03	2D	1	02	—	0	04	—	0	0A	—	0
2	02	—	0	08	—	0	06	—	0	03	—	0
3	07	—	0	03	0D	1	0A	34	1	02	—	0

# Simple Memory System Cache

- 16 lines, 4-byte block size
- **Physically addressed**
- Direct mapped

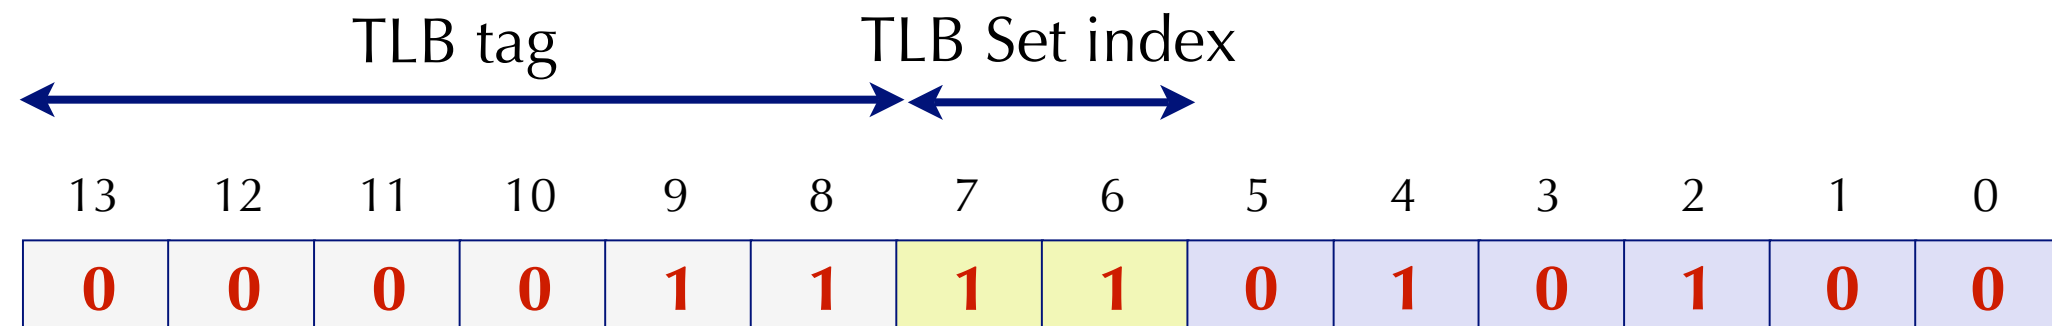


<i>Idx</i>	<i>Tag</i>	<i>Valid</i>	<i>B0</i>	<i>B1</i>	<i>B2</i>	<i>B3</i>
0	19	1	99	11	23	11
1	15	0	—	—	—	—
2	1B	1	00	02	04	08
3	36	0	—	—	—	—
4	32	1	43	6D	8F	09
5	0D	1	36	72	F0	1D
6	31	0	—	—	—	—
7	16	1	11	C2	DF	03

<i>Idx</i>	<i>Tag</i>	<i>Valid</i>	<i>B0</i>	<i>B1</i>	<i>B2</i>	<i>B3</i>
8	24	1	3A	00	51	89
9	2D	0	—	—	—	—
A	2D	1	93	15	DA	3B
B	0B	0	—	—	—	—
C	12	0	—	—	—	—
D	16	1	04	96	34	15
E	13	1	83	77	1B	D3
F	14	0	—	—	—	—

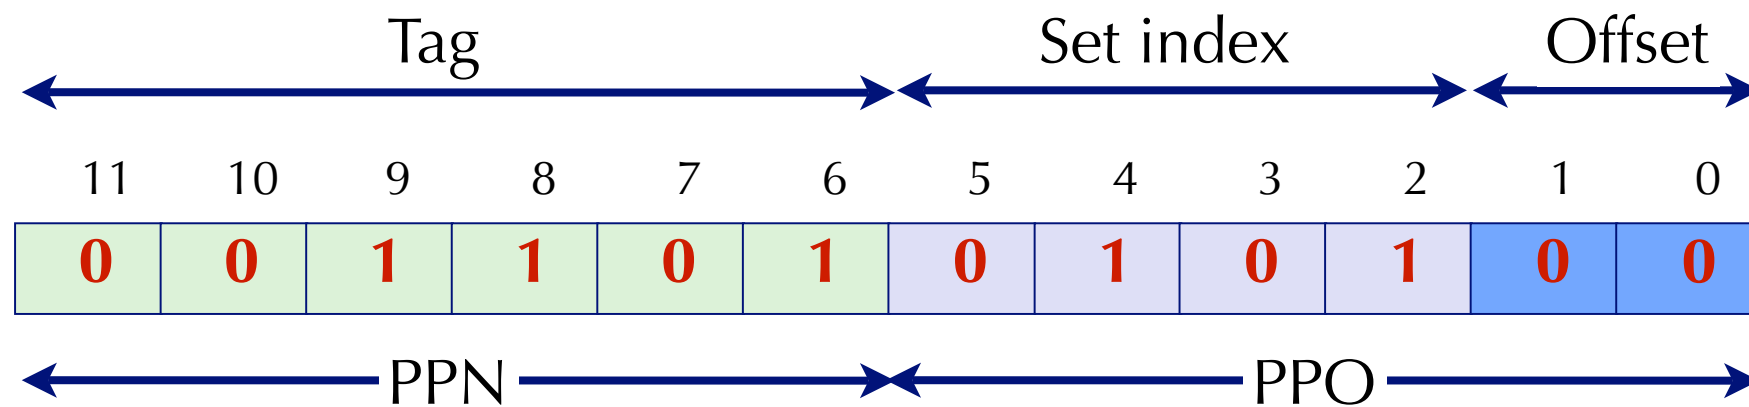
# Address Translation Example #1

Virtual Address: 0x03D4



VPN: 0x0F    TLB Set index: 3    TLB Tag: 0x03    TLB Hit? Y    Page fault? N    PPN: 0x0D

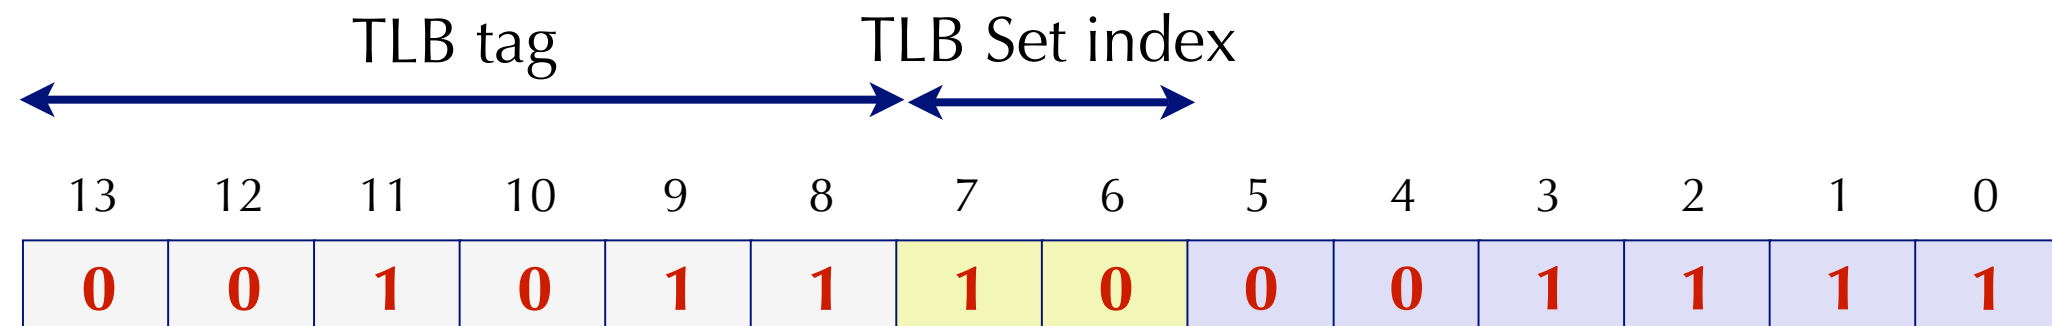
Physical Address



Offset: 0    Set index: 0x5    Tag: 0x0D    Hit? Y    Byte: 0x36

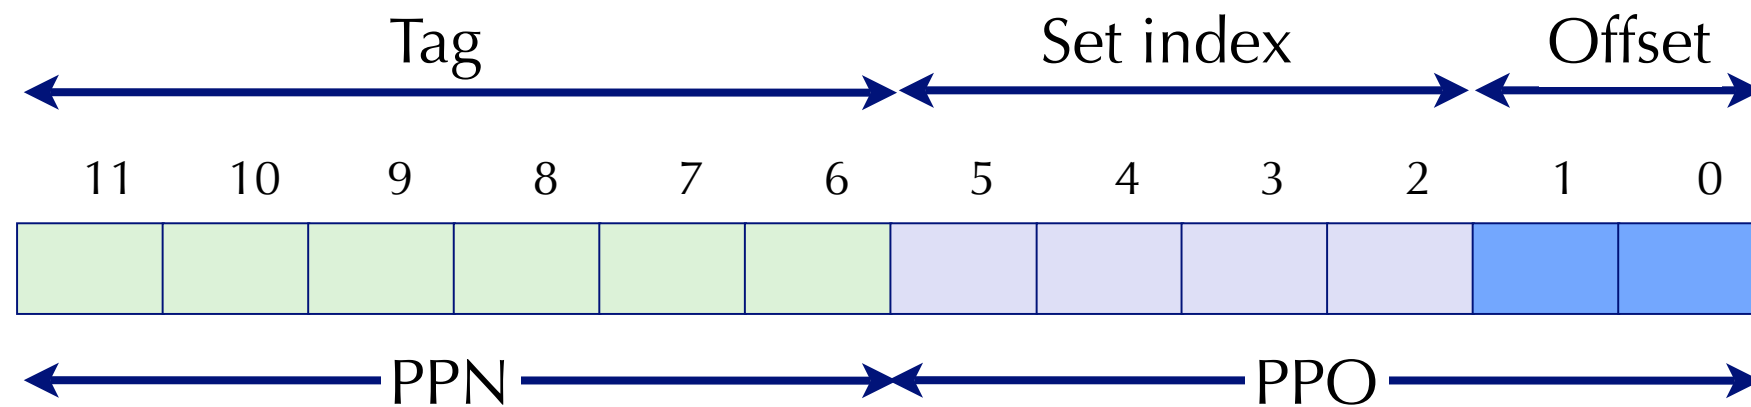
# Address Translation Example #2

Virtual Address: 0x0B8F



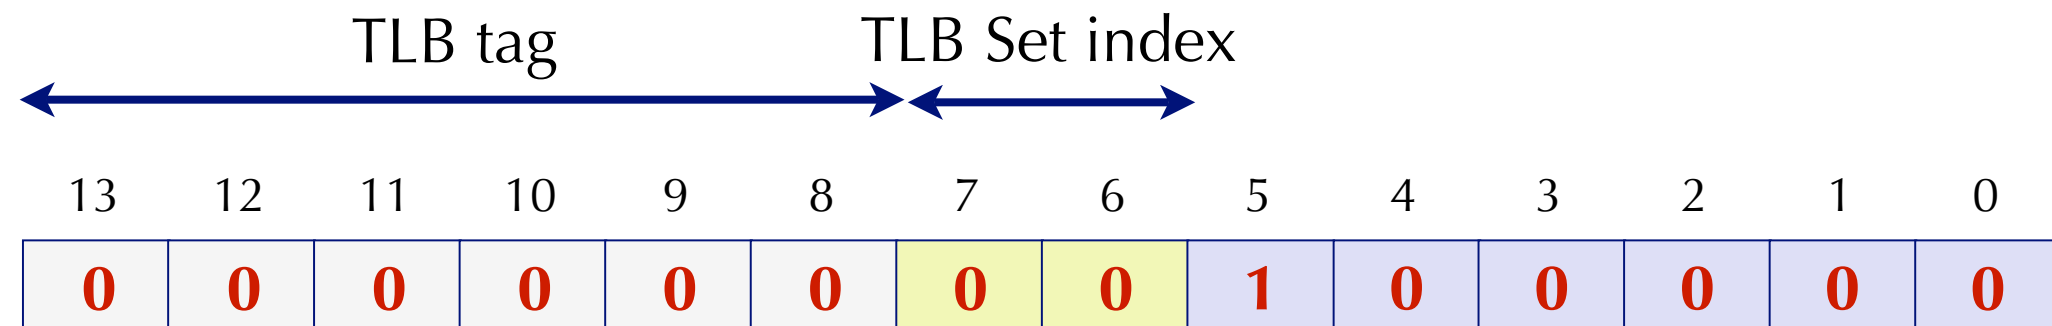
VPN: 0x0E    TLB Set index: 2    TLB Tag: 0x0B    TLB Hit? N    Page fault? Y    PPN: TBD

Physical Address



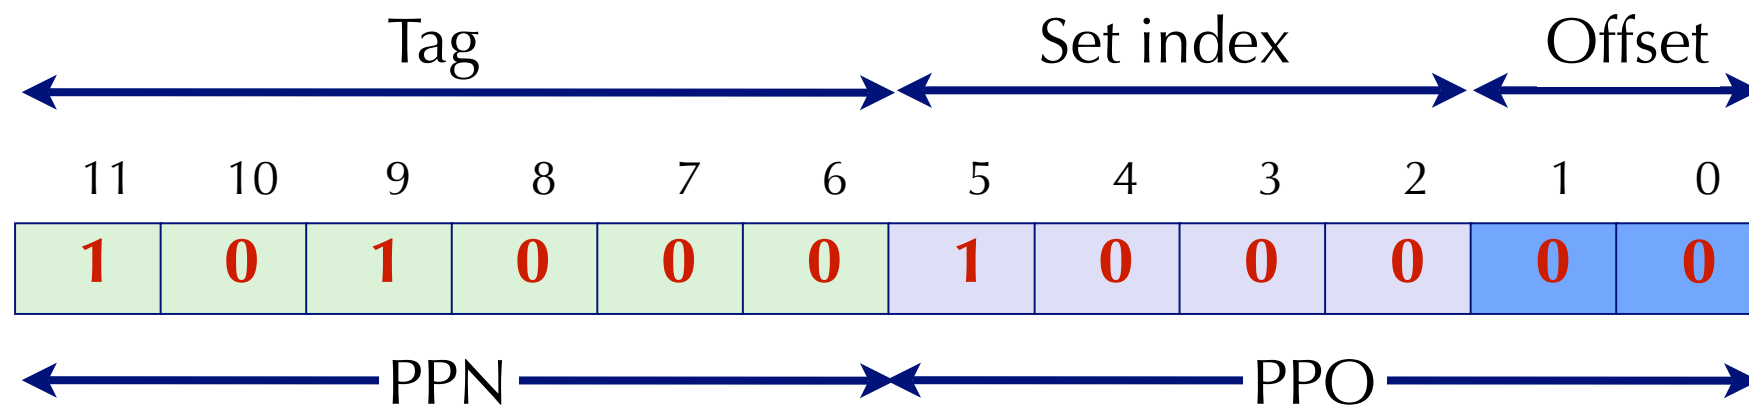
# Address Translation Example #3

Virtual Address: 0x0020



VPN: 0x00    TLB Set index: 0    TLB Tag: 0x00    TLB Hit? N    Page fault? N    PPN: 0x28

Physical Address



Offset: 0    Set index: 0x8    Tag: 0x28    Hit? N    Byte: Mem...

# What happens on a context switch?

- Context switch is fast. Why?
- Only need to change page table base register, and flush the TLB
- Why don't we need to flush L1, L2, L3 cache?
  - CPU caches use *physical addresses*
  - So single physical page with multiple virtual addresses will get reused across processes!

# Summary of VM benefits

- Isolation
  - Each process has its own private linear address space
  - Cannot be corrupted by other processes
- Simplifies memory management and programming
  - Allows multiple concurrent programs to share limited DRAM
- Simplifies protection by providing a convenient method to check permissions
- Efficient
  - Fast translation
    - MMU on CPU, and TLB provides cache of recent translation
  - Fast context switch



# Summary of VM mechanism

- VM is implemented by the MMU and OS working together.
  - MMU does virtual-to-physical address translations.
  - OS sets up the page tables that control those translations.
- The TLB is a cache for recent virtual-to-physical mappings.
  - Avoids lookup in page tables for each memory access.
- The OS handles page faults triggered by the MMU.
  - This is the basis for swapping and demand paging.

# Topics for next time

- Systems programming in UNIX
- File abstraction and I/O operations
- Robust and buffered I/O
- Accessing metadata and directories
- Fun with filehandles