



HARVARD

School of Engineering
and Applied Sciences

UNIX Systems Programming 2: Processes and signals

CS61, Lecture 16

Prof. Stephen Chong

October 26, 2010

Announcements

- In-class **midterm exam**, Thursday 28 Oct
 - Open book, closed note
 - No computers, no smartphones, no internet access
 - Practice exam and quiz solutions on iSite website (accessible only to enrollees)
- Mid-course evaluations
 - Please take the mid-course survey on the course's iSite (linked to from cs61.seas) by end of today.

Topics for today

- The UNIX process abstraction
- Process lifecycle
 - Creating processes: forking
 - Running new programs within a process
 - Terminating and reaping processes
- Signaling processes
- Interprocess communication: pipes

What is a process?

- A **process** is the OS's abstraction for **execution**
 - A process is *an instance of a program in execution*.
 - i.e., each process is running a program; there may be many processes running the same program
- A process provides
 - Private address space
 - Through the mechanism of virtual memory!
 - Illusion of exclusive use of processor

Process context

- Process **context** is the state that the operating system needs to run a process
 - 1) **Address space**
 - The memory that the process can access
 - Consists of various pieces: the program code, static variables, heap, stack, etc.
 - 2) **Processor state**
 - The CPU registers associated with the running process
 - Includes general purpose registers, program counter, stack pointer, etc.
 - 3) **OS resources**
 - Various OS state associated with the process
 - Examples: page table, file table, network sockets, etc.

Context switches

- Multiple processes can run simultaneously.
 - On a single CPU system, only one process is running on the CPU at a time.
 - On a multi-CPU (or multi-core) system, multiple processes really can run concurrently.
 - The OS will timeshare each CPU/core, rapidly switching processes across them all.
- Switching a CPU from running one process to another is called a **context switch**.
 - (1) Save the context of the currently running process,
 - (2) Restore the context of some previously preempted process
 - (3) Resume execution of the newly restored process
- Deciding when to preempt current process and restart previously preempted process is known as **scheduling**
 - Performed by part of the OS called a **scheduler**

Process IDs

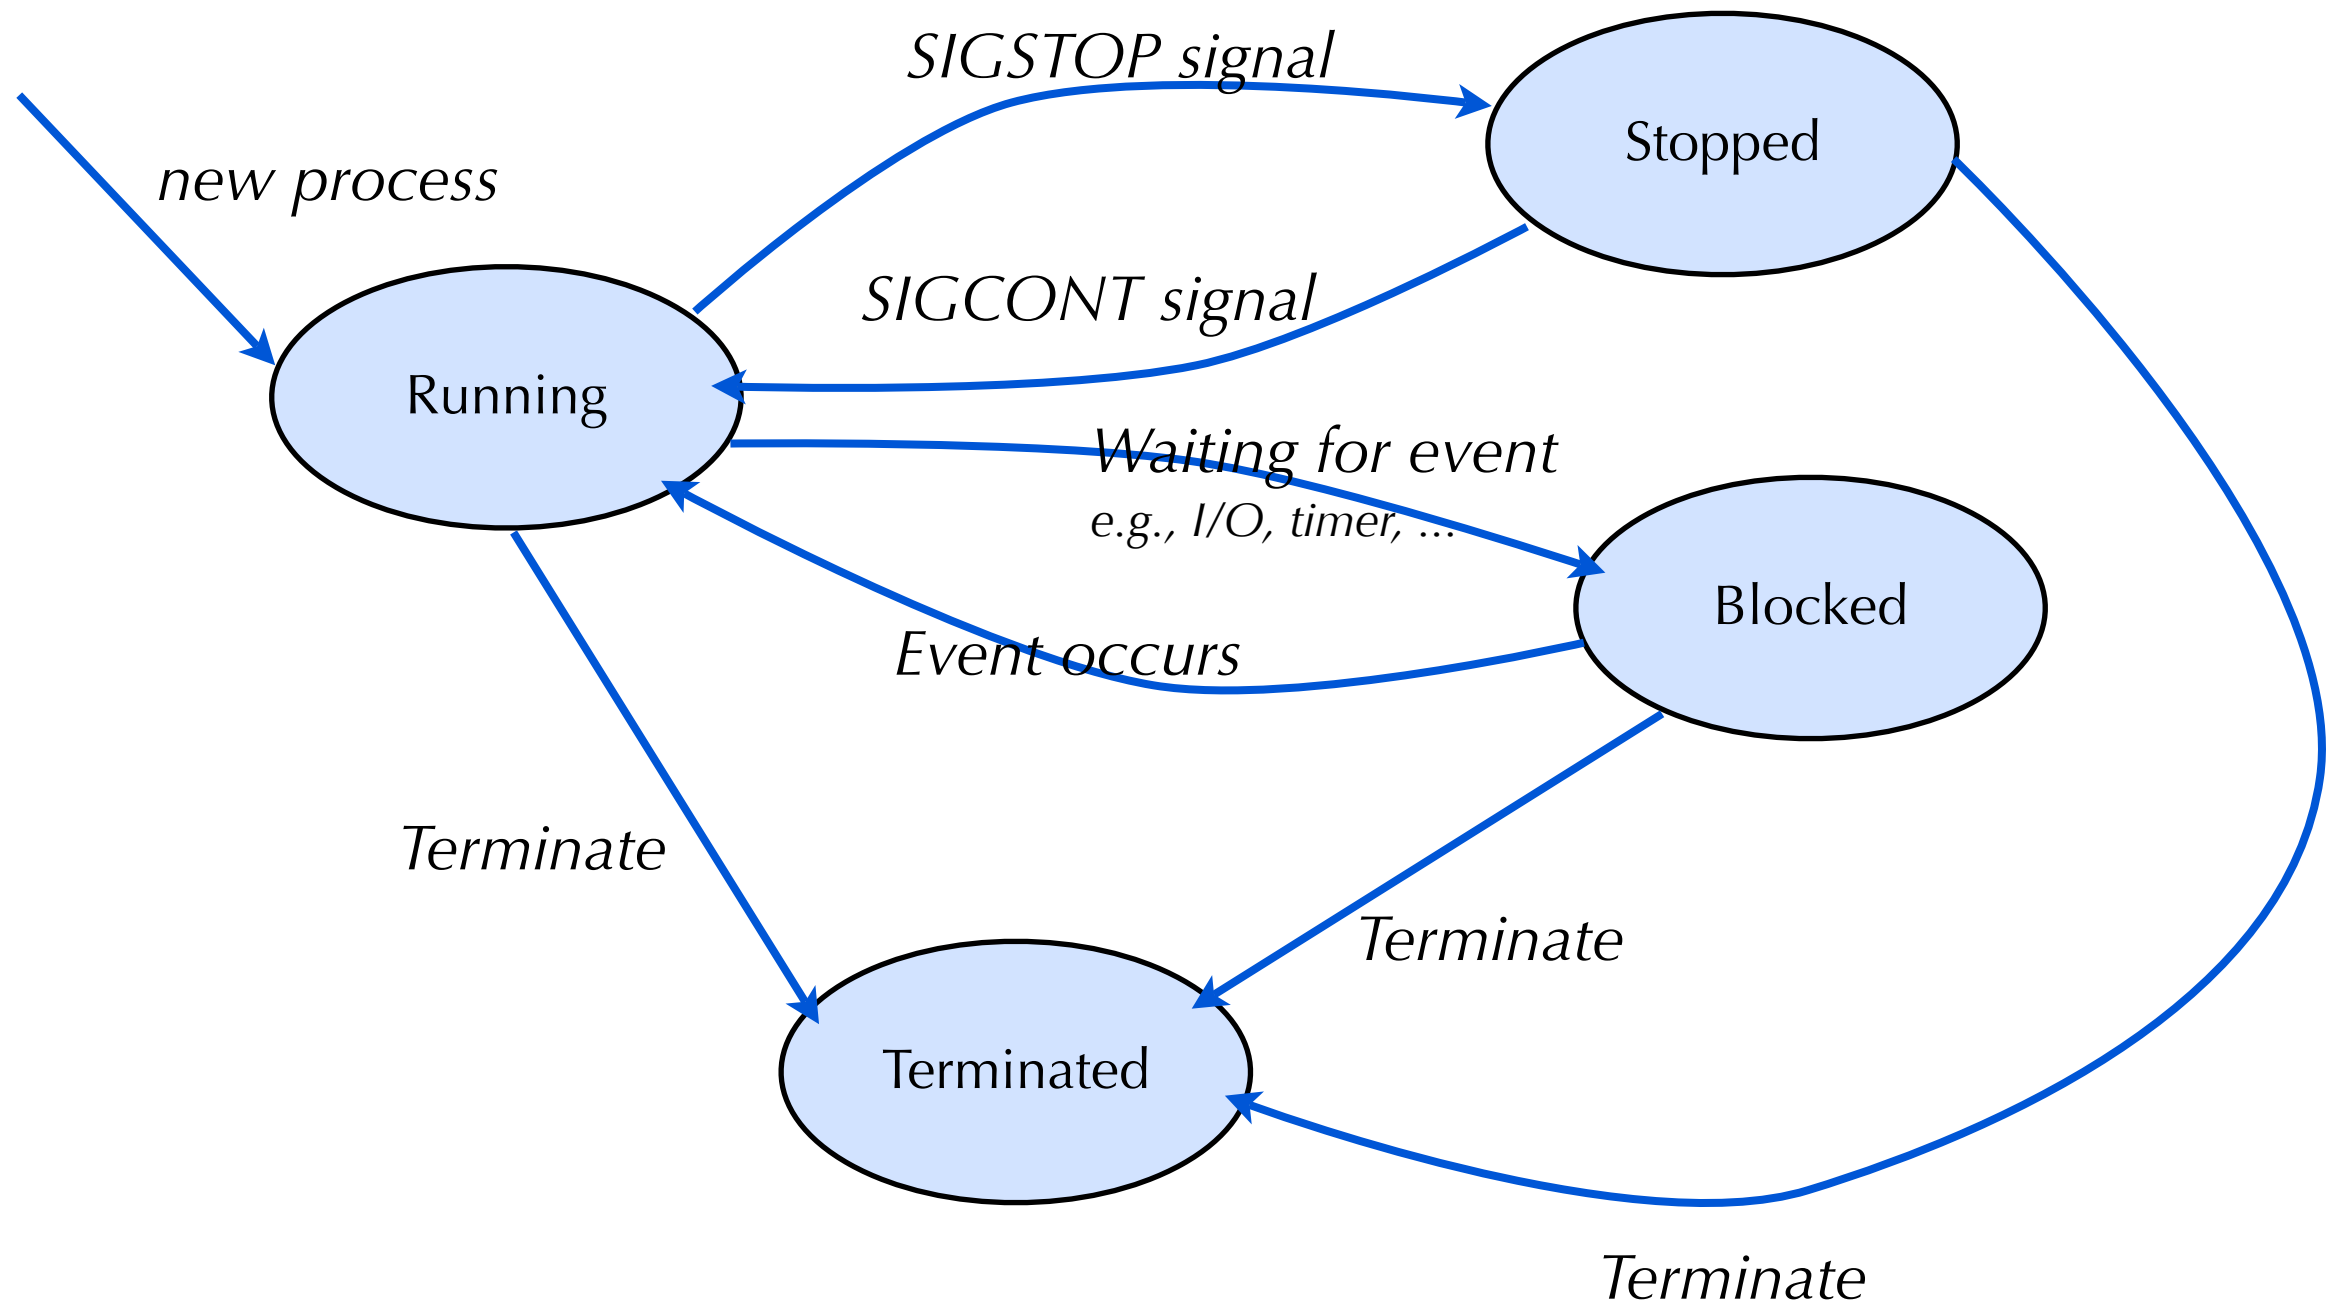
- Each process has a unique positive **process ID (PID)**
- `getpid` returns current process's PID
- `getppid` returns PID of parent of current process

```
pid_t getpid(void);  
pid_t getppid(void);
```


Process states

- At any moment, a process is in one of several states:
 - **Running:**
 - Process is executing on a CPU, or waiting to be executed
 - **Stopped:**
 - Process is *suspended* (due to receiving a certain **signal**) and will not be scheduled
 - More on signals soon...
 - **Waiting (or sleeping or blocked):**
 - Process is waiting for an event to occur, such as completion of I/O, timer, etc.
 - Why is this different than “ready” ?
 - **Terminated:**
 - Process is stopped permanently, e.g., by returning from `main`, or by calling `exit`
- As the process executes, it moves between these states
 - What state is the process in most of the time?

Process lifecycle



Topics for today

- The UNIX process abstraction
- Process lifecycle
 - Creating processes: forking
 - Running new programs within a process
 - Terminating and reaping processes
- Signaling processes
- Interprocess communication: pipes

How are processes created?

- Typically, new process is created when user runs a program
 - E.g., Double-click an application, or type a command at the shell
- In UNIX, starting a new program is done by **some other process**
 - The shell is a process itself!
 - So are the Dock and Finder in MacOS (a variant of UNIX)
- One process (e.g., the shell) is creating another process (the command you want to run)
 - This is called **forking**
 - Every process has a **parent process**
- Chicken-and-egg problem: How does first process get created?
 - At boot time, the OS creates the first process, called `init`, which is responsible for starting up many other processes

fork: Creating New Processes

● `int fork(void)`

- creates a new process (**child process**) that is **identical** to the calling process (**parent process**)
- returns 0 to the child process
- returns child's process ID (pid) to the parent process

```
if (fork() == 0) {  
    printf("hello from child\n");  
} else {  
    printf("hello from parent\n");  
}
```

- Fork is interesting (and often confusing) because it is called once but returns twice

Fork Example #1

- Parent and child process both run the same program.
 - Only difference is the return value from `fork()`
- Child's address space starts as an **exact copy** of parent's
 - They do not share the memory – instead they each have a private copy.
 - Also have the same open files with the same offsets into the files.
 - Includes `stdin`, `stdout`, and `stderr`


```
void fork1()
{
    int x = 1;
    pid_t pid = fork();
    if (pid == 0) {
        printf("Child has x = %d\n", ++x);
    } else {
        printf("Parent has x = %d\n", --x);
    }
    printf("Bye from process %d with x = %d\n", getpid(), x);
}
```

```
Parent has x = 0
Bye from process 9991 with x = 0
Child has x = 2
Bye from process 9992 with x = 2
```

Fork Example #2

- Key Points
 - Both parent and child can continue forking

```
void fork2()
{
    printf("L0\n");
    fork();
    printf("L1\n");
    fork();
    printf("Bye\n");
}
```




Fork Example #3

- Key Points

- Both parent and child can continue forking


```
void fork3()
{
    printf("L0\n");
    fork();
    printf("L1\n");
    fork();
    printf("L2\n");
    fork();
    printf("Bye\n");
}
```



Fork Example #4

- Key Points
 - Both parent and child can continue forking


```
void fork4()
{
    printf("L0\n");
    if (fork() != 0) {
        printf("L1\n");
        if (fork() != 0) {
            printf("L2\n");
            fork();
        }
    }
    printf("Bye\n");
}
```



Fork Example #5

- Key Points
 - Both parent and child can continue forking

```
void fork5()
{
    printf("L0\n");
    if (fork() == 0) {
        printf("L1\n");
        if (fork() == 0) {
            printf("L2\n");
            fork();
        }
    }
    printf("Bye\n");
}
```



Starting new programs

- How do we start a new program, instead of copying the parent?
- Use the UNIX `execve()` system call
- ```
int execve(const char *filename,
 char *const argv[],
 char *const envp[]);
```

  - `filename`: name of executable file to run
  - `argv`: Command line arguments
  - `envp`: environment variable settings (e.g., `$PATH`, `$HOME`, etc.)

# Starting new programs

- `execve()` does not fork a new process!
  - Rather, it replaces the address space and CPU state of the current process
  - Loads the new address space from the executable file and starts it from `main()`
  - So, to start a new program, use `fork()` followed by `execve()`

# Using fork and exec

```
int main(int argc, char **argv) {
 int rv;
 if (fork() == 0) { /* Child process */
 char *newargs[3];
 printf("Hello, I am the child process.\n");
 newargs[0] = "/bin/echo"; /* Convention! Not required!! */
 newargs[1] = "some random string";
 newargs[2] = NULL; /* Indicate end of args array */
 if (execv("/bin/echo", newargs)) {
 printf("warning: execve returned an error.\n");
 exit(-1);
 }
 printf("Child process should never get here\n");
 exit(42);
 }
}
```

# Topics for today

- The UNIX process abstraction
- Process lifecycle
  - Creating processes: forking
  - Running new programs within a process
  - Terminating and reaping processes
- Signaling processes
- Interprocess communication: pipes

# Terminating a process

- A process terminates for one of 3 reasons:
  - (1) return from the `main()` procedure
  - (2) call to the `exit()` function
  - (3) receive a signal whose default action is to terminate



# exit: Destroying Process


- `void exit(int exit_status)`
  - Exits a process with specified exit status.
  - By convention, status of 0 is a “normal” exit, non-zero indicates an error of some kind.
- `atexit()` registers functions to be executed upon exit.

```
void cleanup(void) {
 printf("cleaning up\n");
}

void fork6() {
 atexit(cleanup);
 fork();
 exit(0);
}
```

# Zombies

- When a process terminates (for whatever reason) OS does not remove it from system immediately
- Process stays until it is **reaped** by parent
  - When parent reaps a child process, OS gives the parent the exit status of child, and cleans up child
  - A terminated process that has not been reaped is called a **zombie process**
- How do you reap a child process?




# wait: Synchronizing with Children

- `int wait(int *child_status)`
  - Suspends parent process until one of its children terminates
  - Return value is the pid of the child process that terminated
  - if `child_status != NULL`, it will point to the child's return status
- `child_status` can be accessed using several macros:
  - `WIFEXITED(child_status) == 1` if child exited due to call to `exit()`
  - `WEXITSTATUS(child_status)` gives the return code passed to `exit()`
  - `WCOREDUMP(child_status) == 1` if child dumped core.
  - And others (see “man 2 wait”)

# wait: Synchronizing with Children

```
void fork9() {
 int child_status;

 if (fork() == 0) {
 printf("HC: hello from child\n");
 }
 else {
 printf("HP: hello from parent\n");
 wait(&child_status);
 printf("CT: child has terminated\n");
 }
 printf("Bye\n");
 exit();
}
```



# What if multiple child processes exit?

- `wait()` returns status of exited children in arbitrary order.

```
void fork10()
{
 pid_t pid[N];
 int i;
 int child_status;
 for (i = 0; i < N; i++)
 if ((pid[i] = fork()) == 0)
 exit(100+i); /* Child */
 for (i = 0; i < N; i++) {
 pid_t wpid = wait(&child_status);
 if (WIFEXITED(child_status))
 printf("Child %d terminated with exit status %d\n",
 wpid, WEXITSTATUS(child_status));
 else
 printf("Child %d terminate abnormally\n", wpid);
 }
}
```



# waitpid(): Waiting for a specific process

- `pid_t waitpid(pid_t child_pid, int *status, int options)`
  - Causes parent to wait for a **specific** child process to exit.
- Most general form of wait
  - `child_pid > 0`: wait for specific child to exit
  - `child_pid = -1`: wait for any child to exit
  - return value is PID of child process
  - `options` can be used to specify if call should return immediately (with return value of 0) if no terminated children, and also whether we are interested in stopped processes
  - `status` encodes information about how child exited (or was stopped)

# waitpid(): Waiting for a specific process

- `pid_t waitpid(pid_t child_pid,  
int *status,  
int options)`

```
void fork11()
{
 pid_t pid[N];
 int i;
 int child_status;
 for (i = 0; i < N; i++)
 if ((pid[i] = fork()) == 0)
 exit(100+i); /* Child */
 for (i = 0; i < N; i++) {
 pid_t wpid = waitpid(pid[i], &child_status, 0);
 if (WIFEXITED(child_status))
 printf("Child %d terminated with exit status %d\n",
 wpid, WEXITSTATUS(child_status));
 else
 printf("Child %d terminated abnormally\n", wpid);
 }
}
```



# Back to the zombies...

- Zombie example

```
void zombie()
{
 if (fork() == 0) {
 /* Child */
 printf("Terminating Child, PID = %d\n",
 getpid());
 exit(0);
 } else {
 printf("Running Parent, PID = %d\n",
 getpid());
 while (1)
 ; /* Infinite loop */
 }
}
```

```
linux> ./zombie &
[1] 6639
Running Parent, PID = 6639
Terminating Child, PID = 6640
linux> ps
 PID TTY TIME CMD
 6585 ttyp9 00:00:00 tcsh
 6639 ttyp9 00:00:03 zombie
 6640 ttyp9 00:00:00 zombie <defunct>
 6641 ttyp9 00:00:00 ps
linux>
```

- ps shows child process as “defunct”

# Orphans

- So bad things happen if the parent does not wait for the child...
- If the child exits first, child becomes a zombie
- If the parent exits first, the child becomes an **orphan**.
  - Problem: All processes (except for `init`) need a parent process.
  - Orphan processes “adopted” by `init` (PID 1 on most UNIX systems)
  - If child subsequently terminates, it will be reaped by `init`
    - `init` reaps zombie orphans...

# Nonterminating Child Example

```
void fork8()
{
 if (fork() == 0) {
 /* Child */
 printf("Running Child, PID = %d\n",
 getpid());
 while (1)
 ; /* Infinite loop */
 } else {
 printf("Terminating Parent, PID = %d\n",
 getpid());
 exit(0);
 }
}
```

```
linux> ./forks 8
Terminating Parent, PID = 6675
Running Child, PID = 6676
linux> ps
 PID TTY TIME CMD
 6585 ttys9 00:00:00 tcsh
 6676 ttys9 00:00:06 forks
 6677 ttys9 00:00:00 ps
linux> kill 6676
linux> ps
 PID TTY TIME CMD
 6585 ttys9 00:00:00 tcsh
 6678 ttys9 00:00:00 ps
```

- Child process still active even though parent has terminated
- Must kill explicitly, or else will keep running indefinitely

# Zombie orphan

- Zombie example

```
void zombie()
{
 if (fork() == 0) {
 /* Child */
 printf("Terminating Child, PID = %d\n",
 getpid());
 exit(0);
 } else {
 printf("Running Parent, PID = %d\n",
 getpid());
 while (1)
 ; /* Infinite loop */
 }
}
```

```
linux> ./zombie &
[1] 6639
Running Parent, PID = 6639
Terminating Child, PID = 6640
linux> ps
 PID TTY TIME CMD
 6585 ttyp9 00:00:00 tcsh
 6639 ttyp9 00:00:03 zombie
 6640 ttyp9 00:00:00 zombie <defunct>
 6641 ttyp9 00:00:00 ps
linux> kill 6639
[1] Terminated
linux> ps
 PID TTY TIME CMD
 6585 ttyp9 00:00:00 tcsh
 6642 ttyp9 00:00:00 ps
```

- ps shows child process as “defunct”
- Killing parent allows child to be reaped

# Topics for today

- The UNIX process abstraction
- Process lifecycle
  - Creating processes: forking
  - Running new programs within a process
  - Terminating and reaping processes
- Signaling processes
- Interprocess communication: pipes

# Signals

- Unix provides a mechanism to allow processes and OS to interrupt other processes
- A signal is small message to notify a process of some system event
  - These messages not normally visible to program
  - e.g.,

| ID | Name    | Default Action   | Corresponding Event                      |
|----|---------|------------------|------------------------------------------|
| 2  | SIGINT  | Terminate        | Interrupt (e.g., ctl-c from keyboard)    |
| 9  | SIGKILL | Terminate        | Kill program (cannot override or ignore) |
| 11 | SIGSEGV | Terminate & Dump | Segmentation violation                   |
| 14 | SIGALRM | Terminate        | Timer signal                             |
| 17 | SIGCHLD | Ignore           | Child stopped or terminated              |



# Signal concepts

- Two distinct steps to transfer a signal:
  - (1) OS **sends** (**delivers**) signal to destination process
    - either because of some system event, or because explicitly requested via `kill` function
  - (2) Process **receives** signal (i.e., forced by OS to react to signal in some way)
    - Process can react in one of three ways:
      - ▶ ignore signal (i.e., do nothing)
      - ▶ terminate (maybe dumping core)
      - ▶ **catch** a signal with a **signal handler** function



# Signal concepts

- Signal sent but not yet received is **pending**
  - At most one signal of each type is pending
  - Signals are not queued!
    - If process has pending signal of type  $k$ , then subsequent signals of type  $k$  are discarded
- Process can **block** receipt of certain signals.
  - Blocked signals will be delivered but not received until process unblocks
- Any signal received at most once

# Pending and blocking signals

- OS maintains **pending** and **blocked** bit vectors for each process
  - **pending** represents set of pending signals
    - OS sets bit  $k$  of pending when signal of type  $k$  is delivered
    - OS clears bit  $k$  of pending when signal of type  $k$  is received
  - **blocked** represents set of signals process has blocked
    - Can be set and cleared using `sigprocmask` function
- For a process, OS computes  $\text{pnb} = \text{pending} \ \& \ \sim\text{blocked}$ 
  - If  $\text{pnb} == 0$  then no signals to be received
  - If  $\text{pnb} \neq 0$  then OS chooses a signal to be received, and triggers some action by process

# Sending signals with kill

- kill programs sends an arbitrary signal to a process
  - E.g., `kill -9 24818` sends SIGKILL to process 24818
- Also a function: `kill(pid_t p, int signal)`
- Can send a signal to a specific process, or all processes in a **process group**
  - Every process belongs to a process group
  - Read textbook for more info

# Default actions

- Each signal type has a predefined **default action**
- One of
  - The process terminates
  - The process terminates and dumps core
  - The process stops (until restarted by a SIGCONT signal)
  - The process ignores the action

# Signal handlers

- `signal(int signum, handler_t *handler)`
  - Overrides default action for signals of kind `signum`
- Different values for `handler`
  - `SIG_IGN`: ignore signals of type `signum`
  - `SIG_DFL`: revert to the default action for signals of type `signum`
  - Otherwise, it is a function pointer for a **signal handler**
    - Function will be called on receipt of signal of type `signum`
    - Referred to as **installing** handler
    - Handler execution is called **handling** or **catching** signal
    - When handler returns, control flow of interrupted process continues

# Signal handler example

```
void int_handler(int sig) {
 printf("Process %d received signal %d\n",
 getpid(), sig);
}

int main() {
 signal(SIGINT, int_handler);
 while (1)
 ;
}
```

```
$./signaleg
```

```
^C
```

```
Process 319 received signal 2
```

```
^C
```

```
Process 319 received signal 2
```

```
^C
```

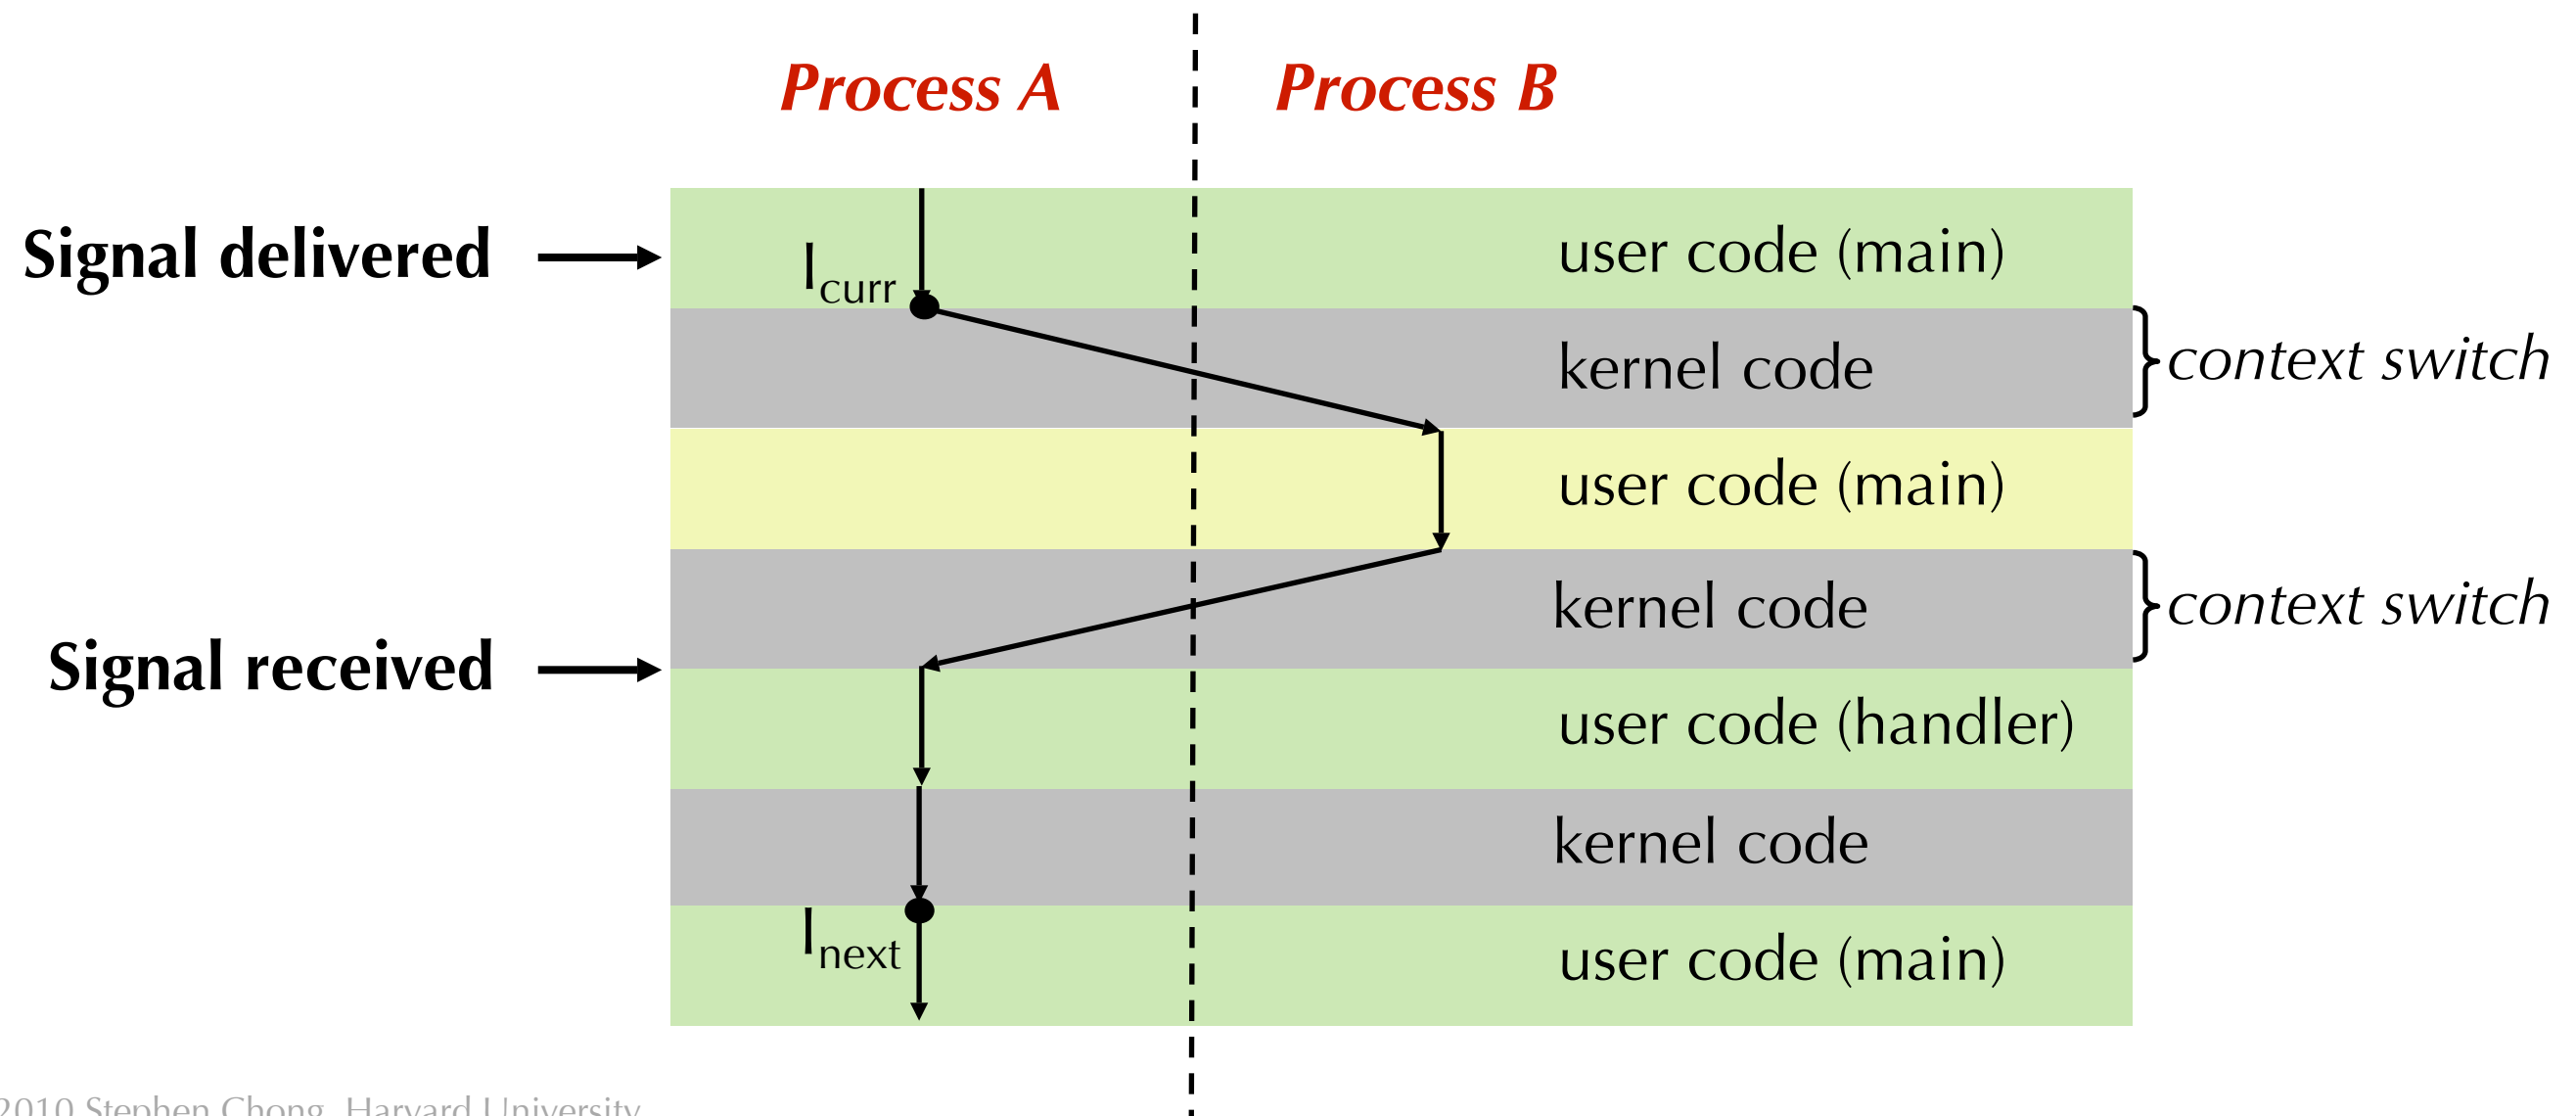
```
Process 319 received signal 2
```

```
Killed
```

```
$ kill -9 319
```

# Signal handlers as concurrent flows

- Signal handlers run **concurrently** with main program
  - Signal handler is not a separate process
  - Concurrent here means “non-sequential”, as opposed to “parallel”






# Topics for today

- The UNIX process abstraction
- Process lifecycle
  - Creating processes: forking
  - Running new programs within a process
  - Terminating and reaping processes
- Signaling processes
- Interprocess communication: pipes

# Communication between processes

- UNIX provides several mechanisms for inter-process communication (IPC)
  - Shared memory regions
  - Sockets (also used for communication over a network).
  - Pipes
- A **pipe** is a pair of file descriptors for communication between two processes.
  - One process can write data to one “end” of the pipe
  - The other process can read data from the other “end” of the pipe.
- `int pipe(int filedes[2]);`
  - `filedes[1]` is the write end of the pipe; `filedes[0]` is the read end of the pipe.



# Using pipes

- But how do we get two processes to use a pipe?
- Idea: Parent process first creates the pipe, then forks the child
  - Since parent and child share open files, they can communicate.
- This is exactly what the UNIX shell does for you when you “pipe” the output of one command into another.

```
$./myprog | grep 'somestring'
```

- Shell creates the pipe, forks both “myprog” and “grep”, and uses `dup2()` to wire the ends of the pipe into `stdout` and `stdin` of each process.
- Somewhat more complex example:

```
$./myprog | grep 'somestring' | sort | uniq | more
```

# Pipe example

```
main() {
 char inbuf[BUFSIZE];
 int p[2], j, pid;

 /* open pipe */
 if(pipe(p) == -1)
 {
 perror("pipe call error");
 exit(1);
 }

 switch(pid = fork()){
 case -1: perror("error: fork failed");
 exit(2);

 case 0: /* if child then write down pipe */
 close(p[0]); /* first close the read end of the pipe */
 write(p[1], "Hello there.", 12);
 write(p[1], "This is a message.", 18);
 write(p[1], "How are you?", 12);
 break;
 default: /* parent reads pipe */
 close(p[1]); /* first close the write end of the pipe */
 read(p[0], inbuf, BUFSIZE); /* What is wrong here?? */
 printf("Parent read: %s\n", inbuf);
 wait(NULL);
 }
 exit(0);
}
```

# Next lecture

- Midterm exam!
- (And then network programming)