



HARVARD

**School of Engineering
and Applied Sciences**

Topics in Systems

CS61, Lecture 24

Hanspeter Pfister, Daniel Margo, Stephen Chong

December 2, 2010

Announcements

- Lab 5 due today
- Final exam Thursday Dec 16, 2pm-5pm
 - Regular office hours will continue until Dec 15, unless otherwise posted
- Q now open
 - Please give feedback!
 - You will receive email regarding course evaluations, or go to my.harvard.

Today and Thursday

- How does what we've learnt so far relate to:
 - The world around us, including technologies we use every day?
 - Current trends in computer systems technology and research?
- Next two lectures will look at a variety of topics
 - Moore's Law
 - Internet-scale
 - Mobile and embedded systems
 - GPUs
 - Operating systems
 - Programming languages

Operating Systems

(CS161)

What is an Operating System?

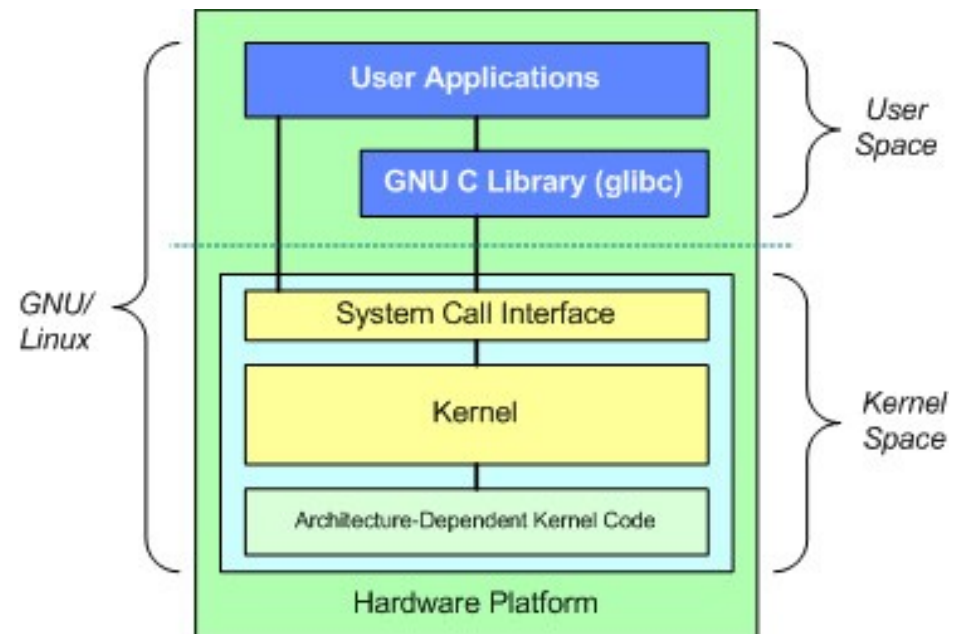
- The code between processes and hardware.
- It provides abstractions and allocates resources
 - CPU (multi-threading)
 - Memory (virtual memory)
 - Hard Disk (filesystem)
- The OS is an illusionist.

What is an Operating System?

- The code between processes and hardware.
- It provides abstractions and allocates resources
 - CPU (multi-threading)
 - Memory (virtual memory)
 - Hard Disk (filesystem)
- The OS is an illusionist.
 - Which is to say, it is a stupendous liar.

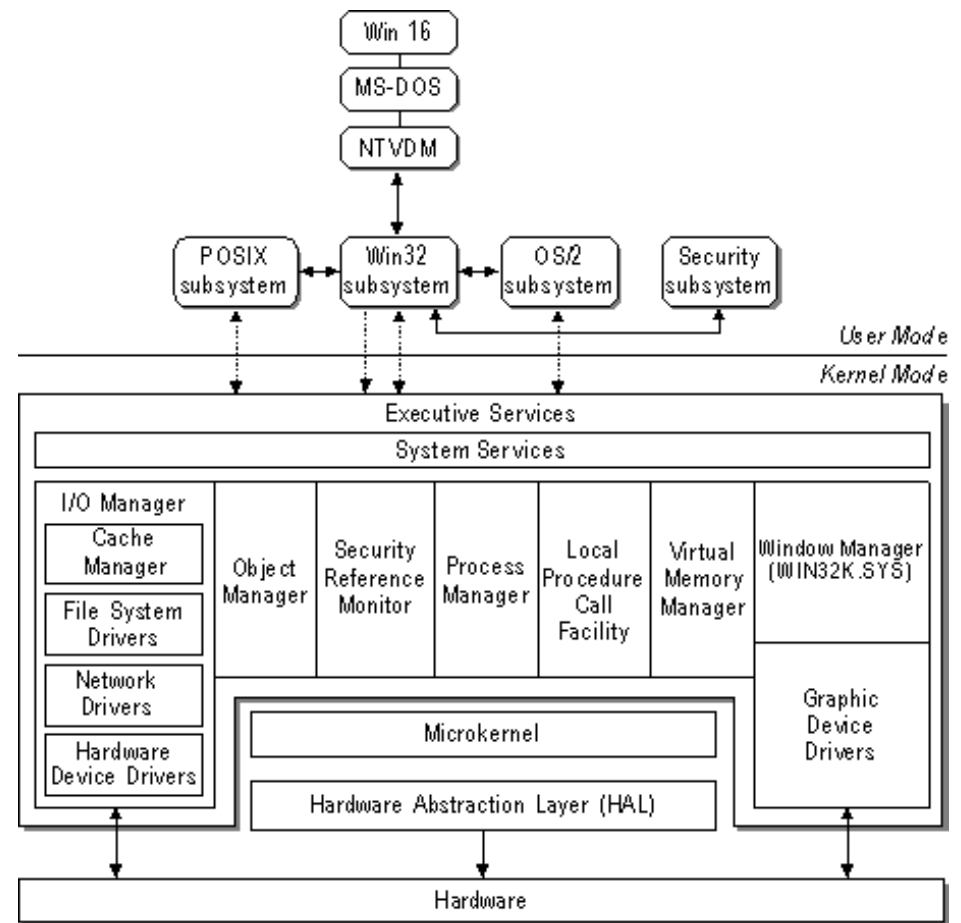
OS Architecture

- Applications make requests of the OS kernel through the system call interface and other traps (more on this shortly).
- There are many user applications, but only one OS; hence, the OS must handle concurrent requests.
- Most of the kernel is written in a portable high-level language (C).
- Small hardware-dependent layer written in machine code.



OS Architecture

- Kernel code is complex; there are many popular designs. NT is a “hybrid monolithic/microkernel”.
- Most designs strive to be modular; one component for each hardware abstraction.
- Lots of design concerns:
 - Good Abstractions!
 - Fairness
 - Speed
 - Security
 - Maintainability



Syscall Interface

- In CS61 we taught you:
 - To run a program: `fork` and `execve`!
 - ...so what does that mean?
- `fork` and `execve` are **syscalls**:
 - A request to the OS to do something. It may fail.
 - A lot like a function call into a shared library.
- But unlike a shared library, the OS is “trusted”:
 - User processes aren't allowed to mess with other processes or hardware, or change file permissions, etc.
 - But the OS has to; it uses special CPU instructions to do so.

Kernel Mode

- The `syscall` instruction causes the CPU to jump to the OS's “syscall handler” in **kernel mode**.
- In kernel mode, the CPU can do things like:
 - Access any process's memory.
 - Manipulate page tables.
 - Access hardware (e.g. the hard disk).
 - Transfer control between processes.
- The syscall handler inspects argument registers to figure out what the user wants, and then calls the appropriate OS component to handle the request.

Traps

- A **trap** is any control transfer to the kernel mode OS.
 - Syscalls are a voluntary trap.
- Involuntary traps happen constantly!
 - Segfault! OS takes control, kills your process.
 - Page fault! OS takes control, loads your page.
 - Device interrupts! Your file write is done/packet has arrived.
 - Timer interrupts! You've had the CPU for long enough.
- Setting traps is a kernel mode privilege.
 - Traps are how the OS maintains control over the machine.
 - The OS must never lose control! *No user code in kernel mode.*

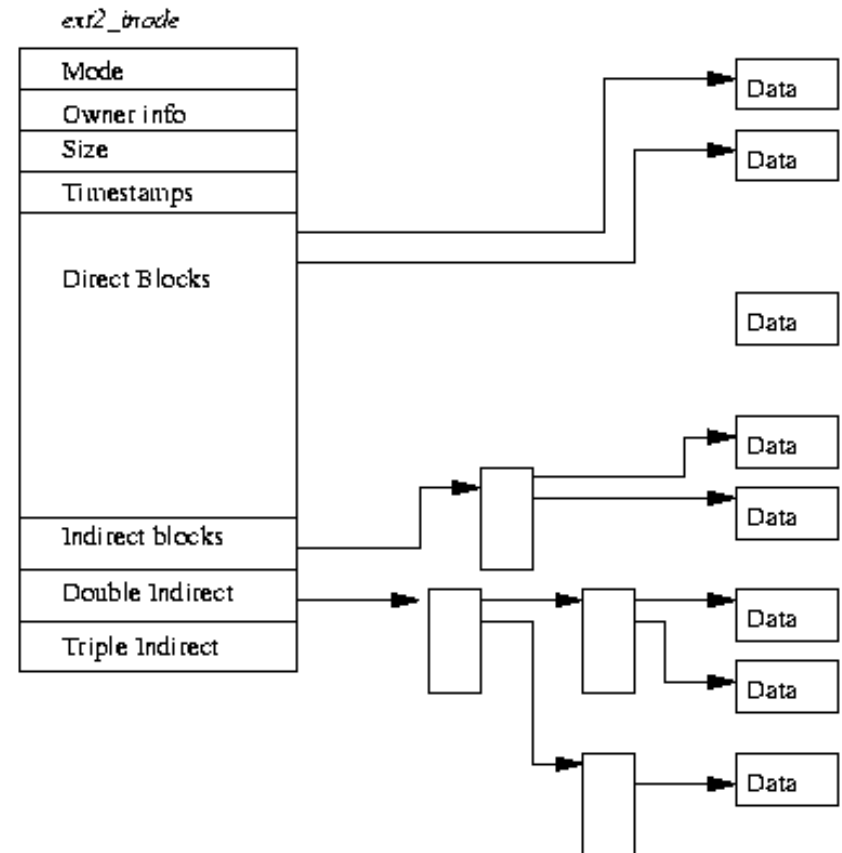
Filesystems

In CS61, we told you...

- Files are:
 - “Infinite-length” byte-addressable character streams.
 - You can open, close, read, and write them.
 - They live in hierarchical directories.
- Hard discs are:
 - Platters, tracks, sectors, and blocks of data. “Logical blocks” are addressable by their ID number.
 - Blocks are fixed-length, and sometimes they just go bad.
- There is a **lot of code** between files and hard disks!

Files (ext2)

- Every file has a corresponding disk structure called an **inode**.
- The inode contains file metadata, and pointers to data blocks.
- Only so much room for pointers...so we also have pointers to blocks full of pointers!
- ext2 has a max file size (2TB).



Directories (ext2)

- A directory is like a file that contains filenames and inode #s.

- Note that filenames are a property of the containing directory, not the file!

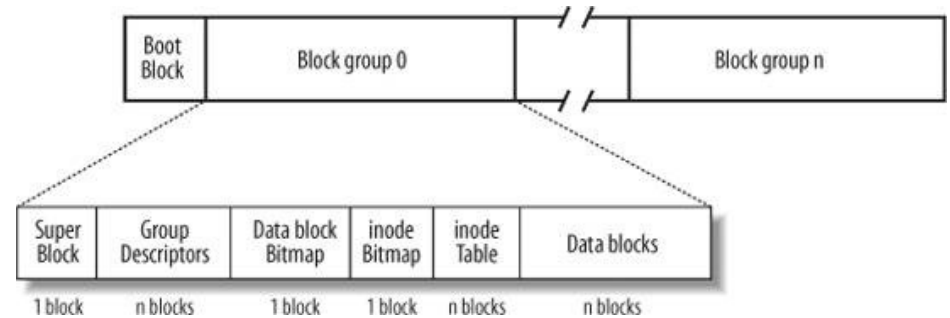
- To find the file “/a/b/c”:
 - Open unique “/” directory.
 - Find “a/” in “/”.
 - Find “b/” in “a/”.
 - Find “c” in “b/”.

	inode	rec_len	name_len	file_type	name			
0	21	12	1	2	.	\0	\0	\0
12	22	12	2	2	.	.	\0	\0
24	53	16	5	2	h	o	m	e
40	67	28	3	2	u	s	r	\0
52	0	16	7	1	o	l	d	f
68	34	12	4	2	s	b	i	n

- A cache speeds this up a lot.

Filesystem (ext2)

- The filesystem consists of logical blocks organized in **block groups**.
- Each group contains:
 - A copy of the **superblock** (filesystem metadata).
 - A **group descriptor** (block group metadata).
 - Bitmaps that tell which blocks are allocated/free.
 - inodes and data blocks.



A Modern Filesystem (ZFS)

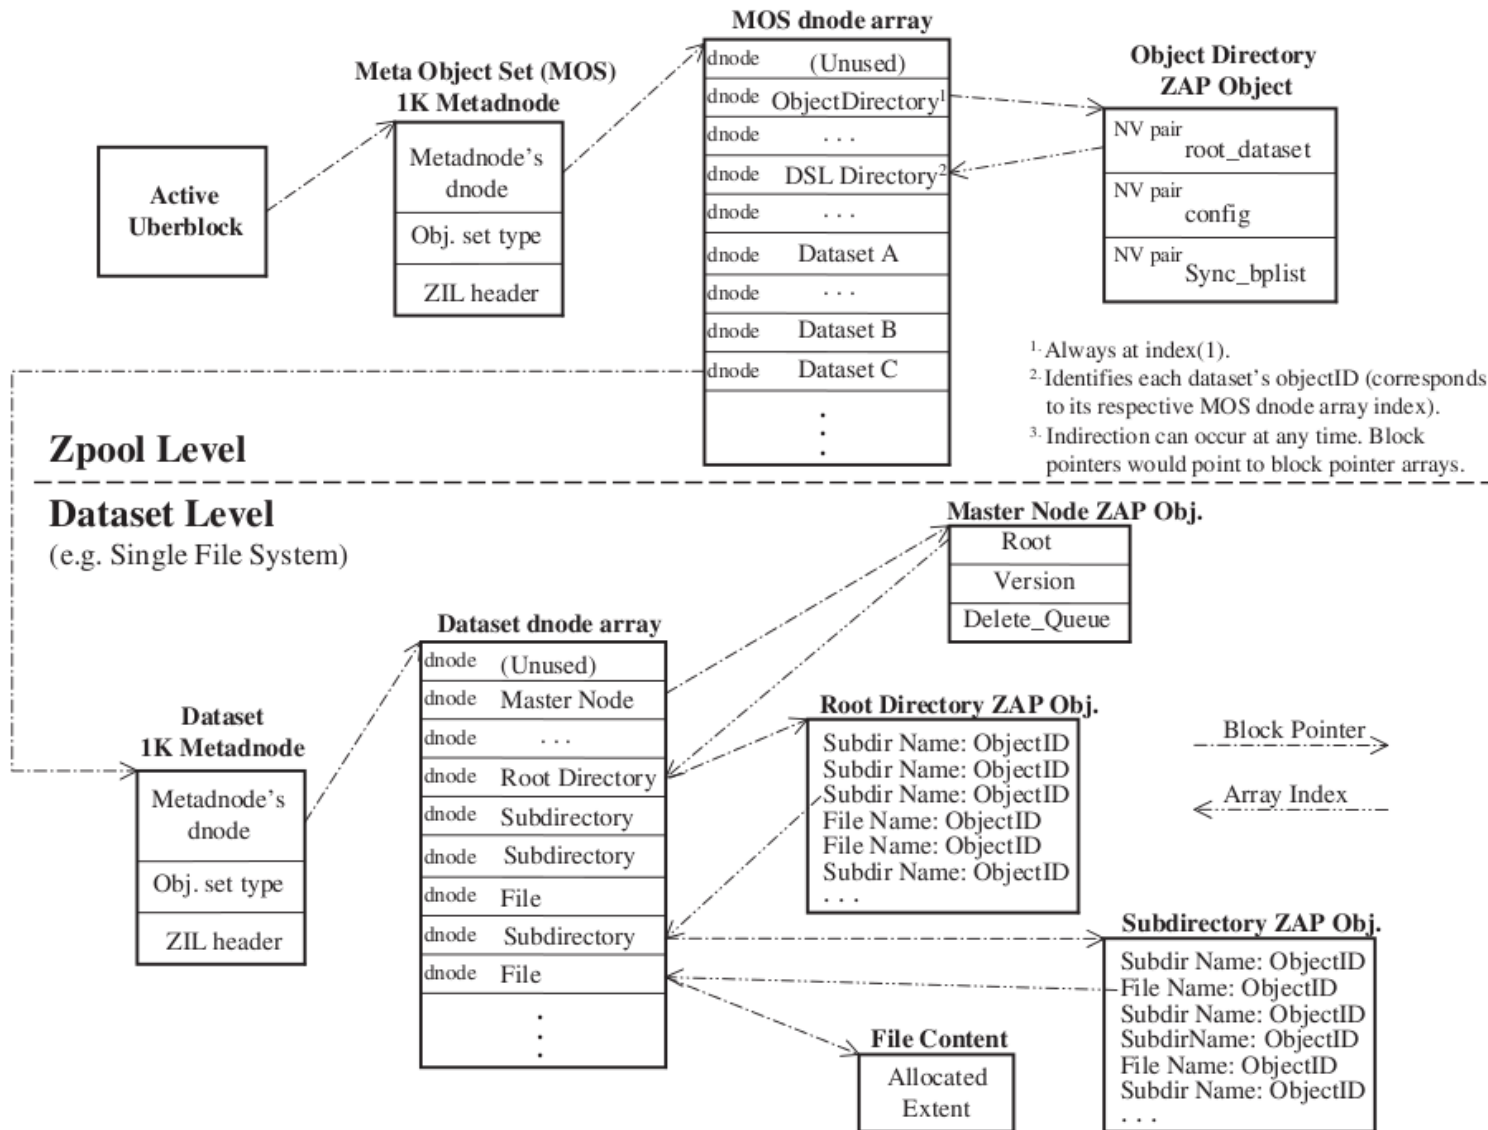


Fig. 2 – On-Disk Data Walk (Figure adapted from Bruning, 2008a).

CS161

- CS161 is a project oriented-course.
 - We give you a half-finished OS...
 - 202 .c files; 84 .h files; 30,514 lines of code.
 - ...and then you have to finish it.
- The OS we give you, OS/161, is “real”:
 - Written at Harvard by David Holland.
 - Inspired by NetBSD, and very similar.
 - Runs on a machine simulator (for sane debugging).

CS161 Projects

- Synchronization:
 - Implement primitives!
- Syscalls:
 - Implement all of them! Every one! Even `fork` and `execve`!
- Virtual Memory:
 - Write it from scratch. And make it thread safe.
 - And don't forget the OS has to page *itself*. Pretty weird stuff.
- Filesystems:
 - Make it thread safe (surprisingly challenging).
 - Performance is important; changes are encouraged.
- You will learn a lot if you put in the effort.

Questions?



HARVARD

School of Engineering
and Applied Sciences

Topics in Systems: Programming Languages

CS61, Lecture 24

Stephen Chong

December 2, 2010

What is PL?

- We use **programming languages** to tell computers what to do
 - Describes computation
 - An interface between humans and computers
- PL is the study of programming languages
 - Principles and limitations of **computing** (programming) **models**
 - Design, implementation, use, and comparison of systems or languages built on these models

Benefits of PL

- Widely-applicable design and implementation techniques
 - Understanding (and efficiently implementing) many models of computation leads to insights and principles for system design
 - e.g., Google's map-reduce framework
 - Inspired by map and reduce constructs in LISP and other functional languages
 - Stateless, effect-free computation over streams of data—allows it to scale
 - e.g., Many systems take complex input and perform computation described by that input
 - Web browsers, printer drivers, PDF renderers, scripted robot control systems, spreadsheets, ...

Benefits of PL

- Create new domain specific languages or Virtual Machines
 - Mathematica, MATLAB, LaTeX, Verilog, VHDL, ...
 - Many programmers create new domain specific APIs, languages, or VMs.
 - How expressive should language be? How to execute programs?
 - Many common mistakes: dynamically scoped variables, non-symmetric functional call and return, ...
- Domain specific languages provide computational model for thinking about data and algorithms specific to problems in a specific domain
 - e.g., How to use map-reduce effectively?
 - Sawzall, Dryad, Pig: languages that give a higher level of abstraction than MapReduce
 - E.g. MLFi


Abstraction

- Programming languages provide **abstraction**
- Abstractions of underlying machinery or mechanisms
 - e.g., assembly language instead of machine opcodes
 - e.g., high-level languages (C, Java, Pascal, ...) instead of assembly
 - e.g., monitors, semaphores, actors, join-patterns, ..., provide abstraction of concurrent execution
- Abstractions of domain specific concepts
 - Swazall, Dryad, Pig, MLFi, ...
- Sapir–Whorf hypothesis: language influences thought



Compilation vs Interpretation

- Compilation: translating higher level languages to lower level languages
 - e.g., C to assembly
- Interpretation: a program (interpreter) executes the source code directly
 - e.g., JavaScript interpreter in browser
- Pros and cons:
 - Efficiency, portability, obfuscation, ease of rewriting
- Blurring the boundaries:
 - Virtual machines
 - JIT compilation



Functional languages

- What values can a program manipulate at runtime?
 - Integers, strings, floats, ...
- Functional languages allow *functions* to be runtime values
 - i.e., functions can be given as arguments, received as return values, etc.
 - e.g.,
 - $\text{addFour} \equiv \lambda x:\text{int}. x+4$ is a function that takes an argument x and adds 4 to it.
 - $\text{foo} \equiv \lambda f:\text{int} \rightarrow \text{int}. f(f(34))$ is a function that takes an argument f , applies f to 34, and applies f to the result
 - $\text{foo}(\text{addFour})$ evaluates to 42
 - Can be mind-bending the first time you see it!
 - $\text{fact} \equiv \lambda f, n. \text{ if } (n=0) \text{ then } 1 \text{ else } n*(f(f, n-1))$
 - $\text{fact}(\text{fact}, 5)$

Functional languages

- Functional languages are typically **pure**
 - Without side-effects, do not affect memory or perform I/O
- Simplifies concurrent programming
 - All computations just perform reads
 - Means concurrent computations do not affect each other so synchronization simpler

Declarative vs. Imperative

- Imperative languages: computation described as *commands*: statements that change program state
 - e.g., `x := 1; while (x < 10) x := x+1;`
- Declarative languages: computation described as *what* the computation should achieve, not *how* it should achieve it
 - e.g., `select avg(salary) from employees where dept = "Accounting"`
- Presents different abstractions to the programmer
 - Declarative programming higher level, and typically more compact
 - SQL one of the big success of declarative programming
- Joe Hellerstein (visiting Harvard this year): BOOM: Orders of Magnitude simpler code for the cloud

Language-based security

- Using programming language techniques and abstractions for security
- Using formal semantics of programming languages to **define** security
 - What does it mean for a system to be secure?
- Mechanisms to enforce language abstractions
 - E.g., buffer overruns are dangerous because they violate the programmer's abstraction. Mechanisms to enforce this abstraction (e.g., error when array accessed out of bounds) improve security
- Mechanisms to enforce/reason about program behavior
 - Program analysis
 - e.g., type system can prove that program never treats an `int` as if it were a pointer
 - e.g., reasoning about correctness of cryptographic protocols
 - Program rewriting
 - Does the program satisfy correct behavior? Who knows? Just rewrite it so that it does...
 - ...

Language-based security

- Greg Morrisett: using sophisticated logics to prove that programs are “correct”, including programs of Unknown Provenance
- Me: application-specific information security: specifying, reasoning about, and enforcing



HARVARD

**School of Engineering
and Applied Sciences**

The end of CS 61

What we've covered

- An introduction to computer systems
 - Machine representation of data and programs
 - Compilation, optimization, linking, loading
 - Memory, storage, caching, virtual memory, dynamic memory management
 - Systems programming
 - I/O, sockets, processes
 - Concurrency
 - Threads, synchronization mechanisms, synchronization problems
- Answered (to some extent!) the questions
 - What happens when I run a program?
 - How do computers work?
 - What affects the performance and reliability of my programs?

Where to from here

- We've just scratched the surface!
- Some courses
 - CS 141: Computing hardware
 - CS 143: Computer networks
 - CS 152: Programming languages
 - CS 153: Compilers
 - CS 161: Operating systems
 - CS 171: Visualization
 - CS 175: Computer graphics
 - CS 189r: Autonomous multi-robot systems
 - Other 100 level courses and many 200 level courses...
- Harvard Computer Society <http://www.hcs.harvard.edu>
- Systems Research at Harvard <http://www.eecs.harvard.edu/~syrah/>

Final exam

- Final exam Thursday Dec 16, 2pm-5pm
- All material covered in lecture (except guest lecture), sections, and labs is examinable
- Similar format to midterm
- Open book, closed note
- Office hours will continue as normal
 - (unless otherwise posted)
- A practice final will be released soon