



HARVARD

School of Engineering
and Applied Sciences

Machine Programming 3: Procedures

CS61, Lecture 5

Prof. Stephen Chong

September 15, 2011

Announcements

- Assignment 2 (Binary bomb) due next week
 - If you haven't yet please create a VM to make sure the infrastructure works for you

Today

- Procedures
 - The stack
 - Stack frames
 - Leave
 - Register conventions
- x86_64

Procedure calls


```
void foo(...) {  
    ...  
    bar();  
    ...  
}
```

```
void bar(...) {  
    int x, y;  
    x = baz();  
    ...  
    y = quux();  
    ...  
}
```

```
int baz(...) {  
    int z;  
    ...  
    z = baz();  
    ...  
    return z;  
}
```

```
int quux(...) {  
    ...  
    return 42;  
}
```

Call chain




- How do we call procedures?
- Where do we store local variables (e.g., x,y,z)?
- How do we return values from procedures?
- How do we support recursion?

Stack

- Stack is used for handling function calls and local storage
 - Stores local variables, return address, saved registers, ...
- Stack pointer **%esp** always holds address of top stack element
- Stack grows **downwards!**

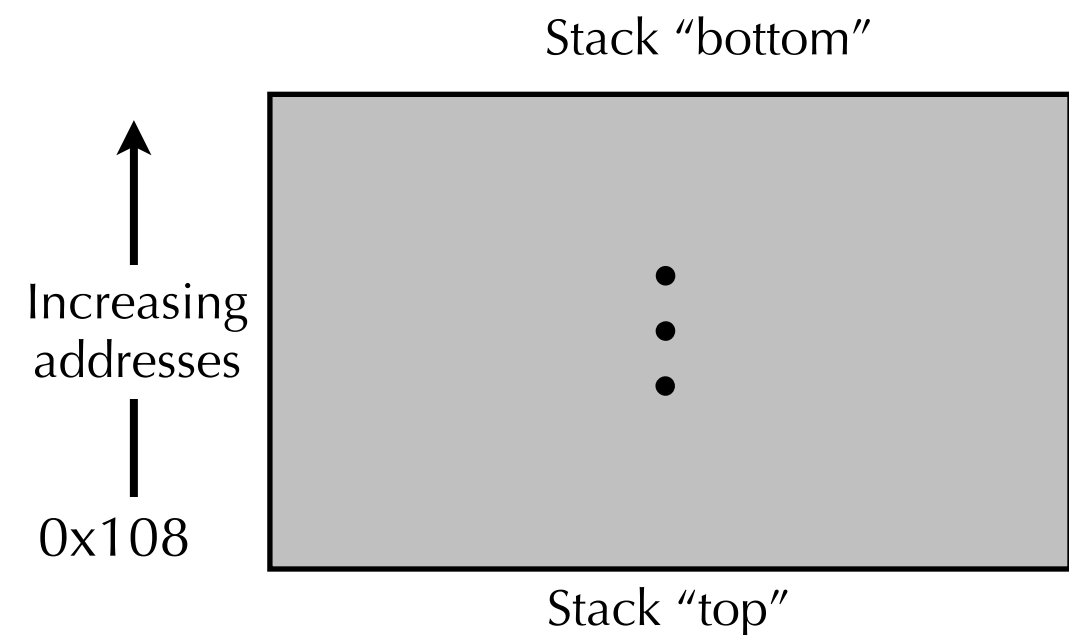
%eax	0x123
%edx	0
%esp	0x108



Pushing and popping

- Two data movement instructions for stack: **pushl** and **popl**
- **pushl** *src*
 - Push four bytes onto stack
 - Effect is
$$R[\%esp] \leftarrow R[\%esp] - 4$$
$$M[R[\%esp]] \leftarrow src$$
- E.g., **pushl %eax**


%eax	0x123
%edx	0
%esp	0x108



Pushing and popping

- Two data movement instructions for stack: **pushl** and **popl**
- **pushl** *src*
 - Push four bytes onto stack
 - Effect is
$$R[\%esp] \leftarrow R[\%esp] - 4$$
$$M[R[\%esp]] \leftarrow src$$
- E.g., **pushl %eax**

%eax	0x123
%edx	0
%esp	0x104



Pushing and popping

- **popl** *dest*


- Pops four bytes from stack
- Effect is

$$dest \leftarrow M[R[\%esp]]$$

$$R[\%esp] \leftarrow R[\%esp] + 4$$

- E.g., **popl** %edx

%eax	0x123
%edx	0
%esp	0x104



Pushing and popping


- **popl** *dest*

- Pops four bytes from stack
- Effect is

$$dest \leftarrow M[R[\%esp]]$$
$$R[\%esp] \leftarrow R[\%esp] + 4$$

- E.g., **popl** %edx


%eax	0x123
%edx	0x123
%esp	0x108



Examining the stack

- Can use `movl` to access and modify arbitrary values on the stack
 - No need to access just top element
 - Can “peek” at stack:
 - `movl 12(%esp), %eax`
 - Can “poke” stack:
 - `movl $0xdeadbeef, 12(%esp)`

%eax	0x123
%edx	0x123
%esp	0x108



Procedure control flow


- Stack is used to implement procedure call and return
- Procedure call
 - x86 instruction: `call address`
 - Pushes **return address** on stack, then jumps to *address*
 - What is the return address?
 - Address of instruction **after** the `call` instruction
 - E.g.,

804854e:	e8 3d 06 00 00	call	8048b90 <main>
8048553:	50	pushl	%eax
 - Return address is 0x8048553
- Procedure return
 - x86 instruction: `ret`
 - Pops return address from stack, and jumps to it

Procedure call example


```
804854e: e8 3d 06 00 00    call 8048b90 <main>
8048553: 50                pushl %eax
```

Before call



%esp	0x108
%eip	0x804854e

After call




%esp	0x104
%eip	0x8048b90

Procedure call example


```
804854e: e8 3d 06 00 00    call 8048b90 <main>
8048553: 50               pushl %eax
```

Before call



%esp	0x108
%eip	0x804854e

After call




%esp	0x104
%eip	0x8048b90

Procedure return example

8048591: c3


ret

Before return



%esp	0x104
%eip	0x8048b91

After return



%esp	0x108
%eip	0x8048553


Stack-based languages

- Languages that support recursion
 - E.g., C, Pascal, Java
 - Must be able to support multiple instantiations of single procedure
 - Code must be **reentrant**
- Each invocation of a procedure has its own local state
 - Arguments to the procedure (e.g., `x`)
 - Local variables within the procedure (e.g., `rval`)
 - Return address
- Where are these stored?

```
int rfact(int x) {  
    int rval;  
    if (x <= 1)  
        return 1;  
    rval = rfact(x-1);  
    return rval * x;  
}
```

Stack frame

- Each procedure invocation has an associated **stack frame**
 - The “chunk” of the stack for that procedure invocation
 - Contains local variables, arguments to functions, and return address
 - Needed from when procedure called to when it returns
- Stack discipline
 - Stack frame released when procedure returns
 - Callee must return before caller does
- Current stack frame described by two registers
 - **%ebp**: frame pointer
 - Points to base (or “bottom”) of current stack frame
 - **%esp**: stack pointer
 - Points to stop of stack (i.e., top of current stack frame)



Stack frame example


```
void foo(...) {  
    ...  
    bar();  
    ...  
}
```

```
void bar(...) {  
    int x, y;  
    x = baz();  
    ...  
    y = quux();  
    ...  
}
```

```
int baz(...) {  
    int z;  
    ...  
    z = baz();  
    ...  
    return z;  
}
```

```
int quux(...) {  
    ...  
    return 42;  
}
```

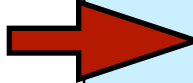
Call chain



Stack frame example

Call chain


foo



```
void foo(...) {  
    ...  
    bar();  
    ...  
}
```

%ebp →



%esp →



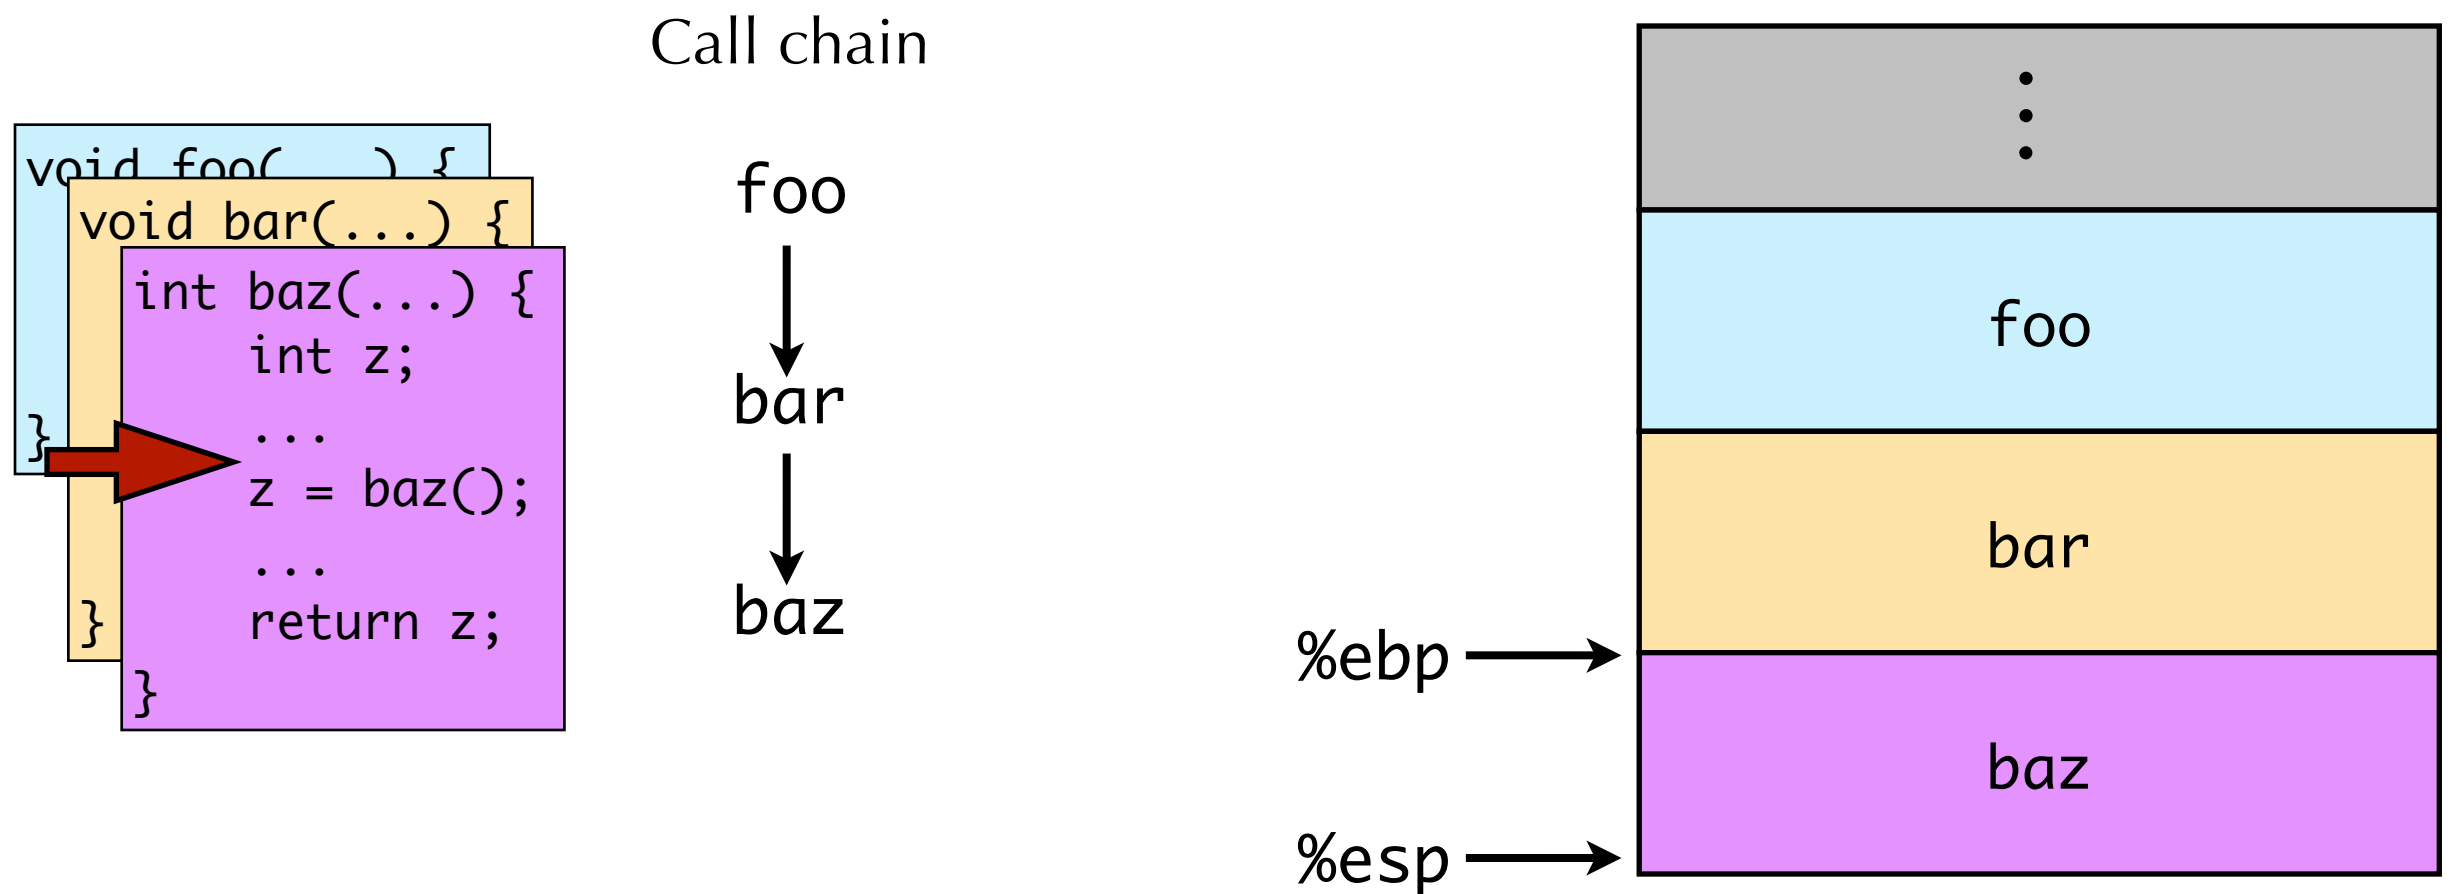
Stack frame example

Call chain

foo
↓
bar




Stack frame example

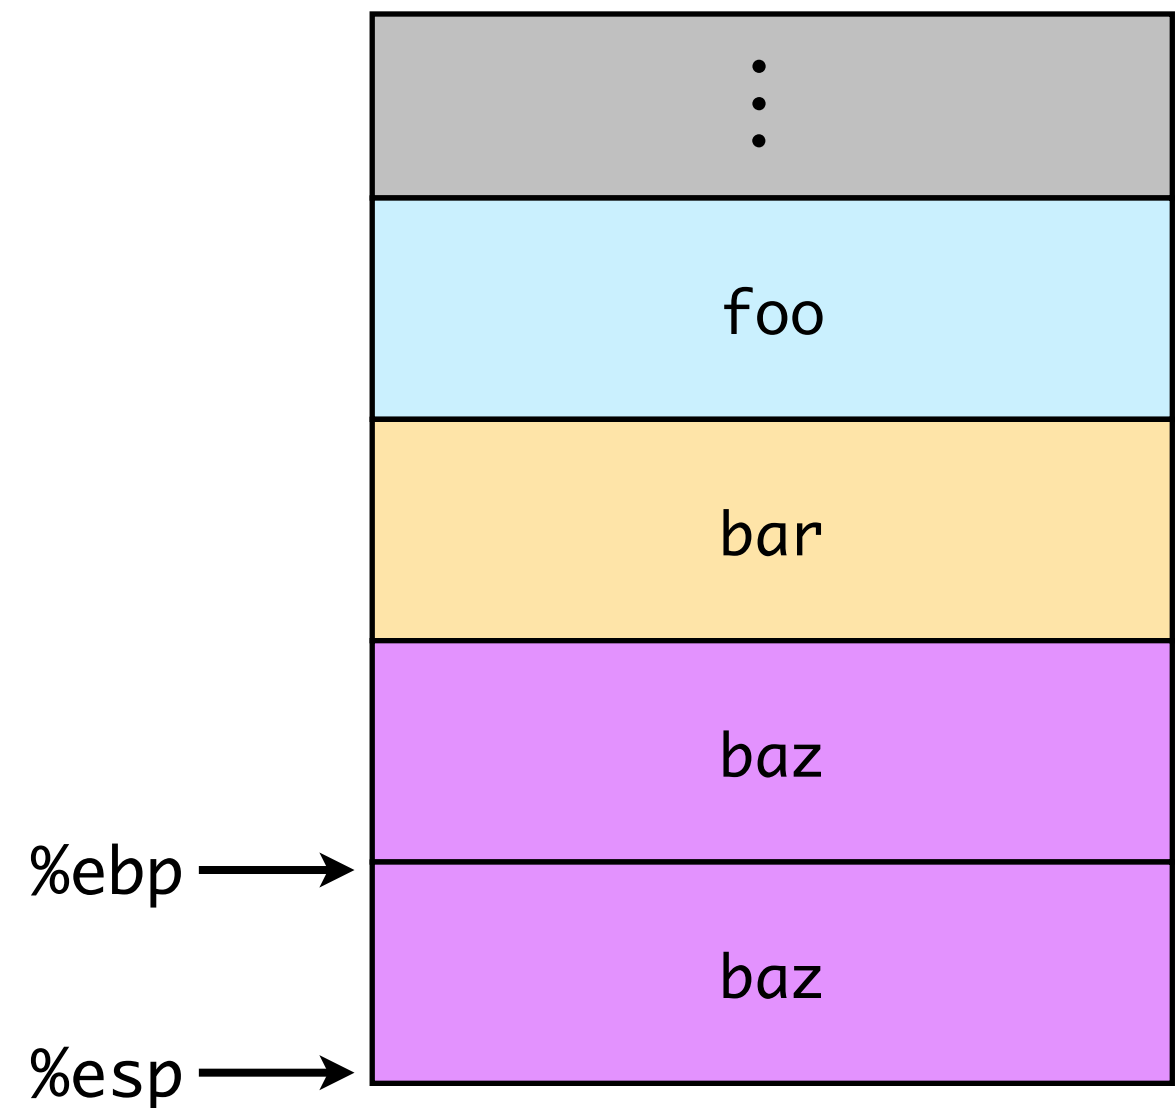


Stack frame example

Call chain




foo
↓
bar
↓
baz
↓
baz

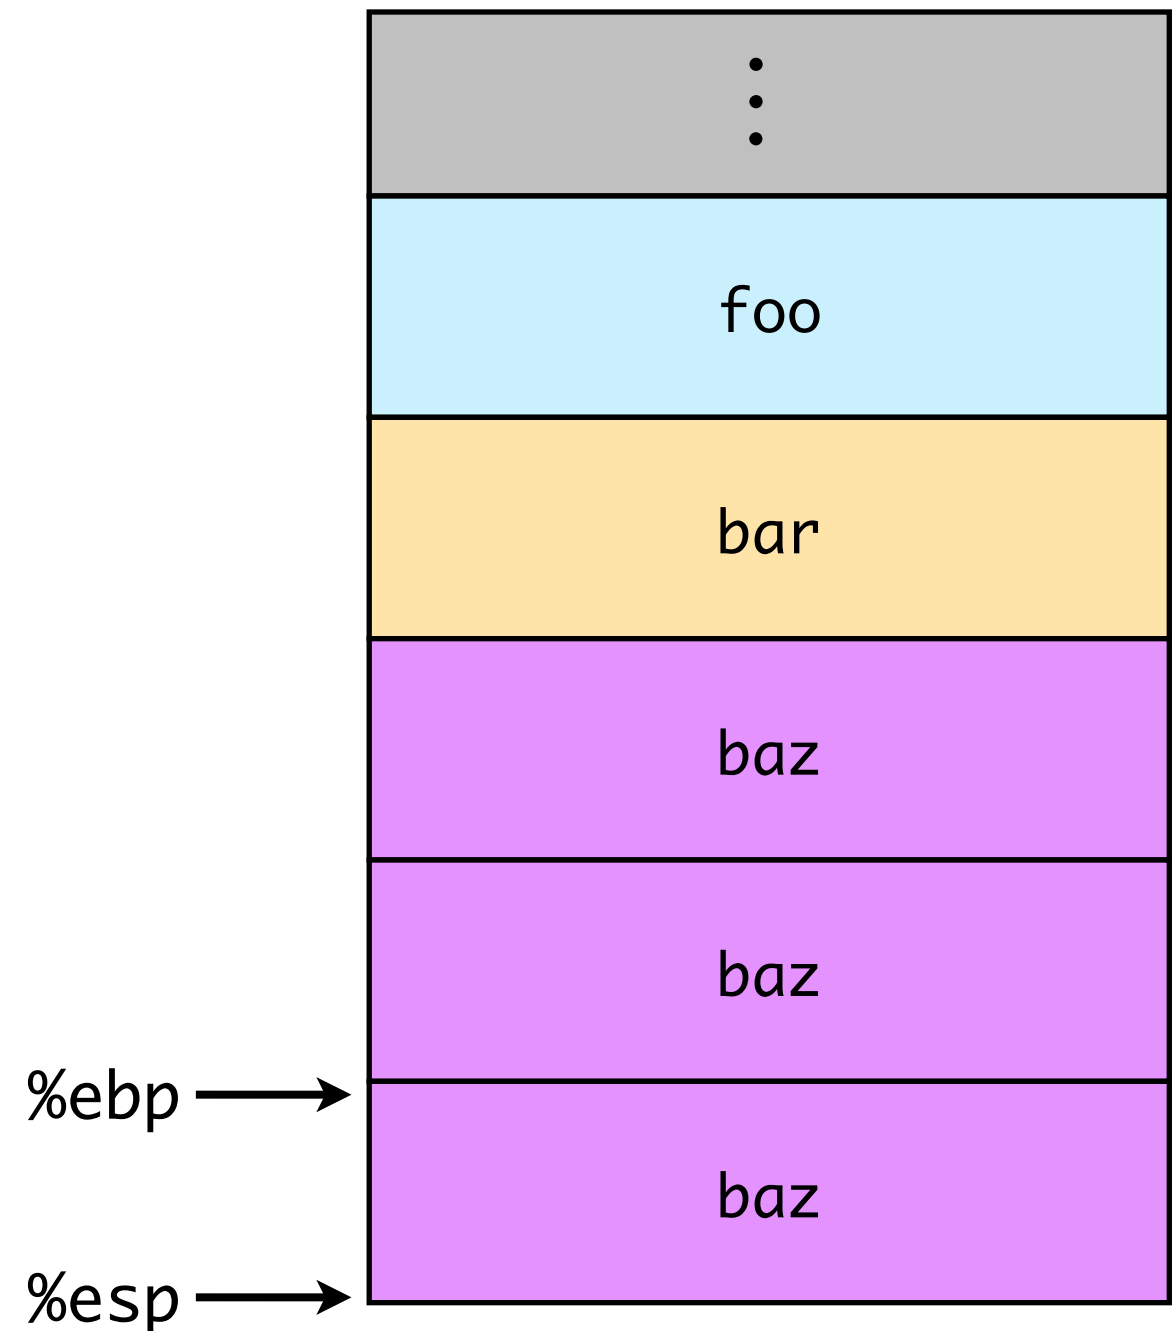


Stack frame example

Call chain




foo
↓
bar
↓
baz
↓
baz
↓
baz

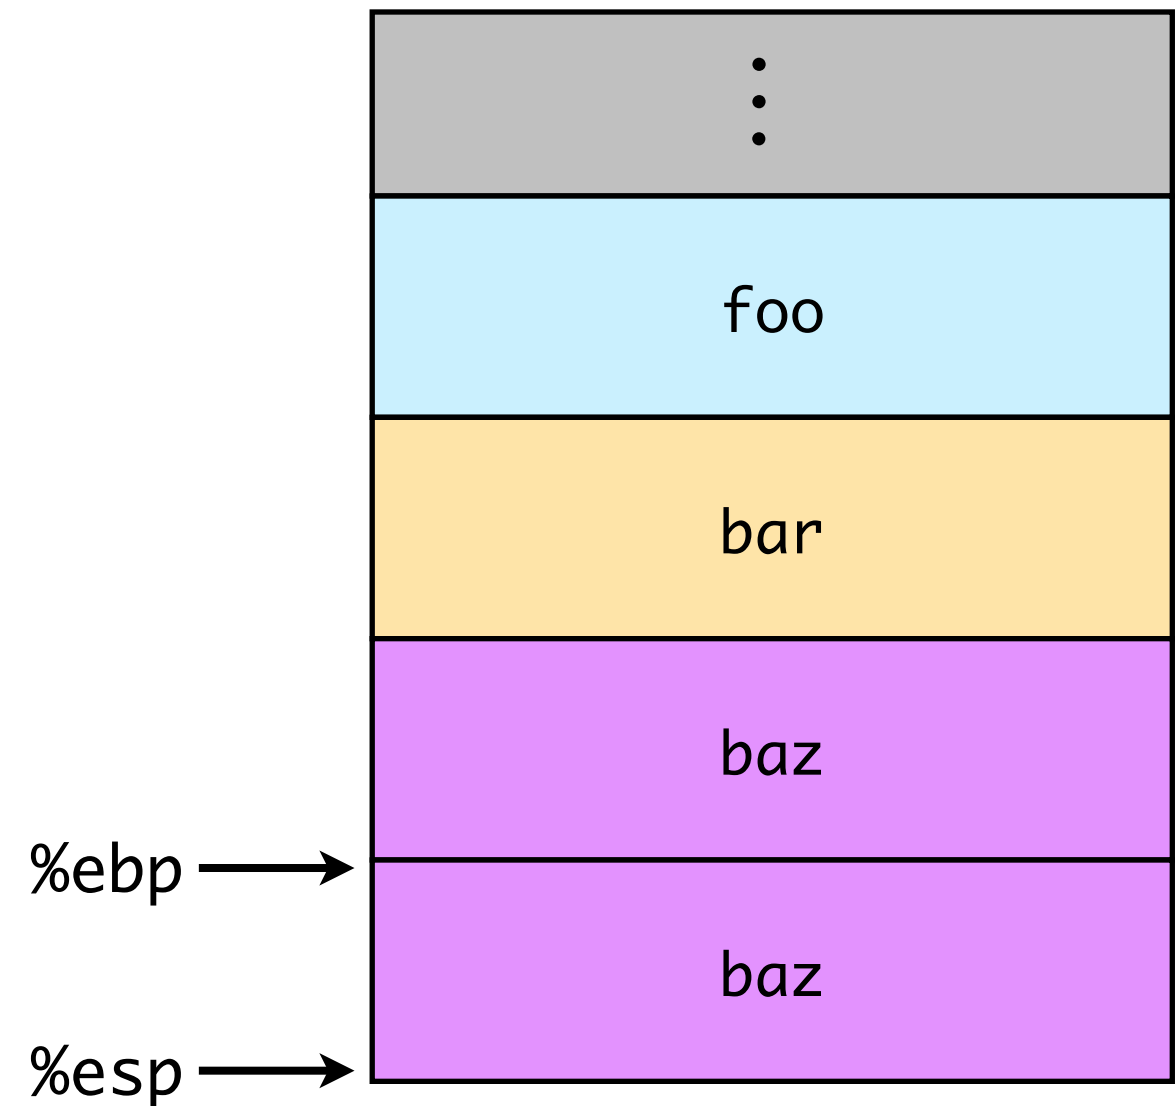


Stack frame example

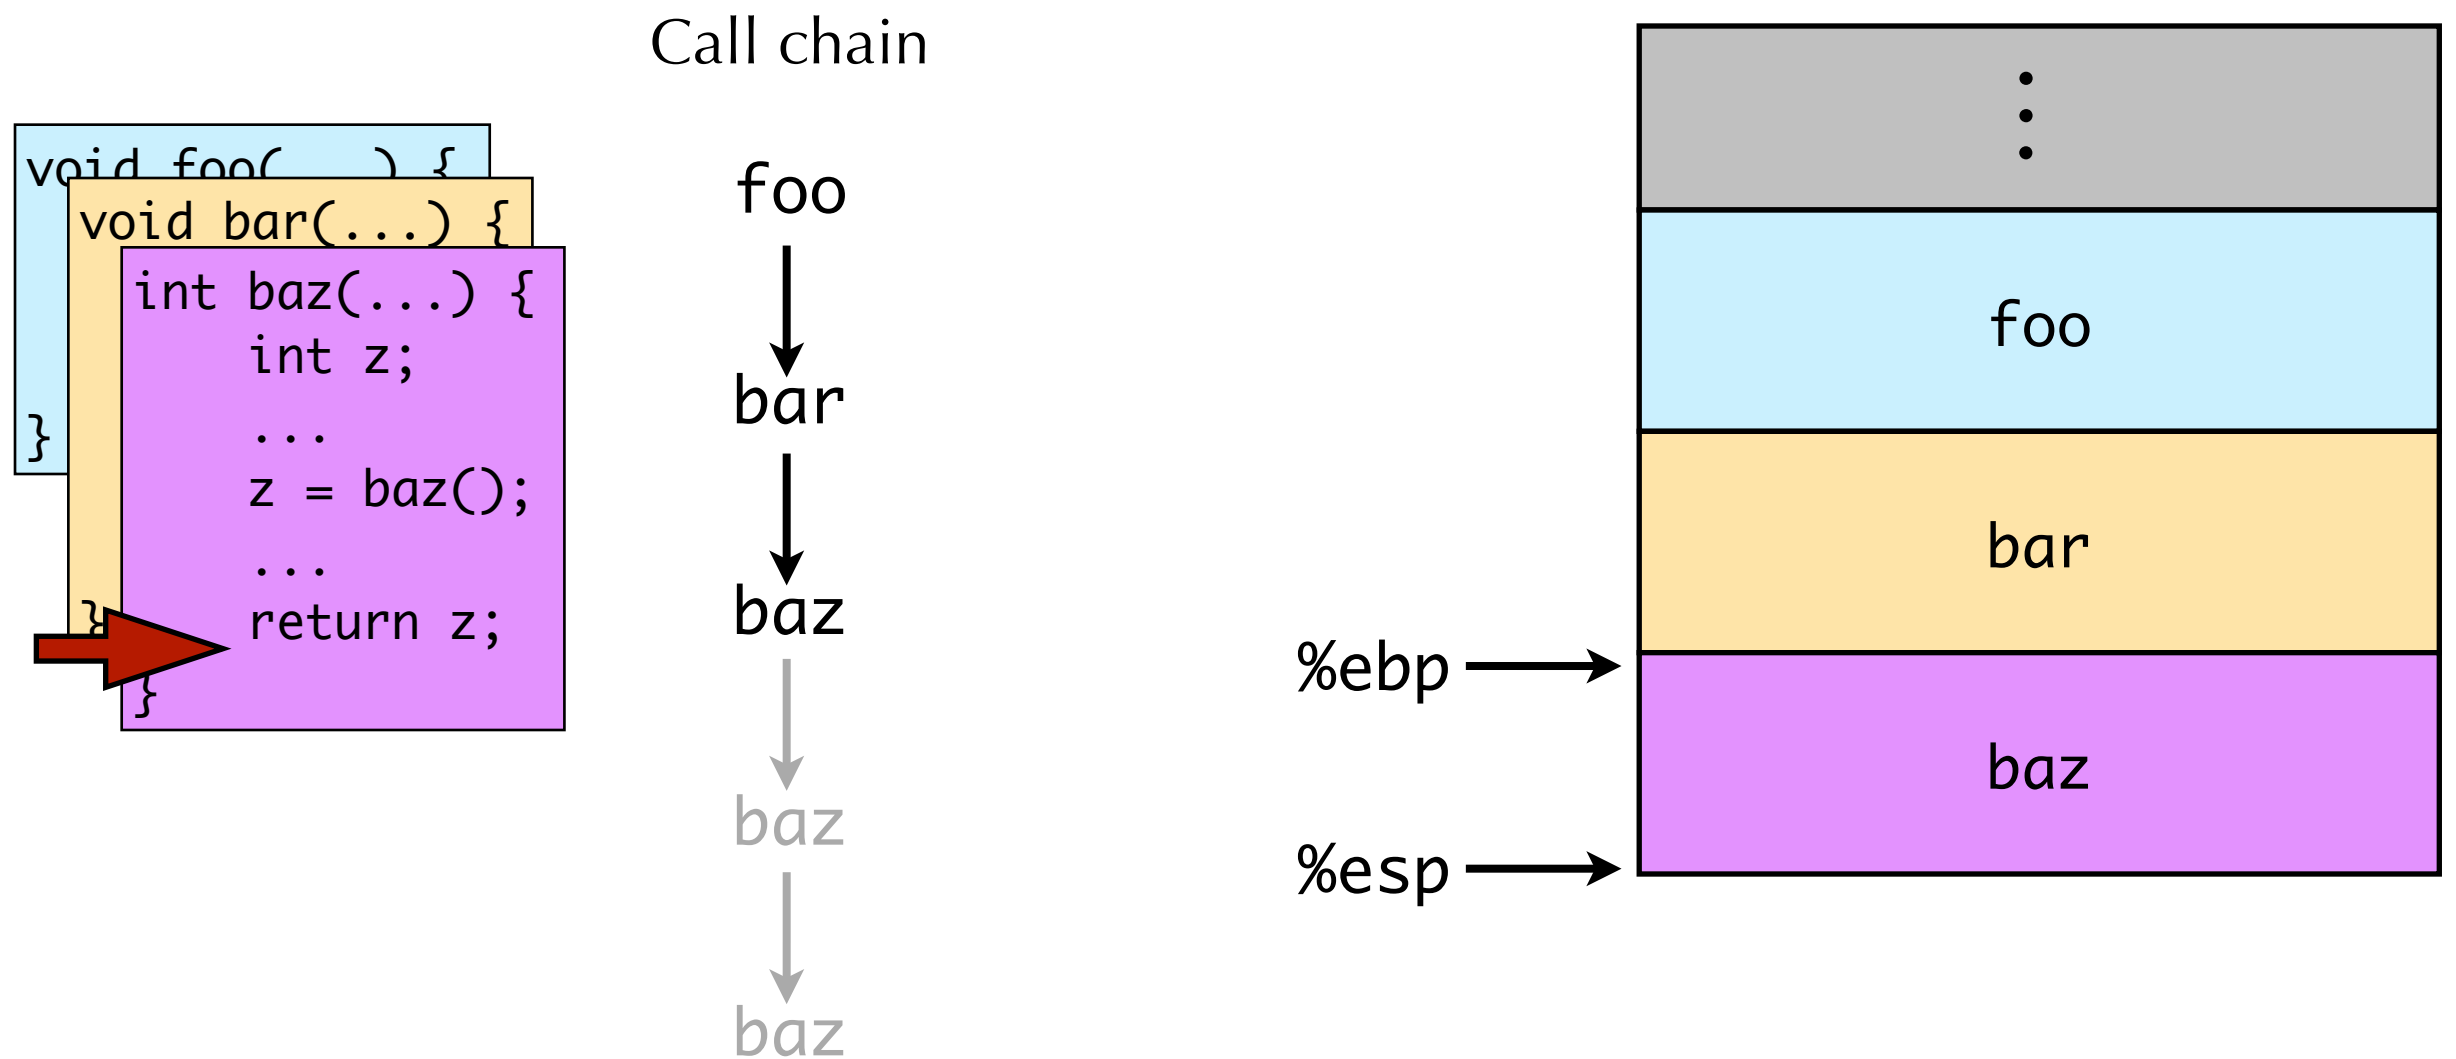
Call chain



foo
↓
bar
↓
baz
↓
baz
↓
baz


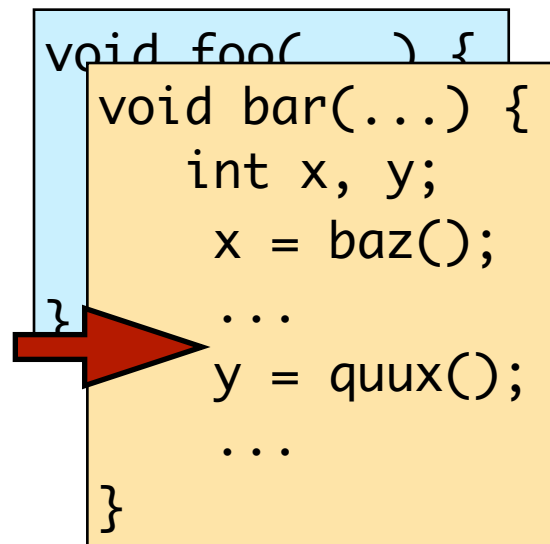


Stack frame example

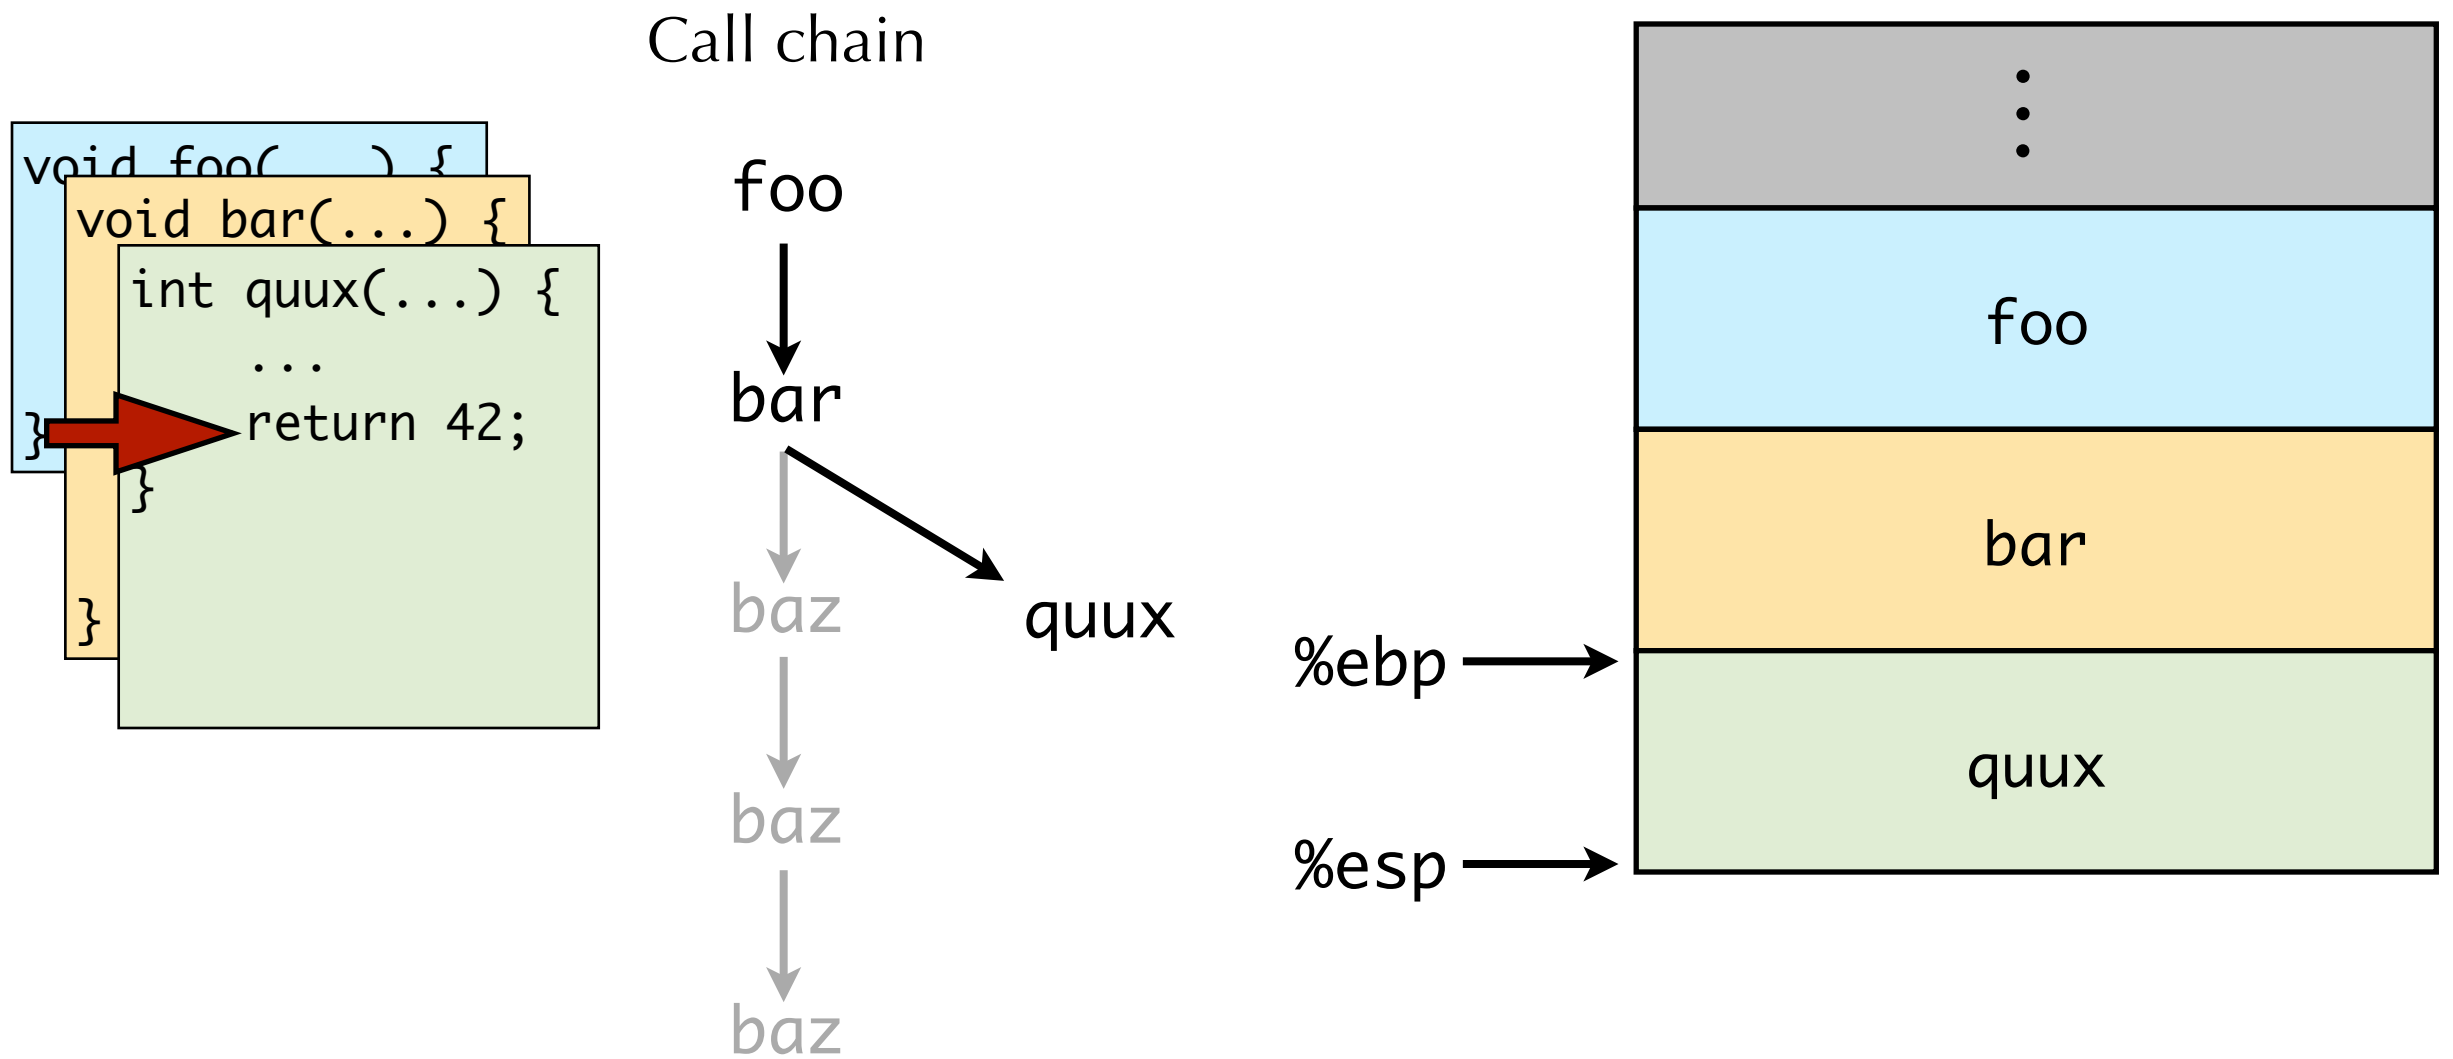


Stack frame example

Call chain



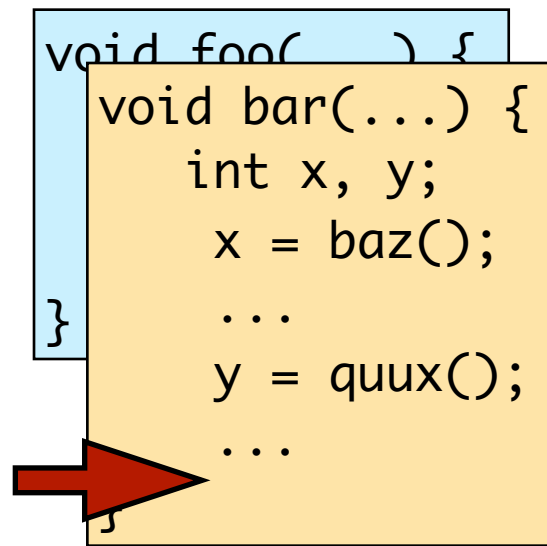


Stack frame example




Stack frame example

Call chain





Stack frame example

```
void foo(...) {  
    ...  
    bar();  
    ...  
}
```

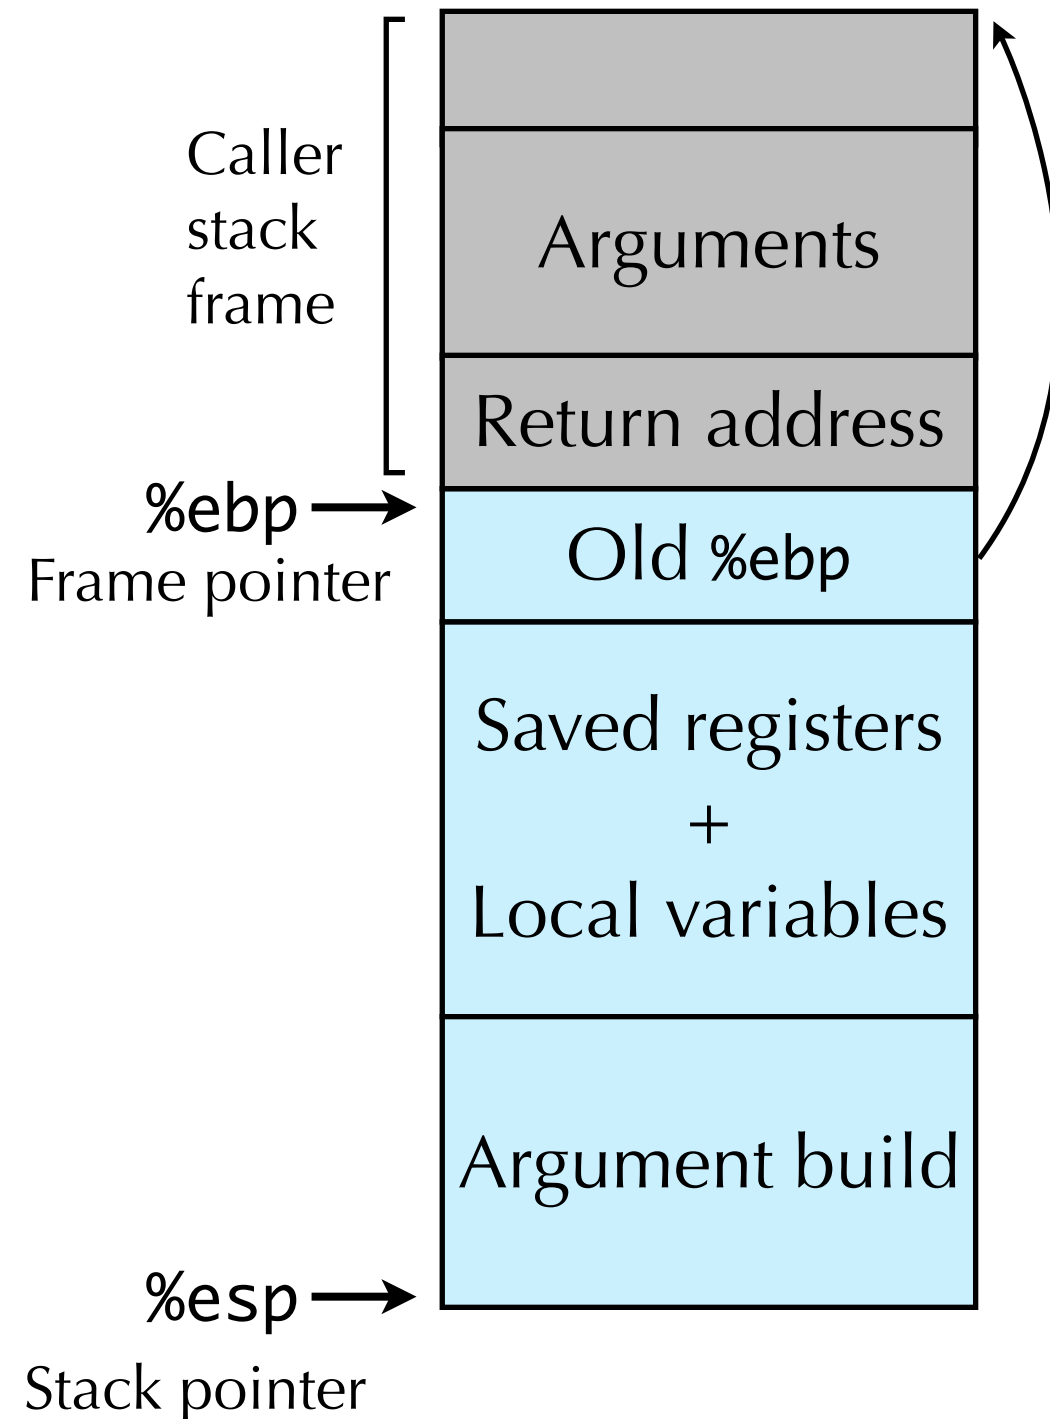


Call chain



x86/Linux stack frame

- The exact layout of a stack frame is a convention.
 - Depends on hardware, OS, and compiler used.
- x86/Linux stack frame contains:
 - Old value of `%ebp` (from previous frame)
 - Any saved registers (more later)
 - Local variables (if not kept in registers)
 - Arguments to function about to be called
- The **caller's** stack frame contains:
 - Return address – pushed by call instruction
 - Arguments for this function call



Swap revisited

```
/* Global vars */  
int zip1 = 15213;  
int zip2 = 91125;  
  
void call_swap() {  
    swap(&zip1, &zip2);  
}
```

```
void swap(int *xp, int *yp) {  
    int t0 = *xp;  
    int t1 = *yp;  
    *xp = t1;  
    *yp = t0;  
}
```

```
call_swap:  
    ...  
    pushl $zip2    # Push args  
    pushl $zip1    #   on stack  
    call swap      # Do the call  
    ...
```

Swap revisited


```
/* Global vars */
int zip1 = 15213;
int zip2 = 91125;

void call_swap() {
    swap(&zip1, &zip2);
}
```

```
void swap(int *xp, int *yp) {
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

call_swap:

```
...
pushl $zip2    # Push args
pushl $zip1    #   on stack
call swap     # Do the call
...
```



Swap revisited

```
/* Global vars */
int zip1 = 15213;
int zip2 = 91125;

void call_swap() {
    swap(&zip1, &zip2);
}
```

```
void swap(int *xp, int *yp) {
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

call_swap:


...

pushl \$zip2 # Push args

pushl \$zip1 # on stack

call swap # Do the call

...




Swap revisited

```
/* Global vars */
int zip1 = 15213;
int zip2 = 91125;

void call_swap() {
    swap(&zip1, &zip2);
}
```

```
void swap(int *xp, int *yp) {
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

```
call_swap:
    ...
    pushl $zip2    # Push args
    pushl $zip1    #   on stack
    call swap      # Do the call
    ...
```




Swap revisited

```
/* Global vars */
int zip1 = 15213;
int zip2 = 91125;

void call_swap() {
    swap(&zip1, &zip2);
}
```

```
void swap(int *xp, int *yp) {
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

```
call_swap:
    ...
    pushl $zip2    # Push args
    pushl $zip1    #   on stack
    call swap      # Do the call
    ...
```




Code for swap

```
void swap(int *xp, int *yp) {  
    int t0 = *xp;  
    int t1 = *yp;  
    *xp = t1;  
    *yp = t0;  
}
```


<pre>swap: pushl %ebp movl %esp,%ebp pushl %ebx</pre>]	Set up
<pre> movl 12(%ebp),%ecx movl 8(%ebp),%edx movl (%ecx),%eax movl (%edx),%ebx movl %eax,(%edx) movl %ebx,(%ecx)</pre>]	Body
<pre> movl -4(%ebp),%ebx movl %ebp,%esp popl %ebp ret</pre>]	Finish

Swap setup

Stack entering swap



Resulting stack




```
pushl %ebp  
movl %esp,%ebp  
pushl %ebx
```


Set up

Swap setup

Stack entering swap




Resulting stack




<code>pushl %ebp</code>] Set up
<code>movl %esp,%ebp</code>	
<code>pushl %ebx</code>	

Swap setup

Stack entering swap




Resulting stack



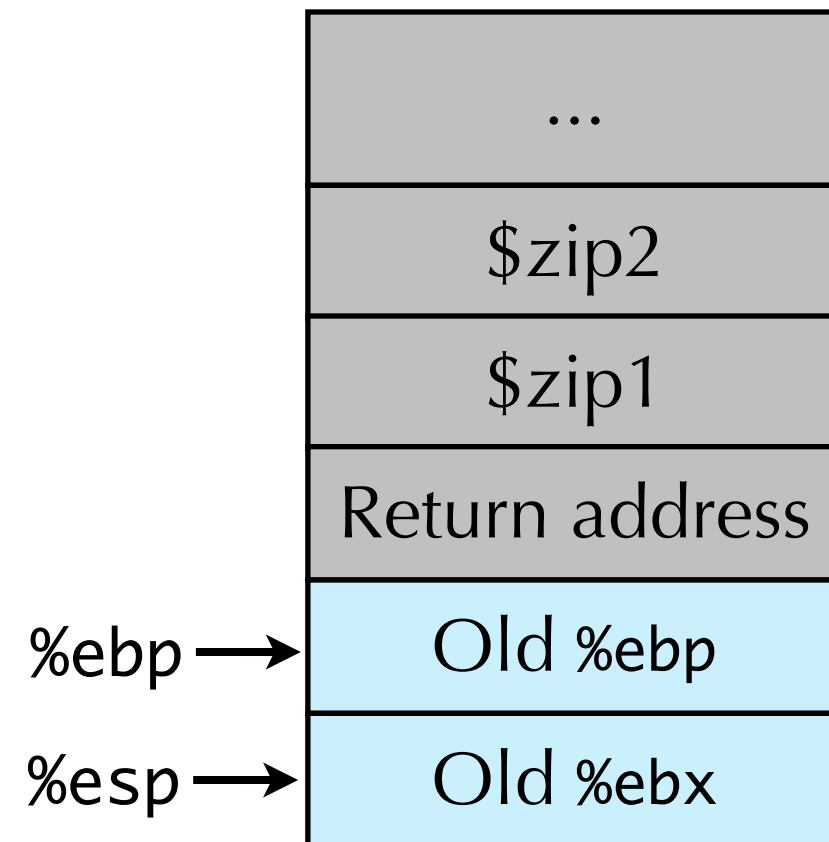
<code>pushl %ebp</code>] Set up
<code>movl %esp,%ebp</code>	
<code>pushl %ebx</code>	

Swap setup

Stack entering swap



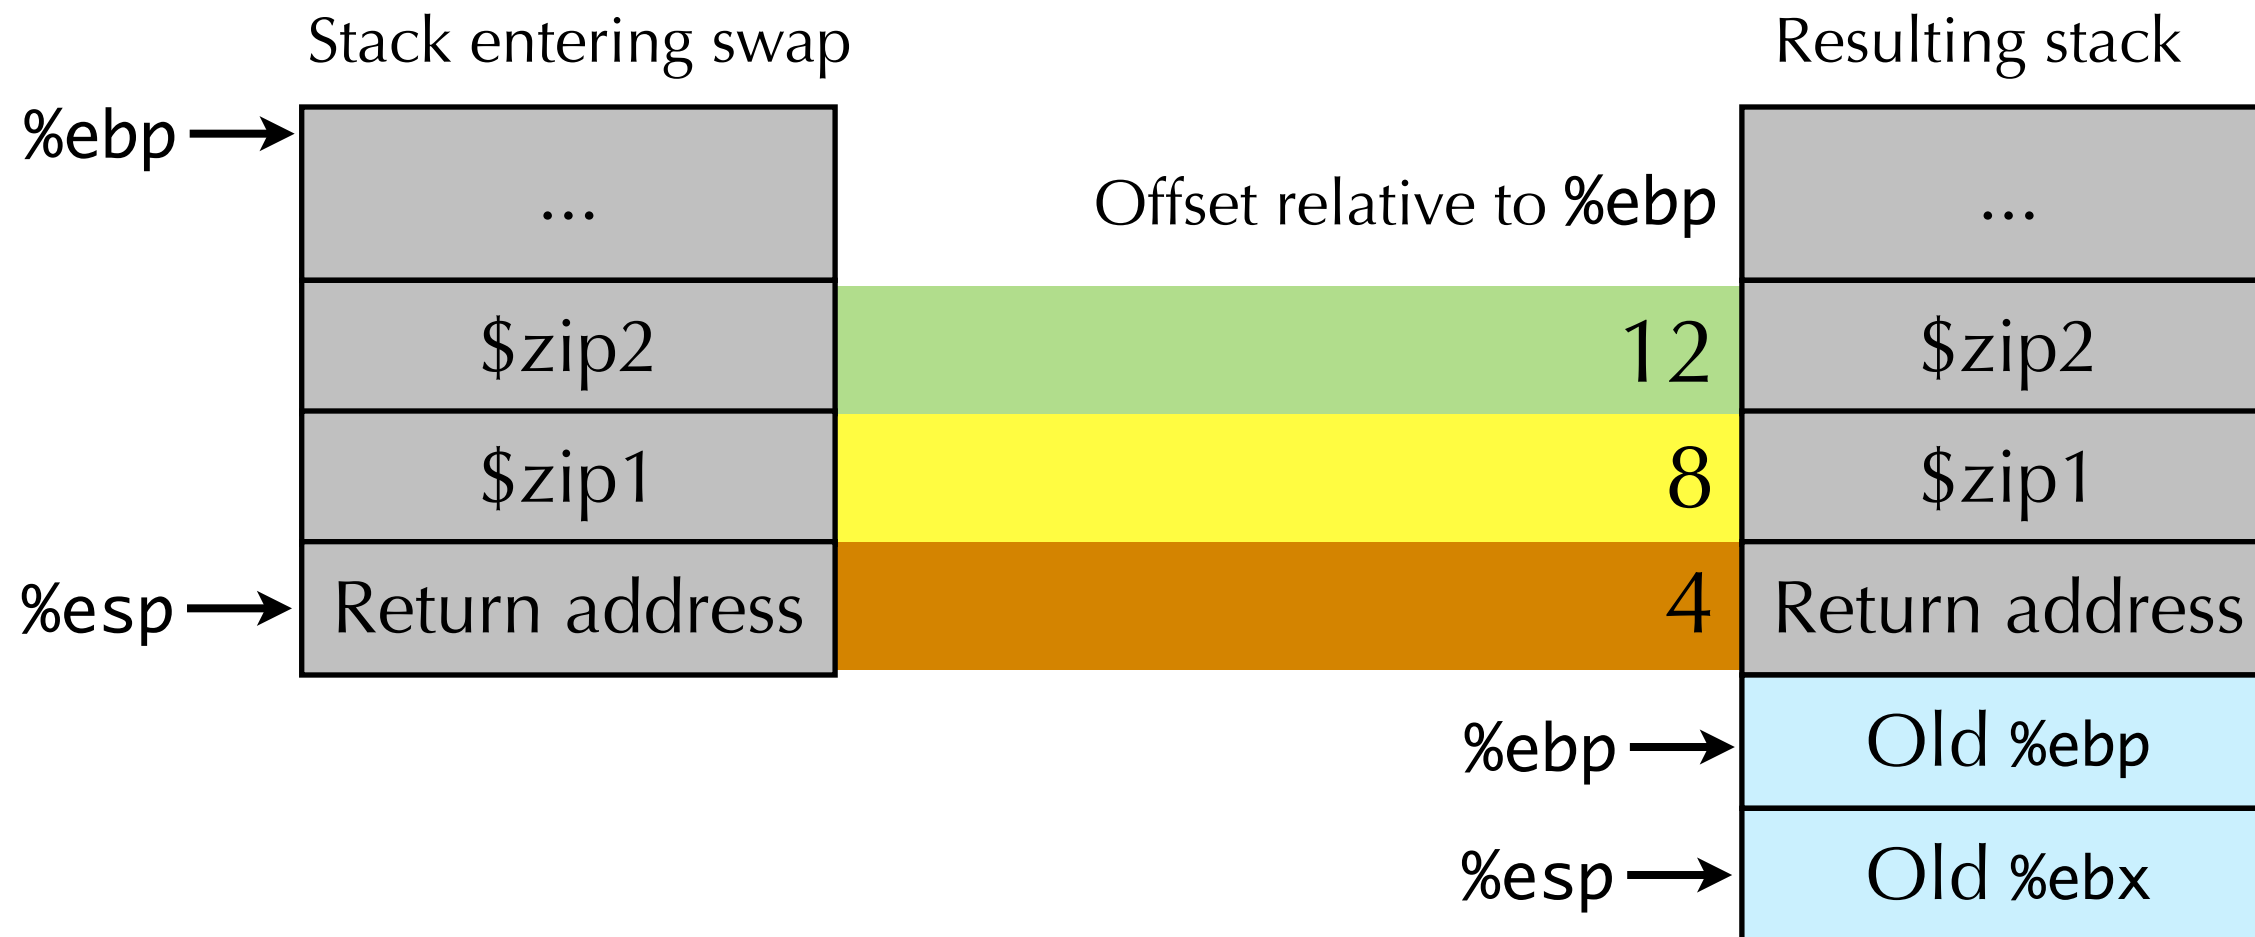
Resulting stack



```
pushl %ebp  
movl %esp,%ebp  
pushl %ebx
```

Set up

Swap body

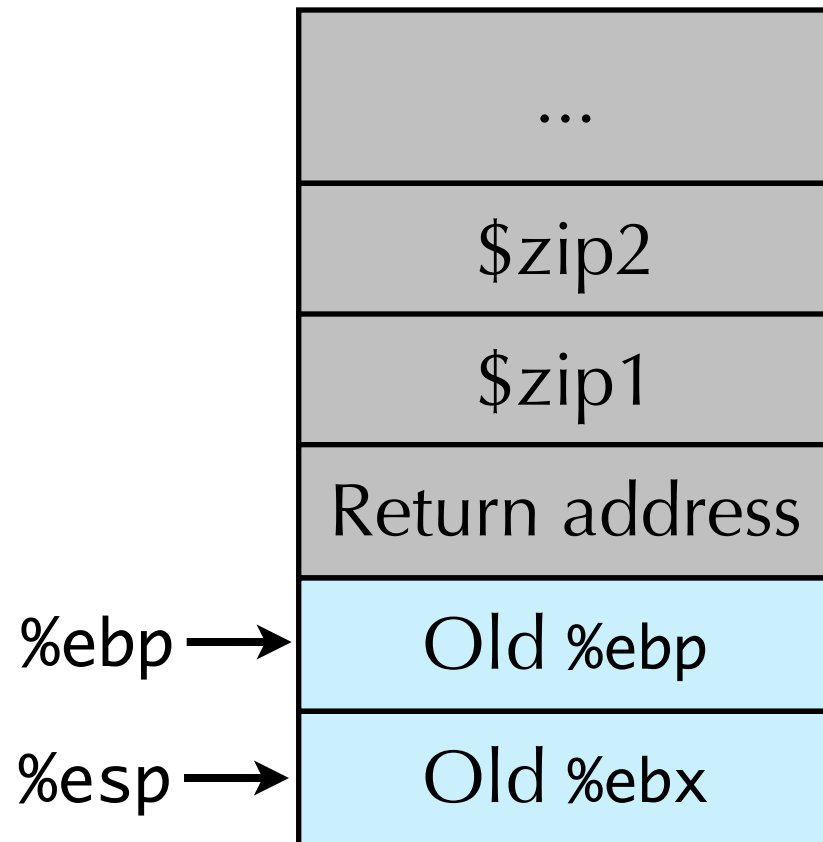


```
movl 12(%ebp),%ecx
movl 8(%ebp),%edx
movl (%ecx),%eax
movl (%edx),%ebx
movl %eax,(%edx)
movl %ebx,(%ecx)
```

Body

Swap finish

Stack at end swap body




```
movl -4(%ebp),%ebx  
movl %ebp,%esp  
popl %ebp  
ret
```

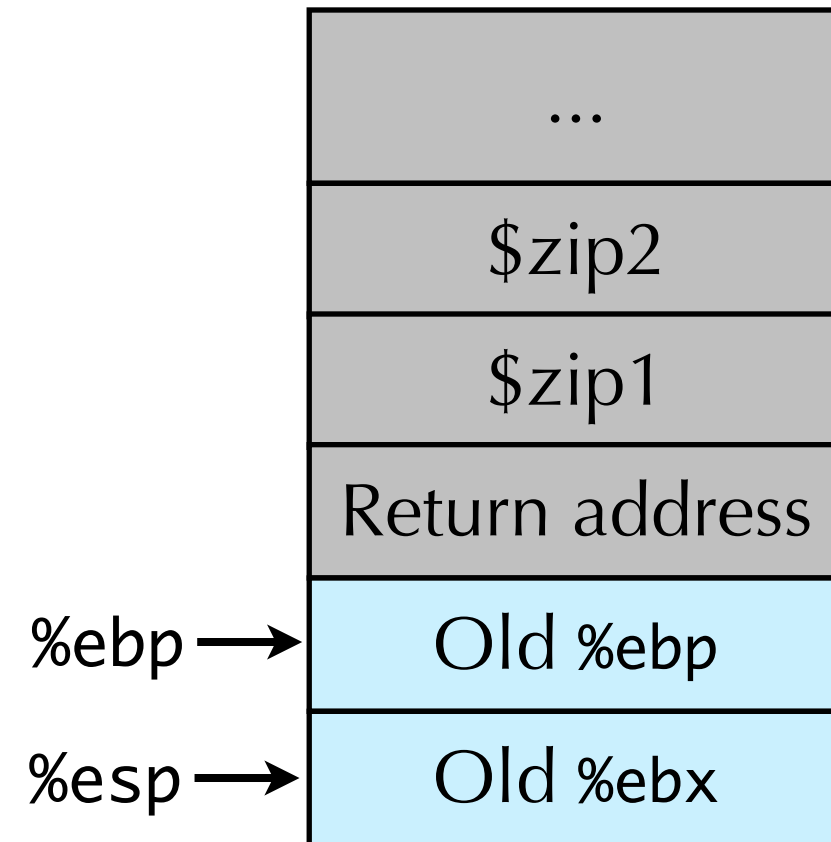
Finish

Swap finish

Stack at end swap body



Resulting stack

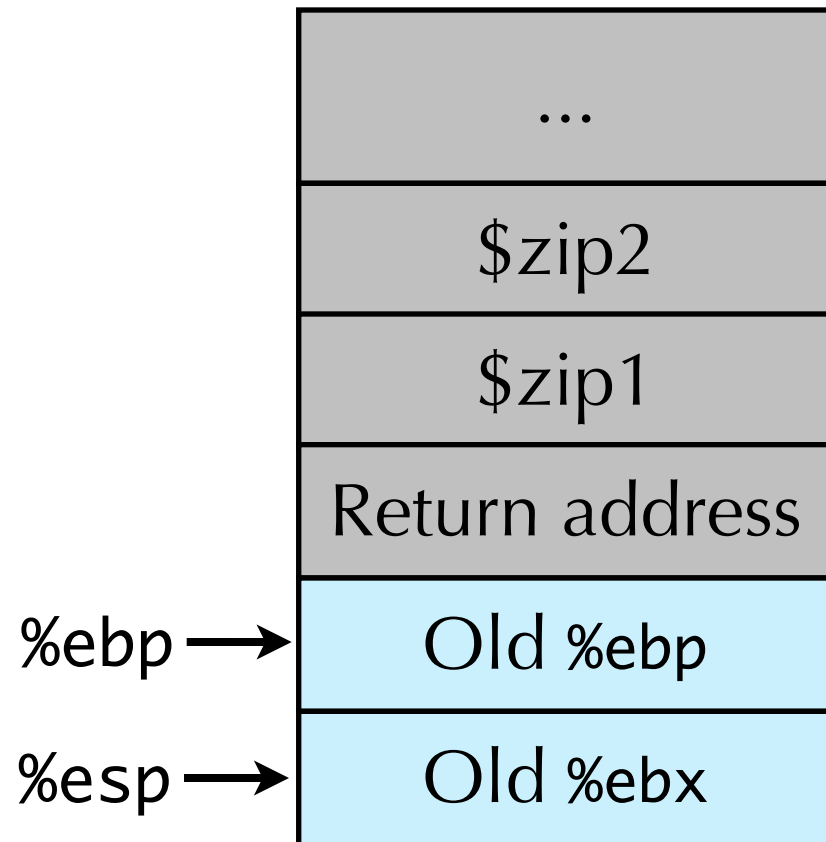


```
movl -4(%ebp),%ebx  
movl %ebp,%esp  
popl %ebp  
ret
```

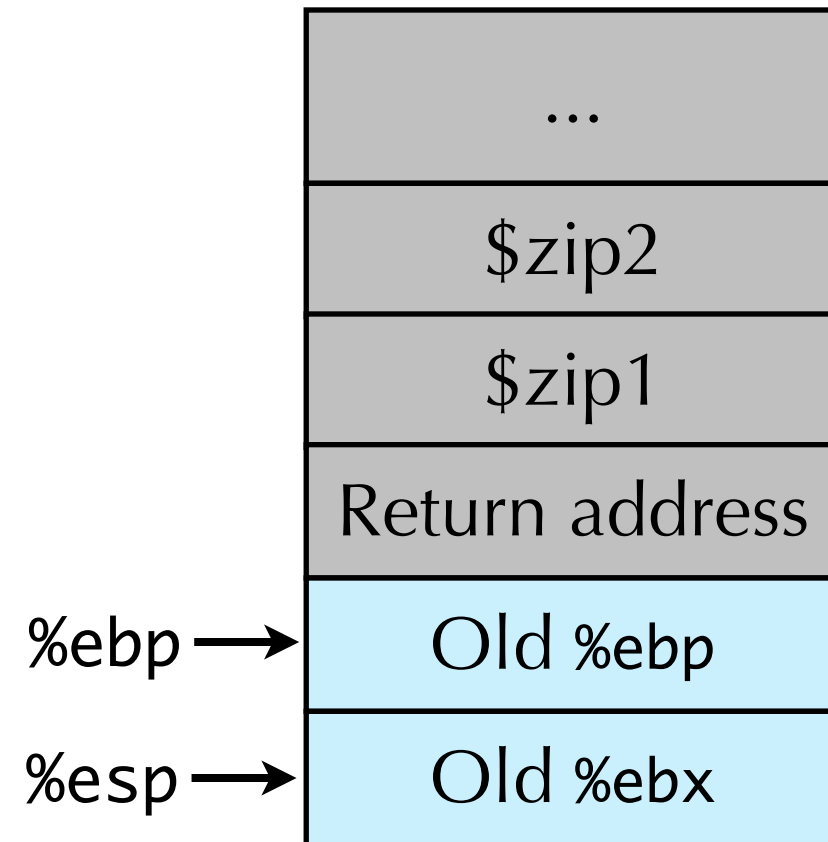
Finish

Swap finish

Stack at end swap body



Resulting stack



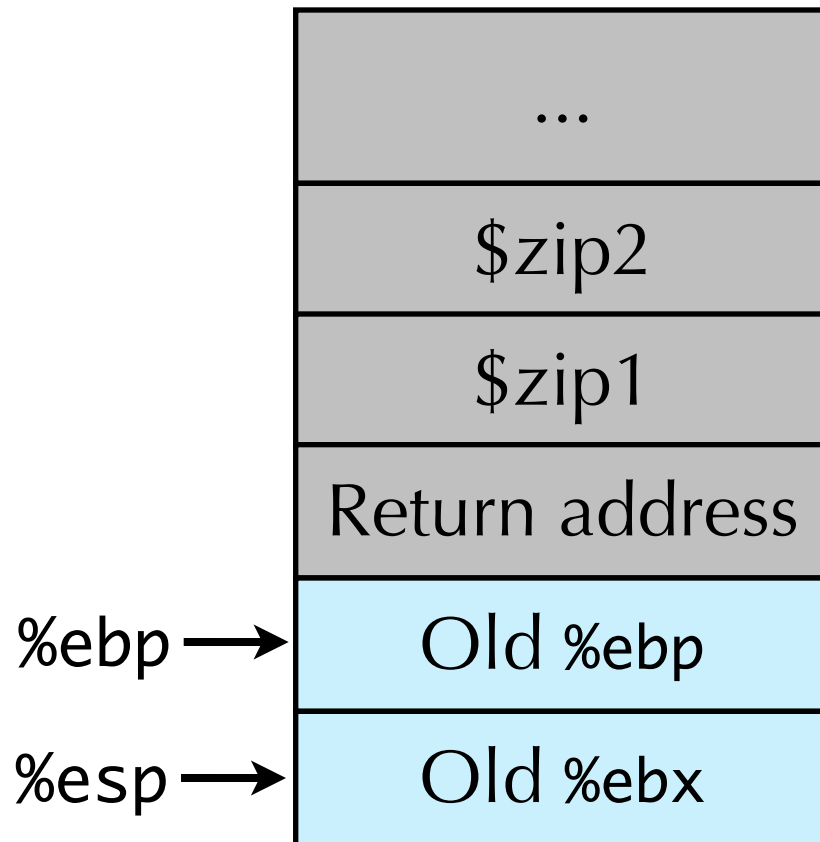
```
movl -4(%ebp),%ebx  
movl %ebp,%esp  
popl %ebp  
ret
```

Finish


Restores old value of %ebx!

Swap finish

Stack at end swap body



Resulting stack




```
movl -4(%ebp),%ebx  
movl %ebp,%esp  
popl %ebp  
ret
```


Finish

Swap finish

Stack at end swap body



Resulting stack




```
movl -4(%ebp),%ebx  
movl %ebp,%esp  
popl %ebp  
ret
```


Finish

Swap finish

Stack at end swap body



Resulting stack



```
movl -4(%ebp),%ebx  
movl %ebp,%esp  
popl %ebp  
ret
```

Finish

leave instruction

- Actual disassembly of swap

080483a4 <swap>:

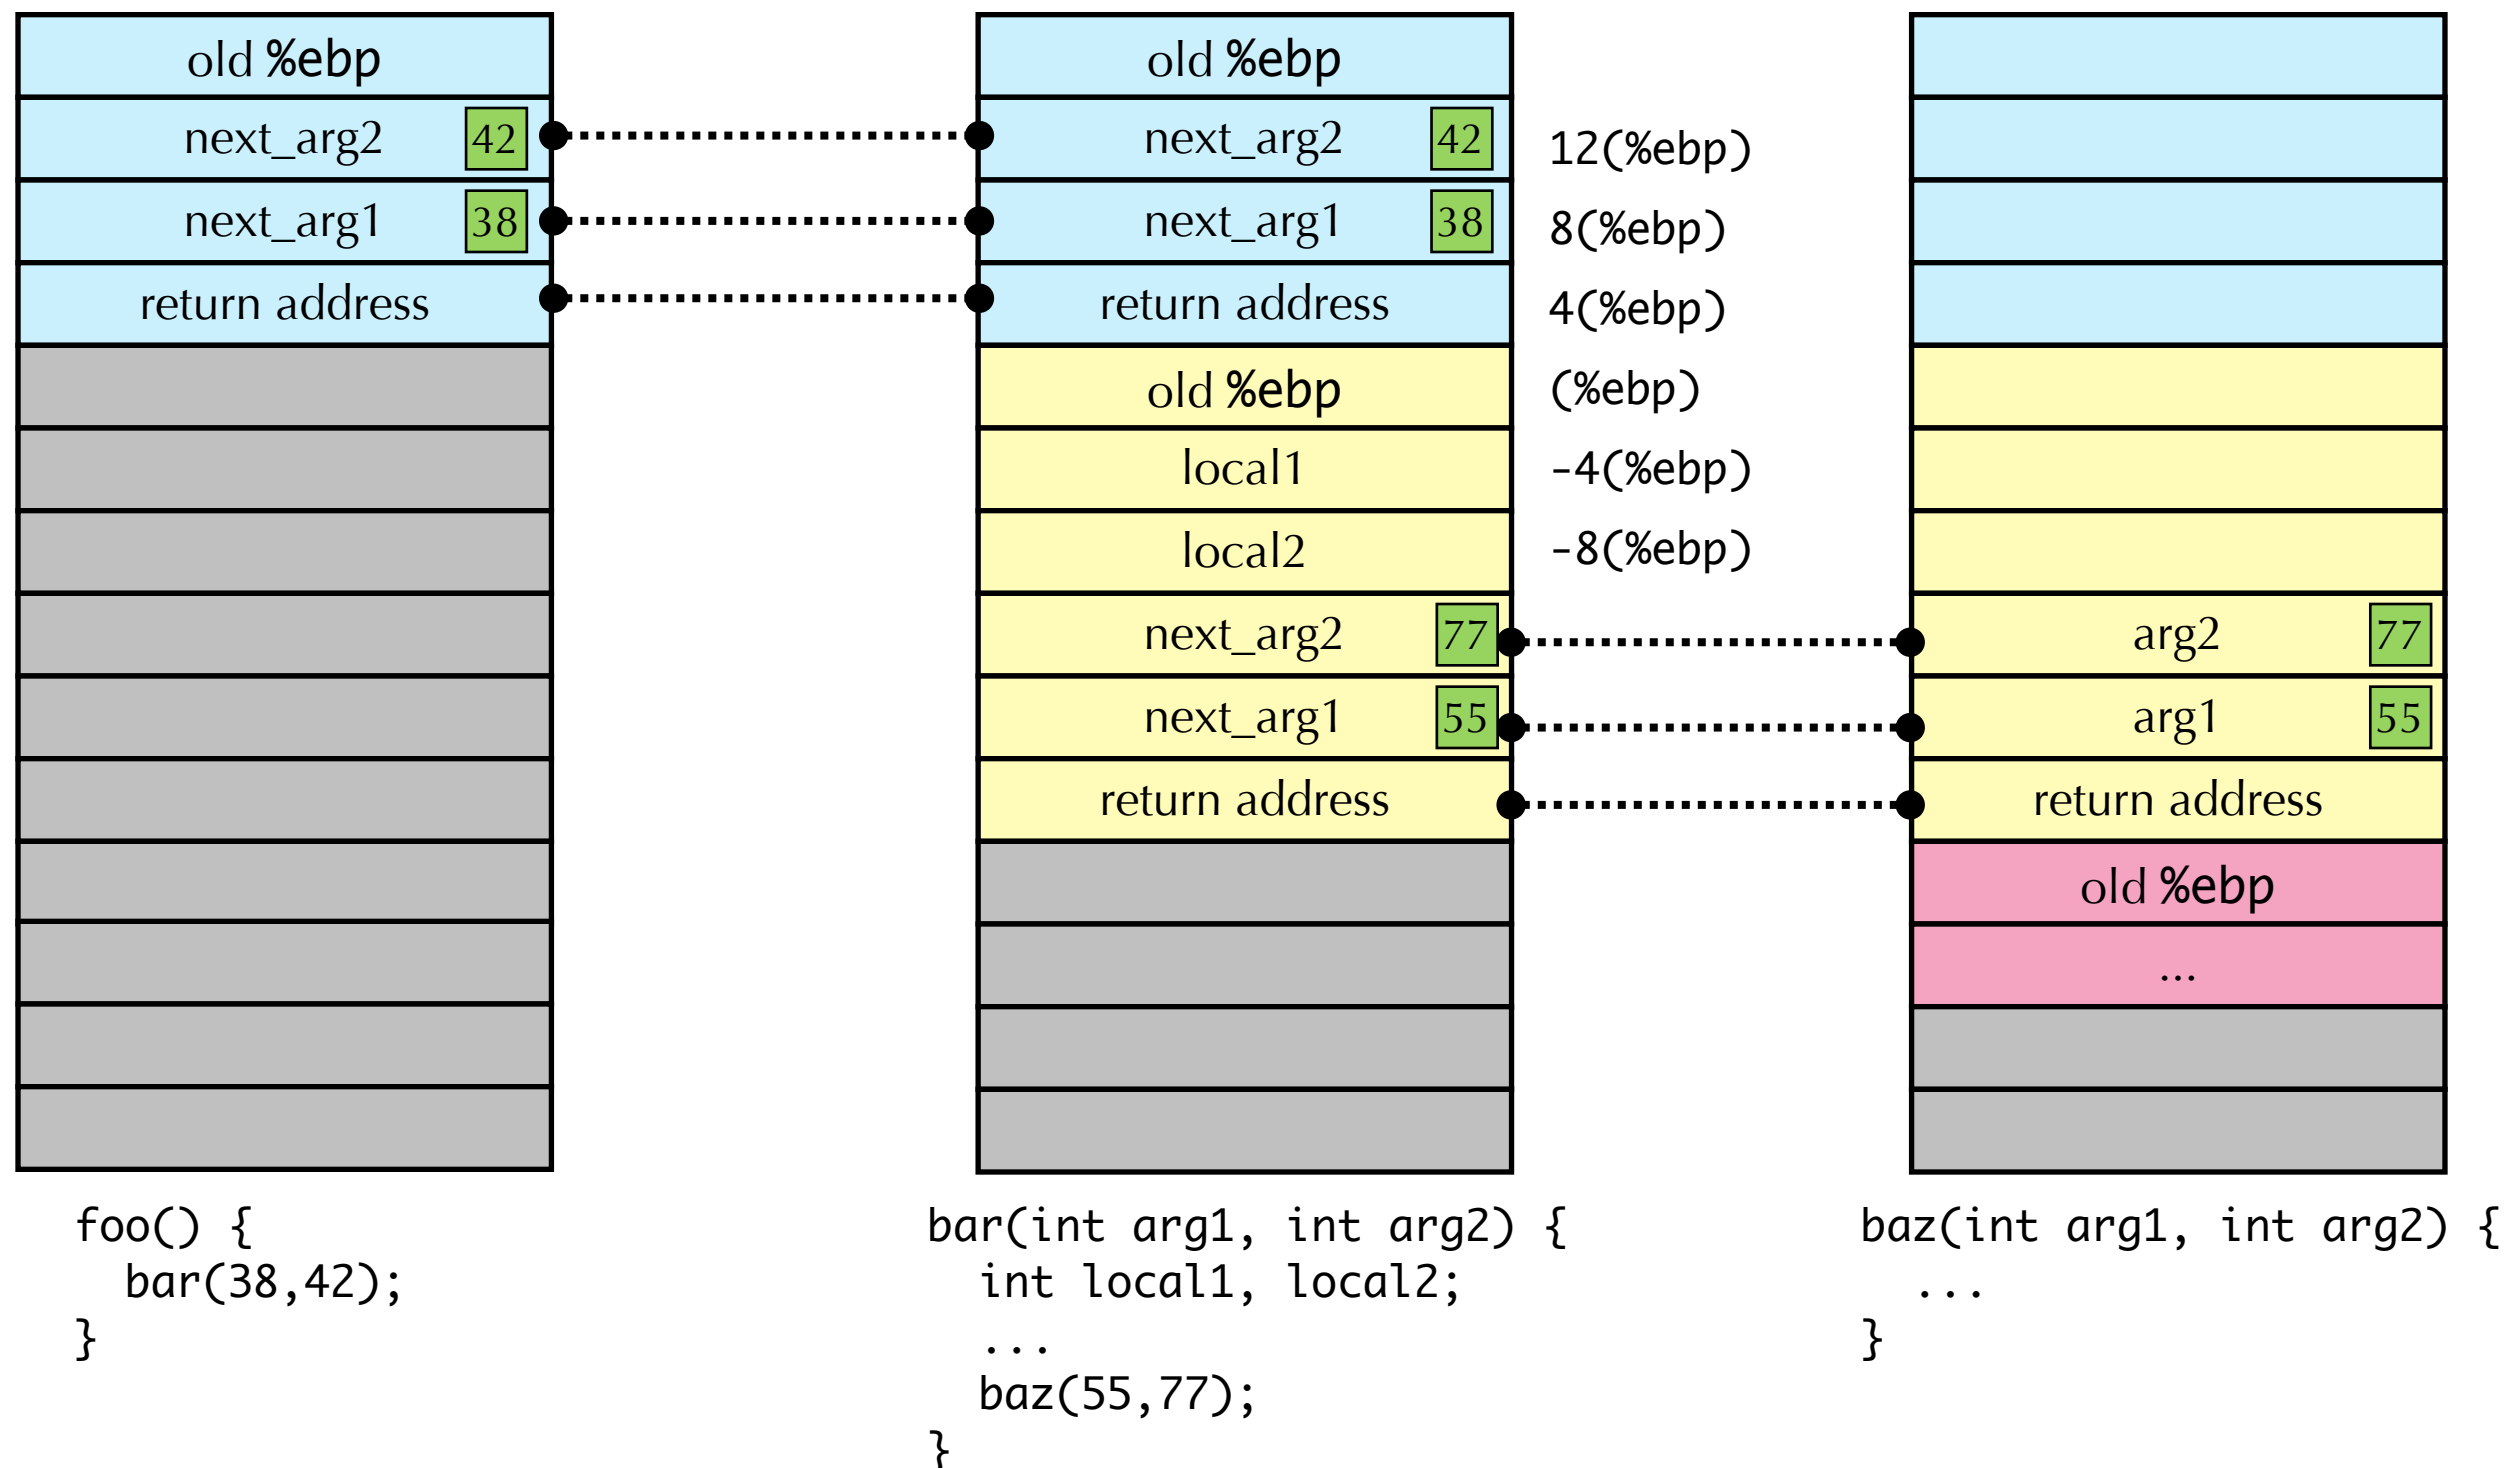
80483a4:	55	push	%ebp
80483a5:	89 e5	mov	%esp,%ebp
80483a7:	53	push	%ebx
80483a8:	8b 55 08	mov	0x8(%ebp),%edx
80483ab:	8b 4d 0c	mov	0xc(%ebp),%ecx
80483ae:	8b 1a	mov	(%edx),%ebx
80483b0:	8b 01	mov	(%ecx),%eax
80483b2:	89 02	mov	%eax,(%edx)
80483b4:	89 19	mov	%ebx,(%ecx)
80483b6:	5b	pop	%ebx
80483b7:	c9	leave	
80483b8:	c3	ret	

```
movl -4(%ebp),%ebx
movl %ebp,%esp
popl %ebp
ret
```

- leave prepares the stack for returning
- leave is equivalent to

```
movl %ebp,%esp
popl %ebp
```

Stack frame cheat sheet



Return values

- By convention, the compiler leaves return value in `%eax`

```
int absdiff(int x, int y)
{
    int result;
    if (x > y) {
        result = x-y;
    } else {
        result = y-x;
    }
    return result;
}
```

```
absdiff:
    pushl    %ebp
    movl     %esp, %ebp
    movl     8(%ebp), %edx
    movl     12(%ebp), %eax
    cmpl     %eax, %edx
    jle      .L7
    subl     %eax, %edx
    movl     %edx, %eax
.L8:
    leave
    ret
.L7:
    subl     %edx, %eax
    jmp      .L8
```

Return values

- By convention, the compiler leaves return value in %eax

```
int logical(int x, int y)
{
    int t1 = x^y;
    int t2 = t1 >> 17;
    int mask = (1<<13) - 7;
    int rval = t2 & mask;
    return rval;
}
```

```
logical:
    pushl %ebp
    movl %esp,%ebp

    movl 8(%ebp),%eax
    xorl 12(%ebp),%eax
    sarl $17,%eax
    andl $8185,%eax

    movl %ebp,%esp
    popl %ebp
    ret
```

- Works fine for 32-bit values
- For floating point values: other registers used
- For structs: return value is left on stack, caller must copy data elsewhere
 - Why must caller copy the data?

Register saving conventions

- When procedure `foo()` calls `bar()`
`foo()` is the **caller**, `bar()` is the **callee**
- Suppose `bar()` needs to modify some registers when it run
 - But `foo()` is using some of the same registers for its own purposes

```
foo:
    ...
    movl $2138, %edx
    call bar
    addl %edx, %eax
    ...
    ret
```

```
bar:
    ...
    movl 8(%ebp), %edx
    addl $14850, %edx
    ...
    ret
```


- Contents of `%edx` clobbered by `bar()`!

Register saving conventions

- Need to save some of the clobbered registers on the stack.
- Who saves the registers? The caller? The callee?
 - **Caller save:** caller saves registers in its stack frame before call
 - **Callee save:** callee saves registers it will clobber in its stack frame, and restores them before return
- What are advantages and disadvantages of each?
 - Caller save: caller must be conservative and save everything, since it doesn't know what callee will clobber.
 - Callee save: callee must be conservative and save everything, since it doesn't know what caller wants preserved.

x86/Linux register conventions

- x86/Linux uses a mixture of caller-save and callee-save!
- Three registers managed as caller-save
 - %eax, %ecx, %edx
- Three registers managed as callee-save
 - %ebx, %esi, %edi
- Frame and stack registers managed specially
 - %esp, %ebp



Procedures summary

- The stack makes function calls work!
 - Private storage for each invocation of a procedure call
 - Multiple function invocations don't clobber each other
 - Addressing of local variables and arguments is relative to stack frame `%ebp`
 - Recursion works too
 - Requires that procedures return in order of invocations (nesting is preserved)
- Procedures implemented using a combination of **hardware support** plus **software conventions**
 - Hardware support: `call`, `ret`, `leave`, `pushl`, `popl`
 - Software conventions: Register saving conventions, managing `%esp` and `%ebp`, managing layout of stack
 - Software conventions defined by the OS and the compiler.
 - No guarantee it will be the same on a different software platform.

Today

- Procedures
 - The stack
 - Stack frames
 - Leave
 - Register conventions
- x86_64

x86-64

- x86 (aka IA32) instruction set defined in about 1985
 - Has been dominant instruction format for many years
- x86-64 extends x86 to 64 bits
 - Originally developed by AMD (Advanced Micro Devices), Intel's competitor
 - Intel originally introduced Itanium (aka IA-64), a 64-bit ISA that was not backwards compatible. Not commercially successful.
 - Also referred to as AMD64, Intel64, and x64
- Currently in transition from 32 bits to 64 bits
 - Most new machines you buy will be 64 bits

Differences between x86 and x86-64

- Data types

C declaration	Intel data type	Assembly code suffix	32-bit	64-bit
char	Byte	b	1	1
short int	Word	w	2	2
int	Double word	l	4	4
long int	Quad word	q	4	8
long long int	Quad word	q	8	8
char *	Quad word	q	4	8
float	Single precision	s	4	4
double	Double precision	d	8	8
long double	Extended precision	t	10/12	10/16

Differences between x86 and x86-64

- Registers

- x86 has 8 registers
- x86-64 has 16 registers
 - Each is 64 bits
 - Extend existing registers and add new ones
 - Make `%ebp/%rbp` general purpose

<code>%rax</code>	<code>%eax</code>
<code>%rbx</code>	<code>%ebx</code>
<code>%rcx</code>	<code>%ecx</code>
<code>%rdx</code>	<code>%edx</code>
<code>%rsi</code>	<code>%esi</code>
<code>%rdi</code>	<code>%edi</code>
<code>%rsp</code>	<code>%esp</code>
<code>%rbp</code>	<code>%ebp</code>

<code>%r8</code>	<code>%r8d</code>
<code>%r9</code>	<code>%r9d</code>
<code>%r10</code>	<code>%r10d</code>
<code>%r11</code>	<code>%r11d</code>
<code>%r12</code>	<code>%r12d</code>
<code>%r13</code>	<code>%r13d</code>
<code>%r14</code>	<code>%r14d</code>
<code>%r15</code>	<code>%r15d</code>

x86-64 instructions

- Long word **l** (4 Bytes) \leftrightarrow Quad word **q** (8 Bytes)
- New instructions:
 - `movl` \rightarrow `movq` `addl` \rightarrow `addq` `sall` \rightarrow `sallq` etc.
- 32-bit instructions generate 32-bit results
 - Set higher order bits of destination register to 0
 - E.g., `addl`
- `gcc` makes more efficient use of x86-64 instructions
 - E.g., more extensive use of conditional move operation
 - `gcc -m32` will produce 32-bit code

Procedure calls

- Up to six (integral) arguments can be passed in registers
 - Instead of on stack
 - `%rdi, %rsi, %rdx, %rcx, %r8, %r9`
- Some procedures do not need a stack frame at all!
 - Few arguments,
few local variables,
no local arrays or structs,
no need to take address of local variables,
no need to pass arguments on stack to another function,
⇒ no need for stack frame
 - Can result in very low overhead for some function calls!

Stack frames

- No frame pointer!
 - x86_64 makes `%rbp/%ebp` general purpose
- Instead, procedures subtract a constant from stack pointer (`%rsp`) at beginning, add constant at procedure return
 - Accesses all stack elements via offsets from `%rsp`
 - No need for `%rbp`
- Stack frame size is constant during procedure call
 - Stack pointer does not fluctuate as in IA32
 - i.e., through pushes and pops

Next week

- Structures and arrays
- Buffer overruns
- Assignment 2 due Thursday