



HARVARD

School of Engineering
and Applied Sciences

Processor Architecture

CS61, Lecture 8

Prof. Stephen Chong

September 27, 2011

Announcements

- HW 3 released
 - Please check you have a directory “hw3” in your CS 61 home directory
 - Contact course staff if you don't
- Infrastructure
 - Please try not to create more than one VM at a time
- Interested in going to grad school in CS?
 - Panel tonight: 6pm in MD 119
 - Pizza

Today


- Processor architecture
 - Logic gates
 - Adders and multiplexors
 - Registers
 - Instruction set encoding
 - A sequential processor
 - Pipelining
 - CISC vs RISC

Logic gates

- Electronic circuits used to compute functions on bits, and store bits
 - Typical technology:
 - Logical value 1 represented by high voltage ($\sim 1.0\text{V}$)
 - Logical value 0 represented by low voltage ($\sim 0.0\text{V}$)
- **Logic gates** are the basic computing elements of digital circuits
 - Implement Boolean functions
 - Can be implemented with transistors, electro-mechanical relays, optical components, ...


Logic gates

AND gate



a	b	out
0	0	0
0	1	0
1	0	0
1	1	1


OR gate



a	b	out
0	0	0
0	1	1
1	0	1
1	1	1


Logic gates

NOT gate



a	out
0	1
1	0


XOR gate



a	b	out
0	0	0
0	1	1
1	0	1
1	1	0

- Other logic gates are possible
 - e.g., NAND, NOR

Logic gate example




- Logic gates are always active
 - Output always being computed from input
- But takes time for signal to propagate
 - If input changes, will take some time before output correctly reflects input

A one-bit adder

- Suppose we are building a circuit to add two bits
 - $a_0 + b_0$
- How many outputs do we need?

$\begin{array}{r} 0 \\ + 0 \\ \hline 0 \end{array}$	$\begin{array}{r} 0 \\ + 1 \\ \hline 1 \end{array}$	$\begin{array}{r} 1 \\ + 0 \\ \hline 1 \end{array}$	$\begin{array}{r} 1 \\ + 1 \\ \hline 10 \end{array}$
---	---	---	--

- Two: one “output” bit, and one “carry” bit




A 4-bit adder

- Suppose we are adding two 4-bit numbers
 - $a_3a_2a_1a_0 + b_3b_2b_1b_0$
- We can construct a 4-bit adder using the 1-bit adder we just developed...


A 4-bit adder

- Suppose we are adding two 4-bit numbers




Multiplexors

- A **multiplexor** (or MUX) selects one of several data inputs based on a control input
- E.g., a single-bit multiplexor
 - If s is 0 then out will equal a
 - If s is 1 then out will equal b



4-way multiplexor

- Can have more than just two inputs...
 - If s_0 is 0 and s_1 is 0 then out will equal a
 - If s_0 is 0 and s_1 is 1 then out will equal b
 - If s_0 is 1 and s_1 is 0 then out will equal c
 - If s_0 is 1 and s_1 is 1 then out will equal d



Word-level multiplexors

- Can also choose an entire word, not just one bit
 - A, B, C, D, and OUT represent 32-bit words
 - s_0 and s_1 select one of the words



Memory



- Combinational circuits do not store information
 - They generate output based on their input
 - Do not have **state**
- Let's consider two classes of memory devices
 - **Clocked registers**
 - **Random-access memories**

Clocked registers

- **Clocked registers** (or **registers**) store individual bits or words
- A clock signal controls loading value into a register
- Note: slightly different meaning of word *register*
 - In hardware: *register* is storage directly connected to rest of circuit by its input and output wires
 - In machine-level programming: *register* refers to small collection of addressable words in the CPU
 - Program registers can be implemented with hardware registers


Register

- Register output is the current register state
- When clock rises, values at register input is captured, and becomes new state



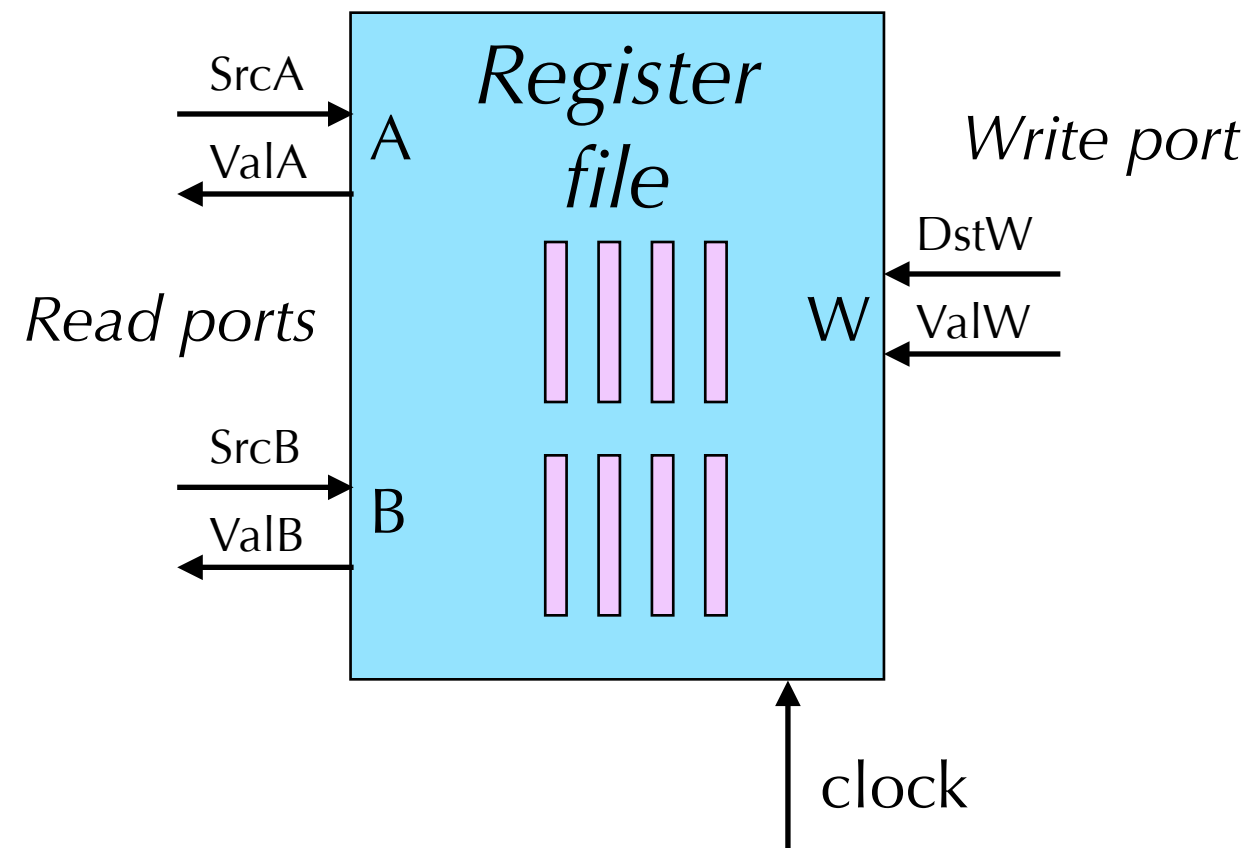
Register

- Register output is the current register state
- When clock rises, values at register input is captured, and becomes new state




Register file

- A **register file** has multiple registers
 - Can perform simultaneous reads and writes
- E.g.,
 - Two read ports A and B
 - SrcA selects one of the registers, the contents of which are output on ValA
 - E.g., SrcA set to 011 selects program register %ebx
 - One write port W
 - DstW selects one of the registers, the contents of which are overwritten with ValW on the clock rise



Data memory

- For the moment, let's assume a simple model for main memory
 - Specify an Address
 - If $\text{read?}=1$ then contents of address will be on DataOut
 - If $\text{write?}=1$ then contents of address will be set to DataIn on clock rise
- Memory is more complicated than this
 - We will see more about memory later in the course



How do they fit together?

- Suppose we encode instruction

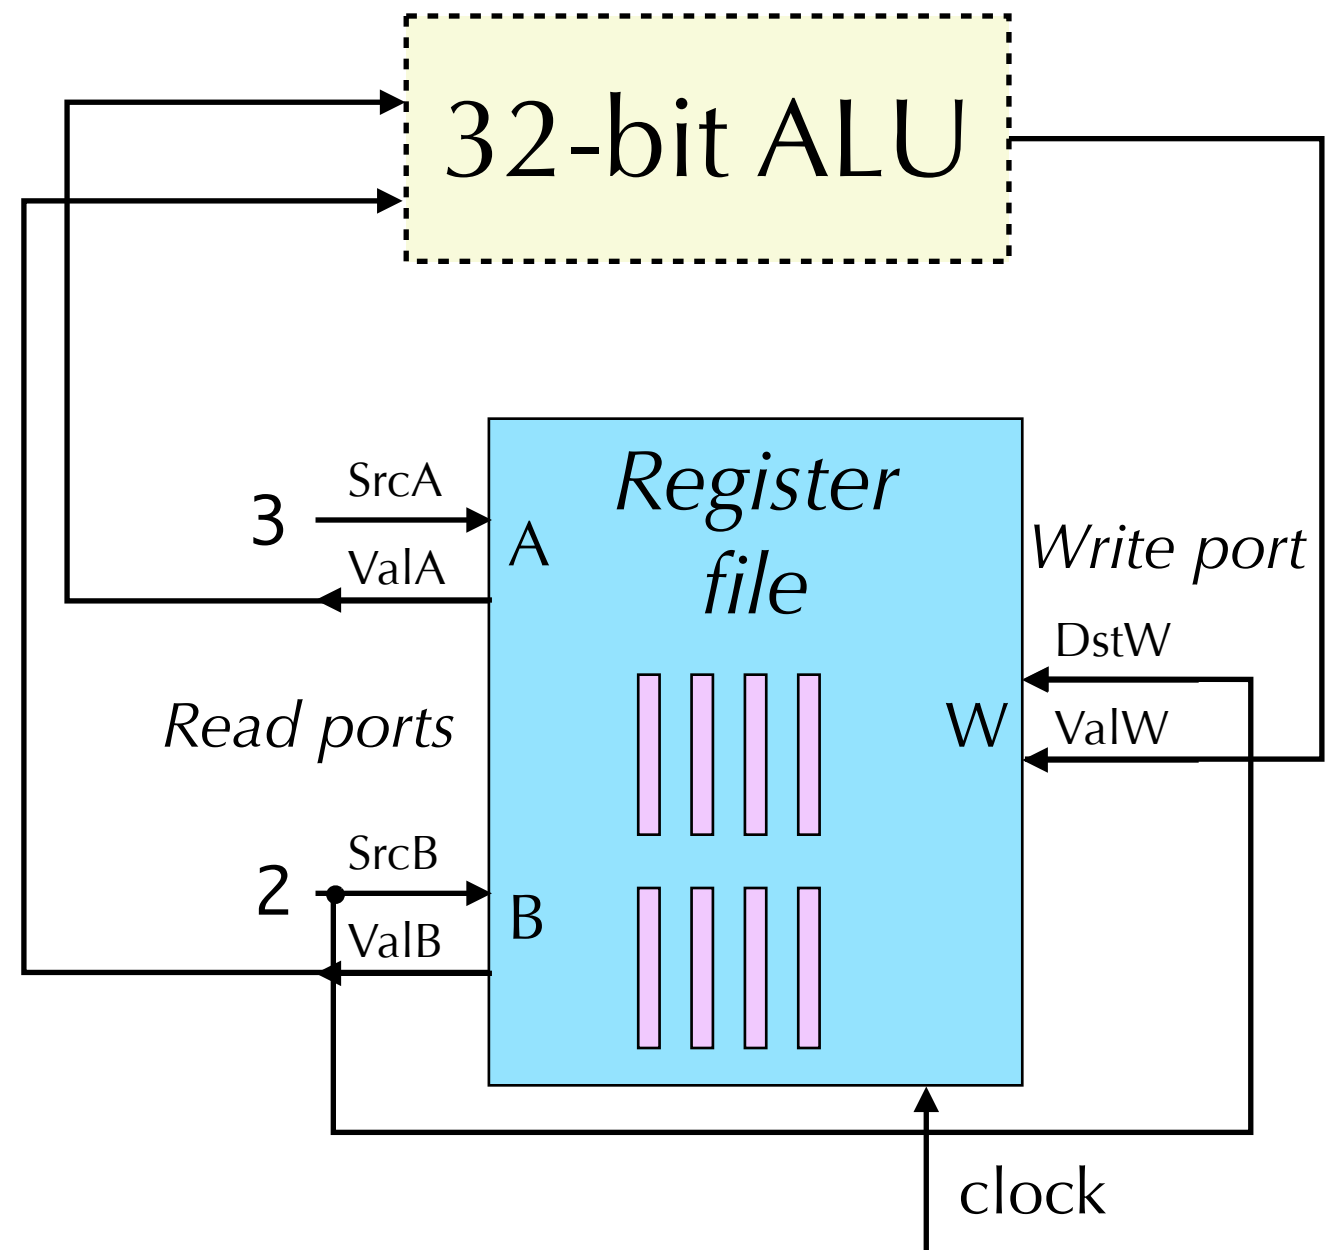
`add %ebx, %edx`

as bytes `0x6032`

Addition
instruction

`%ebx`

`%edx`



Today

- Processor architecture
 - Logic gates
 - Adders and multiplexors
 - Registers
 - Instruction set encoding
 - A sequential processor
 - Pipelining
 - CISC vs RISC

Instruction Set Architectures

- Many different instructions, even in a simple ISA
 - Move instructions
 - Different variants: to and from register, from memory to register, from register to memory
 - Arithmetic/logical instructions
 - add, sub, and, xor, ...
 - Control flow instructions
 - jmp, jz, jle, jl, ...
 - call, ret
 - ...
- What instructions are in the ISA influence how we can implement the ISA
- We'll consider a simple instruction set: Y86
 - Described in text book; a simplification of x86

Instruction encoding

- Machine instructions are encoded as bytes

```
% objdump -d scshell
```

```
...
```

8048404:	55	push	%ebp
8048405:	89 e5	mov	%esp,%ebp
8048407:	83 ec 14	sub	\$0x14,%esp
804840a:	c7 45 fc 00 00 00 00	movl	\$0x0,-0x4(%ebp)
8048411:	eb 2a	jmp	804843d
8048413:	8b 15 4c 97 04 08	mov	0x804974c,%edx
8048419:	8b 45 fc	mov	-0x4(%ebp),%eax
804841c:	8d 04 02	lea	(%edx,%eax,1),%eax
804841f:	0f b6 10	movzbl	(%eax),%edx
8048422:	8b 45 08	mov	0x8(%ebp),%eax
8048425:	0f b6 00	movzbl	(%eax),%eax
8048428:	38 c2	cmp	%al,%dl
804842a:	74 09	je	8048435

```
...
```

Instruction encoding

- Machine instructions are encoded as bytes
- How instructions are encoded affects how we implement an ISA
- We use a *variable length* encoding
 - Different instructions may take a different number of bytes
 - x86 is a variable length encoding
 - More compact representation than fixed length, but more complex to figure out the address of the next instruction

Instruction encoding

- Initial byte specifies instruction type
 - High-order nibble indicates code, and low-order nibble indicates function
 - E.g., 60 is addl, 61 is subl
- Registers are identified using 4 bits

Byte	0	1	2	3	4	5
halt	0	0				
nop	1	0				
rrmovl rA, rB	2	0	rA	rB		
irmovl V, rB	3	0	8	rB	V	
rmmovl rA, D(rB)	4	0	rA	rB	D	
mrmmovl D(rB), rA	5	0	rA	rB	D	
OpI rA, rB	6	fn	rA	rB		
jXX Dest	7	fn	Dest			
cmovXX rA, rB	2	fn	rA	rB		
call Dest	8	0	Dest			
ret	9	0				
pushl rA	A	0	rA	F		
popl rA	B	0	rA	F		

Sequential processor

- Let's first try a processor that executes one instruction each clock cycle
 - Requires a long clock cycle (why?) \Rightarrow inefficient
 - We will improve performance soon...
- Break each instruction up into 6 stages
 - Uniform structure simplifies implementation
 - Not every instruction will have each stage

Execution stages

Fetch

Read bytes of instruction from memory

Decode

Reads up to two operands from the register file

Execute

Perform arithmetic/logic operation
(includes computing effective address of memory reference, incrementing or decrementing stack pointer, testing condition codes for branch conditions, ...)

Memory

Read or write data to/from memory

Write back

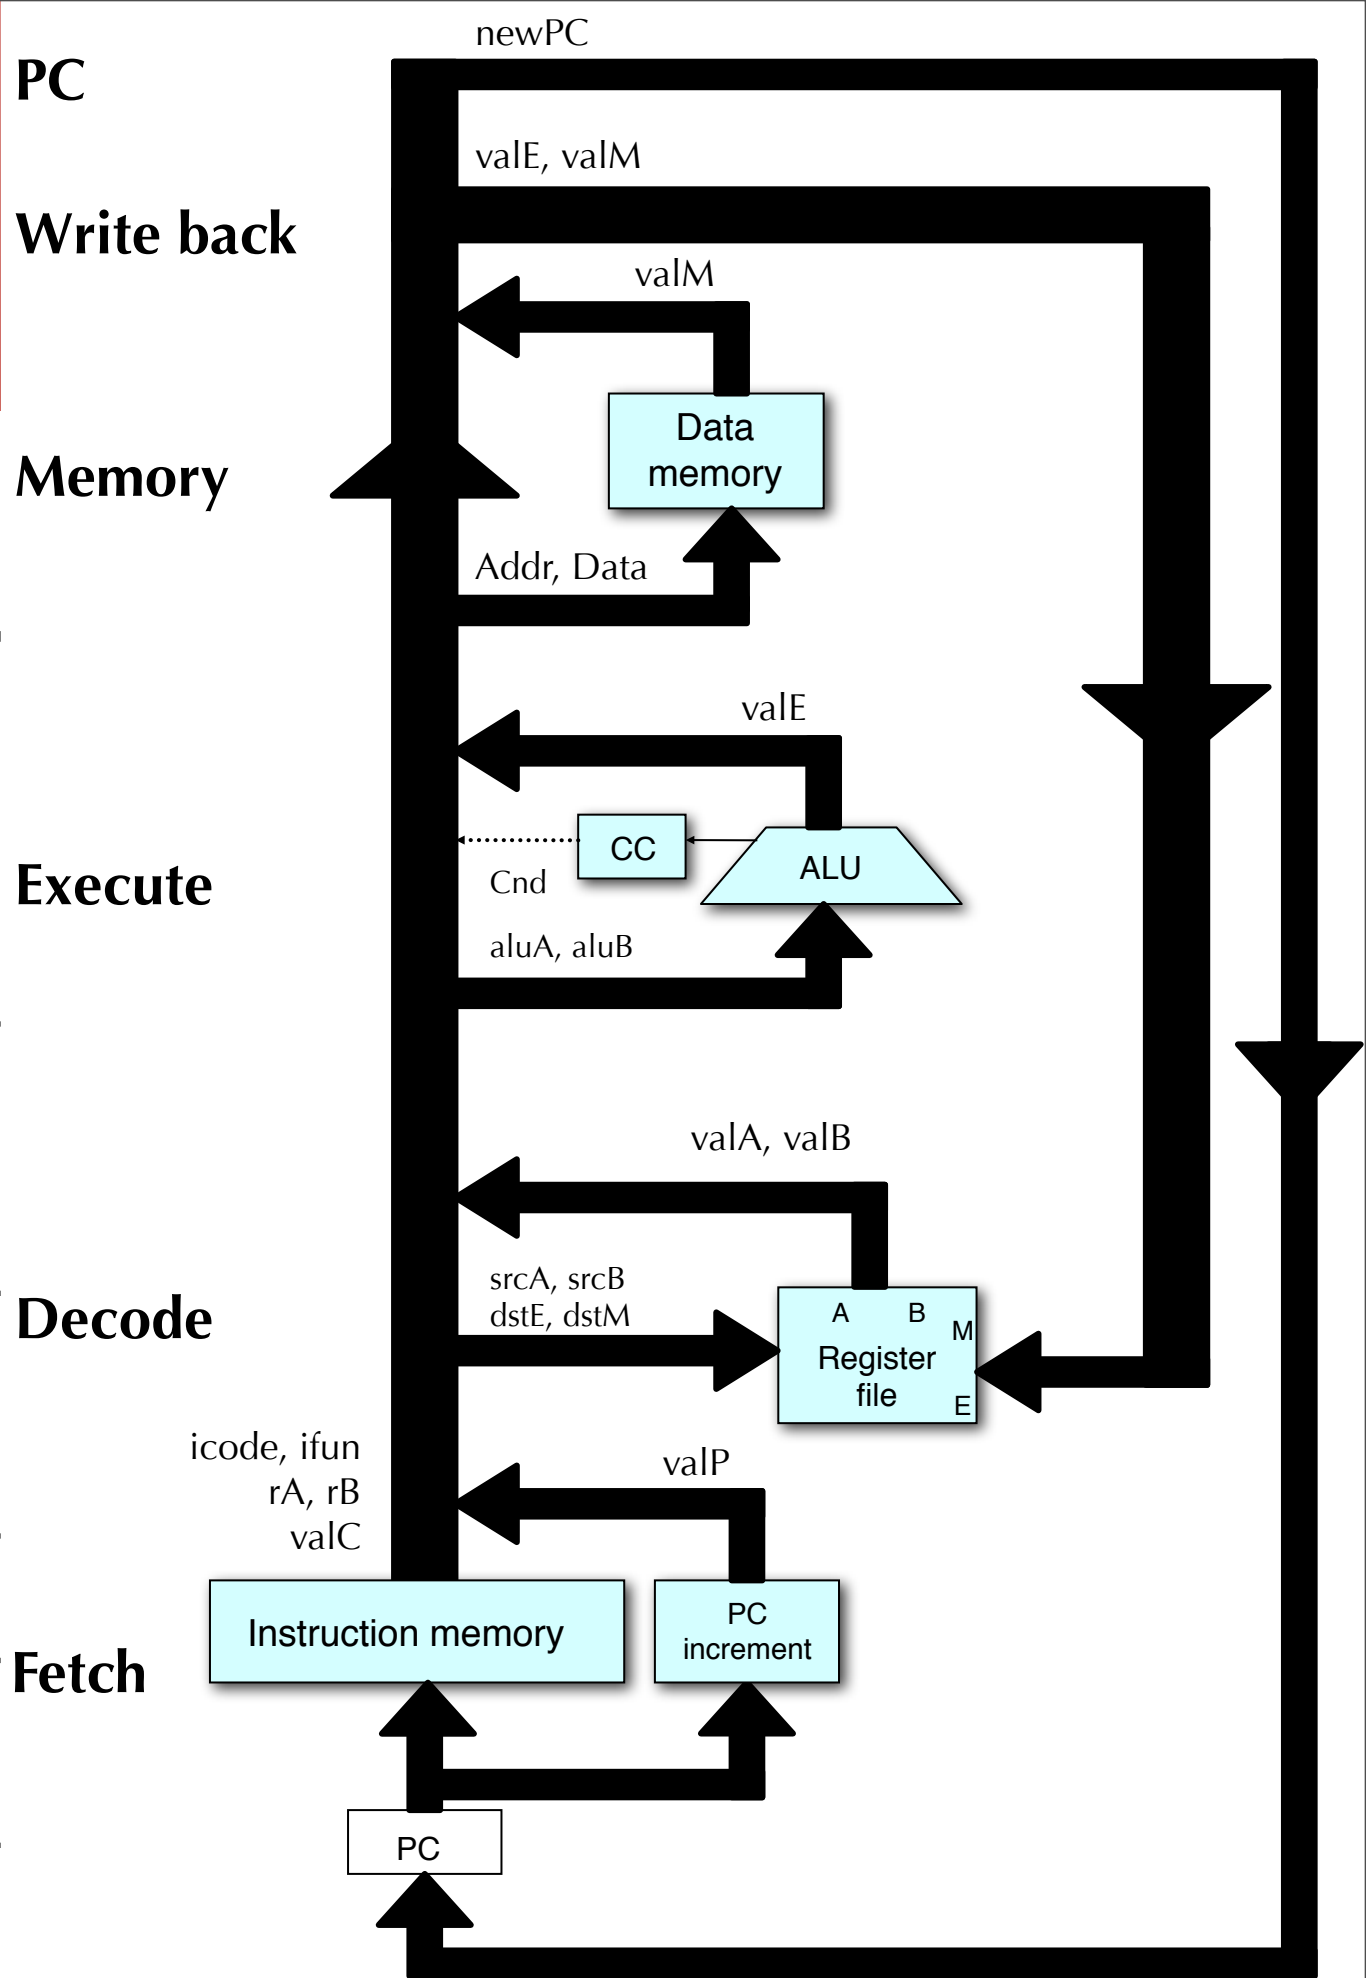
Writes up to two results to the register file

PC update

Set PC to address of next instruction

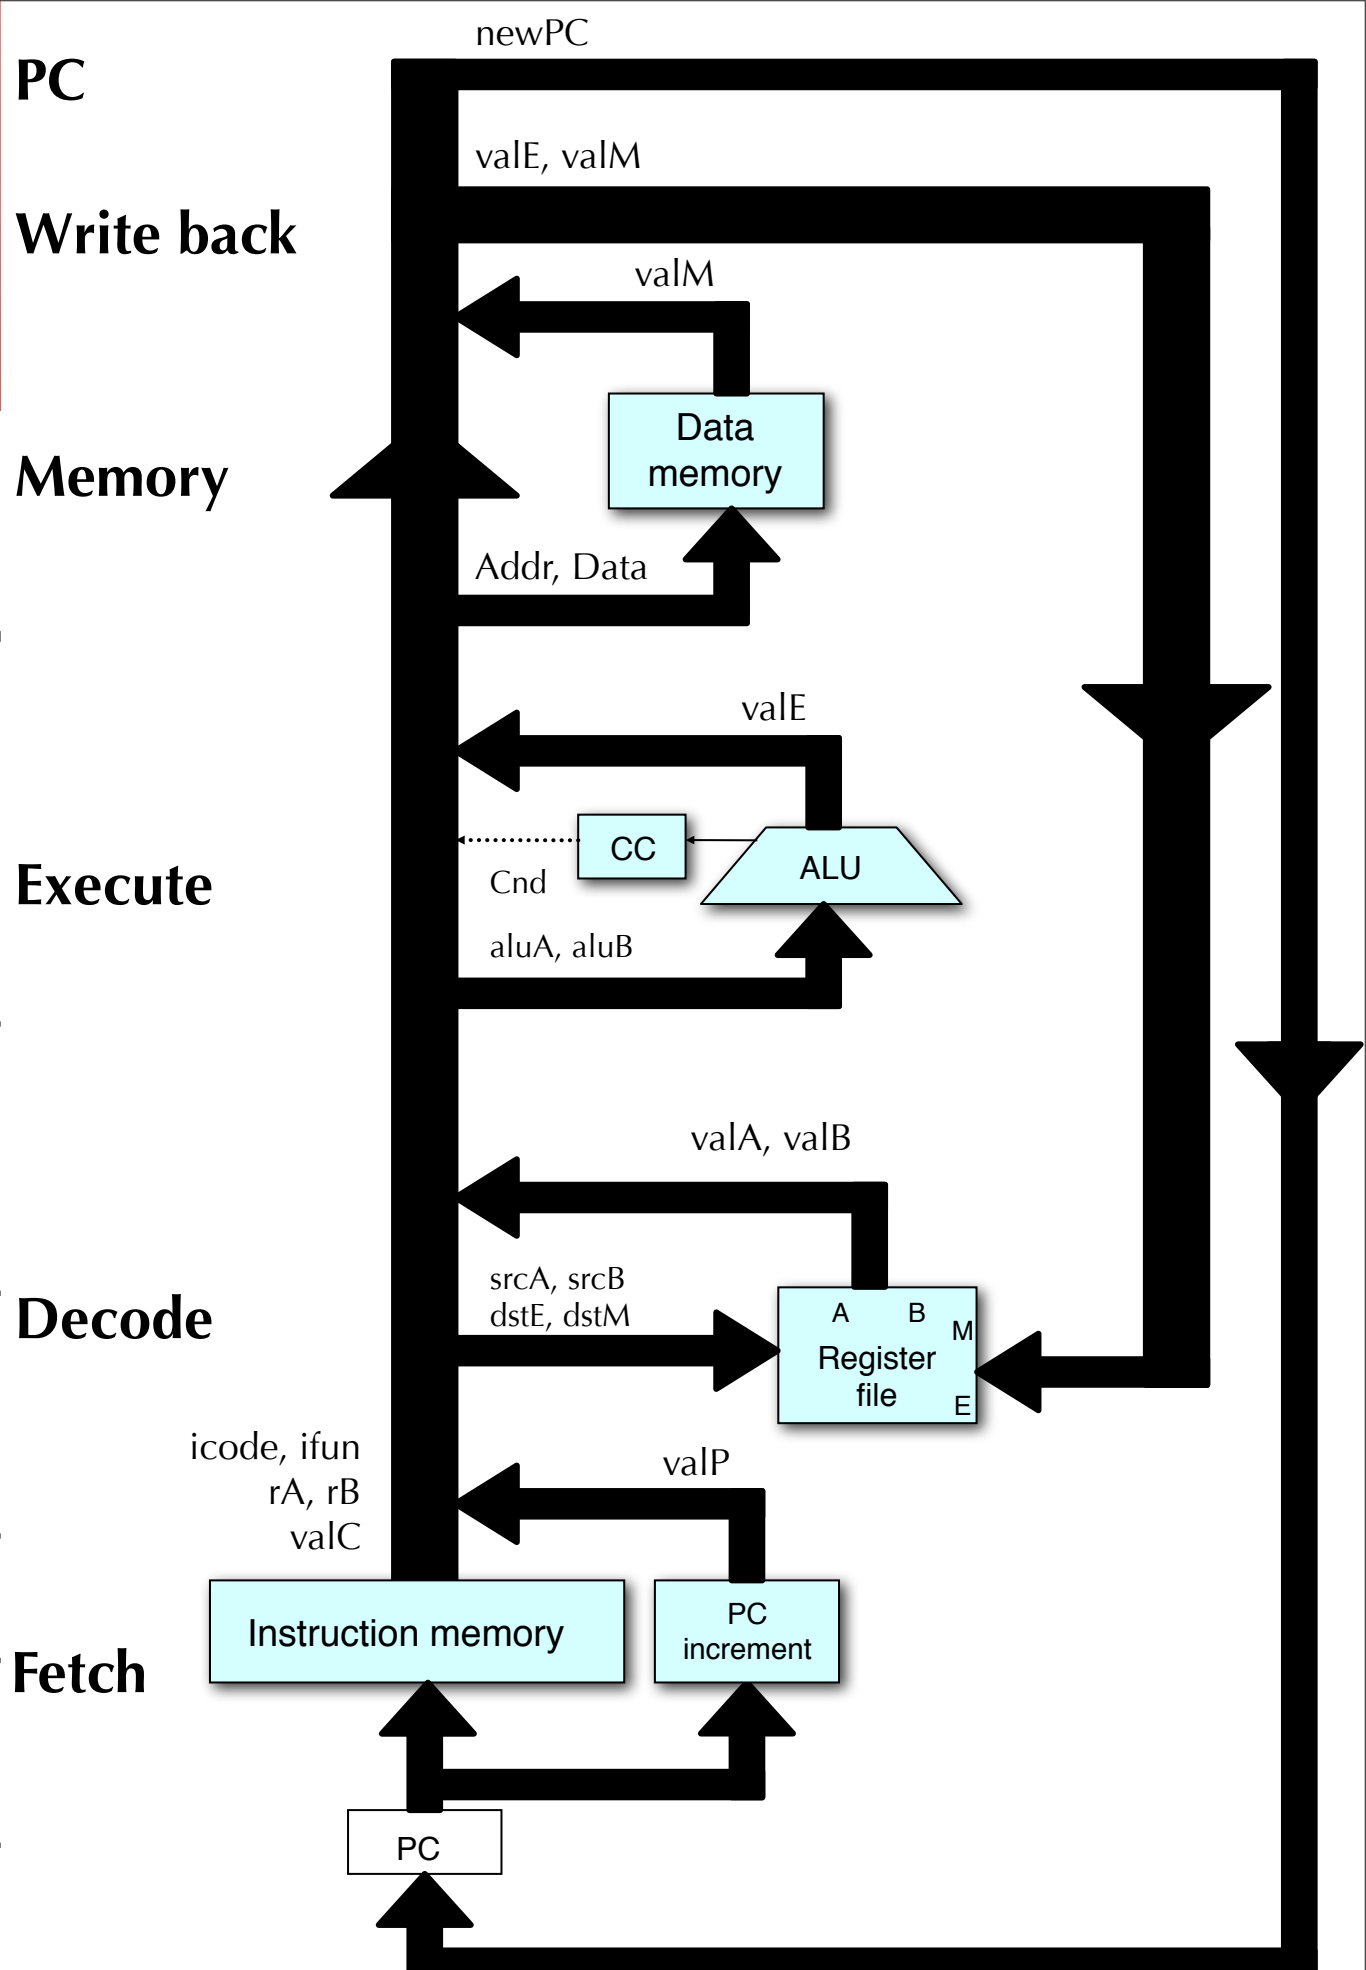
Execution stages

Stage	OPl rA, rB
Fetch	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$ $\text{rA:rB} \leftarrow M_1[\text{PC}+1]$ $\text{valP} \leftarrow \text{PC} + 2$
Decode	$\text{valA} \leftarrow R[\text{rA}]$ $\text{valB} \leftarrow R[\text{rB}]$
Execute	$\text{valE} \leftarrow \text{valB OP valA}$ Set CC
Memory	–
Write back	$R[\text{rB}] \leftarrow \text{valE}$
PC update	$\text{PC} \leftarrow \text{valP}$



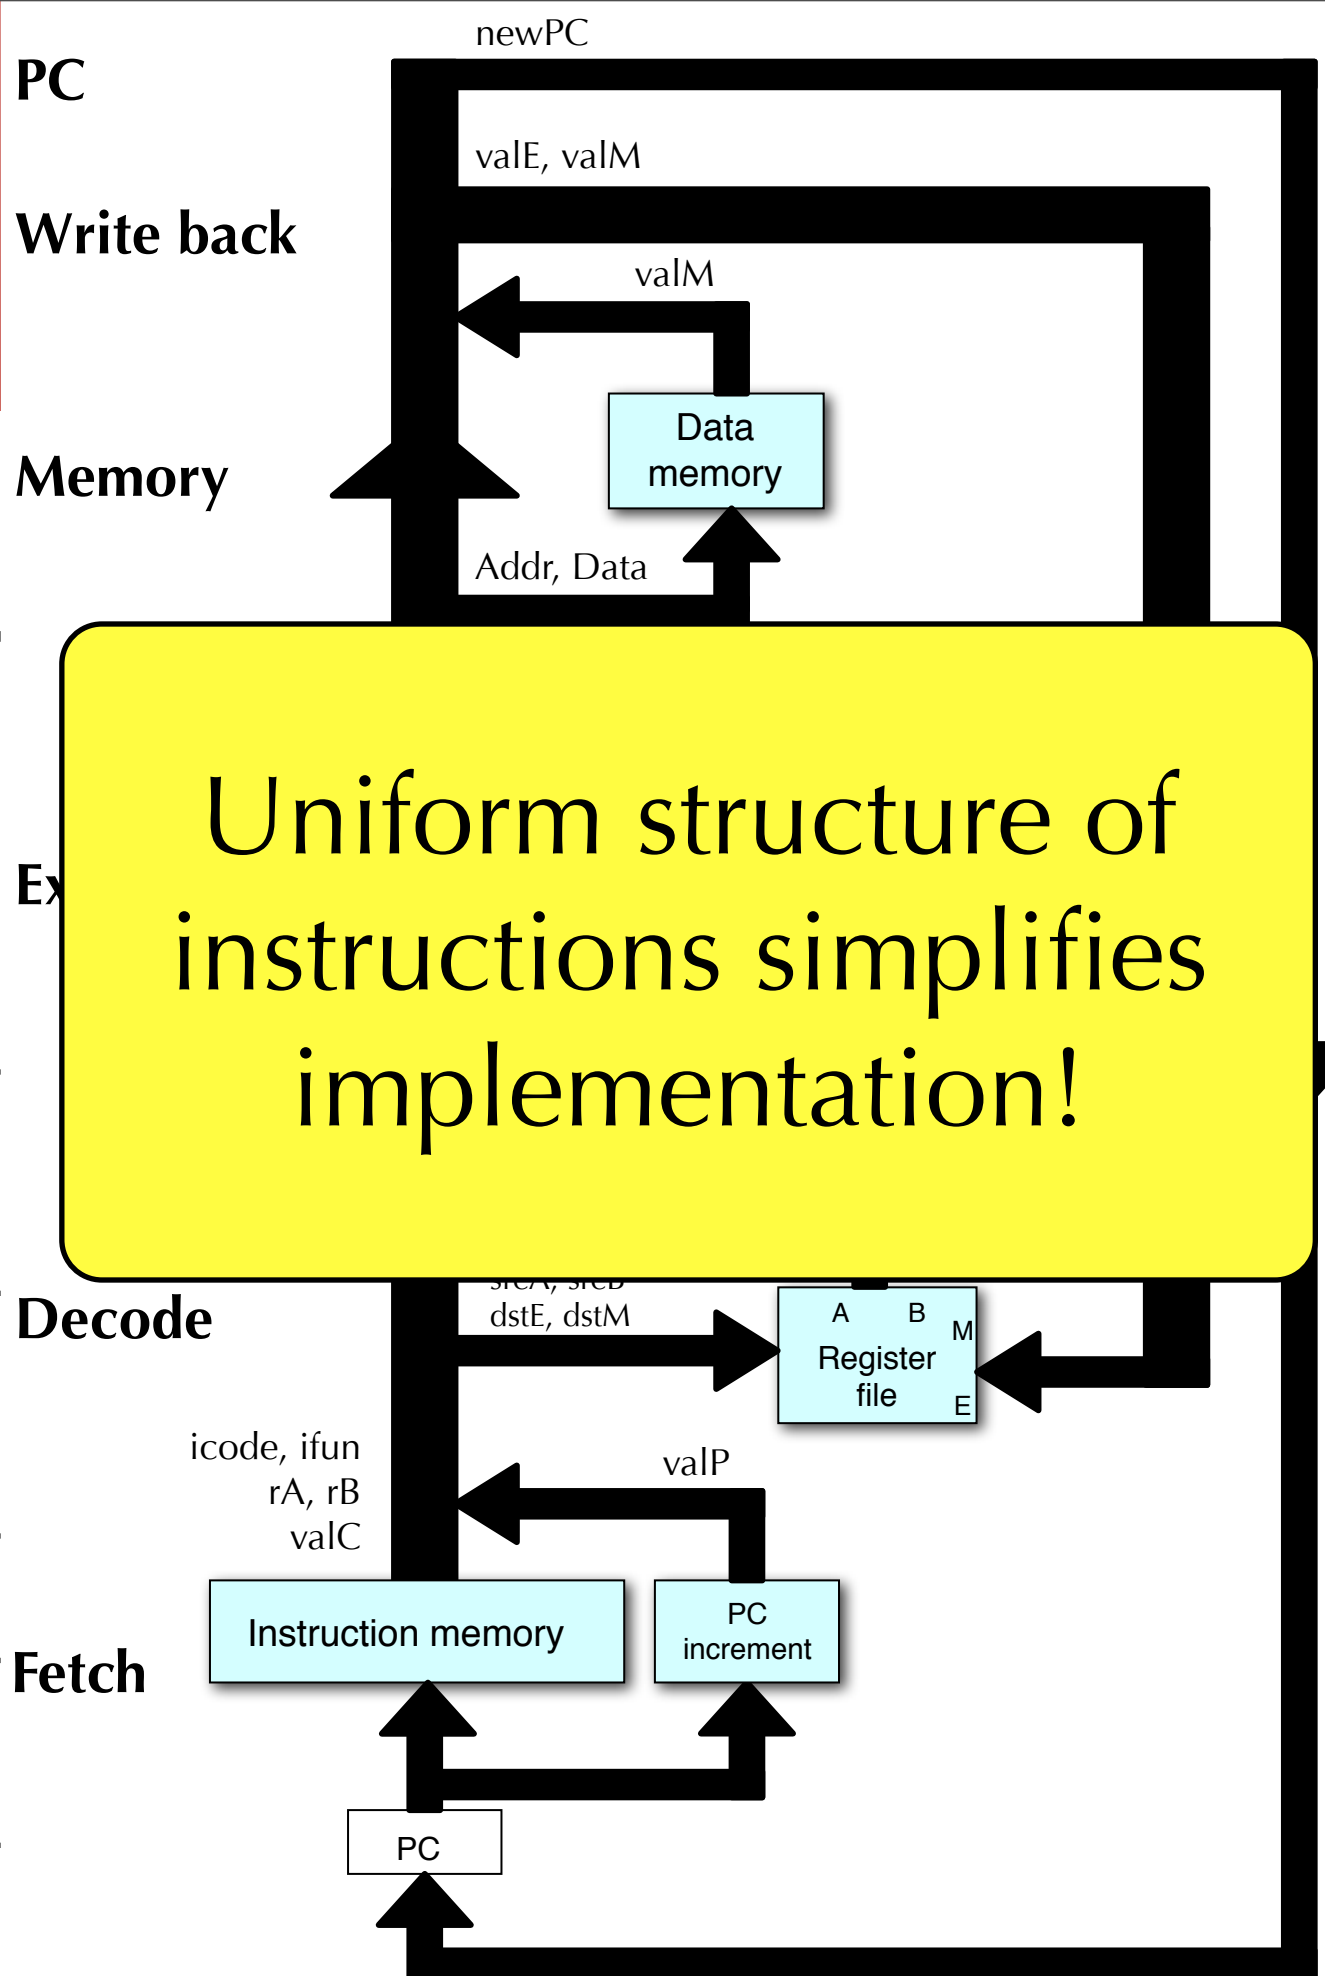
Execution stages

Stage	rrmovl rA, rB
Fetch	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$ $\text{rA:rB} \leftarrow M_1[\text{PC}+1]$ $\text{valP} \leftarrow \text{PC} + 2$
Decode	$\text{valA} \leftarrow R[\text{rA}]$
Execute	$\text{valE} \leftarrow \text{valA}$
Memory	–
Write back	$R[\text{rB}] \leftarrow \text{valE}$
PC update	$\text{PC} \leftarrow \text{valP}$



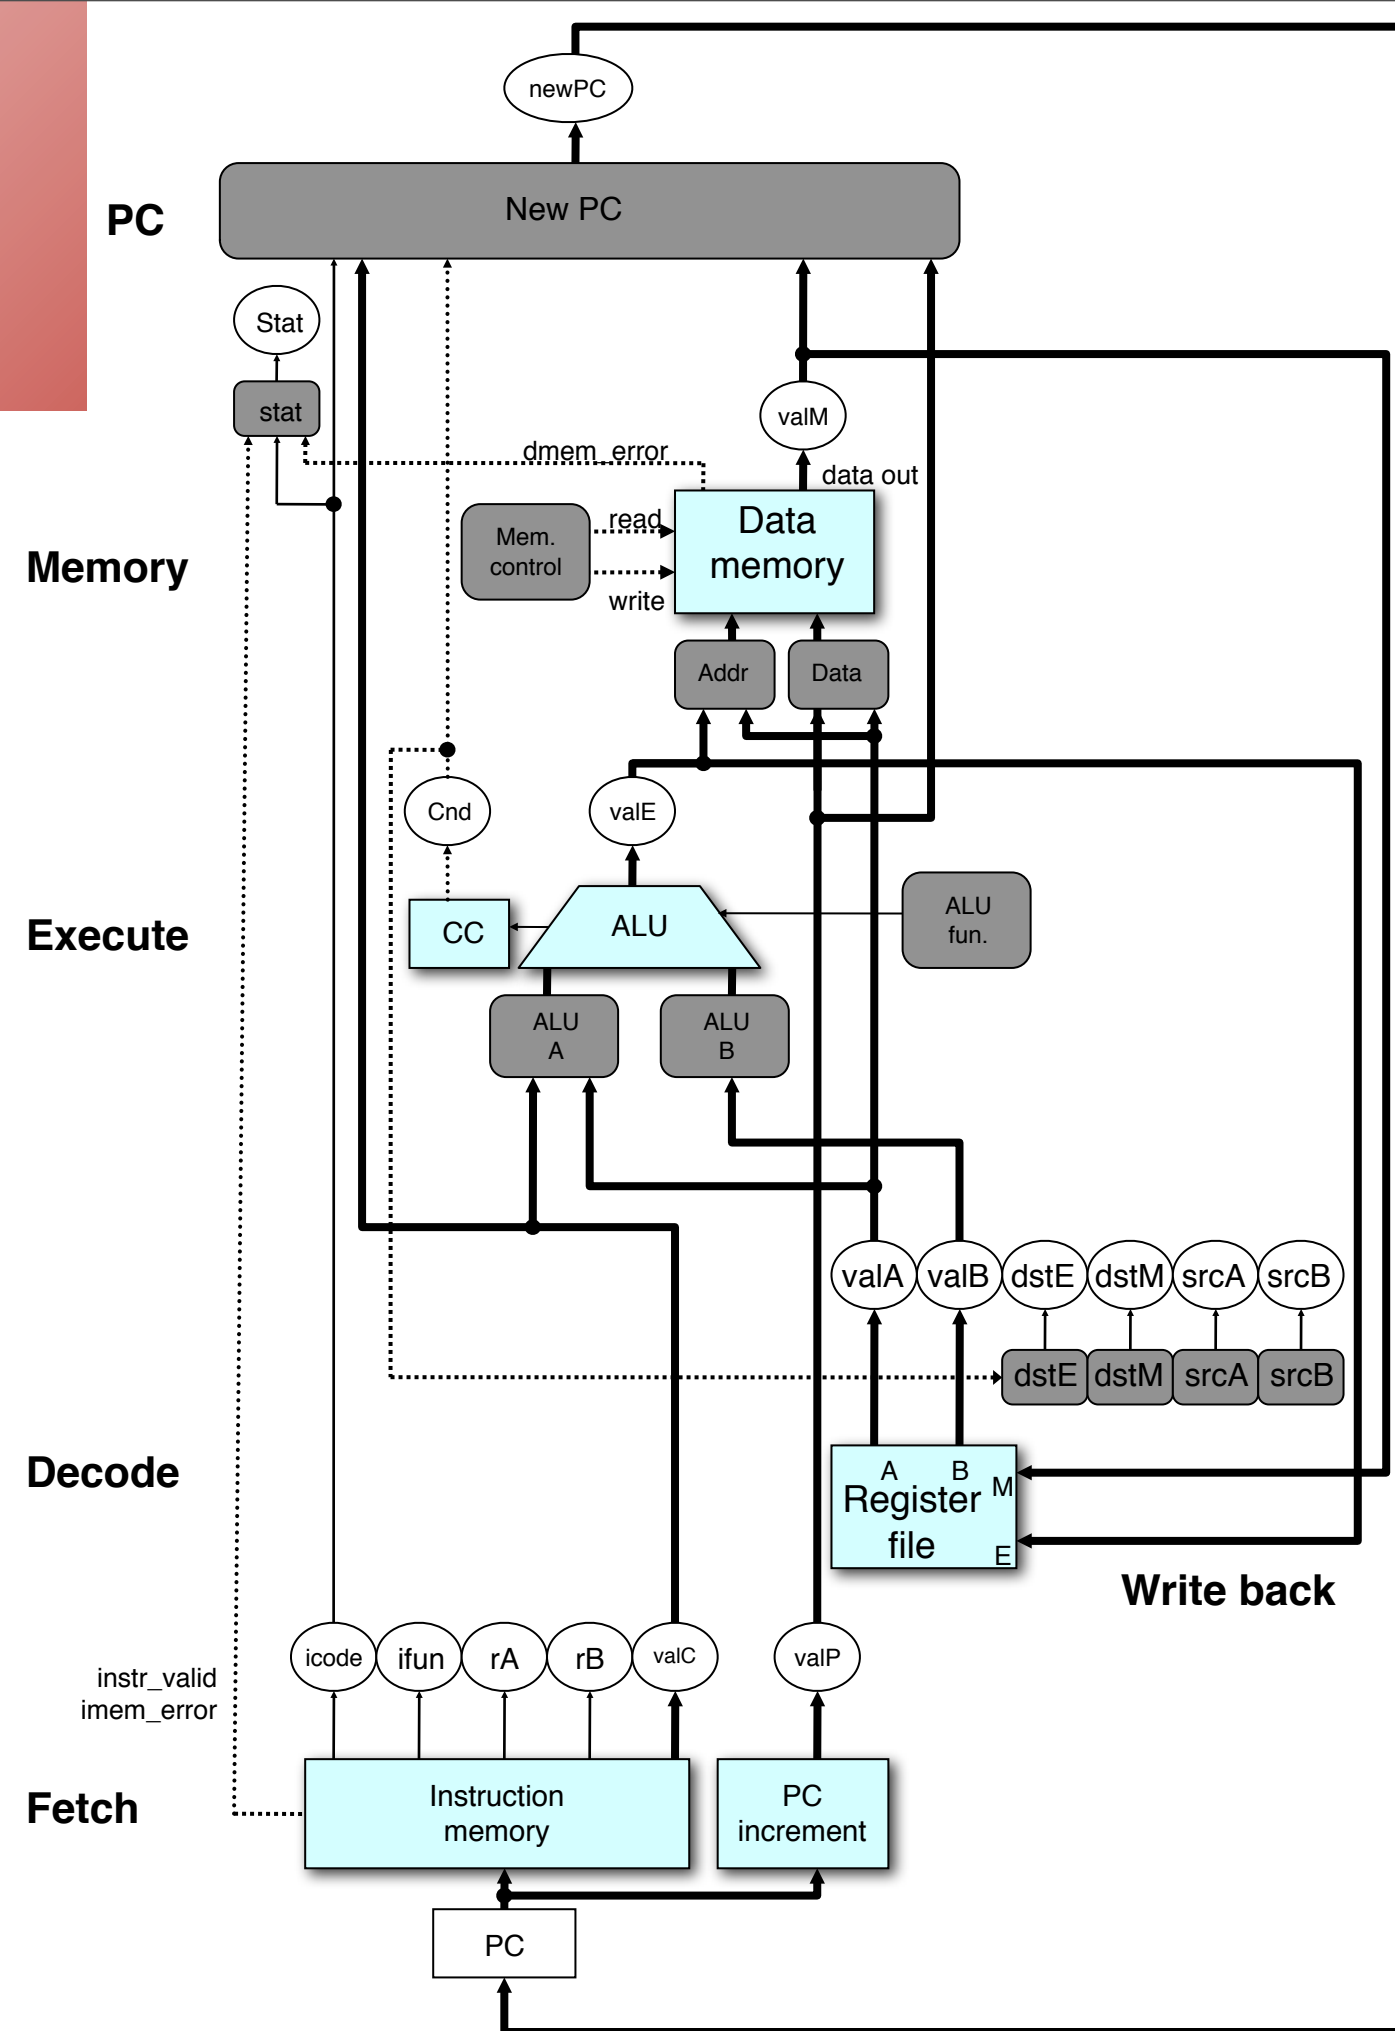
Execution stages

Stage	mrmovl C(rB), rA
Fetch	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$ $\text{rA:rB} \leftarrow M_1[\text{PC}+1]$ $\text{valC} \leftarrow M_4[\text{PC}+2]$ $\text{valP} \leftarrow \text{PC} + 6$
Decode	$\text{valB} \leftarrow R[\text{rB}]$
Execute	$\text{valE} \leftarrow \text{valB} + \text{valC}$
Memory	$\text{valM} \leftarrow M_4[\text{valE}]$
Write back	$R[\text{rA}] \leftarrow \text{valM}$
PC update	$\text{PC} \leftarrow \text{valP}$

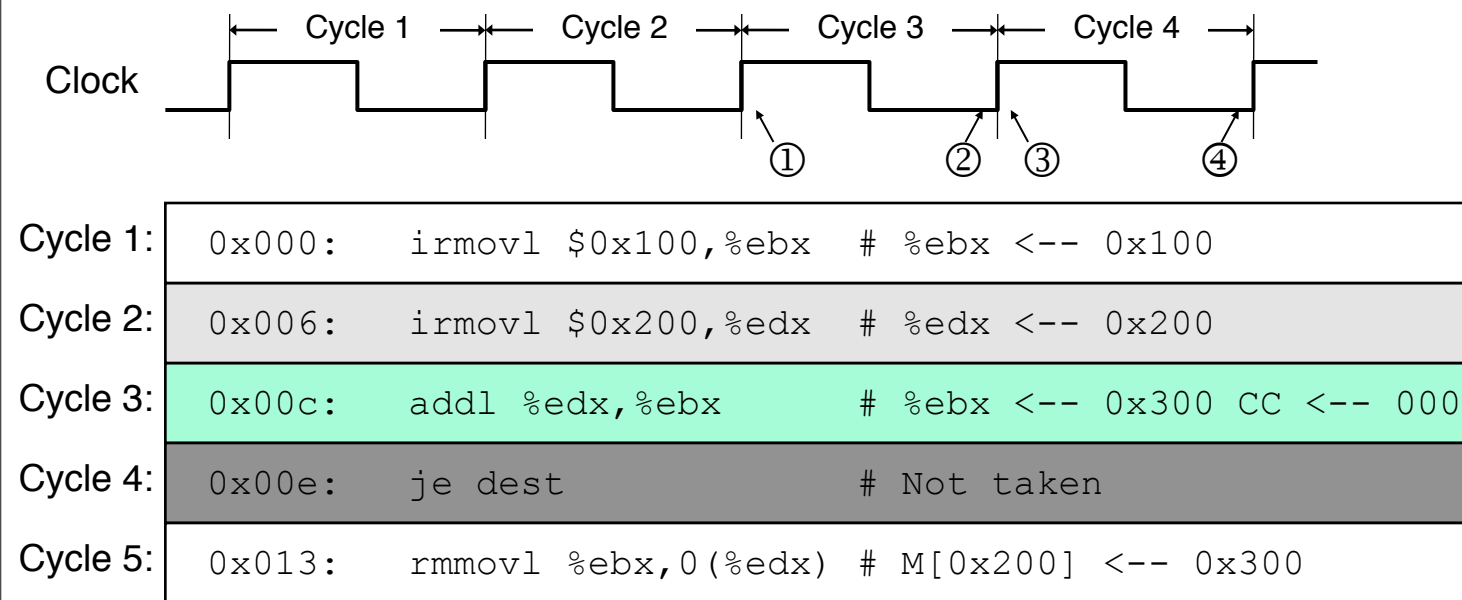


Hardware structure

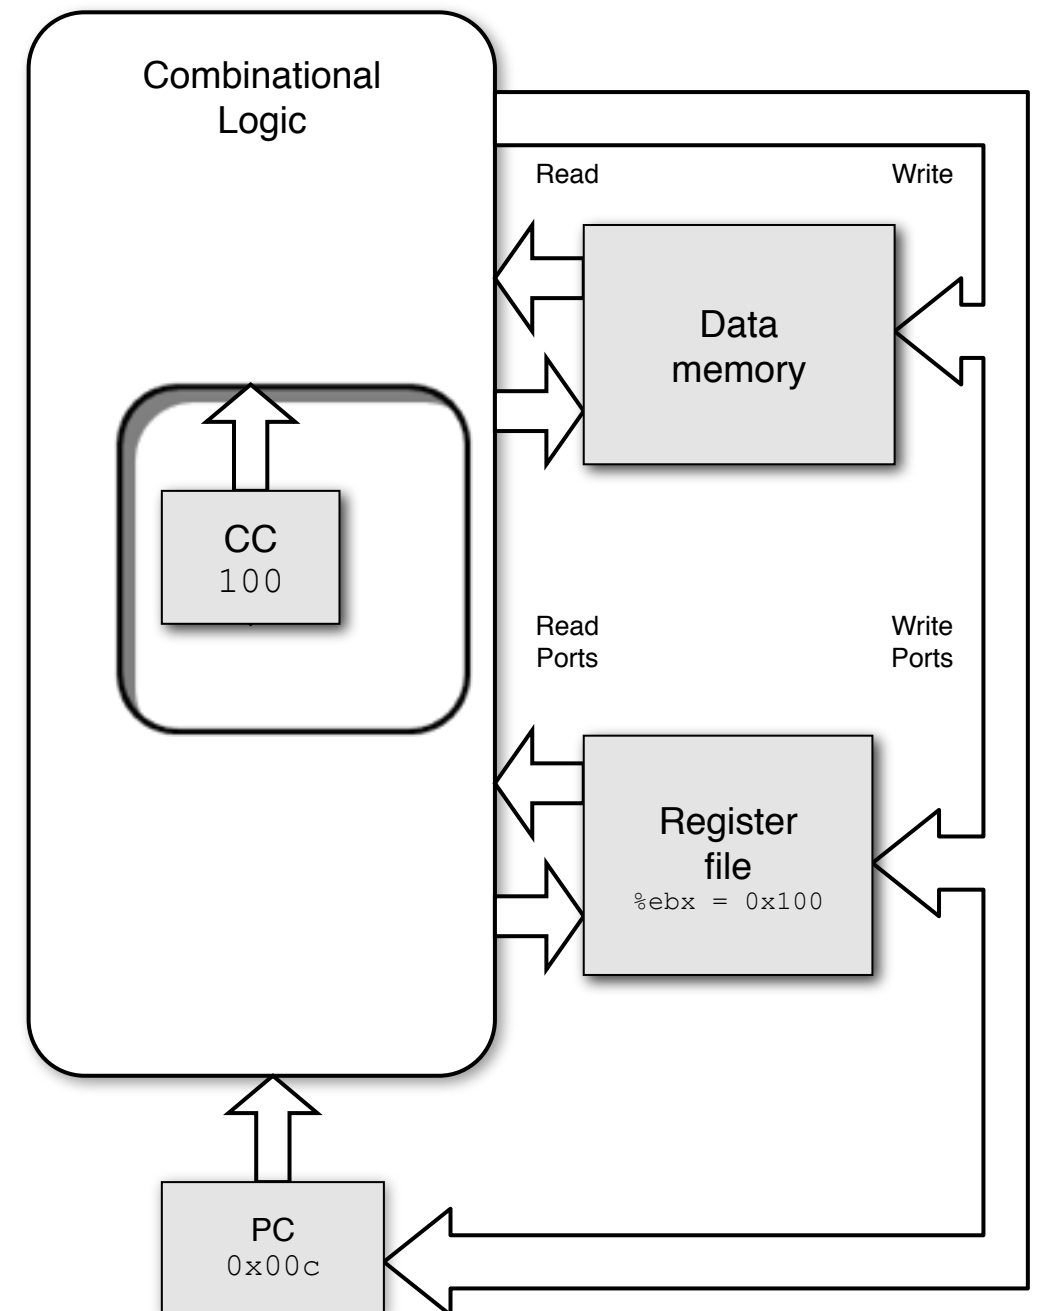
- Processing is performed by **hardware units** associated with the different stages
- Clock cycle needs to be long enough that signal can propagate throughout all units




Tracing cycles



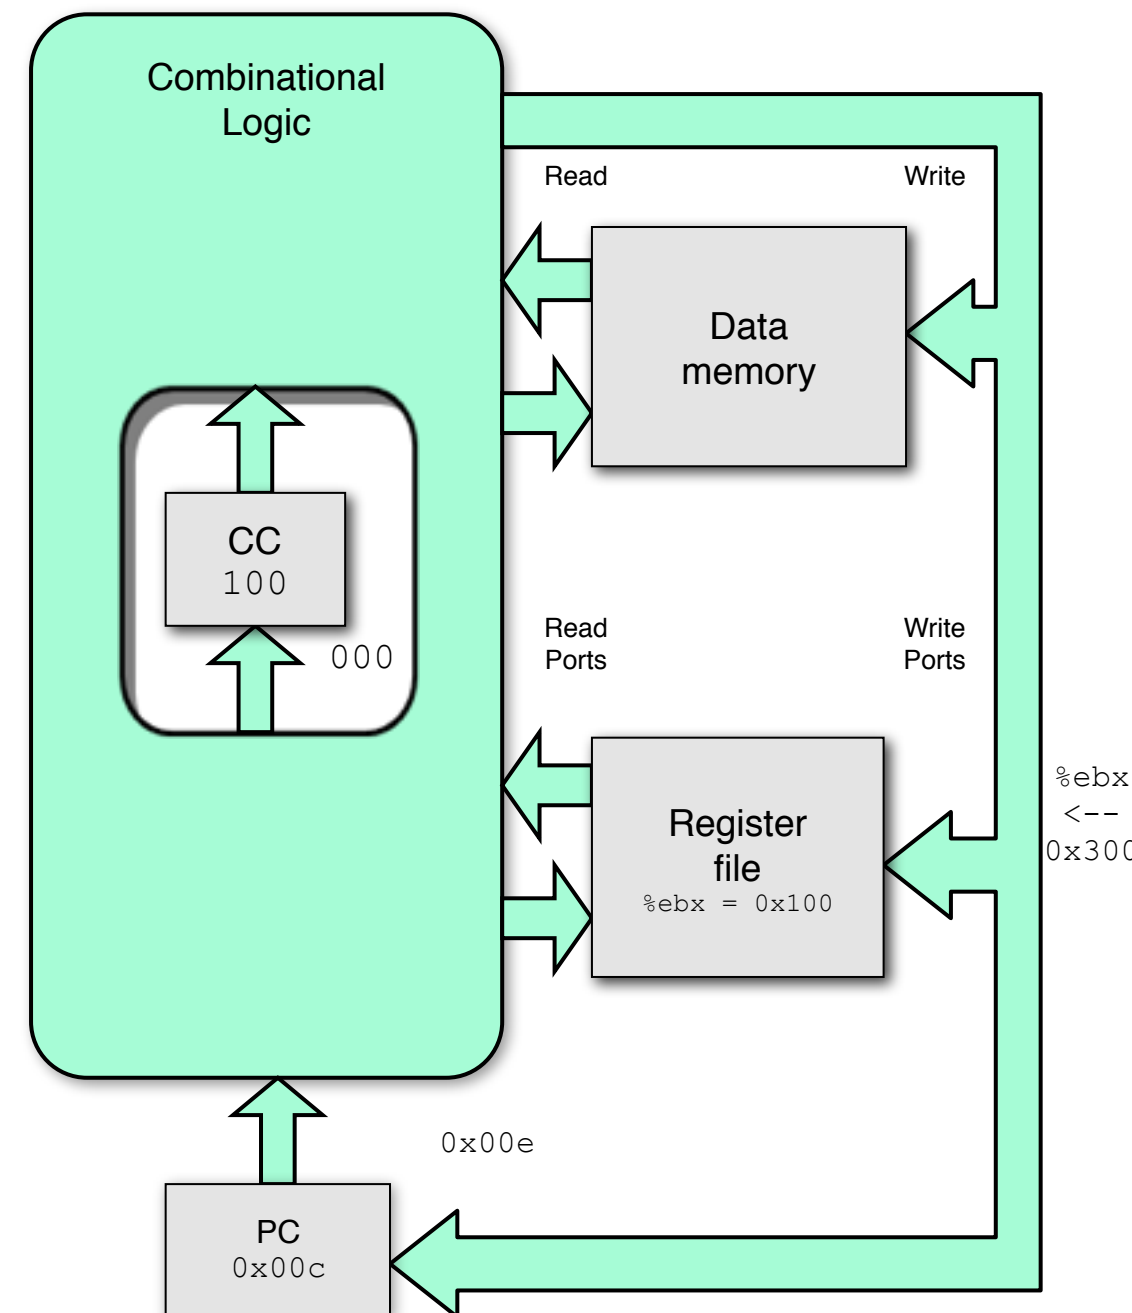
① Beginning of cycle 3




Tracing cycles



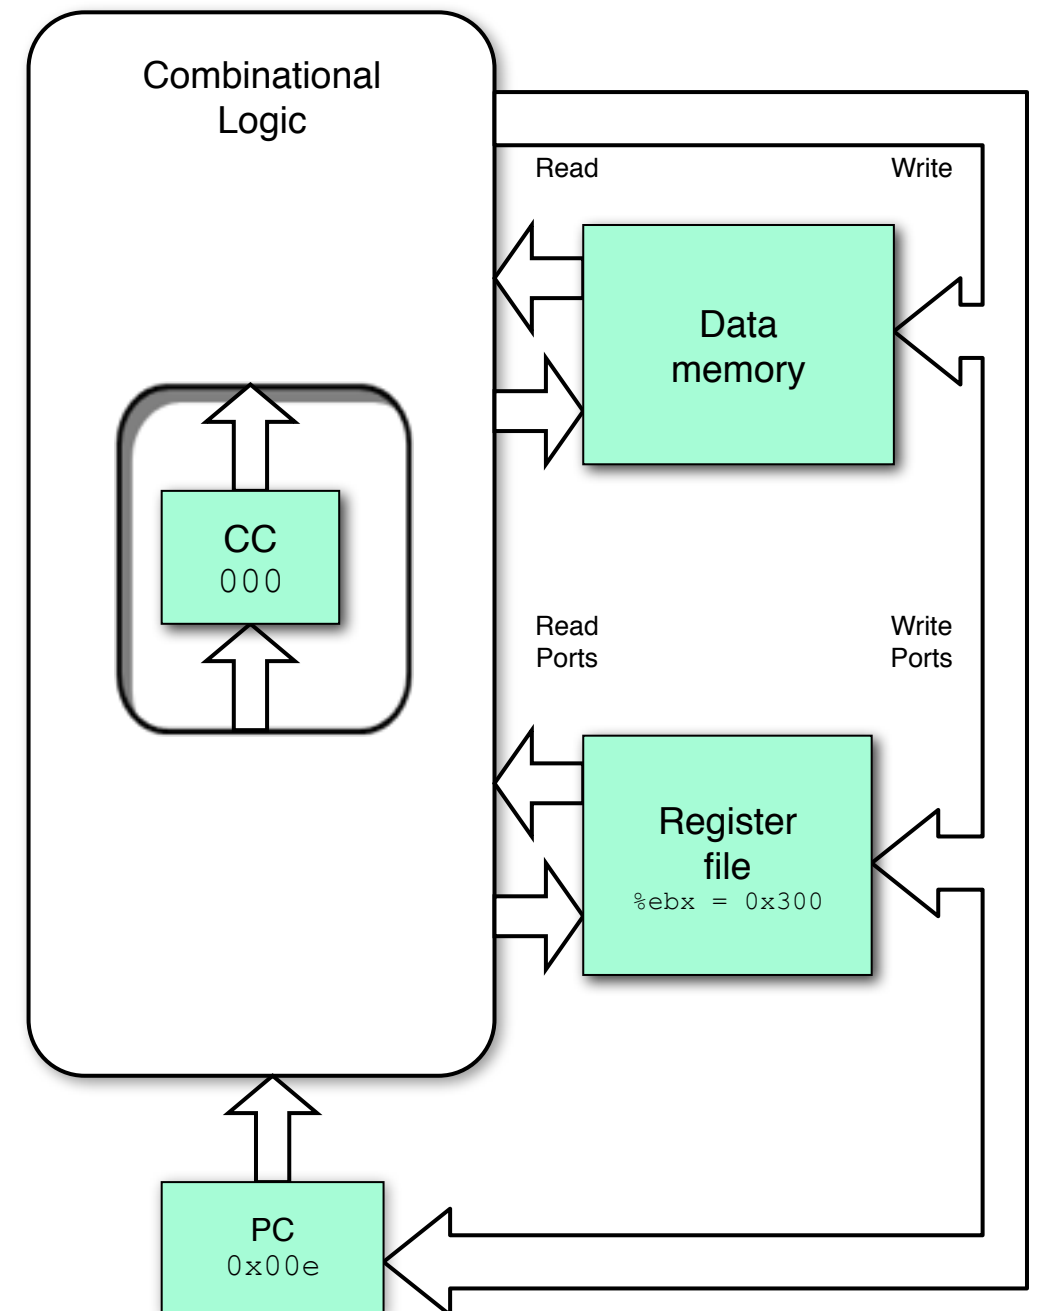
② End of cycle 3




Tracing cycles



③ Beginning of cycle 4

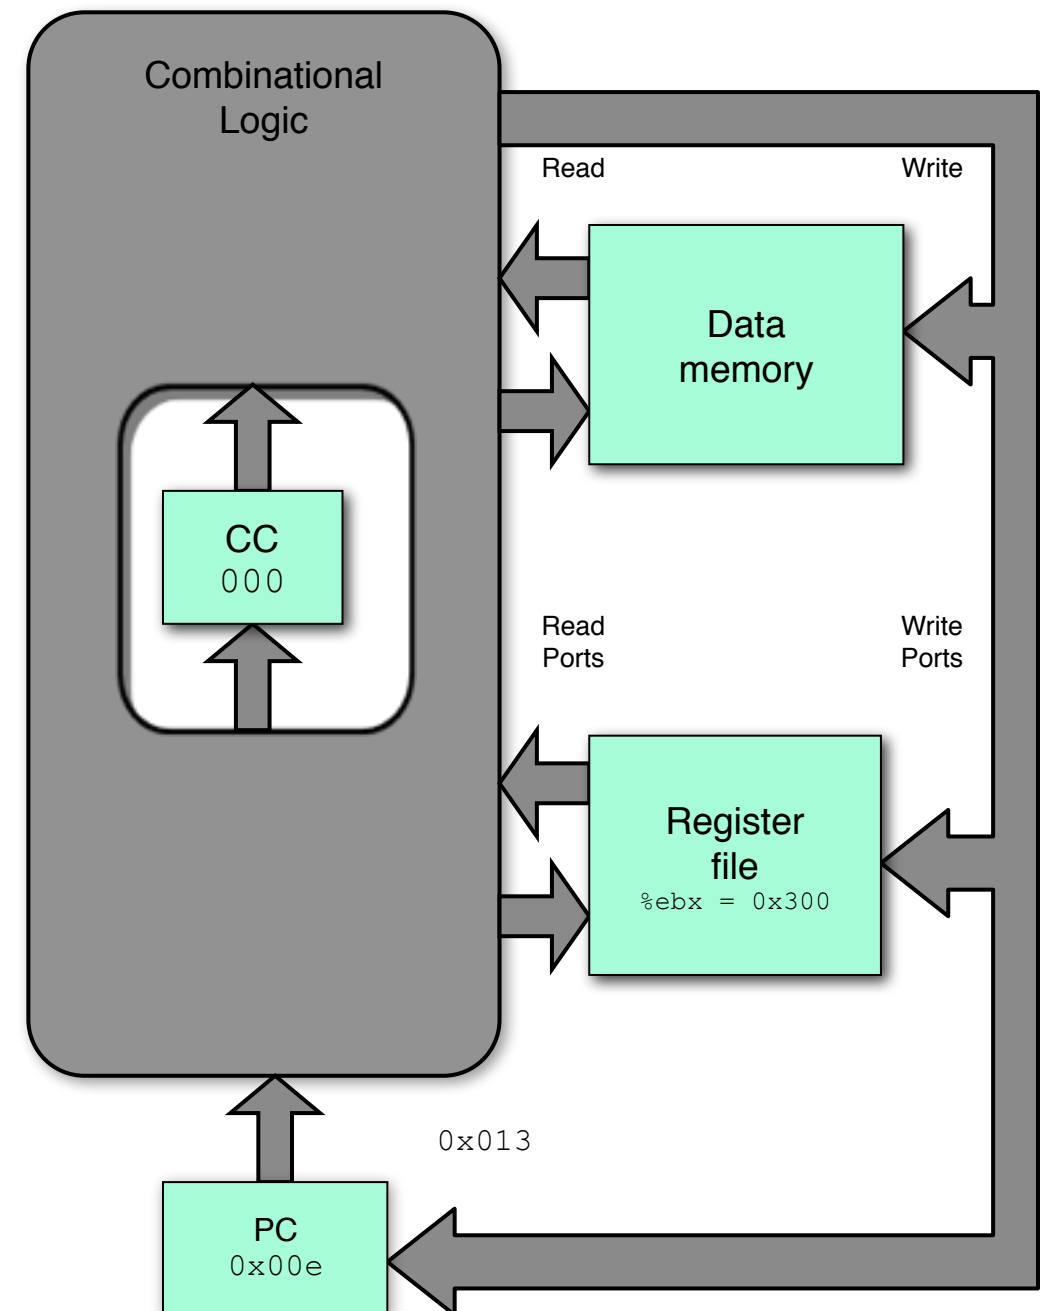


Tracing cycles




Cycle 1:	0x000:	irmovl \$0x100,%ebx	# %ebx <-- 0x100
Cycle 2:	0x006:	irmovl \$0x200,%edx	# %edx <-- 0x200
Cycle 3:	0x00c:	addl %edx,%ebx	# %ebx <-- 0x300 CC <-- 000
Cycle 4:	0x00e:	je dest	# Not taken
Cycle 5:	0x013:	rmmovl %ebx,0(%edx)	# M[0x200] <-- 0x300

④ End of cycle 4



Inefficient!

- Because cycle has to be long enough for signal to propagate, hardware units are often idle!
- Only one instruction is executed at a time



Time


addl \$4, %eax



addl \$4, %edx



addl \$4, %ecx




Pipelining

- But we can start on next instruction while previous instruction is still processing
- This is known as **pipelining**:
 - Task to be performed is divided into discrete stages
 - Multiple tasks can be worked on at same time, each in a different stage
 - E.g., Car wash
 - Suppose 3 stages: Suds, Rinse, and Polish
- Pipelining increases **throughput** of system
 - Number of tasks per unit time is higher
- Pipelining may increase **latency**
 - Any one task may take longer than if no pipeline used

Pipelining

Time



- Clock can be much faster
 - Needs to be long enough for signal to propagate through a single stage
- In order to achieve pipelining, need to add hardware registers between the hardware units
 - Overhead of these registers limits how deep a pipeline can be

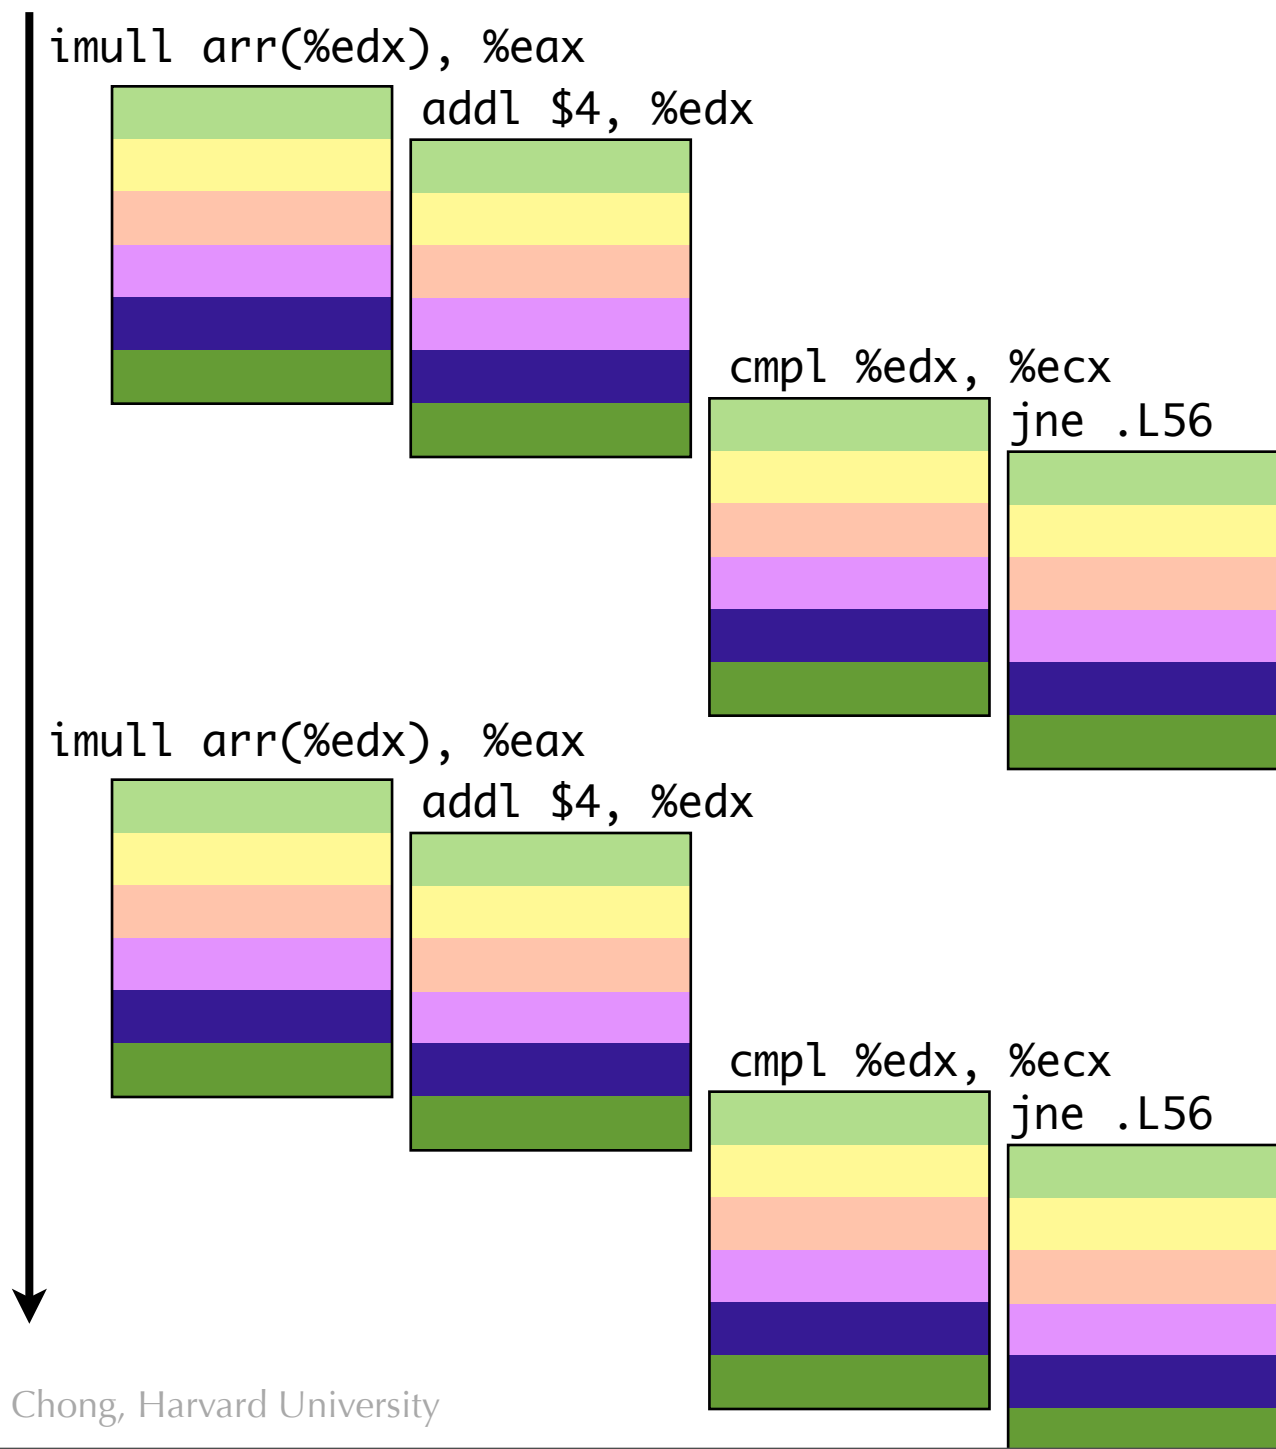
Example

```
int * arr;
int prod_array(int n) {
    int i, result=1;
    for (i = 0; i < n; i++) {
        result *= arr[i];
    }
    return result;
}
```

```
# Main body of loop
.L56:
    imull    arr(%edx), %eax
    addl     $4, %edx
    cmpl     %edx, %ecx
    jne      .L56
```

Example

Time



Main body of loop

.L56:

```
imull    arr(%edx), %eax
addl     $4, %edx
cmpl     %edx, %ecx
jne      .L56
```

- `cmpl` needs to read `%edx`, which isn't written until the write-back stage of `addl $4, %edx`!
 - One solution: Insert "bubble" in the pipeline, waiting until results are available
 - Also known as **stalling**
- What PC should be used for next iteration?
 - Result not available until Update PC stage of `jne`
- Other **hazards** can also complicate pipeline operation

More efficient pipelines

- Instead of inserting bubbles into pipeline, can we do better?
- **Forwarding results**
 - Instead of waiting for write-back stage of previous instruction, previous instruction passes result directly to next instruction
- **Branch prediction**
 - Instead of waiting to know next PC, predict what it will be and start executing
 - May be wrong: need to be able to throw away speculative execution
 - Modern branch prediction techniques have very good prediction rates: 90%+
- **Out-of-order execution**
 - Re-ordering instructions can reduce stalling
 - Processor may execute instructions in different order than specified
- ...

Today

- Processor architecture
 - Logic gates
 - Adders and multiplexors
 - Registers
 - Instruction set encoding
 - A sequential processor
 - Pipelining
 - CISC vs RISC

CISC vs RISC

- CISC (“sisk”): Complex Instruction Set Computer
- RISC (“risk”): Reduced Instruction Set Computer
- Different philosophies regarding the design (and implementation) of ISAs

CISC vs RISC

- CISC (“sisk”): Complex Instruction Set Computer
 - Historically first
 - Large instruction sets (x86 has several hundred)
 - Specialized instructions for high-level tasks
 - Instructions that are closer to what applications are wanting to do
 - Can provide hardware support for application-specific instructions
 - ▶ E.g., x86 contains instructions such as *LOOPZ label*, which decrements %cx (without modifying flags) and jumps to label if %cx is non-zero
 - Presents a clean interface to programmer
 - Hides implementation details such as pipelining

CISC vs RISC

- RISC (“risk”): Reduced Instruction Set Computer
 - Philosophy developed in early 1980s
 - Small, simple, instruction sets (typically <100)
 - E.g., may have only base+displacement memory addressing
 - E.g., memory access only via load and store; ALU operations need register operands
 - E.g., no condition codes, only explicit test instructions
 - Often fixed length encodings for instructions
 - Leads to simple, efficient implementation
 - Reveals implementation details to programmer (e.g., pipe-lining)
 - E.g., certain instruction sequences may be prohibited
 - E.g., jump instruction may not take effect until after following instruction
 - ▶ Compiler must be aware of these restrictions, and can use them to optimize performance
 - ARM (originally “Acorn RISC Machine”) widely used in embedded devices

Modern computers

- In most settings, neither CISC nor RISC clearly better
- RISC machines
 - Exposing implementation details made it difficult to use them, and difficult to evolve the ISA
 - Added more instructions
- CISC machines
 - Take advantage of RISC-like pipelines
 - Essentially translate CISC instructions into simpler RISC-like instructions
 - E.g., `addl %eax, 8(%esp)` broken up into a load from memory, followed by an addition, followed by a store to memory

Next Lecture

- Program optimizations
 - How to greatly improve the performance of your C programs with some simple changes
- Code motion
- Strength reduction
- Common subexpression elimination
- Loop unrolling
- Tail recursion