



HARVARD

School of Engineering
and Applied Sciences

Dynamic Memory Allocation

CS61, Lecture 10

Prof. Stephen Chong

October 4, 2011

Announcements 1/2

- Assignment 4: Malloc
 - Will be released today
 - May work in groups of one or two
 - Please go to website and enter your group by **Sunday 11:59PM**
 - ▶ Every group must do this (i.e., even if you are working individually)
 - Two deadlines:
 - Design checkpoint: Thursday October 13, 10pm
 - Final submission: Thursday October 20, 11:59pm
 - Encourage you to use version control
- Assignment 3 (Buffer bomb) due Thursday

Announcements 2/2

- Give us feedback on when you'd like office hours
 - <http://tinyurl.com/689seas>
 - Link posted on website

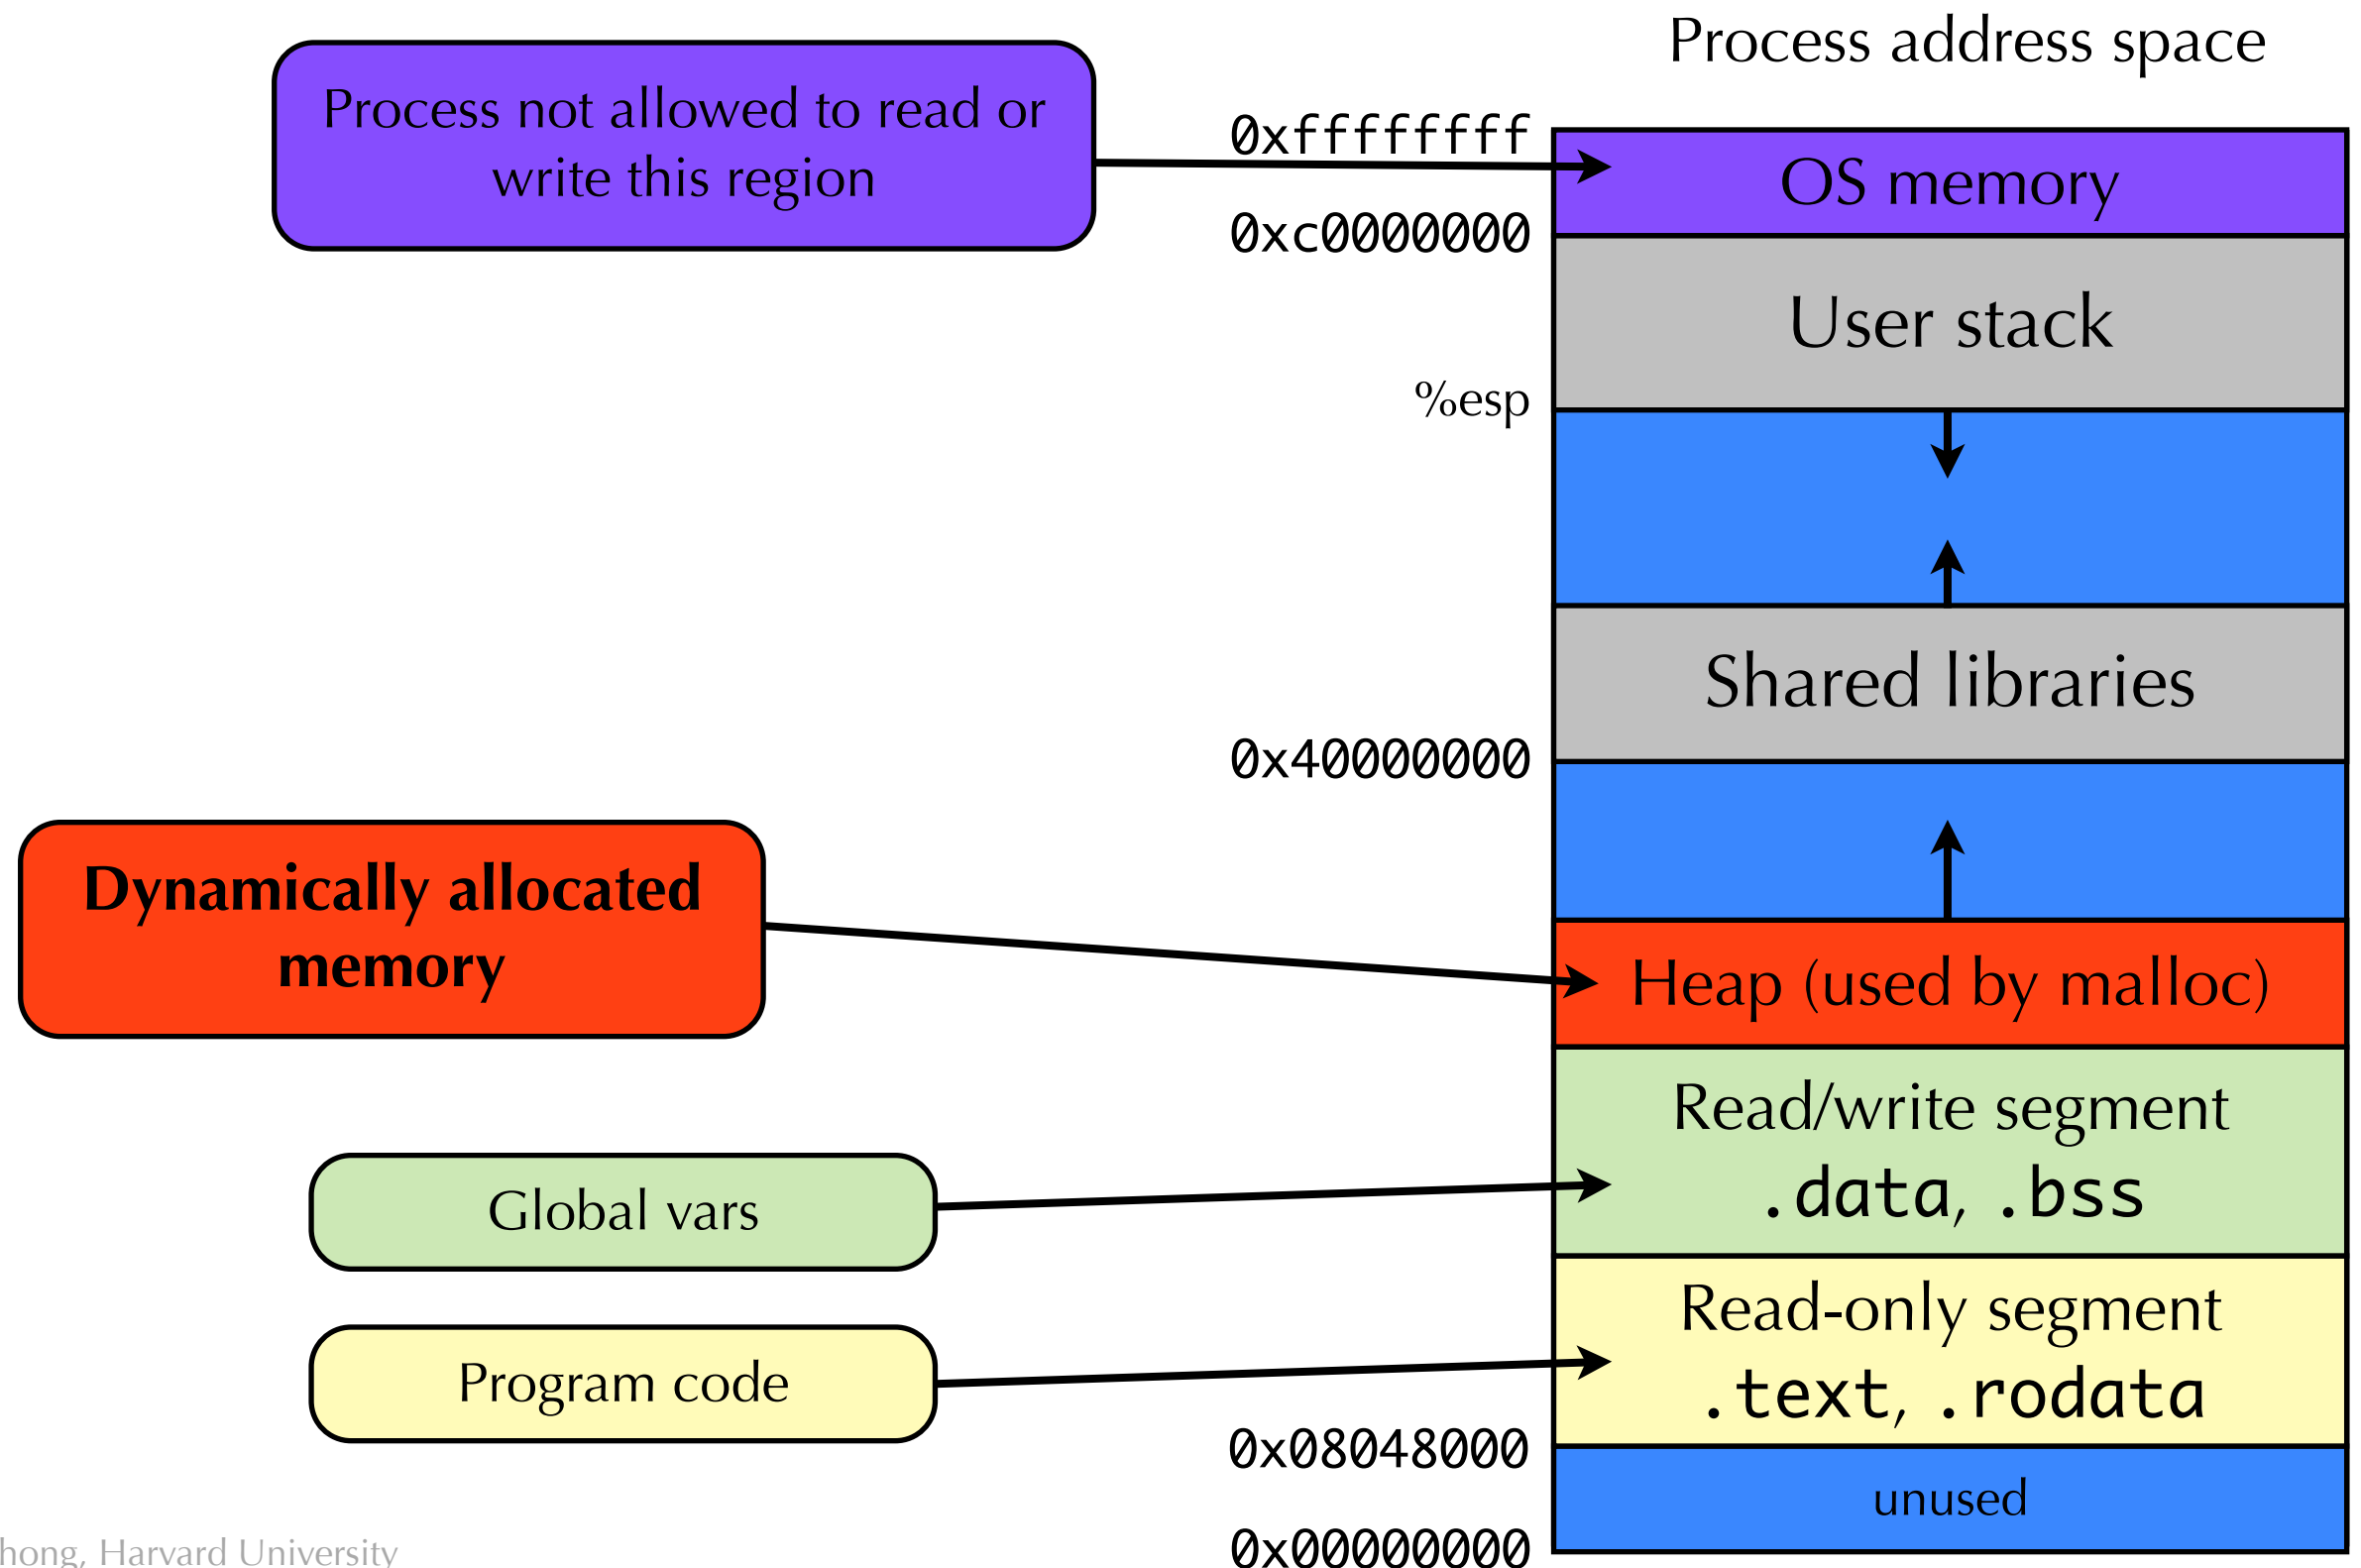
Topics for today

- Dynamic memory allocation
 - Implicit vs. explicit memory management
- Performance goals
- Fragmentation
- Free block list management
 - Implicit free list
 - Explicit free list
 - Segregated lists
 - Tradeoffs

Harsh Reality: Memory Matters!


- Memory is not unlimited!
 - It must be allocated and managed
 - Many applications are memory dominated
 - Especially those based on complex, graph algorithms
- Memory referencing bugs especially pernicious
 - Effects are distant in both time and space
- Memory performance is not uniform
 - Cache and virtual memory effects can greatly affect program performance
 - Adapting program to characteristics of memory system can lead to major speed improvements

A process's view of memory



The heap

- The **heap** is the region of a program's memory used for dynamic allocation.
- Program can allocate and free blocks of memory within the heap.



Free blocks, and allocated blocks

- Heap starts out as a single big “free block” of some fixed size (say a few MB)



- Program may request memory, which splits up the the free space.



- Program may free up some memory, which increases the free space



- Over time the heap will contain a mixture of free and allocated blocks.



- Heap may need to grow in size (but typically never shrinks)

- Program can grow the heap if it is too small to handle an allocation request.
- On UNIX, the `sbrk()` system call is used to expand the size of the heap.




Dynamic Memory Management

- How do we decide when to expand the heap (using `sbrk()`)?
- How do we manage allocating and freeing bytes on the heap?
- There are two broad classes of memory management schemes:
- Explicit memory management
 - Application code responsible for both explicitly **allocating** and **freeing** memory.
 - Example: `malloc()` and `free()`
- Implicit memory management
 - Application code can allocate memory, but **does not explicitly free memory**
 - Rather, rely on **garbage collection** to “clean up” memory objects no longer in use
 - Used in languages like Java, Python, OCaml

Malloc Package

- `#include <stdlib.h>`
- `void *malloc(size_t size)`
 - If successful:
 - Returns a pointer to a memory block of at least `size` bytes
 - If `size == 0`, returns `NULL`
 - If unsuccessful: returns `NULL`.
- `void free(void *p)`
 - Returns the block pointed at by `p` to pool of available memory
 - `p` must come from a previous call to `malloc` or `realloc`.
- `void *realloc(void *p, size_t size)`
 - Changes size of block `p` and returns pointer to new block.
 - Contents of new block unchanged up to min of old and new size.

Allocation Examples



Constraints

- Application code must obey following constraints:
 - Allowed to issue arbitrary sequence of allocation and free requests
 - Free requests must correspond to an allocated block
- Memory management code must obey following constraints:
 - Can't control number or size of requested blocks
 - Must respond immediately to all allocation requests
 - i.e., can't reorder or buffer requests
 - Must allocate blocks from free memory
 - i.e., can only place allocated blocks in free memory
 - Must align blocks so they satisfy all alignment requirements
 - Can't mess around with allocated memory
 - Can only manipulate and modify free memory
 - Can't move the allocated blocks once they are allocated (i.e., compaction is not allowed)

Topics for today

- Dynamic memory allocation
 - Implicit vs. explicit memory management
- Performance goals
- Fragmentation
- Free block list management
 - Implicit free list
 - Explicit free list
 - Segregated lists
 - Tradeoffs

What do we want?

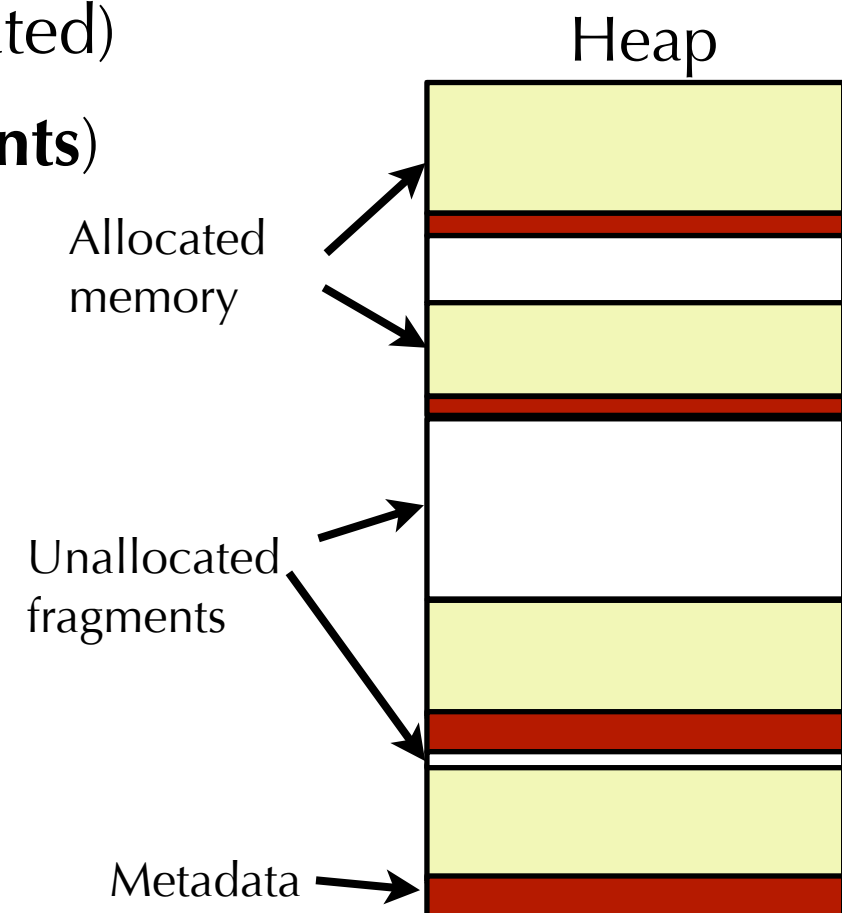
- Want our memory management to be:
 - Fast
 - Minimize overhead of allocation and deallocation operations.
 - Efficient
 - Don't waste memory space

Performance Goals: Allocation overhead

- Want our memory allocator to be fast!
 - Minimize the overhead of both allocation and deallocation operations.
- One useful metric is **throughput**:
 - Given a series of allocate or free requests
 - Maximize the number of completed requests per unit time
- Example:
 - 5,000 malloc calls and 5,000 free calls in 10 seconds
 - Throughput is 1,000 operations/second.

Performance Goals: Memory Utilization

- Allocators rarely do a perfect job of managing memory.
 - Usually there is some “waste” involved in the process.
- Examples of waste...
 - Extra metadata or internal structures used by the allocator itself (example: Keeping track of where free memory is located)
 - Chunks of heap memory that are unallocated (**fragments**)
- We define **memory utilization** as...
 - The **total amount of memory allocated to the application** divided by the **total heap size**
- Ideally, we'd like utilization to be to 100%
 - In practice this is not possible, but want to be as close as possible



Conflicting performance goals


- High throughput and good utilization are difficult to achieve simultaneously.
- A fast allocator may not be efficient in terms of memory utilization.
 - Faster allocators tend to be “sloppier” with their memory usage.
- Likewise, a space-efficient allocator may not be very fast
 - To keep track of memory waste (i.e., tracking fragments), the allocation operations generally take longer to run.
- Trick is to balance these two conflicting goals.

Topics for today

- Dynamic memory allocation
 - Implicit vs. explicit memory management
- Performance goals
- Fragmentation
- Free block list management
 - Implicit free list
 - Explicit free list
 - Segregated lists
 - Tradeoffs
- Alignment

Internal Fragmentation

- Poor memory utilization caused by **fragmentation**.
 - Comes in two forms: **internal** and **external** fragmentation
- **Internal fragmentation**
 - Internal fragmentation is the difference between block size and payload size.



- Caused by overhead of maintaining heap data structures, padding for alignment purposes, or the policy used by the memory allocator
- Example: Say the allocator always “rounds up” to next highest power of 2 when allocating blocks.
 - So `malloc(1025)` will actually allocate 2048 bytes of heap space!

External Fragmentation

- Occurs when there is enough aggregate heap memory, but no single free block is large enough to satisfy a given request.

`p1 = malloc(4)`



`p2 = malloc(5)`



`p3 = malloc(6)`



`free(p2)`



`p4 = malloc(6)`

No free block big enough!

Topics for today

- Dynamic memory allocation
 - Implicit vs. explicit memory management
- Performance goals
- Fragmentation
- Free block list management
 - Implicit free list
 - Explicit free list
 - Segregated lists
 - Tradeoffs

Free block list management


- How do we know how much memory to free just given a pointer?
- How do we keep track of the free blocks?
- What do we do with the extra space when allocating a memory block that is smaller than the free block it is placed in?
- How do we pick which free block to use for allocation?

Knowing how much to free

- Standard method

- Keep the length of a block in a header preceding the block.
- Requires an extra word for every allocated block

`p = malloc(4)`



`free(p)`




Keeping Track of Free Blocks

- One of the biggest jobs of an allocator is knowing where the free memory is.
 - Affects throughput and utilization
- Many approaches to free block management.
 - Today, we will talk about three techniques
 - Implicit free lists
 - Explicit free lists
 - Segregated free lists
 - There are other approaches

Implicit free list

- Idea: Each block contains a **header** with some extra information.
 - **Allocated bit** indicates whether block is allocated or free.
 - **Size field** indicates entire size of block (including the header)
 - Trick: Allocation bit is just the high-order bit of the size word
- For this lecture, assume header size is 1 byte.
 - Makes later pictures easier to understand.
 - This means the block size is only 7 bits, so max. block size is 127 bytes (2^7-1).
 - Clearly a real implementation would want to use a larger header (e.g., 4 bytes).

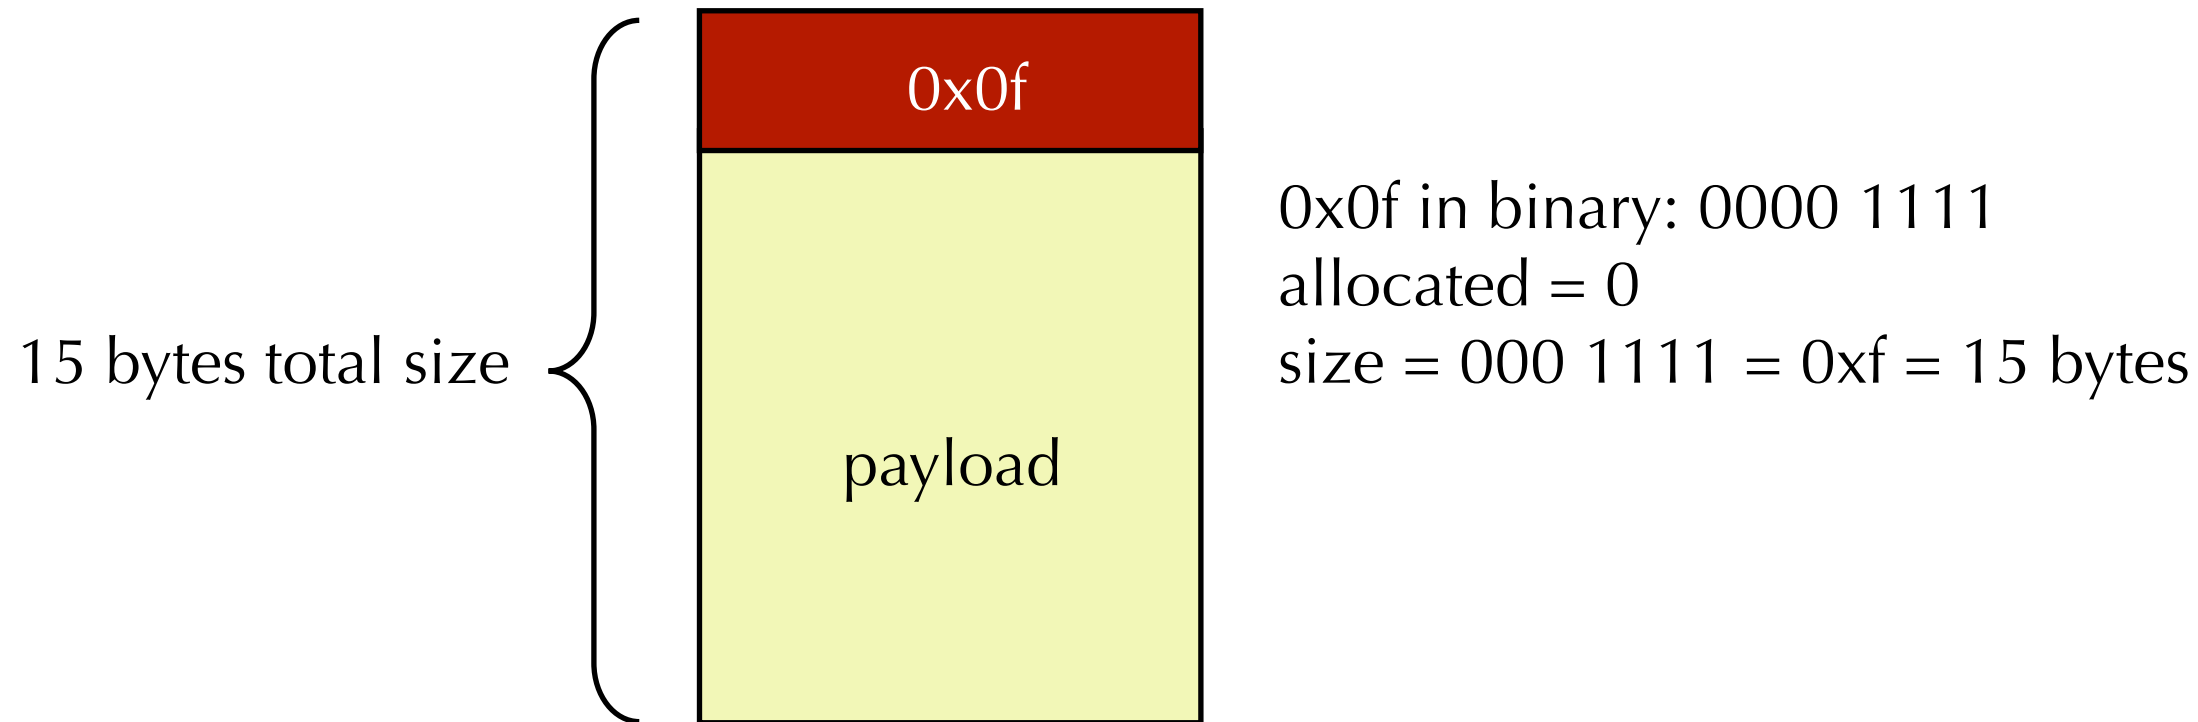
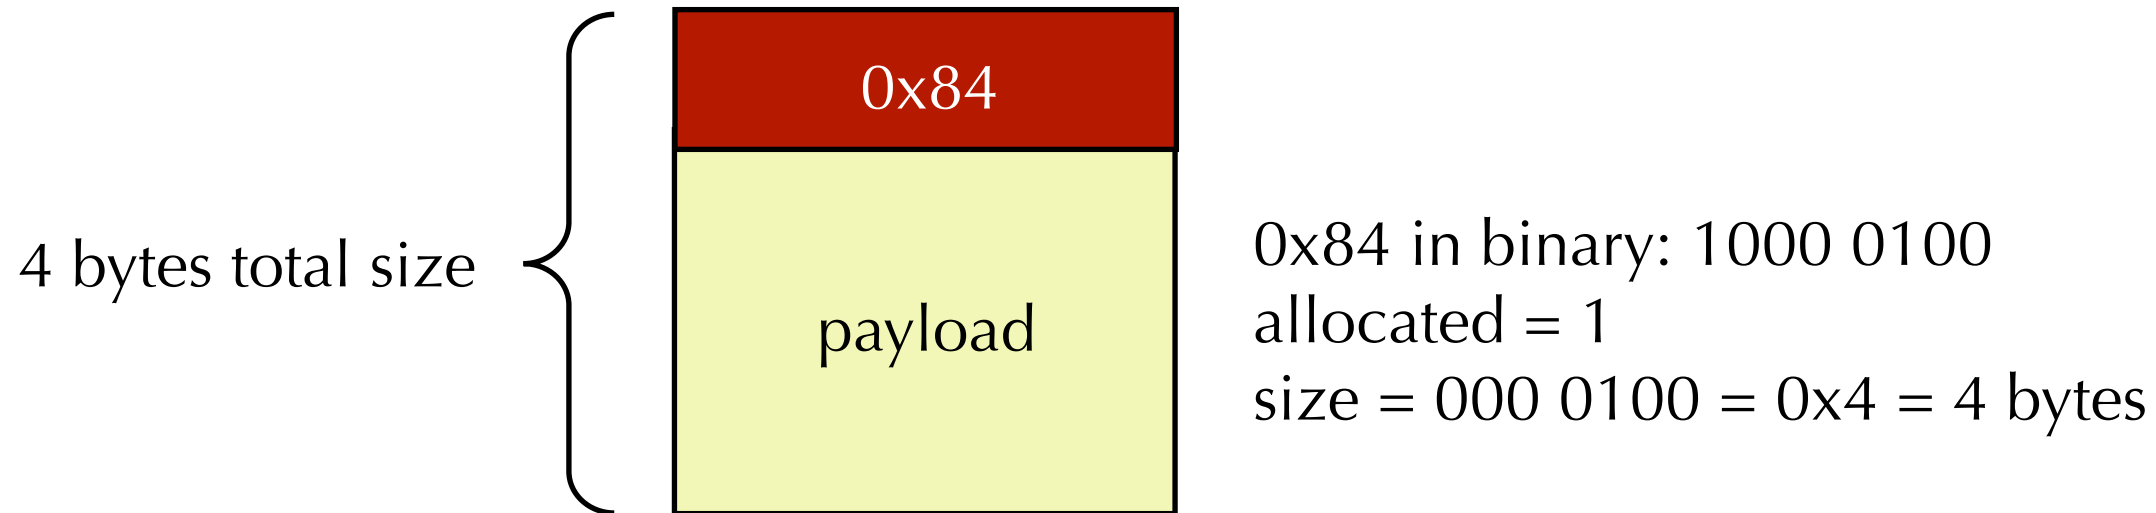


a = 1: block is allocated
a = 0: block is free


size: block size

payload: application data

Examples



Implicit free list




- No **explicit** structure tracking location of free/allocated blocks.
 - Rather, the size word (and allocated bit) in each block form an **implicit** “block list”
- How do we find a free block in the heap?
 - Start scanning from the beginning of the heap.
 - Traverse each block until (a) we find a free block and (b) the block is large enough to handle the request.
 - This is called the **first fit** strategy.

Implicit List: Finding a Free Block

- **First fit** strategy:
 - Search list from beginning, choose first free block that fits
 - Can take linear time in total number of blocks (allocated and free)
 - In practice it can cause “splinters” at beginning of list
- **Next fit** strategy:
 - Like first-fit, but search list from location of end of previous search
 - Research suggests that fragmentation is worse than first-fit
- **Best fit** strategy:
 - Search the list, choose the free block with the closest size that fits
 - Keeps fragments small — usually helps fragmentation
 - Runs slower than first- or next-fit, since the entire list must be searched each time


Implicit List: Allocating in Free Block

- Splitting free blocks
 - Since allocated space might be smaller than free space, we may need to **split** the free block that we're allocating within
 - E.g., `malloc(3)`



```
void addblock(ptr_t p, int len) {  
    int oldsize = *p;           // Get old size of free block  
    *p = len | 0x80;           // Set new size + alloc bit  
    if (len < oldsize)  
        *(p+len) = oldsize - len; // Set size of remaining  
    }                          // part of free block
```


`addblock(p, 4)`




Implicit List: Freeing a Block

- Simplest implementation:
 - Simply clear the allocated bit in the header

```
/* Here, p points to the block header. */  
/* This sets the high-order bit to 0. */  
void free_block(ptr_t p) { *p = *p & 0x7f; }
```




free(p+1)




Implicit List: Freeing a Block

- Simplest implementation:
 - Simply clear the allocated bit in the header
 - /* Here, p points to the block header. */
 - /* This sets the high-order bit to 0. */
 - void free_block(ptr_t p) { *p = *p & ~0x80; }
 - But can lead to “false fragmentation”



free(p+1)




malloc(5)

- There is enough free space, but the allocator won't be able to find it!

Implicit List: Coalescing

- **Coalesce** with adjacent block(s) if they are free

```
void free_block(ptr_t p) {  
    *p = *p & 0x7f;    // Clear allocated bit  
    next = p + *p;      // Find next block  
    if ((*next & 0x80) == 0)  
        *p = *p + *next; // Add to this block if  
                          // next is also free  
}
```




- This is coalescing with the next free block.
- How would we coalesce with the previous free block?

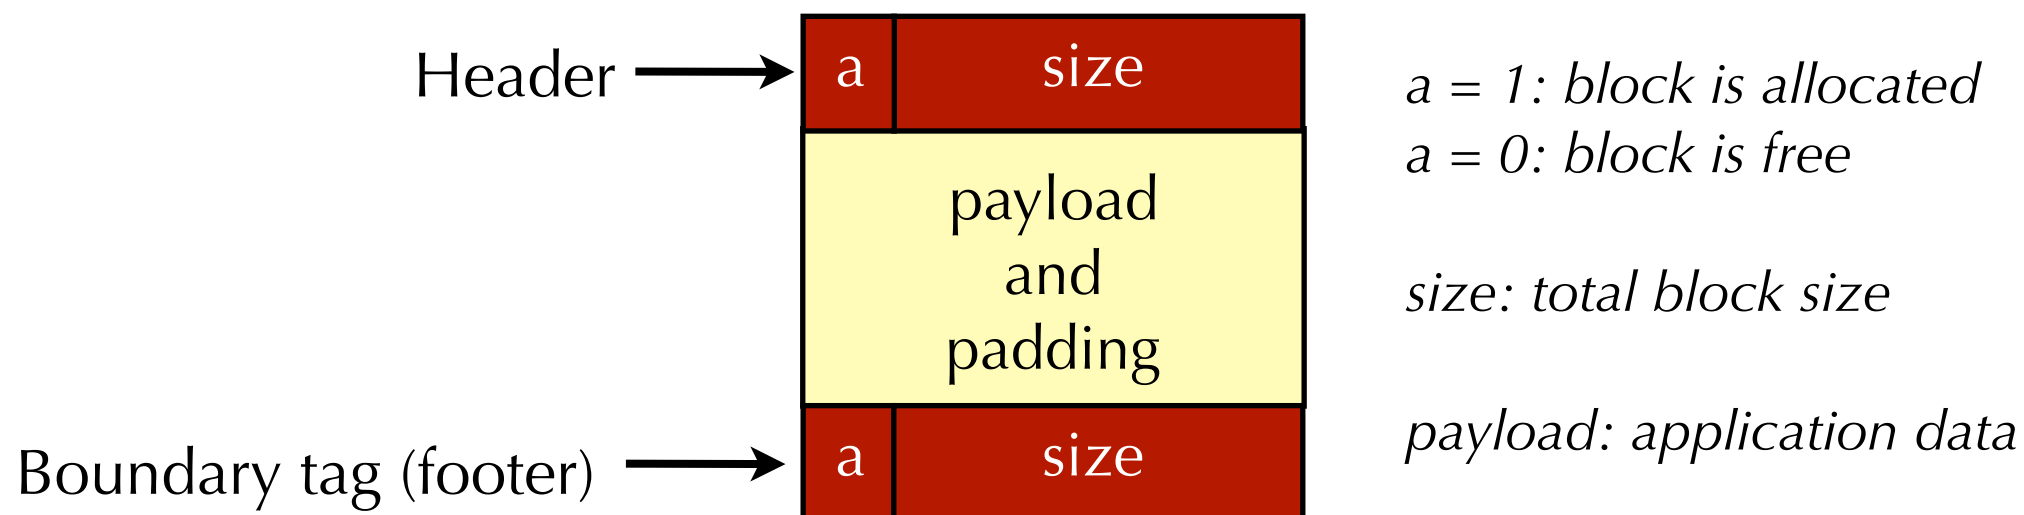
Implicit List: Bidirectional Coalescing

- **Boundary tags** [Knuth73]

- Also maintain the size/allocated word at end of free blocks (a footer)

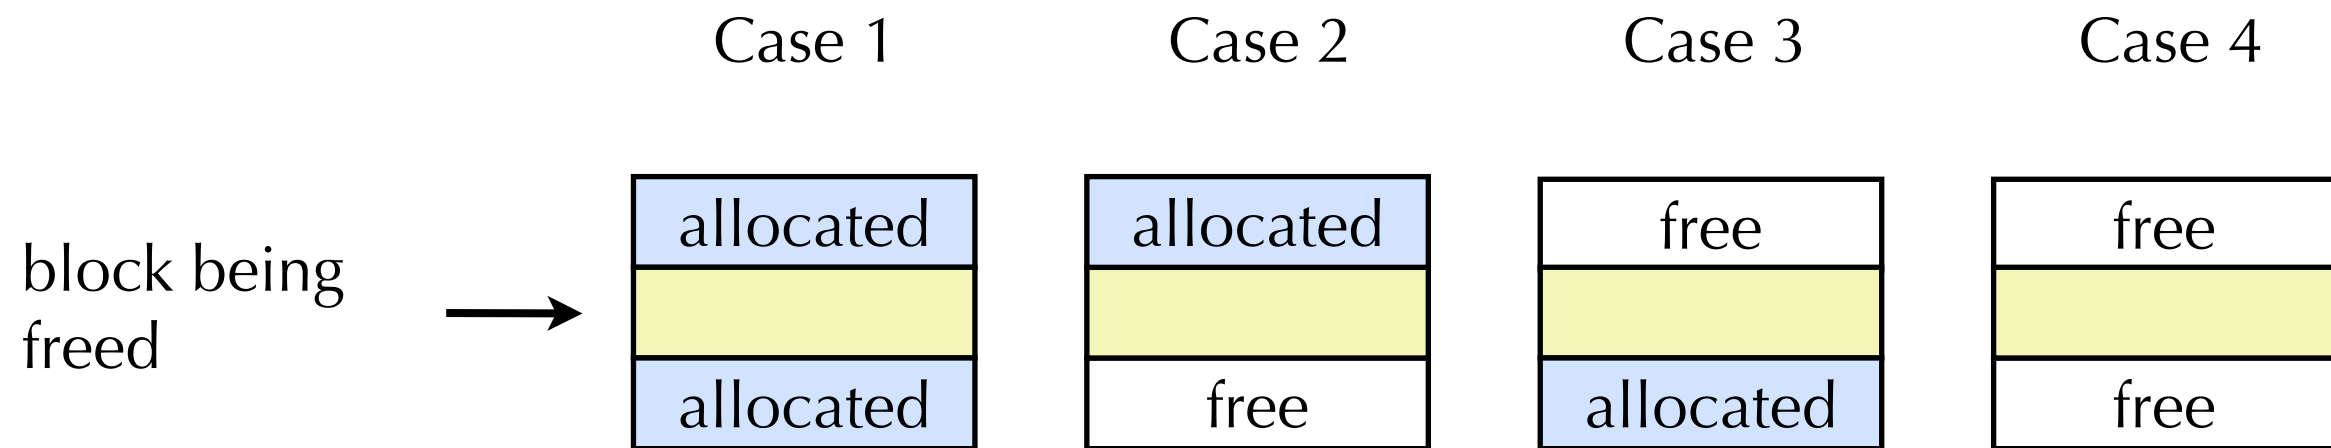


- Allows us to traverse the "block list" backwards, but requires extra space




- Important and general technique!

Constant Time Coalescing



Constant Time Coalescing (Case 1)


m1	1
m1	1
n	1
n	1
m2	1
m2	1



m1	1
m1	1
n	0
n	0
m2	1
m2	1

Constant Time Coalescing (Case 2)


m1	1
m1	1
n	1
n	1
m2	0
m2	0



m1	1
m1	1
n+m2	0
n+m2	0

Constant Time Coalescing (Case 3)


m1	0
m1	0
n	1
n	1
m2	1
m2	1



n+m1	0
n+m1	0
m2	1
m2	1

Constant Time Coalescing (Case 4)

m1	0
m1	0
n	1
n	1
m2	0
m2	0



n	0
n	0

Implicit Lists: Summary


- Implementation: Very simple.
- Allocation cost: Linear time worst case
- Free cost: Constant time, even with coalescing
- Memory usage: Depends on placement policy
 - First fit, next fit, or best fit
- Not used in practice for `malloc/free` because of linear time allocation.
- The concepts of splitting and boundary tag coalescing are general to *all* allocators.

Topics for today

- Dynamic memory allocation
 - Implicit vs. explicit memory management
- Performance goals
- Fragmentation
- Free block list management
 - Implicit free list
 - Explicit free list
 - Segregated lists
 - Tradeoffs


Explicit Free Lists

- Use an **explicit** data structure to track the free blocks
 - Just a doubly-linked list.
 - No pointers to or from allocated blocks: Can put the pointers into the payload!
 - Still need boundary tags, in order to perform free block coalescing.




Explicit Free Lists

- Note that free blocks need not be linked in the same order they appear in memory!
 - Free blocks can be chained together in any order.



Allocation using an explicit free list


Before:



```
p = malloc(size);
```

After:


split 2nd free block




Deallocation with an explicit free list

- Same idea as previously
 - Step 1: Mark block as free
 - Step 2: Coalesce adjacent free blocks
 - Step 3: Insert free block into the free list
- Where in the free list do we put a newly freed block?
 - LIFO (last-in-first-out) policy
 - Insert freed block at the beginning of the free list
 - Simple and constant time
 - Address-ordered policy
 - Insert freed blocks so that free list blocks are always in address order
 - Con: requires search
 - Pro: studies suggest fragmentation is lower than LIFO


Freeing With a LIFO Policy (Case 1)




- Insert the freed block at the root of the free block list



Freeing With a LIFO Policy (Case 2)




- Splice out predecessor block, coalesce both memory blocks and insert the new block at the root of the list




Freeing With a LIFO Policy (Case 3)

Before:




- Splice out successor block, coalesce both memory blocks and insert the new block at the root of the list


After:



Freeing With a LIFO Policy (Case 4)



- Splice out predecessor and successor blocks, coalesce all 3 memory blocks and insert the new block at the root of the list



Explicit List Summary




- Comparison to implicit list:
 - Allocation is linear time in number of **free blocks**
 - Implicit list allocation is linear time in the number of **total blocks**
- Slightly more complicated allocate and free since need to splice blocks in and out of the list
- Need some extra space for the links
 - 2 extra words needed for each free block
 - But these can be stored in the payload, since only needed for free blocks.

Topics for today

- Dynamic memory allocation
 - Implicit vs. explicit memory management
- Performance goals
- Fragmentation
- Free block list management
 - Implicit free list
 - Explicit free list
 - Segregated lists
 - Tradeoffs

Segregated List (seglist) Allocators

- Use a different free list for blocks of different sizes!



- Often have separate size class for every small size (4,5,6,...)
- For larger sizes typically have a size class for each power of 2

Seglist Allocator

- To allocate a block of size n :
 - Determine correct free list to use
 - Search that free list for block of size $m \geq n$
 - If an appropriate block is found:
 - Split block and place fragment on appropriate list
 - If no block is found, try next larger class
 - Repeat until block is found
- If no free block is found:
 - Request additional heap memory from OS (using `sbrk()` system call)
 - Allocate block of n bytes from this new memory
 - Place remainder as a single free block in largest size class.

Freeing with Seglist

- To free a block:
 - Mark block as free
 - Coalesce (if needed)
 - Place free block on appropriate sized list

Seglist advantages

- Advantages of seglist allocators
 - Higher throughput
 - Faster to find appropriate sized block: Look in the right list.
 - Better memory utilization
 - First-fit search of segregated free list approximates a best-fit search of entire heap.
 - Extreme case: Giving each size its own segregated list is equivalent to best-fit.

Topics for today

- Dynamic memory allocation
 - Implicit vs. explicit memory management
- Performance goals
- Fragmentation
- Free block list management
 - Implicit free list
 - Explicit free list
 - Segregated lists
 - Tradeoffs

Allocation Policy Tradeoffs

- Data structure for free lists
 - Implicit lists, explicit lists, segregated lists
 - Other structures possible, e.g. explicit free blocks in binary tree, sorted by size
- Placement policy: First fit, next fit, or best fit
 - Best fit has higher overhead, but less fragmentation.
- Splitting policy: When do we split free blocks?
 - Splitting leads to more internal fragmentation, since each block needs its own header.
- Coalescing policy: When do we coalesce free blocks?
 - **Immediate coalescing:** Coalesce each time free is called
 - **Deferred coalescing:** Improve free performance by deferring coalescing until needed.
 - E.g., While scanning the free list for `malloc()`, or when external fragmentation reaches some threshold.

Topics for next time

- Allocation requirements
- Common memory bugs
- Implicit memory management: Garbage collection