



**HARVARD**

School of Engineering  
and Applied Sciences

# Dynamic Memory Allocation

*CS61, Lecture 11*

Prof. Stephen Chong

October 6, 2011

# Announcements 1/2

- Reminder: No section on Monday
  - Monday sections have been rescheduled
  - See website for details
  - Please attend whichever section you can make
- Homework 3 (Buffer bomb) due today
- Homework 4 (malloc) released
  - Design checkpoint due Thursday 13 October, 10pm
  - Final submission due Thursday 20 October, 11:59pm

# Announcements 2/2

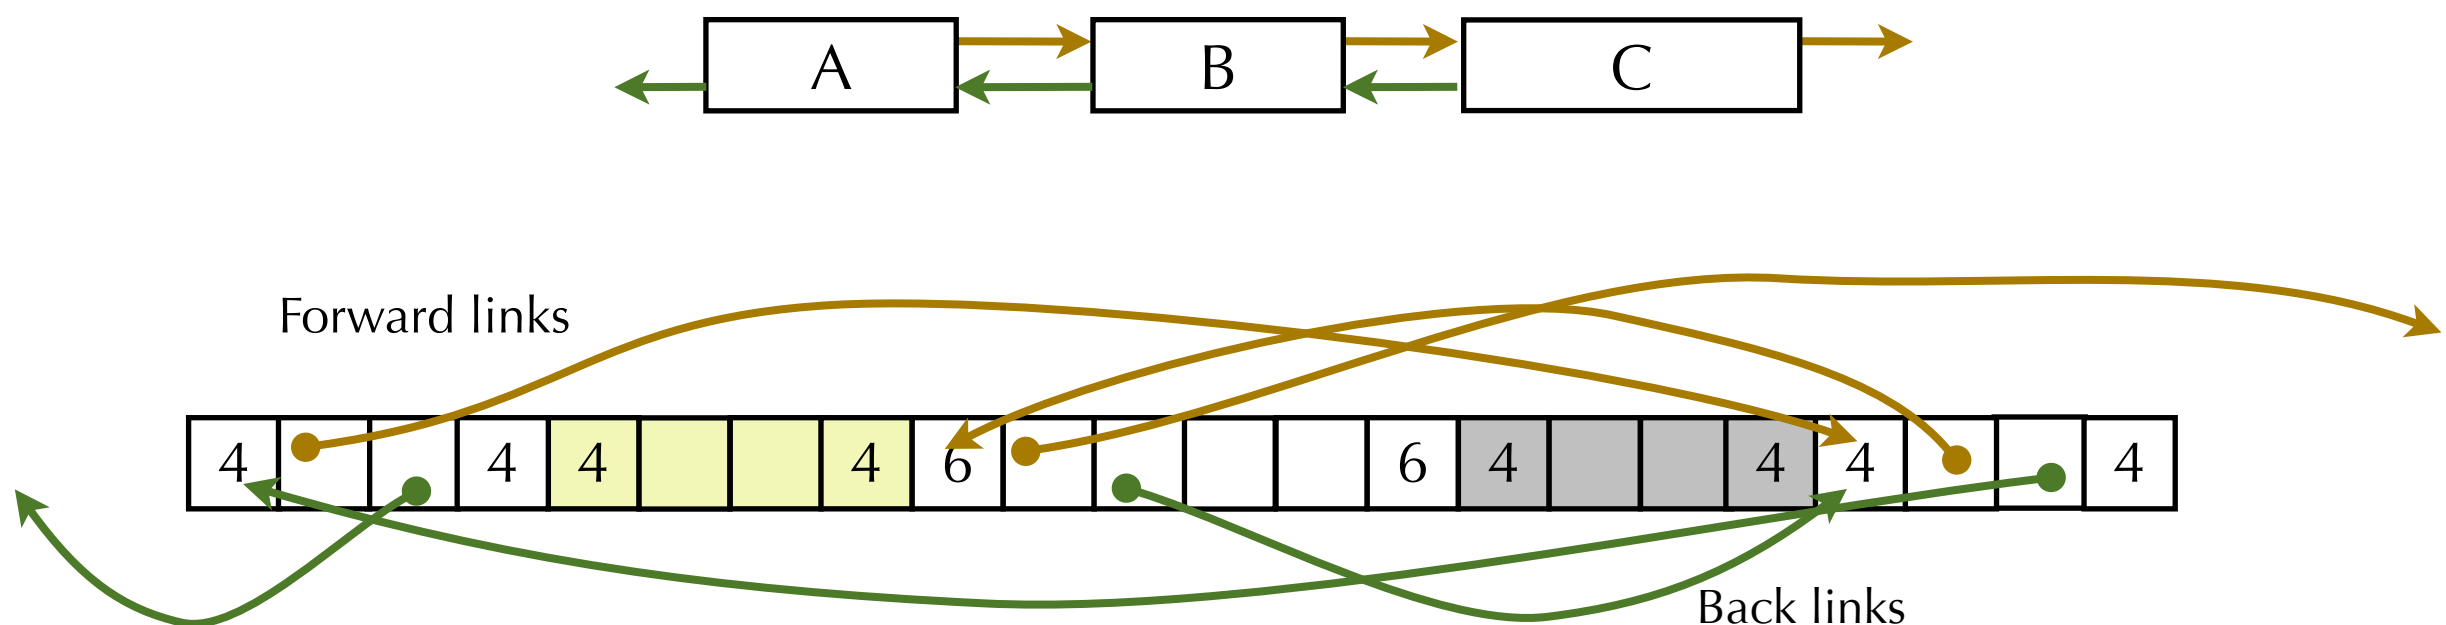
- Interested in concentrating in CS?
  - Don't wait until last minute!
    - Sophomores must declare concentration by Nov 16
  - Prof. Harry Lewis is the Computer Science DUS
    - Office hours posted on his website
    - He is happy to talk about the concentration, or sign you up!
  - Fill in departmental study plan
    - <http://studyplan.seas.harvard.edu/ug/>
    - Going through exercise will help you understand requirements
    - Can always revise it later
  - To actually enter concentration, need to fill out a form on Registrar's website, and get Prof. Lewis' signature

# Today

- Free block list management
  - Implicit free list
  - Explicit free list
  - Segregated lists
  - Tradeoffs
- Alignment issues
- Example allocator implementation: `dlmalloc`
- Common memory problems
- Garbage collection
  - Mark and sweep
  - Generational GC
  - Reference counting



# Explicit Free Lists

- Use an explicit data structure to track the free blocks
  - In this case, a doubly-linked free list
  - Pointers for list just live in the payload
    - Payload is unused for free blocks!



# Segregated List (seglist) Allocators

- Use a different free list for blocks of different sizes!



- Often have separate size class for every small size (4,5,6,...)
- For larger sizes typically have a size class for each power of 2

# Seglist Allocator

- To allocate a block of size  $n$ :
  - Determine correct free list to use
  - Search that free list for block of size  $m \geq n$
  - If an appropriate block is found:
    - Split block and place fragment on appropriate list
  - If no block is found, try next larger class
  - Repeat until block is found
- If no free block is found:
  - Request additional heap memory from OS (using `sbrk()` system call)
  - Allocate block of  $n$  bytes from this new memory
  - Place remainder as a single free block in largest size class.

# Freeing with Seglist

- To free a block:
  - Mark block as free
  - Coalesce (if needed)
  - Place free block on appropriate sized list



# Seglist advantages

- Advantages of seglist allocators
  - Higher throughput
    - Faster to find appropriate sized block: Look in the right list.
  - Better memory utilization
    - First-fit search of segregated free list approximates a best-fit search of entire heap.
    - Extreme case: Giving each size its own segregated list is equivalent to best-fit.

# Allocation Policy Tradeoffs

- Data structure for free lists
  - Implicit lists, explicit lists, segregated lists
  - Other structures possible, e.g. explicit free blocks in binary tree, sorted by size
- Placement policy: First fit, next fit, or best fit
  - Best fit has higher overhead, but less fragmentation.
- Splitting policy: When do we split free blocks?
  - Splitting leads to more internal fragmentation, since each block needs its own header.
- Coalescing policy: When do we coalesce free blocks?
  - **Immediate coalescing:** Coalesce each time free is called
  - **Deferred coalescing:** Improve free performance by deferring coalescing until needed.
    - E.g., While scanning the free list for `malloc()`, or when external fragmentation reaches some threshold.

# Today

- Free block list management
  - Implicit free list
  - Explicit free list
  - Segregated lists
  - Tradeoffs
- Alignment issues
- Example allocator implementation: `dlmalloc`
- Common memory problems
- Garbage collection
  - Mark and sweep
  - Generational GC
  - Reference counting


# Alignment issues

- Most `malloc()` implementations ensure that the returned pointer is **aligned** to an 8-byte boundary.
  - This is to ensure that if the pointer is used to store a struct, it will be properly aligned.

```
struct mystruct *foo;  
void *p;  
  
p = malloc(sizeof(struct mystruct));  
foo = (struct mystruct *)p;
```

# Alignment issues

- Implication: malloc() may have to pad the block that it allocates



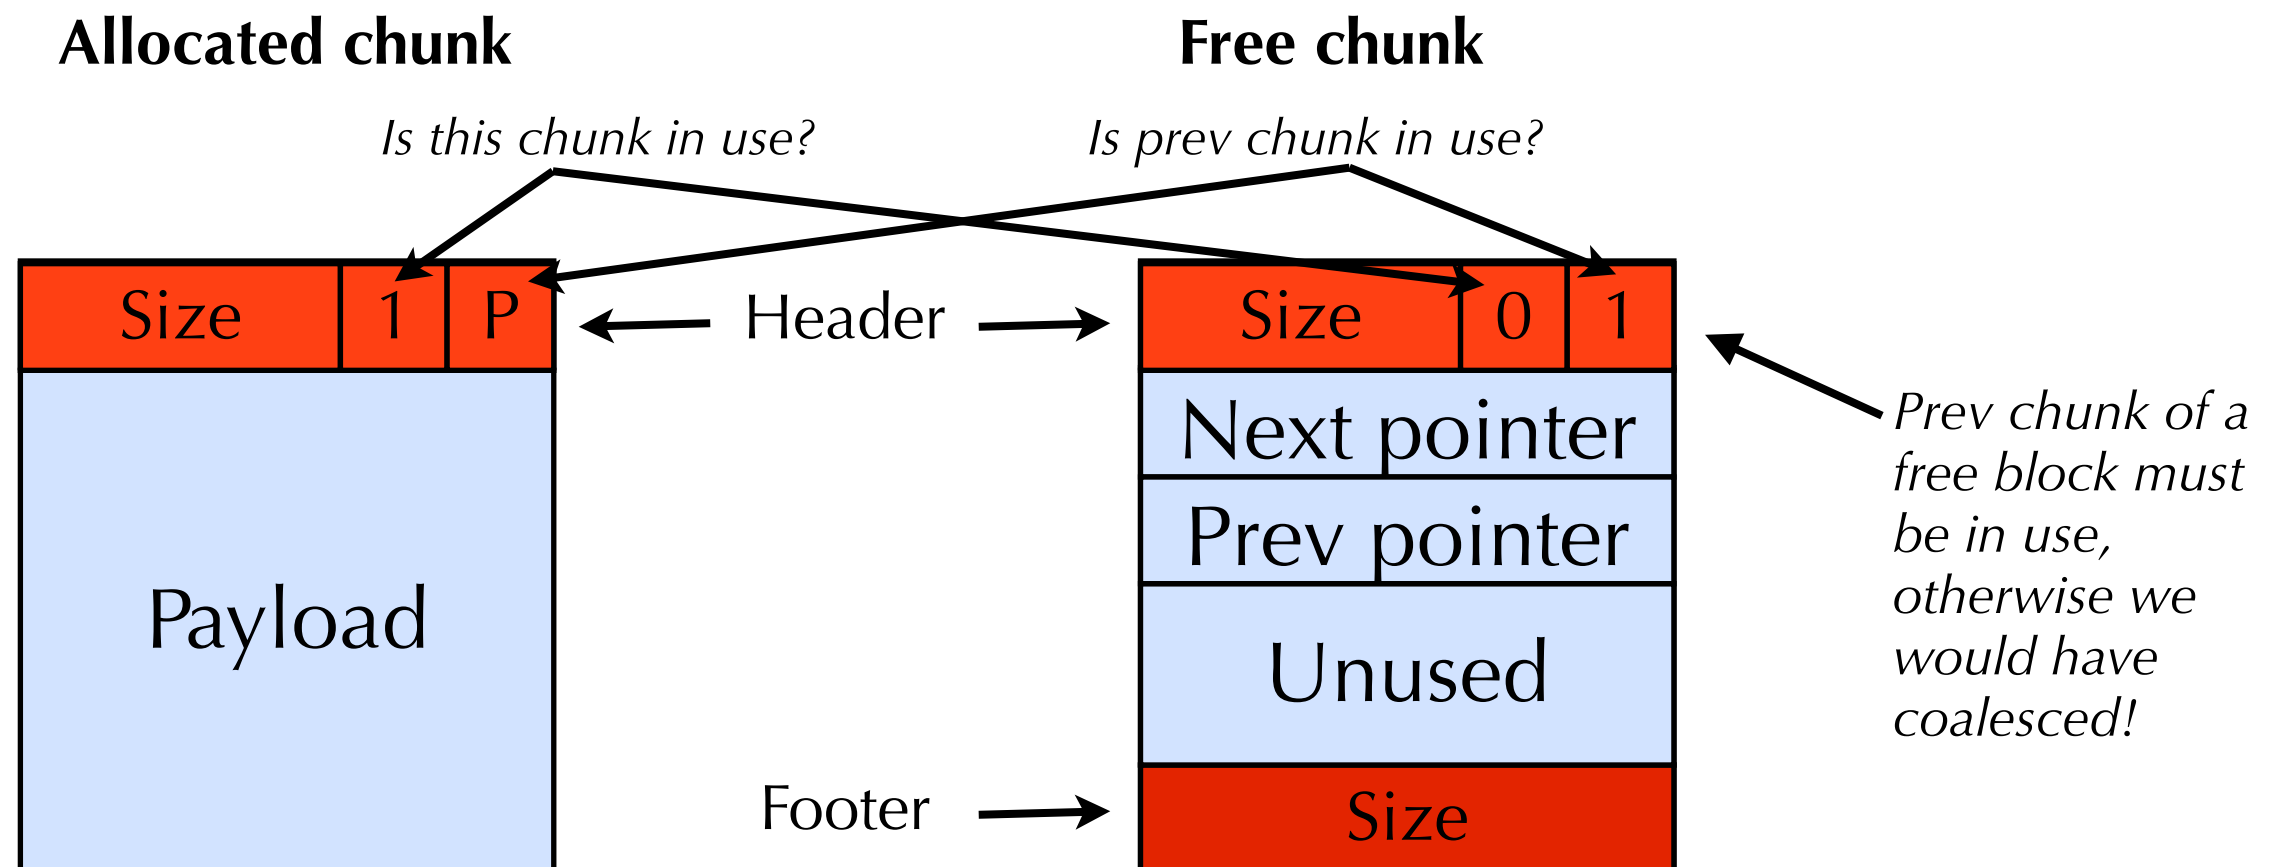
- Padding can often be “recycled” as header, boundary tags, etc.
  - (Not shown in the above example.)
  - Key is that the **payload of the buffer is appropriately aligned.**

# Doug Lea allocator

- `dlmalloc`
  - Fast and efficient general purpose allocator
    - Basis of `glibc` allocator
  - Since 1992
- Essentially best-fit
  - Ties are broken in least-recently-used order
    - Reduces fragmentation
  - Deviates for requests less than 256 bytes
    - Prefer space adjacent to previous small request
    - Break ties in most-recently-used order
      - ▶ Increase locality of series of small allocations
- All operations bounded by constant factor (in # bits of `size_t`)


# Doug Lea allocator

- Uses blocks (aka “chunks”) with boundary tags
  - 8 byte alignment
  - Free chunks have header and footer
  - Allocated chunks have only header
  - A chunk is at least 16 bytes (4 for header + 4 for footer + 8 byte alignment requirement)



# Bins

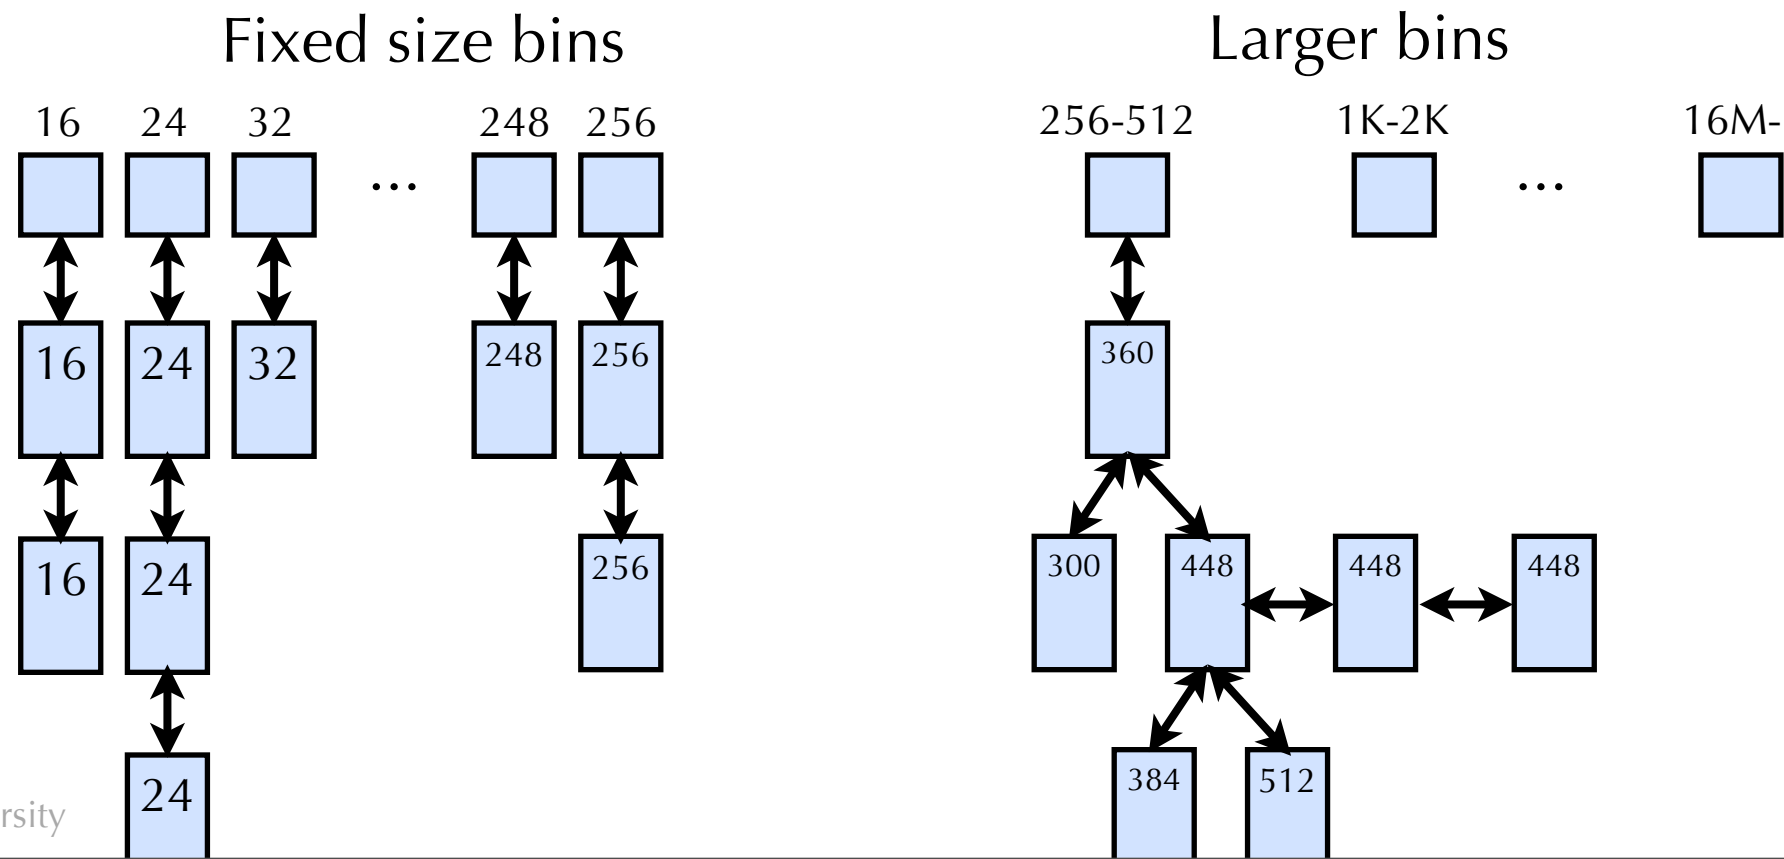
- Segregated lists
  - Bins for small (fixed size) chunks
    - Chunks stored in circular doubly linked list
  - Bins for large chunks ( $> 256$  bytes)
    - Chunks stored as special kind of binary tree
    - Chunks of same size stored as doubly linked list off node





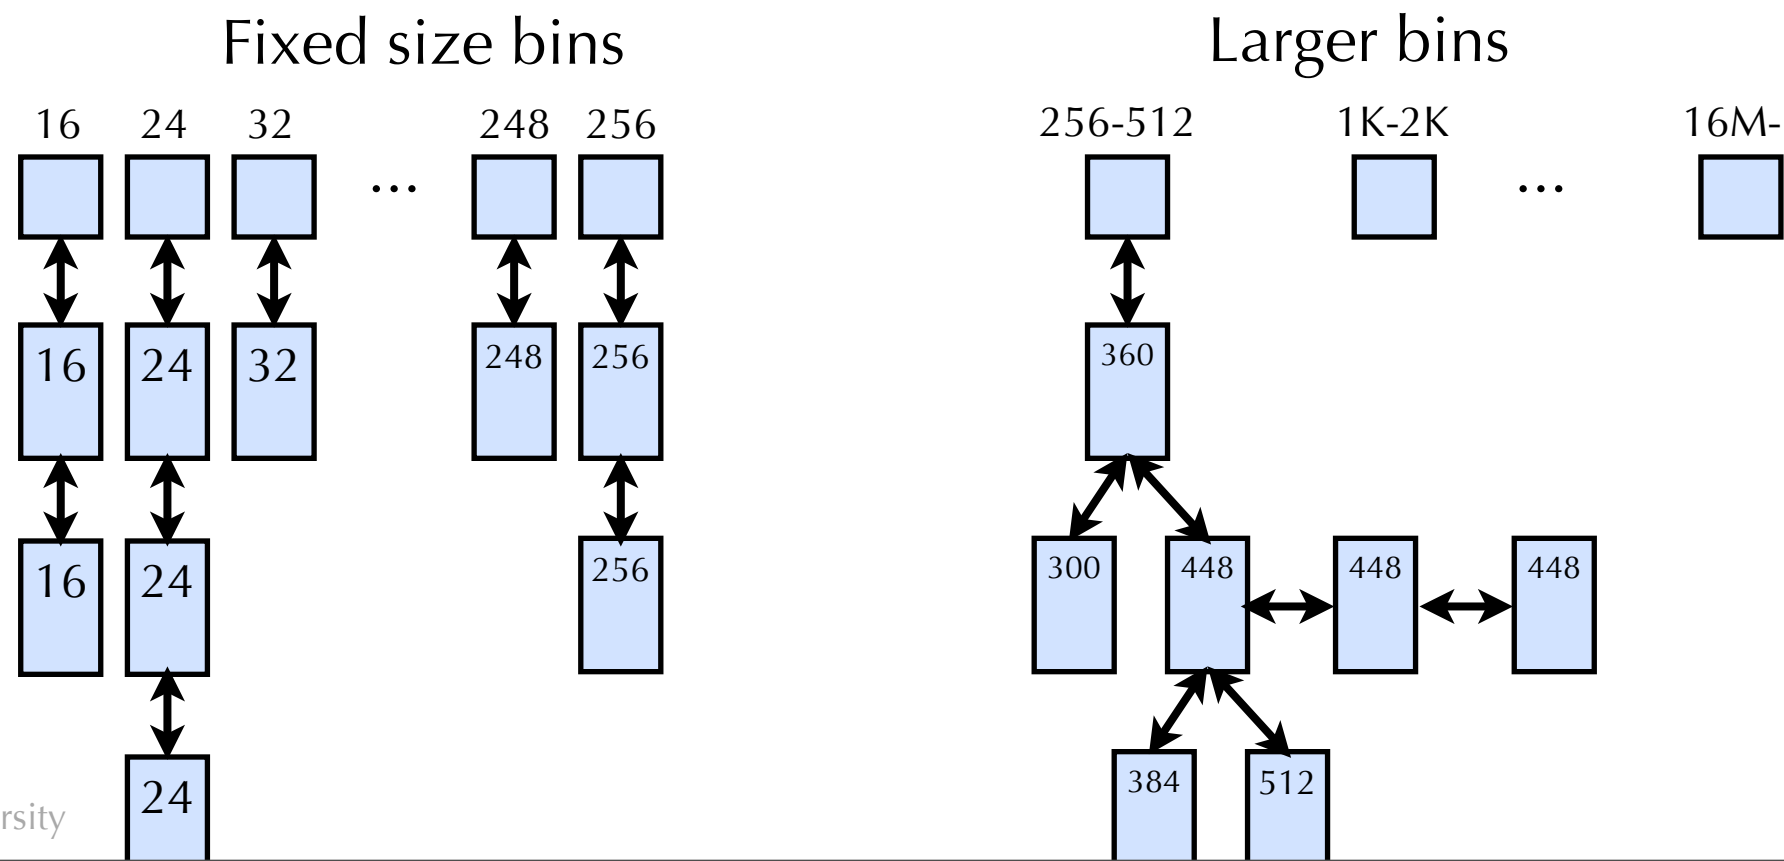
# Allocation in dlmalloc

- Find an appropriate size bin
- Look through the structure for the best fit.
  - Lists are arranged so that
    - bins for small chunks are in most-recently-used order
      - ▶ To increase locality
    - lists for large chunks of same size are in least-recently-used order
      - ▶ To decrease fragmentation



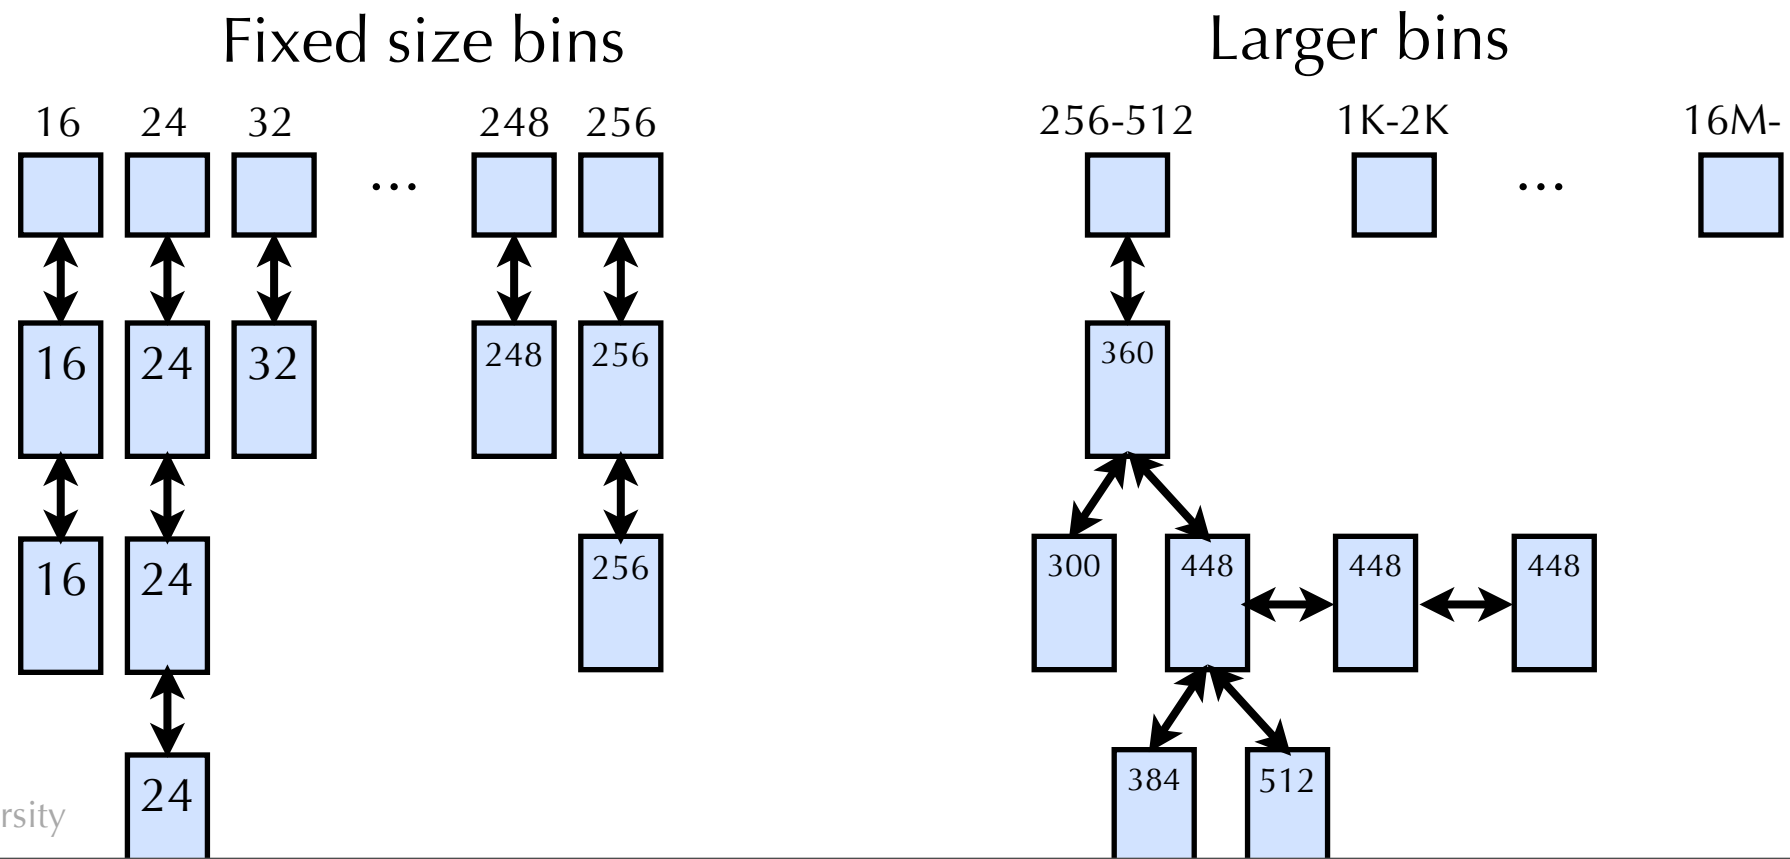
# Allocation in dlmalloc

- Algorithm a bit more complicated
- Some other heuristics used for allocation
  - e.g., maintain a “designated victim”, the remainder from the last split, and use it for small requests if possible
- Last free block of heap treated specially



# Free in dlmalloc

- Always coalesce on free



# Today

- Free block list management
  - Implicit free list
  - Explicit free list
  - Segregated lists
  - Tradeoffs
- Alignment issues
- Example allocator implementation: `dlmalloc`
- Common memory problems
- Garbage collection
  - Mark and sweep
  - Generational GC
  - Reference counting

# Memory-Related Perils and Pitfalls

- Dereferencing bad pointers
- Reading uninitialized memory
- Overwriting memory
- Referencing nonexistent variables
- Freeing blocks multiple times
- Referencing freed blocks
- Failing to free blocks

# Dereferencing Bad Pointers

- The classic `scanf` bug

```
int val;  
scanf("%d", val);
```

**Not a pointer!**

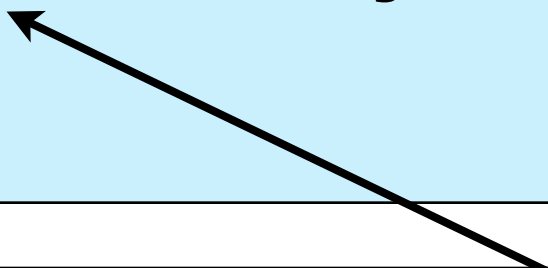


# Reading Uninitialized Memory

- Assuming that heap data is initialized to zero

```
/* return y = Ax */
int *matvec(int **A, int *x) {
    int *y = malloc(N*sizeof(int));
    int i, j;

    for (i=0; i<N; i++)
        for (j=0; j<N; j++)
            y[i] += A[i][j]*x[j];
    return y;
}
```



**y[i] not necessarily initialized to zero**

# Overwriting Memory

- Off-by-one error

```
int **p;  
  
p = malloc(N*sizeof(int *));  
  
for (i=0; i<= N; i++) {  
    p[i] = malloc(M*sizeof(int));  
}
```

**Goes through loop N+1 times!**



# Overwriting Memory

- Not checking the max string size

```
char s[8];  
gets(s);
```

**Buffer overflow vulnerability!**

# Overwriting Memory

- Misunderstanding pointer arithmetic

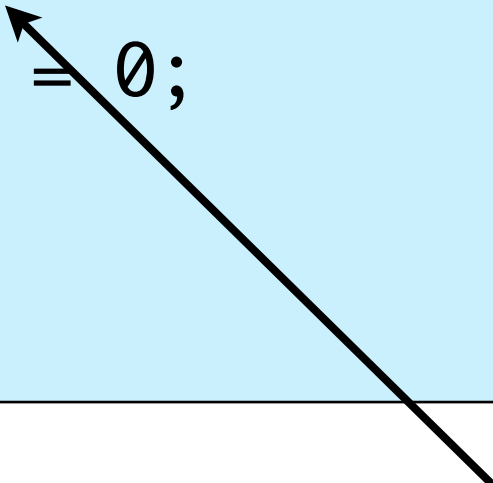
```
int *search(int *p, int val) {  
    while (*p && *p != val)  
        p += sizeof(int);  
    return p;  
}
```

**Should be p++!  
Incorrectly reads every fourth  
element in array**

# Overwriting Memory

- Referencing a pointer instead of the object pointed to

```
void trim(char *s, int *length) {  
    if (*length > 0) {  
        if (s[(*length)-1] == '\n') {  
            *length--;  
            s[*length] = 0;  
        }  
    }  
}
```



**Should be (\*length)--!**  
**\*length-- is \*(length--), which modifies pointer.**

# Referencing Nonexistent Variables

- Forgetting that local variables disappear when a function returns

```
int *foo () {  
    int val;  
  
    return &val;  
}
```

**Pointer into stack frame. Valid memory address, but may point to another function's stack frame.**

# Freeing Blocks Multiple Times

- Nasty!

```
x = malloc(N*sizeof(int));  
...manipulate x...  
free(x);  
  
y = malloc(M*sizeof(int));  
...manipulate y...  
free(x);
```

**Free same pointer twice!**

- What would this do to an explicit free list?
  - Can lead to security vulnerabilities where attacker controls some region of the heap!

# Referencing Freed Blocks

- Evil!

```
x = malloc(N*sizeof(int));  
...manipulate x...  
free(x);  
...  
y = malloc(M*sizeof(int));  
for (i=0; i<M; i++)  
    y[i] = x[i]++;
```

**Accessing a pointer that  
has already been freed!**

# Failing to Free Blocks (Memory Leaks)

- Slow, long-term killer!

```
foo() {  
    int *x = malloc(N*sizeof(int));  
    ...  
    return;  
}
```

**Object can't be accessed  
ever again, but is not freed.**

# Failing to Free Blocks (Memory Leaks)

- Freeing only part of a data structure

```
struct list {  
    int val;  
    struct list *next;  
};  
  
foo() {  
    struct list *head = malloc(sizeof(struct list));  
    head->val = 0;  
    head->next = NULL;  
    ...create and manipulate the rest of the list...  
    free(head);  
    return;  
}
```

**Only head free, not rest  
of structure.**



# Dealing With Memory Bugs

- Conventional debugger (gdb)
  - Good for finding bad pointer dereferences
  - Hard to detect the other memory bugs
- Debugging malloc (CSRI UToronto malloc)
  - Wrapper around conventional malloc
  - Detects memory bugs at malloc and free boundaries
    - Memory overwrites that corrupt heap structures
    - Some instances of freeing blocks multiple times
    - Memory leaks
  - Cannot detect all memory bugs
    - Overwrites into the middle of allocated blocks
    - Freeing block twice that has been reallocated in the interim
    - Referencing freed blocks

# Dealing With Memory Bugs (cont.)

- Binary translator: valgrind, Purify
  - Powerful debugging and analysis technique
  - Rewrites text section of executable object file
  - Can detect all errors as debugging malloc
  - Can also check each individual reference at runtime
    - Bad pointers
    - Overwriting
    - Referencing outside of allocated block
- Use a different language
  - Some languages present a different model of memory to programmer
  - Avoid some-to-all of the common memory problems described here

# Today

- Free block list management
  - Implicit free list
  - Explicit free list
  - Segregated lists
  - Tradeoffs
- Alignment issues
- Example allocator implementation: `dlmalloc`
- Common memory problems
- Garbage collection
  - Mark and sweep
  - Generational GC
  - Reference counting

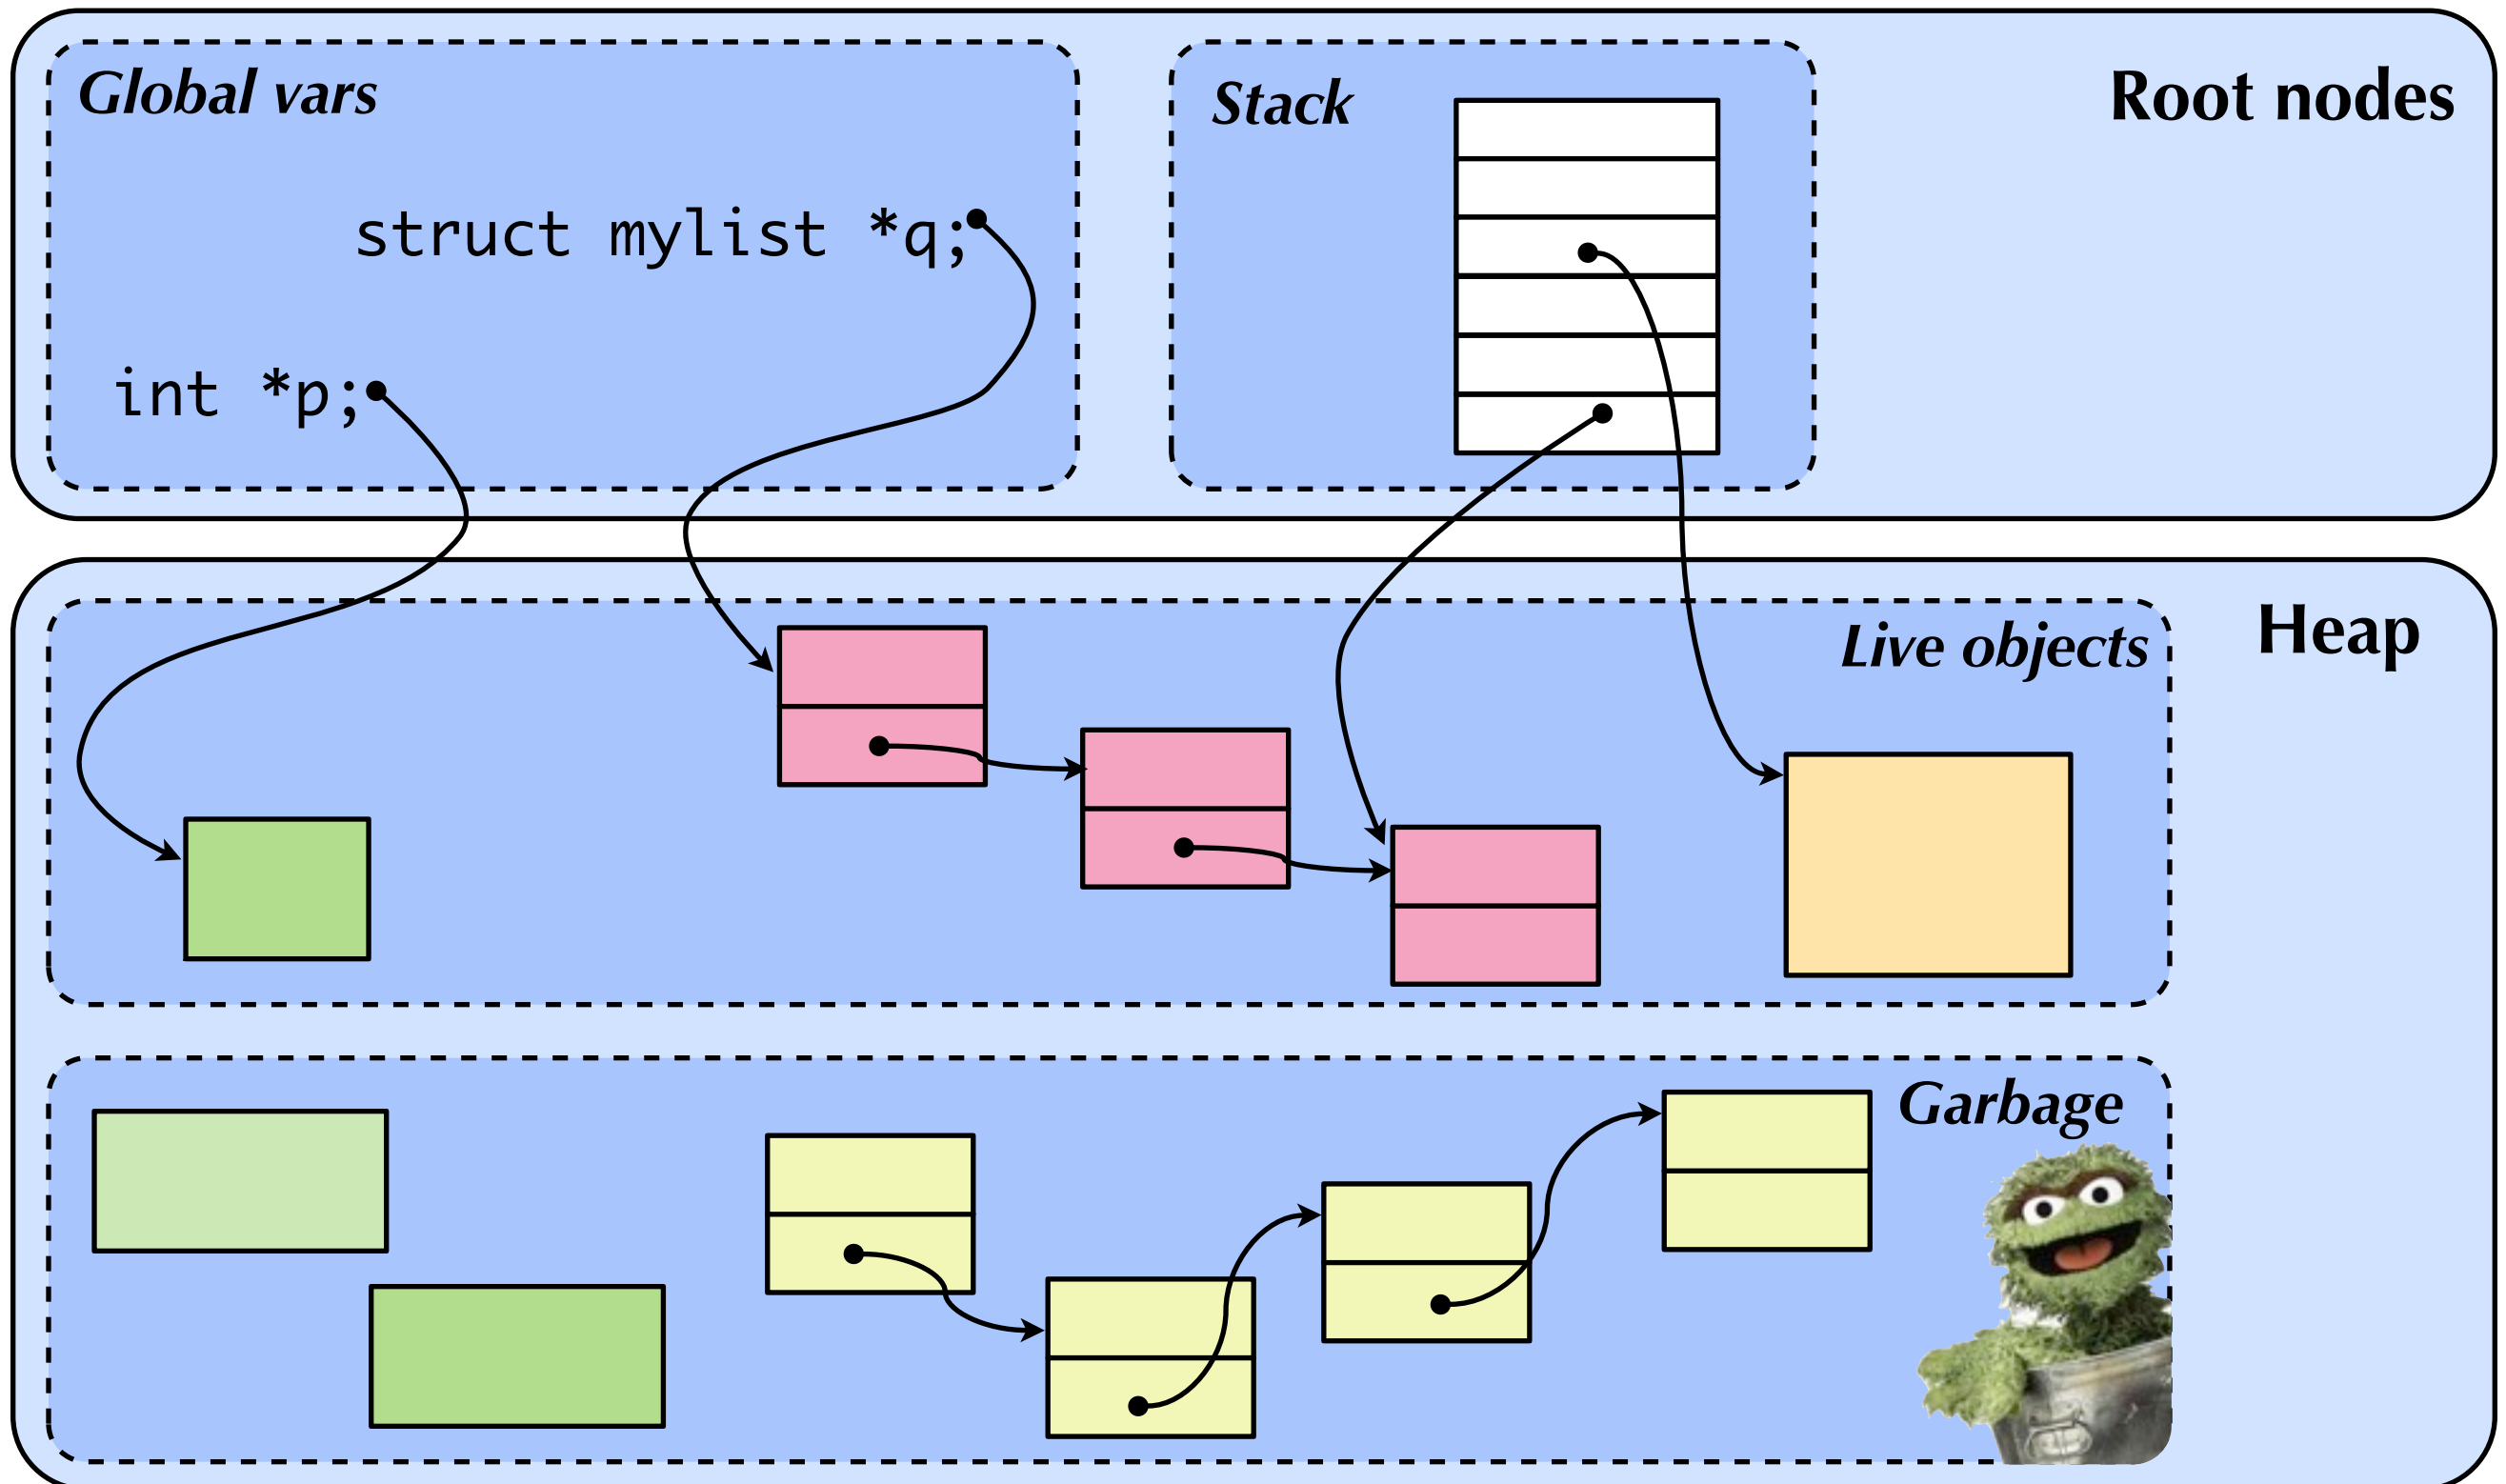
# Implicit Memory Management: Garbage Collection

- Garbage collection: automatic reclamation of heap-allocated storage
  - Application never has to explicitly free memory!

```
void foo() {  
    int *p = malloc(128);  
    return; /* p block is now garbage */  
}
```

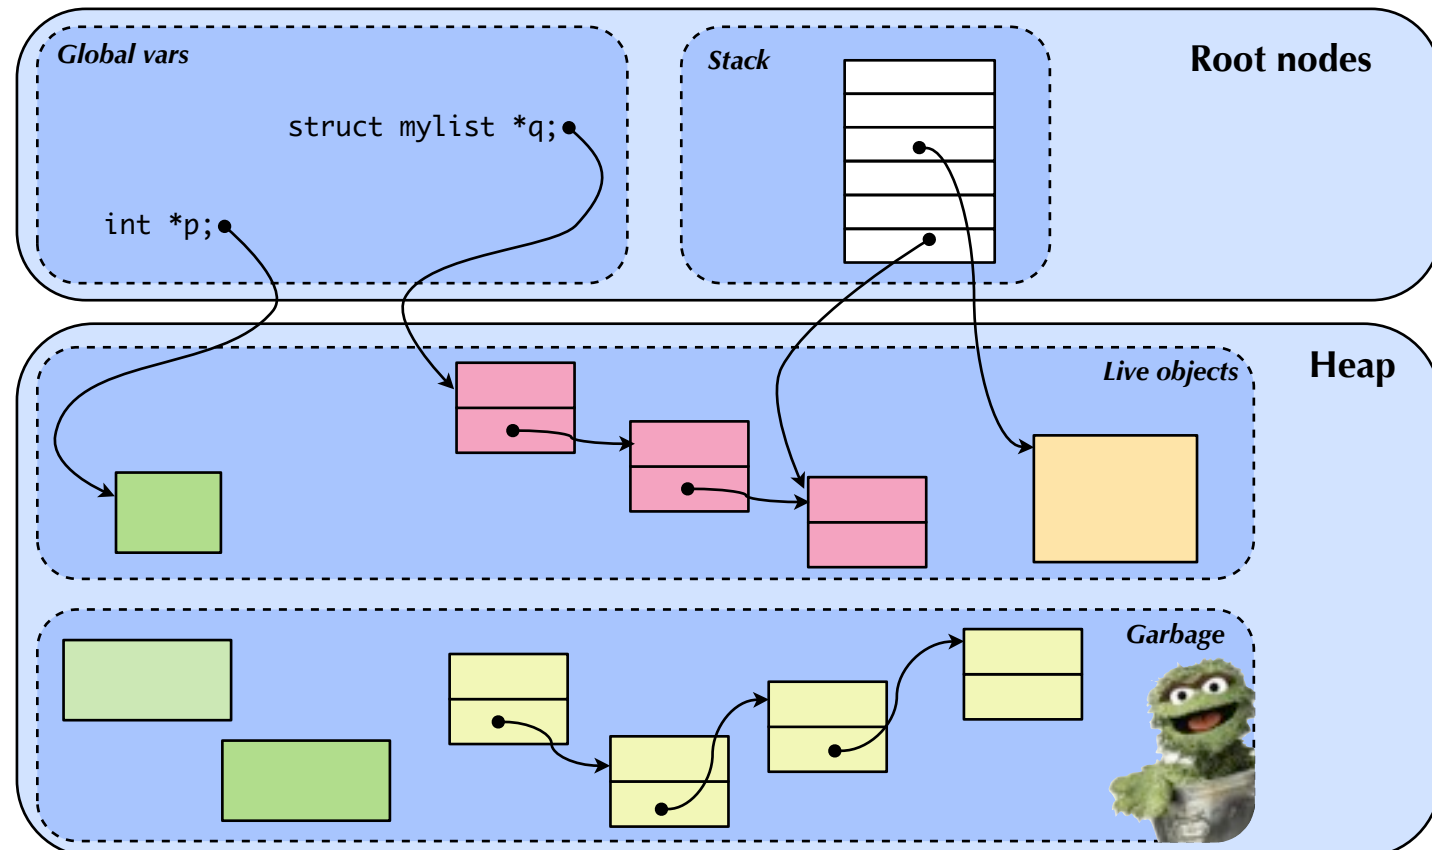
- Common in functional languages, scripting languages, and modern object oriented languages:
  - Lisp, ML, Java, Perl, Python, etc.
- Variants (conservative garbage collectors) also exist for C and C++
  - These do not generally manage to collect all garbage though.

# Garbage collection concepts



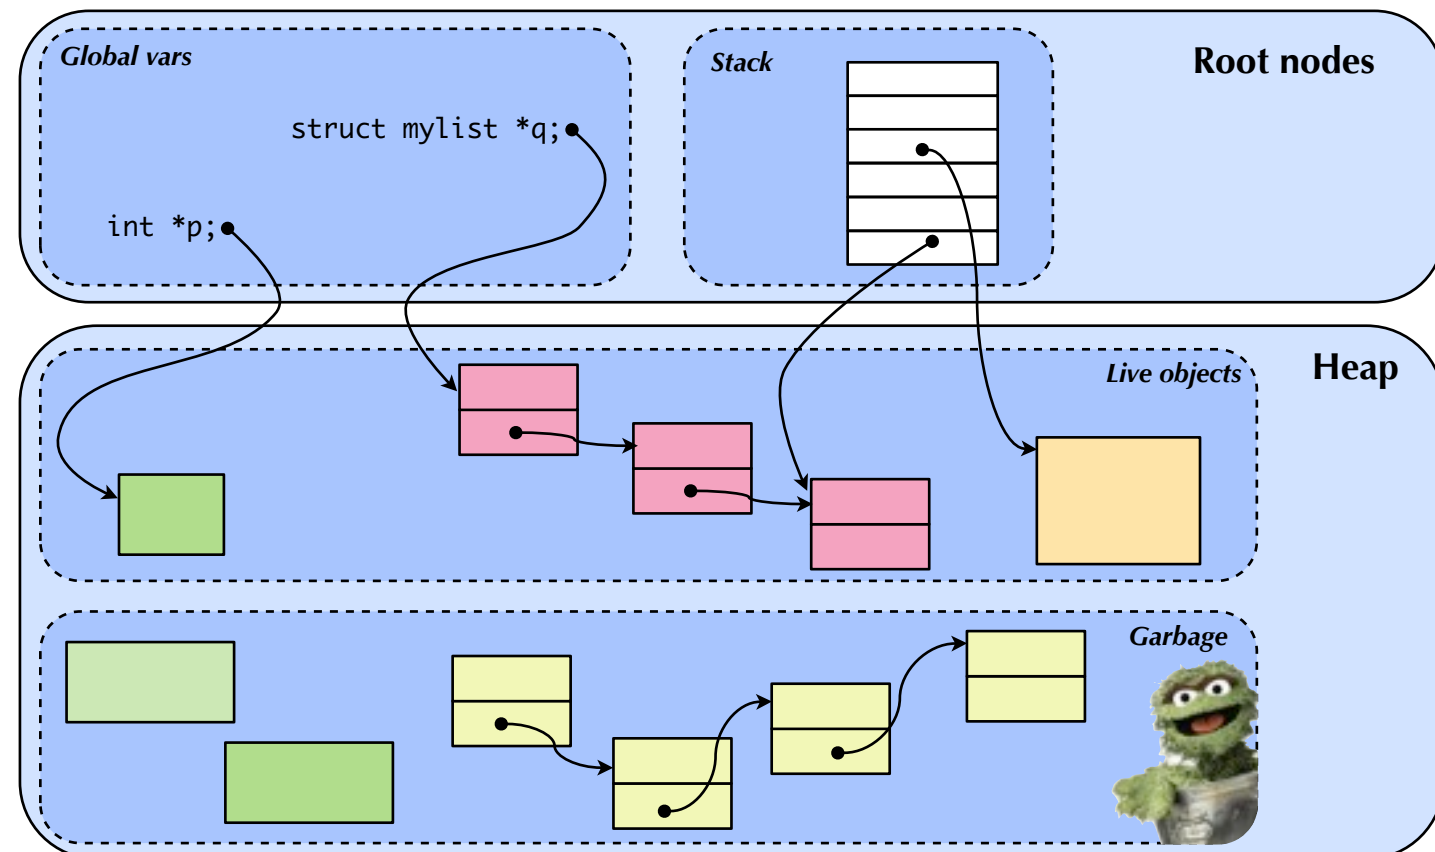
# Garbage collection concepts

- A block of memory is **reachable** if there is some path from a root node to the block.
  - If a block is reachable, it's **live** and should not be reclaimed.
  - If a block is not reachable, it's **garbage**.



# Garbage collection to the rescue

- Solves some of common memory problems
  - Double free, access after free, some memory leaks, ...
- Say we have a block that is reachable, but never used by the program. What happens?
  - **Memory leak!**
  - Garbage collectors assume all reachable objects can be used in the future.





# Challenges

- Challenge: How do we know where the pointers are?
- Global vars are easy: We know the type at compile time.
- What about the stack?
  - GC needs to know which slots on the stack are pointers, and which are not.
  - Depends on call chain of functions at runtime, and how each function stores local vars on the stack.
- What about pointers **between** blocks of memory?
  - Linked lists, trees, etc.
  - GC needs to know where the pointers are internal to each block!



# Fake pointers

- In C, you can create a pointer directly from an integer...

```
int i = 0x40b0cdf;  
*(int *)i = 42;
```

- But type of `i` is `int`! How do we know it is used as a pointer?
- Conservative garbage collectors treat all bit patterns in memory that could be pointers as pointers.
  - If contents of location could be memory address in heap, assume it is
- Problem: Leads to “false” memory leaks
  - If `int i = 0x40b0cdf` really is just an integer value, not a pointer, memory at location `0x40b0cdf` will not be reclaimed!

# Fake pointers

- In C, you can create a pointer directly from an integer...

```
int i = 0x40b0cdf;  
*(int *)i = 42;
```

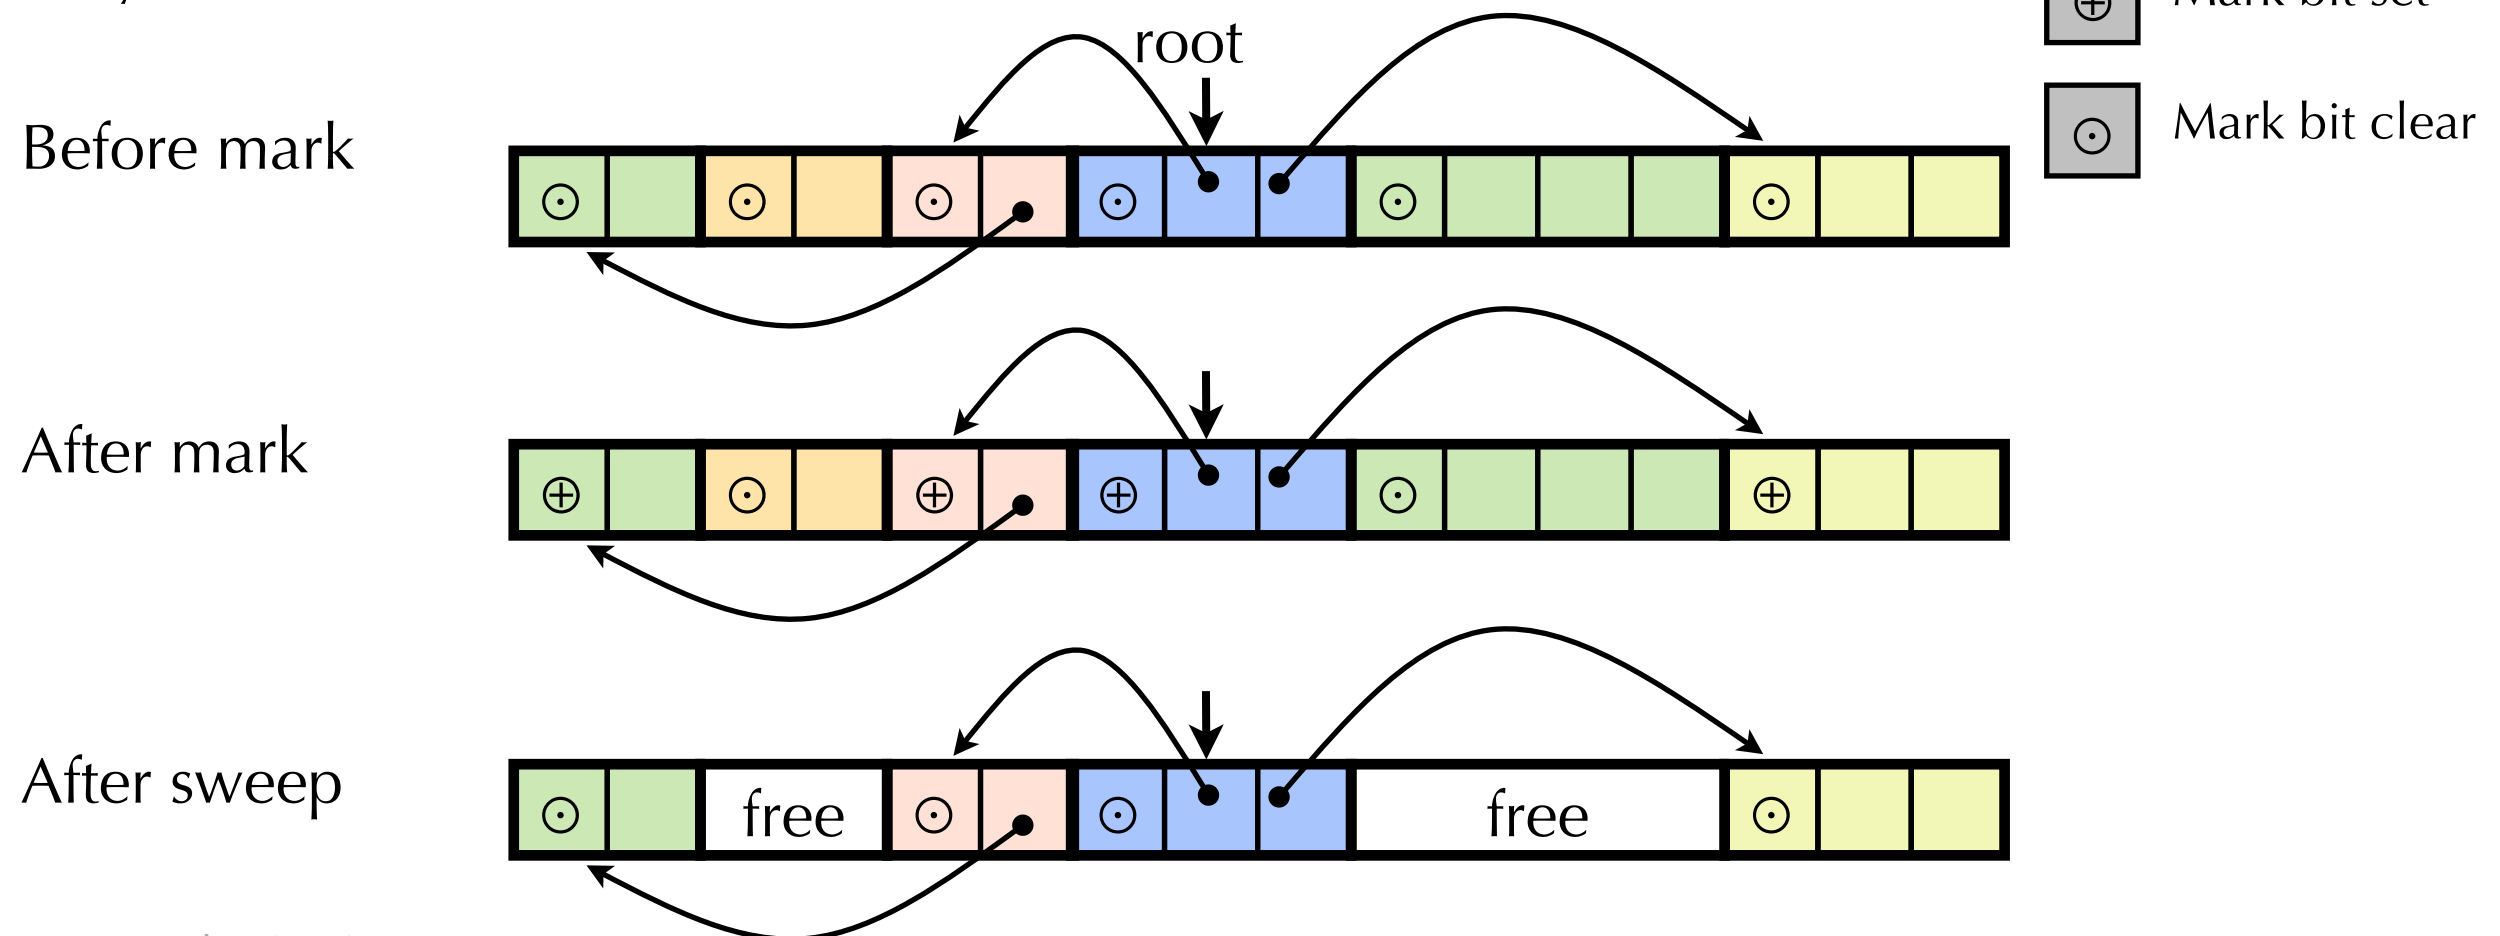
- Solution: Don't let program create pointers from ints.
  - Java (and other languages) never let you “forge” pointers like this.
  - Allows for precise garbage collection: GC always knows where pointers are, since pointers are “special”.

# Garbage collection techniques

- We'll consider three techniques
  - Mark-and-sweep
  - Generational garbage collection
  - Reference counting

# Mark and Sweep Collecting

- Idea: Use a mark bit in the header of each block
  - GC scans all memory objects (starting at roots), and sets mark bit on each reachable memory block.
- Sweeping:
  - Scan over all memory blocks.
  - Any block without the mark bit set is freed



# Mark and Sweep

- Advantages
  - Easily handles cycles
  - Low overhead
- Disadvantages
  - Need to scan all objects
    - Slow if large object set
- Observation: Most objects have short lifespans
  - Most objects created since the last garbage collection
  - If an object survives a few garbage collection cycles, it is likely to survive longer.

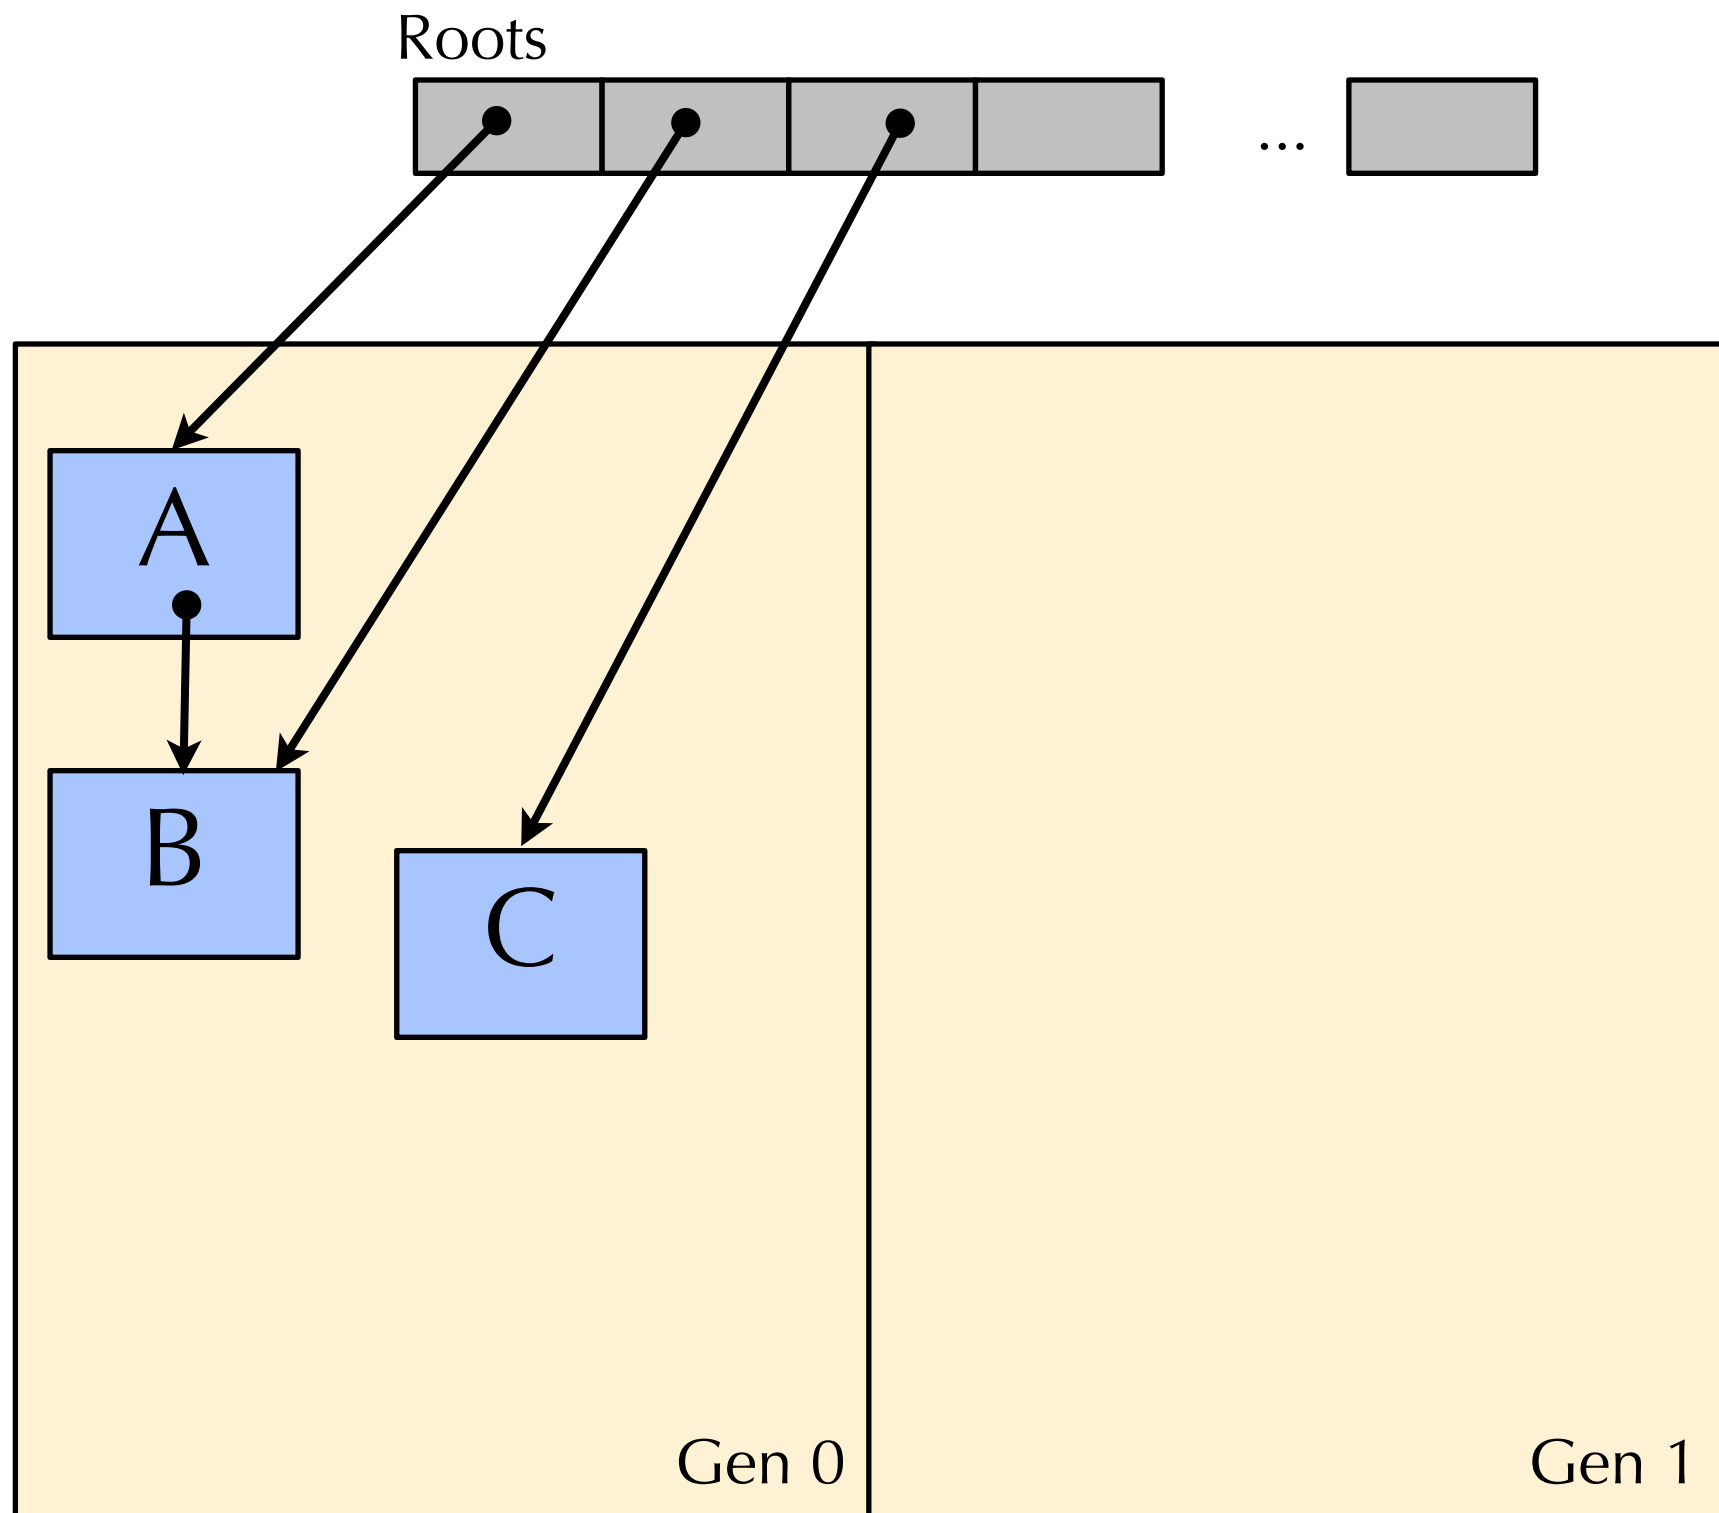
# Generational Garbage Collection

- Key idea: divide objects into **generations**, and garbage collect newer generation more frequently than older generations
- Takes advantage that most objects do not live long
- Used in Java and .NET

# Generational GC Algorithm

- Divide objects into generations (Gen0, Gen1...)
  - New objects go in Gen0
- Garbage collect younger generations aggressively
  - Perform GC on Gen(n)
    - mark and sweep, or something else
  - Promote surviving objects of Gen(n) to Gen(n+1)
- May occasionally need a “full” GC
- What does “promotion” mean?
  - Some implementations assign regions of memory to generations
  - Promotion to older generation means copy the object to a different region of memory
    - This is difficult in C. Why?

# Generational Garbage Collection




*We start executing...*


*New object go in  
Gen0*




# Generational Garbage Collection




# Generational Garbage Collection



# Generational Garbage Collection




# Generational Garbage Collection




*Program continues to execute. Many short lived objects.*

# Generational Garbage Collection



# Generational Garbage Collection



# Today

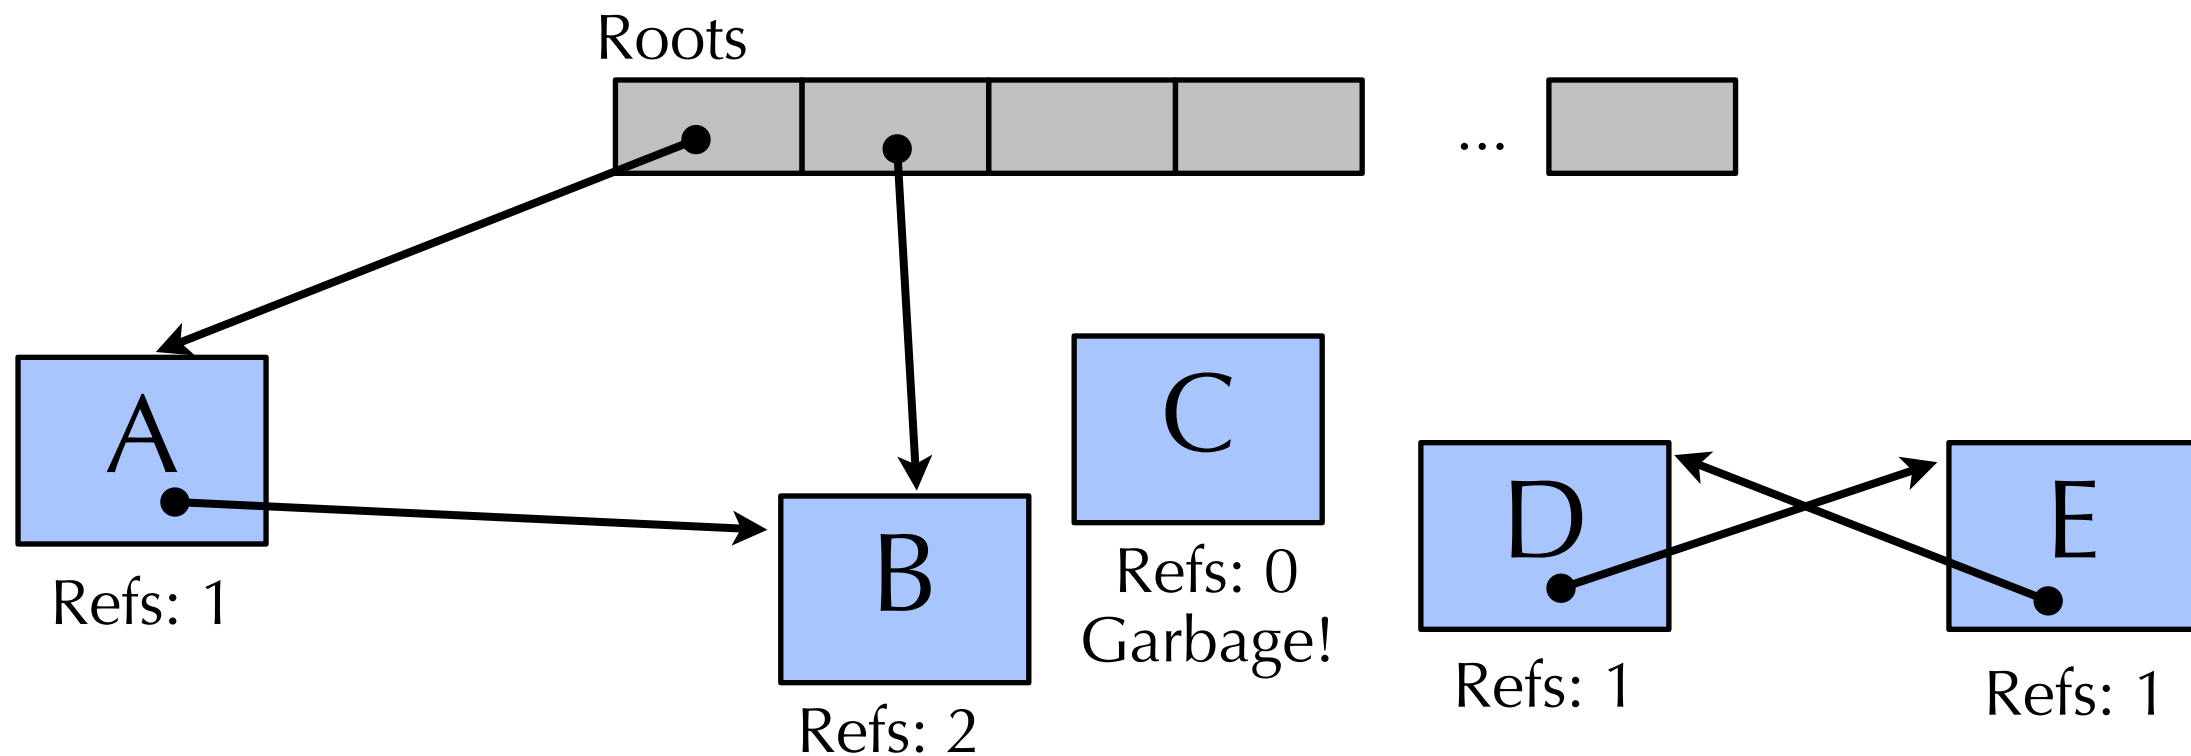
- Free block list management
  - Implicit free list
  - Explicit free list
  - Segregated lists
  - Tradeoffs
- Alignment issues
- Example allocator implementation: `dlmalloc`
- Common memory problems
- Garbage collection
  - Mark and sweep
  - Generational GC
  - Reference counting

# Reference Counting

- An object is garbage if it is not reachable from a root.
- If nothing points to an object (i.e., object is not **referenced**), then it is garbage.
- Problem: How do we decide if an object is no longer referenced?
- Solution: Keep track of number of active references to every object
  - As soon as reference hits 0, can immediately GC object
- Used in: Python, Microsoft COM, Apple Cocoa, etc.
  - Many C hash table implementations use ref. counting to keep track of active entries



# Reference counting example



## • Problems

- Changing pointers (including a variable going out of scope) requires updating reference counts. Can be expensive
- What about cycles? (e.g., D and E)

# Next lecture

- Memory and storage technologies