



**HARVARD**

School of Engineering  
and Applied Sciences

# Memory and Storage Technologies

*CS61, Lecture 12*

Prof. Stephen Chong

October 11, 2011


# Announcements

- HW 4: Malloc
  - If you haven't yet, please fill in the partner form
    - <http://tinyurl.com/CS61-Fa11-malloc-groups>
  - Design checkpoint on Thursday Oct 13
- Mid-course evaluation
  - <http://tinyurl.com/CS61-fa11-midcourse-eval>
  - Opportunity to tell us how the course is going, and how to improve it
  - Responses are anonymous, and do not affect CUE guide scores.


# Today

- Memory
- Disk drives
- I/O and memory buses
- Solid-state disks
- Storage technology trends

# What's inside a computer anyway?

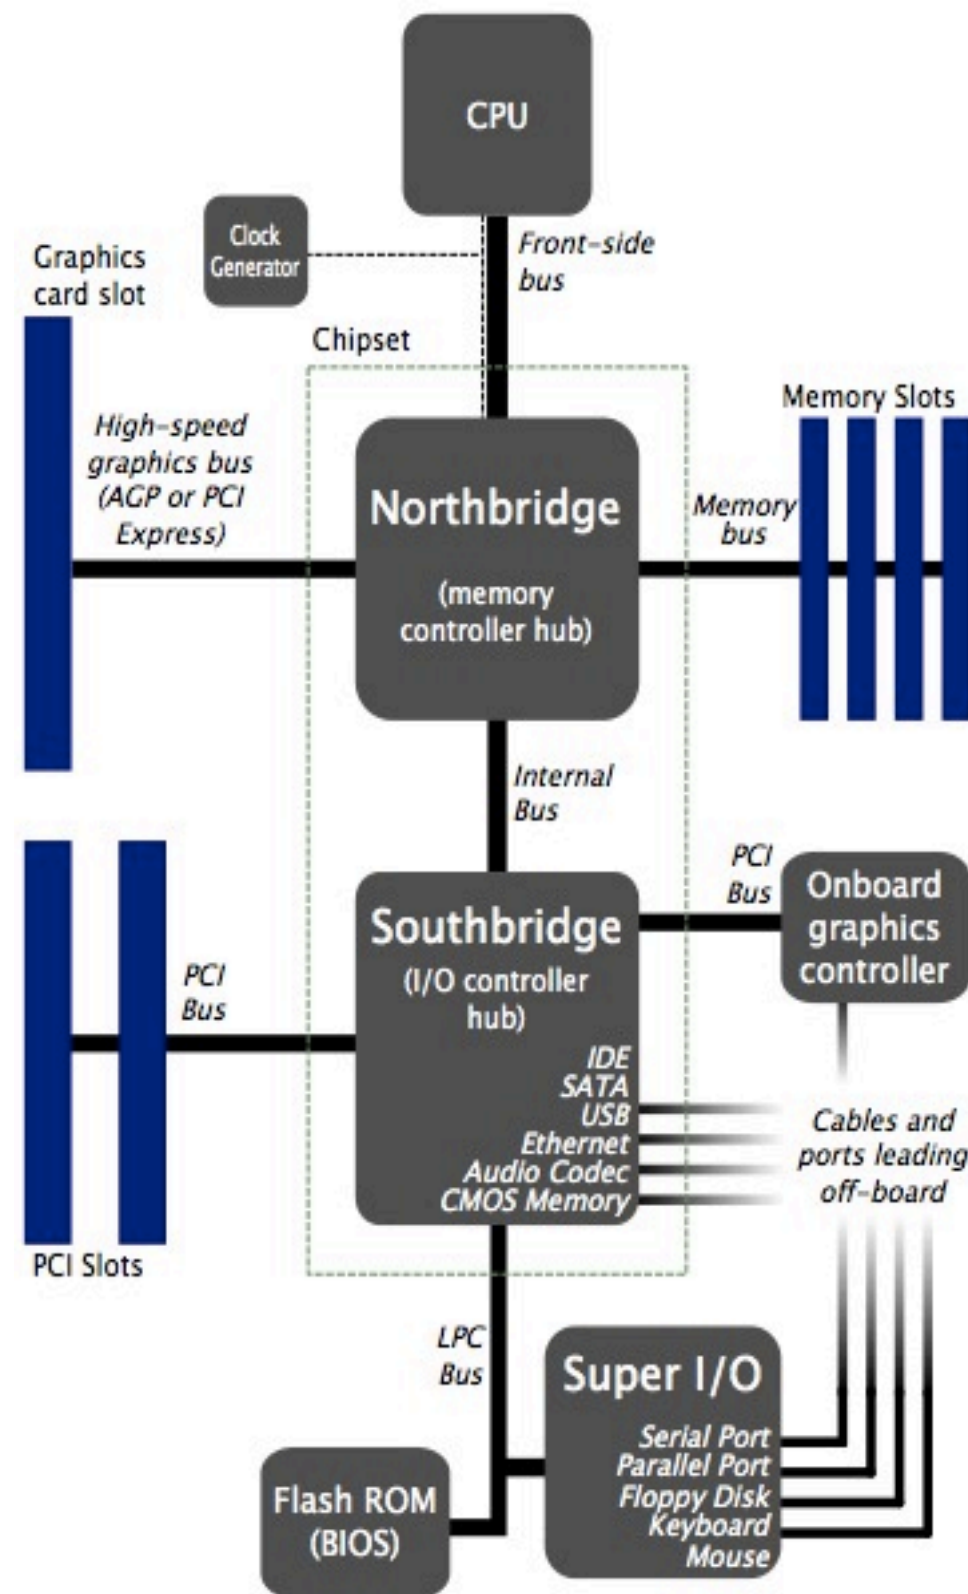








# What's inside a computer anyway?



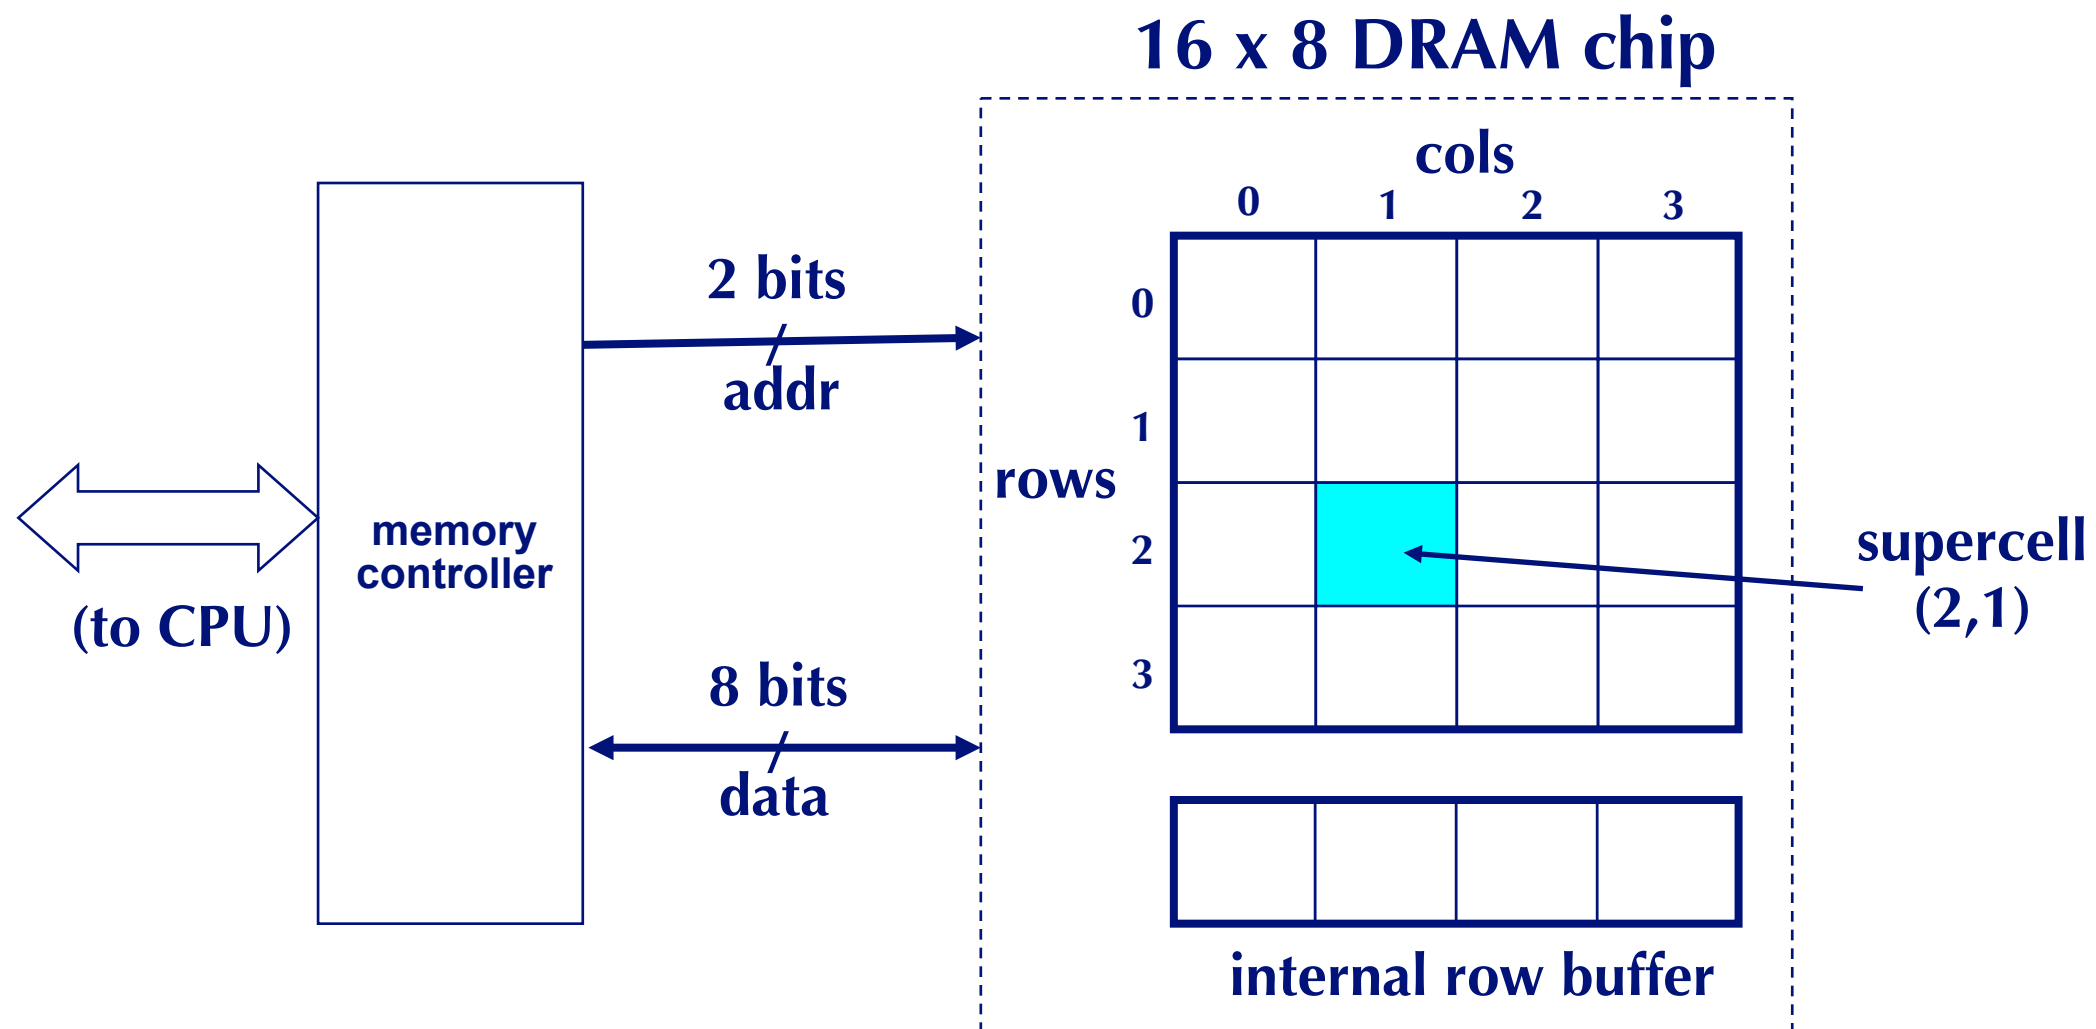
[http://en.wikipedia.org/wiki/File:Motherboard\\_diagram.svg](http://en.wikipedia.org/wiki/File:Motherboard_diagram.svg)

# Random-Access Memory (RAM)

- Key features
  - **RAM** is traditionally packaged as a chip.
  - Basic storage unit is normally a **cell** (one bit per cell).
  - Multiple RAM chips form a memory.
- Static RAM (**SRAM**)
  - Each cell stores a bit with a four or six-transistor circuit.
  - Retains value indefinitely, as long as it is kept powered.
  - *Relatively* insensitive to electrical noise (EMI), radiation, etc.
  - Faster and more expensive than DRAM – can be 100x more expensive!
- Dynamic RAM (**DRAM**)
  - Each cell stores bit with a capacitor. One transistor is used for access
  - Stored electric charge decays over time -- must be refreshed every 10-100 ms.
  - More sensitive to disturbances (EMI, radiation,...) than SRAM.
  - Slower, but cheaper than SRAM.

# Conventional DRAM Organization

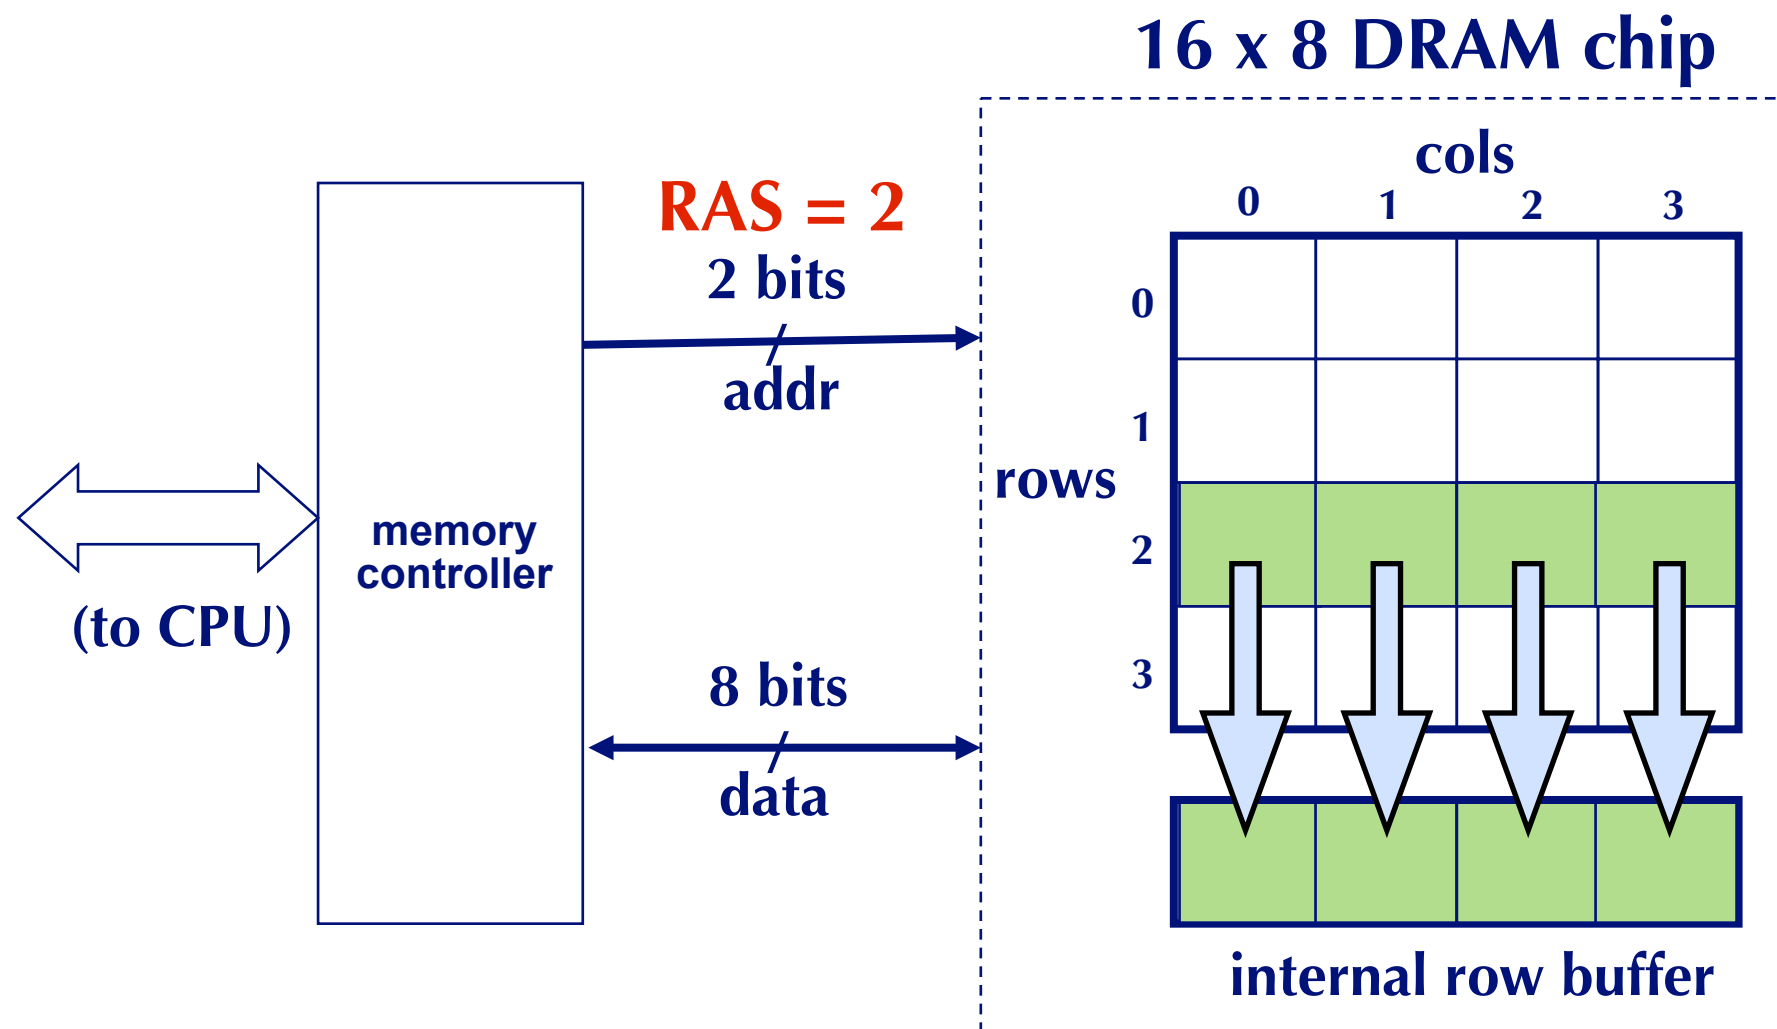
- $d \times w$  DRAM:
  - $d \times w$  total bits organized as  $d$  **supercells** of size  $w$  bits





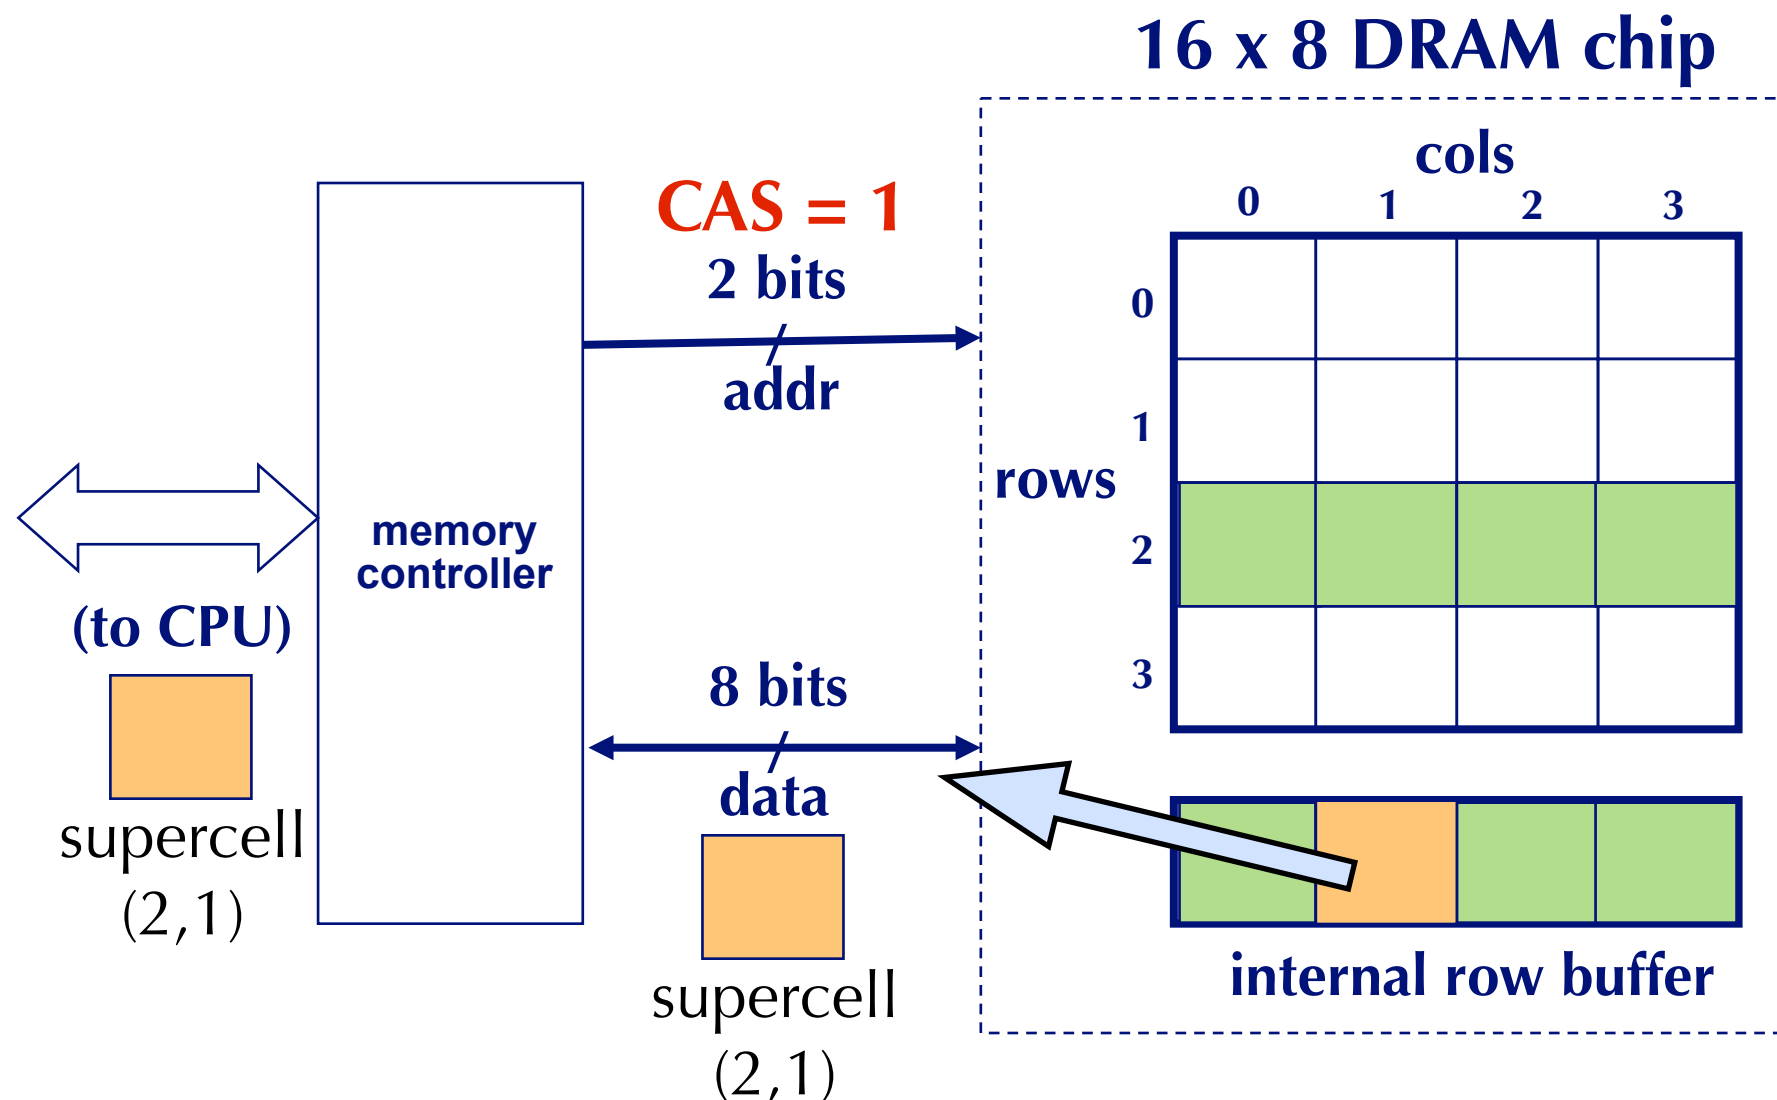
# Reading DRAM supercell (2,1)

- Step 1(a): Row access strobe (**RAS**) selects row 2
- Step 1(b): Row 2 copied from DRAM array to row buffer



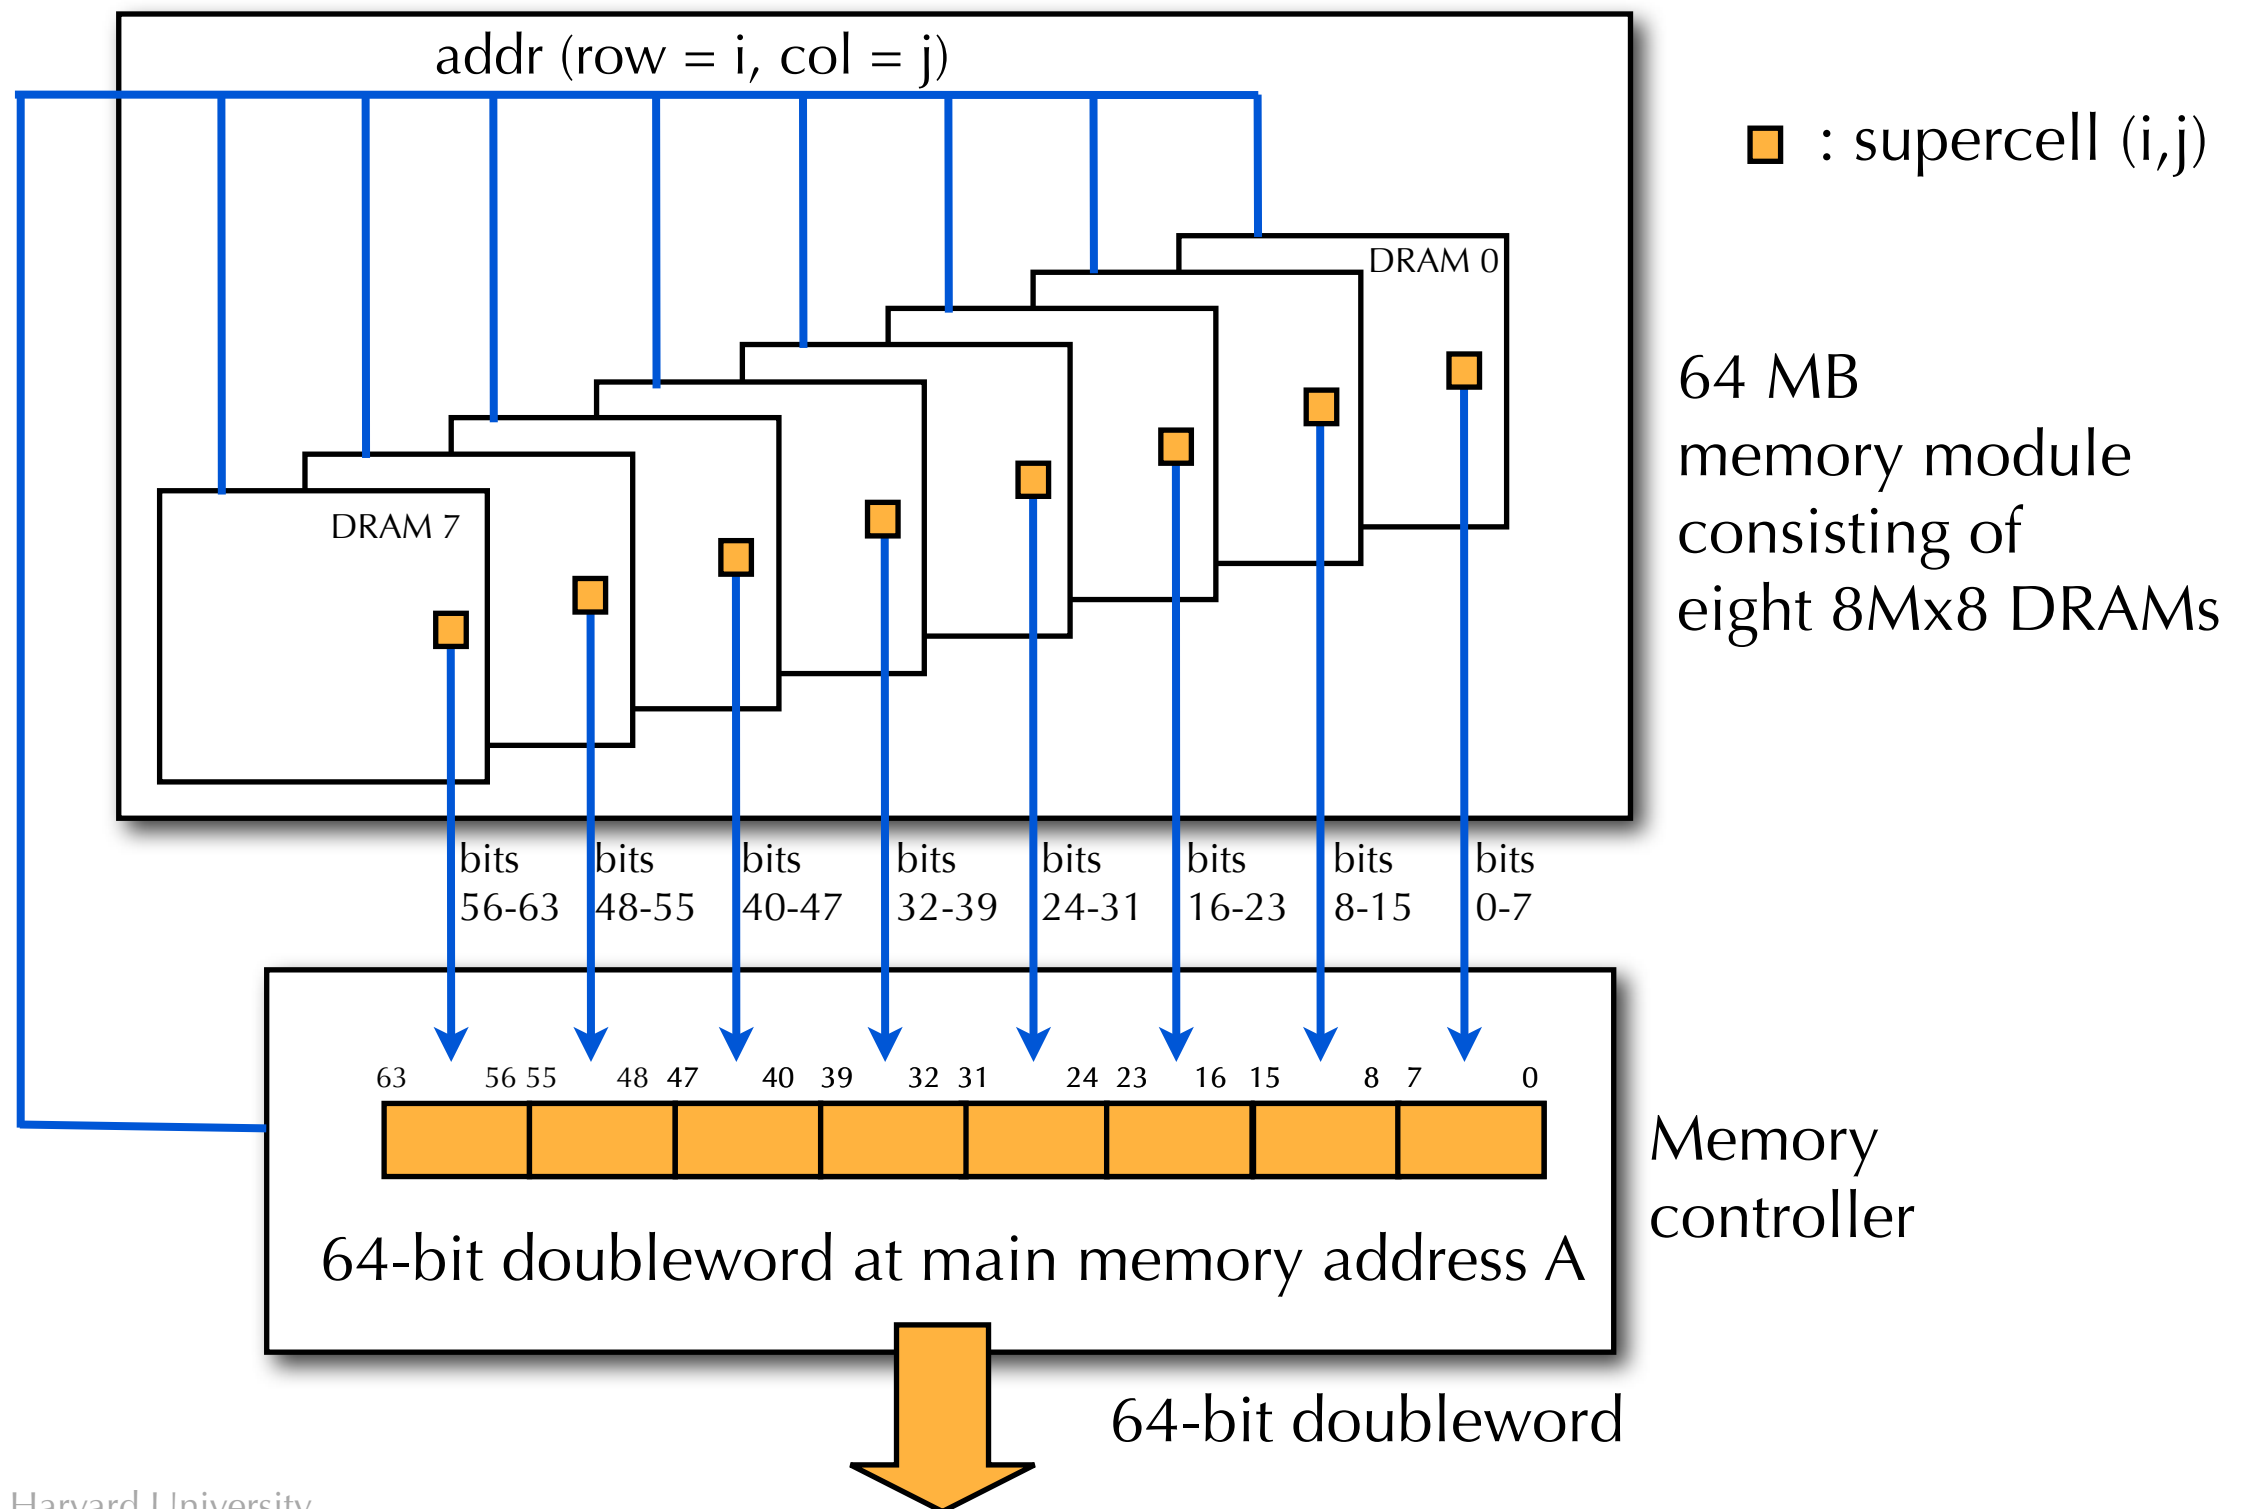
# Reading DRAM supercell (2,1)

- Step 2(a): Column access strobe (**CAS**) selects column 1
- Step 2(b): Supercell (2,1) copied from buffer to data lines, and eventually back to CPU






# Memory Modules



# DRAM packaging



Dual In-Line Package (DIP)



Single In-Line Pin Package (SIPP)



Single In-Line Memory Module (SIMM)  
30-pin



72-pin



Double In-Line Memory Module (DIMM)  
168-pin



Double Data Rate (DDR)  
DIMM (184-pin)

[http://en.wikipedia.org/wiki/File:RAM\\_n.jpg](http://en.wikipedia.org/wiki/File:RAM_n.jpg)



# Enhanced DRAMs

- DRAM Cores with better interface logic and faster I/O :
  - Synchronous DRAM (**SDRAM**)
    - Uses a conventional clock signal instead of asynchronous control
  - Double data-rate synchronous DRAM (**DDR SDRAM**)
    - Double edge clocking sends two bits per cycle per pin
  - RamBus™ DRAM (**RDRAM**)
    - Uses faster signaling over fewer wires with a transaction oriented interface protocol
- Obsolete Technologies :
  - Fast page mode DRAM (**FPM DRAM**)
    - Allowed re-use of row-addresses
  - Extended data out DRAM (**EDO DRAM**)
    - Enhanced FPM DRAM with more closely spaced CAS signals.
  - Video RAM (**VRAM**)
    - Dual ported FPM DRAM with a second, concurrent, serial interface
  - Extra functionality DRAMS (**CDRAM, GDRAM**)
    - Added SRAM (CDRAM) and support for graphics operations (GDRAM)

# Nonvolatile Memories

- DRAM and SRAM are **volatile** memories
  - Lose information if powered off.
- Nonvolatile memories retain value even if powered off
  - Read-only memory (**ROM**): programmed during production
  - Magnetic RAM (**MRAM**): stores bit magnetically (in development)
  - Ferro-electric RAM (**FeRAM**): uses a ferro-electric dielectric
  - Programmable ROM (**PROM**): can be programmed once
  - Erasable PROM (**EPROM**): can be bulk erased (UV, X-Ray)
  - Electrically erasable PROM (**EEPROM**): electronic erase capability
  - **Flash memory**: EEPROMs with partial (sector) erase capability





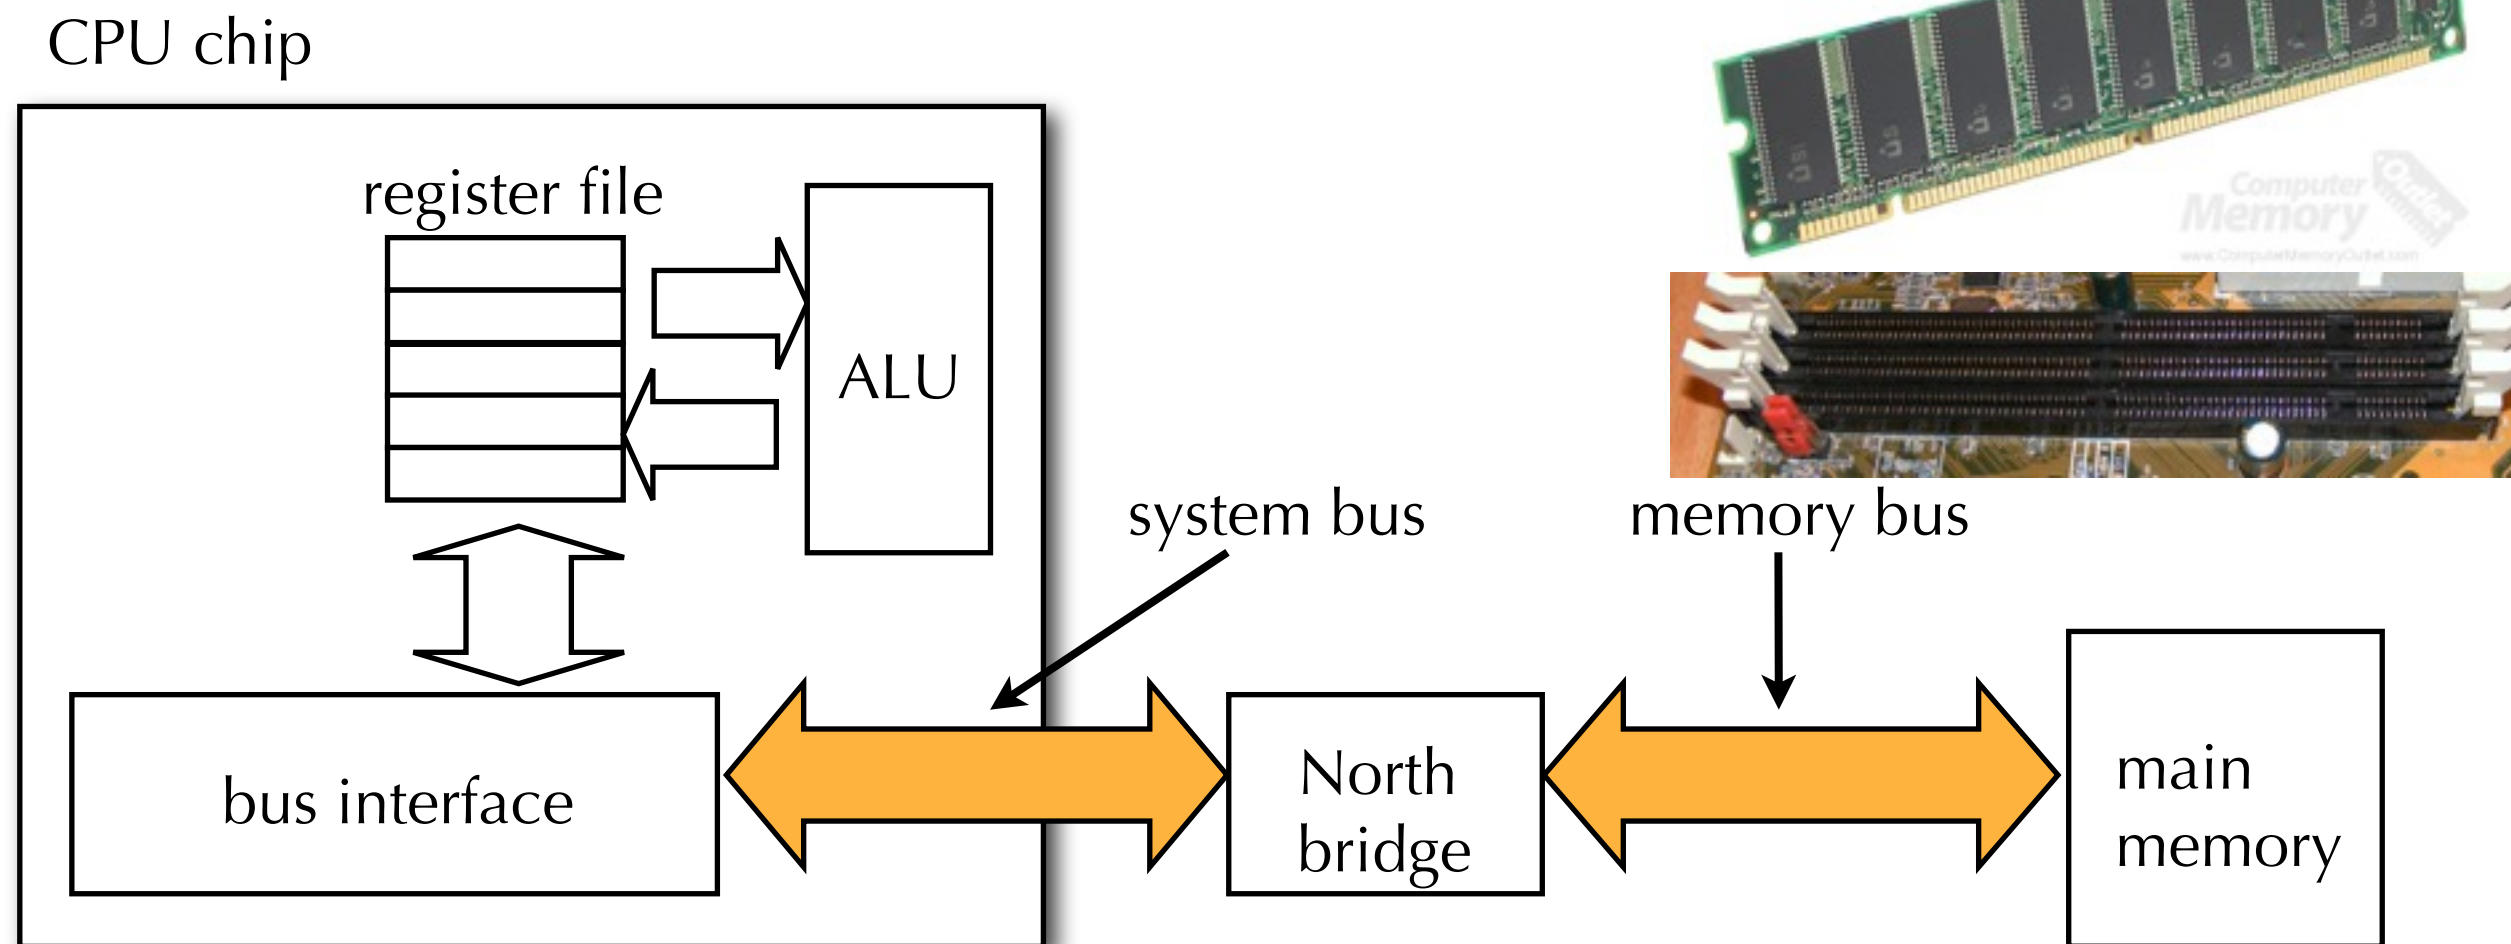
# Nonvolatile Memories

- Uses for Nonvolatile Memories
  - Firmware programs stored in ROM
    - BIOS, controllers for disks, network cards, graphics accelerators, security subsystems,...
  - Solid state disks (flash cards, memory sticks, etc.)
  - Smart cards, embedded systems, appliances
  - Disk caches

# Traditional Bus Structure

## Connecting CPU and Memory

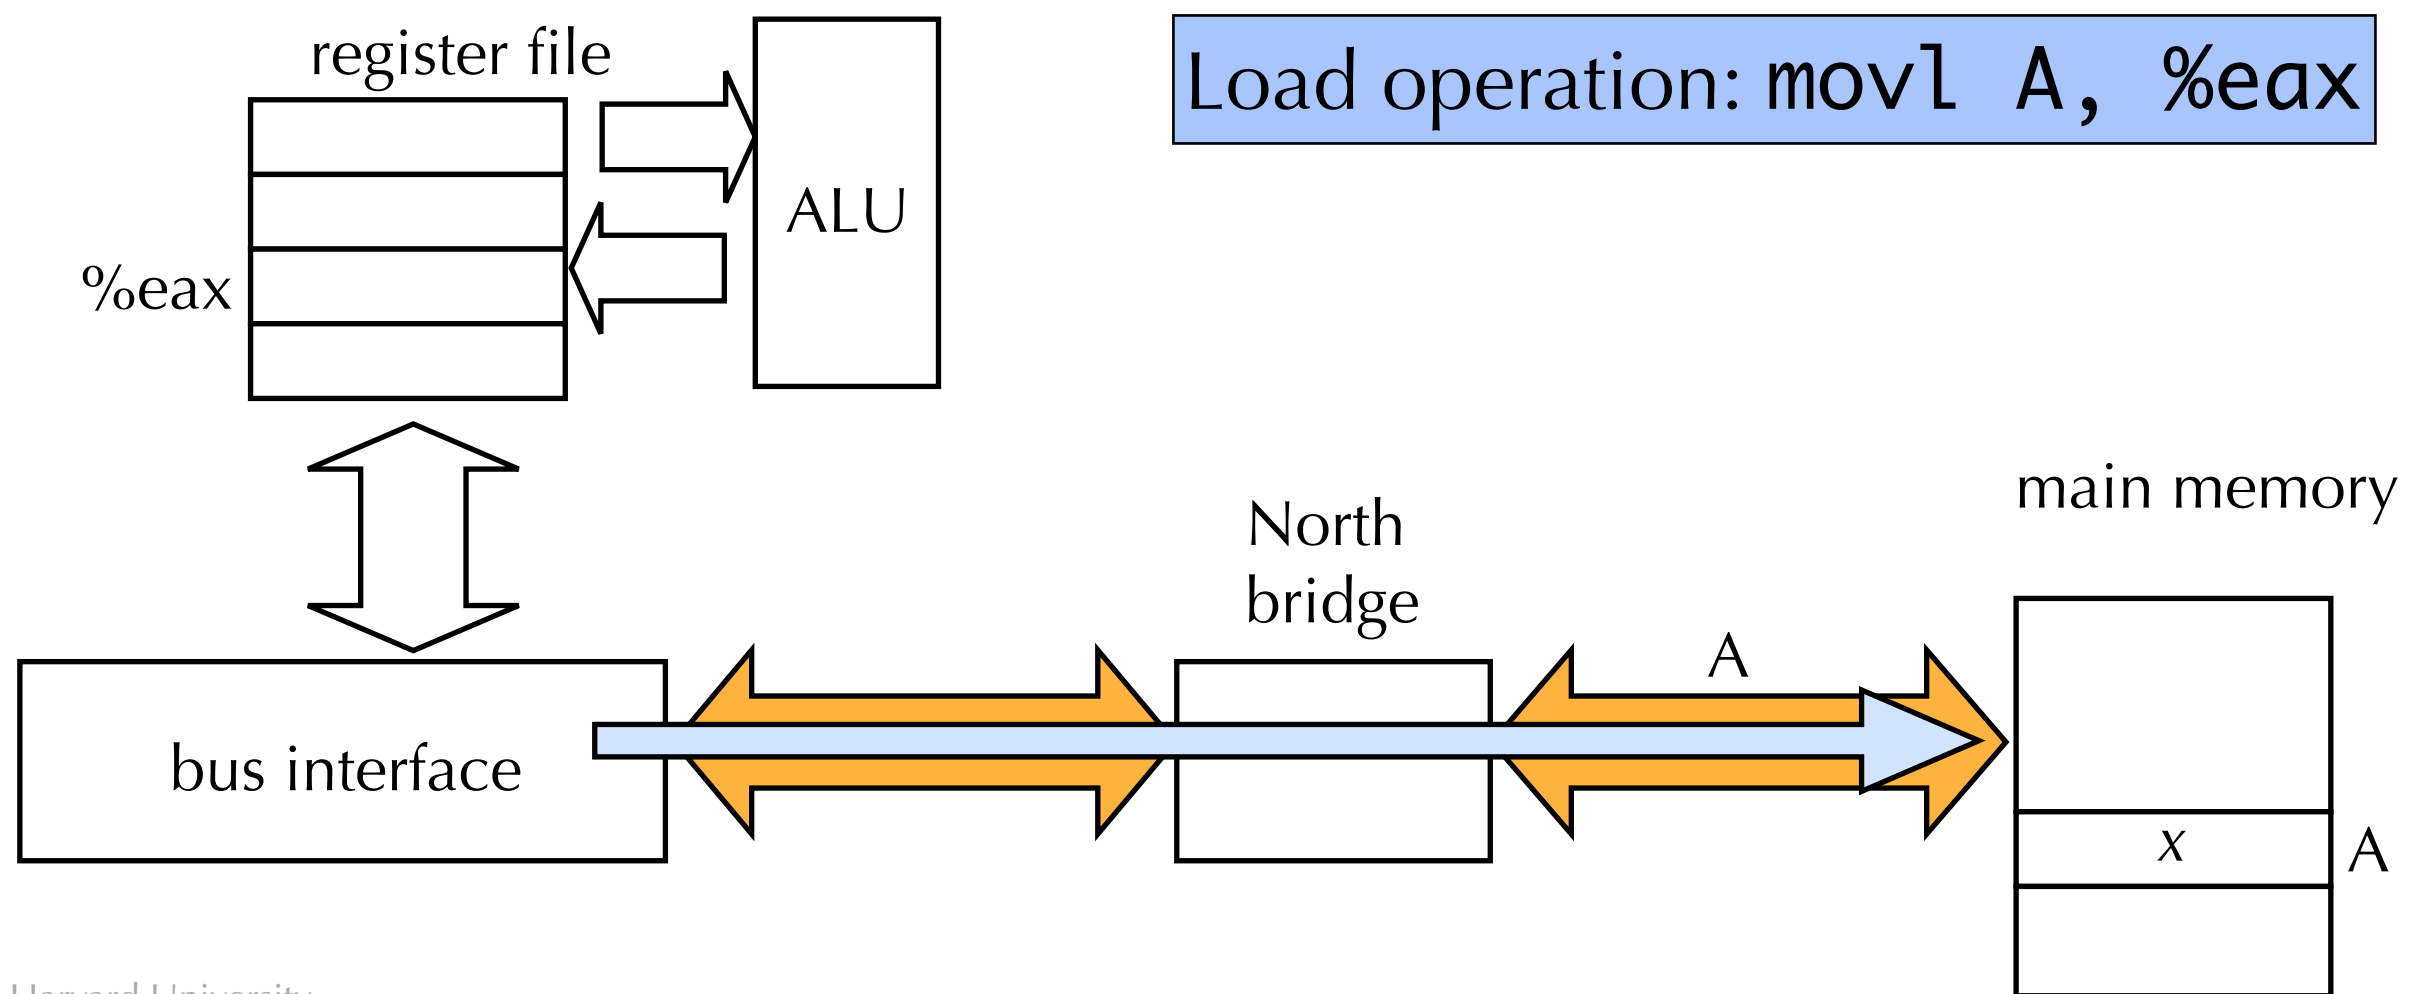
- A **bus** is a collection of parallel wires that carry address, data, and control signals.
- Buses are typically shared by multiple devices.





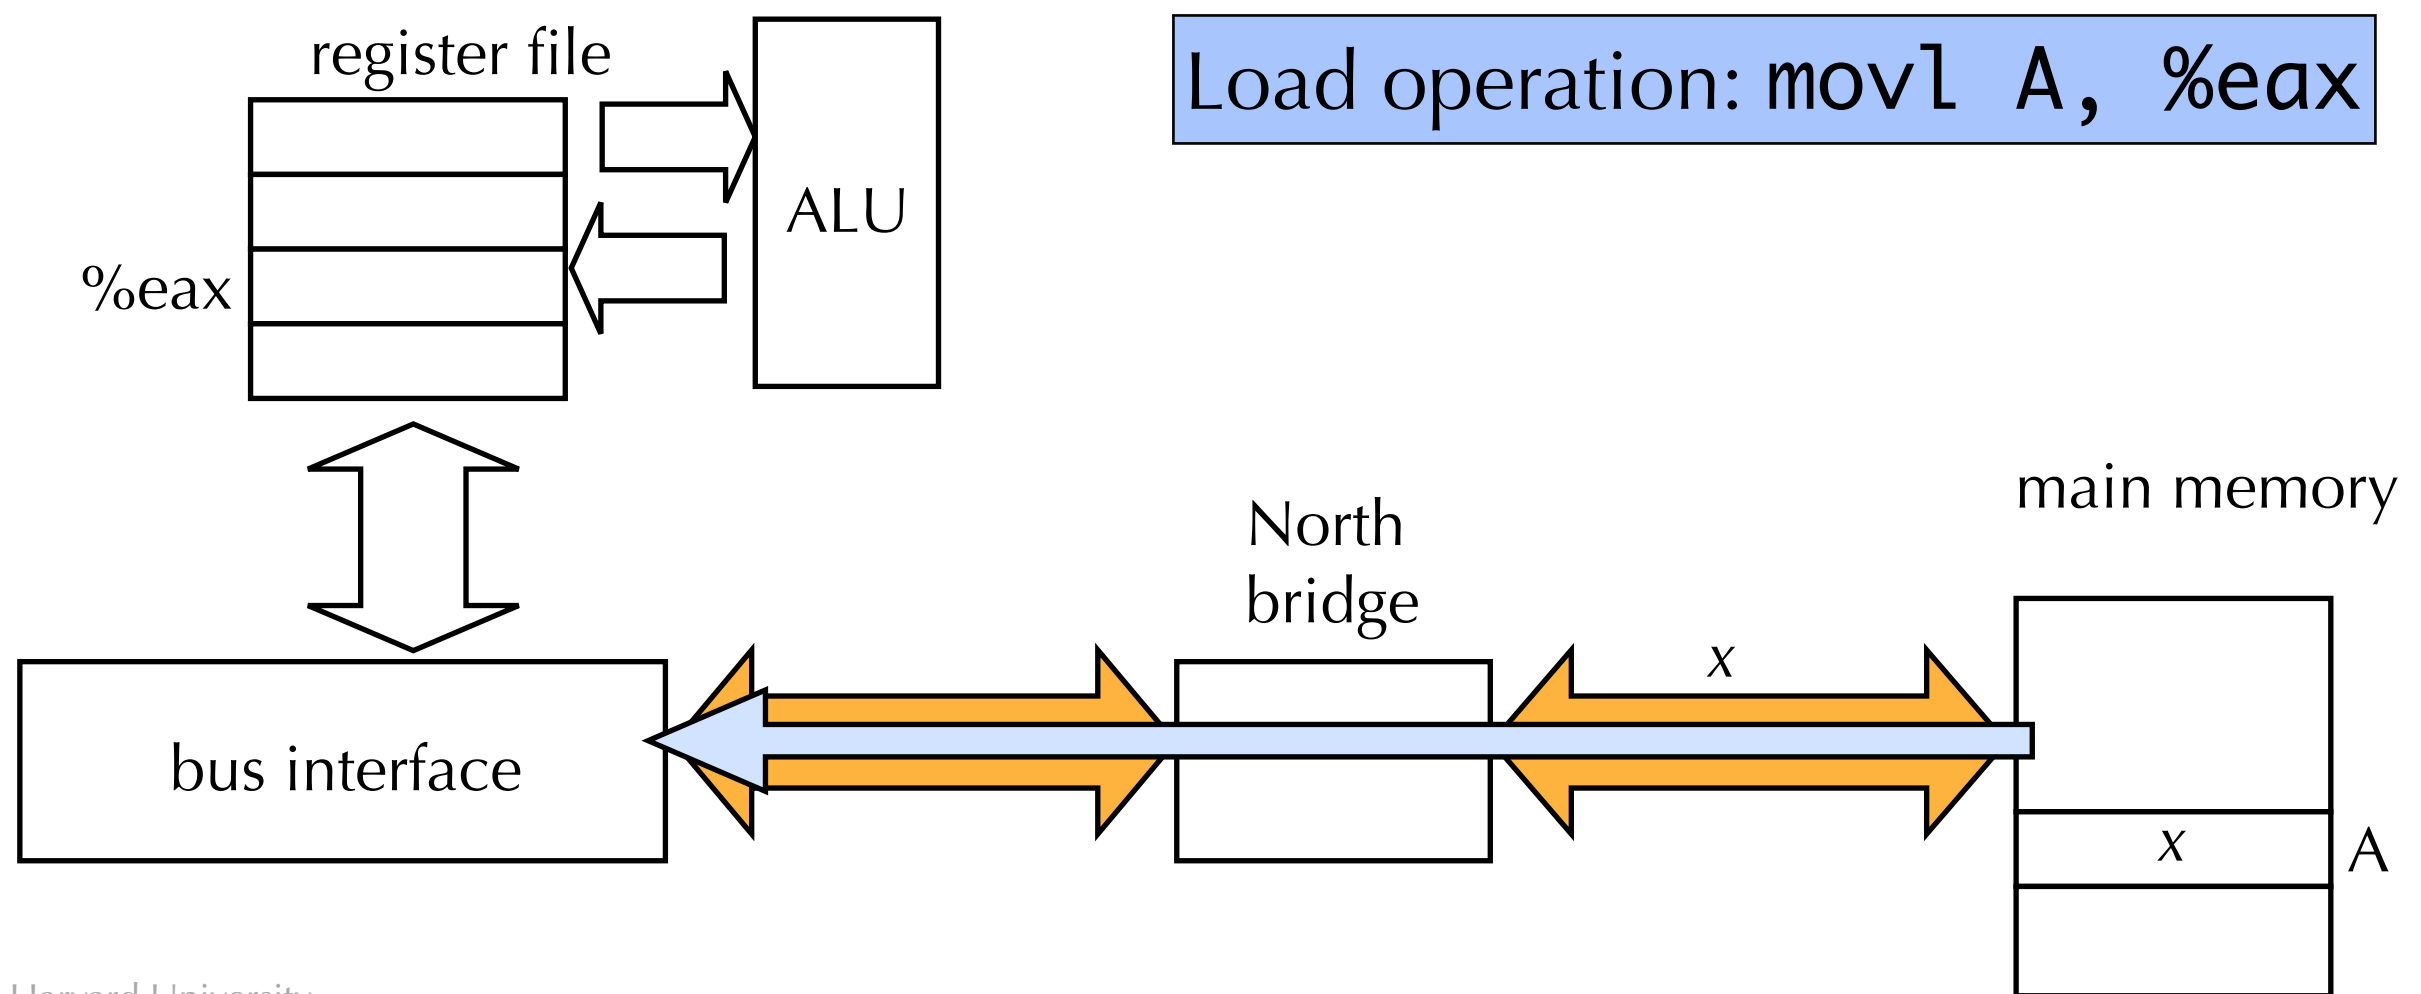
# Memory Read Transaction

- 1. CPU places address A on memory bus



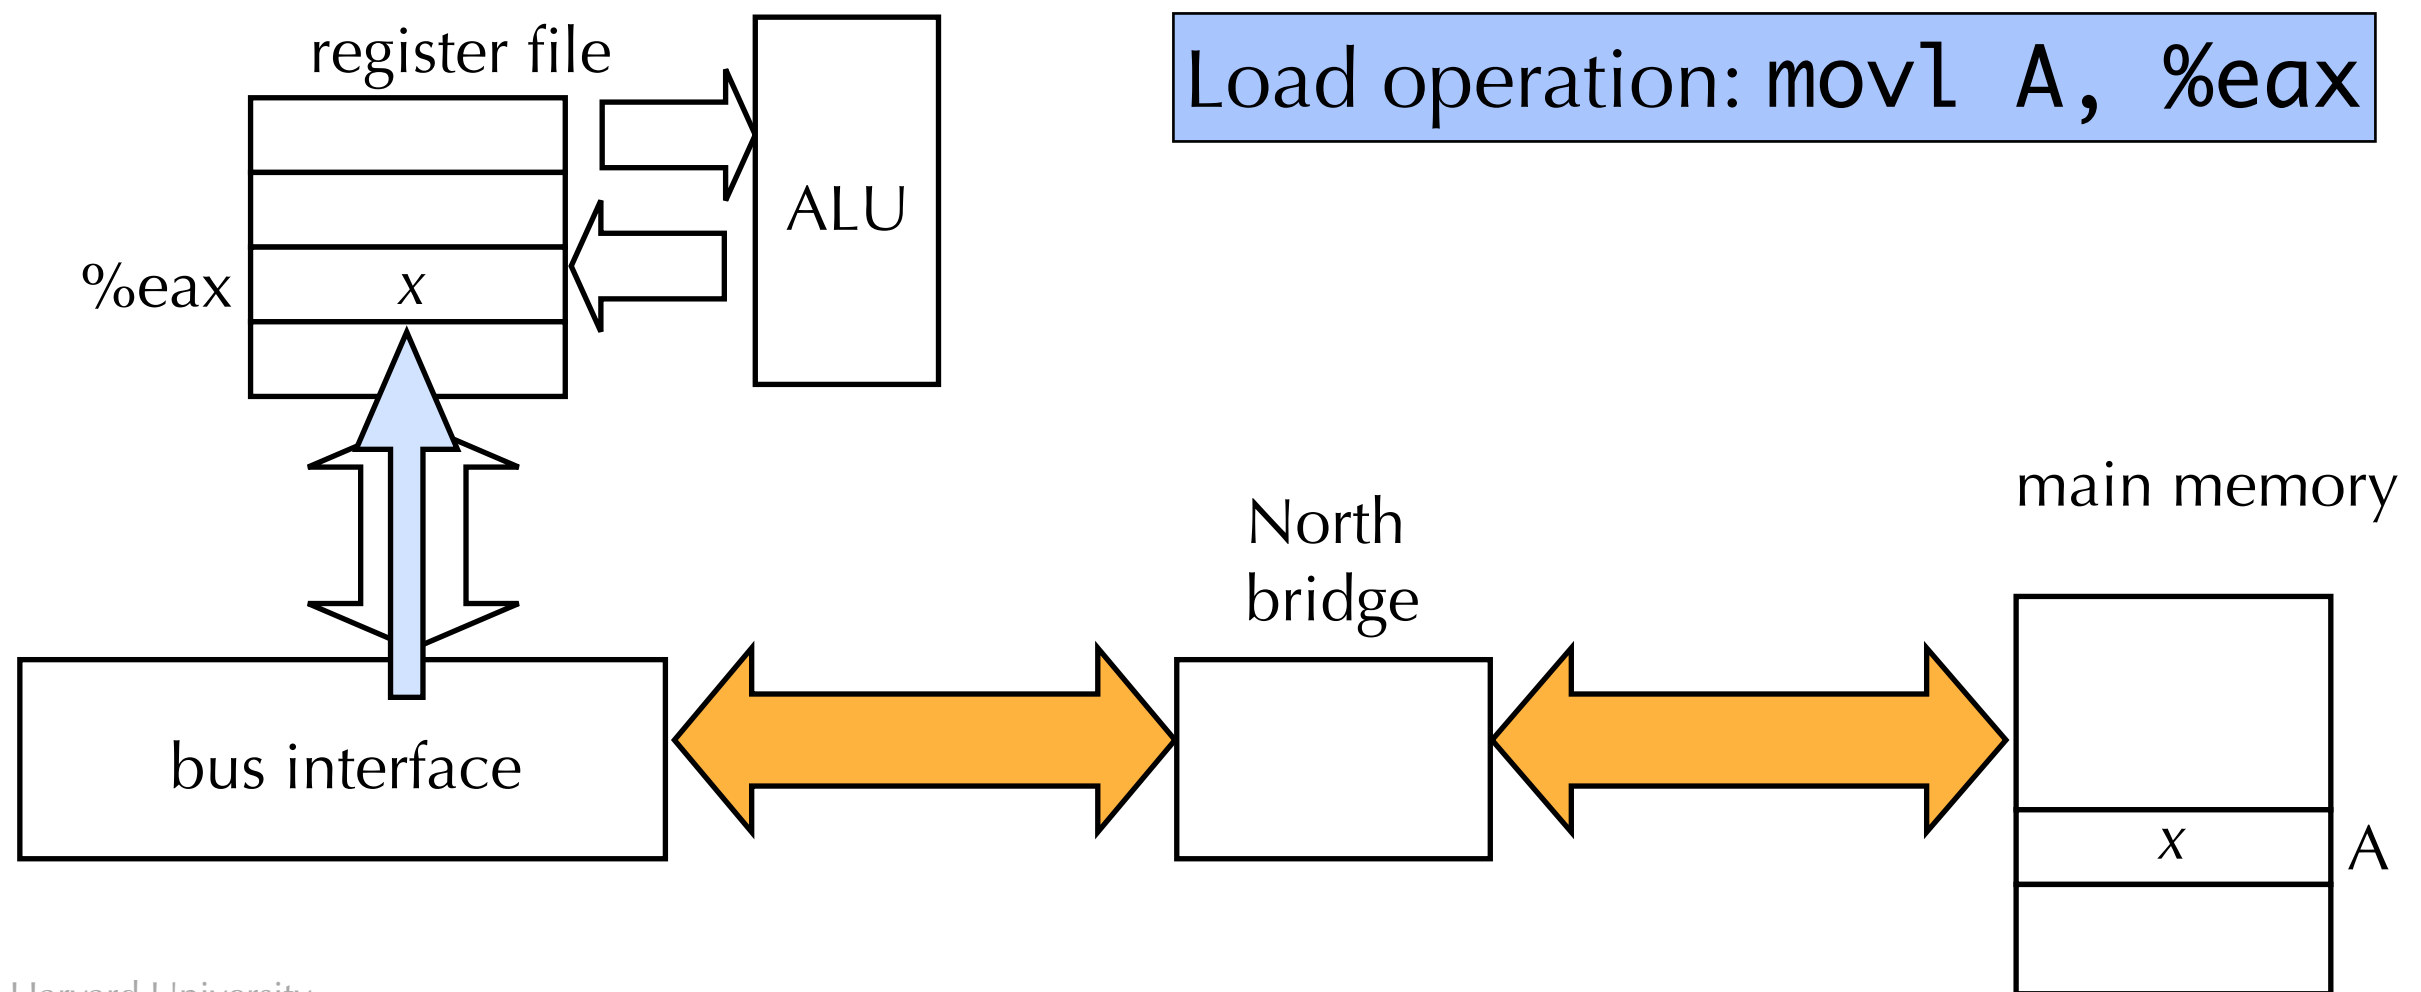
# Memory Read Transaction

- 2. Main memory reads A from memory bus, retrieves word x, and places it on bus



# Memory Read Transaction

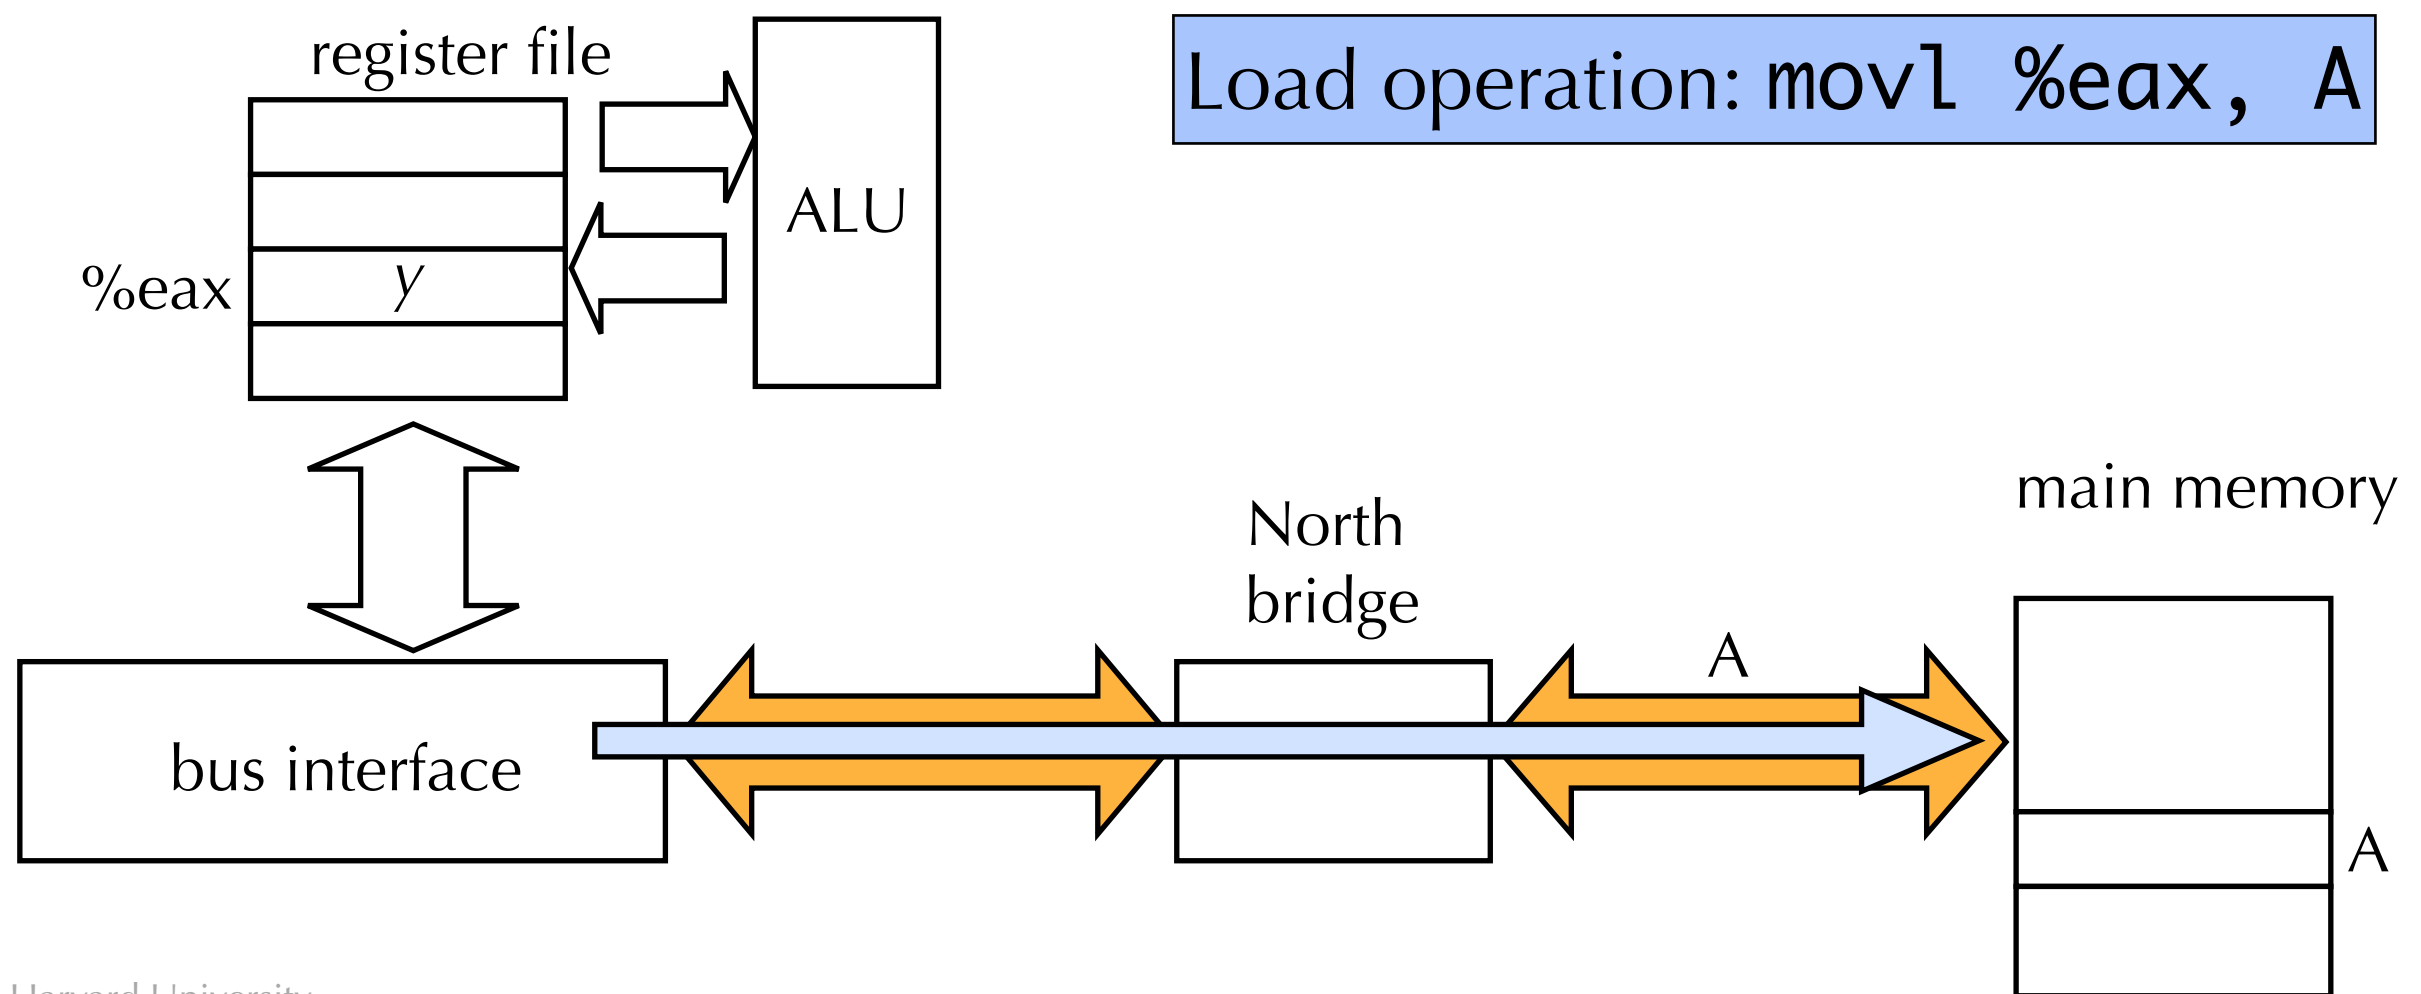
- 3. CPU reads word  $x$  from bus, copies it to register `%eax`





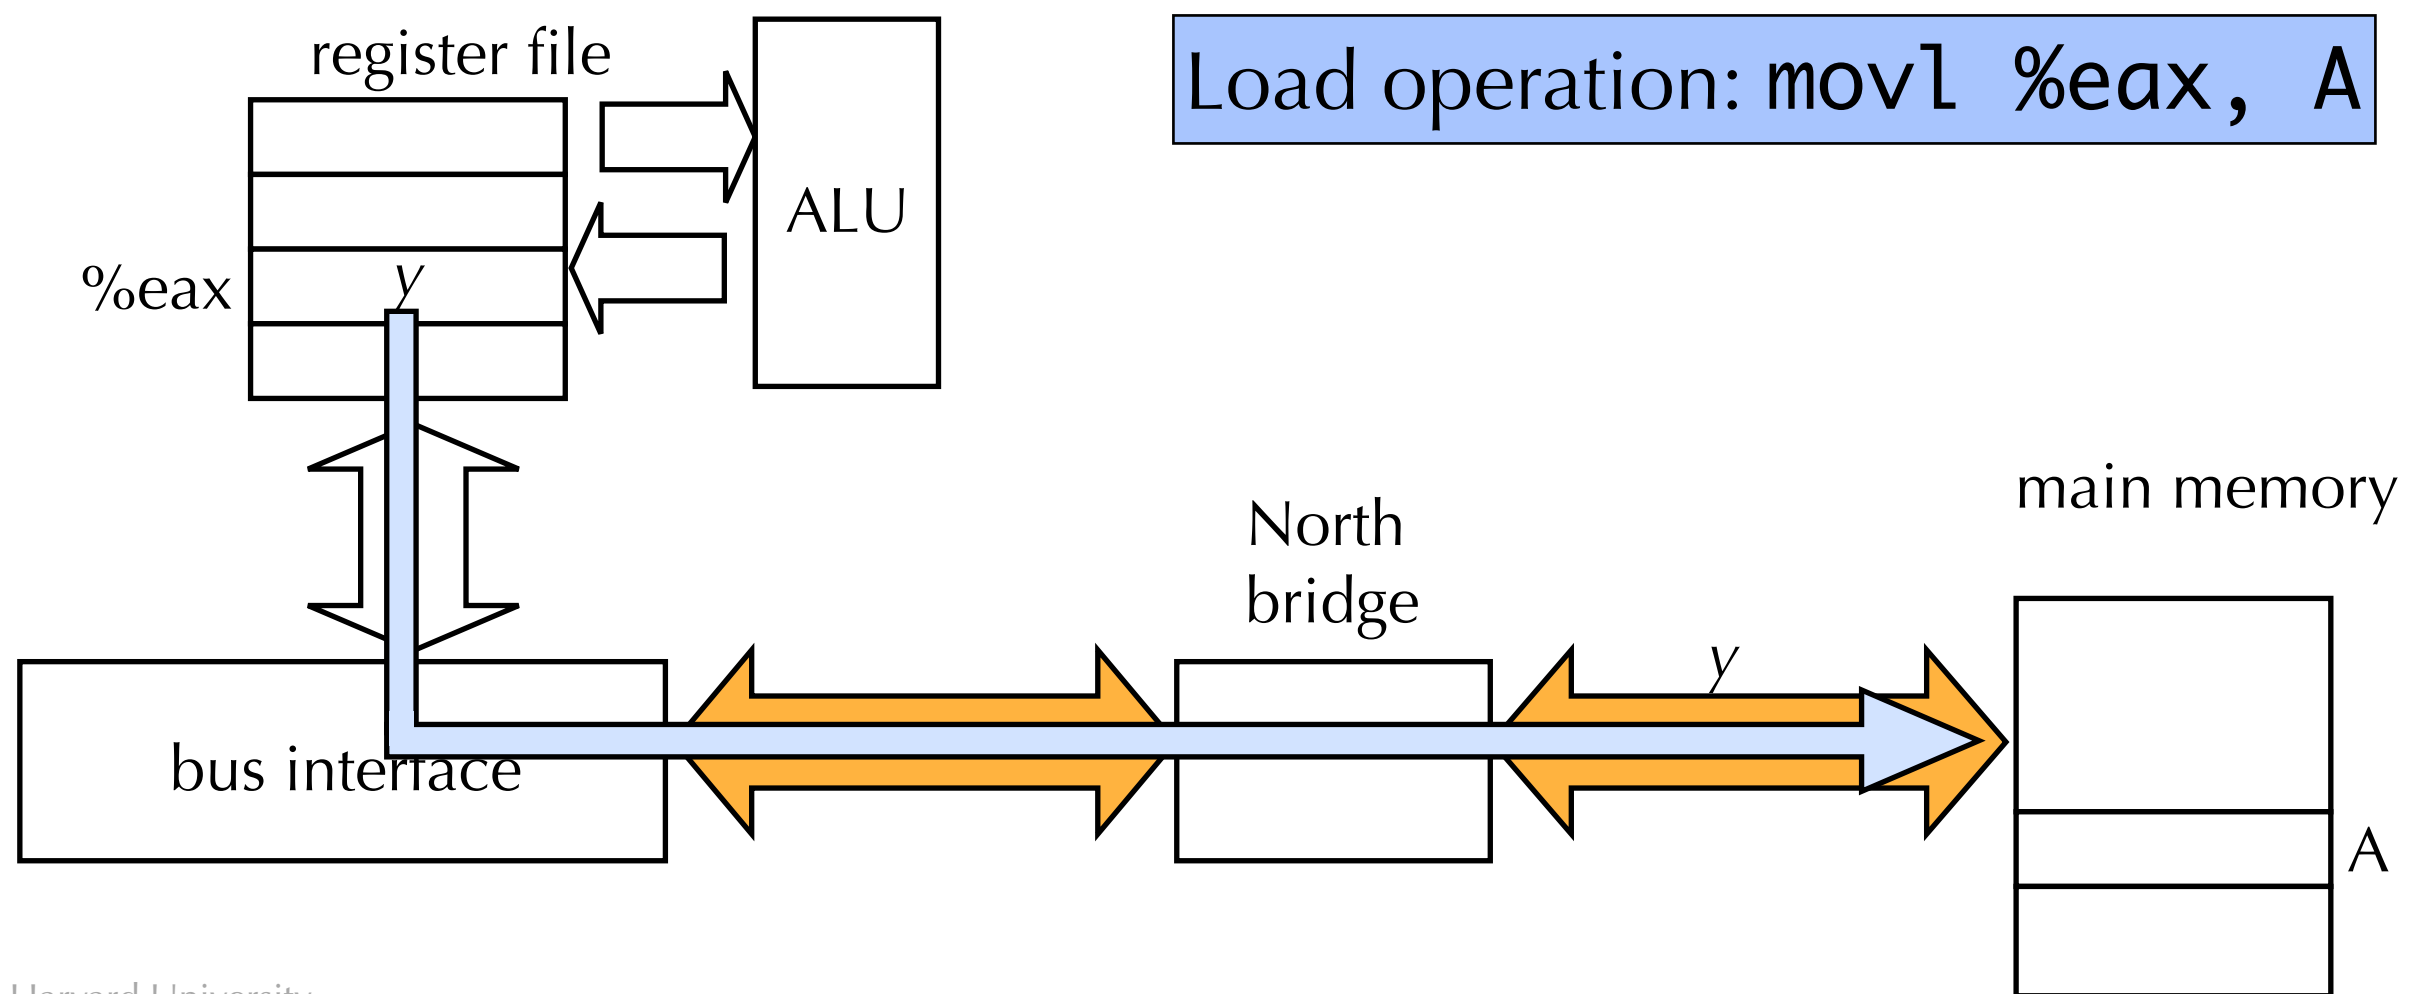
# Memory Write Transaction

- 1. CPU places address  $A$  on memory bus. Main memory reads it and waits for corresponding data word to arrive



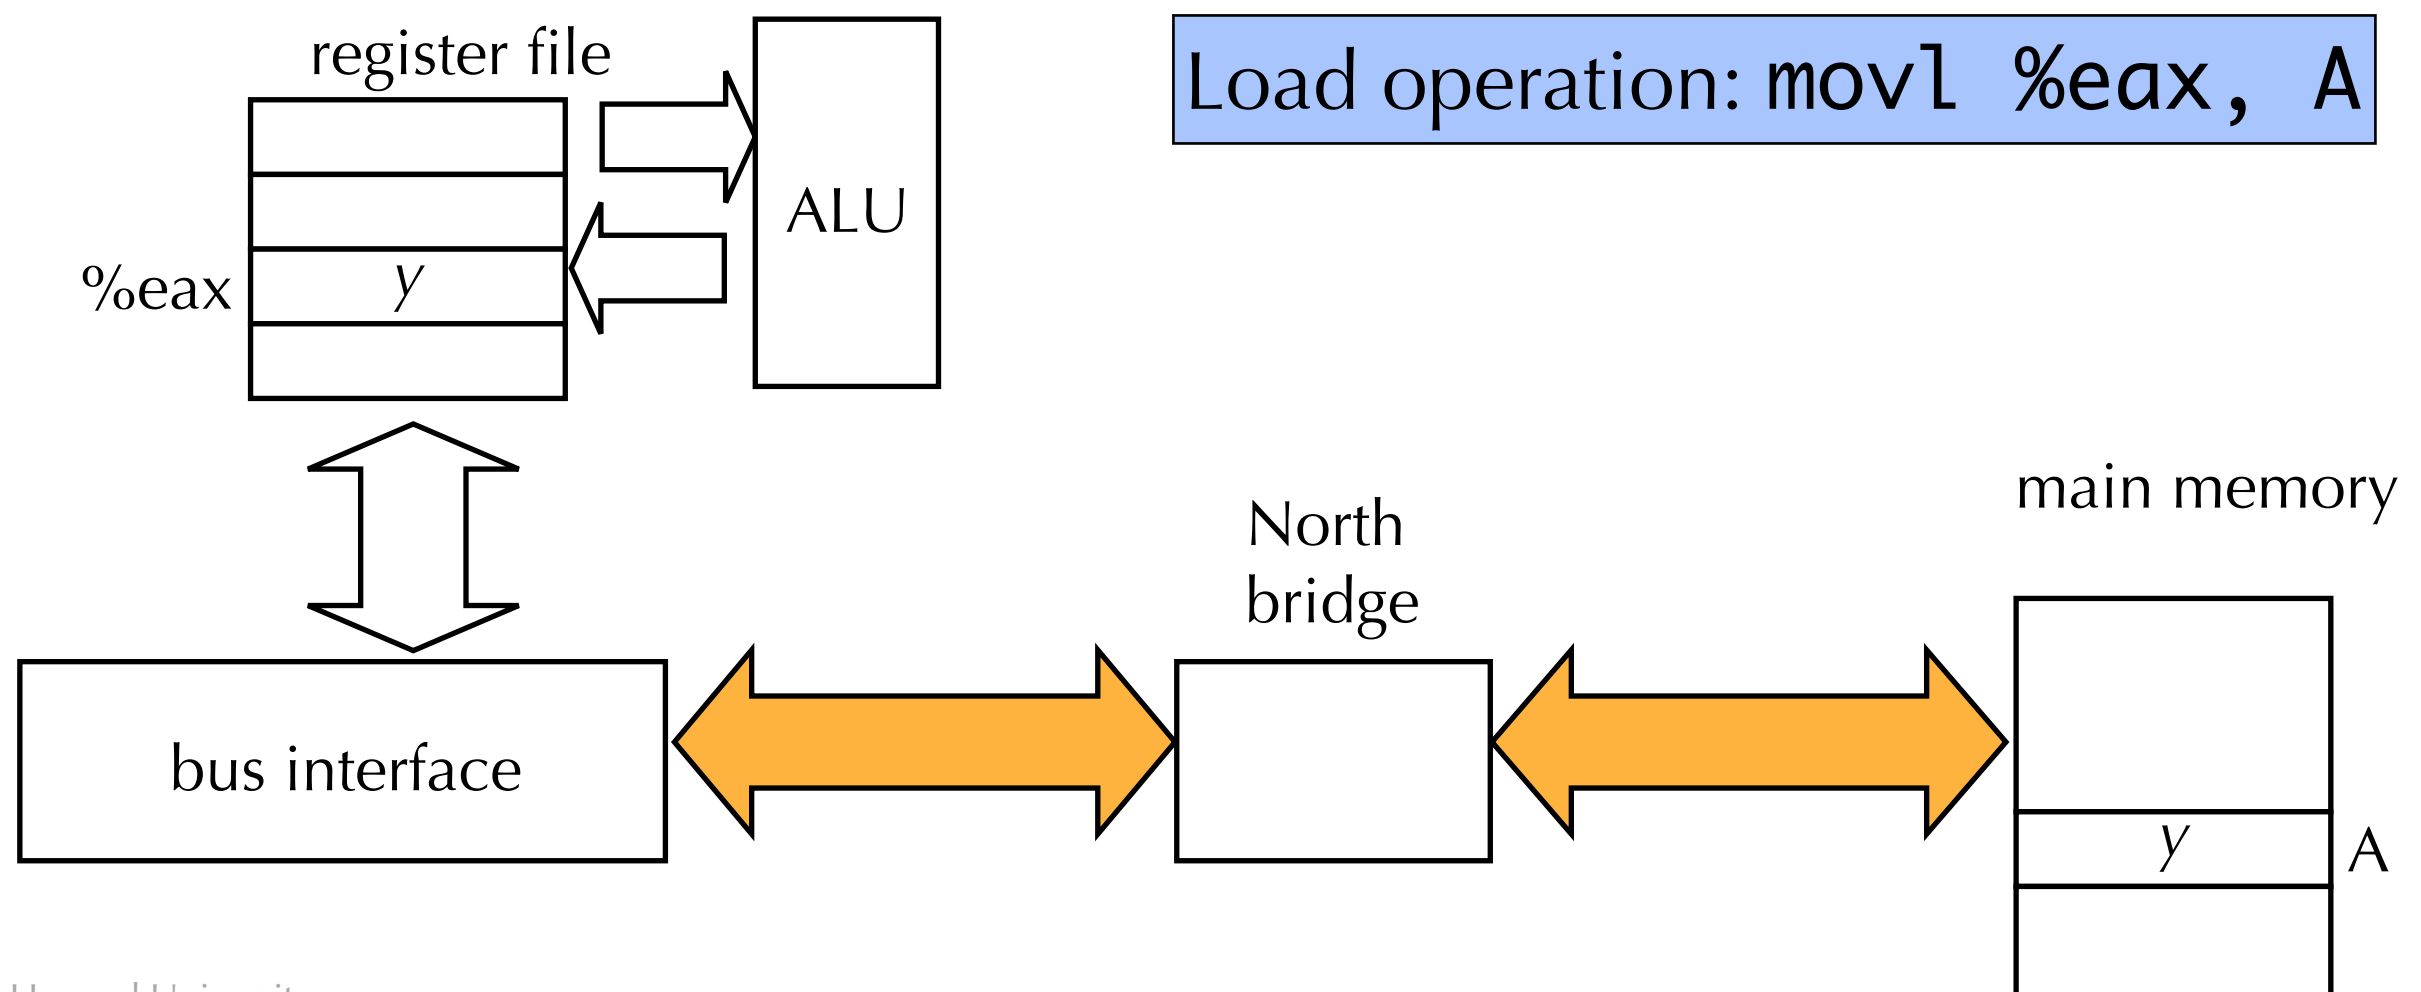
# Memory Write Transaction

- 2. CPU places data word  $y$  on memory bus.



# Memory Write Transaction

- 3. Main memory reads data word  $y$  from bus and stores it at address  $A$ .






# Errors happen!

- Electrical or magnetic interference can cause bits to flip from “1” to “0” (or vice versa)
  - Can be caused by cosmic rays passing through the memory chips
  - Error rates go up as densities increase: Up to one bit error per GB per HOUR
- One solution: Error-Correcting Code (ECC) RAM
  - Contains redundant information used to detect and correct bit errors: Hamming code.
  - Can detect and correct single bit errors; can detect (but not correct) 2-bit errors
  - Costs more: need more memory chips to store ECC information
  - Slower: Requires time to check and correct bit errors

# Forced errors


- Around 80° -100°C, memory starts to have more frequent failures



*Using Memory Errors to Attack a Virtual Machine*  
Govindavajhala and Appel  
2003 IEEE Symposium on Security and Privacy

# What happens to DRAM when you turn off the power?

- The contents are erased... right? right?
- Not so fast.



30 sec

60 sec

5 min

*Lest We Remember: Cold Boot Attacks on Encryption Keys*

Halderman et al.

USENIX Security Symposium 2008



# What happens to DRAM when you turn off the power?

- The contents are erased... right? right?
- Not so fast.



*Lest We Remember: Cold Boot Attacks on Encryption Keys*

Halderman et al.

<http://www.youtube.com/watch?v=JDaicPlgn9U>

USENIX Security Symposium 2008

# Today

- Memory
- Disk drives
- I/O and memory buses
- Solid-state disks
- Storage technology trends


# Disk storage

- Storage devices that hold enormous amounts of data
  - Hundreds to thousands of **gigabytes** ( $=10^9$  bytes)
  - RAM: hundreds to thousands of megabytes




# A Disk Primer

- (Rotating) Disks consist of one or more platters divided into tracks
  - Each **platter** may have one or two **heads** that perform read/write operations
  - Each **track** consists of multiple **sectors**
  - The set of sectors across all platters is a **cylinder**




Toshiba 4.0 GB 0.85" hard drive



# Disk Operation (Single-Platter View)


The disk surface spins at a fixed rotational rate



The read/write head is attached to the end of the **actuator arm** and flies about 10 nanometers above the disk surface on a thin cushion of air.

By moving radially, the arm can position the read/write head over any track.

# Disk Operation (Multi-Platter View)



# Hard disk capacity

- **Capacity** of disk is maximum number of bits that can be stored on disk
- Determined by:
  - Recording density (bits/inch)
    - Number of bits than can be squeezed into a 1-inch segment of track
  - Track density (tracks/inch)
    - Number of tracks per 1-inch segment of radius
  - Areal density (bits/inch<sup>2</sup>)
    - Recording density  $\times$  track density



# Hard-drive teardown


<http://www.engineerguy.com/videos/video-harddrive.htm>

# Hard Disk Evolution

- IBM 305 RAMAC (1956)
  - First commercially produced hard drive
  - 5 MB capacity, 50 platters each 24" in diameter!



# Hard Disk Evolution





# Disk access time

- Command overhead:
  - Time to issue I/O, get the HDD to start responding, select appropriate head
- Seek time:
  - Time to move disk arm to the appropriate track
  - Depends on how fast you can physically move the disk arm
    - These times are not improving rapidly!
- Rotational latency:
  - Time for the appropriate sector to move under the disk arm
  - Depends on the rotation speed of the disk (e.g., 7200 RPM)
- Transfer time
  - Time to transfer a sector to/from the disk controller
  - Depends on density of bits on disk and RPM of disk rotation
  - Faster for tracks near the outer edge of the disk – why?
    - Modern drives have more sectors on the outer tracks!



# Disk access time

- Access time dominated by seek time and rotational latency.
  - First bit in a sector is the most expensive, the rest are free.
- Disk is slooooooow...
  - SRAM access time is about 4 ns, DRAM about 60 ns
  - Disk is about 40,000 times slower than SRAM,
  - 2,500 times slower than DRAM.
  - Requires careful scheduling of I/O requests

# Example disk characteristics

- Seagate Constellation ES SAS 6Gb/s 1-TB Hard Drive

- Form factor: 3.5"
- Capacity: 1.0 TB
- Rotation rate: 7,200 RPM
  - Avg. rotational latency: 4.16 ms
- Platters: 2 (4 surfaces)
- Cylinders: 248,600
- Cache: 32MB
- Average read time: 8.3ms
- Average write time: 9.3ms
- Transfer rate:     Buffer to/from disk: 95-212 MB/s  
                             Host to/from drive (sustained): 60-150 MB/s



# Disks are messy

- Disks provide a low level interface for reading and writing sectors
  - Generally read/write an entire sector at a time
    - So, what do you do if you need to write a single byte to a file?
  - No notion of “files” or “directories”, just raw sectors
  - Disk may have numerous bad sectors that need to be avoided
- Difficult to use the low level interface

# Logical Disk Blocks

- Modern disks present a simpler abstract view of the complex sector geometry:
  - Disk presented as sequence of B **logical blocks**, each of fixed size.
- Mapping between logical blocks and actual (physical) sectors
  - Maintained by hardware/firmware device called **disk controller**.
  - Converts requests for logical blocks into (surface,track,sector) triples.
- Benefits of abstraction:
  - Controller can transparently avoid bad sectors
    - Just change mapping
  - Controller can set aside **spare cylinders**
    - Accounts for the difference in “formatted capacity” and “maximum capacity”

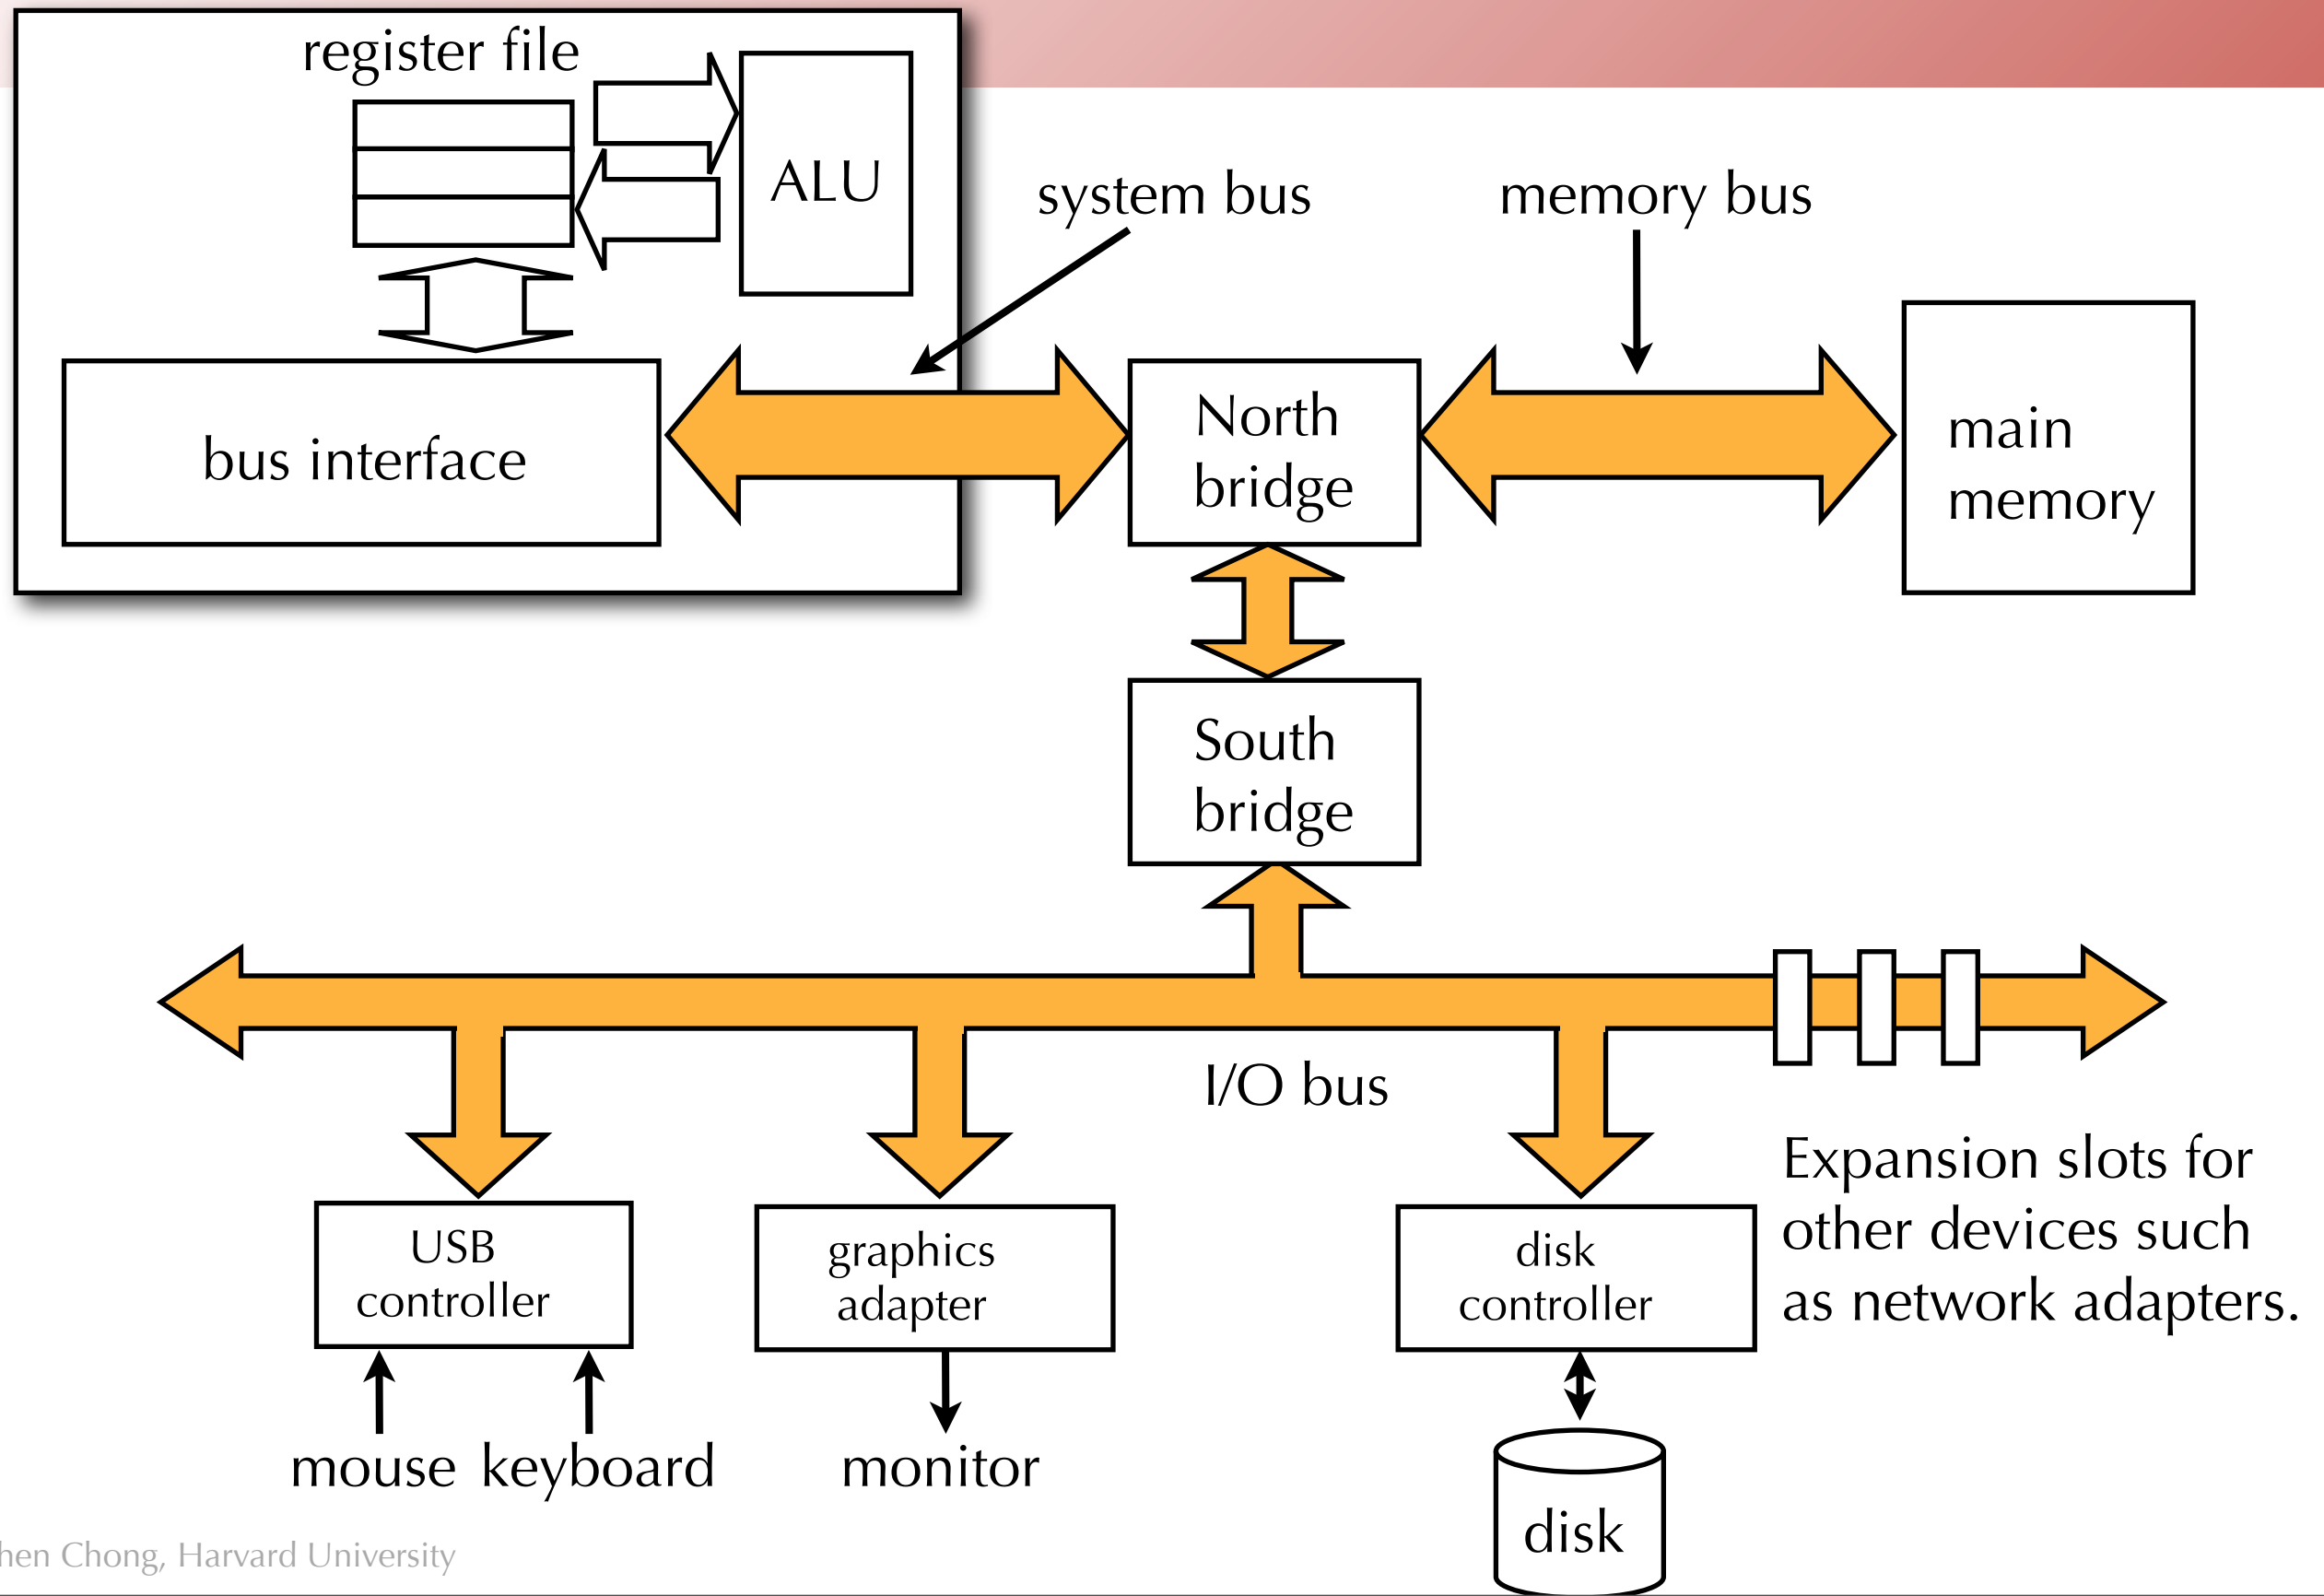


# Today

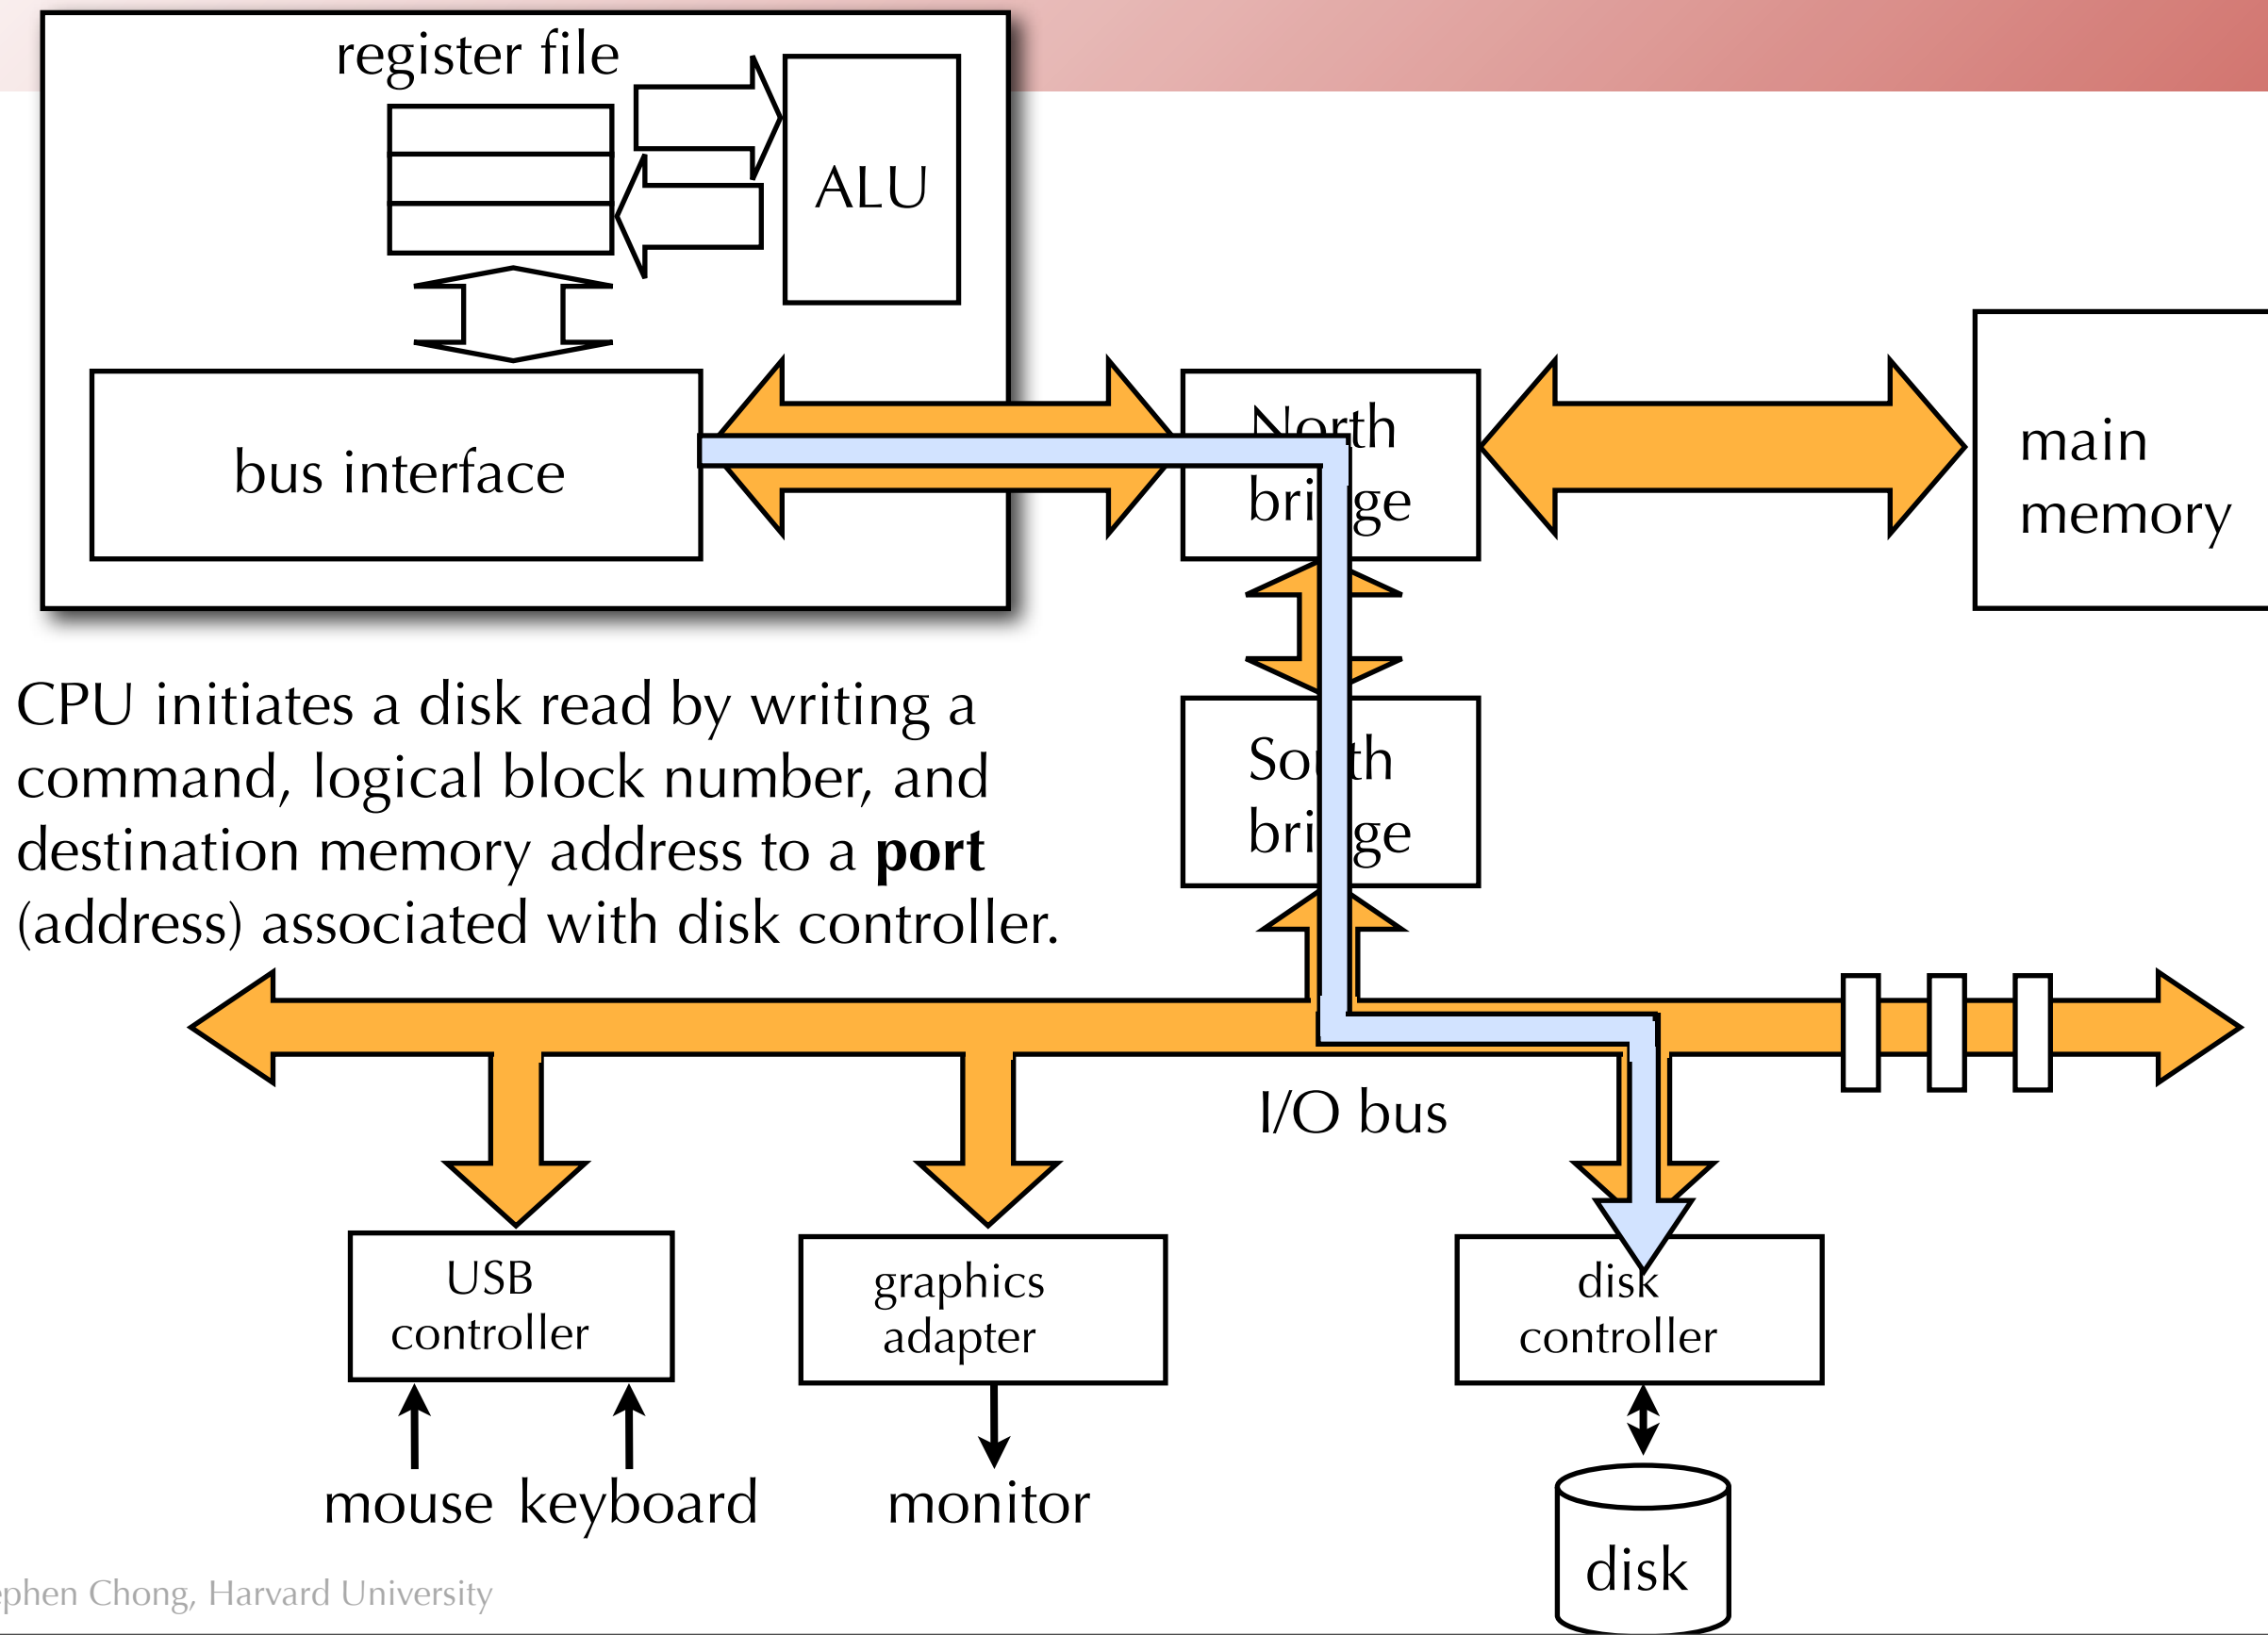
- Memory
- Disk drives
- I/O and memory buses
- Solid-state disks
- Storage technology trends

# I/O Bus

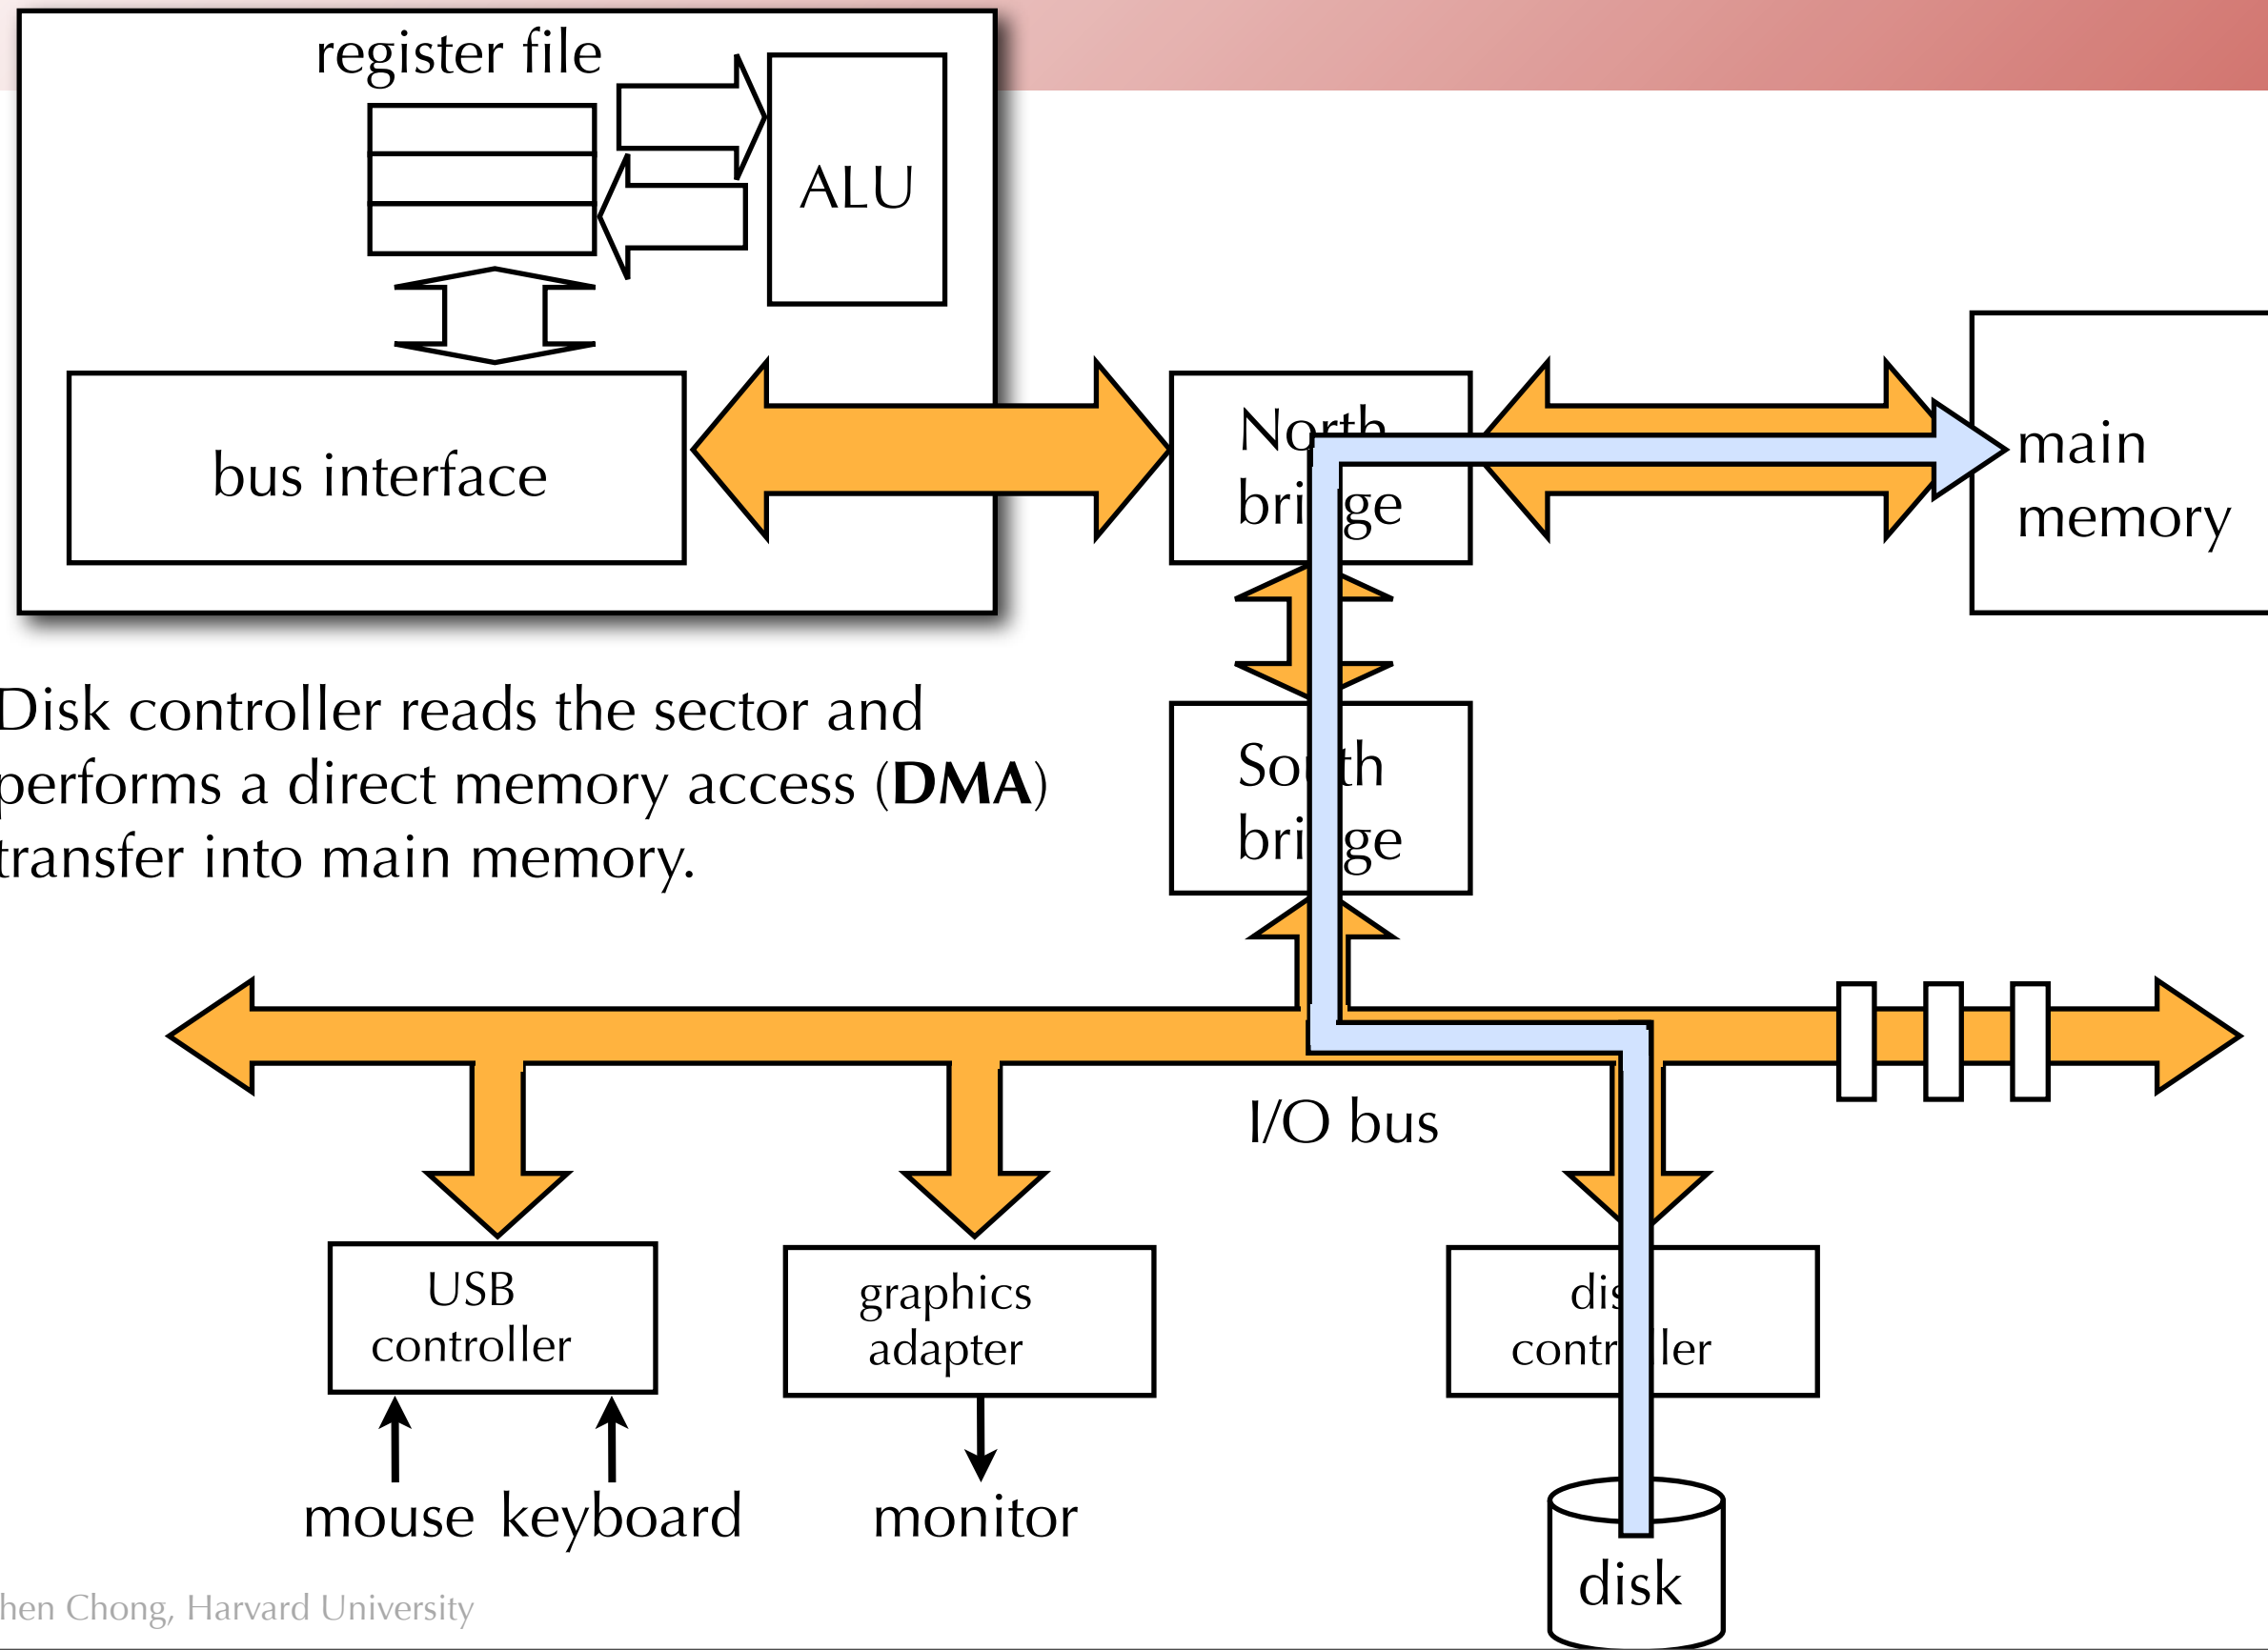
CPU chip



# Reading a disk sector



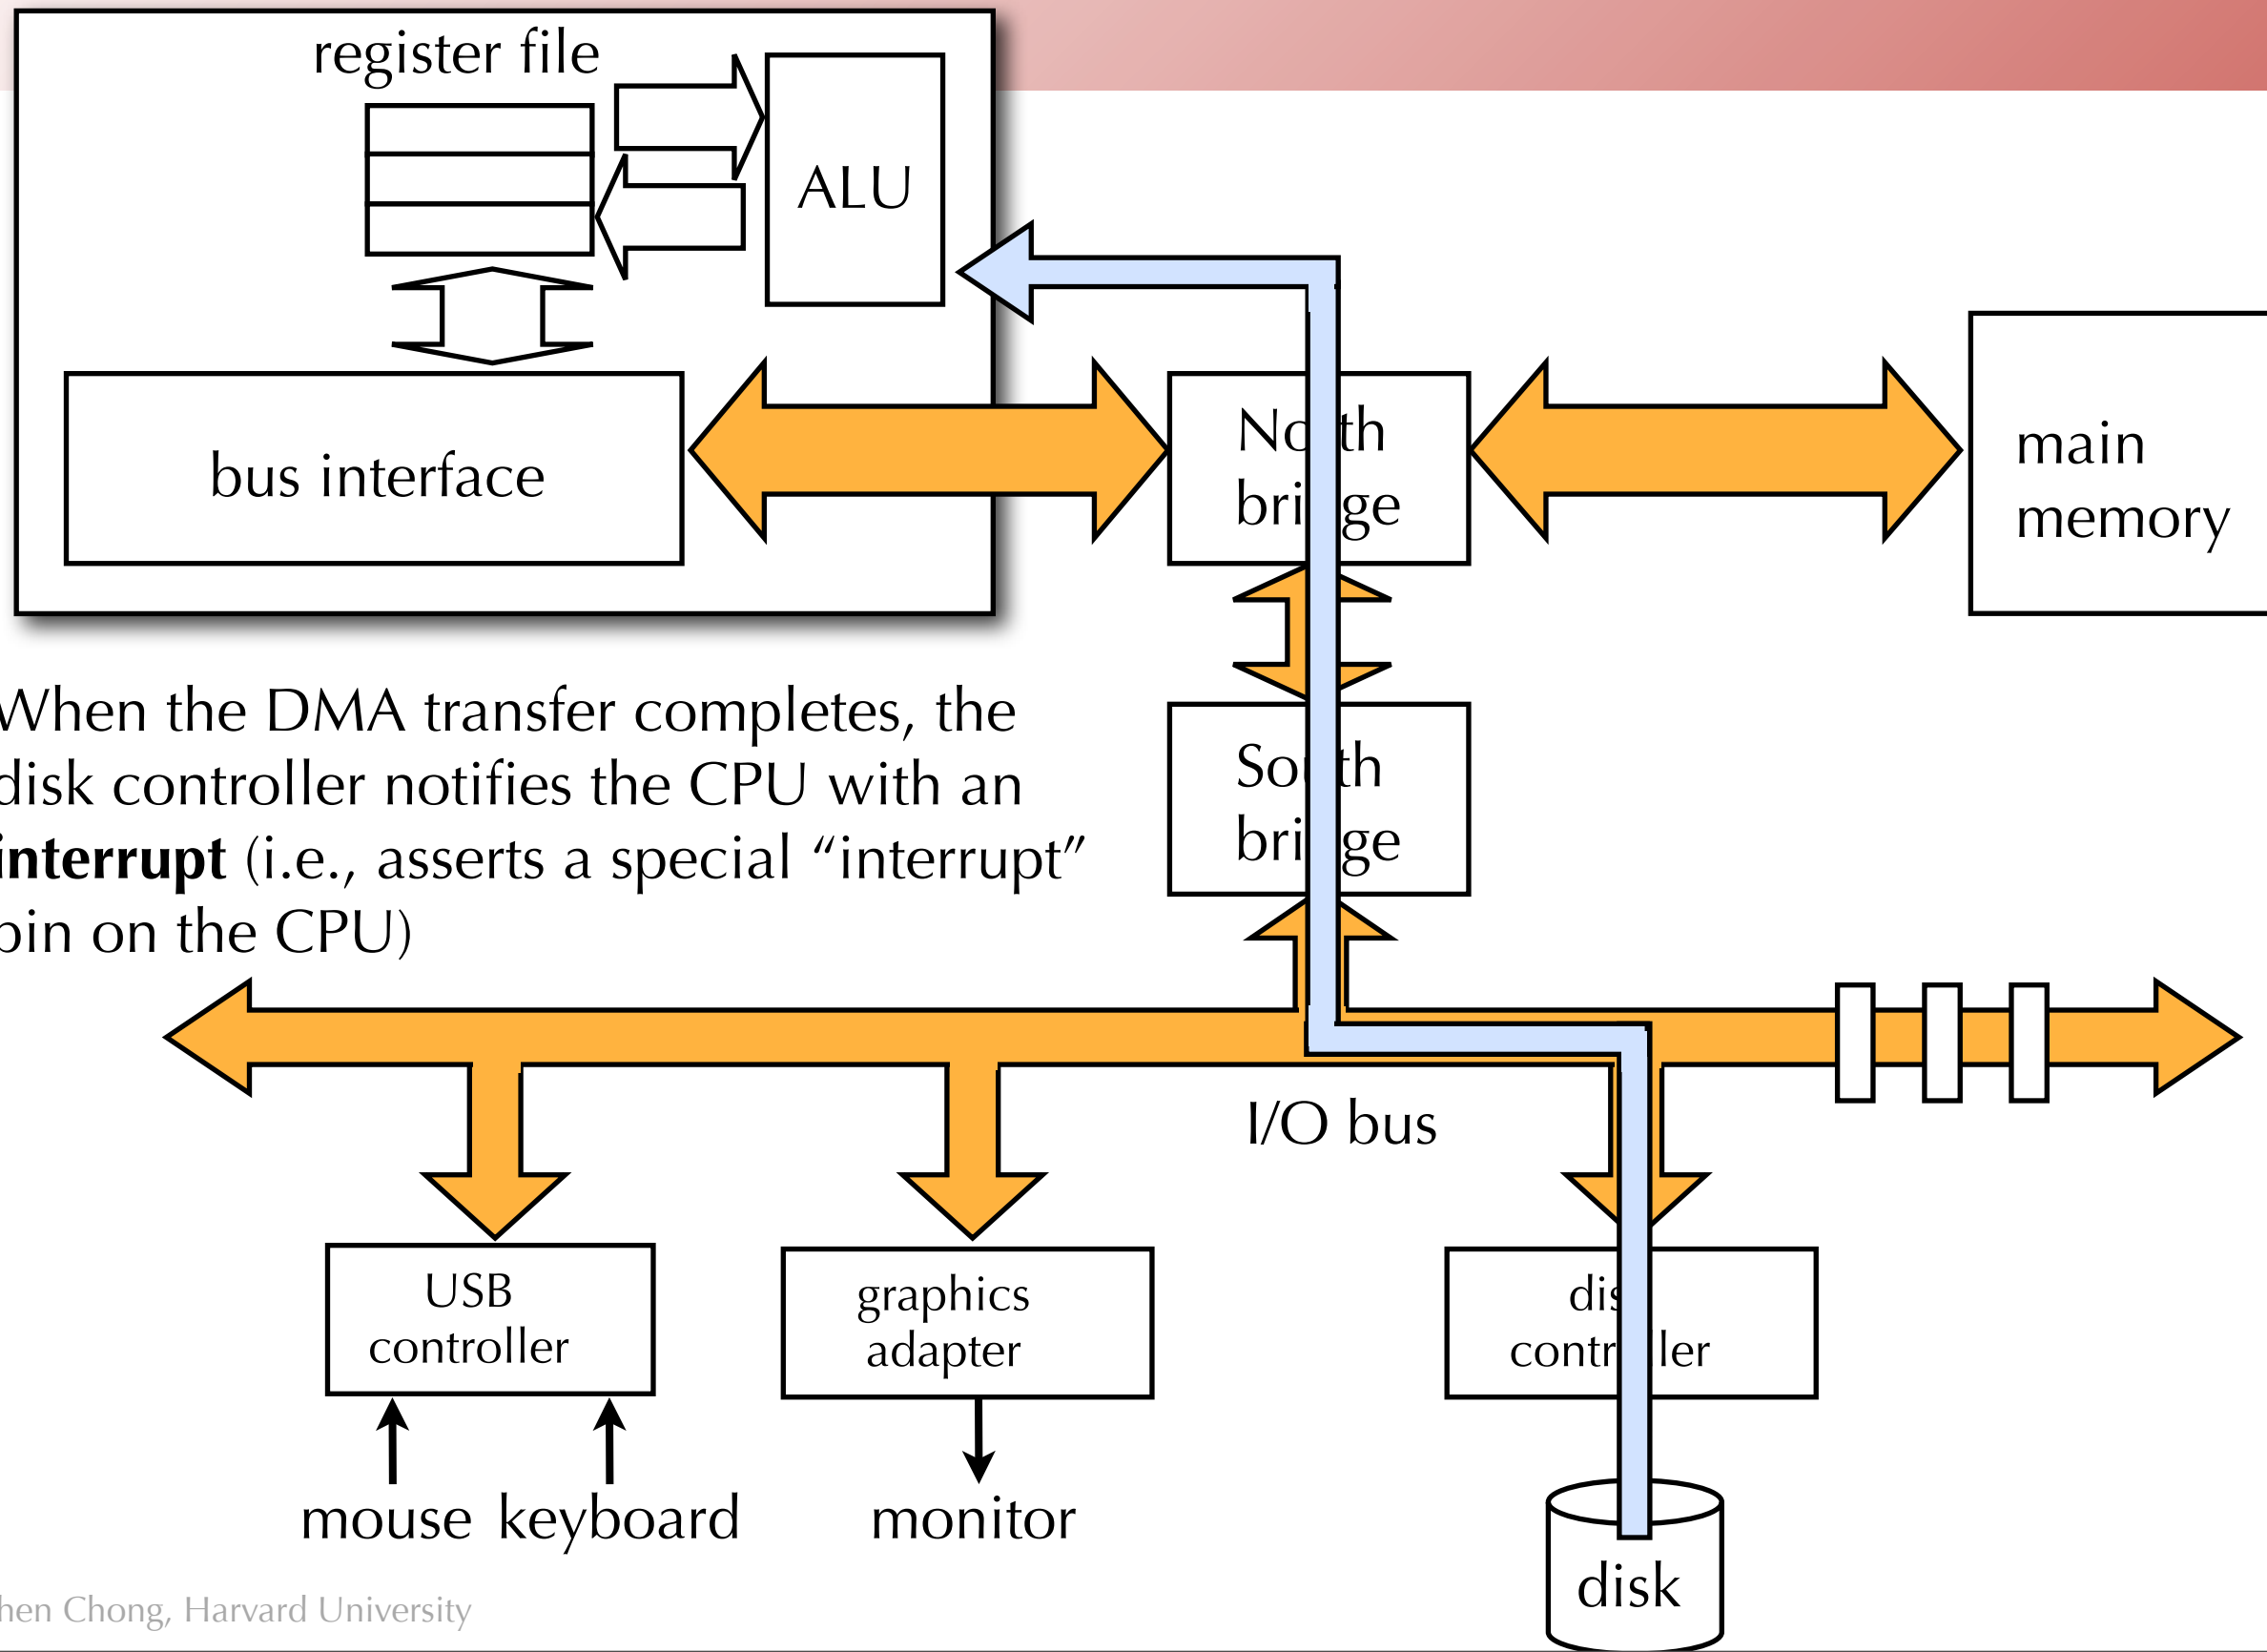
# Reading a disk sector



Disk controller reads the sector and performs a direct memory access (**DMA**) transfer into main memory.



# Reading a disk sector



When the DMA transfer completes, the disk controller notifies the CPU with an **interrupt** (i.e., asserts a special “interrupt” pin on the CPU)

# Today

- Memory
- Disk drives
- I/O and memory buses
- Solid-state disks
- Storage technology trends

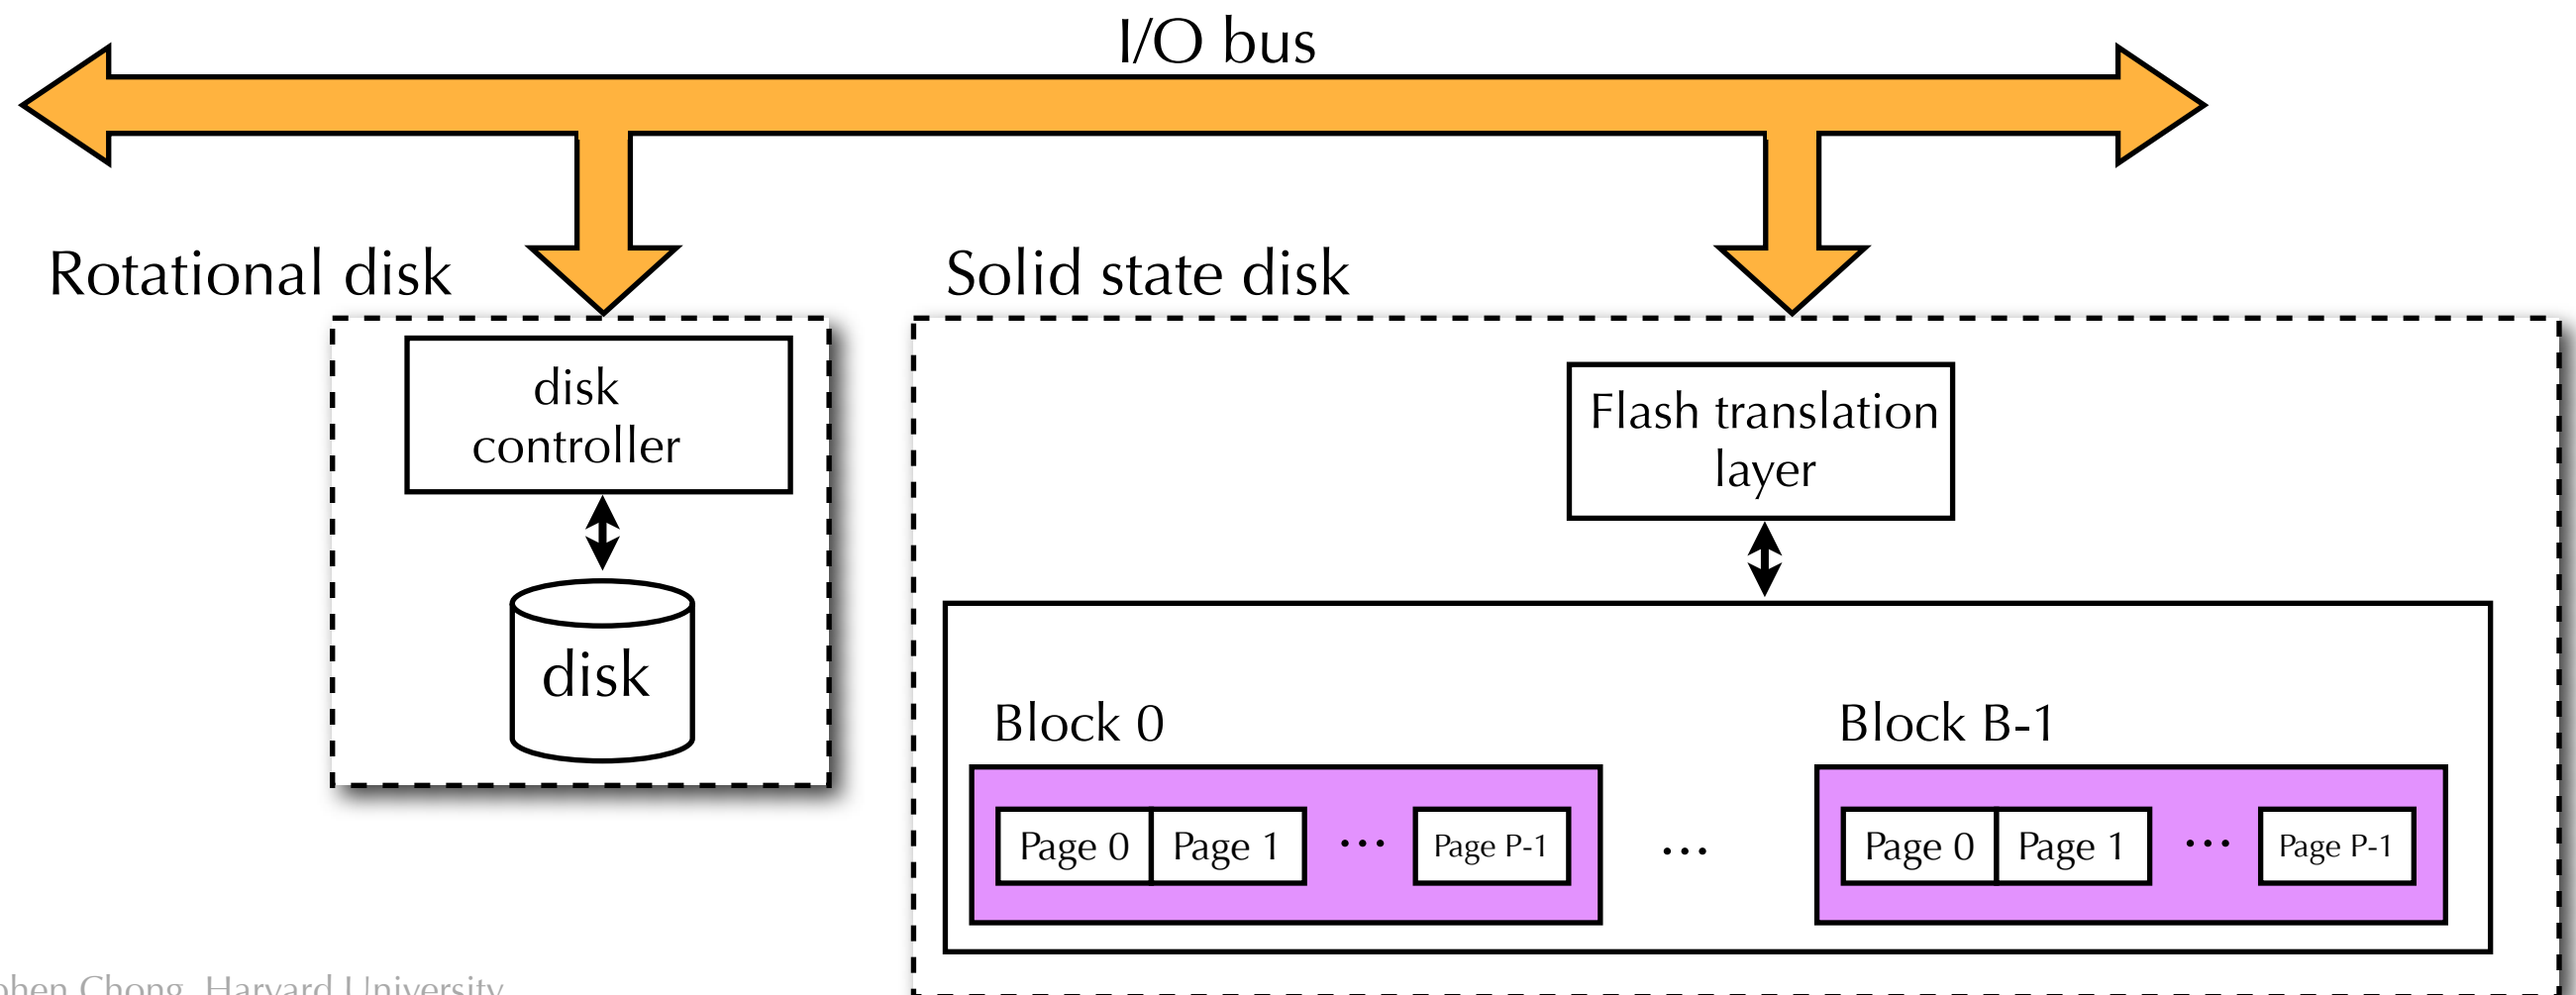
# Solid state disks (flash memory)

- Non-volatile, solid state storage
  - No moving parts!
  - Fast access times (about 0.1 msec!!)
  - Lower power (no moving parts)
- Expensive: about \$2.50/GB versus ~\$0.25/GB for HDD.



# Solid state disks (flash memory)

- Accessed through logical blocks
  - Presents same interface as rotational disks
- Sequence of  $B$  blocks, each with  $P$  pages
  - Page typically 512B – 4KB,  $P$  is typically 32-128





# Solid state access times

- Reads

- Sequential read throughput: 250 MB/s
- Random read throughput: 140 MB/s
- Random read access time: 30μs

- Writes

- Sequential write throughput: 170 MB/s
- Random write throughput: 14 MB/s
- Random write access time: 300μs

- Writes significantly slower!

- To write a page, entire block must first be erased
- Each block can be erased about 100,000 times before wearing out

# Solid state vs. rotating

- Pros:

- faster (no spinup, no seeking)
- less power
- more rugged (no moving parts, handle wider temperature range)
- quieter
- fewer errors

- Cons

- wear out (wear leveling to reduce this)
- wear leveling increases fragmentation
- more expensive
  - But cost coming down

- Stay tuned for innovative uses of flash memory...

# Today

- Memory
- Disk drives
- I/O and memory buses
- Solid-state disks
- Storage technology trends

# Access times in perspective

- 2.26 GHz processor  $\Rightarrow$  1 cycle = 0.44 ns
- Use physical distance as analogy

Data		Distance analogy	
Storage technology	Time	Distance	Intuition
Access register			
SRAM access			
DRAM access			
Flash read access			
Flash write access			
Disk seek			



# Access times in perspective

- 2.26 GHz processor  $\Rightarrow$  1 cycle = 0.44 ns
- Use physical distance as analogy

Data		Distance analogy	
Storage technology	Time	Distance	Intuition
Access register	0.5 ns	0.5 m	<i>Within arms reach</i>
SRAM access	10 ns	10 m	<i>Office next door to you</i>
DRAM access	50 ns	50 m	<i>Office one floor away from you</i>
Flash read access	30 $\mu$ s	30 km	<i>1.5 times length of Manhattan island</i>
Flash write access	300 $\mu$ s	300 km	<i>approx. Boston to New York City</i>
Disk seek	9 ms	9,000 km	<i>approx. Boston to LA and back</i>

# Storage technology trends

- Different storage technologies have different price and performance trade-offs
- Price and performance properties are changing at different rates

# Storage technology trends

## SRAM

metric	1980	1985	1990	1995	2000	2005	2010	2010:1980
\$/MB	19,200	2,900	320	256	100	75	60	320 x
access (ns)	300	150	35	15	3	2	1.5	200 x

## DRAM


metric	1980	1985	1990	1995	2000	2005	2010	2010:1980
\$/MB	8,000	880	100	30	1	0.1	0.06	130,000 x
access (ns)	375	200	100	70	60	50	40	9 x
typical size(MB)	0.064	0.256	4	16	64	2,000	8,000	125,000 x

## Disk

metric	1980	1985	1990	1995	2000	2005	2010	2010:1980
\$/MB	500	100	8	0.30	0.01	0.005	0.0003	1,600,000 x
access (ms)	87	75	28	10	8	5	3	29 x
typical size(MB)	1	10	160	1,000	20,000	160,000	1,500,000	1,500,000 x

# Storage technology trends

- Gap between CPU and memory increasing!





# Next lecture

- Caching!