



HARVARD

School of Engineering
and Applied Sciences

Virtual Memory

CS61, Lecture 15

Prof. Stephen Chong

October 20, 2011

Announcements

- Midterm review session: Monday Oct 24
 - 5:30pm to 7pm, 60 Oxford St. room 330
 - Large and small group interaction

Wall of Flame

Rob Jelinek

Ryan Meitzer

Neal Wu

Andrew Wang

Josh Lee

Julie Zhang

Zhe Lu

Peter Lin

Joseph Ong

Rohit Chaki


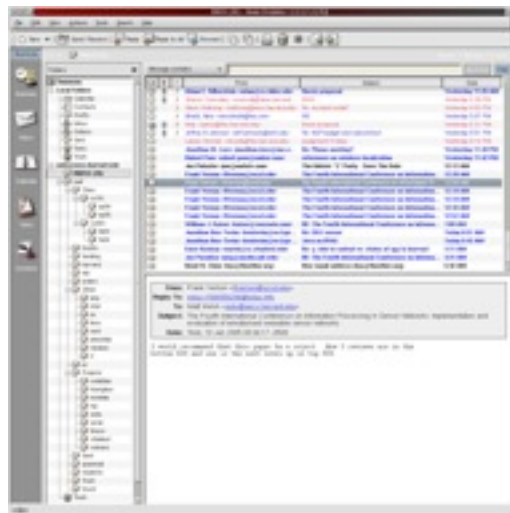
Alex Leigh

Today

- Running multiple programs at once
- Virtual memory
 - Address spaces
 - Benefits
 - Physical memory as a cache for virtual memory
 - Page tables
 - Page hits and page faults
 - Virtual memory as a tool for memory management
 - Virtual memory as a tool for memory protection
 - Address translation
 - Table Lookaside Buffer (TLB)
 - Example

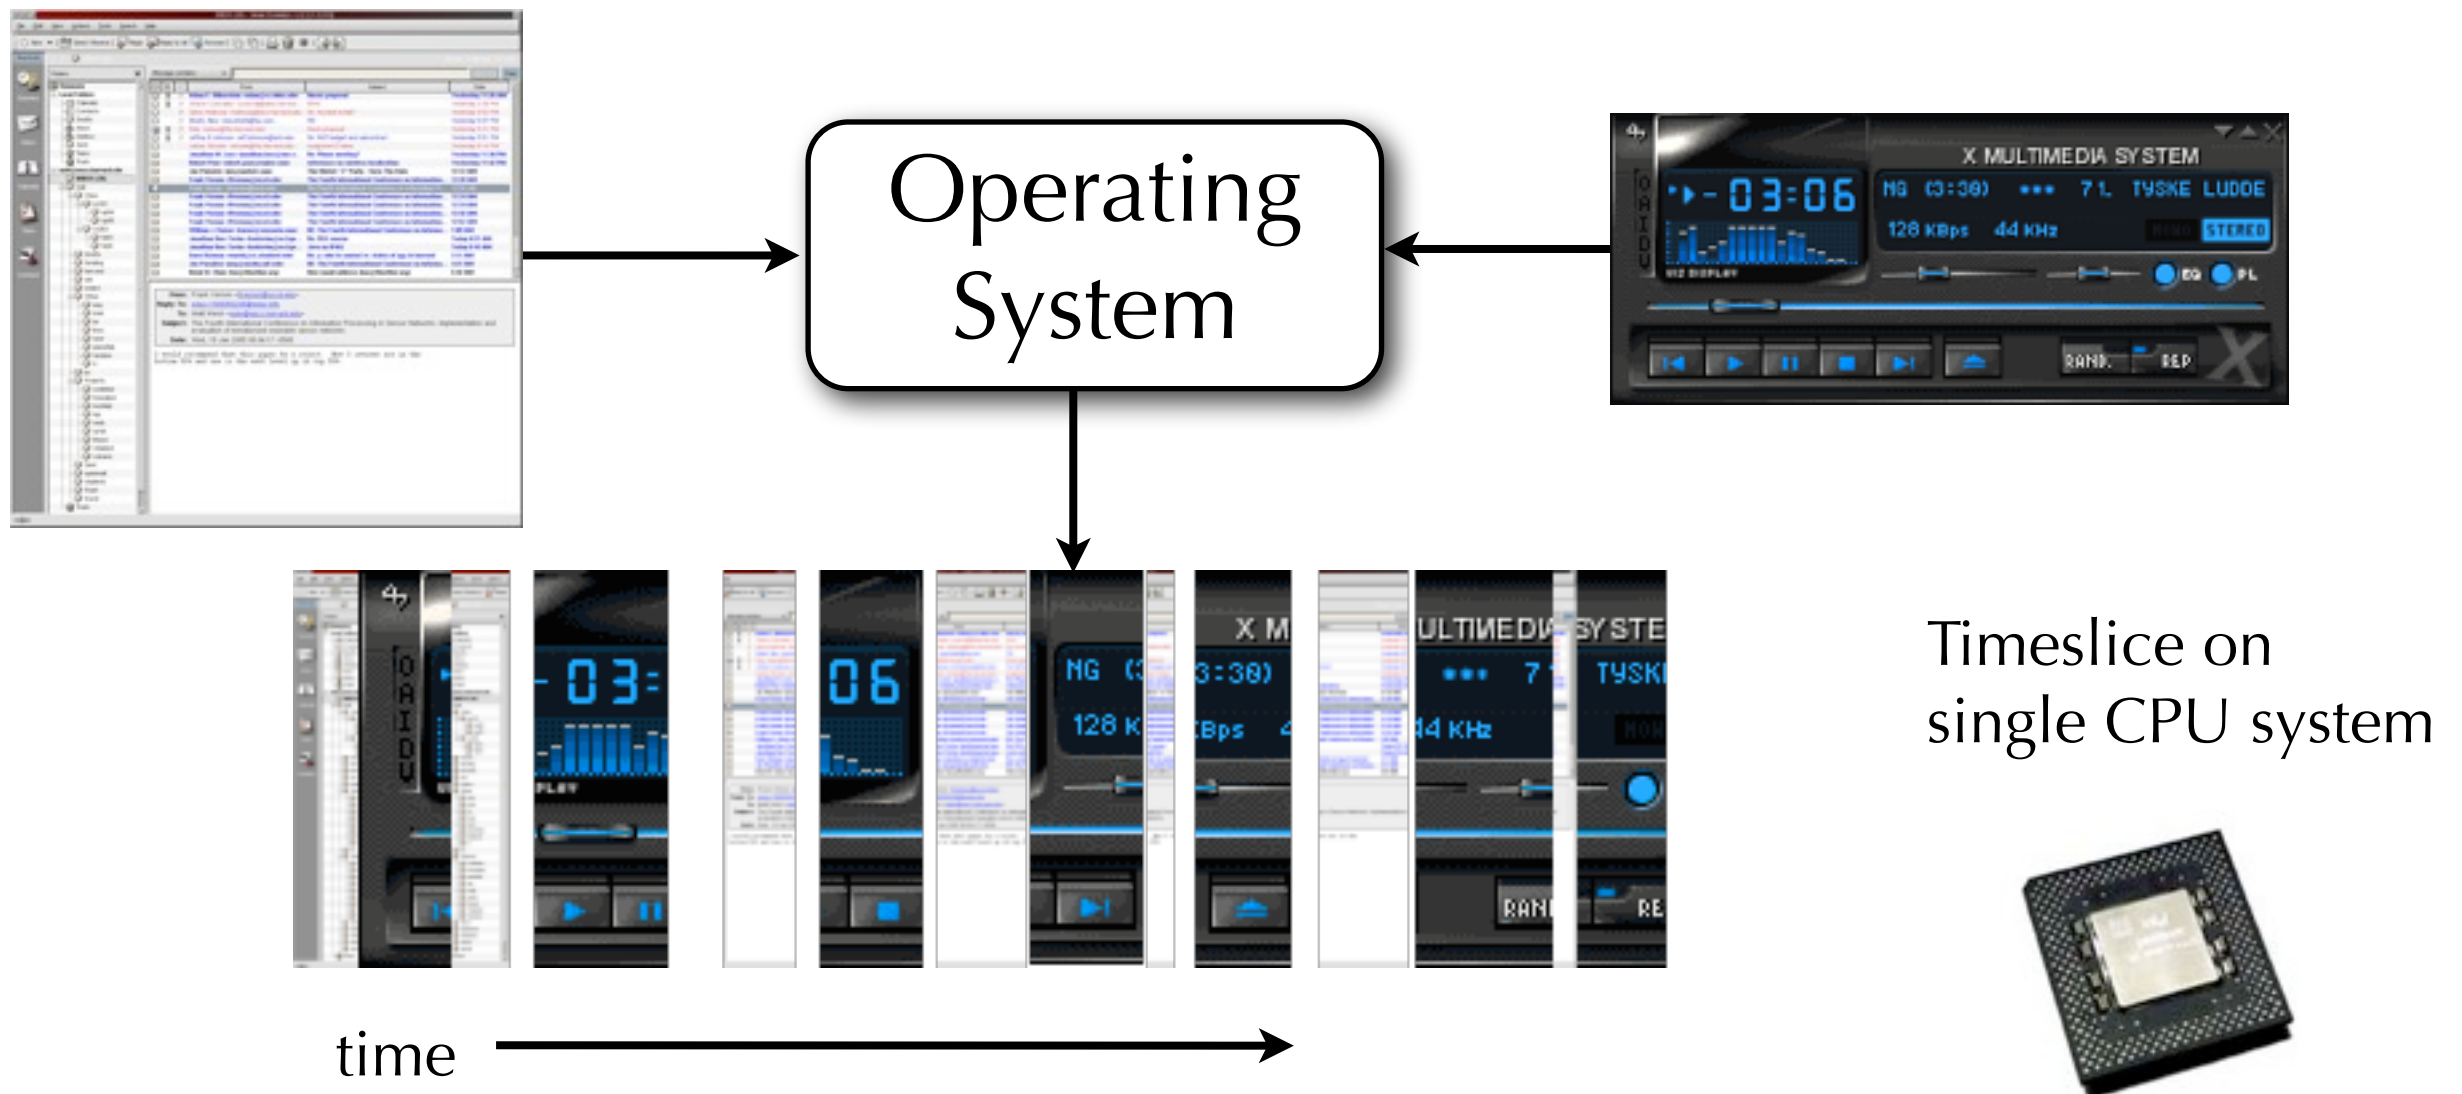
Running multiple programs at once

- So far, we have discussed memory management for a single program running on a computer.
- Most modern computers run multiple programs simultaneously.
 - This is called **multiprogramming** or **multitasking**.




Timeslicing

- The OS **timeslices** multiple applications on a single CPU
 - Switches between applications extremely rapidly, i.e., 100 times/sec
 - Each switch between two applications is called a **context switch**.




Virtual memory

- OS also gives every program illusion of having its own memory!
- Each program running on the machine is called a **process**.
 - Each process has a **process ID** (a number) and an **owner** – the user ID running the process.
 - UNIX command **ps** prints out info about the running processes.
- Each process has its own **address space**
 - Contents of that process' memory, visible only to that process.
 - 4GB on 32 bit machine, ~16 TB on 64 bit machine




Problem 1: How Does Everything Fit?

4GB to 16 EB virtual memory




Physical main memory:
Few Gigabytes



Problem 2: Memory Management

Physical main memory


Process 1
Process 2
Process 3
...
Process n




*What goes
where?*



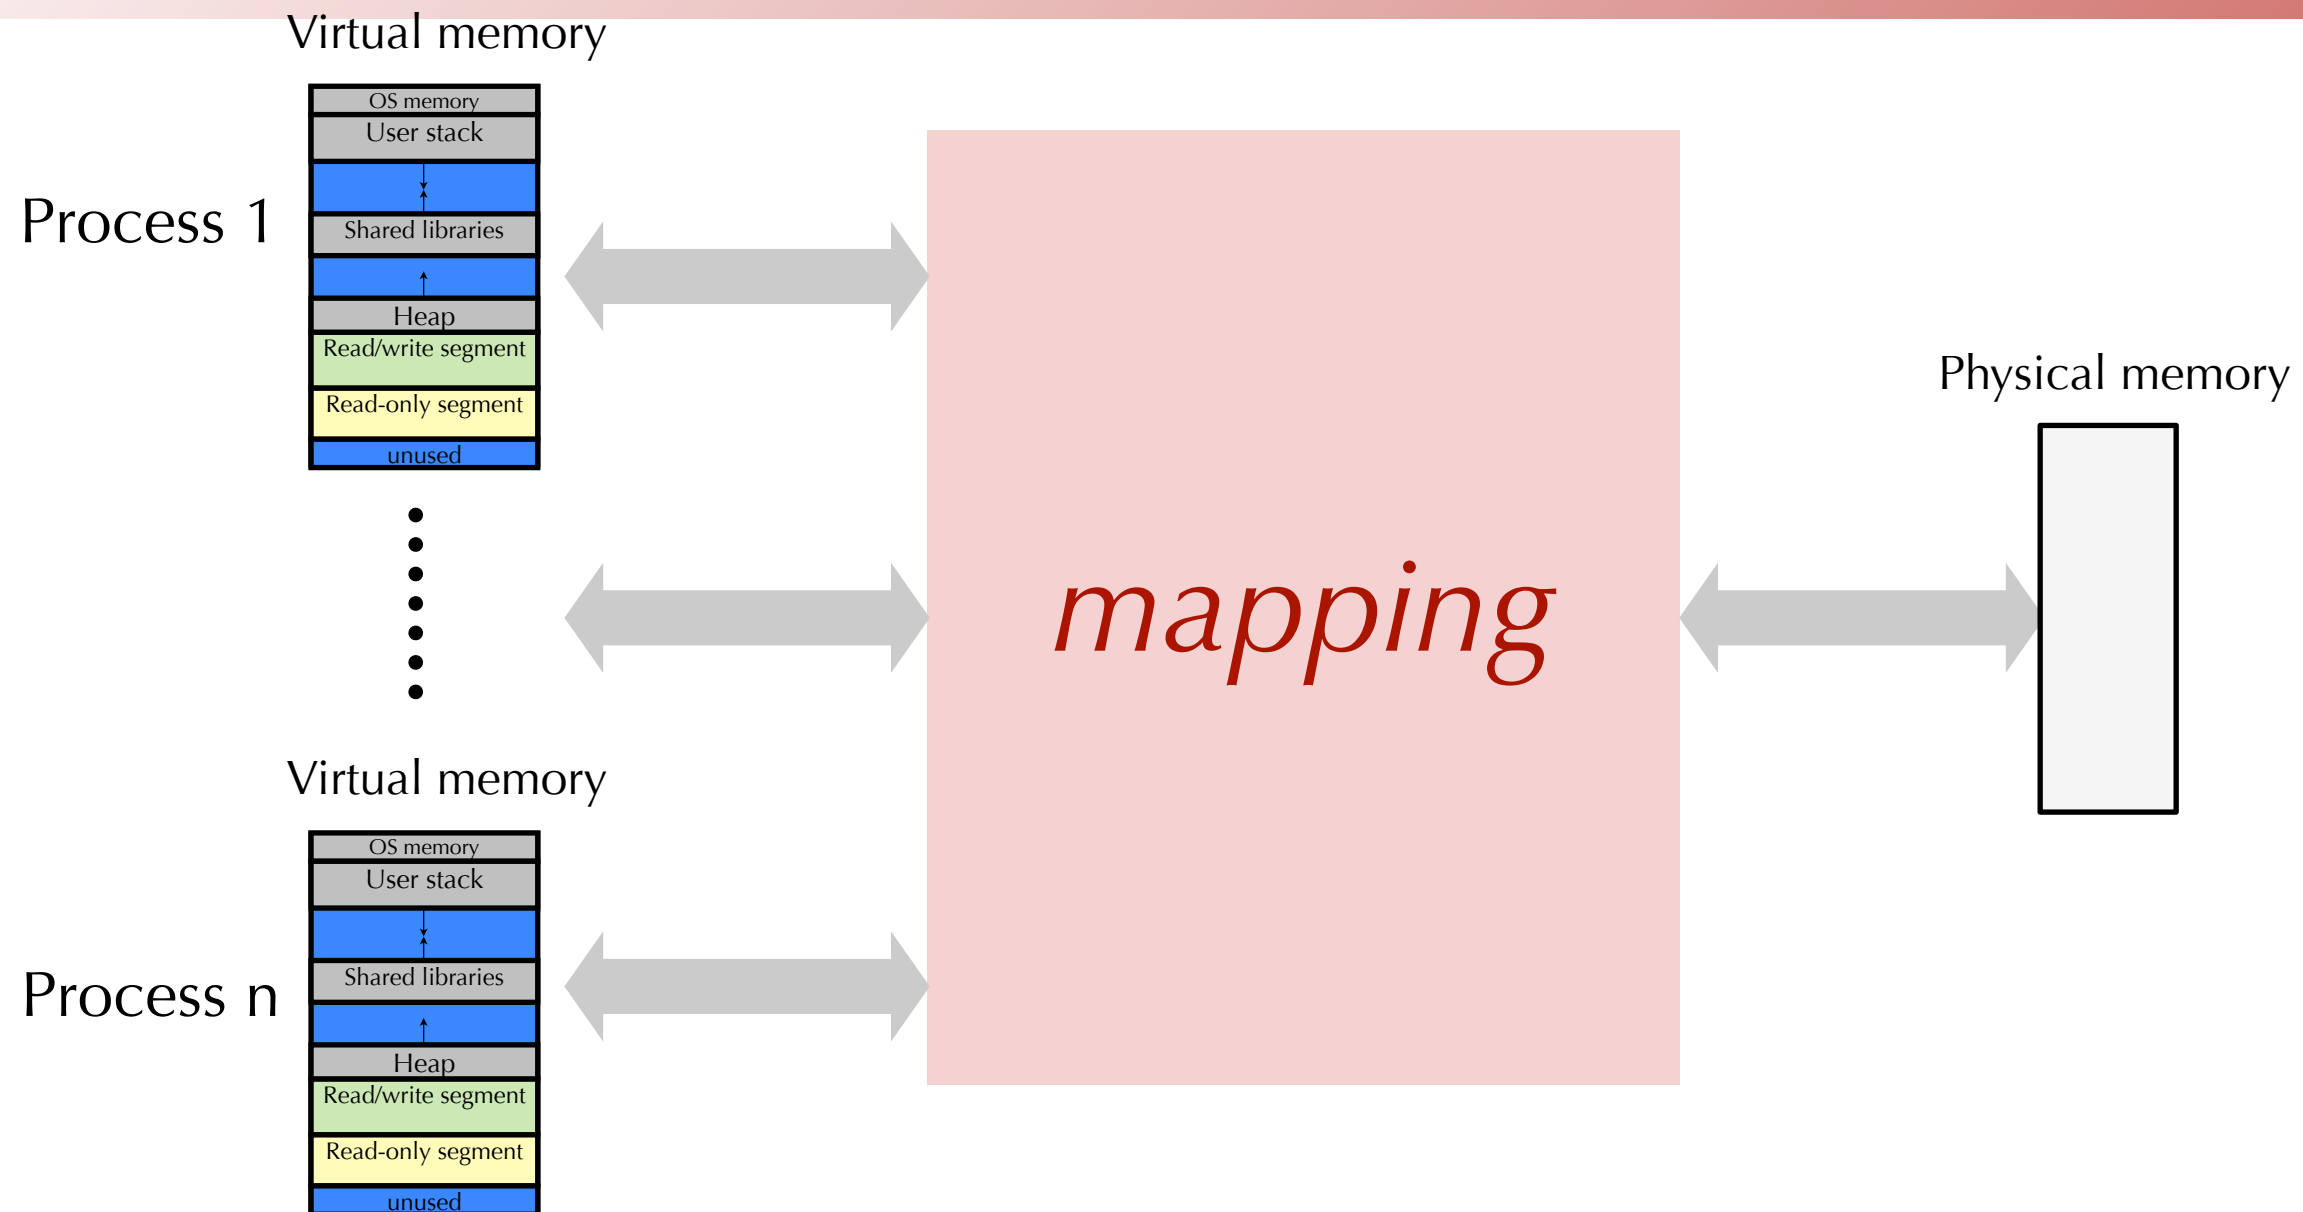
Problem 3: How To Protect



Problem 4: How To Share



Solution: Level Of Indirection



- Each process gets its own private memory space
- Solves the previous problems

Aside: indirection



All problems in computer science can be solved by another level of indirection...

Except for the problem of too many layers of indirection.

David Wheeler

1927–2004

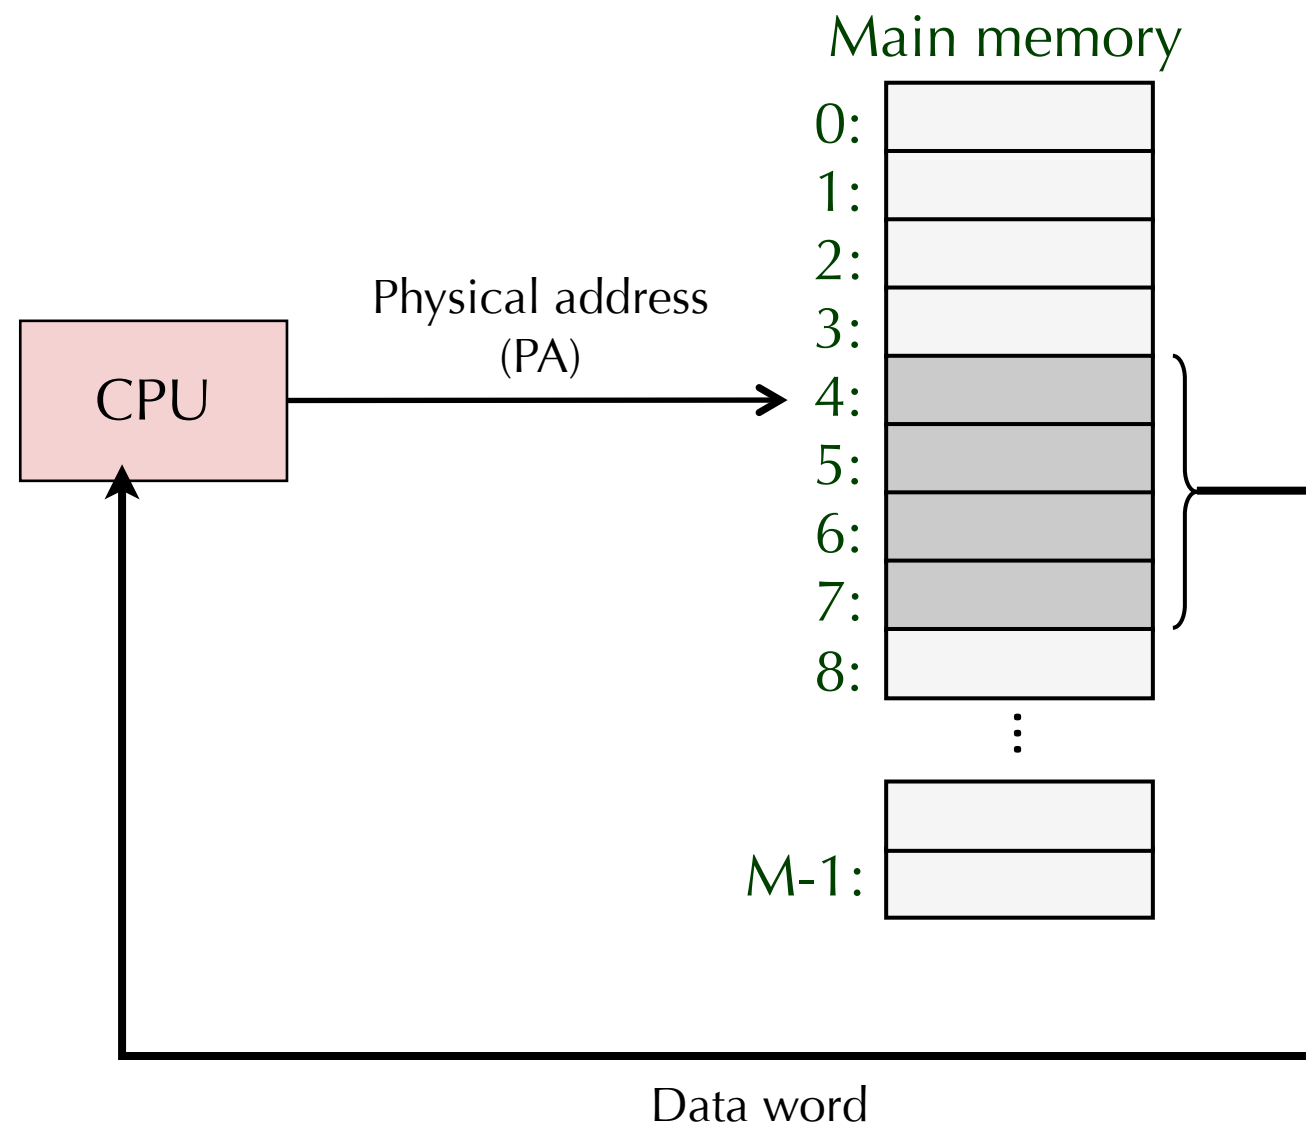
Computer Science Pioneer

World's first PhD in
Computer Science

Address Spaces

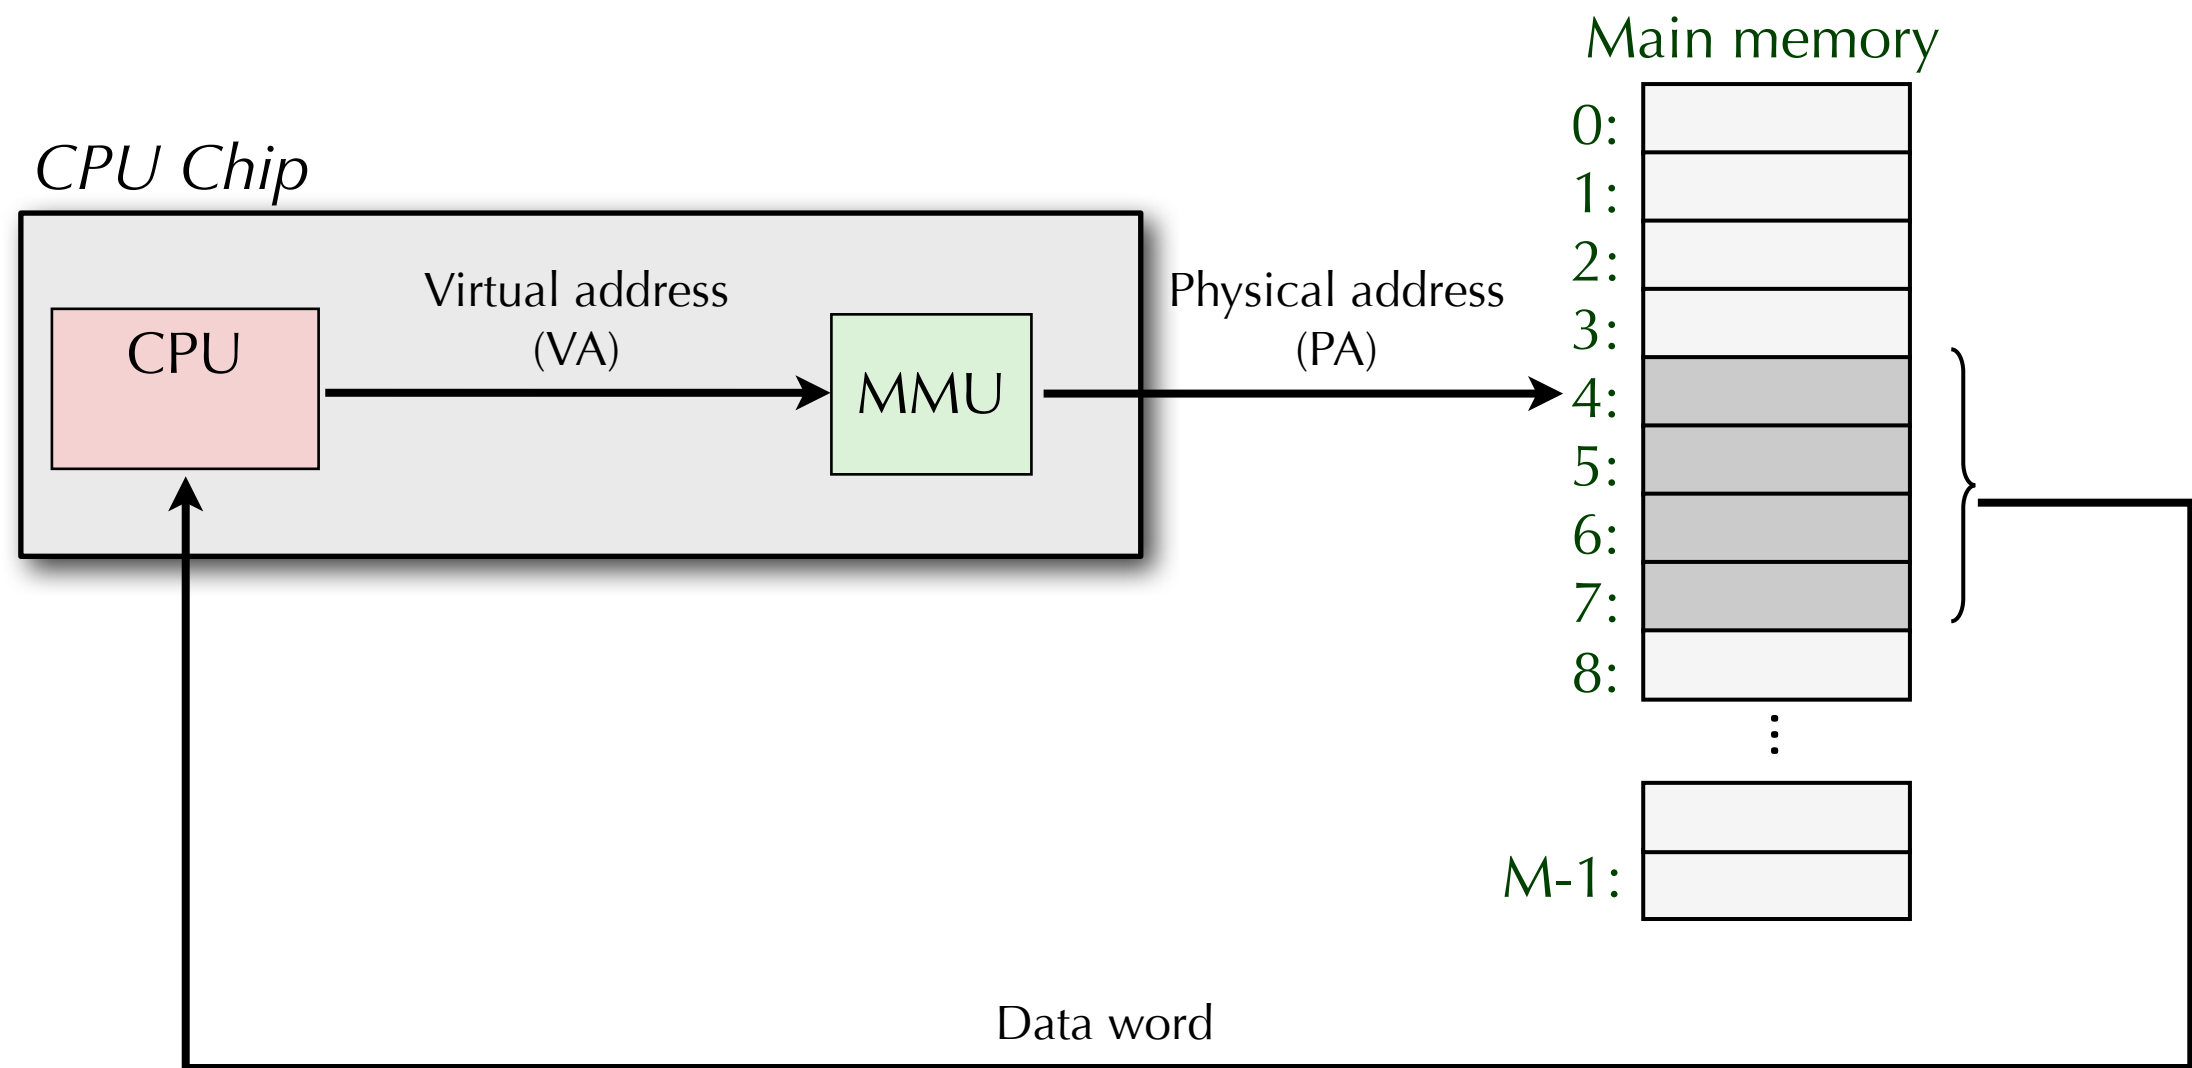
- **Linear address space:**
 - Ordered set of contiguous non-negative integer addresses:
 $\{0, 1, 2, 3 \dots\}$
- **Virtual address space:**
 - Set of $N = 2^n$ virtual addresses
 $\{0, 1, 2, 3, \dots, N-1\}$
- **Physical address space:**
 - Set of $M = 2^m$ physical addresses
 $\{0, 1, 2, 3, \dots, M-1\}$
- Clean distinction between data (bytes) and their attributes (addresses)
- Each object can now have multiple addresses
- Every byte in main memory: one physical address, one (or more) virtual addresses

A System Using Physical Addressing



- Used in “simple” systems like embedded microcontrollers in devices like cars, elevators, and digital picture frames

A System Using Virtual Addressing



- Used in all modern desktops, laptops, workstations
- One of the great ideas in computer science

Benefits of Virtual Memory (VM)

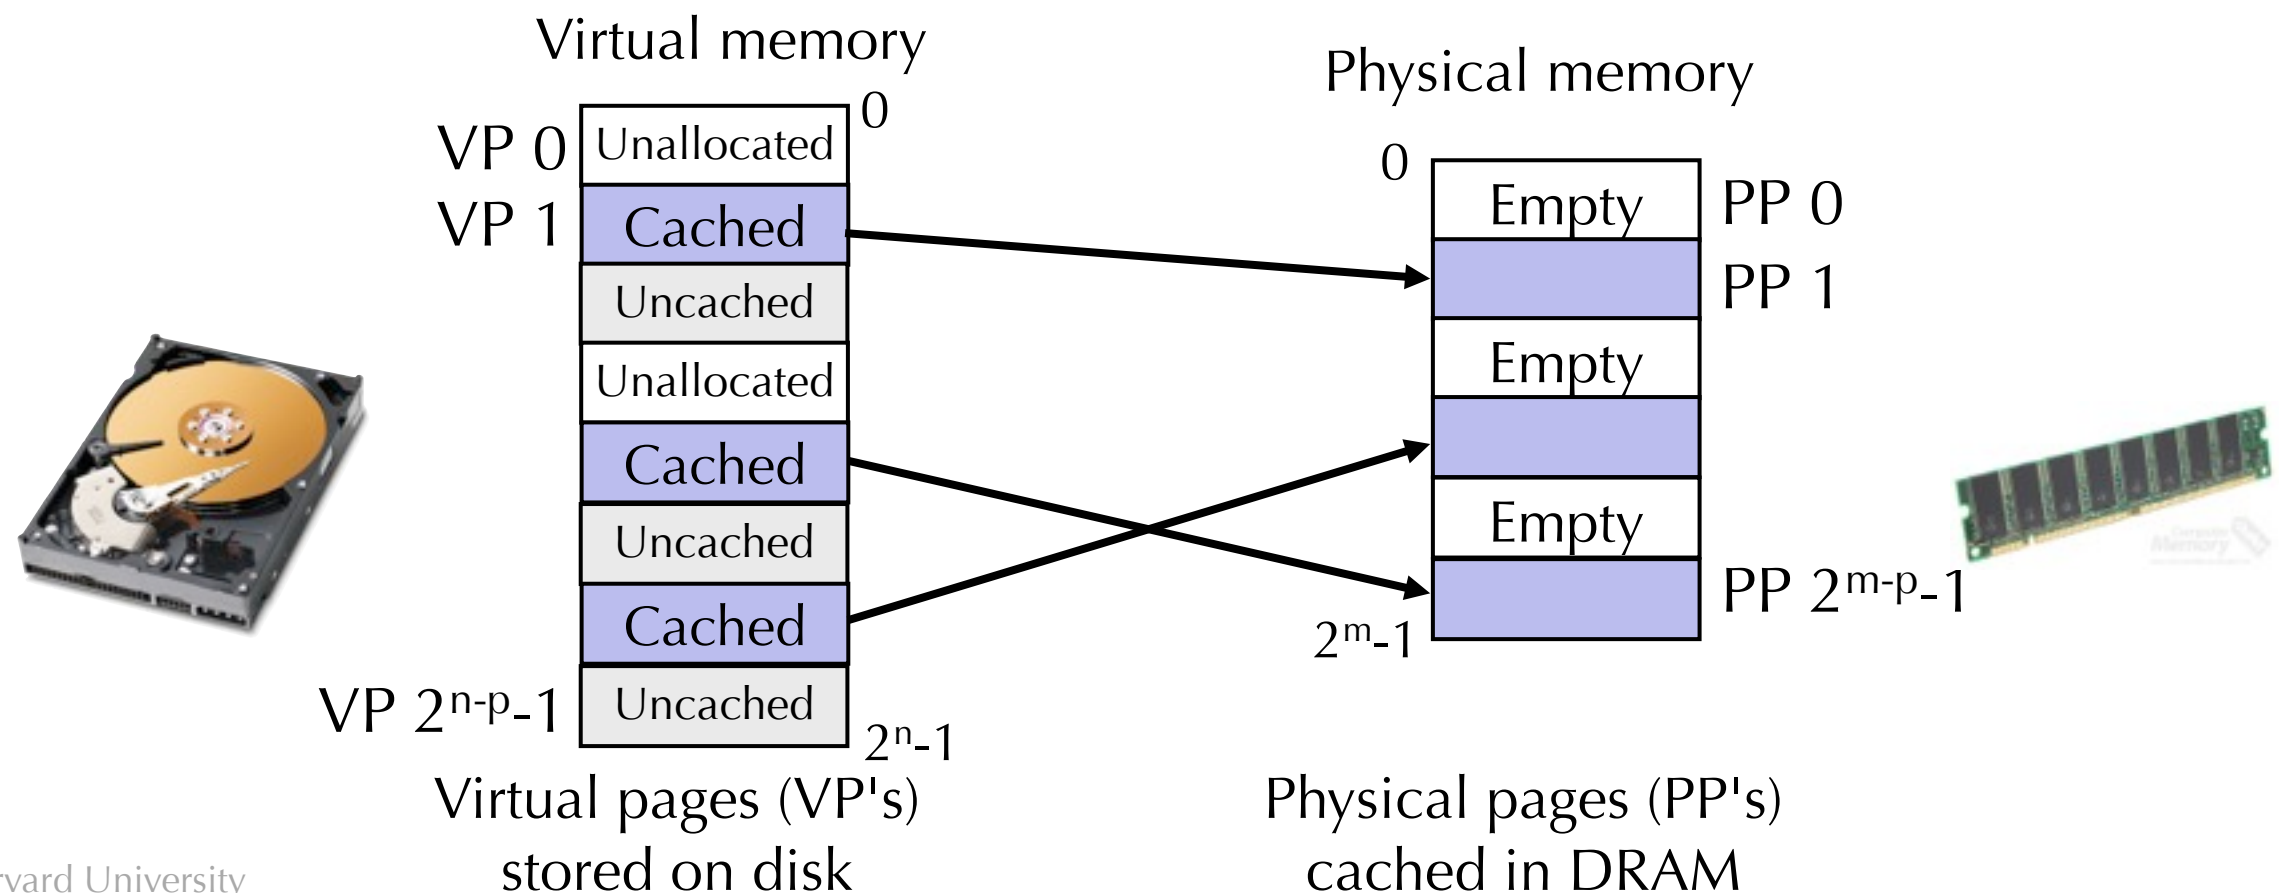
- Efficient use of limited main memory (RAM)
 - Use RAM as a cache for the parts of a virtual address space
 - some non-cached parts stored on disk
 - some (unallocated) non-cached parts stored nowhere
 - Keep only active areas of virtual address space in memory
 - transfer data back and forth as needed
- Simplifies memory management for programmers
 - Each process gets the same full, private linear address space
- Isolates address spaces
 - One process can't interfere with another's memory
 - because they operate in different address spaces
 - User process cannot access privileged information
 - different sections of address spaces have different permissions

Today

- Running multiple programs at once
- Virtual memory
 - Address spaces
 - Benefits
 - Physical memory as a cache for virtual memory
 - Page tables
 - Page hits and page faults
 - Virtual memory as a tool for memory management
 - Virtual memory as a tool for memory protection
 - Address translation
 - Table Lookaside Buffer (TLB)
 - Example


VM as a Tool for Caching

- Virtual memory: array of $N = 2^n$ contiguous bytes
 - think of the array (allocated part) as being stored on disk
- Physical main memory (DRAM) = cache for allocated virtual memory
- Blocks are called **pages**; size = 2^p



Memory hierarchy: Intel Core i7

Processor



DRAM Cache Organization

- DRAM cache organization driven by the enormous miss penalty
 - DRAM is about 10x slower than SRAM
 - Disk is about 100,000x slower than DRAM
 - For first byte, faster for next byte
- Consequences
 - Large page (block) size: typically 4-8 KB, sometimes 4 MB
 - **Fully associative**
 - Fully associate cache: A cache with only one set
 - Any virtual page can be placed in any physical page
 - Requires a “large” mapping function – different from CPU caches
 - Highly sophisticated, expensive replacement algorithms
 - Too complicated and open-ended to be implemented in hardware
 - Write-back rather than write-through

Why does virtual memory work?

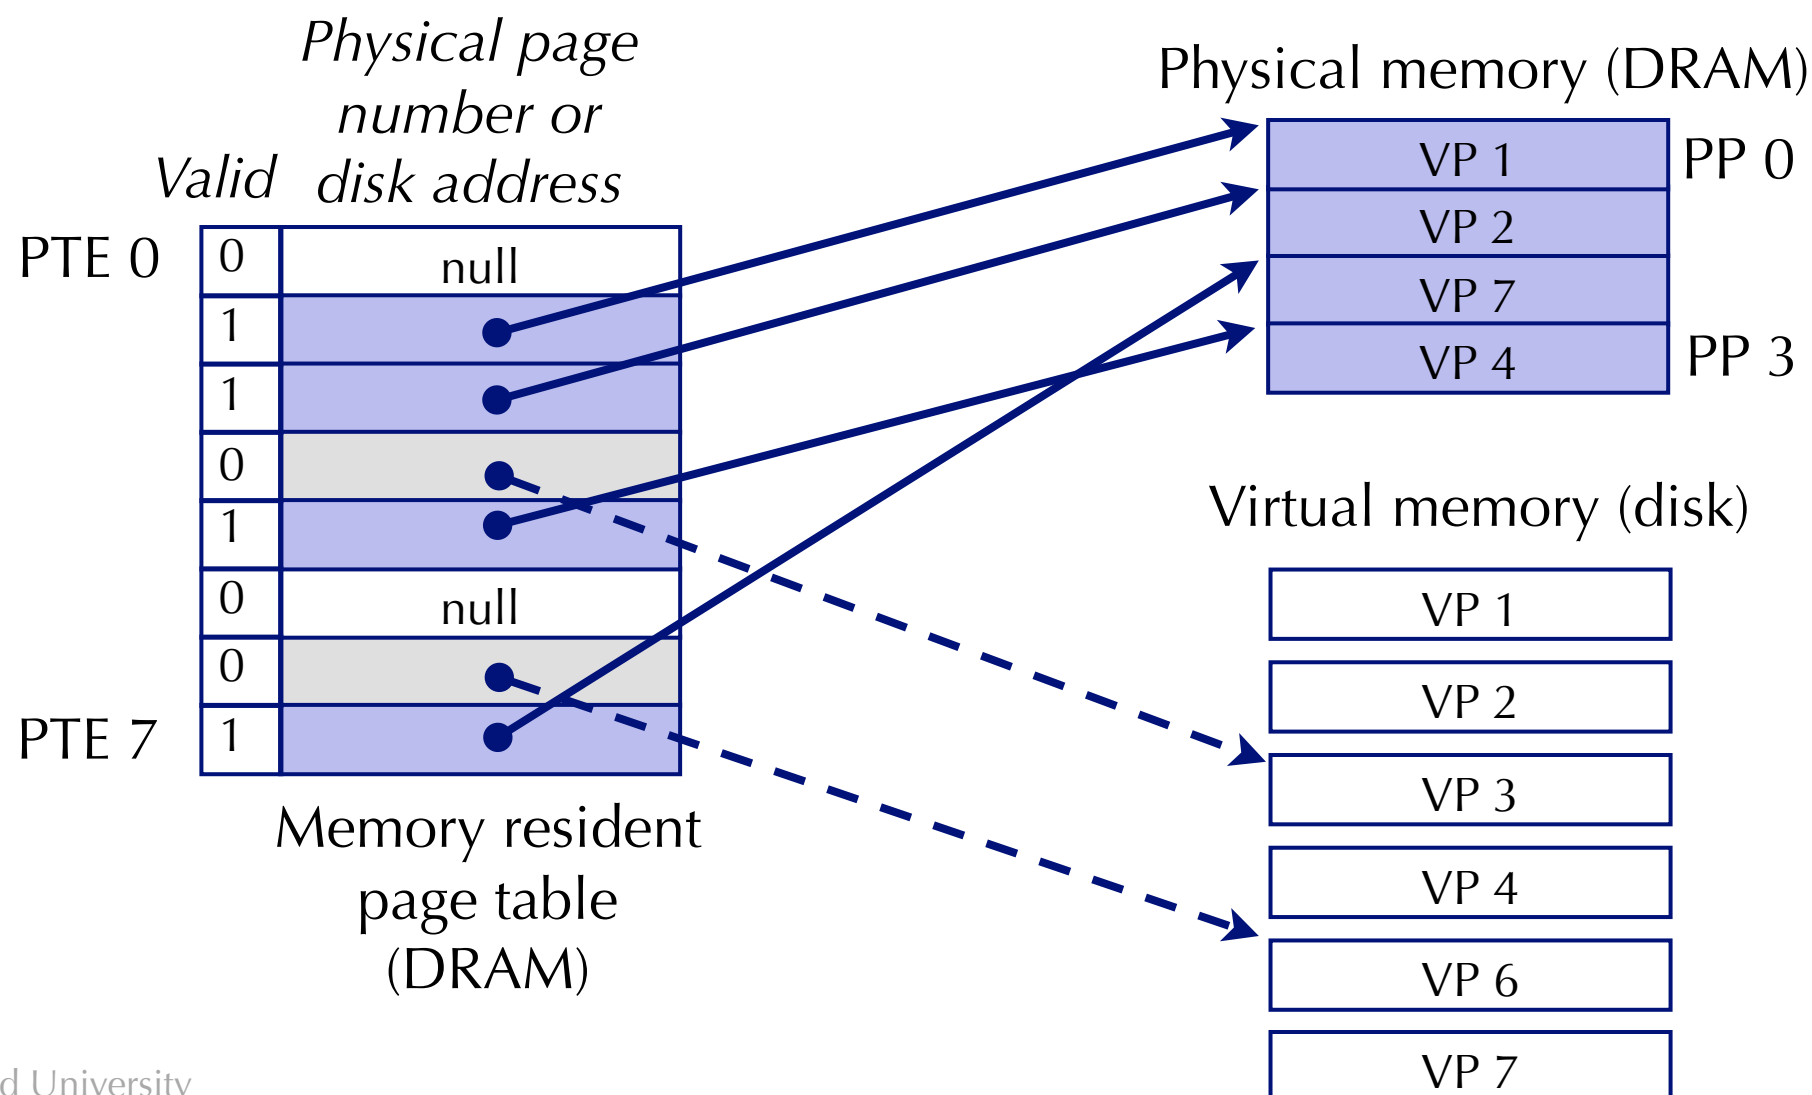
- Virtual memory works because of **locality**
- At any point in time, programs tend to access a set of active virtual pages called the working set
 - Programs with better temporal locality will have smaller working sets
- If (working set size $<$ main memory size)
 - Good performance for one process after compulsory misses
- If ($\text{SUM}(\text{working set sizes}) > \text{main memory size}$)
 - **Thrashing**: Performance meltdown where pages are swapped (copied) in and out continuously

Today

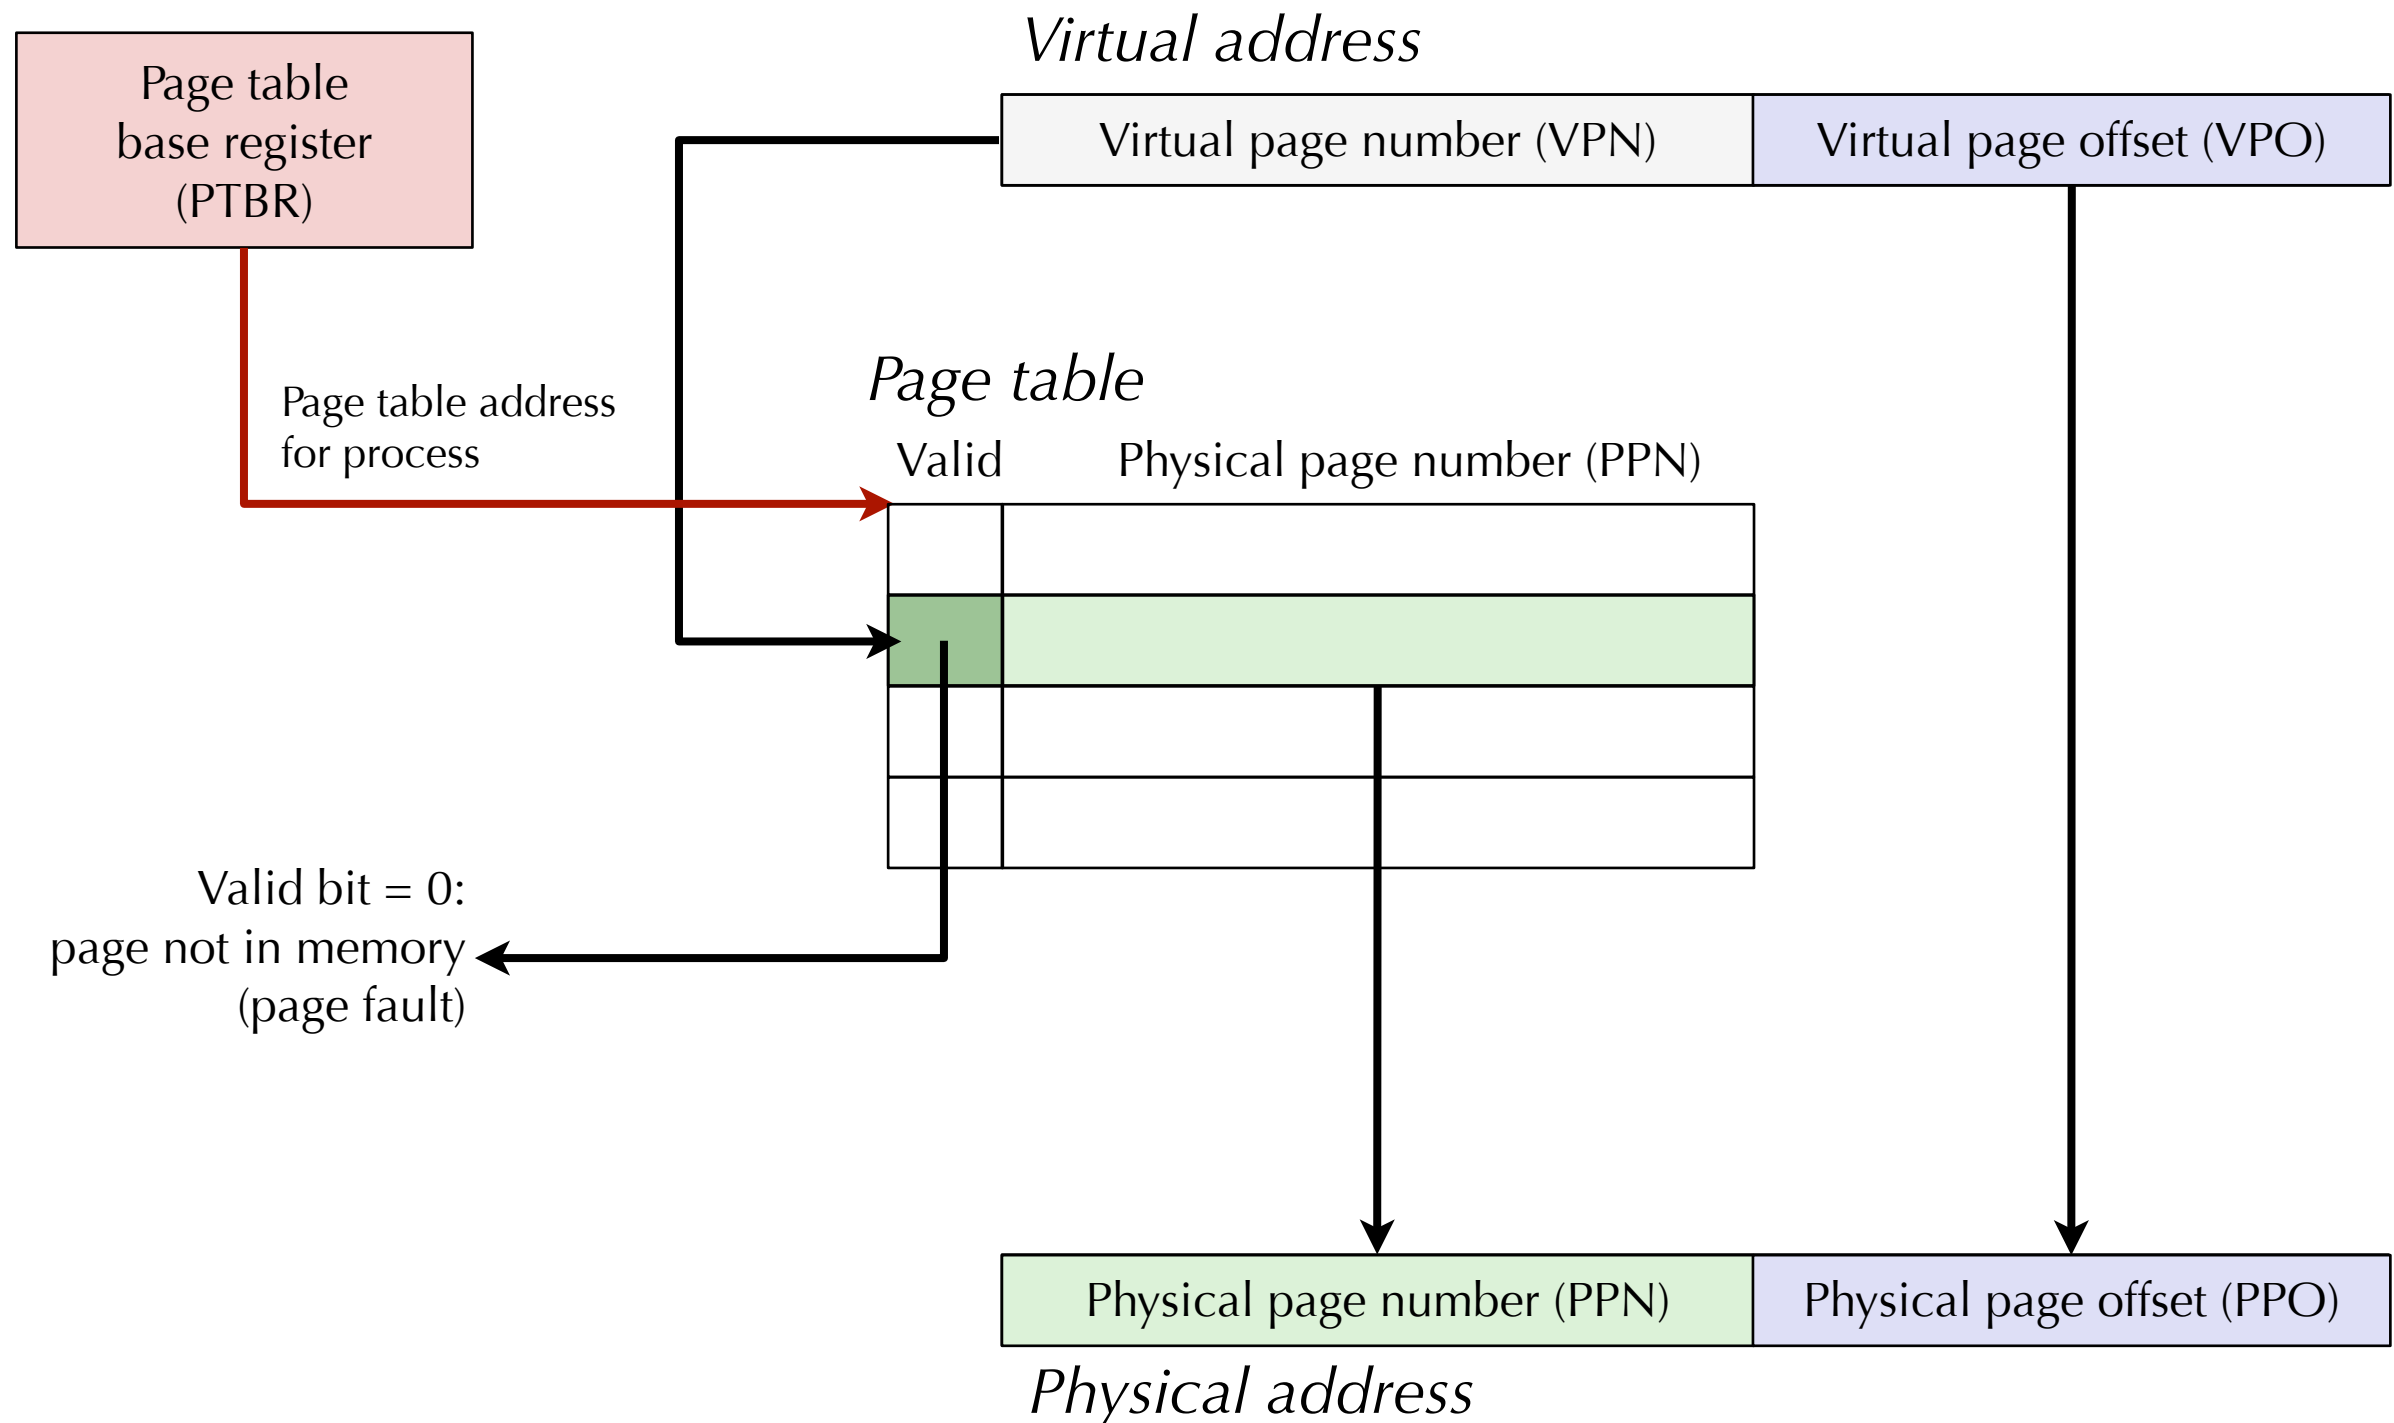
- Running multiple programs at once
- Virtual memory
 - Address spaces
 - Benefits
 - Physical memory as a cache for virtual memory
 - Page tables
 - Page hits and page faults
 - Virtual memory as a tool for memory management
 - Virtual memory as a tool for memory protection
 - Address translation
 - Table Lookaside Buffer (TLB)
 - Example

Page Tables

- A **page table** is an array of page table entries (PTEs) that maps virtual pages to physical pages. Here: 8 VPs
 - Per-process kernel data structure in DRAM



Address Translation With a Page Table



Page table size

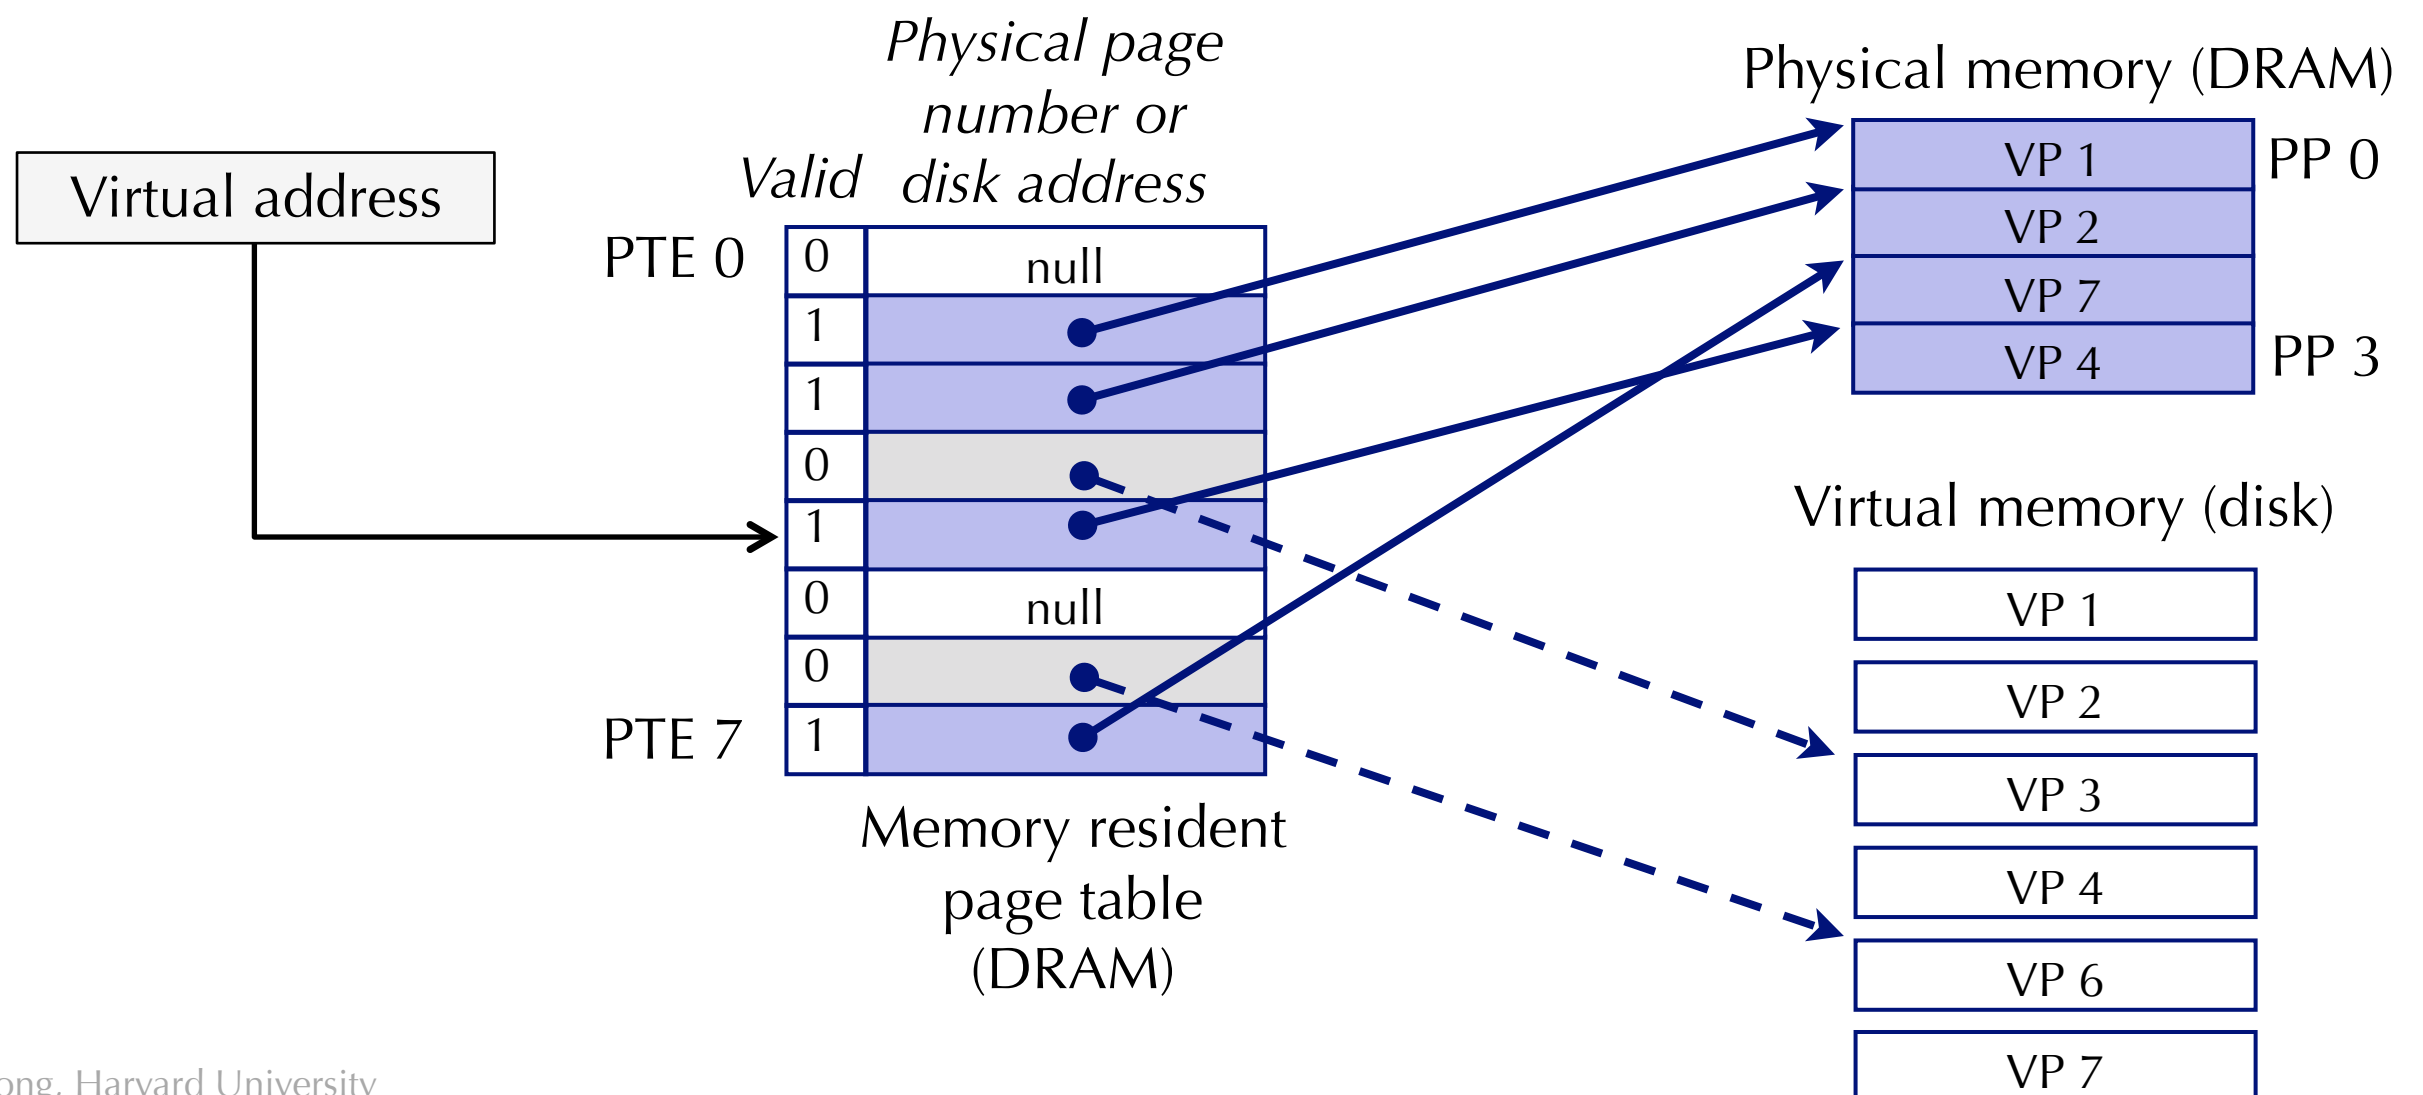
- How big is the page table for a process?
- Well ... we need one PTE per page.
- Say we have a 32-bit address space, and the page size is 4KB
- How many pages?
 - $2^{32} = 4$ GB total memory
 - $4\text{GB} / 4\text{KB} = 1,048,576 (= 1\text{M})$ pages = 1M PTEs
- How big is a PTE?
 - Depends on CPU architecture. On x86 it is 4 bytes
- Page table size for one process: $1\text{M PTEs} \times 4 \text{ bytes/PTE} = 4 \text{ MB}$
- For 100 processes that's 400 MB of memory just for page tables!!!
- Solution: swap page tables out to disk...

Today

- Running multiple programs at once
- Virtual memory
 - Address spaces
 - Benefits
 - Physical memory as a cache for virtual memory
 - Page tables
 - Page hits and page faults
 - Virtual memory as a tool for memory management
 - Virtual memory as a tool for memory protection
 - Address translation
 - Table Lookaside Buffer (TLB)
 - Example

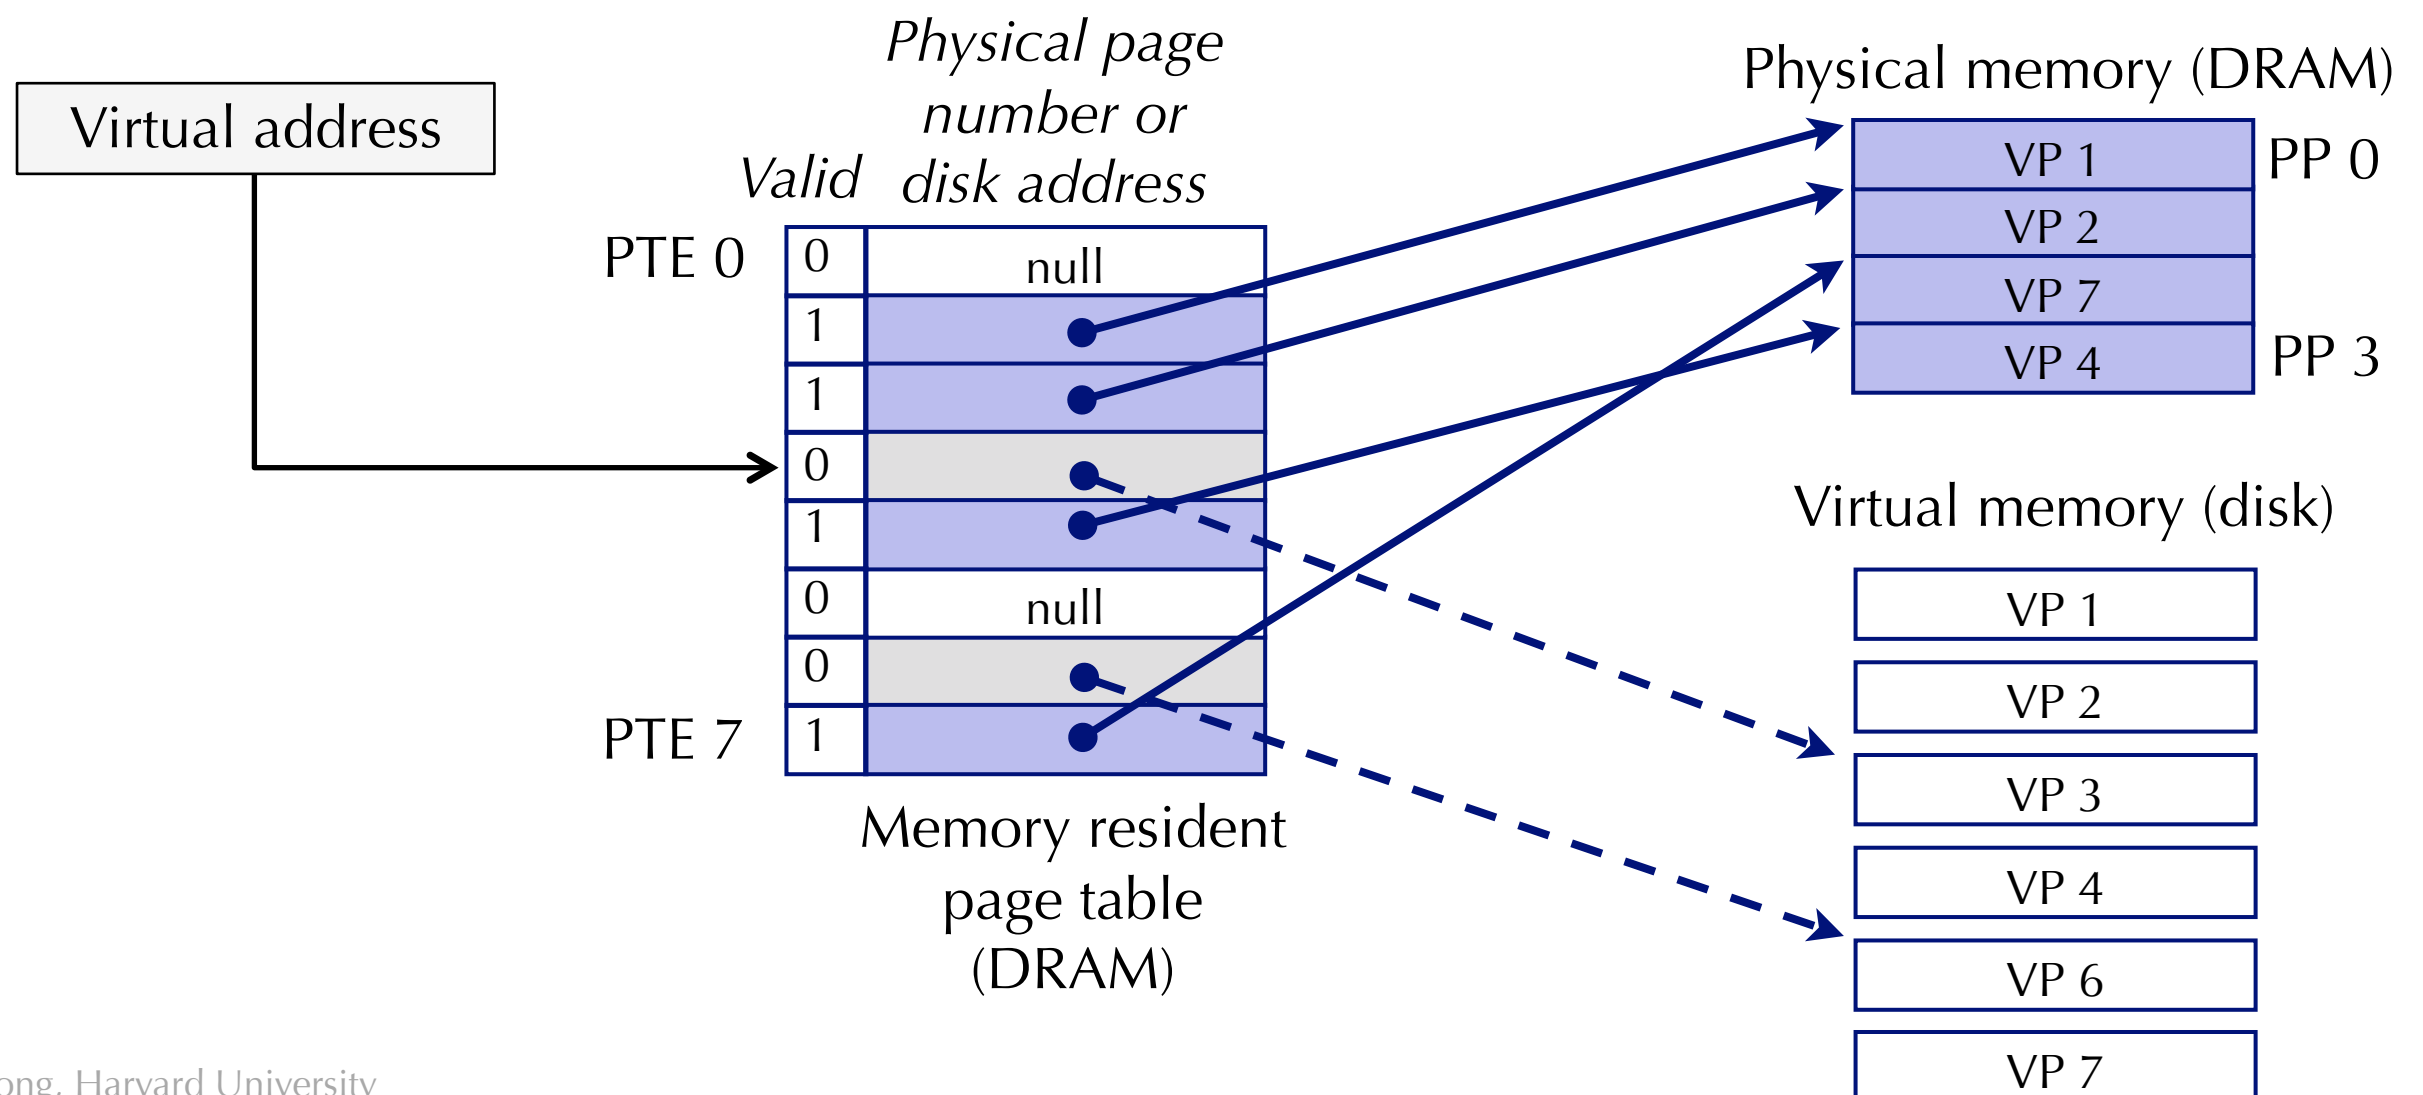
Page Hit

- **Page hit:** reference to VM word that is in physical memory



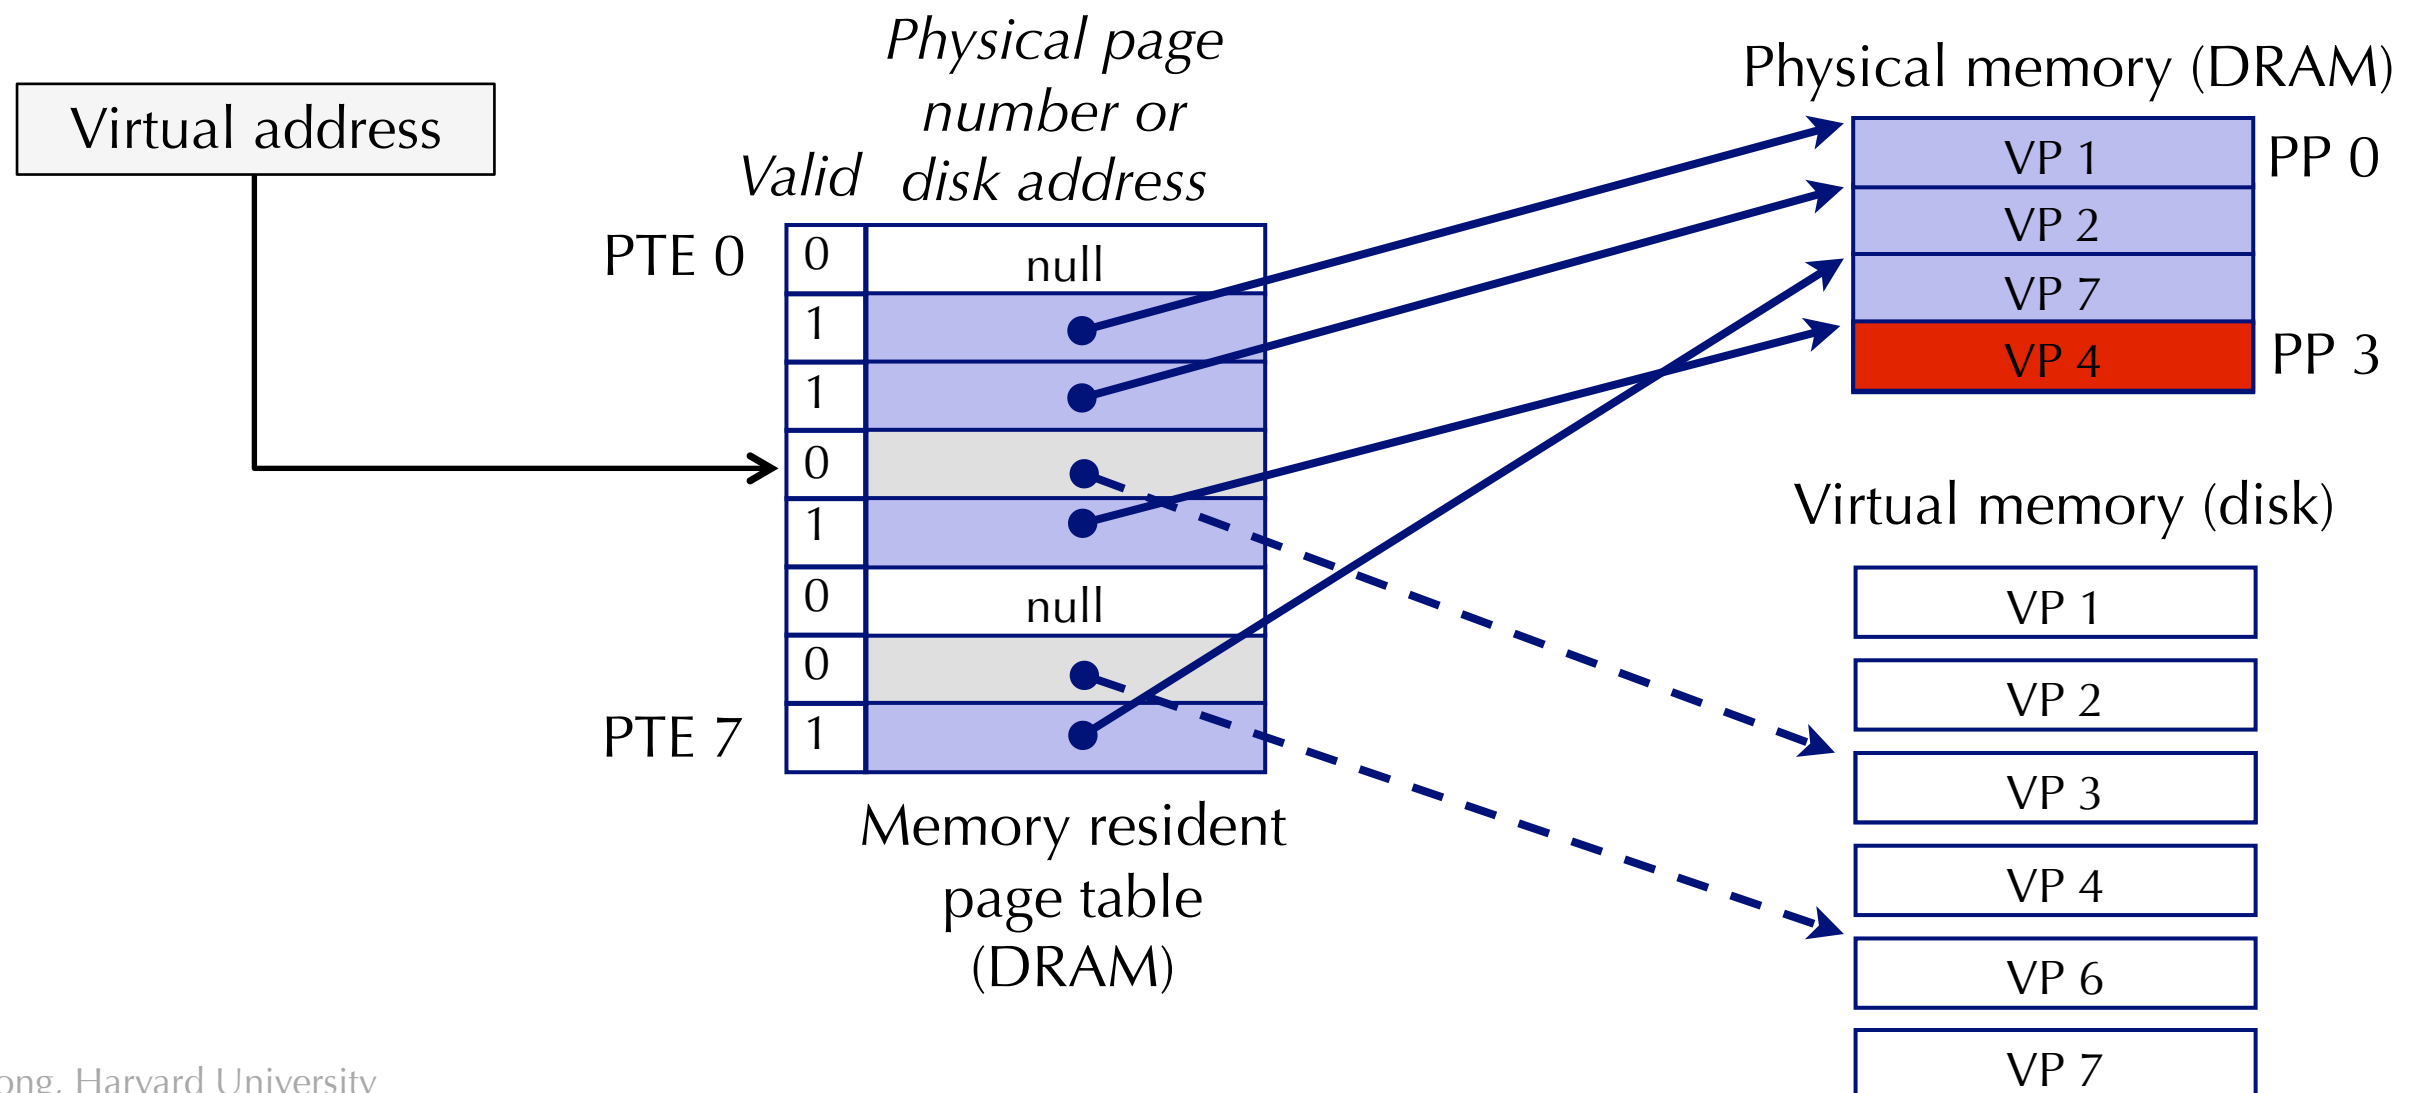
Page fault

- **Page fault:** reference to VM word that is not in physical memory
 - i.e., invalid entry in page table
 - Could be bad memory address, or page table currently on disk



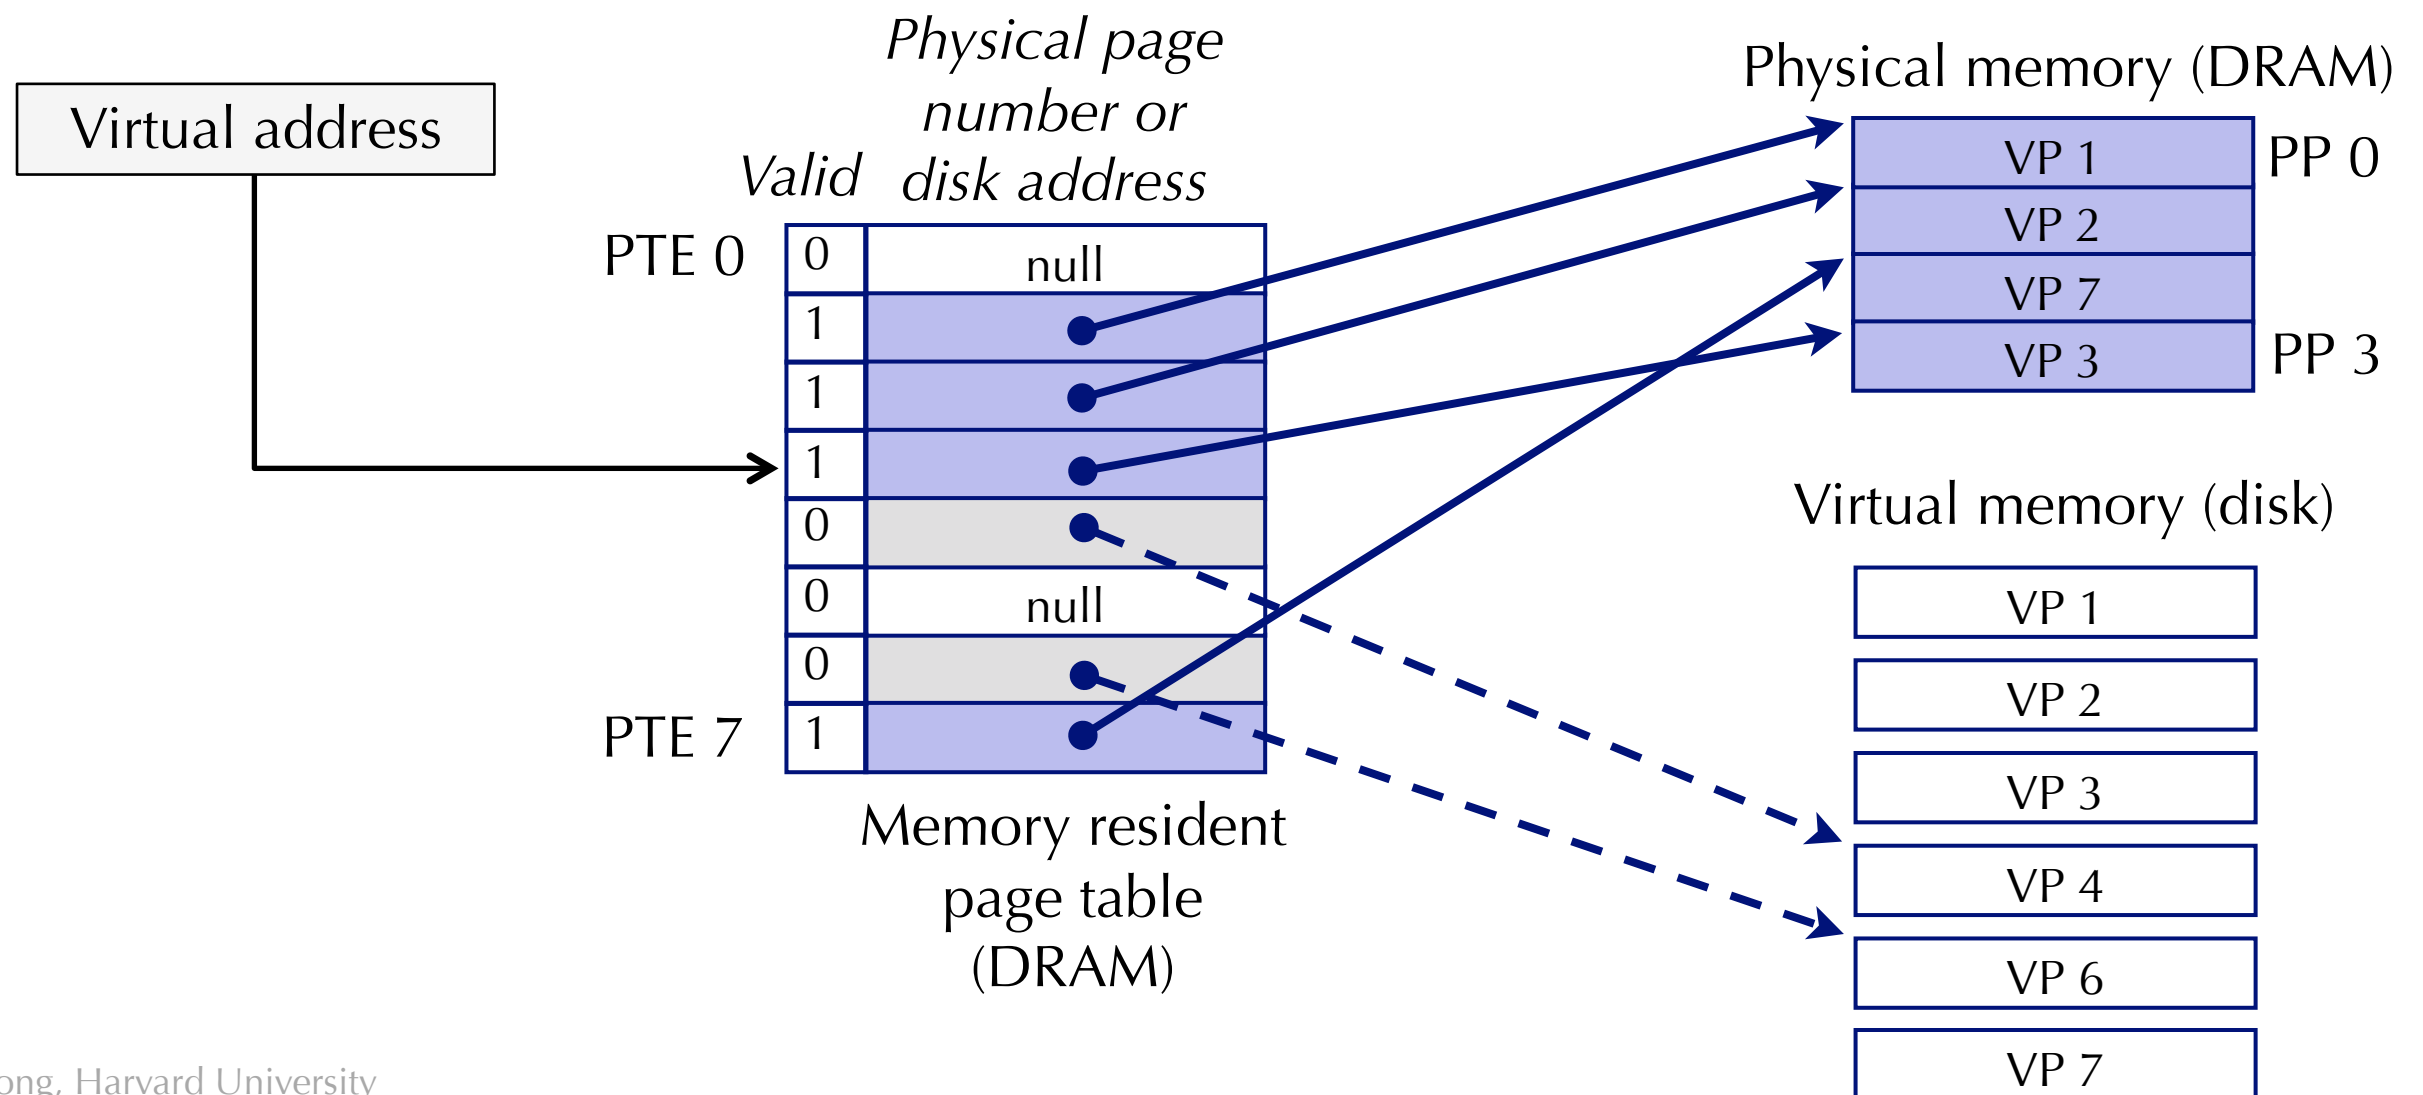
Handling page faults

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)



Handling page faults

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)
- Offending instruction is restarted: page hit!

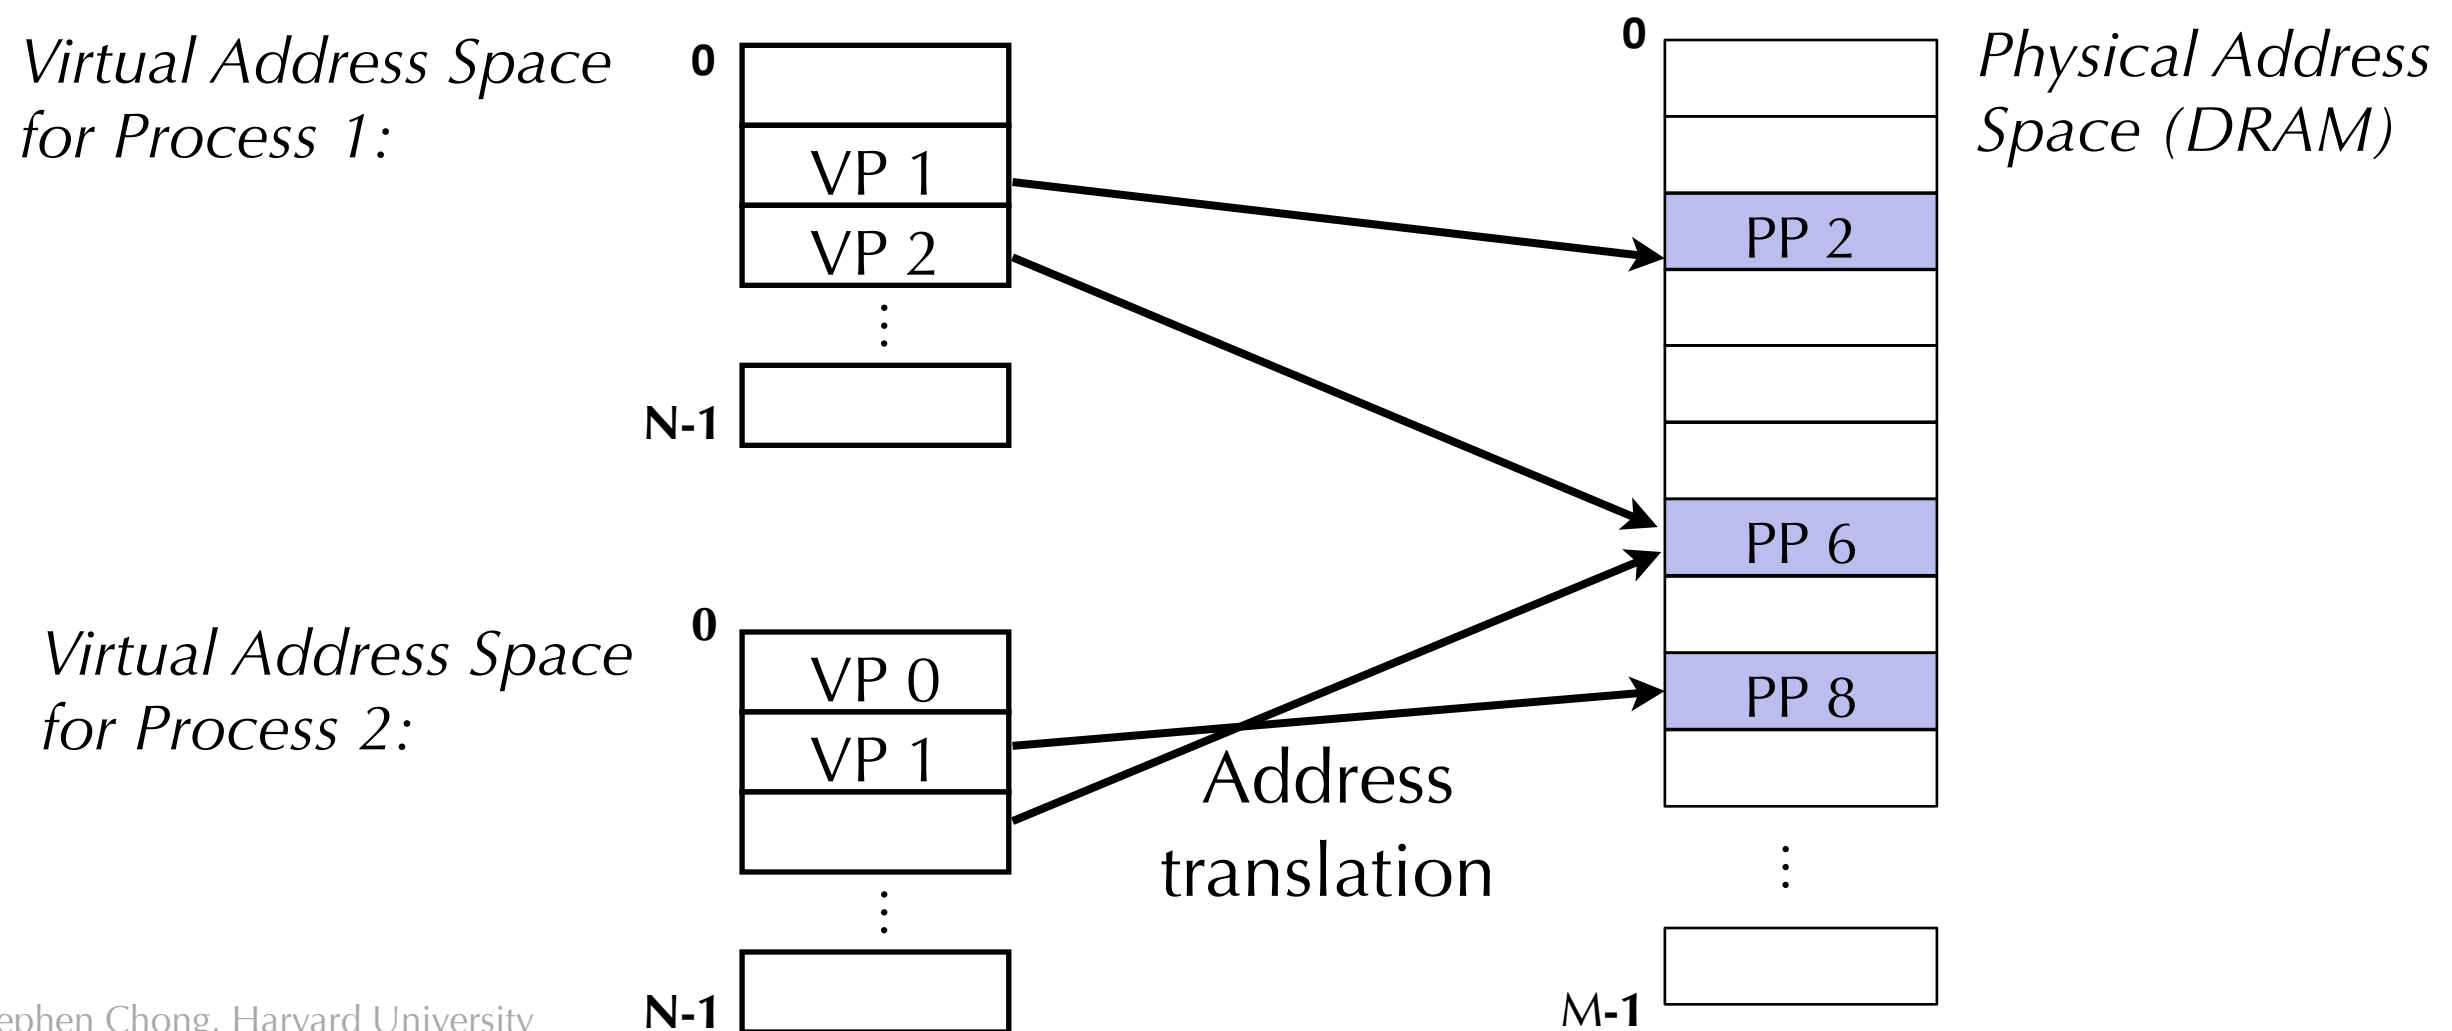


Today

- Running multiple programs at once
- Virtual memory
 - Address spaces
 - Benefits
 - Physical memory as a cache for virtual memory
 - Page tables
 - Page hits and page faults
 - Virtual memory as a tool for memory management
 - Virtual memory as a tool for memory protection
 - Address translation
 - Table Lookaside Buffer (TLB)
 - Example

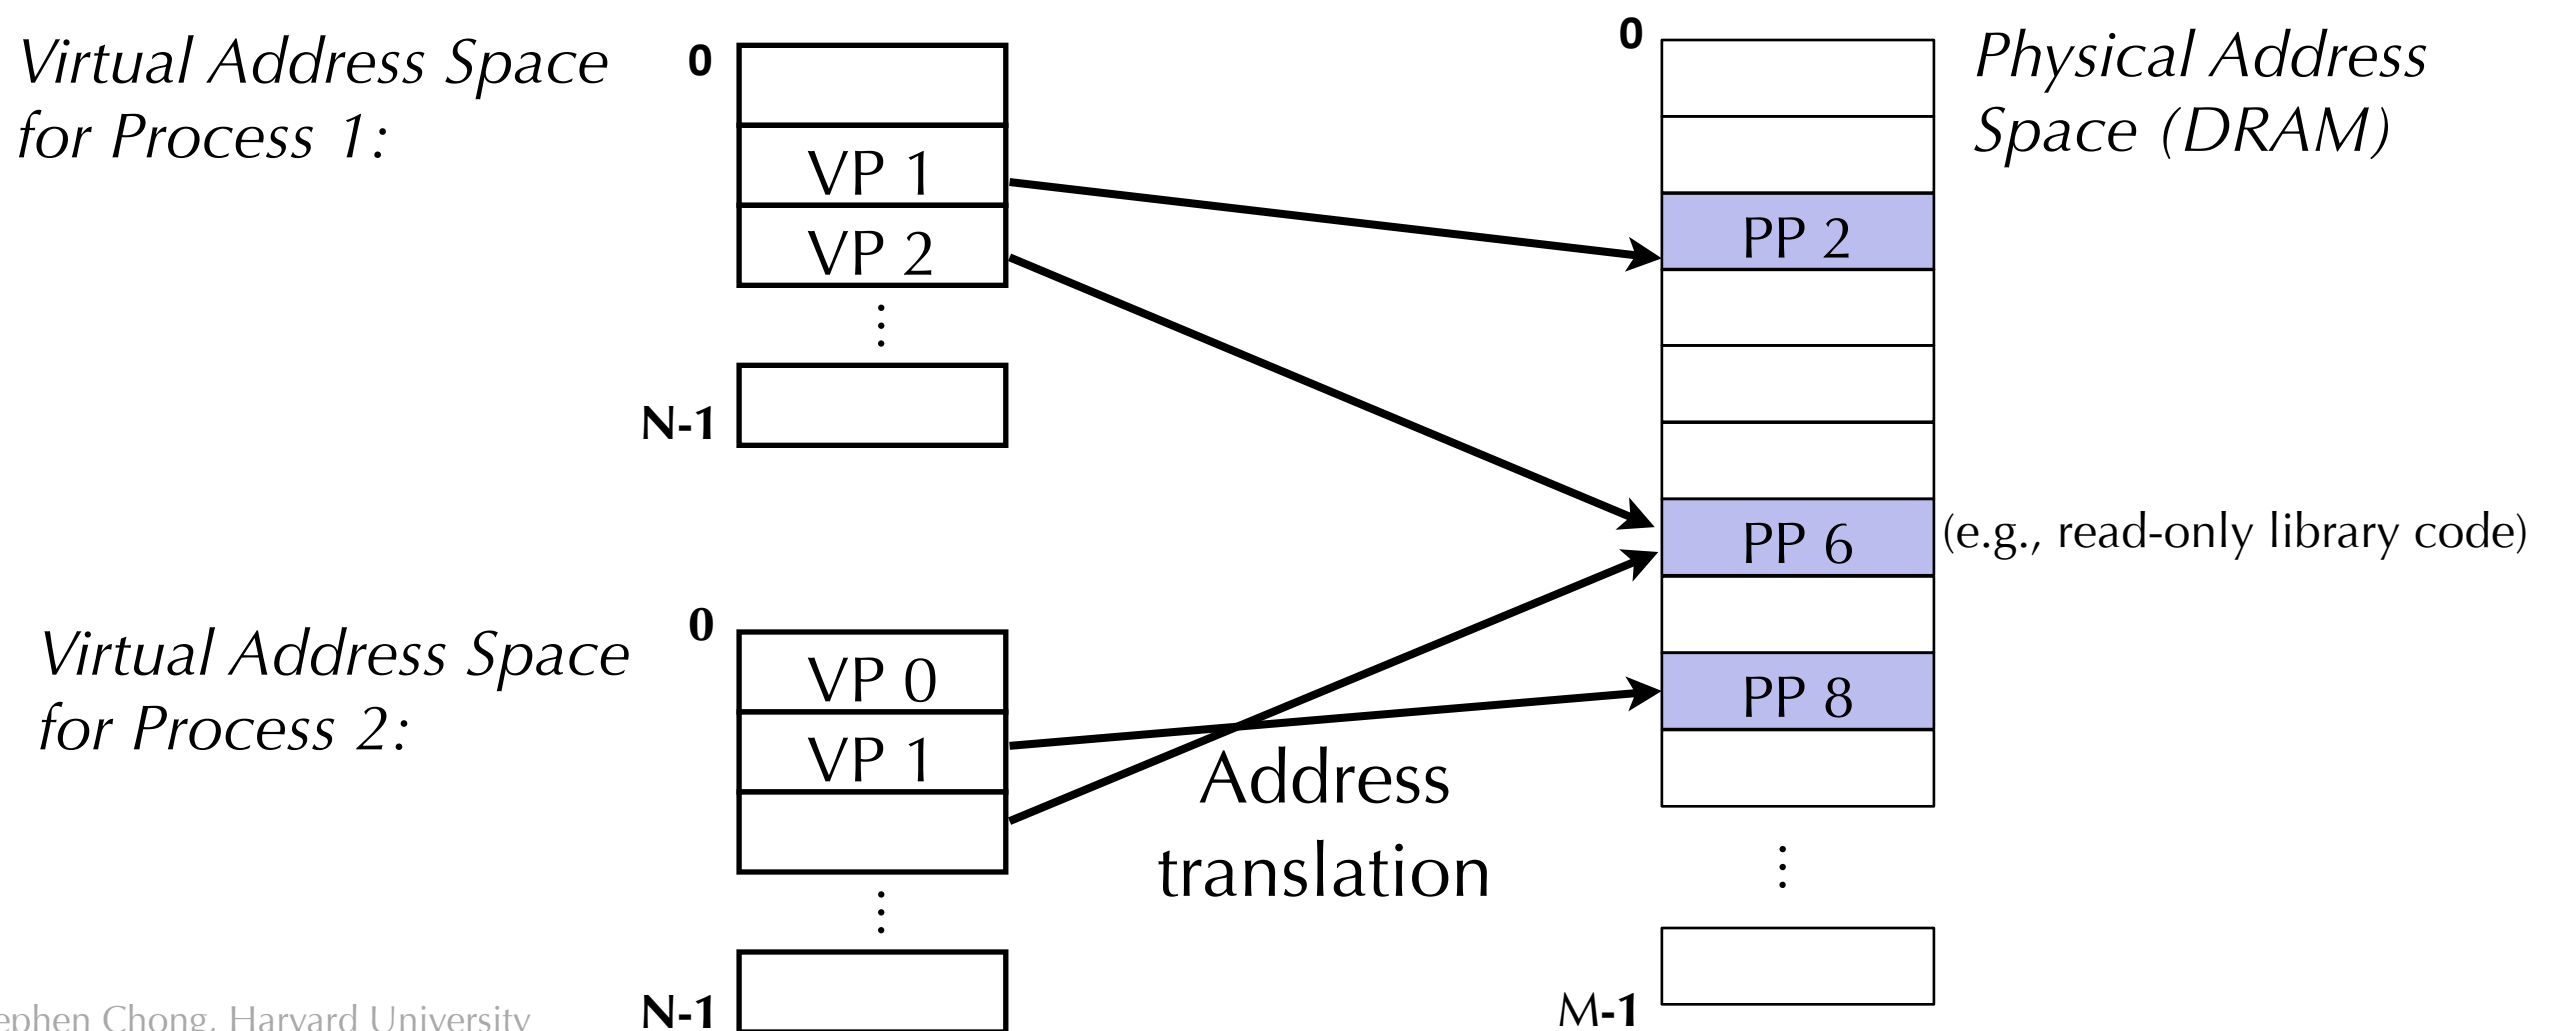
VM as a Tool for Memory Management

- Key idea: each process has its own virtual address space
 - It can view memory as a simple linear array
 - Mapping function scatters addresses through physical memory
 - Well chosen mappings simplify memory allocation and management



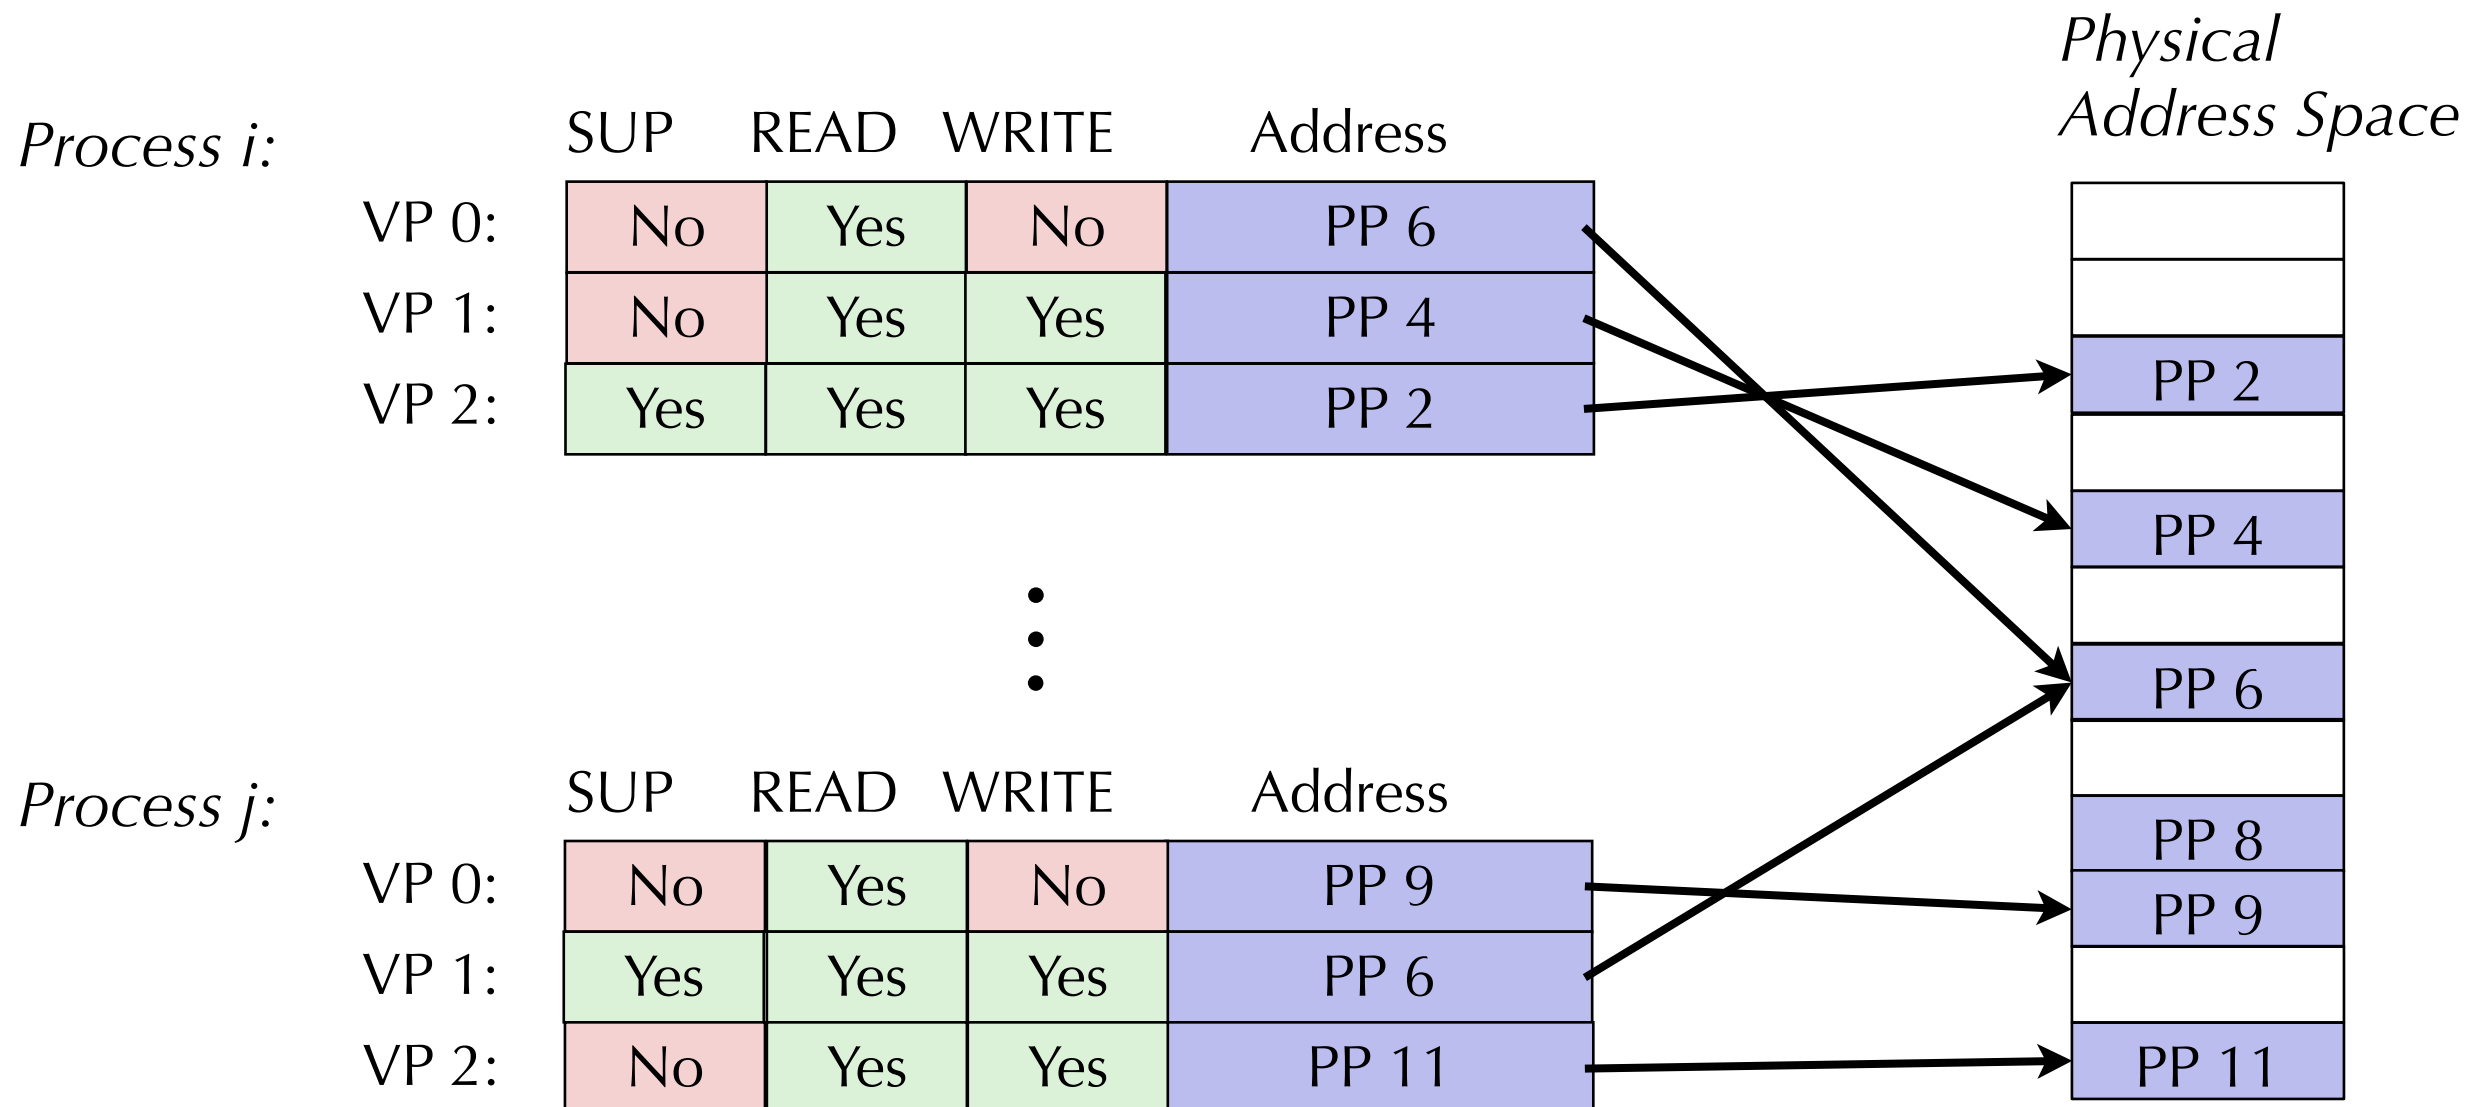
VM as a Tool for Memory Management

- Memory allocation
 - Each virtual page can be mapped to any physical page
 - A virtual page can be stored in different physical pages at different times
- Sharing code and data among processes
 - Map virtual pages to the same physical page (here: PP 6)



VM as a Tool for Memory Protection

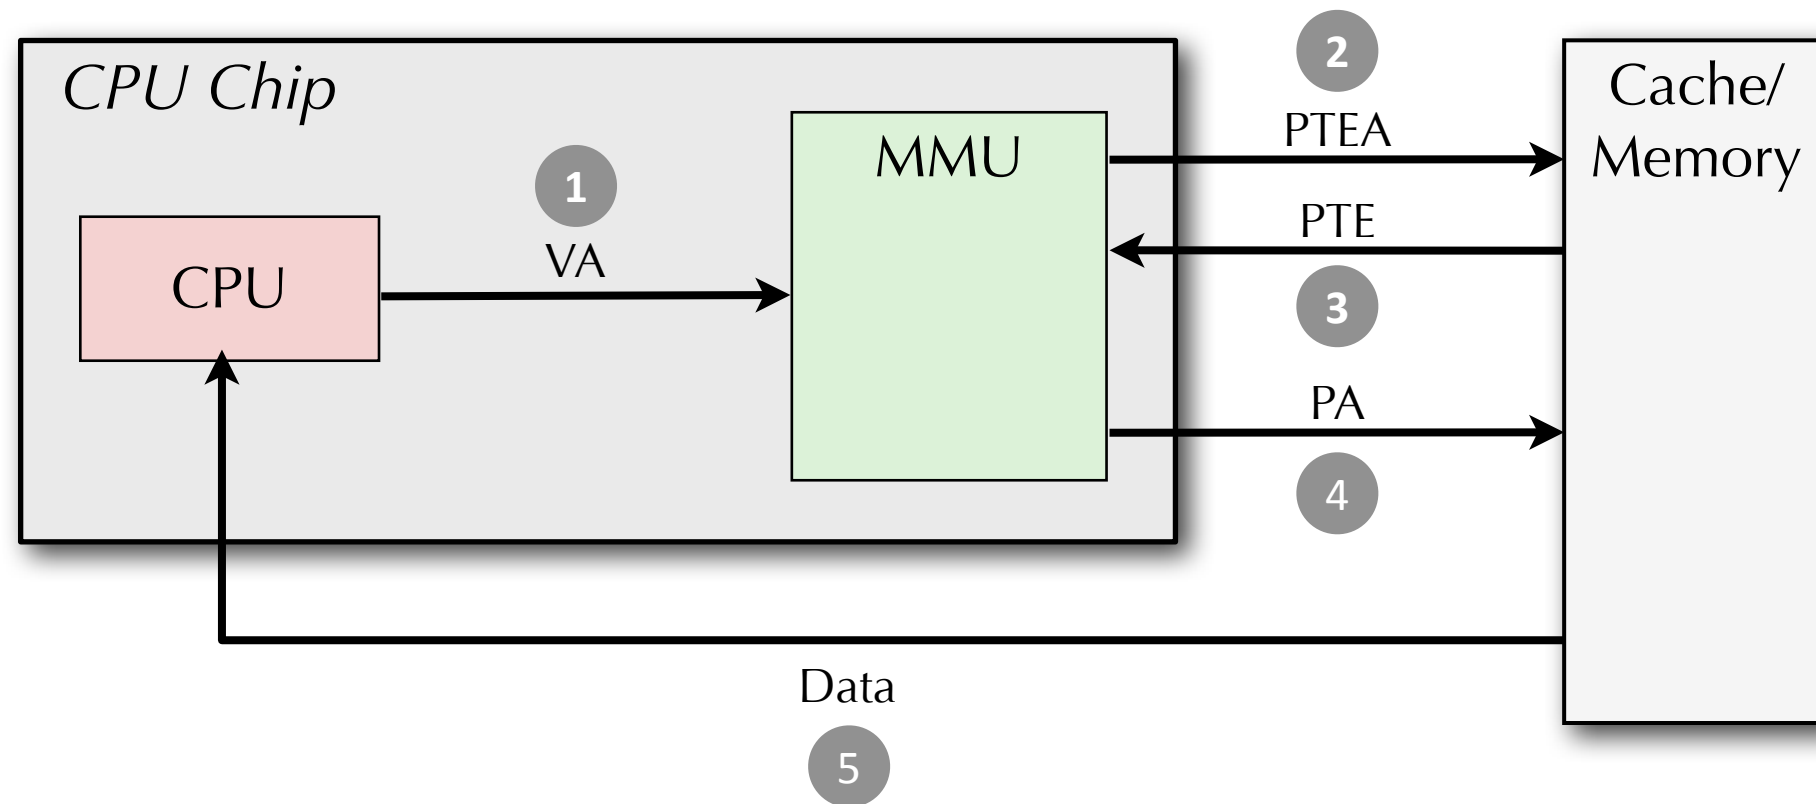
- Extend PTEs with permission bits
- Page fault handler checks these before remapping



Today

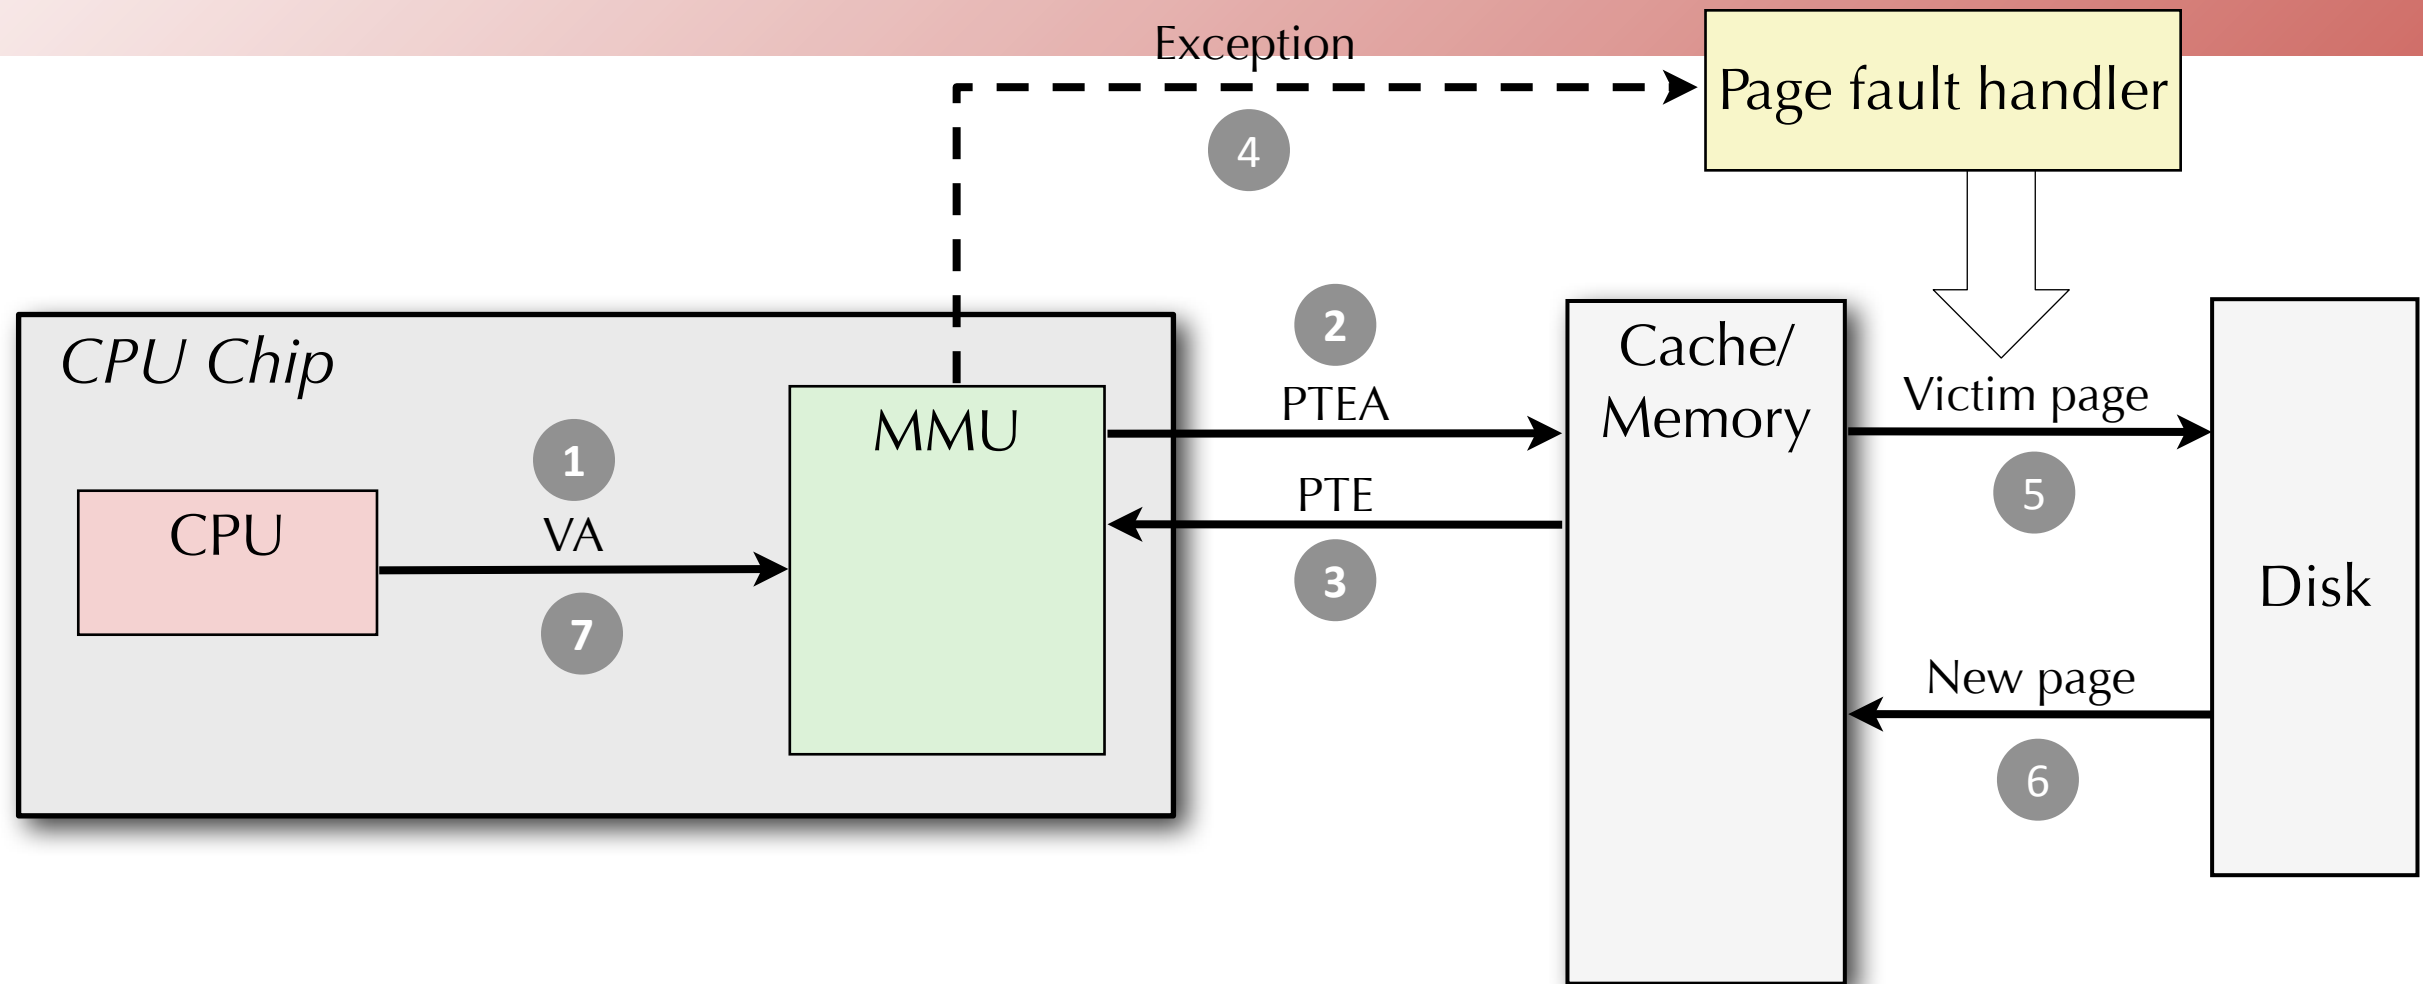
- Running multiple programs at once
- Virtual memory
 - Address spaces
 - Benefits
 - Physical memory as a cache for virtual memory
 - Page tables
 - Page hits and page faults
 - Virtual memory as a tool for memory management
 - Virtual memory as a tool for memory protection
 - Address translation
 - Table Lookaside Buffer (TLB)
 - Example

Address Translation: Page Hit



- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) MMU sends physical address to cache/memory
- 5) Cache/memory sends data word to processor

Address Translation: Page Fault




- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) Valid bit is zero, so MMU triggers page fault exception
- 5) Handler identifies victim (and, if dirty, pages it out to disk)
- 6) Handler pages in new page and updates PTE in memory
- 7) Handler returns to original process, restarting faulting instruction

Speeding up Translation with a TLB


- Page table entries (PTEs) are cached in L1 like any other memory word
 - PTEs may be evicted by other data references
 - PTE hit still requires a 1-cycle delay
- Solution: **Translation Lookaside Buffer (TLB)**
 - Small hardware cache in MMU
 - Maps virtual page numbers to physical page numbers
 - Contains complete page table entries for small number of pages

TLB Hit



- A TLB hit eliminates a memory access

TLB Miss




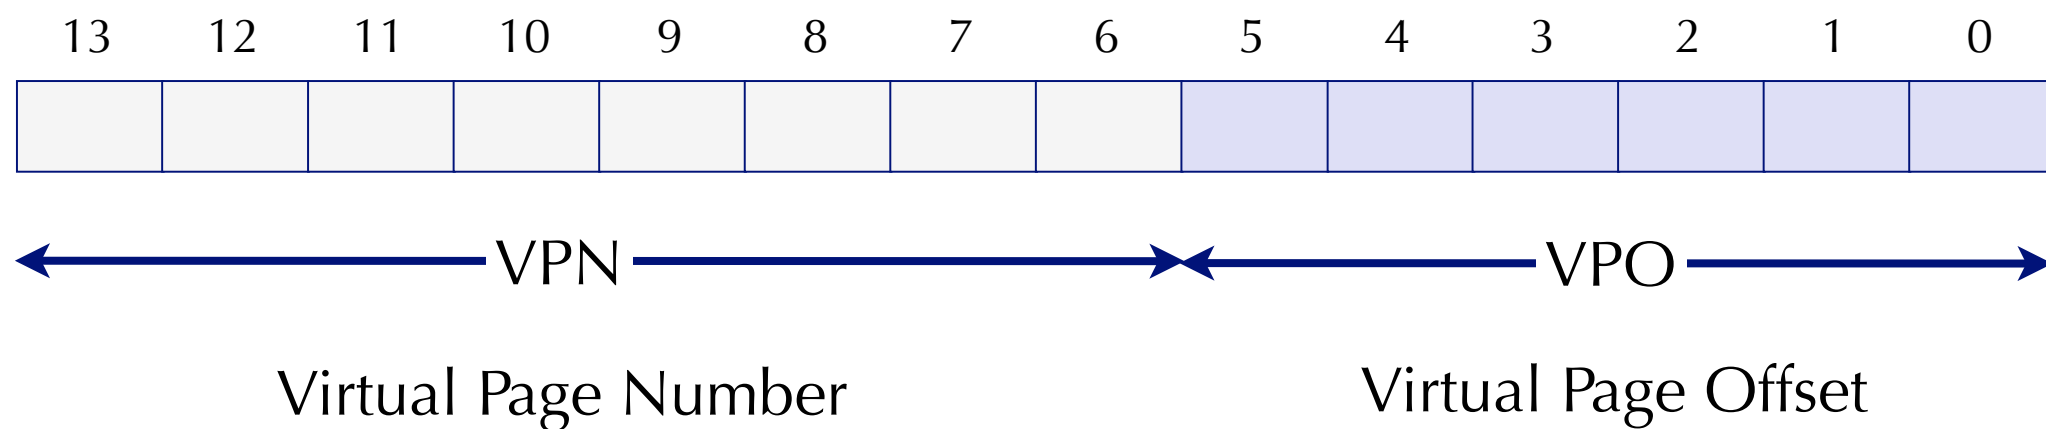
- A TLB miss incurs an additional memory access (the PTE)

Today

- Running multiple programs at once
- Virtual memory
 - Address spaces
 - Benefits
 - Physical memory as a cache for virtual memory
 - Page tables
 - Page hits and page faults
 - Virtual memory as a tool for memory management
 - Virtual memory as a tool for memory protection
 - Address translation
 - Table Lookaside Buffer (TLB)
 - Example

Simple Memory System Example

- Addressing
 - 14-bit virtual addresses
 - 12-bit physical address
 - Page size = 64 bytes



Simple Memory System Page Table


- Only show first 16 entries (out of 256)

<i>VPN</i>	<i>PPN</i>	<i>Valid</i>
00	28	1
01	—	0
02	33	1
03	02	1
04	—	0
05	16	1
06	—	0
07	—	0

<i>VPN</i>	<i>PPN</i>	<i>Valid</i>
08	13	1
09	17	1
0A	09	1
0B	—	0
0C	—	0
0D	2D	1
0E	11	1
0F	0D	1

Simple Memory System TLB

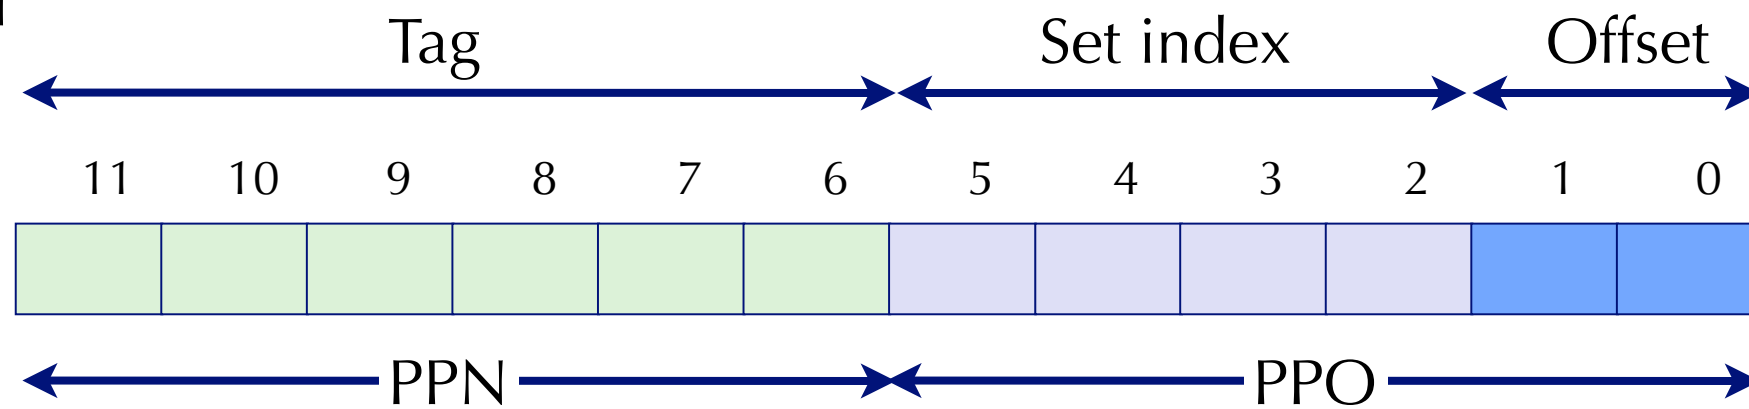
- 16 entries
- 4-way associative



Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	03	—	0	09	0D	1	00	—	0	07	02	1
1	03	2D	1	02	—	0	04	—	0	0A	—	0
2	02	—	0	08	—	0	06	—	0	03	—	0
3	07	—	0	03	0D	1	0A	34	1	02	—	0

Simple Memory System Cache

- 16 lines, 4-byte block size
- **Physically addressed**
- Direct mapped




<i>Idx</i>	<i>Tag</i>	<i>Valid</i>	<i>B0</i>	<i>B1</i>	<i>B2</i>	<i>B3</i>
0	19	1	99	11	23	11
1	15	0	—	—	—	—
2	1B	1	00	02	04	08
3	36	0	—	—	—	—
4	32	1	43	6D	8F	09
5	0D	1	36	72	F0	1D
6	31	0	—	—	—	—
7	16	1	11	C2	DF	03

<i>Idx</i>	<i>Tag</i>	<i>Valid</i>	<i>B0</i>	<i>B1</i>	<i>B2</i>	<i>B3</i>
8	24	1	3A	00	51	89
9	2D	0	—	—	—	—
A	2D	1	93	15	DA	3B
B	0B	0	—	—	—	—
C	12	0	—	—	—	—
D	16	1	04	96	34	15
E	13	1	83	77	1B	D3
F	14	0	—	—	—	—

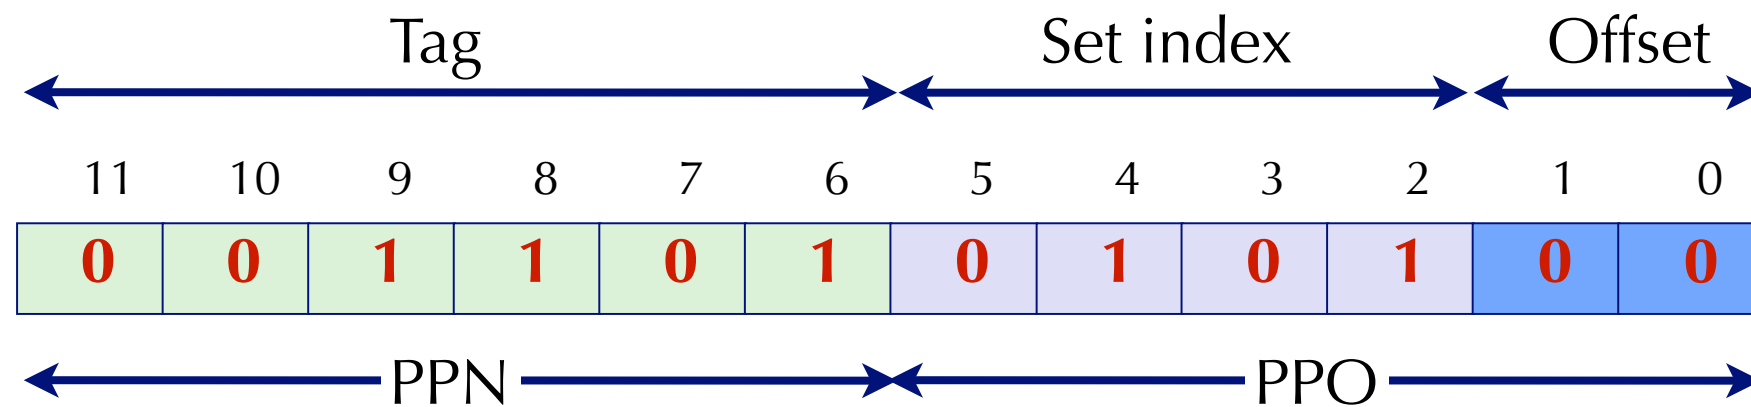
Address Translation Example #1

Virtual Address: 0x03D4



VPN: 0x0F TLB Set index: 3 TLB Tag: 0x03 TLB Hit? Y Page fault? N PPN: 0x0D

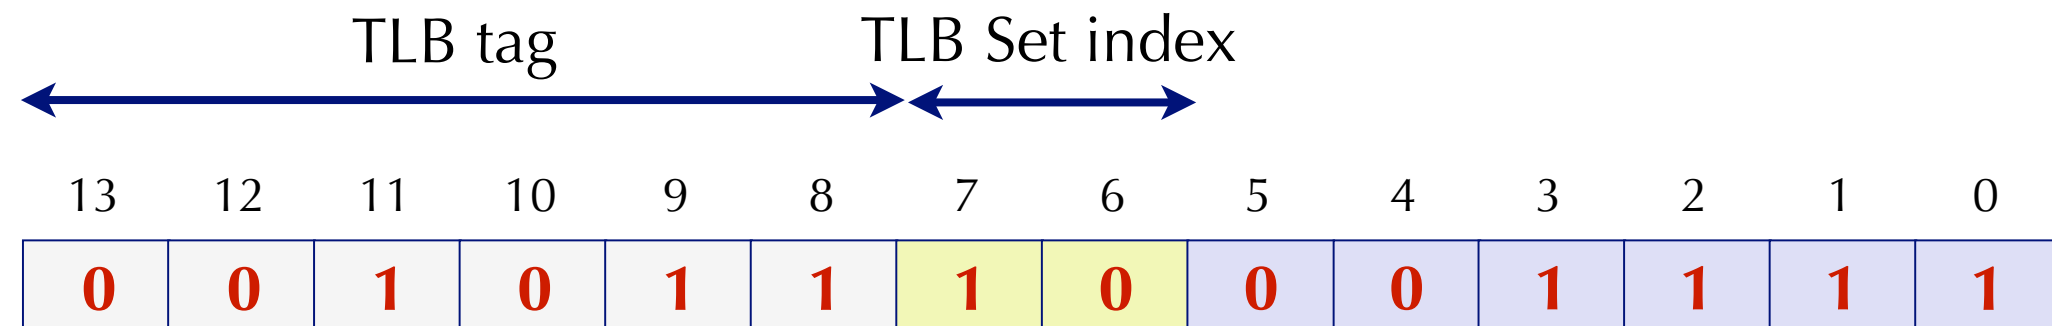
Physical Address



Offset: 0 Set index: 0x5 Tag: 0x0D Hit? Y Byte: 0x36

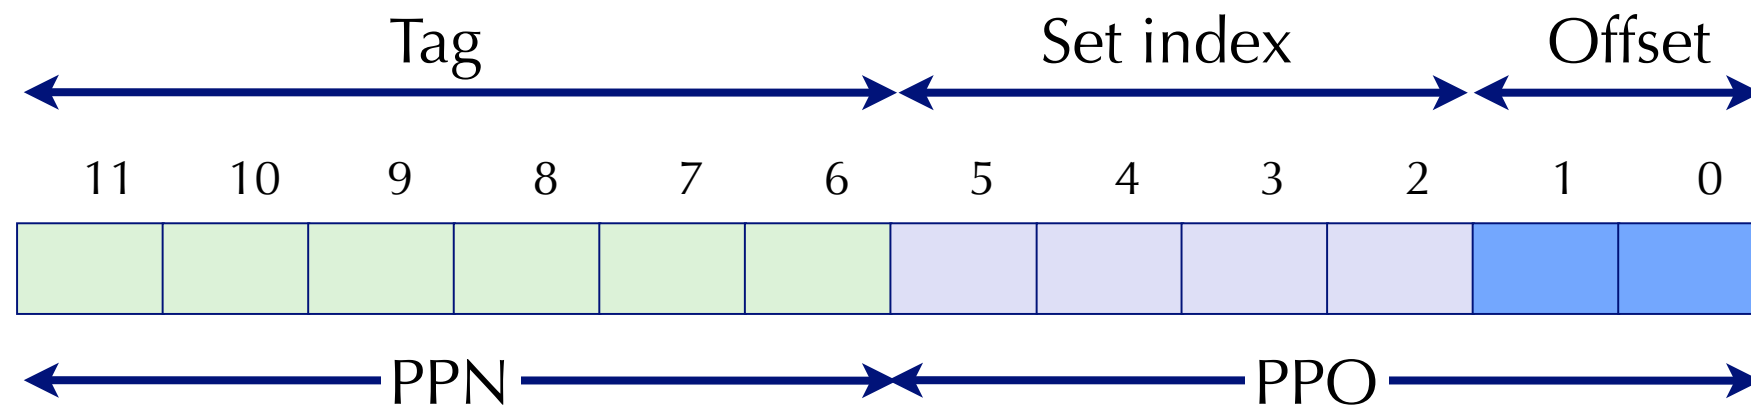
Address Translation Example #2

Virtual Address: 0x0B8F



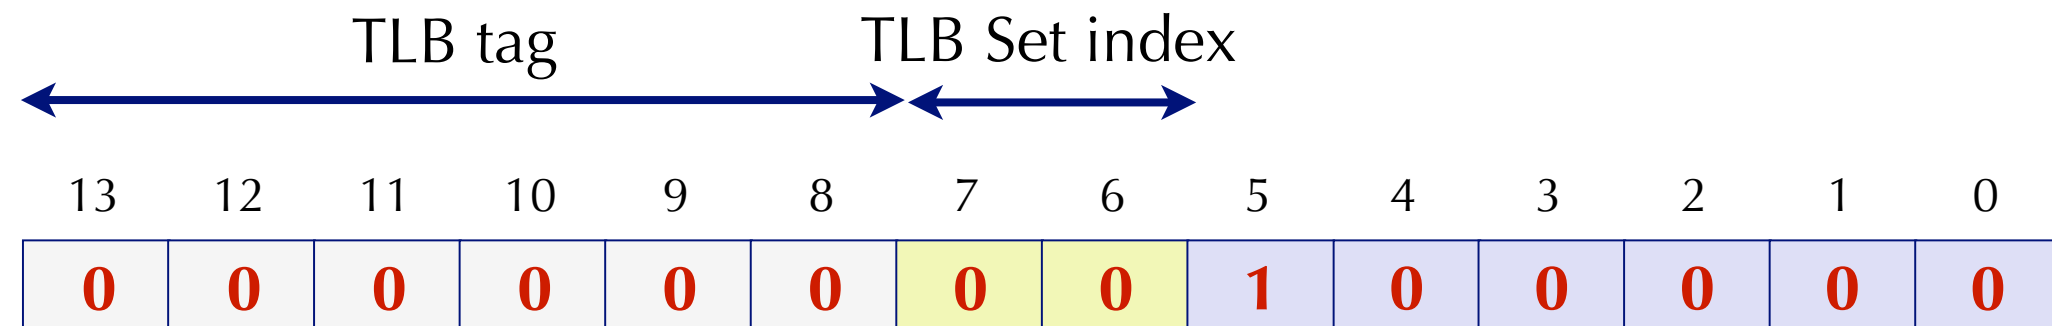
VPN: 0x0E TLB Set index: 2 TLB Tag: 0x0B TLB Hit? N Page fault? Y PPN: TBD

Physical Address



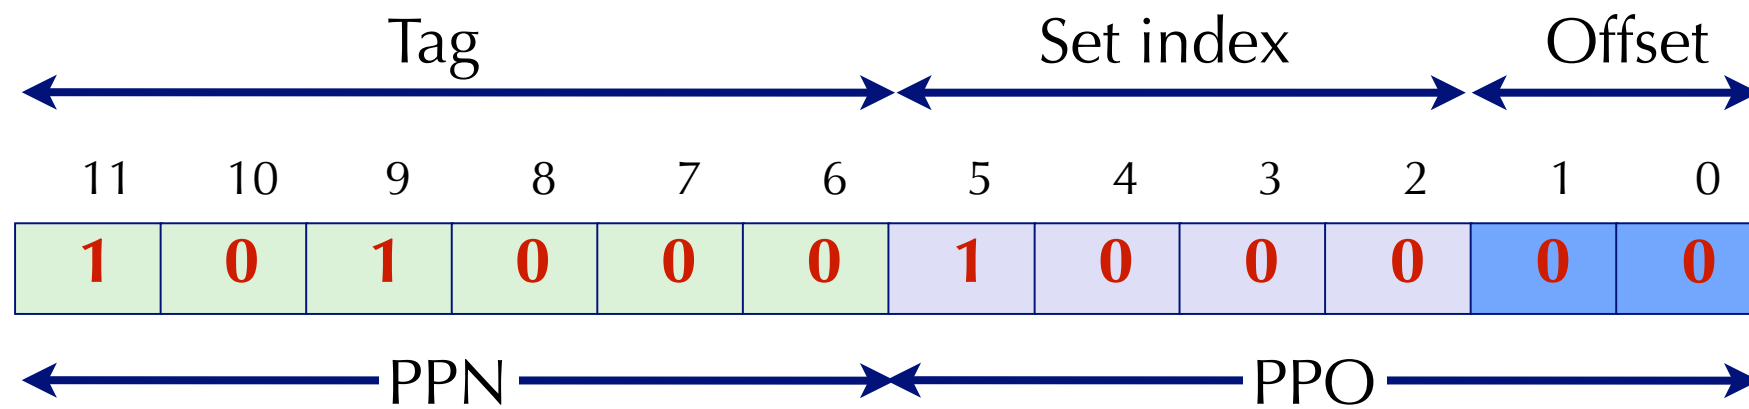
Address Translation Example #3

Virtual Address: 0x0020



VPN: 0x00 TLB Set index: 0 TLB Tag: 0x00 TLB Hit? N Page fault? N PPN: 0x28

Physical Address



Offset: 0 Set index: 0x8 Tag: 0x28 Hit? N Byte: Mem...

What happens on a context switch?

- Context switch is fast. Why?
- Only need to change page table base register, and flush the TLB
- Why don't we need to flush L1, L2, L3 cache?
 - CPU caches use *physical addresses*
 - So single physical page with multiple virtual addresses will get reused across processes!

Summary of VM benefits

- Isolation
 - Each process has its own private linear address space
 - Cannot be corrupted by other processes
- Simplifies memory management and programming
 - Allows multiple concurrent programs to share limited DRAM
- Simplifies protection by providing a convenient method to check permissions
- Efficient
 - Fast translation
 - MMU on CPU, and TLB provides cache of recent translation
 - Fast context switch

Summary of VM mechanism

- VM is implemented by the MMU and OS working together.
 - MMU does virtual-to-physical address translations.
 - OS sets up the page tables that control those translations.
- The TLB is a cache for recent virtual-to-physical mappings.
 - Avoids lookup in page tables for each memory access.
- The OS handles page faults triggered by the MMU.
 - This is the basis for swapping and demand paging.

Next lecture

- Linking and loading